

Table of Contents

[Machine Learning Studio Documentation](#)

[Overview](#)

[What's Machine Learning?](#)

[Machine Learning Studio](#)

[ML Studio capabilities](#)

[ML Studio basics \(infographic\)](#)

[Frequently asked questions](#)

[What's new?](#)

[Get Started](#)

[Create your first experiment](#)

[Example walkthrough](#)

[Create a predictive solution](#)

[1 - Create a workspace](#)

[2 - Upload data](#)

[3 - Create experiment](#)

[4 - Train and evaluate](#)

[5 - Deploy web service](#)

[6 - Access web service](#)

[Data Science for Beginners](#)

[1 - Five questions](#)

[2 - Is your data ready?](#)

[3 - Ask the right question](#)

[4 - Predict an answer](#)

[5 - Copy other people's work](#)

[R quick start](#)

[How To](#)

[Set up tools and utilities](#)

[Manage a workspace](#)

[Acquire and understand data](#)

- [Import training data](#)
- [Develop models](#)
 - [Create and train models](#)
 - [Operationalize models](#)
 - [Overview](#)
 - [Deploy models](#)
 - [Manage web services](#)
 - [Retrain models](#)
 - [Consume models](#)
 - [Examples](#)
 - [Sample experiments](#)
 - [Sample datasets](#)
 - [Customer churn example](#)
 - [Reference](#)
 - [Azure PowerShell module \(New\)](#)
 - [Azure PowerShell module \(Classic\)](#)
 - [Algorithm & Module reference](#)
 - [REST management APIs](#)
 - [Web service error codes](#)
 - [Related](#)
 - [Cortana Intelligence Gallery](#)
 - [Overview](#)
 - [Industries](#)
 - [Solutions](#)
 - [Experiments](#)
 - [Jupyter Notebooks](#)
 - [Competitions](#)
 - [Competitions FAQ](#)
 - [Tutorials](#)
 - [Collections](#)
 - [Custom Modules](#)
 - [Resources](#)

[Azure Roadmap](#)

[Net# Neural Networks Language](#)

[Pricing](#)

[Service updates](#)

[Blog](#)

[MSDN forum](#)

[Stack Overflow](#)

[Videos](#)

[Get help from live chat](#)

Azure Machine Learning Studio is a collaborative, drag-and-drop tool you can use to build, test, and deploy predictive analytics solutions on your data. Tutorials, videos, and example models show you how to use Studio to build and deploy machine learning models.

-

[Learn about Azure Machine Learning Studio](#)

-

[Azure Machine Learning Video Library](#)

-

[Get started with Data Science](#)

-

[Get started creating an Azure Machine Learning Studio experiment](#)

Reference

- **Command-Line**

[PowerShell \(New\)](#)

[PowerShell \(Classic\)](#)

- **Languages**

[Algorithm & Module reference](#)

- **REST**

[REST API reference](#)

[Web service error codes](#)

Introduction to Machine Learning in the Azure cloud

9/25/2017 • 6 min to read • [Edit Online](#)

What is machine learning?

Machine learning is a data science technique that allows computers to use existing data to forecast future behaviors, outcomes, and trends. Using machine learning, computers learn without being explicitly programmed.

Forecasts or predictions from machine learning can make apps and devices smarter. When you shop online, machine learning helps recommend other products you might like based on what you've purchased. When your credit card is swiped, machine learning compares the transaction to a database of transactions and helps detect fraud. When your robot vacuum cleaner vacuums a room, machine learning helps it decide whether the job is done.

For a brief overview, try the video series [Data Science for Beginners](#). Without using jargon or math, Data Science for Beginners introduces machine learning and steps you through a simple predictive model.

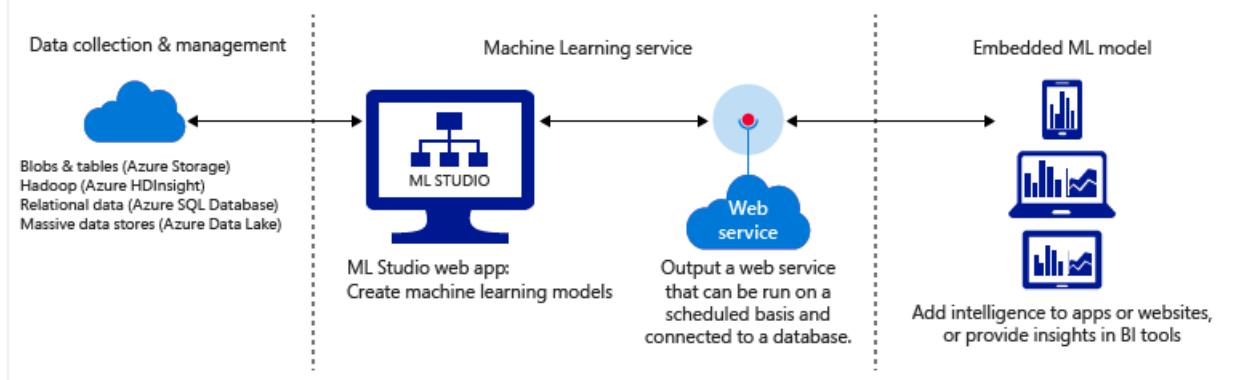
What is Machine Learning in the Microsoft Azure cloud?

Azure Machine Learning is a cloud predictive analytics service that makes it possible to quickly create and deploy predictive models as analytics solutions.

You can work from a ready-to-use library of algorithms, use them to create models on an internet-connected PC, and deploy your predictive solution quickly. Start from ready-to-use examples and solutions in the [Cortana Intelligence Gallery](#).

Azure Machine Learning: Basic workflow

Build models from data and operationalize a machine learning solution



Azure Machine Learning not only provides tools to model predictive analytics, but also provides a fully managed service you can use to deploy your predictive models as ready-to-consume web services.

What is predictive analytics?

Predictive analytics uses math formulas called algorithms that analyze historical or current data to identify patterns or trends in order to forecast future events.

Tools to build complete machine learning solutions in the cloud

Azure Machine Learning has everything you need to create complete predictive analytics solutions in the cloud, from a large algorithm library, to a studio for building models, to an easy way to deploy your model as a web

service. Quickly create, test, operationalize, and manage predictive models.

Machine Learning Studio: Create predictive models

In [Machine Learning Studio](#), you can quickly create predictive models by dragging, dropping, and connecting modules. You can experiment with different combinations, and [try it out for free](#).

- In [Cortana Intelligence Gallery](#), you can try analytics solutions authored by others or contribute your own. Post questions or comments about experiments to the community, or share links to experiments via social networks such as LinkedIn and Twitter.

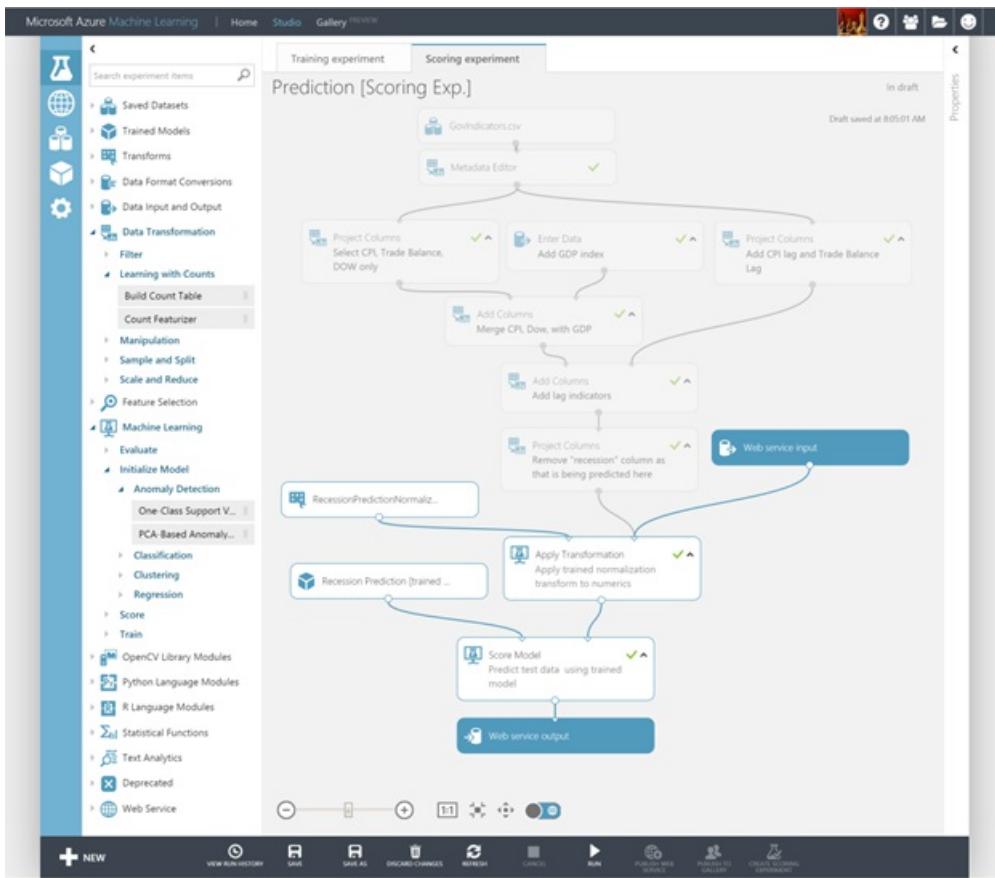
The screenshot shows the Cortana Intelligence Gallery homepage. At the top, there's a navigation bar with 'Cortana Intelligence Gallery' and icons for search, user profile, and sign in. Below the header, a message encourages users to contribute their analytics solutions. The main content area displays several cards representing different machine learning projects:

- COMPETITION**: Decoding Brain Signals (Microsoft) - A card featuring a brain icon and a network graph.
- COLLECTION**: Cognitive Services (Microsoft) - A card featuring a brain icon and the Microsoft logo.
- NOTEBOOK**: Introduction to Azure ML R notebooks (Microsoft) - A card featuring a globe icon and the letter 'R'.
- NOTEBOOK**: Predicting breast cancer using data from AzureML (Microsoft) - A card featuring a pink ribbon icon.
- TUTORIAL**: Energy Demand Forecast Template with SQL Server R Services (Microsoft) - A card featuring a bar chart and an upward arrow.

Below this, there's a section titled 'Recently added' with four more items:

COLLECTION	EXPERIMENT	COLLECTION	EXPERIMENT
Cognitive Services	Decoding Brain Signals - RD H Attempt 2	SuperFamily	lab City-Market Clustering - Copy

- Use a large library of [Machine Learning algorithms and modules](#) in Machine Learning Studio to jump-start your predictive models. Choose from sample experiments, R and Python packages, and best-in-class algorithms from Microsoft businesses like Xbox and Bing. Extend Studio modules with your own custom [R](#) and [Python](#) scripts.



Operationalize predictive analytics solutions by publishing your own

The following tutorials show you how to operationalize your predictive analytics models:

- [Deploy web services](#)
- [Retrain models through APIs](#)
- [Manage web service endpoints](#)
- [Scale a web service](#)
- [Consume web services](#)

Key machine learning terms and concepts

Machine learning terms can be confusing. Here are definitions of key terms to help you. Use comments following to tell us about any other term you'd like defined.

Data exploration, descriptive analytics, and predictive analytics

Data exploration is the process of gathering information about a large and often unstructured data set in order to find characteristics for focused analysis.

Data mining refers to automated data exploration.

Descriptive analytics is the process of analyzing a data set in order to summarize what happened. The vast majority of business analytics - such as sales reports, web metrics, and social networks analysis - are descriptive.

Predictive analytics is the process of building models from historical or current data in order to forecast future outcomes.

Supervised and unsupervised learning

Supervised learning algorithms are trained with labeled data - in other words, data comprised of examples of the answers wanted. For instance, a model that identifies fraudulent credit card use would be trained from a data set with labeled data points of known fraudulent and valid charges. Most machine learning is supervised.

Unsupervised learning is used on data with no labels, and the goal is to find relationships in the data. For

instance, you might want to find groupings of customer demographics with similar buying habits.

Model training and evaluation

A machine learning model is an abstraction of the question you are trying to answer or the outcome you want to predict. Models are trained and evaluated from existing data.

Training data

When you train a model from data, you use a known data set and make adjustments to the model based on the data characteristics to get the most accurate answer. In Azure Machine Learning, a model is built from an algorithm module that processes training data and functional modules, such as a scoring module.

In supervised learning, if you're training a fraud detection model, you use a set of transactions that are labeled as either fraudulent or valid. You split your data set randomly, and use part to train the model and part to test or evaluate the model.

Evaluation data

Once you have a trained model, evaluate the model using the remaining test data. You use data you already know the outcomes for, so that you can tell whether your model predicts accurately.

Other common machine learning terms

- **algorithm:** A self-contained set of rules used to solve problems through data processing, math, or automated reasoning.
- **anomaly detection:** A model that flags unusual events or values and helps you discover problems. For example, credit card fraud detection looks for unusual purchases.
- **categorical data:** Data that is organized by categories and that can be divided into groups. For example a categorical data set for autos could specify year, make, model, and price.
- **classification:** A model for organizing data points into categories based on a data set for which category groupings are already known.
- **feature engineering:** The process of extracting or selecting features related to a data set in order to enhance the data set and improve outcomes. For instance, airfare data could be enhanced by days of the week and holidays. See [Feature selection and engineering in Azure Machine Learning](#).
- **module:** A functional part in a Machine Learning Studio model, such as the Enter Data module that enables entering and editing small data sets. An algorithm is also a type of module in Machine Learning Studio.
- **model:** A supervised learning model is the product of a machine learning experiment comprised of training data, an algorithm module, and functional modules, such as a Score Model module.
- **numerical data:** Data that has meaning as measurements (continuous data) or counts (discrete data). Also referred to as *quantitative data*.
- **partition:** The method by which you divide data into samples. See [Partition and Sample](#) for more information.
- **prediction:** A prediction is a forecast of a value or values from a machine learning model. You might also see the term "predicted score." However, predicted scores are not the final output of a model. An evaluation of the model follows the score.
- **regression:** A model for predicting a value based on independent variables, such as predicting the price of a car based on its year and make.
- **score:** A predicted value generated from a trained classification or regression model, using the [Score Model module](#) in Machine Learning Studio. Classification models also return a score for the probability of the predicted value. Once you've generated scores from a model, you can evaluate the model's accuracy using the [Evaluate Model module](#).
- **sample:** A part of a data set intended to be representative of the whole. Samples can be selected randomly or based on specific features of the data set.

Next steps

You can learn the basics of predictive analytics and machine learning using a [step-by-step tutorial](#) and by [building on samples](#).

What is Azure Machine Learning Studio?

9/25/2017 • 5 min to read • [Edit Online](#)

Microsoft Azure Machine Learning Studio is a collaborative, drag-and-drop tool you can use to build, test, and deploy predictive analytics solutions on your data. Machine Learning Studio publishes models as web services that can easily be consumed by custom apps or BI tools such as Excel.

Machine Learning Studio is where data science, predictive analytics, cloud resources, and your data meet.

NOTE

Try Azure Machine Learning for free

No credit card or Azure subscription required. [Get started now.](#)

The Machine Learning Studio interactive workspace

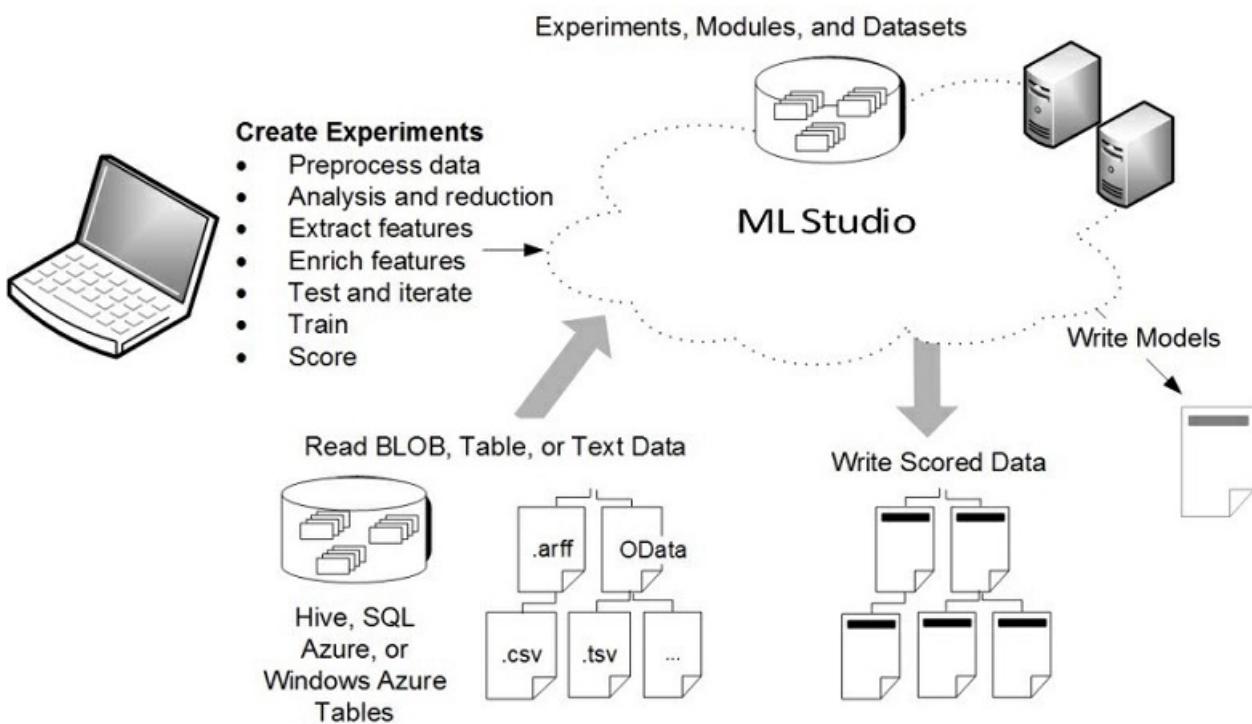
To develop a predictive analysis model, you typically use data from one or more sources, transform and analyze that data through various data manipulation and statistical functions, and generate a set of results. Developing a model like this is an iterative process. As you modify the various functions and their parameters, your results converge until you are satisfied that you have a trained, effective model.

Azure Machine Learning Studio gives you an interactive, visual workspace to easily build, test, and iterate on a predictive analysis model. You drag-and-drop **datasets** and analysis **modules** onto an interactive canvas, connecting them together to form an **experiment**, which you run in Machine Learning Studio. To iterate on your model design, you edit the experiment, save a copy if desired, and run it again. When you're ready, you can convert your **training experiment** to a **predictive experiment**, and then publish it as a **web service** so that your model can be accessed by others.

There is no programming required, just visually connecting datasets and modules to construct your predictive analysis model.

TIP

To download and print a diagram that gives an overview of the capabilities of Machine Learning Studio, see [Overview diagram of Azure Machine Learning Studio capabilities](#).



Get started with Machine Learning Studio

When you first enter [Machine Learning Studio](#) you see the **Home** page. From here you can view documentation, videos, webinars, and find other valuable resources.

Click the upper-left menu  and you'll see several options.

Cortana Intelligence

Click **Cortana Intelligence** and you'll be taken to the home page of the [Cortana Intelligence Suite](#). The Cortana Intelligence Suite is a fully managed big data and advanced analytics suite to transform your data into intelligent action. See the Suite home page for full documentation, including customer stories.

Azure Machine Learning

There are two options here, **Home**, the page where you started, and **Studio**.

Click **Studio** and you'll be taken to the [Azure Machine Learning Studio](#). First you'll be asked to sign in using your Microsoft account, or your work or school account. Once signed in, you'll see the following tabs on the left:

- **PROJECTS** - Collections of experiments, datasets, notebooks, and other resources representing a single project
- **EXPERIMENTS** - Experiments that you have created and run or saved as drafts
- **WEB SERVICES** - Web services that you have deployed from your experiments
- **NOTEBOOKS** - Jupyter notebooks that you have created
- **DATASETS** - Datasets that you have uploaded into Studio
- **TRAINED MODELS** - Models that you have trained in experiments and saved in Studio
- **SETTINGS** - A collection of settings that you can use to configure your account and resources.

Gallery

Click **Gallery** and you'll be taken to the [Cortana Intelligence Gallery](#). The Gallery is a place where a community of data scientists and developers share solutions created using components of the Cortana Intelligence Suite.

For more information about the Gallery, see [Share and discover solutions in the Cortana Intelligence Gallery](#).

Components of an experiment

An experiment consists of datasets that provide data to analytical modules, which you connect together to construct a predictive analysis model. Specifically, a valid experiment has these characteristics:

- The experiment has at least one dataset and one module
- Datasets may be connected only to modules
- Modules may be connected to either datasets or other modules
- All input ports for modules must have some connection to the data flow
- All required parameters for each module must be set

You can create an experiment from scratch, or you can use an existing sample experiment as a template. For more information, see [Copy example experiments to create new machine learning experiments](#).

For an example of creating a simple experiment, see [Create a simple experiment in Azure Machine Learning Studio](#).

For a more complete walkthrough of creating a predictive analytics solution, see [Develop a predictive solution with Azure Machine Learning](#).

Datasets

A dataset is data that has been uploaded to Machine Learning Studio so that it can be used in the modeling process. A number of sample datasets are included with Machine Learning Studio for you to experiment with, and you can upload more datasets as you need them. Here are some examples of included datasets:

- **MPG data for various automobiles** - Miles per gallon (MPG) values for automobiles identified by number of cylinders, horsepower, etc.
- **Breast cancer data** - Breast cancer diagnosis data.
- **Forest fires data** - Forest fire sizes in northeast Portugal.

As you build an experiment you can choose from the list of datasets available to the left of the canvas.

For a list of sample datasets included in Machine Learning Studio, see [Use the sample data sets in Azure Machine Learning Studio](#).

Modules

A module is an algorithm that you can perform on your data. Machine Learning Studio has a number of modules ranging from data ingress functions to training, scoring, and validation processes. Here are some examples of included modules:

- [Convert to ARFF](#) - Converts a .NET serialized dataset to Attribute-Relation File Format (ARFF).
- [Compute Elementary Statistics](#) - Calculates elementary statistics such as mean, standard deviation, etc.
- [Linear Regression](#) - Creates an online gradient descent-based linear regression model.
- [Score Model](#) - Scores a trained classification or regression model.

As you build an experiment you can choose from the list of modules available to the left of the canvas.

A module may have a set of parameters that you can use to configure the module's internal algorithms. When you select a module on the canvas, the module's parameters are displayed in the **Properties** pane to the right of the canvas. You can modify the parameters in that pane to tune your model.

For some help navigating through the large library of machine learning algorithms available, see [How to choose algorithms for Microsoft Azure Machine Learning](#).

Deploying a predictive analytics web service

Once your predictive analytics model is ready, you can deploy it as a web service right from Machine Learning Studio. For more details on this process, see [Deploy an Azure Machine Learning web service](#).

Overview diagram of Azure Machine Learning Studio capabilities

9/25/2017 • 1 min to read • [Edit Online](#)

The **Microsoft Azure Machine Learning Studio Capabilities Overview** diagram gives you a high-level overview of how you can use Machine Learning Studio to develop a predictive analytics model and operationalize it in the Azure cloud.

Azure Machine Learning Studio has available a large number of machine learning algorithms, along with modules that help with data input, output, preparation, and visualization. Using these components you can develop a predictive analytics experiment, iterate on it, and use it to train your model. Then with one click you can operationalize your model in the Azure cloud so that it can be used to score new data.

This diagram demonstrates how all those pieces fit together.

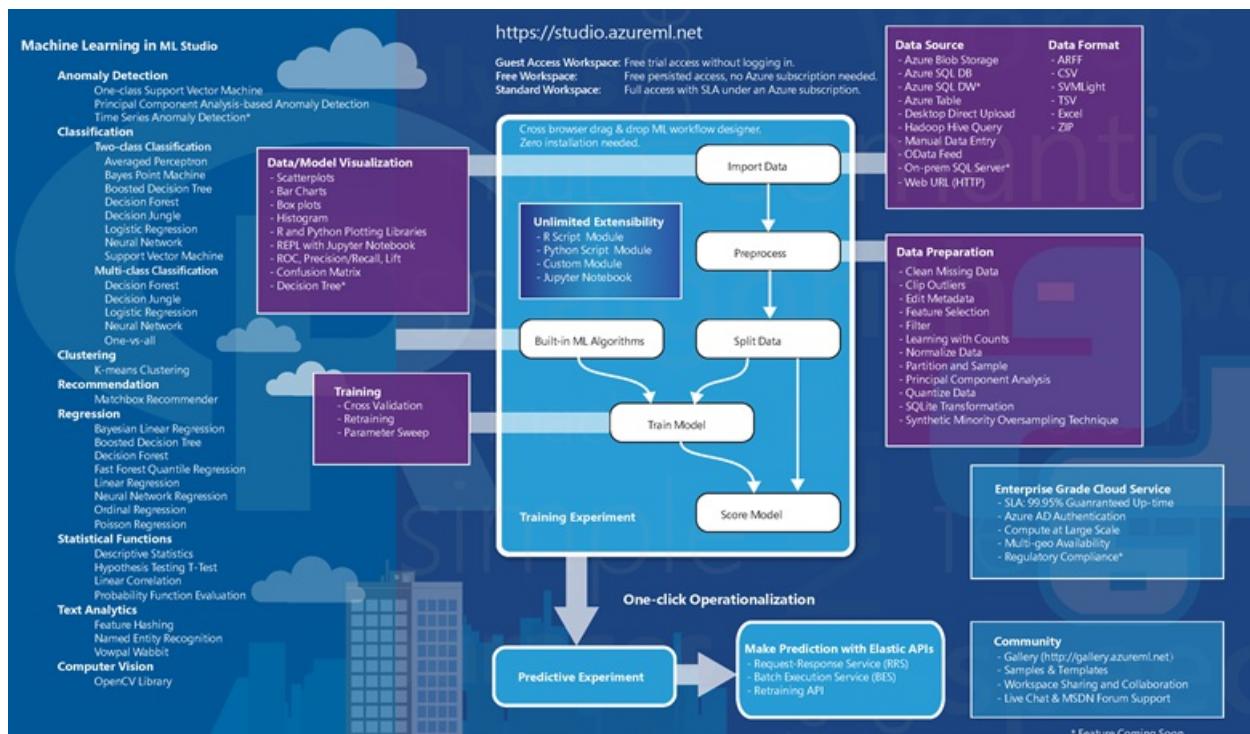
NOTE

See [Machine learning algorithm cheat sheet for Microsoft Azure Machine Learning Studio](#) for additional help in navigating through and choosing the machine learning algorithms available in Machine Learning Studio.

Download the Machine Learning Studio overview diagram

Download the **Microsoft Azure Machine Learning Studio Capabilities Overview** diagram and get a high-level view of the capabilities of Machine Learning Studio. To keep it nearby, you can print the diagram in tabloid size (11 x 17 in.).

Download the diagram here: [Microsoft Azure Machine Learning Studio Capabilities Overview](#)



More help with Machine Learning Studio

- For an overview of Microsoft Azure Machine Learning, see [Introduction to machine learning on Microsoft Azure](#).
- For an overview of Machine Learning Studio, see [What is Azure Machine Learning Studio?](#).
- For a detailed discussion of the machine learning algorithms available in Machine Learning Studio, see [How to choose algorithms for Microsoft Azure Machine Learning](#).

NOTE

Try Azure Machine Learning for free

No credit card or Azure subscription required. [Get started now.](#)

Downloadable Infographic: Machine learning basics with algorithm examples

9/25/2017 • 1 min to read • [Edit Online](#)

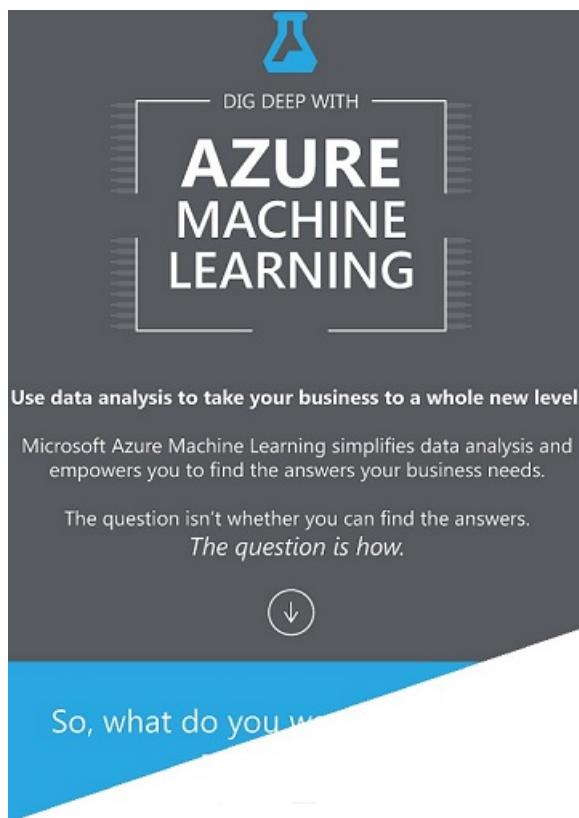
Download this easy-to-understand infographic overview of machine learning basics to learn about popular algorithms used to answer common machine learning questions. Algorithm examples help the machine learning beginner understand which algorithms to use and what they're used for.

Popular algorithms in Machine Learning Studio

Azure Machine Learning Studio comes with a large library of algorithms for predictive analytics. The infographic identifies four popular families of algorithms - regression, anomaly detection, clustering, and classification - and provides links to working examples in the [Cortana Intelligence Gallery](#). The Gallery contains example experiments and tutorials that demonstrate how these algorithms can be applied in many real-world solutions.

Download the infographic with algorithm examples

[Download: Infographic of machine learning basics with links to algorithm examples \(PDF\)](#)



More help with algorithms for beginners and advanced users

- For a deeper discussion of the different types of machine learning algorithms, how they're used, and how to choose the right one for your solution, see [How to choose algorithms for Microsoft Azure Machine Learning](#).
- For a list by category of all the machine learning algorithms available in Machine Learning Studio, see [Initialize Model](#) in the Machine Learning Studio Algorithm and Module Help.
- For a complete alphabetical list of algorithms and modules in Machine Learning Studio, see [A-Z list of Machine Learning Studio modules](#) in Machine Learning Studio Algorithm and Module Help.

- To download and print a diagram that gives an overview of the capabilities of Machine Learning Studio, see [Overview diagram of Azure Machine Learning Studio capabilities](#).
- For an overview of the Cortana Intelligence Gallery and the many community-generated resources available there, see [Share and discover resources in the Cortana Intelligence Gallery](#).

NOTE

Try Azure Machine Learning for free

No credit card or Azure subscription required. [Get started now.](#)

Azure Machine Learning frequently asked questions: Billing, capabilities, limitations, and support

11/30/2017 • 29 min to read • [Edit Online](#)

Here are some frequently asked questions (FAQs) and corresponding answers about Azure Machine Learning, a cloud service for developing predictive models and operationalizing solutions through web services. These FAQs provide questions about how to use the service, which includes the billing model, capabilities, limitations, and support.

Have a question you can't find here?

Azure Machine Learning has a forum on MSDN where members of the data science community can ask questions about Azure Machine Learning. The Azure Machine Learning team monitors the forum. Go to the [Azure Machine Learning Forum](#) to search for answers or to post a new question of your own.

General questions

What is Azure Machine Learning?

Azure Machine Learning is a fully managed service that you can use to create, test, operate, and manage predictive analytic solutions in the cloud. With only a browser, you can sign in, upload data, and immediately start machine-learning experiments. Drag-and-drop predictive modeling, a large pallet of modules, and a library of starting templates make common machine-learning tasks simple and quick. For more information, see the [Azure Machine Learning service overview](#). For an introduction to machine learning that explains key terminology and concepts, see [Introduction to Azure Machine Learning](#).

NOTE

[Try Azure Machine Learning for free](#)

No credit card or Azure subscription required. [Get started now](#).

What is Machine Learning Studio?

Machine Learning Studio is a workbench environment that you access by using a web browser. Machine Learning Studio hosts a pallet of modules in a visual composition interface that helps you build an end-to-end, data-science workflow in the form of an experiment.

For more information about Machine Learning Studio, see [What is Machine Learning Studio?](#)

What is the Machine Learning API service?

The Machine Learning API service enables you to deploy predictive models, like those that are built into Machine Learning Studio, as scalable, fault-tolerant, web services. The web services that the Machine Learning API service creates are REST APIs that provide an interface for communication between external applications and your predictive analytics models.

For more information, see [How to consume an Azure Machine Learning Web service](#).

Where are my Classic web services listed? Where are my New (Azure Resource Manager-based) web services listed?

Web services created using the Classic deployment model and web services created using the New Azure Resource

Manager deployment model are listed in the [Microsoft Azure Machine Learning Web Services](#) portal.

Classic web services are also listed in [Machine Learning Studio](#) on the **Web services** tab.

Azure Machine Learning questions

What are Azure Machine Learning web services?

Machine Learning web services provide an interface between an application and a Machine Learning workflow scoring model. An external application can use Azure Machine Learning to communicate with a Machine Learning workflow scoring model in real time. A call to a Machine Learning web service returns prediction results to an external application. To make a call to a web service, you pass an API key that was created when you deployed the web service. A Machine Learning web service is based on REST, a popular architecture choice for web programming projects.

Azure Machine Learning has two types of web services:

- Request-Response Service (RRS): A low latency, highly scalable service that provides an interface to the stateless models created and deployed by using Machine Learning Studio.
- Batch Execution Service (BES): An asynchronous service that scores a batch for data records.

There are several ways to consume the REST API and access the web service. For example, you can write an application in C#, R, or Python by using the sample code that's generated for you when you deployed the web service.

The sample code is available on:

- The Consume page for the web service in the Azure Machine Learning Web Services portal
- The API Help Page in the web service dashboard in Machine Learning Studio

You can also use the sample Microsoft Excel workbook that's created for you and is available in the web service dashboard in Machine Learning Studio.

What are the main updates to Azure Machine Learning?

For the latest updates, see [What's new in Azure Machine Learning](#).

Machine Learning Studio questions

Import and export data for Machine Learning

What data sources does Machine Learning support?

You can download data to a Machine Learning Studio experiment in three ways:

- Upload a local file as a dataset
- Use a module to import data from cloud data services
- Import a dataset saved from another experiment

To learn more about supported file formats, see [Import training data into Machine Learning Studio](#).

How large can the data set be for my modules?

Modules in Machine Learning Studio support datasets of up to 10 GB of dense numerical data for common use cases. If a module takes more than one input, the 10 GB value is the total of all input sizes. You can also sample larger datasets by using queries from Hive or Azure SQL Database, or you can use Learning by Counts preprocessing before ingestion.

The following types of data can expand to larger datasets during feature normalization and are limited to less than 10 GB:

- Sparse
- Categorical
- Strings
- Binary data

The following modules are limited to datasets less than 10 GB:

- Recommender modules
- Synthetic Minority Oversampling Technique (SMOTE) module
- Scripting modules: R, Python, SQL
- Modules where the output data size can be larger than input data size, such as Join or Feature Hashing
- Cross-validation, Tune Model Hyperparameters, Ordinal Regression, and One-vs-All Multiclass, when the number of iterations is very large

What are the limits for data upload?

For datasets that are larger than a couple GBs, upload data to Azure Storage or Azure SQL Database, or use Azure HDInsight rather than directly uploading from a local file.

Can I read data from Amazon S3?

If you have a small amount of data and want to expose it via an HTTP URL, then you can use the [Import Data](#) module. For larger amounts of data, transfer it to Azure Storage first, and then use the [Import Data](#) module to bring it into your experiment.

Is there a built-in image input capability?

You can learn about image input capability in the [Import Images](#) reference.

Modules

The algorithm, data source, data format, or data transformation operation that I am looking for isn't in Azure Machine Learning Studio. What are my options?

You can go to the [user feedback forum](#) to see feature requests that we are tracking. Add your vote to a request if a capability that you're looking for has already been requested. If the capability that you're looking for doesn't exist, create a new request. You can view the status of your request in this forum, too. We track this list closely and update the status of feature availability frequently. In addition, you can use the built-in support for R and Python to create custom transformations when needed.

Can I bring my existing code into Machine Learning Studio?

Yes, you can bring your existing R or Python code into Machine Learning Studio, run it in the same experiment with Azure Machine Learning learners, and deploy the solution as a web service via Azure Machine Learning. For more information, see [Extend your experiment with R](#) and [Execute Python machine learning scripts in Azure Machine Learning Studio](#).

Is it possible to use something like PMML to define a model?

No, Predictive Model Markup Language (PMML) is not supported. You can use custom R and Python code to define a module.

How many modules can I execute in parallel in my experiment?

You can execute up to four modules in parallel in an experiment.

Data processing

Is there an ability to visualize data (beyond R visualizations) interactively within the experiment?

Click the output of a module to visualize the data and get statistics.

When previewing results or data in a browser, the number of rows and columns is limited. Why?

Because large amounts of data might be sent to a browser, data size is limited to prevent slowing down Machine Learning Studio. To visualize all the data/results, it's better to download the data and use Excel or another tool.

Algorithms

What existing algorithms are supported in Machine Learning Studio?

Machine Learning Studio provides state-of-the-art algorithms, such as Scalable Boosted Decision trees, Bayesian Recommendation systems, Deep Neural Networks, and Decision Jungles developed at Microsoft Research. Scalable open-source machine learning packages, like Vowpal Wabbit, are also included. Machine Learning Studio supports machine learning algorithms for multiclass and binary classification, regression, and clustering. See the complete list of [Machine Learning Modules](#).

Do you automatically suggest the right Machine Learning algorithm to use for my data?

No, but Machine Learning Studio has various ways to compare the results of each algorithm to determine the right one for your problem.

Do you have any guidelines on picking one algorithm over another for the provided algorithms?

See [How to choose an algorithm](#).

Are the provided algorithms written in R or Python?

No, these algorithms are mostly written in compiled languages to provide better performance.

Are any details of the algorithms provided?

The documentation provides some information about the algorithms and parameters for tuning are described to optimize the algorithm for your use.

Is there any support for online learning?

No, currently only programmatic retraining is supported.

Can I visualize the layers of a Neural Net Model by using the built-in module?

No.

Can I create my own modules in C# or some other language?

Currently, you can only use R to create new custom modules.

R module

What R packages are available in Machine Learning Studio?

Machine Learning Studio supports more than 400 CRAN R packages today, and here is the [current list](#) of all included packages. Also, see [Extend your experiment with R](#) to learn how to retrieve this list yourself. If the package that you want is not in this list, provide the name of the package at the [user feedback forum](#).

Is it possible to build a custom R module?

Yes, see [Author custom R modules in Azure Machine Learning](#) for more information.

Is there a REPL environment for R?

No, there is no Read-Eval-Print-Loop (REPL) environment for R in the studio.

Python module

Is it possible to build a custom Python module?

Not currently, but you can use one or more [Execute Python Script](#) modules to get the same result.

Is there a REPL environment for Python?

You can use the Jupyter Notebooks in Machine Learning Studio. For more information, see [Introducing Jupyter Notebooks in Azure Machine Learning Studio](#).

Web service

Retrain

How do I retrain Azure Machine Learning models programmatically?

Use the retraining APIs. For more information, see [Retrain Machine Learning models programmatically](#). Sample code is also available in the [Microsoft Azure Machine Learning Retraining Demo](#).

Create

Can I deploy the model locally or in an application that doesn't have an Internet connection?

No.

Is there a baseline latency that is expected for all web services?

See the [Azure subscription limits](#).

Use

When would I want to run my predictive model as a Batch Execution service versus a Request Response service?

The Request Response service (RRS) is a low-latency, high-scale web service that is used to provide an interface to stateless models that are created and deployed from the experimentation environment. The Batch Execution service (BES) is a service that asynchronously scores a batch of data records. The input for BES is like data input that RRS uses. The main difference is that BES reads a block of records from a variety of sources, such as Azure Blob storage, Azure Table storage, Azure SQL Database, HDInsight (hive query), and HTTP sources. For more information, see [How to consume an Azure Machine Learning Web service](#).

How do I update the model for the deployed web service?

To update a predictive model for an already deployed service, modify and rerun the experiment that you used to author and save the trained model. After you have a new version of the trained model available, Machine Learning Studio asks you if you want to update your web service. For details about how to update a deployed web service, see [Deploy a Machine Learning web service](#).

You can also use the Retraining APIs. For more information, see [Retrain Machine Learning models programmatically](#). Sample code is also available in the [Microsoft Azure Machine Learning Retraining Demo](#).

How do I monitor my web service deployed in production?

After you deploy a predictive model, you can monitor it from the Azure Machine Learning Web Services portal. Each deployed service has its own dashboard where you can see monitoring information for that service. For more information about how to manage your deployed web services, see [Manage a Web service using the Azure Machine Learning Web Services portal](#) and [Manage an Azure Machine Learning workspace](#).

Is there a place where I can see the output of my RRS/BES?

For RRS, the web service response is typically where you see the result. You can also write it to Azure Blob storage. For BES, the output is written to a blob by default. You can also write the output to a database or table by using the [Export Data](#) module.

Can I create web services only from models that were created in Machine Learning Studio?

No, you can also create web services directly by using Jupyter Notebooks and RStudio.

Where can I find information about error codes?

See [Machine Learning Module Error Codes](#) for a list of error codes and descriptions.

Scalability

What is the scalability of the web service?

Currently, the default endpoint is provisioned with 20 concurrent RRS requests per endpoint. You can scale this to 200 concurrent requests per endpoint, and you can scale each web service to 10,000 endpoints per web service as described in [Scaling a Web Service](#). For BES, each endpoint can process 40 requests at a time, and additional requests beyond 40 requests are queued. These queued requests run automatically as the queue drains.

Are R jobs spread across nodes?

No.

How much data can I use for training?

Modules in Machine Learning Studio support datasets of up to 10 GB of dense numerical data for common use cases. If a module takes more than one input, the total size for all inputs is 10 GB. You can also sample larger datasets via Hive queries, via Azure SQL Database queries, or by preprocessing with [Learning with Counts](#) modules before ingestion.

The following types of data can expand to larger datasets during feature normalization and are limited to less than 10 GB:

- Sparse
- Categorical
- Strings
- Binary data

The following modules are limited to datasets less than 10 GB:

- Recommender modules
- Synthetic Minority Oversampling Technique (SMOTE) module
- Scripting modules: R, Python, SQL
- Modules where the output data size can be larger than input data size, such as Join or Feature Hashing
- Cross-Validate, Tune Model Hyperparameters, Ordinal Regression, and One-vs-All Multiclass, when number of iterations is very large

For datasets that are larger than a few GBs, upload data to Azure Storage or Azure SQL Database, or use HDInsight rather than directly uploading from a local file.

Are there any vector size limitations?

Rows and columns are each limited to the .NET limitation of Max Int: 2,147,483,647.

Can I adjust the size of the virtual machine that runs the web service?

No.

Security and availability

Who can access the http endpoint for the web service by default? How do I restrict access to the endpoint?

After a web service is deployed, a default endpoint is created for that service. The default endpoint can be called by using its API key. You can add more endpoints with their own keys from the Web Services portal or programmatically by using the Web Service Management APIs. Access keys are needed to make calls to the web service. For more information, see [How to consume an Azure Machine Learning Web service](#).

What happens if my Azure storage account can't be found?

Machine Learning Studio relies on a user-supplied Azure storage account to save intermediary data when it executes the workflow. This storage account is provided to Machine Learning Studio when a workspace is created. After the workspace is created, if the storage account is deleted and can no longer be found, the workspace will stop functioning, and all experiments in that workspace will fail.

If you accidentally deleted the storage account, recreate the storage account with the same name in the same region as the deleted storage account. After that, resync the access key.

What happens if my storage account access key is out of sync?

Machine Learning Studio relies on a user-supplied Azure storage account to store intermediary data when it executes the workflow. This storage account is provided to Machine Learning Studio when a workspace is created, and the access keys are associated with that workspace. If the access keys are changed after the workspace is created, the workspace can no longer access the storage account. It will stop functioning and all experiments in that workspace will fail.

If you changed storage account access keys, resync the access keys in the workspace by using the Azure portal.

Support and training

Where can I get training for Azure Machine Learning?

The [Azure Machine Learning Documentation Center](#) hosts video tutorials and how-to guides. These step-by-step guides introduce the services and explain the data science life cycle of importing data, cleaning data, building predictive models, and deploying them in production by using Azure Machine Learning.

We add new material to the Machine Learning Center on an ongoing basis. You can submit requests for additional learning material on Machine Learning Center at the [user feedback forum](#).

You can also find training at [Microsoft Virtual Academy](#).

How do I get support for Azure Machine Learning?

To get technical support for Azure Machine Learning, go to [Azure Support](#), and select **Machine Learning**.

Azure Machine Learning also has a community forum on MSDN where you can ask questions about Azure Machine Learning. The Azure Machine Learning team monitors the forum. Go to [Azure Forum](#).

Billing questions

How does Machine Learning billing work?

Azure Machine Learning has two components: Machine Learning Studio and Machine Learning web services.

While you are evaluating Machine Learning Studio, you can use the Free billing tier. The Free tier also lets you deploy a Classic web service that has limited capacity.

If you decide that Azure Machine Learning meets your needs, you can sign up for the Standard tier. To sign up, you must have a Microsoft Azure subscription.

In the Standard tier, you are billed monthly for each workspace that you define in Machine Learning Studio. When you run an experiment in the studio, you are billed for compute resources when you are running an experiment.

When you deploy a Classic web service, transactions and compute hours are billed on the Pay As You Go basis.

New (Resource Manager-based) web services introduce billing plans that allow for more predictability in costs. Tiered pricing offers discounted rates to customers who need a large amount of capacity.

When you create a plan, you commit to a fixed cost that comes with an included quantity of API compute hours and API transactions. If you need more included quantities, you can add instances to your plan. If you need a lot more included quantities, you can choose a higher tier plan that provides considerably more included quantities and a better discounted rate.

After the included quantities in existing instances are used up, additional usage is charged at the overage rate that's associated with the billing plan tier.

NOTE

Included quantities are reallocated every 30 days, and unused included quantities do not roll over to the next period.

For additional billing and pricing information, see [Machine Learning Pricing](#).

Does Machine Learning have a free trial?

Azure Machine Learning has a free subscription option that's explained in [Machine Learning Pricing](#). Machine Learning Studio has an eight-hour quick evaluation trial that's available when you sign in to [Machine Learning Studio](#).

In addition, when you sign up for an Azure free trial, you can try any Azure services for a month. To learn more about the Azure free trial, visit [Azure free trial FAQ](#).

What is a transaction?

A transaction represents an API call that Azure Machine Learning responds to. Transactions from Request-Response Service (RRS) and Batch Execution Service (BES) calls are aggregated and charged against your billing plan.

Can I use the included transaction quantities in a plan for both RRS and BES transactions?

Yes, your transactions from your RRS and BES are aggregated and charged against your billing plan.

What is an API compute hour?

An API compute hour is the billing unit for the time that API calls take to run by using Machine Learning compute resources. All your calls are aggregated for billing purposes.

How long does a typical production API call take?

Production API call times can vary significantly, generally ranging from hundreds of milliseconds to a few seconds. Some API calls might require minutes depending on the complexity of the data processing and machine-learning model. The best way to estimate production API call times is to benchmark a model on the Machine Learning service.

What is a Studio compute hour?

A Studio compute hour is the billing unit for the aggregate time that your experiments use compute resources in studio.

In New (Azure Resource Manager-based) web services, what is the Dev/Test tier meant for?

Resource Manager-based web services provide multiple tiers that you can use to provision your billing plan. The Dev/Test pricing tier provides limited, included quantities that allow you to test your experiment as a web service without incurring costs. You have the opportunity to see how it works.

Are there separate storage charges?

The Machine Learning Free tier does not require or allow separate storage. The Machine Learning Standard tier requires users to have an Azure storage account. Azure Storage is [billed separately](#).

Does Machine Learning support high availability?

Yes. For details, see [Machine Learning Pricing](#) for a description of the service level agreement (SLA).

What specific kind of compute resources will my production API calls be run on?

The Machine Learning service is a multitenant service. Actual compute resources that are used on the back end vary and are optimized for performance and predictability.

Management of New (Resource Manager-based) web services

What happens if I delete my plan?

The plan is removed from your subscription, and you are billed for prorated usage.

NOTE

You cannot delete a plan that a web service is using. To delete the plan, you must either assign a new plan to the web service or delete the web service.

What is a plan instance?

A plan instance is a unit of included quantities that you can add to your billing plan. When you select a billing tier for your billing plan, it comes with one instance. If you need more included quantities, you can add instances of the selected billing tier to your plan.

How many plan instances can I add?

You can have one instance of the Dev/Test pricing tier in a subscription.

For Standard S1, Standard S2, and Standard S3 tiers, you can add as many as necessary.

NOTE

Depending on your anticipated usage, it might be more cost effective to upgrade to a tier that has more included quantities rather than add instances to the current tier.

What happens when I change plan tiers (upgrade / downgrade)?

The old plan is deleted and the current usage is billed on a prorated basis. A new plan with the full included quantities of the upgraded/downgraded tier is created for the rest of the period.

NOTE

Included quantities are allocated per period, and unused quantities do not roll over.

What happens when I increase the instances in a plan?

Quantities are included on a prorated basis and may take 24 hours to be effective.

What happens when I delete an instance of a plan?

The instance is removed from your subscription, and you are billed for prorated usage.

Sign up for New (Resource Manager-based) web services plans

How do I sign up for a plan?

You have two ways to create billing plans.

When you first deploy a Resource Manager-based web service, you can choose an existing plan or create a new plan.

Plans that you create in this manner are in your default region, and your web service will be deployed to that region.

If you want to deploy services to regions other than your default region, you may want to define your billing plans before you deploy your service.

In that case, you can sign in to the Azure Machine Learning Web Services portal, and go to the Plans page. From there, you can add plans, delete plans, and modify existing plans.

Which plan should I choose to start off with?

We recommend you that you start with the Standard S1 tier and monitor your service for usage. If you find that you are using your included quantities rapidly, you can add instances or move to a higher tier and get better discounted rates. You can adjust your billing plan as needed throughout your billing cycle.

Which regions are the new plans available in?

The new billing plans are available in the three production regions in which we support the new web services:

- South Central US
- West Europe
- South East Asia

I have web services in multiple regions. Do I need a plan for every region?

Yes. Plan pricing varies by region. When you deploy a web service to another region, you need to assign it a plan that is specific to that region. For more information, see [Products available by region](#).

New web services: Overages

How do I check if I exceeded my web service usage?

You can view the usage on all your plans on the Plans page in the Azure Machine Learning Web Services portal. Sign in to the portal, and then click the **Plans** menu option.

In the **Transactions** and **Compute** columns of the table, you can see the included quantities of the plan and the percentage used.

What happens when I use up the include quantities in the Dev/Test pricing tier?

Services that have a Dev/Test pricing tier assigned to them are stopped until the next period or until you move them to a paid tier.

For Classic web services and overages of New (Resource Manager-based) web services, how are prices calculated for Request Response (RRS) and Batch (BES) workloads?

For an RRS workload, you are charged for every API transaction call that you make and for the compute time that's associated with those requests. Your RRS production API transaction costs are calculated as the total number of API calls that you make multiplied by the price per 1,000 transactions (prorated by individual transaction). Your RRS API production API compute hour costs are calculated as the amount of time required for each API call to run, multiplied by the total number of API transactions, multiplied by the price per production API compute hour.

For example, for Standard S1 overage, 1,000,000 API transactions that take 0.72 seconds each to run would result

in $(1,000,000 * \$0.50/1K \text{ API transactions})$ in \$500 in production API transaction costs and $(1,000,000 * 0.72 \text{ sec} * \$2/\text{hr})$ \$400 in production API compute hours, for a total of \$900.

For a BES workload, you are charged in the same manner. However, the API transaction costs represent the number of batch jobs that you submit, and the compute costs represent the compute time that's associated with those batch jobs. Your BES production API transaction costs are calculated as the total number of jobs submitted multiplied by the price per 1,000 transactions (prorated by individual transaction). Your BES API production API compute hour costs are calculated as the amount of time required for each row in your job to run multiplied by the total number of rows in your job multiplied by the total number of jobs multiplied by the price per production API compute hour. When you use the Machine Learning calculator, the transaction meter represents the number of jobs that you plan to submit, and the time-per-transaction field represents the combined time that's needed for all rows in each job to run.

For example, assume Standard S1 overage, and you submit 100 jobs per day that each consist of 500 rows that take 0.72 seconds each. Your monthly overage costs would be $(100 \text{ jobs per day} = 3,100 \text{ jobs/mo} * \$0.50/1K \text{ API transactions})$ \$1.55 in production API transaction costs and $(500 \text{ rows} * 0.72 \text{ sec} * 3,100 \text{ Jobs} * \$2/\text{hr})$ \$620 in production API compute hours, for a total of \$621.55.

Azure Machine Learning Classic web services

Is Pay As You Go still available?

Yes, Classic web services are still available in Azure Machine Learning.

Azure Machine Learning Free and Standard tier

What is included in the Azure Machine Learning Free tier?

The Azure Machine Learning Free tier is intended to provide an in-depth introduction to the Azure Machine Learning Studio. All you need is a Microsoft account to sign up. The Free tier includes free access to one Azure Machine Learning Studio workspace per [Microsoft account](#). In this tier, you can use up to 10 GB of storage and operationalize models as staging APIs. Free tier workloads are not covered by an SLA and are intended for development and personal use only.

Free tier workspaces have the following limitations:

- Workloads can't access data by connecting to an on-premises server that runs SQL Server.
- You cannot deploy New Resource Manager base web services.

What is included in the Azure Machine Learning Standard tier and plans?

The Azure Machine Learning Standard tier is a paid production version of Azure Machine Learning Studio. The Azure Machine Learning Studio monthly fee is billed on a per workspace per month basis and prorated for partial months. Azure Machine Learning Studio experiment hours are billed per compute hour for active experimentation. Billing is prorated for partial hours.

The Azure Machine Learning API service is billed depending on whether it's a Classic web service or a New (Resource Manager-based) web service.

The following charges are aggregated per workspace for your subscription.

- Machine Learning Workspace Subscription: The Machine Learning workspace subscription is a monthly fee that provides access to a Machine Learning Studio workspace. The subscription is required to run experiments in the studio and to utilize the production APIs.
- Studio Experiment hours: This meter aggregates all compute charges that are accrued by running experiments in Machine Learning Studio and running production API calls in the staging environment.
- Access data by connecting to an on-premises server that runs SQL Server in your models for your training and scoring.
- For Classic web services:

- Production API Compute Hours: This meter includes compute charges that are accrued by web services running in production.
- Production API Transactions (in 1000s): This meter includes charges that are accrued per call to your production web service.

Apart from the preceding charges, in the case of Resource Manager-based web service, charges are aggregated to the selected plan:

- Standard S1/S2/S3 API Plan (Units): This meter represents the type of instance that's selected for Resource Manager-based web services.
- Standard S1/S2/S3 Overage API Compute Hours: This meter includes compute charges that are accrued by Resource Manager-based web services that run in production after the included quantities in existing instances are used up. The additional usage is charged at the overage rate that's associated with S1/S2/S3 plan tier.
- Standard S1/S2/S3 Overage API Transactions (in 1,000s): This meter includes charges that are accrued per call to your production Resource Manager-based web service after the included quantities in existing instances are used up. The additional usage is charged at the overage rate associated with S1/S2/S3 plan tier.
- Included Quantity API Compute Hours: With Resource Manager-based web services, this meter represents the included quantity of API compute hours.
- Included Quantity API Transactions (in 1,000s): With Resource Manager-based web services, this meter represents the included quantity of API transactions.

How do I sign up for Azure Machine Learning Free tier?

All you need is a Microsoft account. Go to [Azure Machine Learning home](#), and then click **Start Now**. Sign in with your Microsoft account and a workspace in Free tier is created for you. You can start to explore and create Machine Learning experiments right away.

How do I sign up for Azure Machine Learning Standard tier?

You must first have access to an Azure subscription to create a Standard Machine Learning workspace. You can sign up for a 30-day free trial Azure subscription and later upgrade to a paid Azure subscription, or you can purchase a paid Azure subscription outright. You can then create a Machine Learning workspace from the Microsoft Azure portal after you gain access to the subscription. View the [step-by-step instructions](#).

Alternatively, you can be invited by a Standard Machine Learning workspace owner to access the owner's workspace.

Can I specify my own Azure Blob storage account to use with the Free tier?

No, the Standard tier is equivalent to the version of the Machine Learning service that was available before the tiers were introduced.

Can I deploy my machine learning models as APIs in the Free tier?

Yes, you can operationalize machine learning models to staging API services as part of the Free tier. To put the staging API service into production and get a production endpoint for the operationalized service, you must use the Standard tier.

What is the difference between Azure free trial and Azure Machine Learning Free tier?

The [Microsoft Azure free trial](#) offers credits that you can apply to any Azure service for one month. The Azure Machine Learning Free tier offers continuous access specifically to Azure Machine Learning for non-production workloads.

How do I move an experiment from the Free tier to the Standard tier?

To copy your experiments from the Free tier to the Standard tier:

1. Sign in to Azure Machine Learning Studio, and make sure that you can see both the Free workspace and the Standard workspace in the workspace selector in the top navigation bar.
2. Switch to Free workspace if you are in the Standard workspace.
3. In the experiment list view, select an experiment that you'd like to copy, and then click the **Copy** command button.
4. Select the Standard workspace from the dialog box that opens, and then click the **Copy** button. All the associated datasets, trained model, etc. are copied together with the experiment into the Standard workspace.
5. You need to rerun the experiment and republish your web service in the Standard workspace.

Studio workspace

Will I see different bills for different workspaces?

Workspace charges are broken out separately for each applicable meter on a single bill.

What specific kind of compute resources will my experiments be run on?

The Machine Learning service is a multitenant service. Actual compute resources that are used on the back end vary and are optimized for performance and predictability.

Guest Access

What is Guest Access to Azure Machine Learning Studio?

Guest Access is a restricted trial experience. You can create and run experiments in Azure Machine Learning Studio at no cost and without authentication. Guest sessions are non-persistent (cannot be saved) and limited to eight hours. Other limitations include lack of support for R and Python, lack of staging APIs, and restricted dataset size and storage capacity. By comparison, users who choose to sign in with a Microsoft account have full access to the Free tier of Machine Learning Studio that's described previously, which includes a persistent workspace and more comprehensive capabilities. To choose your free Machine Learning experience, click **Get started** on <https://studio.azureml.net>, and then select **Guest Access** or sign in with a Microsoft account.

What's New in Azure Machine Learning

12/8/2017 • 1 min to read • [Edit Online](#)

The March 2017 release of Microsoft Azure Machine Learning updates provides the following feature:

- Dedicated Capacity for Azure Machine Learning BES Jobs

Machine Learning Batch Pool processing uses the [Azure Batch](#) service to provide customer-managed scale for the Azure Machine Learning Batch Execution Service. Batch Pool processing allows you to create Azure Batch pools on which you can submit batch jobs and have them execute in a predictable manner.

For more information, see [Azure Batch service for Machine Learning jobs](#).

The August 2016 release of Microsoft Azure Machine Learning updates provide the following features:

- Classic Web services can now be managed in the new [Microsoft Azure Machine Learning Web Services](#) portal that provides one place to manage all aspects of your Web service.
 - Which provides web service [usage statistics](#).
 - Simplifies testing of Azure Machine Learning Remote-Request calls using sample data.
 - Provides a new Batch Execution Service test page with sample data and job submission history.
 - Provides easier endpoint management.

The July 2016 release of Microsoft Azure Machine Learning updates provide the following features:

- Web services are now managed as Azure resources managed through [Azure Resource Manager](#) interfaces, allowing for the following enhancements:
 - There are new [REST APIs](#) to deploy and manage your Resource Manager based Web services.
 - There is a new [Microsoft Azure Machine Learning Web Services](#) portal that provides one place to manage all aspects of your Web service.
- Incorporates a new subscription-based, multi-region web service deployment model using Resource Manager based APIs leveraging the Resource Manager Resource Provider for Web Services.
- Introduces new [pricing plans](#) and plan management capabilities using the new Resource Manager RP for Billing.
 - You can now [deploy your web service to multiple regions](#) without needing to create a subscription in each region.
- Provides web service [usage statistics](#).
- Simplifies testing of Azure Machine Learning Remote-Request calls using sample data.
- Provides a new Batch Execution Service test page with sample data and job submission history.

In addition, the Machine Learning Studio has been updated to allow you to deploy to the new Web service model or continue to deploy to the classic Web service model.

Machine learning tutorial: Create your first data science experiment in Azure Machine Learning Studio

9/25/2017 • 16 min to read • [Edit Online](#)

If you've never used **Azure Machine Learning Studio** before, this tutorial is for you.

In this tutorial, we'll walk through how to use Studio for the first time to create a machine learning experiment. The experiment will test an analytical model that predicts the price of an automobile based on different variables such as make and technical specifications.

NOTE

This tutorial shows you the basics of how to drag-and-drop modules onto your experiment, connect them together, run the experiment, and look at the results. We're not going to discuss the general topic of machine learning or how to select and use the 100+ built-in algorithms and data manipulation modules included in Studio.

If you're new to machine learning, the video series [Data Science for Beginners](#) might be a good place to start. This video series is a great introduction to machine learning using everyday language and concepts.

If you're familiar with machine learning, but you're looking for more general information about Machine Learning Studio, and the machine learning algorithms it contains, here are some good resources:

- [What is Machine Learning Studio?](#) - This is a high-level overview of Studio.
- [Machine learning basics with algorithm examples](#) - This infographic is useful if you want to learn more about the different types of machine learning algorithms included with Machine Learning Studio.
- [Machine Learning Guide](#) - This guide covers similar information as the infographic above, but in an interactive format.
- [Machine learning algorithm cheat sheet](#) and [How to choose algorithms for Microsoft Azure Machine Learning](#) - This downloadable poster and accompanying article discuss the Studio algorithms in depth.
- [Machine Learning Studio: Algorithm and Module Help](#) - This is the complete reference for all Studio modules, including machine learning algorithms,

NOTE

[Try Azure Machine Learning for free](#)

No credit card or Azure subscription required. [Get started now.](#)

How does Machine Learning Studio help?

Machine Learning Studio makes it easy to set up an experiment using drag-and-drop modules preprogrammed with predictive modeling techniques.

Using an interactive, visual workspace, you drag-and-drop **datasets** and analysis **modules** onto an interactive canvas. You connect them together to form an **experiment** that you run in Machine Learning Studio. You **create a model, train the model, and score and test the model**.

You can iterate on your model design, editing the experiment and running it until it gives you the results you're looking for. When your model is ready, you can publish it as a **web service** so that others can send it new data and get predictions in return.

Open Machine Learning Studio

To get started with Studio, go to <https://studio.azureml.net>. If you've signed into Machine Learning Studio before, click **Sign In**. Otherwise, click **Sign up here** and choose between free and paid options.

Sign in to Machine Learning Studio

Five steps to create an experiment

In this machine learning tutorial, you'll follow five basic steps to build an experiment in Machine Learning Studio to create, train, and score your model:

- **Create a model**
 - [Step 1: Get data](#)
 - [Step 2: Prepare the data](#)
 - [Step 3: Define features](#)
- **Train the model**
 - [Step 4: Choose and apply a learning algorithm](#)
- **Score and test the model**
 - [Step 5: Predict new automobile prices](#)

TIP

You can find a working copy of the following experiment in the [Cortana Intelligence Gallery](#). Go to [Your first data science experiment - Automobile price prediction](#) and click **Open in Studio** to download a copy of the experiment into your Machine Learning Studio workspace.

Step 1: Get data

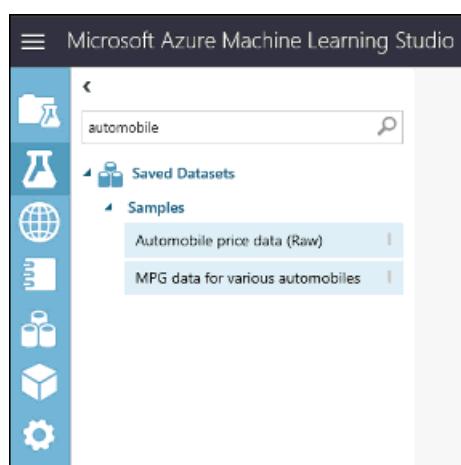
The first thing you need to perform machine learning is data. There are several sample datasets included with Machine Learning Studio that you can use, or you can import data from many sources. For this example, we'll use the sample dataset, **Automobile price data (Raw)**, that's included in your workspace. This dataset includes entries for various individual automobiles, including information such as make, model, technical specifications, and price.

Here's how to get the dataset into your experiment.

1. Create a new experiment by clicking **+NEW** at the bottom of the Machine Learning Studio window, select **EXPERIMENT**, and then select **Blank Experiment**.
2. The experiment is given a default name that you can see at the top of the canvas. Select this text and rename it to something meaningful, for example, **Automobile price prediction**. The name doesn't need to be unique.

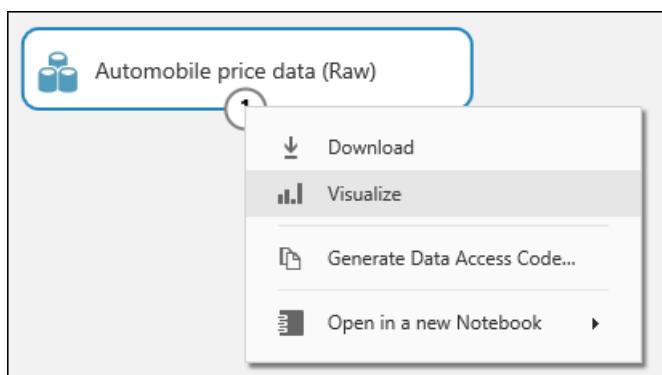


3. To the left of the experiment canvas is a palette of datasets and modules. Type **automobile** in the Search box at the top of this palette to find the dataset labeled **Automobile price data (Raw)**. Drag this dataset to the experiment canvas.



Find the automobile dataset and drag it onto the experiment canvas

To see what this data looks like, click the output port at the bottom of the automobile dataset, and then select **Visualize**.



Click the output port and select "Visualize"

TIP

Datasets and modules have input and output ports represented by small circles - input ports at the top, output ports at the bottom. To create a flow of data through your experiment, you'll connect an output port of one module to an input port of another. At any time, you can click the output port of a dataset or module to see what the data looks like at that point in the data flow.

In this sample dataset, each instance of an automobile appears as a row, and the variables associated with each automobile appear as columns. Given the variables for a specific automobile, we're going to try to predict the price in far-right column (column 26, titled "price").

Automobile price prediction ➤ Automobile price data (Raw) ➤ dataset												
rows	columns											
205	26											
	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-type	peak-rpm	city-mpg	highway-mpg	price
view as												
3			alfa-romero	gas	std	two	convertible		5000	21	27	13495
3			alfa-romero	gas	std	two	convertible		5000	21	27	16500
1			alfa-romero	gas	std	two	hatchback	54	5000	19	26	16500
2	164		audi	gas	std	four	sedan	102	5500	24	30	13950
2	164		audi	gas	std	four	sedan	115	5500	18	22	17450
2			audi	gas	std	two	sedan	110	5500	19	25	15250
1	158		audi	gas	std	four		110	5500	19	25	17710
1			audi	gas	std	four		110	5500	19	25	18920
1	158		audi	gas	turbo	four		140	5500	17	20	23875
0			audi	gas	turbo	two		160	5500	16	22	
2	192		bmw	gas	std	two		101	5800	23	29	16430
0	192		bmw	gas	std	four		101	5800	23	29	16925
0	188		bmw	gas	std	two		121	4250	21	28	20970
0	188		bmw	gas	std	four		121	4250	21	28	21105
1			bmw	gas	std	four		121	4250	20	25	24565

View the automobile data in the data visualization window

Close the visualization window by clicking the "x" in the upper-right corner.

Step 2: Prepare the data

A dataset usually requires some preprocessing before it can be analyzed. For example, you might have noticed the missing values present in the columns of various rows. These missing values need to be cleaned so the model can analyze the data correctly. In our case, we'll remove any rows that have missing values. Also, the **normalized-losses** column has a large proportion of missing values, so we'll exclude that column from the model altogether.

TIP

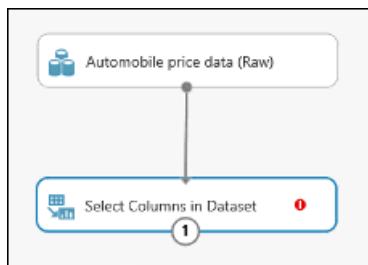
Cleaning the missing values from input data is a prerequisite for using most of the modules.

First we add a module that removes the **normalized-losses** column completely, and then we add another module that removes any row that has missing data.

1. Type **select columns** in the Search box at the top of the module palette to find the [Select Columns in](#)

[Dataset](#) module, then drag it to the experiment canvas. This module allows us to select which columns of data we want to include or exclude in the model.

2. Connect the output port of the **Automobile price data (Raw)** dataset to the input port of the [Select Columns in Dataset](#) module.



Add the "Select Columns in Dataset" module to the experiment canvas and connect it

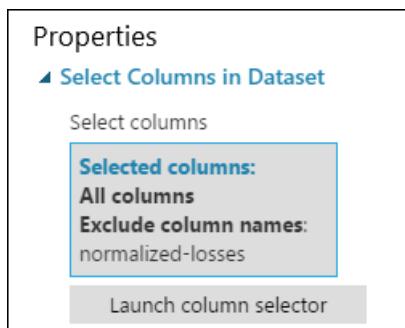
3. Click the [Select Columns in Dataset](#) module and click **Launch column selector** in the **Properties** pane.

- On the left, click **With rules**
- Under **Begin With**, click **All columns**. This directs [Select Columns in Dataset](#) to pass through all the columns (except those columns we're about to exclude).
- From the drop-downs, select **Exclude** and **column names**, and then click inside the text box. A list of columns is displayed. Select **normalized-losses**, and it's added to the text box.
- Click the check mark (OK) button to close the column selector (on the lower-right).



Launch the column selector and exclude the "normalized-losses" column

Now the properties pane for **Select Columns in Dataset** indicates that it will pass through all columns from the dataset except **normalized-losses**.



The properties pane shows that the "normalized-losses" column is excluded

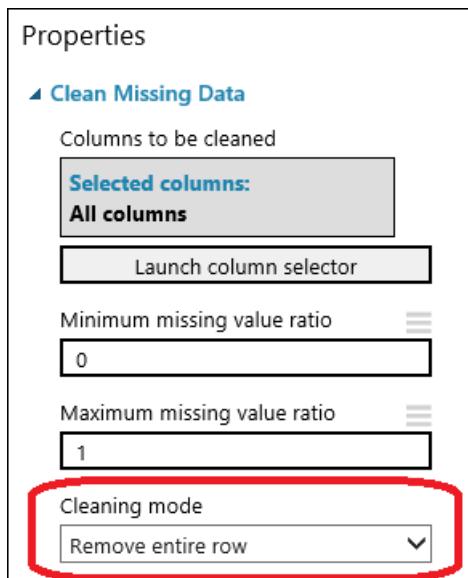
TIP

You can add a comment to a module by double-clicking the module and entering text. This can help you see at a glance what the module is doing in your experiment. In this case double-click the [Select Columns in Dataset](#) module and type the comment "Exclude normalized losses."

![Double-click a module to add a comment][add-comment]

Double-click a module to add a comment

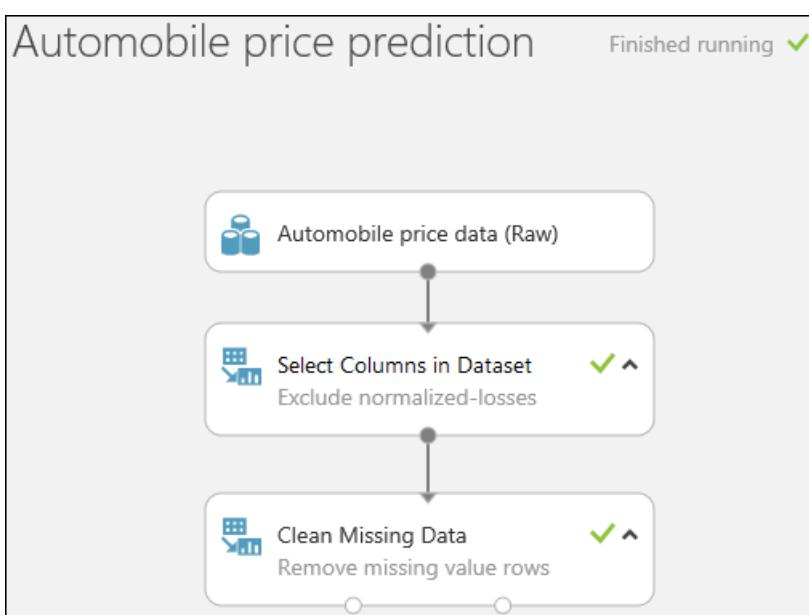
1. Drag the [Clean Missing Data](#) module to the experiment canvas and connect it to the [Select Columns in Dataset](#) module. In the **Properties** pane, select **Remove entire row** under **Cleaning mode**. This directs [Clean Missing Data](#) to clean the data by removing rows that have any missing values. Double-click the module and type the comment "Remove missing value rows."



Set the cleaning mode to "Remove entire row" for the "Clean Missing Data" module

2. Run the experiment by clicking **RUN** at the bottom of the page.

When the experiment has finished running, all the modules have a green check mark to indicate that they finished successfully. Notice also the **Finished running** status in the upper-right corner.



After running it, the experiment should look something like this

TIP

Why did we run the experiment now? By running the experiment, the column definitions for our data pass from the dataset, through the [Select Columns in Dataset](#) module, and through the [Clean Missing Data](#) module. This means that any modules we connect to [Clean Missing Data](#) will also have this same information.

All we have done in the experiment up to this point is clean the data. If you want to view the cleaned dataset, click the left output port of the [Clean Missing Data](#) module and select **Visualize**. Notice that the **normalized-losses** column is no longer included, and there are no missing values.

Now that the data is clean, we're ready to specify what features we're going to use in the predictive model.

Step 3: Define features

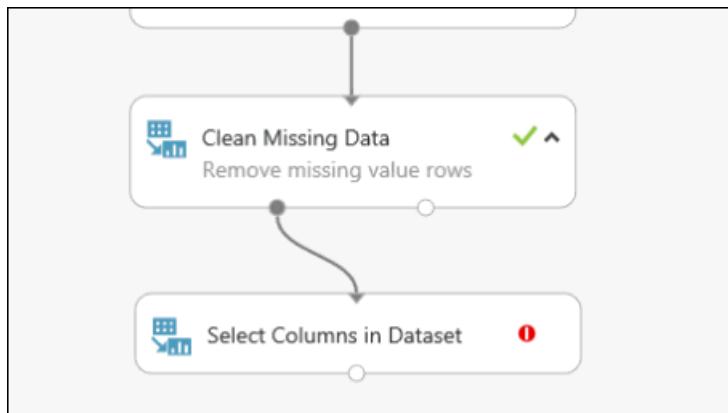
In machine learning, *features* are individual measurable properties of something you're interested in. In our dataset, each row represents one automobile, and each column is a feature of that automobile.

Finding a good set of features for creating a predictive model requires experimentation and knowledge about the problem you want to solve. Some features are better for predicting the target than others. Also, some features have a strong correlation with other features and can be removed. For example, city-mpg and highway-mpg are closely related so we can keep one and remove the other without significantly affecting the prediction.

Let's build a model that uses a subset of the features in our dataset. You can come back later and select different features, run the experiment again, and see if you get better results. But to start, let's try the following features:

make, body-style, wheel-base, engine-size, horsepower, peak-rpm, highway-mpg, price

1. Drag another [Select Columns in Dataset](#) module to the experiment canvas. Connect the left output port of the [Clean Missing Data](#) module to the input of the [Select Columns in Dataset](#) module.



Connect the "Select Columns in Dataset" module to the "Clean Missing Data" module

2. Double-click the module and type "Select features for prediction."
3. Click **Launch column selector** in the **Properties** pane.
4. Click **With rules**.
5. Under **Begin With**, click **No columns**. In the filter row, select **Include** and **column names** and select our list of column names in the text box. This directs the module to not pass through any columns (features) except the ones that we specify.
6. Click the check mark (OK) button.

The dialog shows a list of columns: make, body-style, wheel-base, engine-size, horsepower, peak-rpm, highway-mpg, and price.

Select the columns (features) to include in the prediction

This produces a filtered dataset containing only the features we want to pass to the learning algorithm we'll use in the next step. Later, you can return and try again with a different selection of features.

Step 4: Choose and apply a learning algorithm

Now that the data is ready, constructing a predictive model consists of training and testing. We'll use our data to train the model, and then we'll test the model to see how closely it's able to predict prices.

Classification and *regression* are two types of supervised machine learning algorithms. Classification predicts an answer from a defined set of categories, such as a color (red, blue, or green). Regression is used to predict a number.

Because we want to predict price, which is a number, we'll use a regression algorithm. For this example, we'll use a simple *linear regression* model.

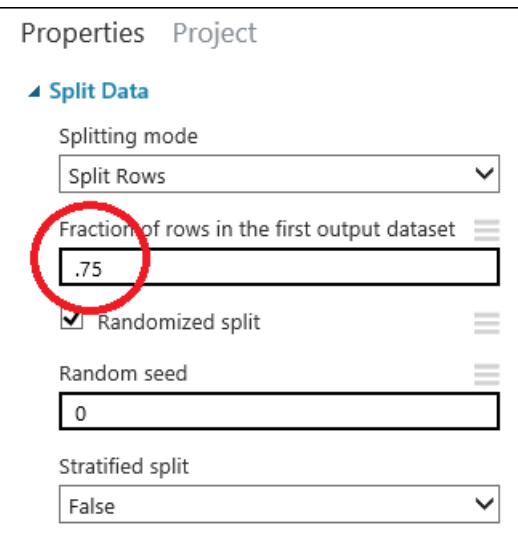
TIP

If you want to learn more about different types of machine learning algorithms and when to use them, you might view the first video in the Data Science for Beginners series, [The five questions data science answers](#). You might also look at the infographic [Machine learning basics with algorithm examples](#), or check out the [Machine learning algorithm cheat sheet](#).

We train the model by giving it a set of data that includes the price. The model scans the data and look for correlations between an automobile's features and its price. Then we'll test the model - we'll give it a set of features for automobiles we're familiar with and see how close the model comes to predicting the known price.

We'll use our data for both training the model and testing it by splitting the data into separate training and testing datasets.

1. Select and drag the [Split Data](#) module to the experiment canvas and connect it to the last [Select Columns in Dataset](#) module.
2. Click the [Split Data](#) module to select it. Find the **Fraction of rows in the first output dataset** (in the **Properties** pane to the right of the canvas) and set it to 0.75. This way, we'll use 75 percent of the data to train the model, and hold back 25 percent for testing (later, you can experiment with using different percentages).

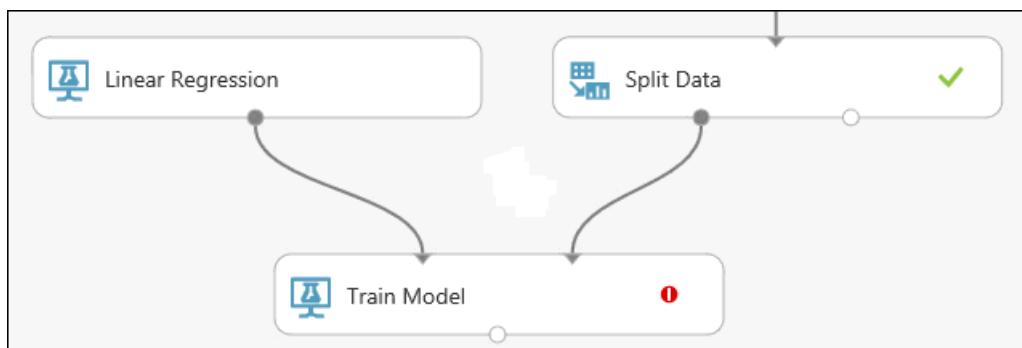


Set the split fraction of the "Split Data" module to 0.75

TIP

By changing the **Random seed** parameter, you can produce different random samples for training and testing. This parameter controls the seeding of the pseudo-random number generator.

3. Run the experiment. When the experiment is run, the [Select Columns in Dataset](#) and [Split Data](#) modules pass column definitions to the modules we'll be adding next.
4. To select the learning algorithm, expand the **Machine Learning** category in the module palette to the left of the canvas, and then expand **Initialize Model**. This displays several categories of modules that can be used to initialize machine learning algorithms. For this experiment, select the [Linear Regression](#) module under the **Regression** category, and drag it to the experiment canvas. (You can also find the module by typing "linear regression" in the palette Search box.)
5. Find and drag the [Train Model](#) module to the experiment canvas. Connect the output of the [Linear Regression](#) module to the left input of the [Train Model](#) module, and connect the training data output (left port) of the [Split Data](#) module to the right input of the [Train Model](#) module.



Connect the "Train Model" module to both the "Linear Regression" and "Split Data" modules

6. Click the [Train Model](#) module, click **Launch column selector** in the **Properties** pane, and then select the **price** column. This is the value that our model is going to predict.

You select the **price** column in the column selector by moving it from the **Available columns** list to the **Selected columns** list.

Select a single column

BY NAME
WITH RULES

AVAILABLE COLUMNS

All Types search columns

SELECTED COLUMNS

All Types search columns

make
body-style
wheel-base
engine-size
horsepower
peak-rpm
highway-mpg

>

<

7 columns available

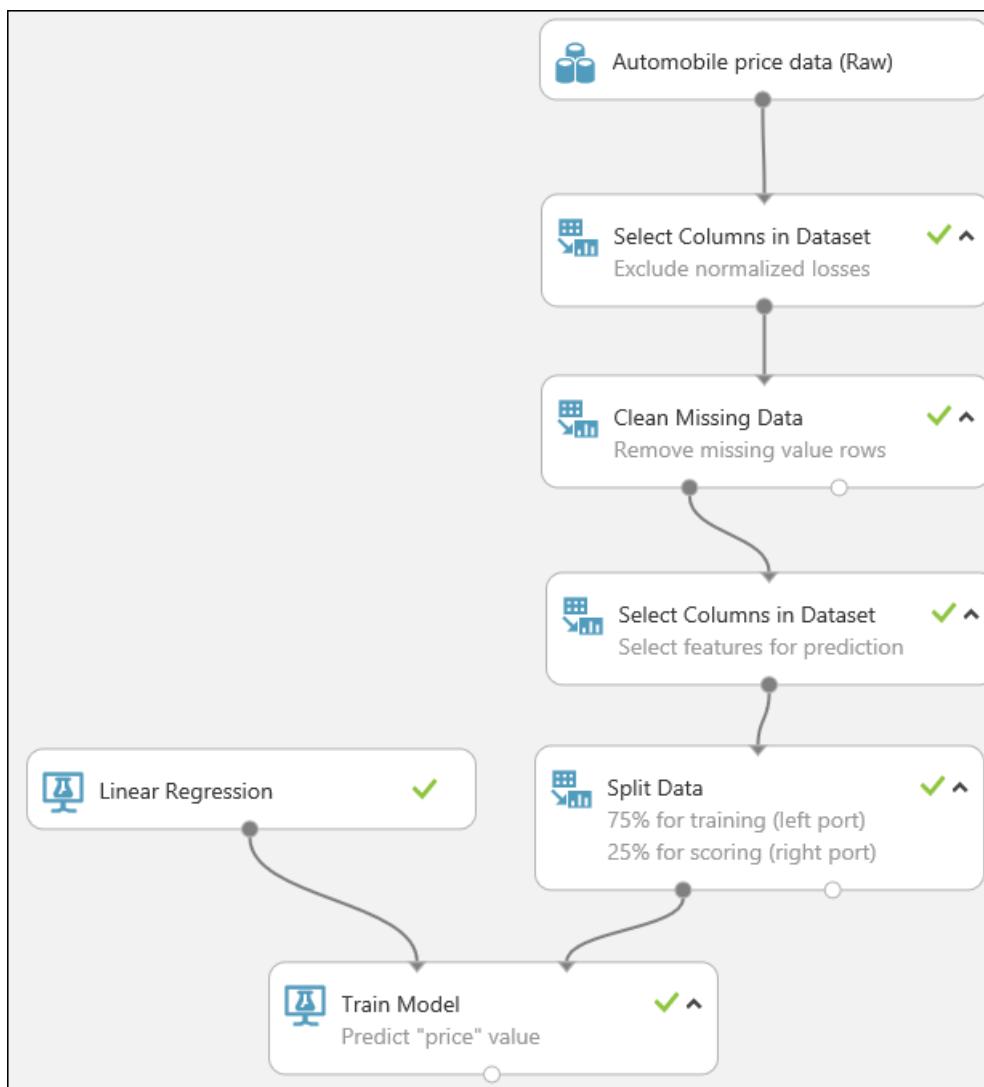
1 columns selected

✓

Select the price column for the "Train Model" module

7. Run the experiment.

We now have a trained regression model that can be used to score new automobile data to make price predictions.

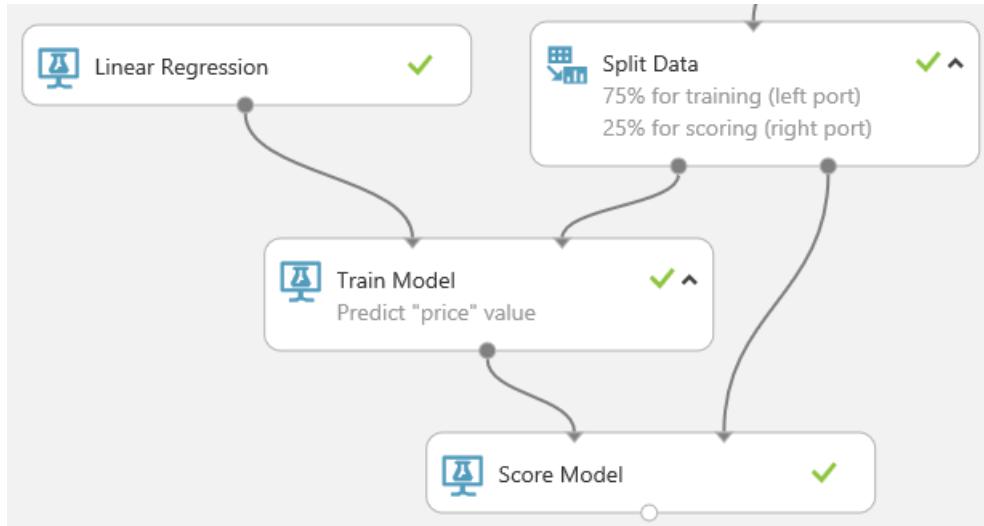


After running, the experiment should now look something like this

Step 5: Predict new automobile prices

Now that we've trained the model using 75 percent of our data, we can use it to score the other 25 percent of the data to see how well our model functions.

1. Find and drag the [Score Model](#) module to the experiment canvas. Connect the output of the [Train Model](#) module to the left input port of [Score Model](#). Connect the test data output (right port) of the [Split Data](#) module to the right input port of [Score Model](#).



Connect the "Score Model" module to both the "Train Model" and "Split Data" modules

2. Run the experiment and view the output from the [Score Model](#) module (click the output port of [Score Model](#) and select **Visualize**). The output shows the predicted values for price and the known values from the test data.

Simple test experiment - Copy > Score Model > Scored dataset										
rows	columns								Known values	Predicted values
48	9	make	body-style	wheel-base	engine-size	horsepower	peak-rpm	highway-mpg	price	Scored Labels
view as	grid	subaru	sedan	97	108	111	4800	29	11259	10286.204819
	list	mitsubishi	hatchback	93.7	92	68	5500	38	6669	5446.847864
	grid	dodge	hatchback	93.7	90	68	5500	38	6229	6344.800711
	list	honda	hatchback	86.6	92	76	6000	38	6855	5528.302953
	grid	alfa-romero	convertible	88.6	130	111	5000	27	16500	13498.476233
	list	volvo	wagon	104.3	141	114	5400	28	16515	16097.608038
	grid	isuzu	hatchback	96	119	90	5000	29	11048	8315.257218
	list	dodge	hatchback	93.7	90	68	5500	41	5572	6630.154608
	grid	bmw	sedan	101.2	109	101	5800	29	16420	10012.408695

Output of the "Score Model" module

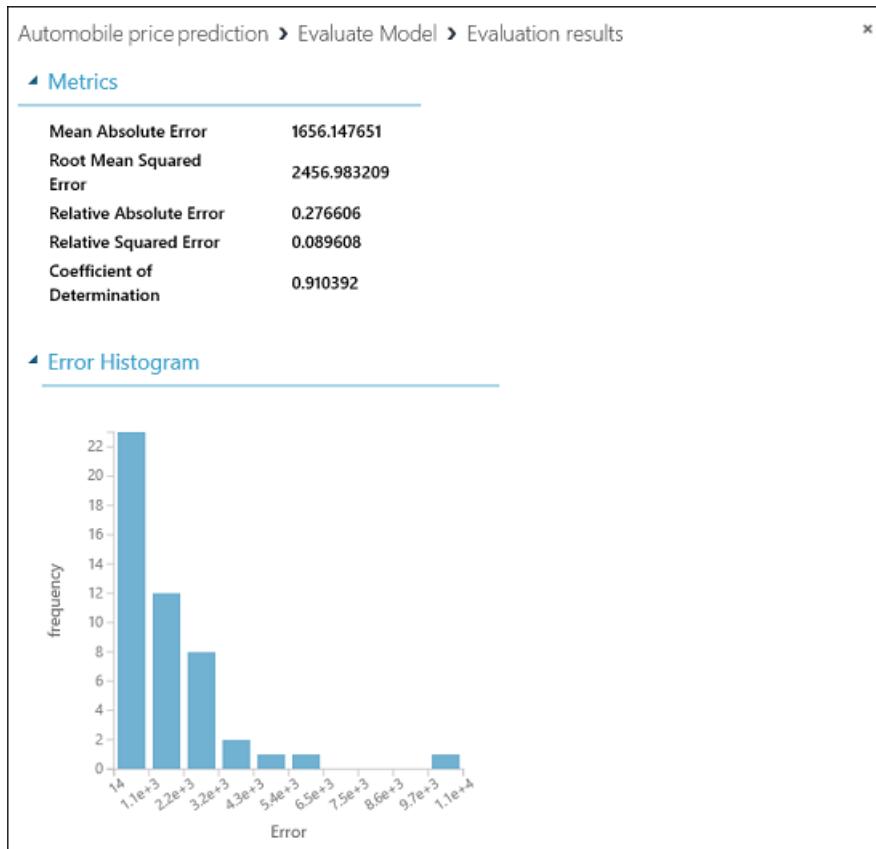
3. Finally, we test the quality of the results. Select and drag the [Evaluate Model](#) module to the experiment canvas, and connect the output of the [Score Model](#) module to the left input of [Evaluate Model](#).

TIP

There are two input ports on the [Evaluate Model](#) module because it can be used to compare two models side by side. Later, you can add another algorithm to the experiment and use [Evaluate Model](#) to see which one gives better results.

4. Run the experiment.

To view the output from the **Evaluate Model** module, click the output port, and then select **Visualize**.



Evaluation results for the experiment

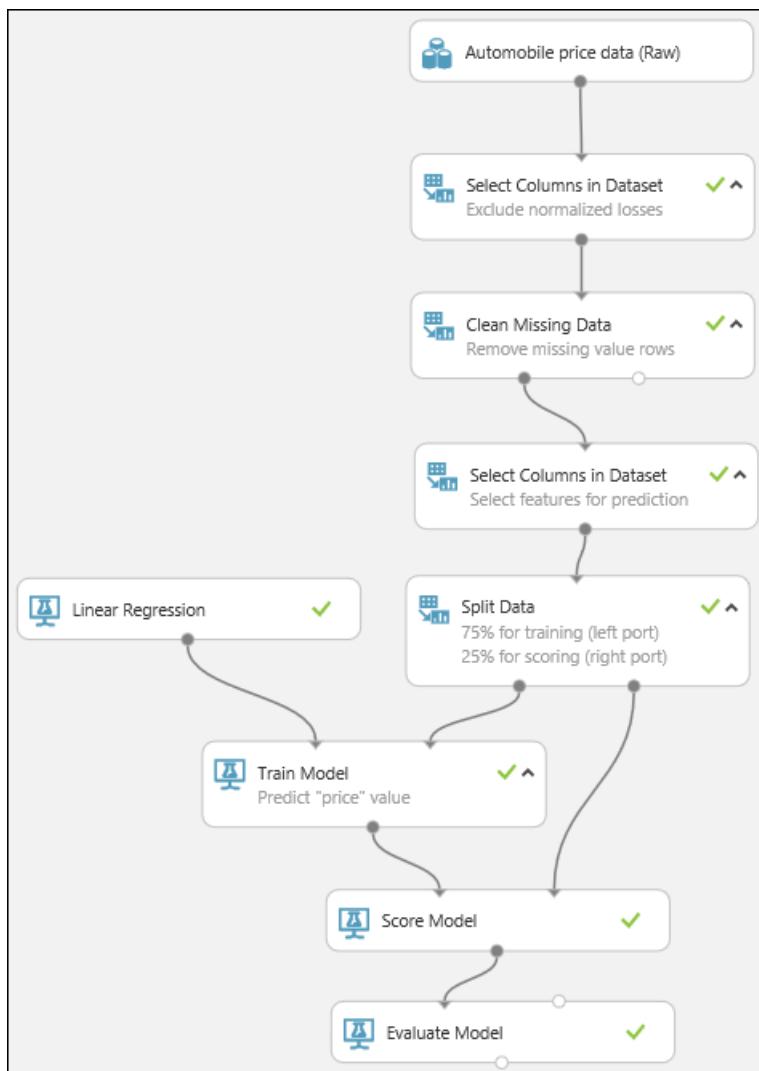
The following statistics are shown for our model:

- **Mean Absolute Error** (MAE): The average of absolute errors (an *error* is the difference between the predicted value and the actual value).
- **Root Mean Squared Error** (RMSE): The square root of the average of squared errors of predictions made on the test dataset.
- **Relative Absolute Error**: The average of absolute errors relative to the absolute difference between actual values and the average of all actual values.
- **Relative Squared Error**: The average of squared errors relative to the squared difference between the actual values and the average of all actual values.
- **Coefficient of Determination**: Also known as the **R squared value**, this is a statistical metric indicating how well a model fits the data.

For each of the error statistics, smaller is better. A smaller value indicates that the predictions more closely match the actual values. For **Coefficient of Determination**, the closer its value is to one (1.0), the better the predictions.

Final experiment

The final experiment should look something like this:



The final experiment

Next steps

Now that you've completed the first machine learning tutorial and have your experiment set up, you can continue to improve the model and then deploy it as a predictive web service.

- **Iterate to try to improve the model** - For example, you can change the features you use in your prediction. Or you can modify the properties of the [Linear Regression](#) algorithm or try a different algorithm altogether. You can even add multiple machine learning algorithms to your experiment at one time and compare two of them by using the [Evaluate Model](#) module. For an example of how to compare multiple models in a single experiment, see [Compare Regressors](#) in the Cortana Intelligence Gallery.

TIP

To copy any iteration of your experiment, use the **SAVE AS** button at the bottom of the page. You can see all the iterations of your experiment by clicking **VIEW RUN HISTORY** at the bottom of the page. For more details, see [Manage experiment iterations in Azure Machine Learning Studio](#).

- **Deploy the model as a predictive web service** - When you're satisfied with your model, you can deploy it as a web service to be used to predict automobile prices by using new data. For more details, see [Deploy an Azure Machine Learning web service](#).

Want to learn more? For a more extensive and detailed walkthrough of the process of creating, training, scoring, and deploying a model, see [Develop a predictive solution by using Azure Machine Learning](#).

Walkthrough: Develop a predictive analytics solution for credit risk assessment in Azure Machine Learning

9/25/2017 • 2 min to read • [Edit Online](#)

In this walkthrough, we take an extended look at the process of developing a predictive analytics solution in Machine Learning Studio. We develop a simple model in Machine Learning Studio, and then deploy it as an Azure Machine Learning web service where the model can make predictions using new data.

This walkthrough assumes that you've used Machine Learning Studio at least once before, and that you have some understanding of machine learning concepts. But it doesn't assume you're an expert in either.

If you've never used **Azure Machine Learning Studio** before, you might want to start with the tutorial, [Create your first data science experiment in Azure Machine Learning Studio](#). That tutorial takes you through Machine Learning Studio for the first time. It shows you the basics of how to drag-and-drop modules onto your experiment, connect them together, run the experiment, and look at the results. Another tool that may be helpful for getting started is a diagram that gives an overview of the capabilities of Machine Learning Studio. You can download and print it from here: [Overview diagram of Azure Machine Learning Studio capabilities](#).

If you're new to the field of machine learning in general, there's a video series that might be helpful to you. It's called [Data Science for Beginners](#) and it can give you a great introduction to machine learning using everyday language and concepts.

NOTE

Try Azure Machine Learning for free

No credit card or Azure subscription required. [Get started now.](#)

The problem

Suppose you need to predict an individual's credit risk based on the information they gave on a credit application.

Credit risk assessment is a complex problem, but we can simplify it a bit for this walkthrough. We'll use it as an example of how you can create a predictive analytics solution using Microsoft Azure Machine Learning. To do this, we use Azure Machine Learning Studio and a Machine Learning web service.

The solution

In this detailed walkthrough, we start with publicly available credit risk data and develop and train a predictive model based on that data. Then we deploy the model as a web service so it can be used by others for credit risk assessment.

To create this credit risk assessment solution, we follow these steps:

1. [Create a Machine Learning workspace](#)
2. [Upload existing data](#)
3. [Create an experiment](#)
4. [Train and evaluate the models](#)
5. [Deploy the web service](#)

6. Access the web service

TIP

You can find a working copy of the experiment that we develop in this walkthrough in the [Cortana Intelligence Gallery](#). Go to [Walkthrough - Credit risk prediction](#) and click **Open in Studio** to download a copy of the experiment into your Machine Learning Studio workspace.

This walkthrough is based on a simplified version of the sample experiment, [Binary Classification: Credit risk prediction](#), also available in the [Gallery](#).

Walkthrough Step 1: Create a Machine Learning workspace

11/30/2017 • 1 min to read • [Edit Online](#)

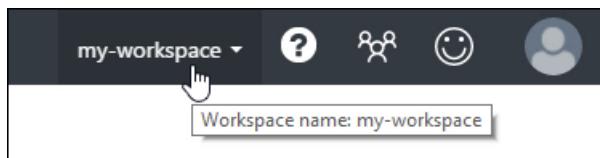
This is the first step of the walkthrough, [Develop a predictive analytics solution in Azure Machine Learning](#).

1. **Create a Machine Learning workspace**
2. [Upload existing data](#)
3. [Create a new experiment](#)
4. [Train and evaluate the models](#)
5. [Deploy the Web service](#)
6. [Access the Web service](#)

To use Machine Learning Studio, you need to have a Microsoft Azure Machine Learning workspace. This workspace contains the tools you need to create, manage, and publish experiments.

The administrator for your Azure subscription needs to create the workspace and then add you as an owner or contributor. For details, see [Create and share an Azure Machine Learning workspace](#).

After your workspace is created, open Machine Learning Studio (<https://studio.azureml.net/Home>). If you have more than one workspace, you can select the workspace in the toolbar in the upper-right corner of the window.



TIP

If you were made an owner of the workspace, you can share the experiments you're working on by inviting others to the workspace. You can do this in Machine Learning Studio on the **SETTINGS** page. You just need the Microsoft account or organizational account for each user.

On the **SETTINGS** page, click **USERS**, then click **INVITE MORE USERS** at the bottom of the window.

Next: Upload existing data

Walkthrough Step 2: Upload existing data into an Azure Machine Learning experiment

9/25/2017 • 3 min to read • [Edit Online](#)

This is the second step of the walkthrough, [Develop a predictive analytics solution in Azure Machine Learning](#)

1. [Create a Machine Learning workspace](#)
2. **Upload existing data**
3. [Create a new experiment](#)
4. [Train and evaluate the models](#)
5. [Deploy the Web service](#)
6. [Access the Web service](#)

To develop a predictive model for credit risk, we need data that we can use to train and then test the model. For this walkthrough, we'll use the "UCI Statlog (German Credit Data) Data Set" from the UC Irvine Machine Learning repository. You can find it here:

[http://archive.ics.uci.edu/ml/datasets/Statlog+\(German+Credit+Data\)](http://archive.ics.uci.edu/ml/datasets/Statlog+(German+Credit+Data))

We'll use the file named **german.data**. Download this file to your local hard drive.

The **german.data** dataset contains rows of 20 variables for 1000 past applicants for credit. These 20 variables represent the dataset's set of features (the *feature vector*), which provides identifying characteristics for each credit applicant. An additional column in each row represents the applicant's calculated credit risk, with 700 applicants identified as a low credit risk and 300 as a high risk.

The UCI website provides a description of the attributes of the feature vector for this data. This includes financial information, credit history, employment status, and personal information. For each applicant, a binary rating has been given indicating whether they are a low or high credit risk.

We'll use this data to train a predictive analytics model. When we're done, our model should be able to accept a feature vector for a new individual and predict whether he or she is a low or high credit risk.

Here's an interesting twist. The description of the dataset on the UCI website mentions what it costs if we misclassify a person's credit risk. If the model predicts a high credit risk for someone who is actually a low credit risk, the model has made a misclassification. But the reverse misclassification is five times more costly to the financial institution: if the model predicts a low credit risk for someone who is actually a high credit risk.

So, we want to train our model so that the cost of this latter type of misclassification is five times higher than misclassifying the other way. One simple way to do this when training the model in our experiment is by duplicating (five times) those entries that represent someone with a high credit risk. Then, if the model misclassifies someone as a low credit risk when they're actually a high risk, the model does that same misclassification five times, once for each duplicate. This will increase the cost of this error in the training results.

Convert the dataset format

The original dataset uses a blank-separated format. Machine Learning Studio works better with a comma-separated value (CSV) file, so we'll convert the dataset by replacing spaces with commas.

There are many ways to convert this data. One way is by using the following Windows PowerShell command:

```
cat german.data | %{$_ -replace " ",","} | sc german.csv
```

Another way is by using the Unix sed command:

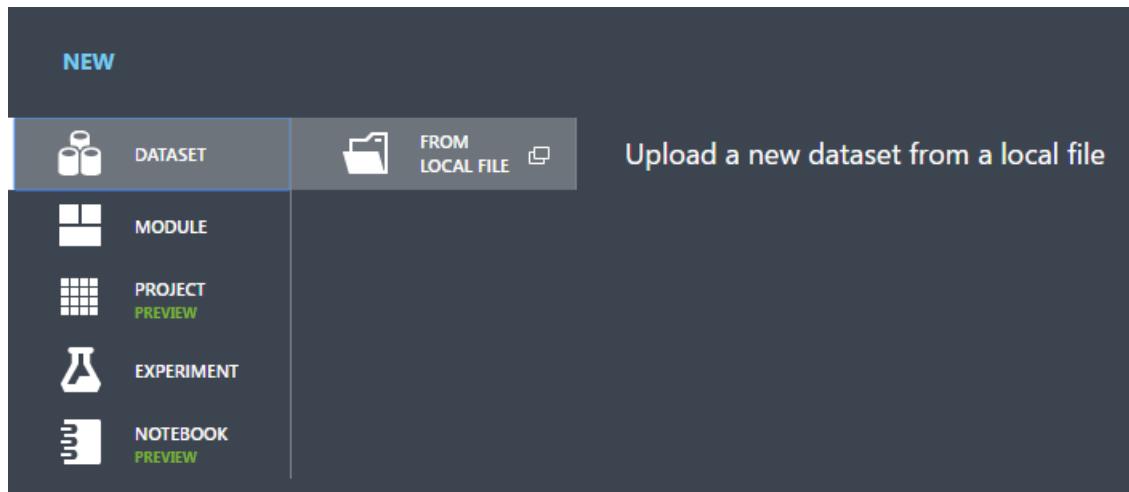
```
sed 's//,/g' german.data > german.csv
```

In either case, we have created a comma-separated version of the data in a file named **german.csv** that we can use in our experiment.

Upload the dataset to Machine Learning Studio

Once the data has been converted to CSV format, we need to upload it into Machine Learning Studio.

1. Open the Machine Learning Studio home page (<https://studio.azureml.net>).
2. Click the menu  in the upper-left corner of the window, click **Azure Machine Learning**, select **Studio**, and sign in.
3. Click **+NEW** at the bottom of the window.
4. Select **DATASET**.
5. Select **FROM LOCAL FILE**.



6. In the **Upload a new dataset** dialog, click **Browse** and find the **german.csv** file you created.
7. Enter a name for the dataset. For this walkthrough, call it "UCI German Credit Card Data".
8. For data type, select **Generic CSV File With no header (.nh.csv)**.
9. Add a description if you'd like.
10. Click the **OK** check mark.

Upload a new dataset

SELECT THE DATA TO UPLOAD:

german.csv

This is the new version of an existing dataset

ENTER A NAME FOR THE NEW DATASET:

UCI German Credit Card Data

SELECT A TYPE FOR THE NEW DATASET:

Generic CSV File With no header (.nh.csv)

PROVIDE AN OPTIONAL DESCRIPTION:

Downloaded from
<http://archive.ics.uci.edu/ml/datasets>

This uploads the data into a dataset module that we can use in an experiment.

You can manage datasets that you've uploaded to Studio by clicking the **DATASETS** tab to the left of the Studio window.

NAME	SUBMITTED BY	DESCRIPTION	DATA TYPE	CREATED	SIZE	PROJECT
UCI German Cred...	john	Downloaded fr...	GenericCSVNoHeader	12/15/2016 ...	78.9 KB	None

For more information about importing other types of data into an experiment, see [Import your training data into Azure Machine Learning Studio](#).

Next: Create a new experiment

Walkthrough Step 3: Create a new Azure Machine Learning experiment

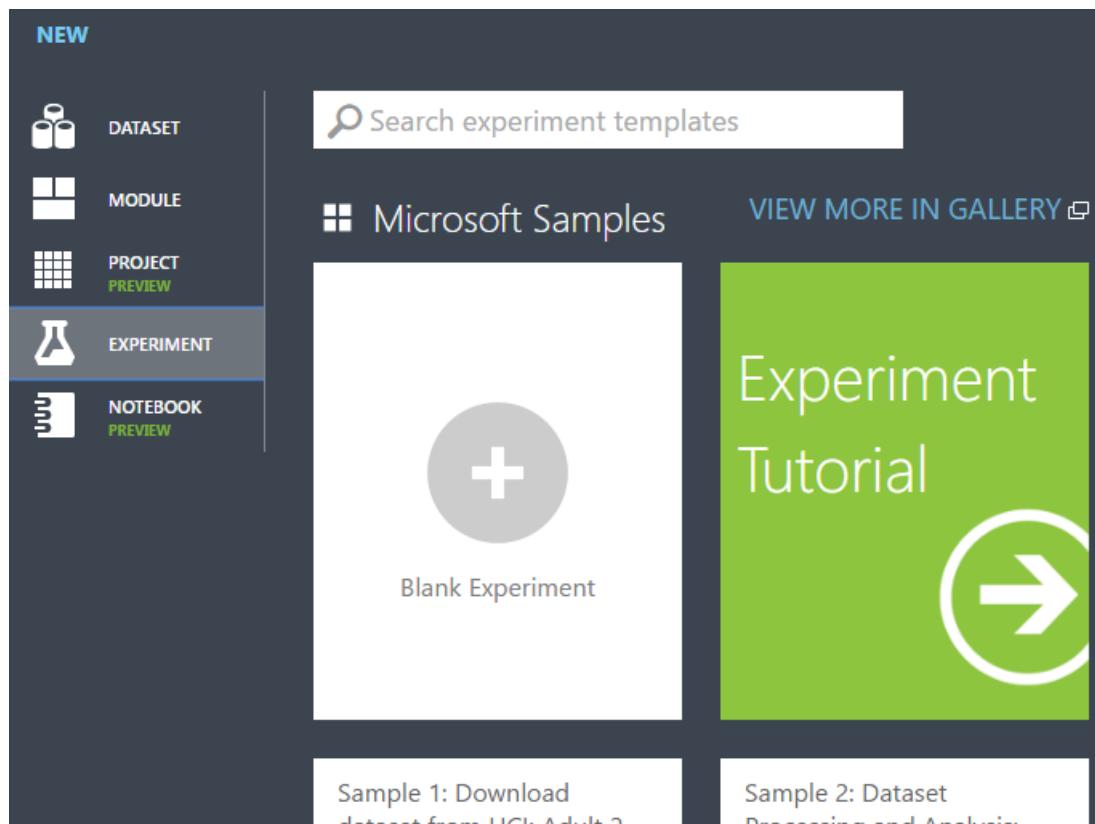
9/25/2017 • 6 min to read • [Edit Online](#)

This is the third step of the walkthrough, [Develop a predictive analytics solution in Azure Machine Learning](#)

1. [Create a Machine Learning workspace](#)
2. [Upload existing data](#)
3. **Create a new experiment**
4. [Train and evaluate the models](#)
5. [Deploy the Web service](#)
6. [Access the Web service](#)

The next step in this walkthrough is to create an experiment in Machine Learning Studio that uses the dataset we uploaded.

1. In Studio, click **+NEW** at the bottom of the window.
2. Select **EXPERIMENT**, and then select "Blank Experiment".



3. Select the default experiment name at the top of the canvas and rename it to something meaningful.



TIP

It's a good practice to fill in **Summary** and **Description** for the experiment in the **Properties** pane. These properties give you the chance to document the experiment so that anyone who looks at it later will understand your goals and methodology.

Properties Project

Experiment Properties

STATUS CODE InDraft

Summary

Credit risk prediction from customer credit application data

Description

This experiment predicts credit risk based on information provided on a credit application. The training data was derived from the "UCI Statlog (German Credit Data) Data Set" from the UCI Machine Learning repository: <[>](http://archive.ics.uci.edu/ml/datasets/Statlog+(German+Credit+Data)).

4. In the module palette to the left of the experiment canvas, expand **Saved Datasets**.
5. Find the dataset you created under **My Datasets** and drag it onto the canvas. You can also find the dataset by entering the name in the **Search** box above the palette.

credit risk experiment

Search experiment items

Saved Datasets

My Datasets

UCI German Credit Card Data

1

Samples

Data Format Conversions

Prepare the data

You can view the first 100 rows of the data and some statistical information for the whole dataset: Click the

output port of the dataset (the small circle at the bottom) and select **Visualize**.

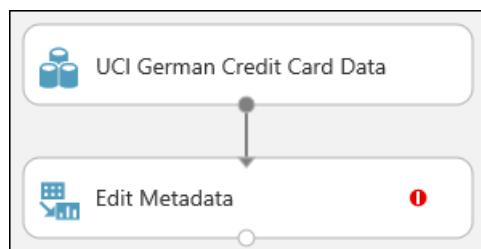
Because the data file didn't come with column headings, Studio has provided generic headings (Col1, Col2, etc.). Good headings aren't essential to creating a model, but they make it easier to work with the data in the experiment. Also, when we eventually publish this model in a web service, the headings help identify the columns to the user of the service.

We can add column headings using the [Edit Metadata](#) module. You use the [Edit Metadata](#) module to change metadata associated with a dataset. In this case, we use it to provide more friendly names for column headings.

To use [Edit Metadata](#), you first specify which columns to modify (in this case, all of them.) Next, you specify the action to be performed on those columns (in this case, changing column headings.)

1. In the module palette, type "metadata" in the **Search** box. The [Edit Metadata](#) appears in the module list.
2. Click and drag the [Edit Metadata](#) module onto the canvas and drop it below the dataset we added earlier.
3. Connect the dataset to the [Edit Metadata](#): click the output port of the dataset (the small circle at the bottom of the dataset), drag to the input port of [Edit Metadata](#) (the small circle at the top of the module), then release the mouse button. The dataset and module remain connected even if you move either around on the canvas.

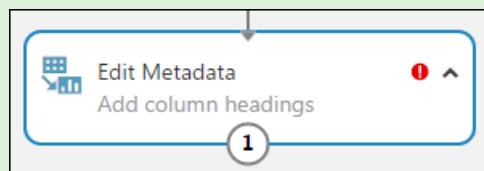
The experiment should now look something like this:



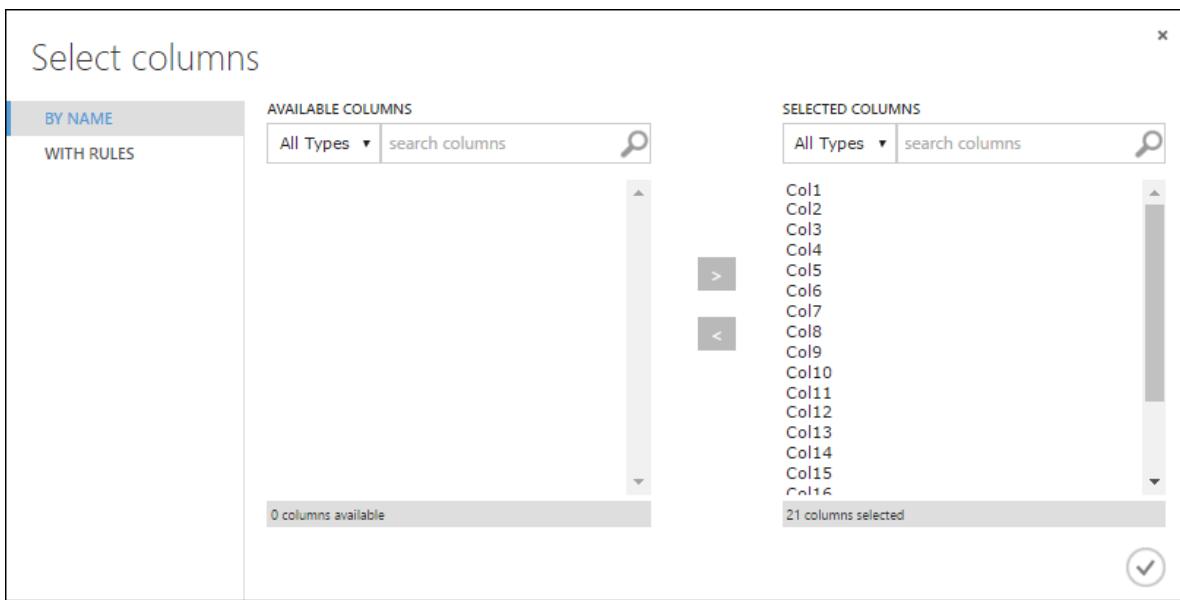
The red exclamation mark indicates that we haven't set the properties for this module yet. We'll do that next.

TIP

You can add a comment to a module by double-clicking the module and entering text. This can help you see at a glance what the module is doing in your experiment. In this case, double-click the [Edit Metadata](#) module and type the comment "Add column headings". Click anywhere else on the canvas to close the text box. To display the comment, click the down-arrow on the module.



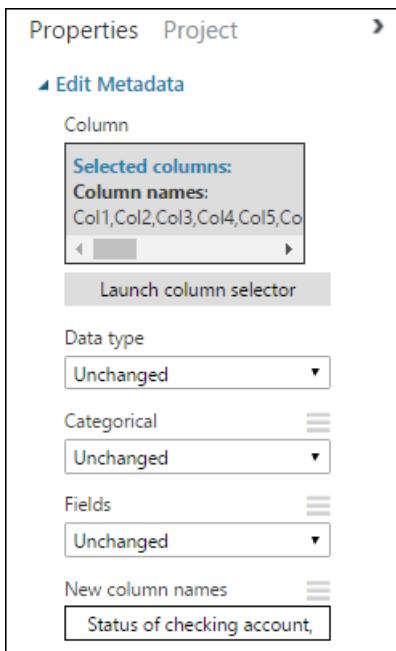
4. Select [Edit Metadata](#), and in the **Properties** pane to the right of the canvas, click **Launch column selector**.
5. In the **Select columns** dialog, select all the rows in **Available Columns** and click **>** to move them to **Selected Columns**. The dialog should look like this:



6. Click the **OK** check mark.
7. Back in the **Properties** pane, look for the **New column names** parameter. In this field, enter a list of names for the 21 columns in the dataset, separated by commas and in column order. You can obtain the columns names from the dataset documentation on the UCI website, or for convenience you can copy and paste the following list:

Status of checking account, Duration in months, Credit history, Purpose, Credit amount, Savings account/bond, Present employment since, Installment rate in percentage of disposable income, Personal status and sex, Other debtors, Present residence since, Property, Age in years, Other installment plans, Housing, Number of existing credits, Job, Number of people providing maintenance for, Telephone, Foreign worker, Credit risk

The Properties pane looks like this:



TIP

If you want to verify the column headings, run the experiment (click **RUN** below the experiment canvas). When it finishes running (a green check mark appears on [Edit Metadata](#)), click the output port of the [Edit Metadata](#) module, and select **Visualize**. You can view the output of any module in the same way to view the progress of the data through the experiment.

Create training and test datasets

We need some data to train the model and some to test it. So in the next step of the experiment, we split the dataset into two separate datasets: one for training our model and one for testing it.

To do this, we use the [Split Data](#) module.

1. Find the [Split Data](#) module, drag it onto the canvas, and connect it to the [Edit Metadata](#) module.
2. By default, the split ratio is 0.5 and the **Randomized split** parameter is set. This means that a random half of the data is output through one port of the [Split Data](#) module, and half through the other. You can adjust these parameters, as well as the **Random seed** parameter, to change the split between training and testing data. For this example, we leave them as-is.

TIP

The property **Fraction of rows in the first output dataset** determines how much of the data is output through the *left* output port. For instance, if you set the ratio to 0.7, then 70% of the data is output through the left port and 30% through the right port.

3. Double-click the [Split Data](#) module and enter the comment, "Training/testing data split 50%".

We can use the outputs of the [Split Data](#) module however we like, but let's choose to use the left output as training data and the right output as testing data.

As mentioned in the [previous step](#), the cost of misclassifying a high credit risk as low is five times higher than the cost of misclassifying a low credit risk as high. To account for this, we generate a new dataset that reflects this cost function. In the new dataset, each high risk example is replicated five times, while each low risk example is not replicated.

We can do this replication using R code:

1. Find and drag the [Execute R Script](#) module onto the experiment canvas.
2. Connect the left output port of the [Split Data](#) module to the first input port ("Dataset1") of the [Execute R Script](#) module.
3. Double-click the [Execute R Script](#) module and enter the comment, "Set cost adjustment".
4. In the **Properties** pane, delete the default text in the **R Script** parameter and enter this script:

```
dataset1 <- maml.mapInputPort(1)
data.set<-dataset1[dataset1[,21]==1,]
pos<-dataset1[dataset1[,21]==2,]
for (i in 1:5) data.set<-rbind(data.set,pos)
maml.mapOutputPort("data.set")
```

Properties Project

Execute R Script

R Script

```
1 dataset1 <- maml.mapInputPort(1)
2 data.set<-dataset1[dataset1[,21]==1,]
3 pos<-dataset1[dataset1[,21]==2,]
4 for (i in 1:5) data.set<-rbind(data.set,pos)
5 maml.mapOutputPort("data.set")
```

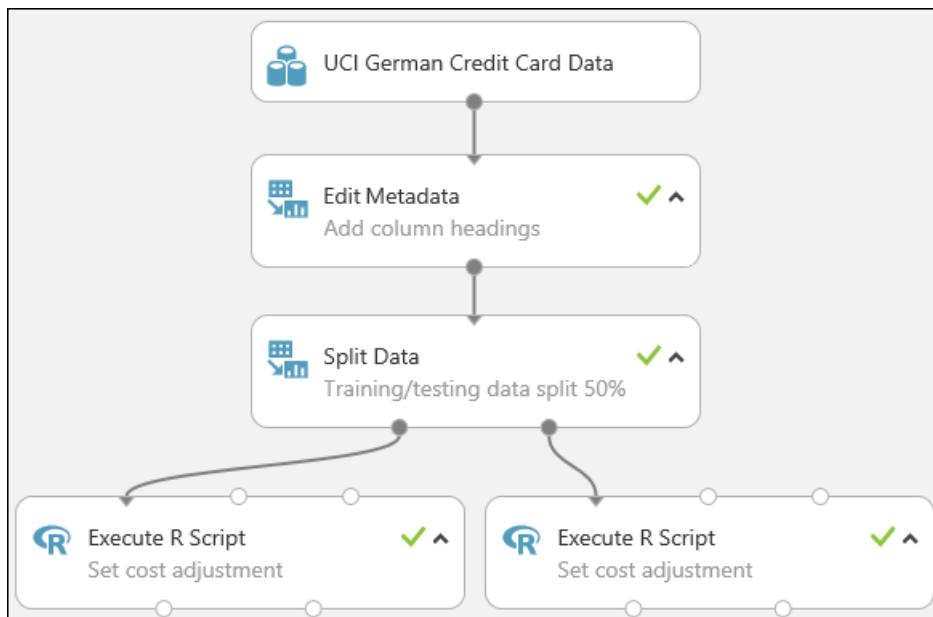
We need to do this same replication operation for each output of the [Split Data](#) module so that the training and testing data have the same cost adjustment. The easiest way to do this is by duplicating the [Execute R Script](#) module we just made and connecting it to the other output port of the [Split Data](#) module.

1. Right-click the [Execute R Script](#) module and select **Copy**.
2. Right-click the experiment canvas and select **Paste**.
3. Drag the new module into position, and then connect the right output port of the [Split Data](#) module to the first input port of this new [Execute R Script](#) module.
4. At the bottom of the canvas, click **Run**.

TIP

The copy of the Execute R Script module contains the same script as the original module. When you copy and paste a module on the canvas, the copy retains all the properties of the original.

Our experiment now looks something like this:



For more information on using R scripts in your experiments, see [Extend your experiment with R](#).

Next: Train and evaluate the models

Walkthrough Step 4: Train and evaluate the predictive analytic models

9/25/2017 • 8 min to read • [Edit Online](#)

This topic contains the fourth step of the walkthrough, [Develop a predictive analytics solution in Azure Machine Learning](#)

1. [Create a Machine Learning workspace](#)
2. [Upload existing data](#)
3. [Create a new experiment](#)
4. **Train and evaluate the models**
5. [Deploy the Web service](#)
6. [Access the Web service](#)

One of the benefits of using Azure Machine Learning Studio for creating machine learning models is the ability to try more than one type of model at a time in a single experiment and compare the results. This type of experimentation helps you find the best solution for your problem.

In the experiment we're developing in this walkthrough, we'll create two different types of models and then compare their scoring results to decide which algorithm we want to use in our final experiment.

There are various models we could choose from. To see the models available, expand the **Machine Learning** node in the module palette, and then expand **Initialize Model** and the nodes beneath it. For the purposes of this experiment, we'll select the [Two-Class Support Vector Machine](#) (SVM) and the [Two-Class Boosted Decision Tree](#) modules.

TIP

To get help deciding which Machine Learning algorithm best suits the particular problem you're trying to solve, see [How to choose algorithms for Microsoft Azure Machine Learning](#).

Train the models

We'll add both the [Two-Class Boosted Decision Tree](#) module and [Two-Class Support Vector Machine](#) module in this experiment.

Two-Class Boosted Decision Tree

First, let's set up the boosted decision tree model.

1. Find the [Two-Class Boosted Decision Tree](#) module in the module palette and drag it onto the canvas.
2. Find the [Train Model](#) module, drag it onto the canvas, and then connect the output of the [Two-Class Boosted Decision Tree](#) module to the left input port of the [Train Model](#) module.

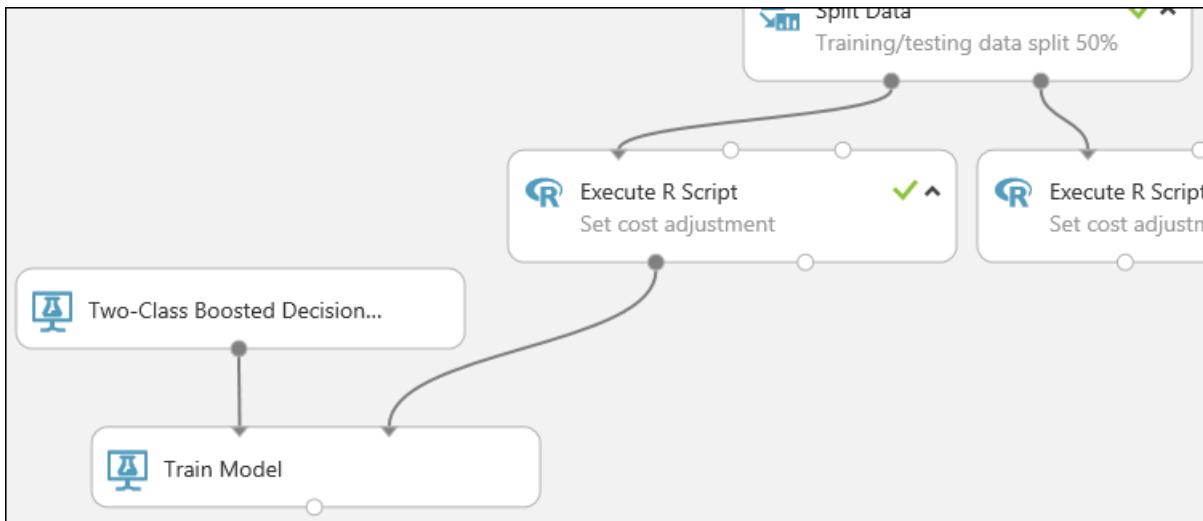
The [Two-Class Boosted Decision Tree](#) module initializes the generic model, and [Train Model](#) uses training data to train the model.

3. Connect the left output of the left [Execute R Script](#) module to the right input port of the [Train Model](#) module (we decided in [Step 3](#) of this walkthrough to use the data coming from the left side of the Split Data module for training).

TIP

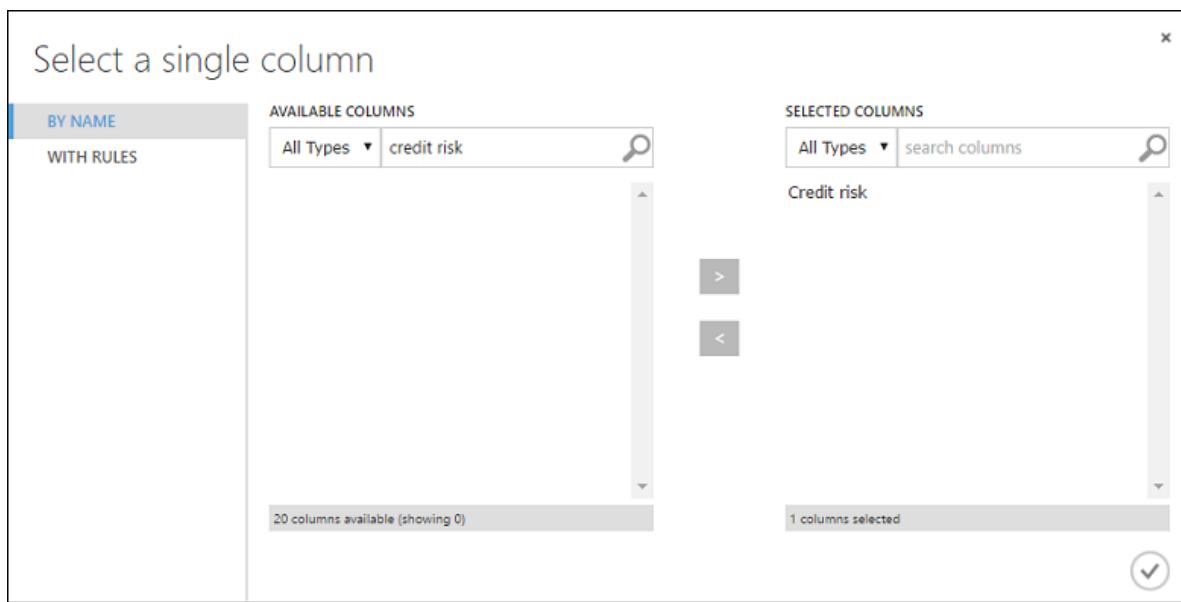
We don't need two of the inputs and one of the outputs of the [Execute R Script](#) module for this experiment, so we can leave them unattached.

This portion of the experiment now looks something like this:



Now we need to tell the [Train Model](#) module that we want the model to predict the Credit Risk value.

1. Select the [Train Model](#) module. In the **Properties** pane, click **Launch column selector**.
2. In the **Select a single column** dialog, type "credit risk" in the search field under **Available Columns**, select "Credit risk" below, and click the right arrow button (>) to move "Credit risk" to **Selected Columns**.



3. Click the **OK** check mark.

Two-Class Support Vector Machine

Next, we set up the SVM model.

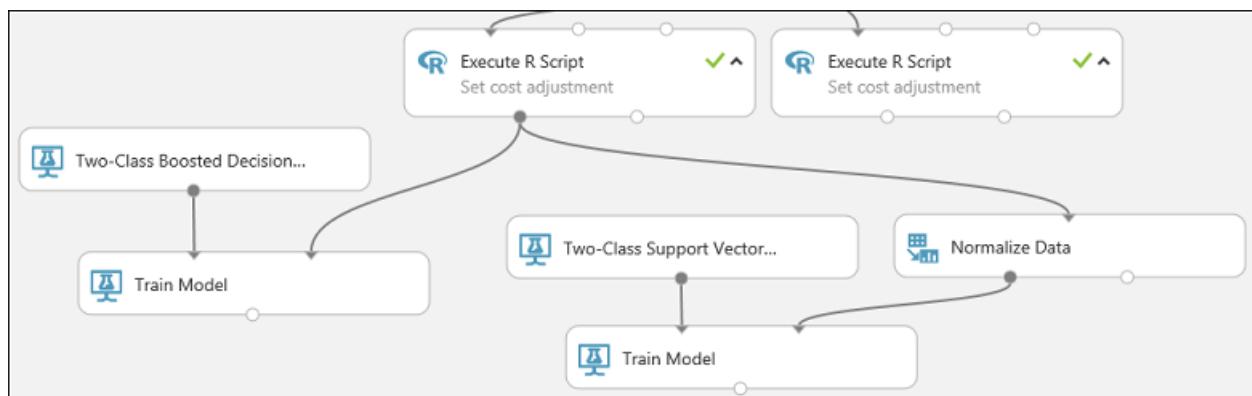
First, a little explanation about SVM. Boosted decision trees work well with features of any type. However, since the SVM module generates a linear classifier, the model that it generates has the best test error when all numeric features have the same scale. To convert all numeric features to the same scale, we use a "Tanh" transformation (with the [Normalize Data](#) module). This transforms our numbers into the [0,1] range. The SVM module converts string features to categorical features and then to binary 0/1 features, so we don't need to manually transform

string features. Also, we don't want to transform the Credit Risk column (column 21) - it's numeric, but it's the value we're training the model to predict, so we need to leave it alone.

To set up the SVM model, do the following:

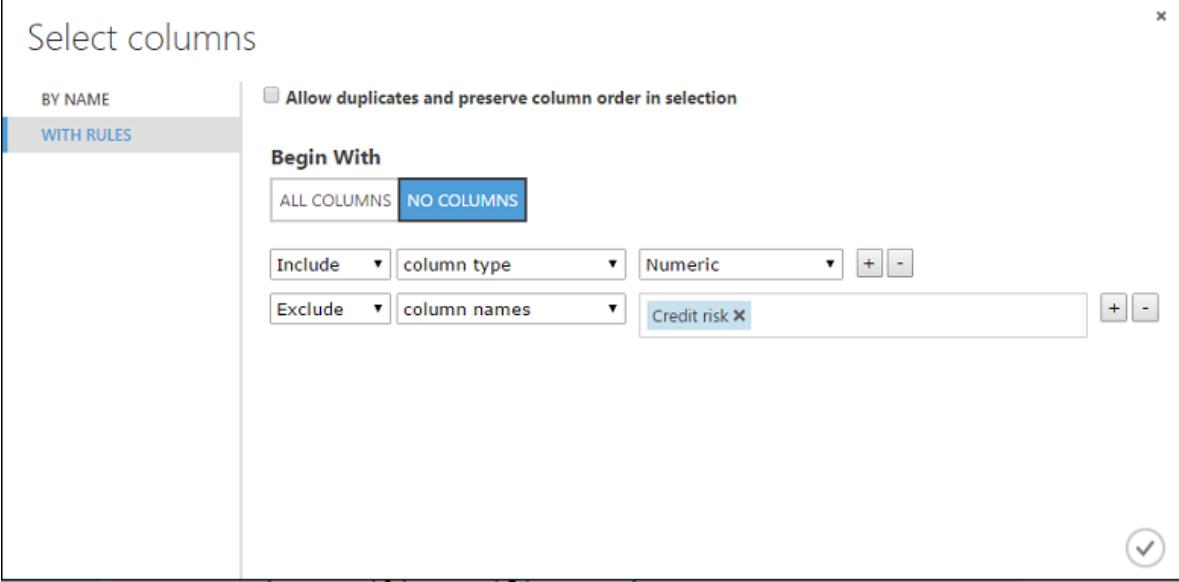
1. Find the [Two-Class Support Vector Machine](#) module in the module palette and drag it onto the canvas.
2. Right-click the [Train Model](#) module, select **Copy**, and then right-click the canvas and select **Paste**. The copy of the [Train Model](#) module has the same column selection as the original.
3. Connect the output of the [Two-Class Support Vector Machine](#) module to the left input port of the second [Train Model](#) module.
4. Find the [Normalize Data](#) module and drag it onto the canvas.
5. Connect the left output of the left [Execute R Script](#) module to the input of this module (notice that the output port of a module may be connected to more than one other module).
6. Connect the left output port of the [Normalize Data](#) module to the right input port of the second [Train Model](#) module.

This portion of our experiment should now look something like this:



Now configure the [Normalize Data](#) module:

1. Click to select the [Normalize Data](#) module. In the **Properties** pane, select **Tanh** for the **Transformation method** parameter.
2. Click **Launch column selector**, select "No columns" for **Begin With**, select **Include** in the first dropdown, select **column type** in the second dropdown, and select **Numeric** in the third dropdown. This specifies that all the numeric columns (and only numeric) are transformed.
3. Click the plus sign (+) to the right of this row - this creates a row of dropdowns. Select **Exclude** in the first dropdown, select **column names** in the second dropdown, and enter "Credit risk" in the text field. This specifies that the Credit Risk column should be ignored (we need to do this because this column is numeric and so would be transformed if we didn't exclude it).
4. Click the **OK** check mark.



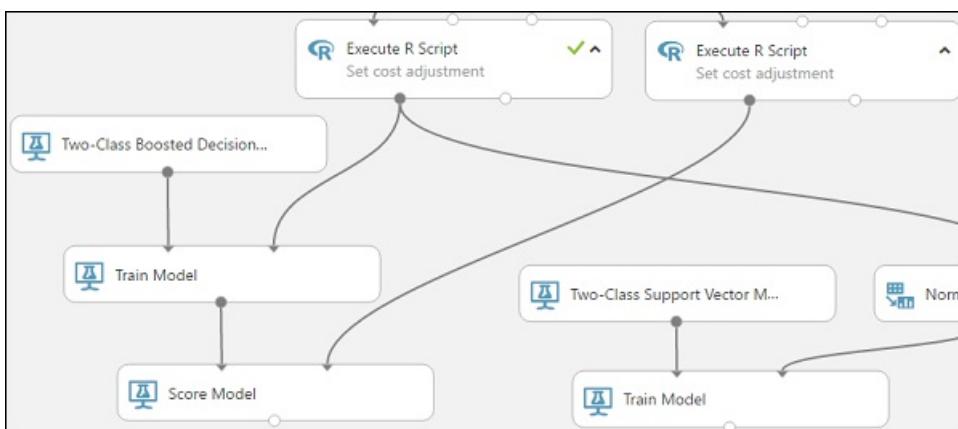
The [Normalize Data](#) module is now set to perform a Tanh transformation on all numeric columns except for the Credit Risk column.

Score and evaluate the models

We use the testing data that was separated out by the [Split Data](#) module to score our trained models. We can then compare the results of the two models to see which generated better results.

Add the Score Model modules

1. Find the [Score Model](#) module and drag it onto the canvas.
2. Connect the [Train Model](#) module that's connected to the [Two-Class Boosted Decision Tree](#) module to the left input port of the [Score Model](#) module.
3. Connect the right [Execute R Script](#) module (our testing data) to the right input port of the [Score Model](#) module.

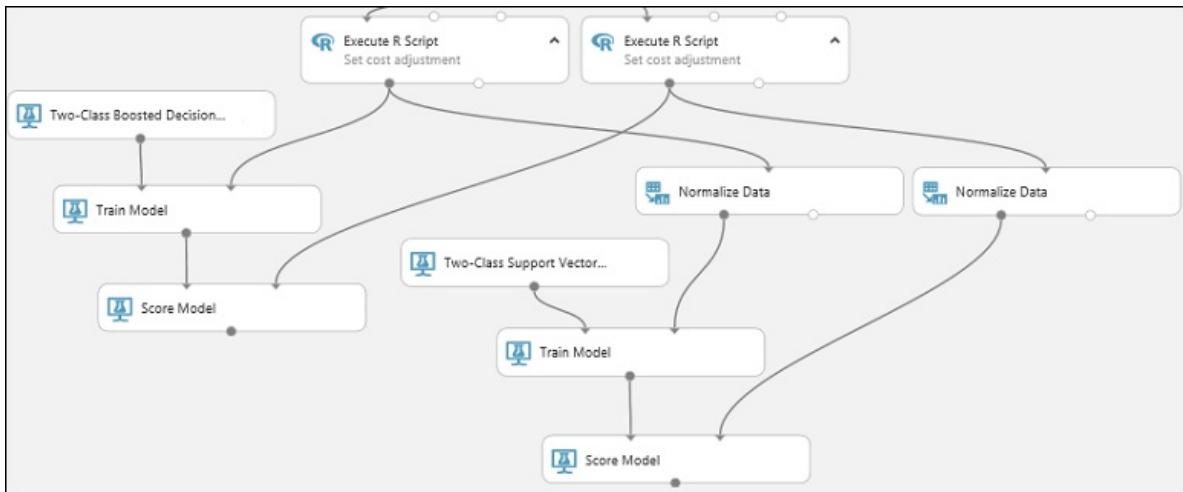


The [Score Model](#) module can now take the credit information from the testing data, run it through the model, and compare the predictions the model generates with the actual credit risk column in the testing data.

4. Copy and paste the [Score Model](#) module to create a second copy.
5. Connect the output of the SVM model (that is, the output port of the [Train Model](#) module that's connected to the [Two-Class Support Vector Machine](#) module) to the input port of the second [Score Model](#) module.
6. For the SVM model, we have to do the same transformation to the test data as we did to the training data. So copy and paste the [Normalize Data](#) module to create a second copy and connect it to the right [Execute R](#)

Script module.

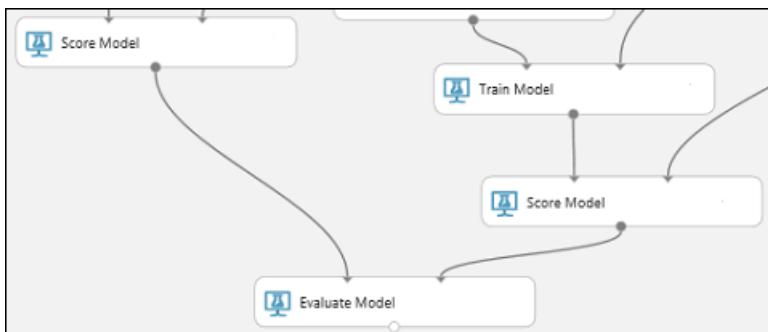
7. Connect the left output of the second Normalize Data module to the right input port of the second Score Model module.



Add the Evaluate Model module

To evaluate the two scoring results and compare them, we use an [Evaluate Model](#) module.

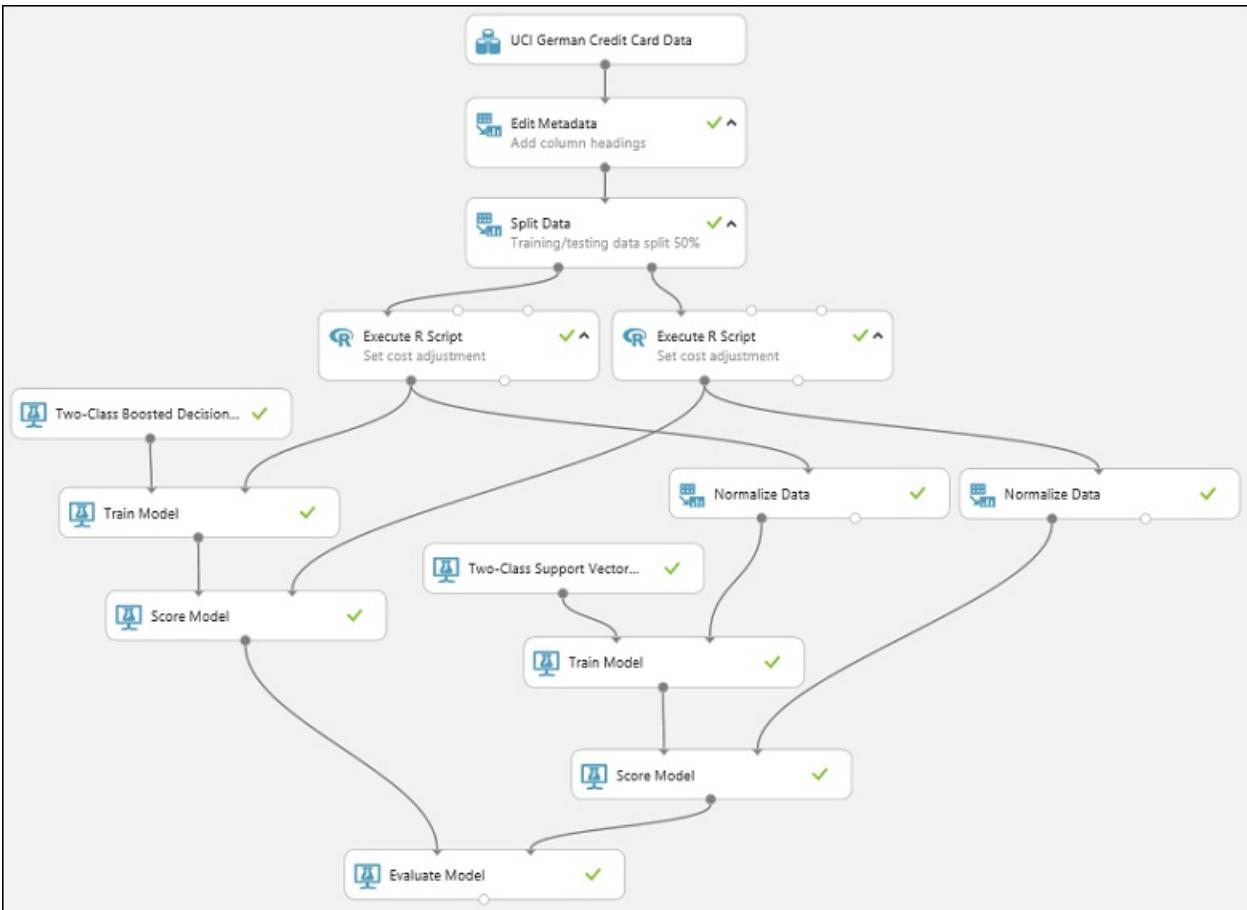
1. Find the [Evaluate Model](#) module and drag it onto the canvas.
2. Connect the output port of the [Score Model](#) module associated with the boosted decision tree model to the left input port of the [Evaluate Model](#) module.
3. Connect the other [Score Model](#) module to the right input port.



Run the experiment and check the results

To run the experiment, click the **RUN** button below the canvas. It may take a few minutes. A spinning indicator on each module shows that it's running, and then a green check mark shows when the module is finished. When all the modules have a check mark, the experiment has finished running.

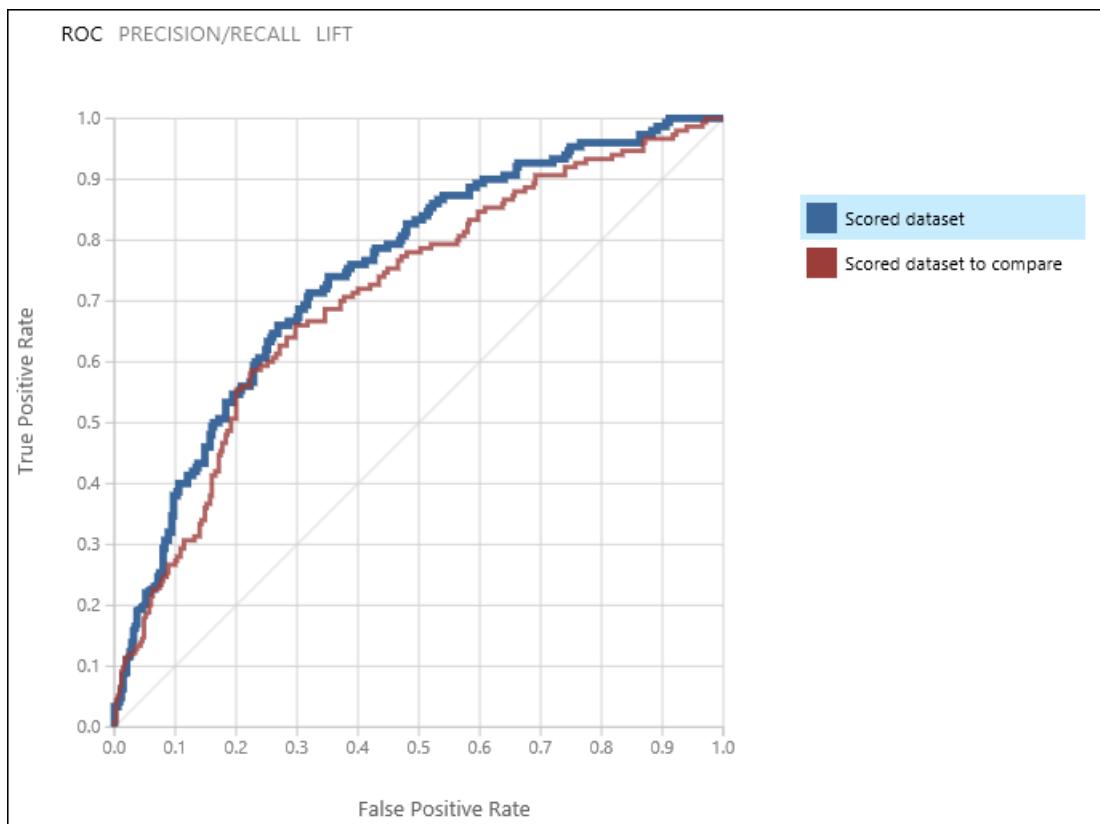
The experiment should now look something like this:



To check the results, click the output port of the [Evaluate Model](#) module and select **Visualize**.

The [Evaluate Model](#) module produces a pair of curves and metrics that allow you to compare the results of the two scored models. You can view the results as Receiver Operator Characteristic (ROC) curves, Precision/Recall curves, or Lift curves. Additional data displayed includes a confusion matrix, cumulative values for the area under the curve (AUC), and other metrics. You can change the threshold value by moving the slider left or right and see how it affects the set of metrics.

To the right of the graph, click **Scored dataset** or **Scored dataset to compare** to highlight the associated curve and to display the associated metrics below. In the legend for the curves, "Scored dataset" corresponds to the left input port of the [Evaluate Model](#) module - in our case, this is the boosted decision tree model. "Scored dataset to compare" corresponds to the right input port - the SVM model in our case. When you click one of these labels, the curve for that model is highlighted and the corresponding metrics are displayed, as shown in the following graphic.



True Positive	False Negative	Accuracy	Precision	Threshold	Cumulative AUC
441	309	0.646	0.846	0.5	0.748
False Positive	True Negative	Recall	F1 Score		
80	270	0.588	0.694		

By examining these values, you can decide which model is closest to giving you the results you're looking for. You can go back and iterate on your experiment by changing parameter values in the different models.

The science and art of interpreting these results and tuning the model performance is outside the scope of this walkthrough. For additional help, you might read the following articles:

- [How to evaluate model performance in Azure Machine Learning](#)
- [Choose parameters to optimize your algorithms in Azure Machine Learning](#)
- [Interpret model results in Azure Machine Learning](#)

TIP

Each time you run the experiment a record of that iteration is kept in the Run History. You can view these iterations, and return to any of them, by clicking **VIEW RUN HISTORY** below the canvas. You can also click **Prior Run** in the **Properties** pane to return to the iteration immediately preceding the one you have open.

You can make a copy of any iteration of your experiment by clicking **SAVE AS** below the canvas. Use the experiment's **Summary** and **Description** properties to keep a record of what you've tried in your experiment iterations.

For more details, see [Manage experiment iterations in Azure Machine Learning Studio](#).

Next: Deploy the web service

Walkthrough Step 5: Deploy the Azure Machine Learning web service

11/30/2017 • 9 min to read • [Edit Online](#)

This is the fifth step of the walkthrough, [Develop a predictive analytics solution in Azure Machine Learning](#)

1. [Create a Machine Learning workspace](#)
 2. [Upload existing data](#)
 3. [Create a new experiment](#)
 4. [Train and evaluate the models](#)
 5. **Deploy the web service**
 6. [Access the web service](#)
-

To give others a chance to use the predictive model we've developed in this walkthrough, we can deploy it as a web service on Azure.

Up to this point we've been experimenting with training our model. But the deployed service is no longer going to do training - it's going to generate new predictions by scoring the user's input based on our model. So we're going to do some preparation to convert this experiment from a ***training*** experiment to a ***predictive*** experiment.

This is a three-step process:

1. Remove one of the models
2. Convert the *training experiment* we've created into a *predictive experiment*
3. Deploy the predictive experiment as a web service

Remove one of the models

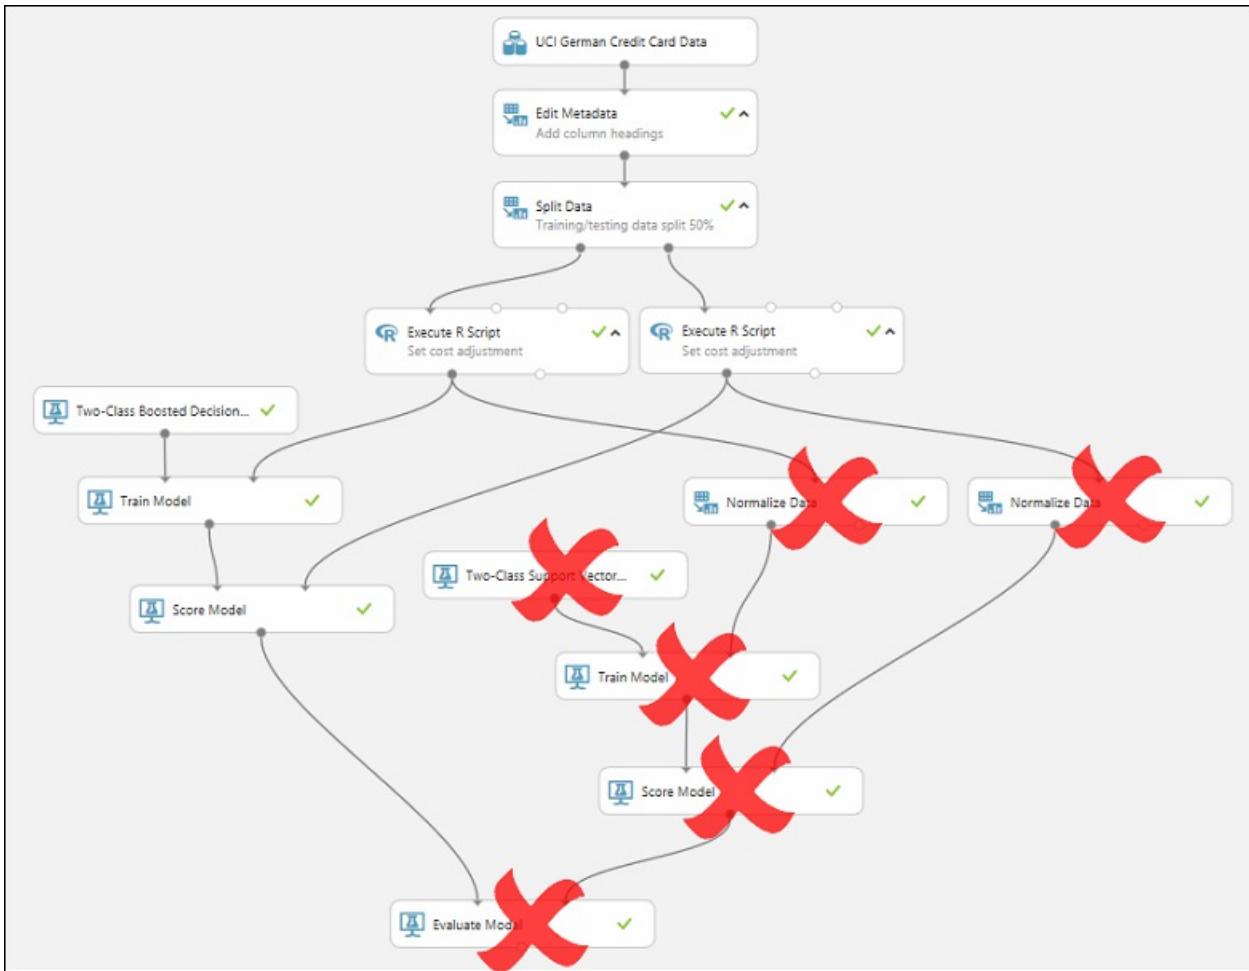
First, we need to trim this experiment down a little. We currently have two different models in the experiment, but we only want to use one model when we deploy this as a web service.

Let's say we've decided that the boosted tree model performed better than the SVM model. So the first thing to do is remove the [Two-Class Support Vector Machine](#) module and the modules that were used for training it. You may want to make a copy of the experiment first by clicking **Save As** at the bottom of the experiment canvas.

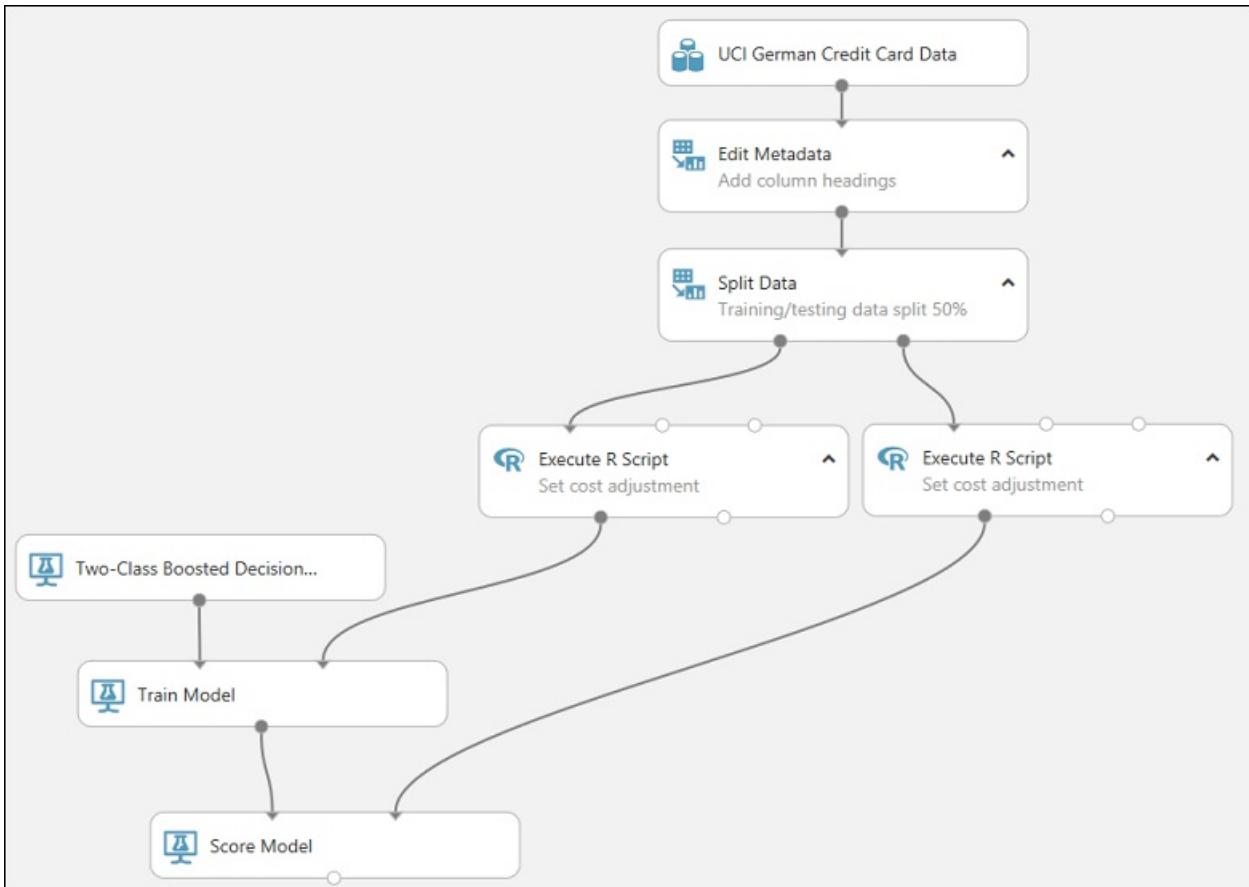
We need to delete the following modules:

- [Two-Class Support Vector Machine](#)
- Train Model and [Score Model](#) modules that were connected to it
- [Normalize Data](#) (both of them)
- [Evaluate Model](#) (because we're finished evaluating the models)

Select each module and press the Delete key, or right-click the module and select **Delete**.



Our model should now look something like this:



Now we're ready to deploy this model using the [Two-Class Boosted Decision Tree](#).

Convert the training experiment to a predictive experiment

To get this model ready for deployment, we need to convert this training experiment to a predictive experiment. This involves three steps:

1. Save the model we've trained and then replace our training modules
2. Trim the experiment to remove modules that were only needed for training
3. Define where the web service will accept input and where it generates the output

We could do this manually, but fortunately all three steps can be accomplished by clicking **Set Up Web Service** at the bottom of the experiment canvas (and selecting the **Predictive Web Service** option).

TIP

If you want more details on what happens when you convert a training experiment to a predictive experiment, see [How to prepare your model for deployment in Azure Machine Learning Studio](#).

When you click **Set Up Web Service**, several things happen:

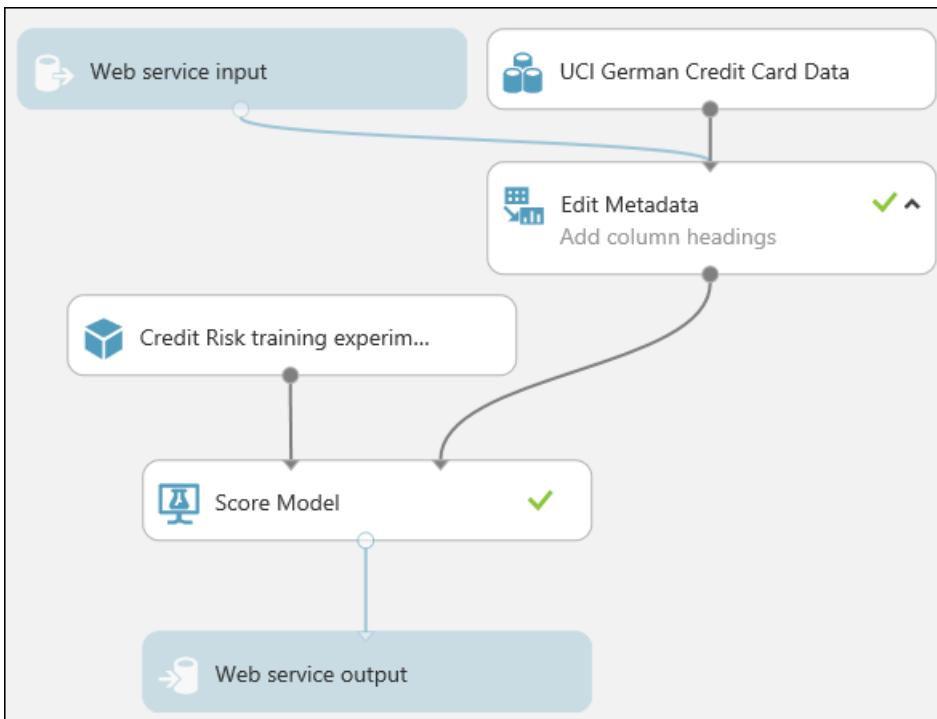
- The trained model is converted to a single **Trained Model** module and stored in the module palette to the left of the experiment canvas (you can find it under **Trained Models**)
- Modules that were used for training are removed; specifically:
 - [Two-Class Boosted Decision Tree](#)
 - [Train Model](#)
 - [Split Data](#)
 - the second [Execute R Script](#) module that was used for test data
- The saved trained model is added back into the experiment
- **Web service input** and **Web service output** modules are added (these identify where the user's data will enter the model, and what data is returned, when the web service is accessed)

NOTE

You can see that the experiment is saved in two parts under tabs that have been added at the top of the experiment canvas. The original training experiment is under the tab **Training experiment**, and the newly created predictive experiment is under **Predictive experiment**. The predictive experiment is the one we'll deploy as a web service.

We need to take one additional step with this particular experiment. We added two [Execute R Script](#) modules to provide a weighting function to the data. That was just a trick we needed for training and testing, so we can take out those modules in the final model. Machine Learning Studio removed one [Execute R Script](#) module when it removed the [Split](#) module. Now we can remove the other and connect [Metadata Editor](#) directly to [Score Model](#).

Our experiment should now look like this:



NOTE

You may be wondering why we left the UCI German Credit Card Data dataset in the predictive experiment. The service is going to score the user's data, not the original dataset, so why leave the original dataset in the model?

It's true that the service doesn't need the original credit card data. But it does need the schema for that data, which includes information such as how many columns there are and which columns are numeric. This schema information is necessary to interpret the user's data. We leave these components connected so that the scoring module has the dataset schema when the service is running. The data isn't used, just the schema.

Run the experiment one last time (click **Run**.) If you want to verify that the model is still working, click the output of the **Score Model** module and select **View Results**. You can see that the original data is displayed, along with the credit risk value ("Scored Labels") and the scoring probability value ("Scored Probabilities").

Deploy the web service

You can deploy the experiment as either a Classic web service, or as a New web service that's based on Azure Resource Manager.

Deploy as a Classic web service

To deploy a Classic web service derived from our experiment, click **Deploy Web Service** below the canvas and select **Deploy Web Service [Classic]**. Machine Learning Studio deploys the experiment as a web service and takes you to the dashboard for that web service. From this page you can return to the experiment (**View snapshot** or **View latest**) and run a simple test of the web service (see **Test the web service** below). There is also information here for creating applications that can access the web service (more on that in the next step of this walkthrough).

credit risk experiment

DASHBOARD CONFIGURATION

General

Published experiment

[View snapshot](#) [View latest](#)

Description

No description provided for this web service.

API key

XXXXXXXXXXXXXXXXXXXXXX [copy]

Default Endpoint

API HELP PAGE	TEST	APPS	LAST UPDATED	SEARCH
REQUEST/RESPONSE	Test	Excel 2013 or later Excel 2010 or earlier workbook	2/5/2016 5:43:22 PM	
BATCH EXECUTION		Excel 2013 or later workbook	2/5/2016 5:43:22 PM	

Additional endpoints

Number of additional endpoints created for this web service: 0

[Manage endpoints in Azure management portal](#)

You can configure the service by clicking the **CONFIGURATION** tab. Here you can modify the service name (it's given the experiment name by default) and give it a description. You can also give more friendly labels for the input and output data.

DASHBOARD **CONFIGURATION**

settings

GENERAL

Display Name	Credit Risk experiment
Description	No description provided for this web service.

ENABLE LOGGING

YES NO [Learn more](#)

ENABLE SAMPLE DATA?

YES NO

INPUT SCHEMA

Col1 (String)
Col2 (Numerical)

[Deploy as a New web service](#)

NOTE

To deploy a New web service you must have sufficient permissions in the subscription to which you are deploying the web service. For more information, see [Manage a web service using the Azure Machine Learning Web Services portal](#).

To deploy a New web service derived from our experiment:

1. Click **Deploy Web Service** below the canvas and select **Deploy Web Service [New]**. Machine Learning Studio transfers you to the Azure Machine Learning web services **Deploy Experiment** page.
2. Enter a name for the web service.
3. For **Price Plan**, you can select an existing pricing plan, or select "Create new" and give the new plan a name and select the monthly plan option. The plan tiers default to the plans for your default region and your web service is deployed to that region.
4. Click **Deploy**.

After a few minutes, the **Quickstart** page for your web service opens.

You can configure the service by clicking the **Configure** tab. Here you can modify the service title and give it a description.

To test the web service, click the **Test** tab (see **Test the web service** below). For information on creating applications that can access the web service, click the **Consume** tab (the next step in this walkthrough will go into more detail).

TIP

You can update the web service after you've deployed it. For example, if you want to change your model, then you can edit the training experiment, tweak the model parameters, and click **Deploy Web Service**, selecting **Deploy Web Service [Classic]** or **Deploy Web Service [New]**. When you deploy the experiment again, it replaces the web service, now using your updated model.

Test the web service

When the web service is accessed, the user's data enters through the **Web service input** module where it's passed to the **Score Model** module and scored. The way we've set up the predictive experiment, the model expects data in the same format as the original credit risk dataset. The results are returned to the user from the web service through the **Web service output** module.

TIP

The way we have the predictive experiment configured, the entire results from the **Score Model** module are returned. This includes all the input data plus the credit risk value and the scoring probability. But you can return something different if you want - for example, you could return just the credit risk value. To do this, insert a **Project Columns** module between **Score Model** and the **Web service output** to eliminate columns you don't want the web service to return.

You can test a Classic web service either in **Machine Learning Studio** or in the **Azure Machine Learning Web Services** portal. You can test a New web service only in the **Machine Learning Web Services** portal.

TIP

When testing in the Azure Machine Learning Web Services portal, you can have the portal create sample data that you can use to test the Request-Response service. On the **Configure** page, select "Yes" for **Sample Data Enabled?**. When you open the Request-Response tab on the **Test** page, the portal fills in sample data taken from the original credit risk dataset.

Test a Classic web service

You can test a Classic web service in Machine Learning Studio or in the Machine Learning Web Services portal.

Test in Machine Learning Studio

1. On the **DASHBOARD** page for the web service, click the **Test** button under **Default Endpoint**. A dialog pops up and asks you for the input data for the service. These are the same columns that appeared in the original credit risk dataset.
2. Enter a set of data and then click **OK**.

Test in the Machine Learning Web Services portal

1. On the **DASHBOARD** page for the web service, click the **Test preview** link under **Default Endpoint**. The test page in the Azure Machine Learning Web Services portal for the web service endpoint opens and asks you for the input data for the service. These are the same columns that appeared in the original credit risk dataset.
2. Click **Test Request-Response**.

Test a New web service

You can test a New web service only in the Machine Learning Web Services portal.

1. In the [Azure Machine Learning Web Services](#) portal, click **Test** at the top of the page. The **Test** page opens and you can input data for the service. The input fields displayed correspond to the columns that appeared in the original credit risk dataset.
2. Enter a set of data and then click **Test Request-Response**.

The results of the test are displayed on the right-hand side of the page in the output column.

Manage the web service

Once you've deployed your web service, whether Classic or New, you can manage it from the [Microsoft Azure Machine Learning Web Services](#) portal.

To monitor the performance of your web service:

1. Sign in to the [Microsoft Azure Machine Learning Web Services](#) portal
2. Click **Web services**
3. Click your web service
4. Click the **Dashboard**

Next: [Access the web service](#)

Walkthrough Step 6: Access the Azure Machine Learning web service

9/25/2017 • 1 min to read • [Edit Online](#)

This is the last step of the walkthrough, [Develop a predictive analytics solution in Azure Machine Learning](#)

1. [Create a Machine Learning workspace](#)
2. [Upload existing data](#)
3. [Create a new experiment](#)
4. [Train and evaluate the models](#)
5. [Deploy the Web service](#)
6. **Access the Web service**

In the previous step in this walkthrough we deployed a web service that uses our credit risk prediction model. Now users are able to send data to it and receive results.

The Web service is an Azure web service that can receive and return data using REST APIs in one of two ways:

- **Request/Response** - The user sends one or more rows of credit data to the service by using an HTTP protocol, and the service responds with one or more sets of results.
- **Batch Execution** - The user stores one or more rows of credit data in an Azure blob and then sends the blob location to the service. The service scores all the rows of data in the input blob, stores the results in another blob, and returns the URL of that container.

The quickest and easiest way to access a Classic web service is through the [Azure ML Request-Response Service Web App](#) or [Azure ML Batch Execution Service Web App Template](#).

These web app templates can build a custom web app that knows your web service's input data and what it will return. All you need to do is provide access to your web service and data, and the template does the rest.

For more information on using the web app templates, see [Consume an Azure Machine Learning Web service with a web app template](#).

You can also develop a custom application to access the web service using starter code provided for you in R, C#, and Python programming languages.

You can find complete details in [How to consume an Azure Machine Learning Web service](#).

Data Science for Beginners video 1: The 5 questions data science answers

9/25/2017 • 4 min to read • [Edit Online](#)

Get a quick introduction to data science from *Data Science for Beginners* in five short videos from a top data scientist. These videos are basic but useful, whether you're interested in doing data science or you work with data scientists.

This first video is about the kinds of questions that data science can answer. To get the most out of the series, watch them all. [Go to the list of videos](#)

Other videos in this series

Data Science for Beginners is a quick introduction to data science taking about 25 minutes total. Check out all five videos:

- Video 1: The 5 questions data science answers
- Video 2: [Is your data ready for data science? \(4 min 56 sec\)](#)
- Video 3: [Ask a question you can answer with data \(4 min 17 sec\)](#)
- Video 4: [Predict an answer with a simple model \(7 min 42 sec\)](#)
- Video 5: [Copy other people's work to do data science \(3 min 18 sec\)](#)

Transcript: The 5 questions data science answers

Hi! Welcome to the video series *Data Science for Beginners*.

Data Science can be intimidating, so I'll introduce the basics here without any equations or computer programming jargon.

In this first video, we'll talk about "The 5 questions data science answers."

Data Science uses numbers and names (also known as categories or labels) to predict answers to questions.

It might surprise you, but *there are only five questions that data science answers*:

- Is this A or B?
- Is this weird?
- How much – or – How many?
- How is this organized?
- What should I do next?

Each one of these questions is answered by a separate family of machine learning methods, called algorithms.

It's helpful to think about an algorithm as a recipe and your data as the ingredients. An algorithm tells how to

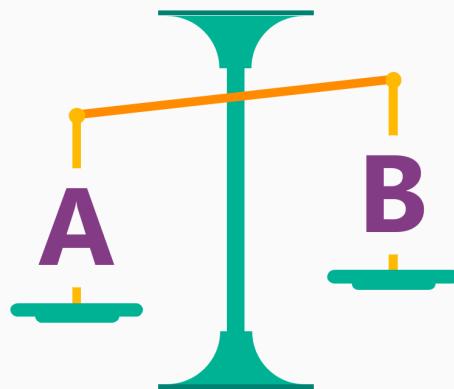
combine and mix the data in order to get an answer. Computers are like a blender. They do most of the hard work of the algorithm for you and they do it pretty fast.

Question 1: Is this A or B? uses classification algorithms

Let's start with the question: Is this A or B?

Is this A or B?

Classification algorithms



This family of algorithms is called two-class classification.

It's useful for any question that has just two possible answers.

For example:

- Will this tire fail in the next 1,000 miles: Yes or no?
- Which brings in more customers: a \$5 coupon or a 25% discount?

This question can also be rephrased to include more than two options: Is this A or B or C or D, etc.? This is called multiclass classification and it's useful when you have several — or several thousand — possible answers.

Multiclass classification chooses the most likely one.

Question 2: Is this weird? uses anomaly detection algorithms

The next question data science can answer is: Is this weird? This question is answered by a family of algorithms called anomaly detection.

Is this weird?

Anomaly detection algorithms



If you have a credit card, you've already benefited from anomaly detection. Your credit card company analyzes your purchase patterns, so that they can alert you to possible fraud. Charges that are "weird" might be a purchase at a store where you don't normally shop or buying an unusually pricey item.

This question can be useful in lots of ways. For instance:

- If you have a car with pressure gauges, you might want to know: Is this pressure gauge reading normal?
- If you're monitoring the internet, you'd want to know: Is this message from the internet typical?

Anomaly detection flags unexpected or unusual events or behaviors. It gives clues where to look for problems.

Question 3: How much? or How many? uses regression algorithms

Machine learning can also predict the answer to How much? or How many? The algorithm family that answers this question is called regression.

How much? How many?

Regression algorithms

Monday	Tuesday
	
72°	

Regression algorithms make numerical predictions, such as:

- What will the temperature be next Tuesday?
- What will my fourth quarter sales be?

They help answer any question that asks for a number.

Question 4: How is this organized? uses clustering algorithms

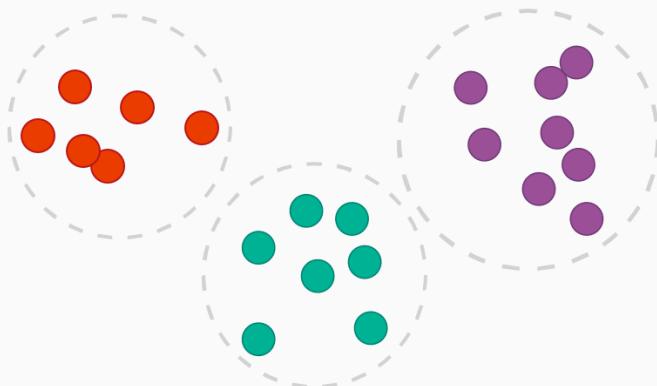
Now the last two questions are a bit more advanced.

Sometimes you want to understand the structure of a data set - How is this organized? For this question, you don't have examples that you already know outcomes for.

There are a lot of ways to tease out the structure of data. One approach is clustering. It separates data into natural "clumps," for easier interpretation. With clustering, there is no one right answer.

How is this organized?

Clustering Algorithms



Common examples of clustering questions are:

- Which viewers like the same types of movies?
- Which printer models fail the same way?

By understanding how data is organized, you can better understand - and predict - behaviors and events.

Question 5: What should I do now? uses reinforcement learning algorithms

The last question – What should I do now? – uses a family of algorithms called reinforcement learning.

Reinforcement learning was inspired by how the brains of rats and humans respond to punishment and rewards. These algorithms learn from outcomes, and decide on the next action.

Typically, reinforcement learning is a good fit for automated systems that have to make lots of small decisions without human guidance.

What should I do now?

Reinforcement Learning Algorithms



Questions it answers are always about what action should be taken - usually by a machine or a robot. Examples are:

- If I'm a temperature control system for a house: Adjust the temperature or leave it where it is?
- If I'm a self-driving car: At a yellow light, brake or accelerate?
- For a robot vacuum: Keep vacuuming, or go back to the charging station?

Reinforcement learning algorithms gather data as they go, learning from trial and error.

So that's it - The 5 questions data science can answer.

Next steps

- [Try a first data science experiment with Machine Learning Studio](#)
- [Get an introduction to Machine Learning on Microsoft Azure](#)

Is your data ready for data science?

9/25/2017 • 4 min to read • [Edit Online](#)

Video 2: Data Science for Beginners series

Learn how to evaluate your data to make sure it meets basic criteria to be ready for data science.

To get the most out of the series, watch them all. [Go to the list of videos](#)

Other videos in this series

Data Science for Beginners is a quick introduction to data science in five short videos.

- Video 1: [The 5 questions data science answers](#) (5 min 14 sec)
- Video 2: Is your data ready for data science?
- Video 3: [Ask a question you can answer with data](#) (4 min 17 sec)
- Video 4: [Predict an answer with a simple model](#) (7 min 42 sec)
- Video 5: [Copy other people's work to do data science](#) (3 min 18 sec)

Transcript: Is your data ready for data science?

Welcome to "Is your data ready for data science?" the second video in the series *Data Science for Beginners*.

Before data science can give you the answers you want, you have to give it some high-quality raw materials to work with. Just like making a pizza, the better the ingredients you start with, the better the final product.

Criteria for data

So, in the case of data science, there are some ingredients that we need to pull together.

We need data that is:

- Relevant
- Connected
- Accurate
- Enough to work with

Is your data relevant?

So the first ingredient - we need data that's relevant.

Irrelevant Data

Price of milk (\$/gal)	Red Sox batting avg.	Blood alcohol content (%)
3.79	.304	.03
3.45	.320	.09
4.06	.259	.01
3.89	.298	.05
4.12	.332	.13
3.92	.270	.06
3.23	.294	.10

Relevant Data

Body mass (kg)	Margaritas	Blood alcohol content (%)
103	3	.03
67	5	.09
87	1	.01
52	2	.05
73	5	.13
79	3	.06
110	7	.10

Look at the table on the left. We met seven people outside of Boston bars, measured their blood alcohol level, the Red Sox batting average in their last game, and the price of milk in the nearest convenience store.

This is all perfectly legitimate data. Its only fault is that it isn't relevant. There's no obvious relationship between these numbers. If I gave you the current price of milk and the Red Sox batting average, there's no way you could guess my blood alcohol content.

Now look at the table on the right. This time we measured each person's body mass and counted the number of drinks they've had. The numbers in each row are now relevant to each other. If I gave you my body mass and the number of Margaritas I've had, you could make a guess at my blood alcohol content.

Do you have connected data?

The next ingredient is connected data.

Disconnected Data

Grill temp. (Fahrenheit)	Weight of beef patty (lb)	Burger rating (out of 10)
	.33	8.2
	.24	5.6
550		7.8
725	.45	9.4
600		8.2
625		6.8
	.49	4.2

Connected Data

Grill temp. (Fahrenheit)	Weight of beef patty (lb)	Burger rating (out of 10)
575	.33	8.2
550	.24	5.6
550	.69	7.8
725	.45	9.4
600	.57	8.2
625	.36	6.8
550	.49	4.2

Here is some relevant data on the quality of hamburgers: grill temperature, patty weight, and rating in the local food magazine. But notice the gaps in the table on the left.

Most data sets are missing some values. It's common to have holes like this and there are ways to work around them. But if there's too much missing, your data begins to look like Swiss cheese.

If you look at the table on the left, there's so much missing data, it's hard to come up with any kind of relationship

between grill temperature and patty weight. This is an example of disconnected data.

The table on the right, though, is full and complete - an example of connected data.

Is your data accurate?

The next ingredient we need is accuracy. Here are four targets that we'd like to hit with arrows.



Look at the target in the upper right. We've got a tight grouping right around the bullseye. That, of course, is accurate. Oddly, in the language of data science, our performance on the target right below it is also considered accurate.

If you were to map out the center of these arrows, you'd see that it's very close to the bullseye. The arrows are spread out all around the target, so they're considered imprecise, but they're centered around the bullseye, so they're considered accurate.

Now look at the upper-left target. Here our arrows hit very close together, a tight grouping. They're precise, but they're inaccurate because the center is way off the bullseye. And, of course, the arrows in the bottom-left target are both inaccurate and imprecise. This archer needs more practice.

Do you have enough data to work with?

Finally, ingredient #4 - we need to have enough data.

Barely enough data



Think of each data point in your table as being a brush stroke in a painting. If you have only a few of them, the painting can be pretty fuzzy - it's hard to tell what it is.

If you add some more brush strokes, then your painting starts to get a little sharper.

When you have barely enough strokes, you can see just enough to make some broad decisions. Is it somewhere I might want to visit? It looks bright, that looks like clean water – yes, that's where I'm going on vacation.

As you add more data, the picture becomes clearer and you can make more detailed decisions. Now I can look at the three hotels on the left bank. You know, I really like the architectural features of the one in the foreground. I'll stay there, on the third floor.

With data that's relevant, connected, accurate, and enough, we have all the ingredients we need to do some high-quality data science.

Be sure to check out the other four videos in *Data Science for Beginners* from Microsoft Azure Machine Learning.

Next steps

- [Try a first data science experiment with Machine Learning Studio](#)
- [Get an introduction to Machine Learning on Microsoft Azure](#)

Ask a question you can answer with data

9/25/2017 • 4 min to read • [Edit Online](#)

Video 3: Data Science for Beginners series

Learn how to formulate a data science problem into a question in Data Science for Beginners video 3. This video includes a comparison of questions for classification and regression algorithms.

To get the most out of the series, watch them all. [Go to the list of videos](#)

Other videos in this series

Data Science for Beginners is a quick introduction to data science in five short videos.

- Video 1: [The 5 questions data science answers](#) (5 min 14 sec)
- Video 2: [Is your data ready for data science?](#) (4 min 56 sec)
- Video 3: Ask a question you can answer with data
- Video 4: [Predict an answer with a simple model](#) (7 min 42 sec)
- Video 5: [Copy other people's work to do data science](#) (3 min 18 sec)

Transcript: Ask a question you can answer with data

Welcome to the third video in the series "Data Science for Beginners."

In this one, you'll get some tips for formulating a question you can answer with data.

You might get more out of this video, if you first watch the two earlier videos in this series: "The 5 questions data science can answer" and "Is your data is ready for data science?"

Ask a sharp question

We've talked about how data science is the process of using names (also called categories or labels) and numbers to predict an answer to a question. But it can't be just any question; it has to be a *sharp question*.

A vague question doesn't have to be answered with a name or a number. A sharp question must.

Imagine you found a magic lamp with a genie who will truthfully answer any question you ask. But it's a mischievous genie, and he'll try to make his answer as vague and confusing as he can get away with. You want to pin him down with a question so airtight that he can't help but tell you what you want to know.

If you were to ask a vague question, like "What's going to happen with my stock?", the genie might answer, "The price will change". That's a truthful answer, but it's not very helpful.

But if you were to ask a sharp question, like "What will my stock's sale price be next week?", the genie can't help but give you a specific answer and predict a sale price.

Examples of your answer: Target data

Once you formulate your question, check to see whether you have examples of the answer in your data.

If our question is "What will my stock's sale price be next week?" then we have to make sure our data includes the stock price history.

If our question is "Which car in my fleet is going to fail first?" then we have to make sure our data includes information about previous failures.

Examples of the answer: Target data



These examples of answers are called a target. A target is what we are trying to predict about future data points, whether it's a category or a number.

If you don't have any target data, you'll need to get some. You won't be able to answer your question without it.

Reformulate your question

Sometimes you can reword your question to get a more useful answer.

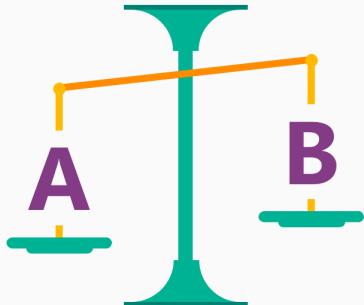
The question "Is this data point A or B?" predicts the category (or name or label) of something. To answer it, we use a *classification algorithm*.

The question "How much?" or "How many?" predicts an amount. To answer it we use a *regression algorithm*.

To see how we can transform these, let's look at the question, "Which news story is the most interesting to this reader?" It asks for a prediction of a single choice from many possibilities - in other words "Is this A or B or C or D?" - and would use a classification algorithm.

But, this question may be easier to answer if you reword it as "How interesting is each story on this list to this reader?" Now you can give each article a numerical score, and then it's easy to identify the highest-scoring article. This is a rephrasing of the classification question into a regression question or How much?

Reformulate your question



How you ask a question is a clue to which algorithm can give you an answer.

You'll find that certain families of algorithms - like the ones in our news story example - are closely related. You can reformulate your question to use the algorithm that gives you the most useful answer.

But, most important, ask that sharp question - the question that you can answer with data. And be sure you have the right data to answer it.

We've talked about some basic principles for asking a question you can answer with data.

Be sure to check out the other videos in "Data Science for Beginners" from Microsoft Azure Machine Learning.

Next steps

- [Try a first data science experiment with Machine Learning Studio](#)
- [Get an introduction to Machine Learning on Microsoft Azure](#)

Predict an answer with a simple model

9/25/2017 • 5 min to read • [Edit Online](#)

Video 4: Data Science for Beginners series

Learn how to create a simple regression model to predict the price of a diamond in Data Science for Beginners video 4. We'll draw a regression model with target data.

To get the most out of the series, watch them all. [Go to the list of videos](#)

Other videos in this series

Data Science for Beginners is a quick introduction to data science in five short videos.

- Video 1: [The 5 questions data science answers](#) (5 min 14 sec)
- Video 2: [Is your data ready for data science?](#) (4 min 56 sec)
- Video 3: [Ask a question you can answer with data](#) (4 min 17 sec)
- Video 4: Predict an answer with a simple model
- Video 5: [Copy other people's work to do data science](#) (3 min 18 sec)

Transcript: Predict an answer with a simple model

Welcome to the fourth video in the "Data Science for Beginners" series. In this one, we'll build a simple model and make a prediction.

A *model* is a simplified story about our data. I'll show you what I mean.

Collect relevant, accurate, connected, enough data

Say I want to shop for a diamond. I have a ring that belonged to my grandmother with a setting for a 1.35 carat diamond, and I want to get an idea of how much it will cost. I take a notepad and pen into the jewelry store, and I write down the price of all of the diamonds in the case and how much they weigh in carats. Starting with the first diamond - it's 1.01 carats and \$7,366.

Now I go through and do this for all the other diamonds in the store.

<u>Carats</u>	<u>Price</u>
1.01	7,366
.49	985
.31	544
1.51	9,140
.37	493
.73	3,011
1.53	11,413
.56	1,814
.41	876
.74	2,690
.63	1,190
.6	4,172
2.06	11,764
1.1	4,682
1.31	6,171

Notice that our list has two columns. Each column has a different attribute - weight in carats and price - and each row is a single data point that represents a single diamond.

We've actually created a small data set here - a table. Notice that it meets our criteria for quality:

- The data is **relevant** - weight is definitely related to price
- It's **accurate** - we double-checked the prices that we write down
- It's **connected** - there are no blank spaces in either of these columns
- And, as we'll see, it's **enough** data to answer our question

Ask a sharp question

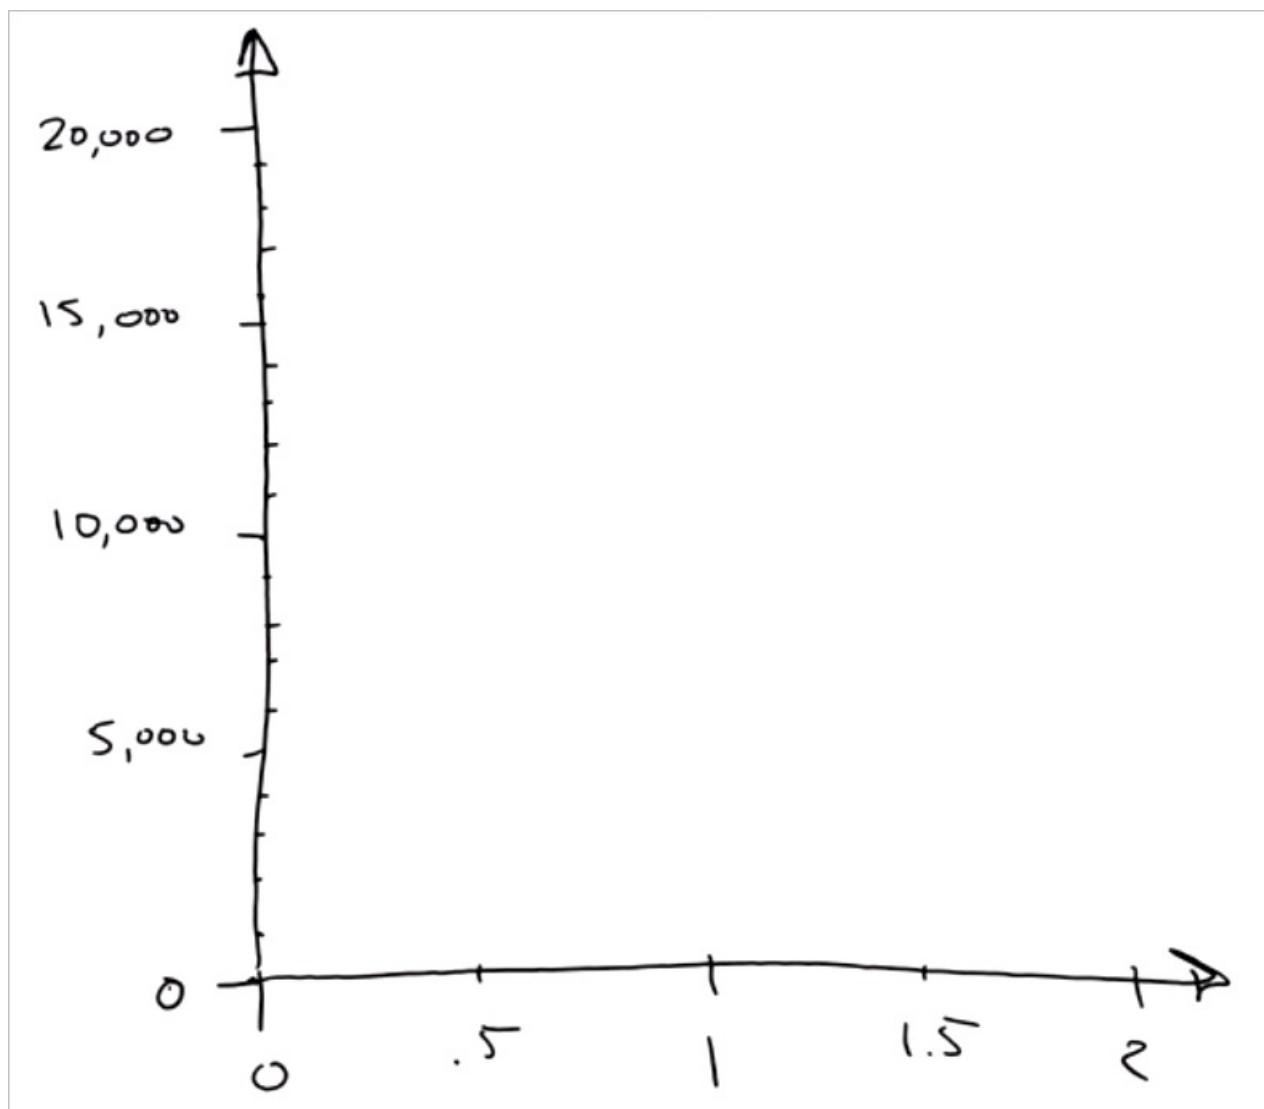
Now we'll pose our question in a sharp way: "How much will it cost to buy a 1.35 carat diamond?"

Our list doesn't have a 1.35 carat diamond in it, so we'll have to use the rest of our data to get an answer to the question.

Plot the existing data

The first thing we'll do is draw a horizontal number line, called an axis, to chart the weights. The range of the weights is 0 to 2, so we'll draw a line that covers that range and put ticks for each half carat.

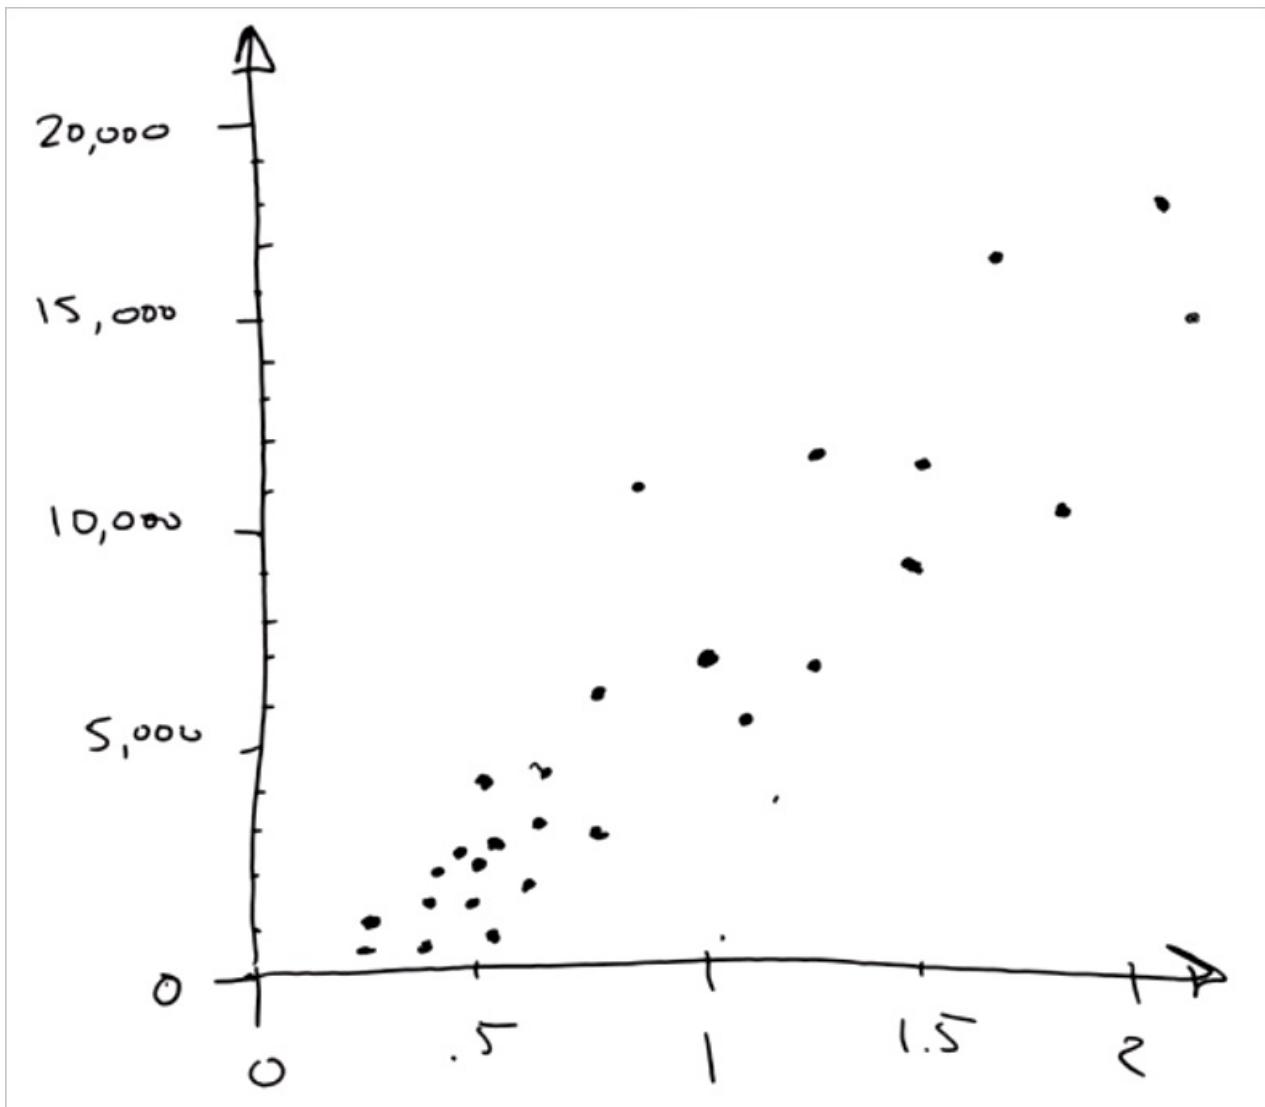
Next we'll draw a vertical axis to record the price and connect it to the horizontal weight axis. This will be in units of dollars. Now we have a set of coordinate axes.



We're going to take this data now and turn it into a *scatter plot*. This is a great way to visualize numerical data sets.

For the first data point, we eyeball a vertical line at 1.01 carats. Then, we eyeball a horizontal line at \$7,366. Where they meet, we draw a dot. This represents our first diamond.

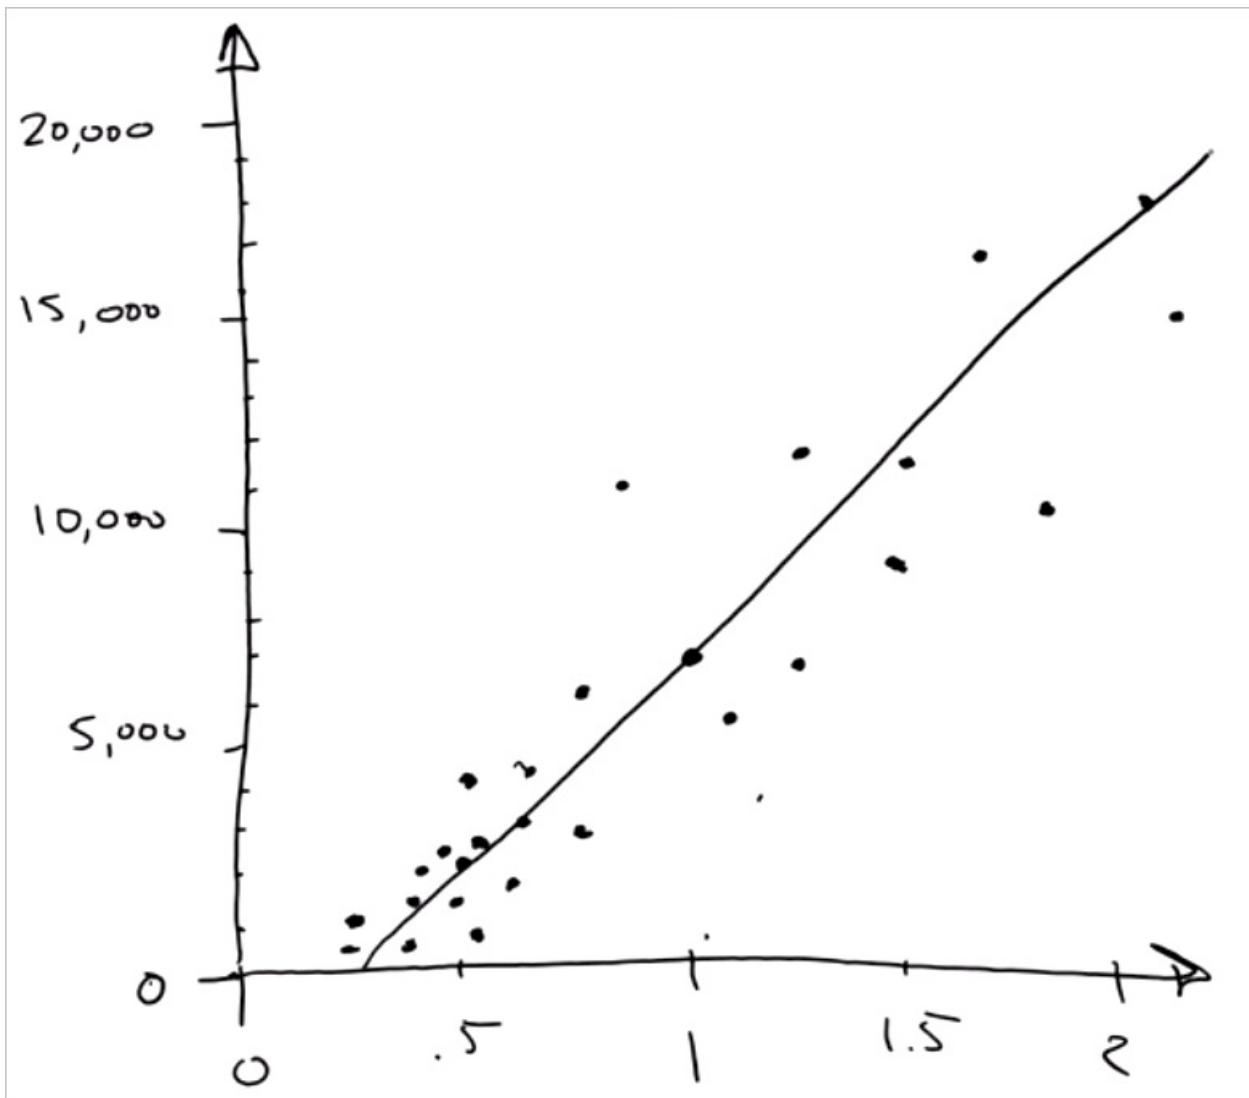
Now we go through each diamond on this list and do the same thing. When we're through, this is what we get: a bunch of dots, one for each diamond.



Draw the model through the data points

Now if you look at the dots and squint, the collection looks like a fat, fuzzy line. We can take our marker and draw a straight line through it.

By drawing a line, we created a *model*. Think of this as taking the real world and making a simplistic cartoon version of it. Now the cartoon is wrong - the line doesn't go through all the data points. But, it's a useful simplification.



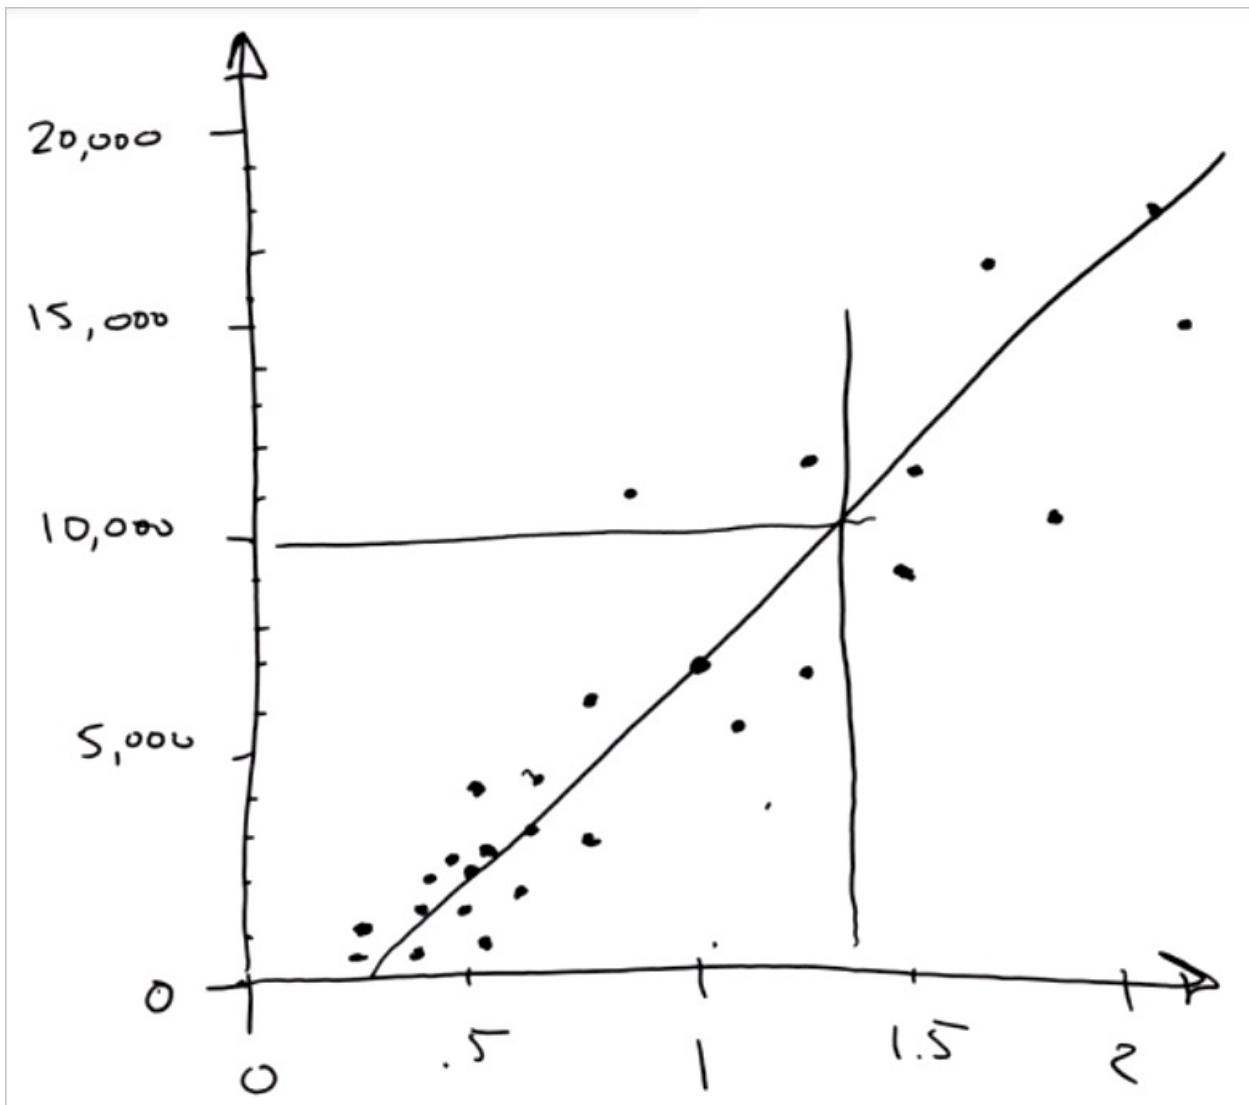
The fact that all the dots don't go exactly through the line is OK. Data scientists explain this by saying that there's the model - that's the line - and then each dot has some *noise* or *variance* associated with it. There's the underlying perfect relationship, and then there's the gritty, real world that adds noise and uncertainty.

Because we're trying to answer the question *How much?* this is called a *regression*. And because we're using a straight line, it's a *linear regression*.

Use the model to find the answer

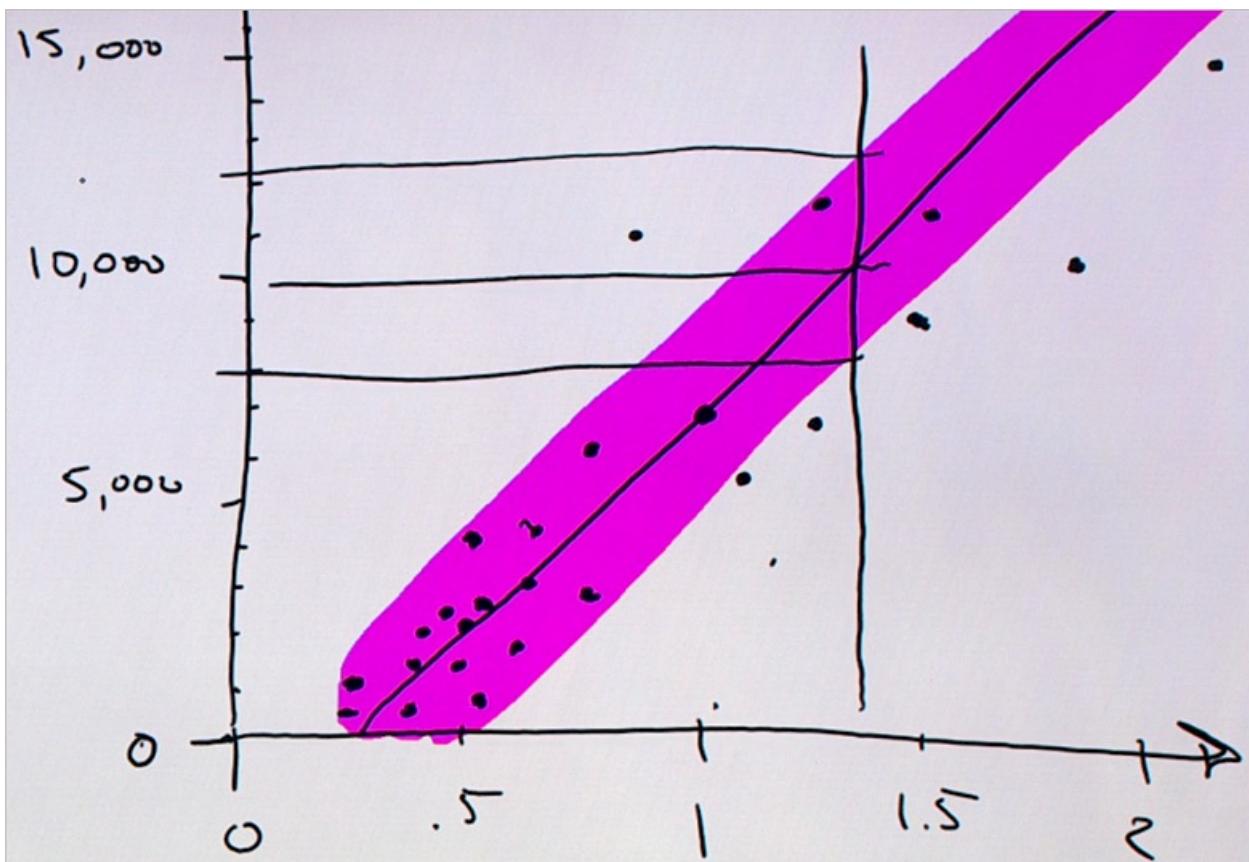
Now we have a model and we ask it our question: How much will a 1.35 carat diamond cost?

To answer our question, we eyeball 1.35 carats and draw a vertical line. Where it crosses the model line, we eyeball a horizontal line to the dollar axis. It hits right at 10,000. Boom! That's the answer: A 1.35 carat diamond costs about \$10,000.



Create a confidence interval

It's natural to wonder how precise this prediction is. It's useful to know whether the 1.35 carat diamond will be very close to \$10,000, or a lot higher or lower. To figure this out, let's draw an envelope around the regression line that includes most of the dots. This envelope is called our *confidence interval*: We're pretty confident that prices fall within this envelope, because in the past most of them have. We can draw two more horizontal lines from where the 1.35 carat line crosses the top and the bottom of that envelope.



Now we can say something about our confidence interval: We can say confidently that the price of a 1.35 carat diamond is about \$10,000 - but it might be as low as \$8,000 and it might be as high as \$12,000.

We're done, with no math or computers

We did what data scientists get paid to do, and we did it just by drawing:

- We asked a question that we could answer with data
- We built a *model* using *linear regression*
- We made a *prediction*, complete with a *confidence interval*

And we didn't use math or computers to do it.

Now if we'd had more information, like...

- the cut of the diamond
- color variations (how close the diamond is to being white)
- the number of inclusions in the diamond

...then we would have had more columns. In that case, math becomes helpful. If you have more than two columns, it's hard to draw dots on paper. The math lets you fit that line or that plane to your data very nicely.

Also, if instead of just a handful of diamonds, we had two thousand or two million, then you can do that work much faster with a computer.

Today, we've talked about how to do linear regression, and we made a prediction using data.

Be sure to check out the other videos in "Data Science for Beginners" from Microsoft Azure Machine Learning.

Next steps

- Try a first data science experiment with Machine Learning Studio
- Get an introduction to Machine Learning on Microsoft Azure

Copy other people's work to do data science

9/25/2017 • 3 min to read • [Edit Online](#)

Video 5: Data Science for Beginners series

One of the trade secrets of data science is getting other people to do your work for you. Find a clustering algorithm example in Cortana Analytics Gallery to use for your own machine learning experiment.

To get the most out of the series, watch them all. [Go to the list of videos](#)

Other videos in this series

Data Science for Beginners is a quick introduction to data science in five short videos.

- Video 1: [The 5 questions data science answers](#) (5 min 14 sec)
- Video 2: [Is your data ready for data science?](#) (4 min 56 sec)
- Video 3: [Ask a question you can answer with data](#) (4 min 17 sec)
- Video 4: [Predict an answer with a simple model](#) (7 min 42 sec)
- Video 5: Copy other people's work to do data science

Transcript: Copy other people's work to do data science

Welcome to the fifth video in the series "Data Science for Beginners."

In this one, you'll discover a place to find examples that you can borrow from as a starting point for your own work. You might get the most out of this video if you first watch the earlier videos in this series.

One of the trade secrets of data science is getting other people to do your work for you.

Find examples in the Cortana Intelligence Gallery

Microsoft has a cloud-based service called [Azure Machine Learning](#) that you're welcome to try for free. It provides you with a workspace where you can experiment with different machine learning algorithms, and, when you've got your solution worked out, you can launch it as a web service.

Part of this service is something called the [Cortana Intelligence Gallery](#). It contains resources, including a collection of Azure Machine Learning experiments, or models, that people have built and contributed for others to use. These experiments are a great way to leverage the thought and hard work of others to get you started on your own solutions.

You can find the gallery at aka.ms/CortanaIntelligenceGallery. Everyone is welcome to browse through it.

Cortana Intelligence Gallery

Browse all Industries Solutions Experiments Machine Learning APIs Custom Modules Learning More

Cortana Intelligence Gallery enables our growing community of developers and data scientists to share their analytics solutions. [Learn how to contribute.](#)

The screenshot shows the Cortana Intelligence Gallery homepage. At the top, there's a navigation bar with links for Browse all, Industries, Solutions, Experiments, Machine Learning APIs, Custom Modules, Learning, and More. Below the navigation is a main content area with a teal header. On the left, there's a large card for a 'Personalized Offers' solution by Microsoft, featuring a computer monitor icon and a person icon. To the right, there are four smaller cards: a 'TUTORIAL' for exploring Azure data services, an 'EXPERIMENT' for reading data from Azure DocumentDB in Azure Machine, a 'CUSTOM MODULE' for adding mean and deviation, and another 'SOLUTION' for demand forecasting. Below these cards is a section titled 'Recently added' with four more cards: 'Kafka for HDInsight', 'User Segmentation', 'Traffic Prediction', and a tree diagram. A 'See all' button is located in the top right corner of this section.

If you click **Experiments** at the top, you'll see a number of the most recent and popular experiments in the gallery. You can search through the rest of experiments by clicking **Browse All** at the top of the screen, and there you can enter search terms and choose search filters.

Find and use a clustering algorithm example

So, for instance, let's say you want to see an example of how clustering works, so you search for "**clustering sweep**" experiments.

The screenshot shows the Cortana Intelligence Gallery search results for the term "clustering sweep". The search bar at the top contains the text "clustering sweep", which is circled in red. Below the search bar, the navigation menu is visible with "Browse all" highlighted in green. The search results show a single experiment card for "clustering sweep" by Gopi Kumar, which includes a thumbnail of a city map with colored dots, the title, author, and a brief description.

Here's an interesting one that someone contributed to the gallery.

EXPERIMENT

Clustering sweep: diabetes dataset

This experiment demonstrates how to use the new Sweep Clustering module to select the best number of centroids and optimize other para...

K-Means Clustering

🕒 1430 ⬇ 733 9 months ago

Jeannine Takaki

Click on that experiment and you get a web page that describes the work that this contributor did, along with some of their results.

Cortana Intelligence Gallery

Browse all Industries Solution Templates Experiments Machine Learning APIs Notebooks Competitions More

EXPERIMENT

Clustering sweep: diabetes dataset

Jeannine Takaki • published on October 28, 2015

Summary

This experiment demonstrates how to use the new Sweep Clustering module to select the best number of centroids and optimize other parameters

Description

Summary

This experiment uses a parameter sweep with the K-means clustering algorithm to select the best number of clusters and the best initial centroid, based on a clustering metric you select. In the experiment, the original dataset is divided into two parts, to handle missing data and to compare the effect of variables on clustering. The results of the models are compared using a principal components graph. The experiment also demonstrates how you can use the [Sweep Clustering](#) module to fill in values for a label column.

Understanding the Data

The dataset contains 768 rows, from a larger dataset that studies the incidence of diabetes among different populations.

The following clinical values are included, which are often used in diagnosing diabetes.

Triceps skin fold measurements (TSF) and body mass index (BMI) are thought to be correlated with

[Open in Studio](#)

+ Add to Collection

🕒 1428 views ⬇ 730 downloads

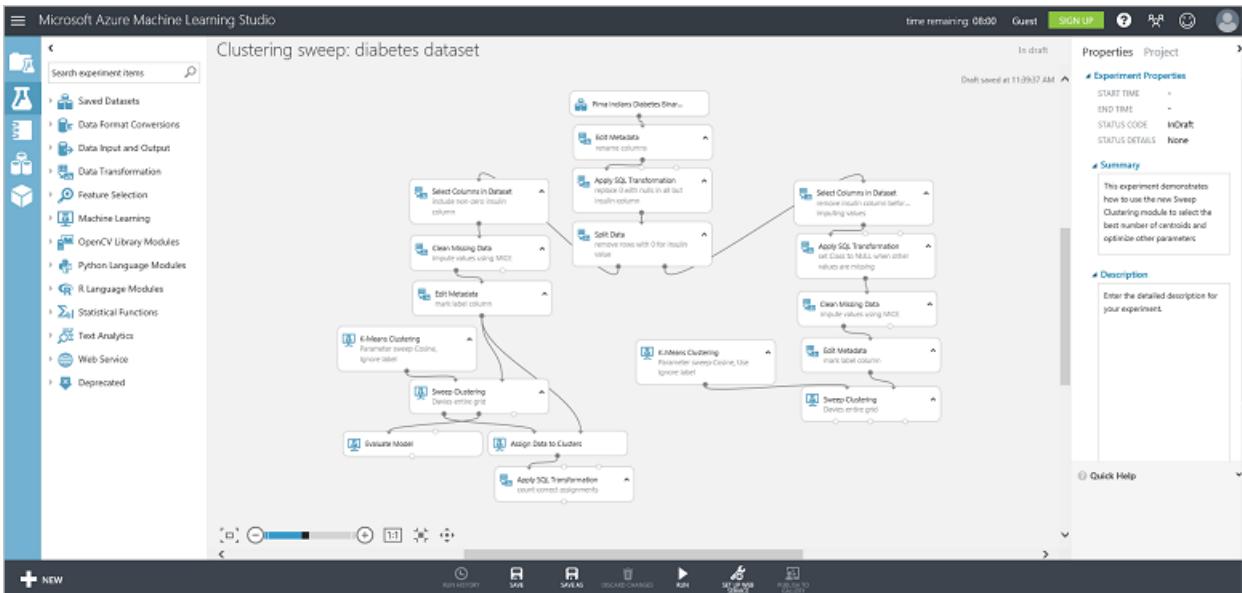
Tweet Share

ALGORITHMS
K-Means Clustering

Notice the link that says **Open in Studio**.



I can click on that and it takes me right to **Azure Machine Learning Studio**. It creates a copy of the experiment and puts it in my own workspace. This includes the contributor's dataset, all the processing that they did, all of the algorithms that they used, and how they saved out the results.



And now I have a starting point. I can swap out their data for my own and do my own tweaking of the model. This gives me a running start, and it lets me build on the work of people who really know what they're doing.

Find experiments that demonstrate machine learning techniques

There are other experiments in the [Cortana Intelligence Gallery](#) that were contributed specifically to provide how-to examples for people new to data science. For instance, there's an experiment in the gallery that demonstrates how to handle missing values ([Methods for handling missing values](#)). It walks you through 15 different ways of substituting empty values, and talks about the benefits of each method and when to use it.

Cortana Intelligence Gallery

Browse all Industries Solution Templates Experiments Machine Learning APIs Notebooks Competitions More

EXPERIMENT

Methods for handling missing values

Brandon Rohrer • published on September 28, 2015

Summary

This experiment illustrates a variety methods for handling missing data on a sample data set.

Description

Real world data is usually missing values, which trip up a lot of machine learning algorithms. There are lots of tricks for dealing with these, but you have to be careful. The way in which you fill them can change the result dramatically. Being explicit and thoughtful about how you handle missing values will get you the very best results.

I've illustrated a large handful of approaches to missing values here in a fake data set. The data shows a group of employees, some of their personal data, and some data regarding an upcoming office party. In every case, knowing what the data means is the most important part of handling it well.

Column 0	age	years_seniority	income	parking_space	attending_party	entree	pets	emergency_contact
Tony	48	27		1	5	shrimp		Pepper
Donald	67	25	86	10	2	beef		Jane
Henry	69	21	95	6	1	chicken	62	Janet
Janet	62	21	110	3	1	beef		Henry

Open in Studio

+ Add to Collection

863 views 95 downloads

[Tweet](#) [Share](#)

TAGS

missing_values data_quality method

[Cortana Intelligence Gallery](#) is a place to find working experiments that you can use as a starting point for your own solutions.

Be sure to check out the other videos in "Data Science for Beginners" from Microsoft Azure Machine Learning.

Next steps

- Try your first data science experiment with Azure Machine Learning
- Get an introduction to Machine Learning on Microsoft Azure

Quickstart tutorial for the R programming language for Azure Machine Learning

9/25/2017 • 53 min to read • [Edit Online](#)

Introduction

This quickstart tutorial helps you quickly start extending Azure Machine Learning by using the R programming language. Follow this R programming tutorial to create, test and execute R code within Azure Machine Learning. As you work through tutorial, you will create a complete forecasting solution by using the R language in Azure Machine Learning.

Microsoft Azure Machine Learning contains many powerful machine learning and data manipulation modules. The powerful R language has been described as the lingua franca of analytics. Happily, analytics and data manipulation in Azure Machine Learning can be extended by using R. This combination provides the scalability and ease of deployment of Azure Machine Learning with the flexibility and deep analytics of R.

NOTE

Try Azure Machine Learning for free

No credit card or Azure subscription required. [Get started now.](#)

Forecasting and the dataset

Forecasting is a widely employed and quite useful analytical method. Common uses range from predicting sales of seasonal items, determining optimal inventory levels, to predicting macroeconomic variables. Forecasting is typically done with time series models.

Time series data is data in which the values have a time index. The time index can be regular, e.g. every month or every minute, or irregular. A time series model is based on time series data. The R programming language contains a flexible framework and extensive analytics for time series data.

In this quickstart guide we will be working with California dairy production and pricing data. This data includes monthly information on the production of several dairy products and the price of milk fat, a benchmark commodity.

The data used in this article, along with R scripts, can be [downloaded here](#). This data was originally synthesized from information available from the University of Wisconsin at <http://future.aae.wisc.edu/tab/production.html>.

Organization

We will progress through several steps as you learn how to create, test and execute analytics and data manipulation R code in the Azure Machine Learning environment.

- First we will explore the basics of using the R language in the Azure Machine Learning Studio environment.
- Then we progress to discussing various aspects of I/O for data, R code and graphics in the Azure Machine Learning environment.
- We will then construct the first part of our forecasting solution by creating code for data cleaning and transformation.
- With our data prepared we will perform an analysis of the correlations between several of the variables in our dataset.
- Finally, we will create a seasonal time series forecasting model for milk production.

Interact with R language in Machine Learning Studio

This section takes you through some basics of interacting with the R programming language in the Machine Learning Studio environment. The R language provides a powerful tool to create customized analytics and data manipulation modules within the Azure Machine Learning environment.

I will use RStudio to develop, test and debug R code on a small scale. This code is then cut and paste into an [Execute R Script](#) module in Machine Learning Studio ready to run.

The Execute R Script module

Within Machine Learning Studio, R scripts are run within the [Execute R Script](#) module. An example of the [Execute R Script](#) module in Machine Learning Studio is shown in Figure 1.

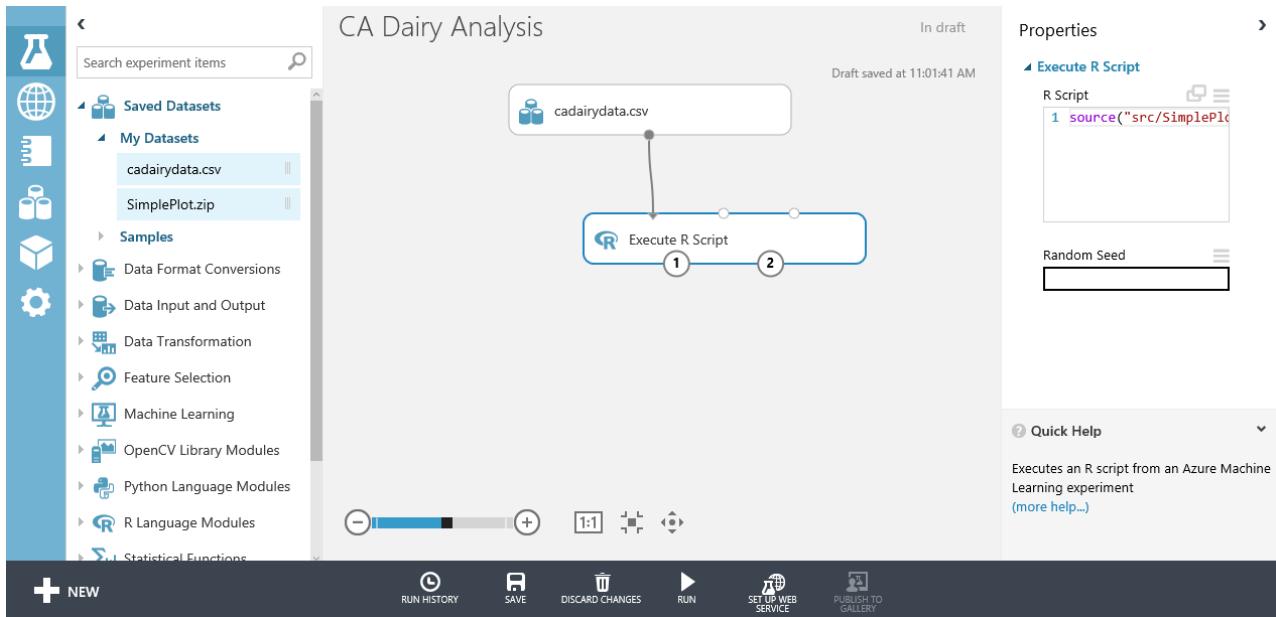


Figure 1. The Machine Learning Studio environment showing the [Execute R Script](#) module selected.

Referring to Figure 1, let's look at some of the key parts of the Machine Learning Studio environment for working with the [Execute R Script](#) module.

- The modules in the experiment are shown in the center pane.
- The upper part of the right pane contains a window to view and edit your R scripts.
- The lower part of right pane shows some properties of the [Execute R Script](#). You can view the error and output logs by clicking on the appropriate spots of this pane.

We will, of course, be discussing the [Execute R Script](#) in greater detail in the rest of this document.

When working with complex R functions, I recommend that you edit, test and debug in RStudio. As with any software development, extend your code incrementally and test it on small simple test cases. Then cut and paste your functions into the R script window of the [Execute R Script](#) module. This approach allows you to harness both the RStudio integrated development environment (IDE) and the power of Azure Machine Learning.

Execute R code

Any R code in the [Execute R Script](#) module will execute when you run the experiment by clicking on the **Run** button. When execution has completed, a check mark will appear on the [Execute R Script](#) icon.

Defensive R coding for Azure Machine Learning

If you are developing R code for, say, a web service by using Azure Machine Learning, you should definitely plan how your code will deal with an unexpected data input and exceptions. To maintain clarity, I have not included much in the way of checking or exception handling in most of the code examples shown. However, as we proceed I will give you several examples of functions by using R's exception handling capability.

If you need a more complete treatment of R exception handling, I recommend you read the applicable sections of the book by Wickham listed in [Appendix B - Further Reading](#).

Debug and test R in Machine Learning Studio

To reiterate, I recommend you test and debug your R code on a small scale in RStudio. However, there are cases where you will need to track down R code problems in the [Execute R Script](#) itself. In addition, it is good practice to check your results in Machine Learning Studio.

Output from the execution of your R code and on the Azure Machine Learning platform is found primarily in output.log. Some additional information will be seen in error.log.

If an error occurs in Machine Learning Studio while running your R code, your first course of action should be to look at error.log. This file can contain useful error messages to help you understand and correct your error. To view error.log, click on **View error log** on the **properties pane** for the [Execute R Script](#) containing the error.

For example, I ran the following R code, with an undefined variable y, in an [Execute R Script](#) module:

```
x<- 1.0  
z<- x+y
```

This code fails to execute, resulting in an error condition. Clicking on **View error log** on the **properties pane** produces the display shown in Figure 2.

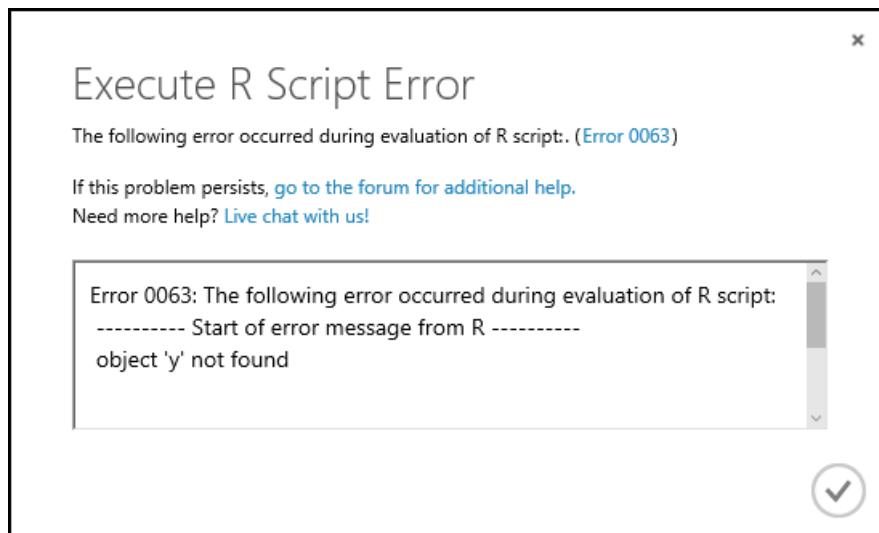


Figure 2. Error message pop-up.

It looks like we need to look in output.log to see the R error message. Click on the [Execute R Script](#) and then click on the **View output.log** item on the **properties pane** to the right. A new browser window opens, and I see the following.

```
[Critical] Error: Error 0063: The following error occurred during evaluation of R script:  
----- Start of error message from R -----  
object 'y' not found  
  
object 'y' not found  
----- End of error message from R -----
```

This error message contains no surprises and clearly identifies the problem.

To inspect the value of any object in R, you can print these values to the output.log file. The rules for examining object values are essentially the same as in an interactive R session. For example, if you type a variable name on a line, the value of the object will be printed to the output.log file.

Packages in Machine Learning Studio

Azure Machine Learning comes with over 350 preinstalled R language packages. You can use the following code in the [Execute R Script](#) module to retrieve a list of the preinstalled packages.

```
data.set <- data.frame(installed.packages())
maml.mapOutputPort("data.set")
```

If you don't understand the last line of this code at the moment, read on. In the rest of this document we will extensively discuss using R in the Azure Machine Learning environment.

Introduction to RStudio

RStudio is a widely used IDE for R. I will use RStudio for editing, testing and debugging some of the R code used in this quick start guide. Once R code is tested and ready, you simply cut and paste from the RStudio editor into a Machine Learning Studio [Execute R Script](#) module.

If you do not have the R programming language installed on your desktop machine, I recommend you do so now. Free downloads of open source R language are available at the Comprehensive R Archive Network (CRAN) at <http://www.r-project.org/>. There are downloads available for Windows, Mac OS, and Linux/UNIX. Choose a nearby mirror and follow the download directions. In addition, CRAN contains a wealth of useful analytics and data manipulation packages.

If you are new to RStudio, you should download and install the desktop version. You can find the RStudio downloads for Windows, Mac OS, and Linux/UNIX at <http://www.rstudio.com/products/RStudio/>. Follow the directions provided to install RStudio on your desktop machine.

A tutorial introduction to RStudio is available at <https://support.rstudio.com/hc/sections/200107586-Using-RStudio>.

I provide some additional information on using RStudio in [Appendix A](#).

Get data in and out of the Execute R Script module

In this section we will discuss how you get data into and out of the [Execute R Script](#) module. We will review how to handle various data types read into and out of the [Execute R Script](#) module.

The complete code for this section is in the zip file you downloaded earlier.

Load and check data in Machine Learning Studio

Load the dataset

We will start by loading the **csdairydata.csv** file into Azure Machine Learning Studio.

- Start your Azure Machine Learning Studio environment.
- Click on **+ NEW** at the lower left of your screen and select **Dataset**.
- Select **From Local File**, and then **Browse** to select the file.
- Make sure you have selected **Generic CSV file with header (.csv)** as the type for the dataset.
- Click the check mark.
- After the dataset has been uploaded, you should see the new dataset by clicking on the **Datasets** tab.

Create an experiment

Now that we have some data in Machine Learning Studio, we need to create an experiment to do the analysis.

- Click on **+ NEW** at the lower left and select **Experiment**, then **Blank Experiment**.
- You can name your experiment by selecting, and modifying, the **Experiment created on ...** title at the top of the page. For example, changing it to **CA Dairy Analysis**.
- On the left of the experiment page, expand **Saved Datasets**, and then **My Datasets**. You should see the

cadairydata.csv that you uploaded earlier.

- Drag and drop the **csdairydata.csv dataset** onto the experiment.
- In the **Search experiment items** box on the top of the left pane, type **Execute R Script**. You will see the module appear in the search list.
- Drag and drop the **Execute R Script** module onto your pallet.
- Connect the output of the **csdairydata.csv dataset** to the leftmost input (**Dataset1**) of the **Execute R Script**.
- **Don't forget to click on 'Save'!**

At this point your experiment should look something like Figure 3.

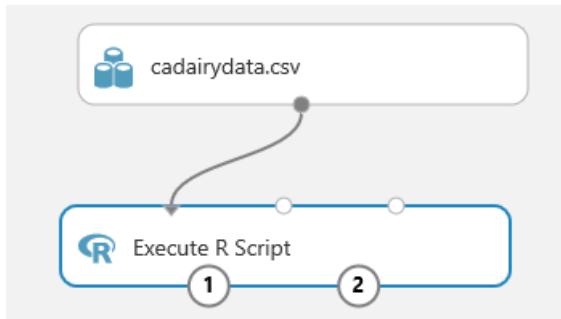


Figure 3. The CA Dairy Analysis experiment with dataset and Execute R Script module.

Check on the data

Let's have a look at the data we have loaded into our experiment. In the experiment, click on the output of the **cadairydata.csv dataset** and select **visualize**. You should see something like Figure 4.

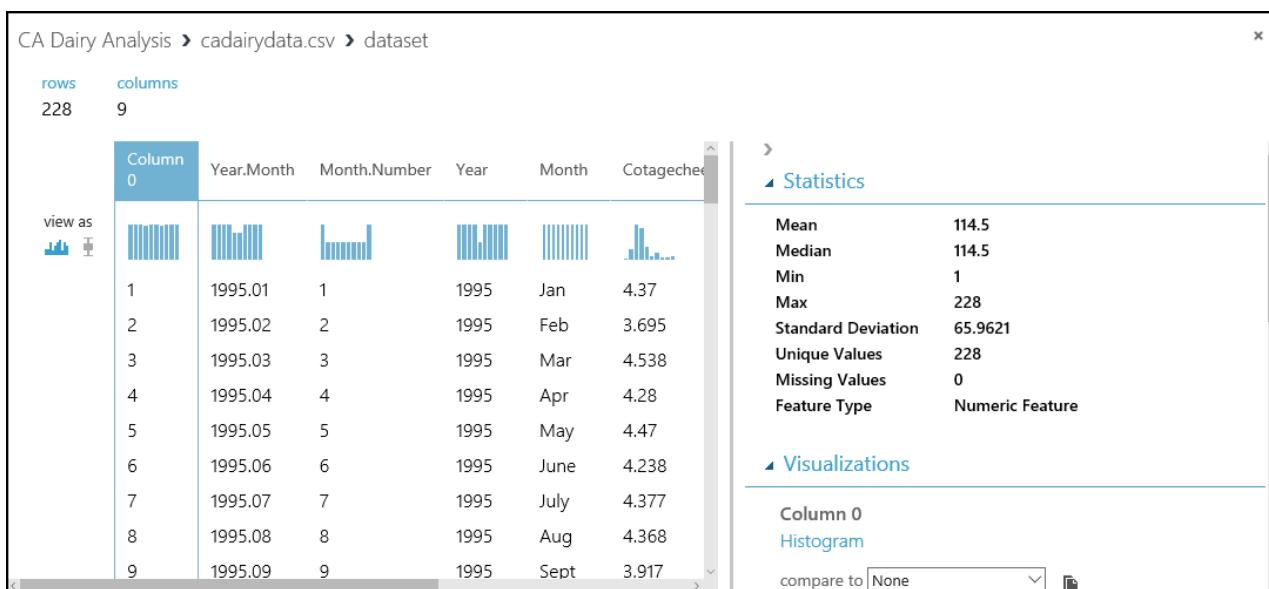


Figure 4. Summary of the *cadairydata.csv* dataset.

In this view we see a lot of useful information. We can see the first several rows of that dataset. If we select a column, the Statistics section shows more information about the column. For example, the Feature Type row shows us what data types Azure Machine Learning Studio assigned to the column. Having a quick look like this is a good sanity check before we start to do any serious work.

First R script

Let's create a simple first R script to experiment with in Azure Machine Learning Studio. I have created and tested the following script in RStudio.

```

## Only one of the following two lines should be used
## If running in Machine Learning Studio, use the first line with maml.mapInputPort()
## If in RStudio, use the second line with read.csv()
cadairydata <- maml.mapInputPort(1)
# cadairydata <- read.csv("cadairydata.csv", header = TRUE, stringsAsFactors = FALSE)
str(cadairydata)
pairs(~ Cottagecheese.Prod + Icecream.Prod + Milk.Prod + N.CA.Fat.Price, data = cadairydata)
## The following line should be executed only when running in
## Azure Machine Learning Studio
maml.mapOutputPort('cadairydata')

```

Now I need to transfer this script to Azure Machine Learning Studio. I could simply cut and paste. However, in this case, I will transfer my R script via a zip file.

Data input to the Execute R Script module

Let's have a look at the inputs to the [Execute R Script](#) module. In this example we will read the California dairy data into the [Execute R Script](#) module.

There are three possible inputs for the [Execute R Script](#) module. You may use any one or all of these inputs, depending on your application. It is also perfectly reasonable to use an R script that takes no input at all.

Let's look at each of these inputs, going from left to right. You can see the names of each of the inputs by placing your cursor over the input and reading the tooltip.

Script Bundle

The Script Bundle input allows you to pass the contents of a zip file into [Execute R Script](#) module. You can use one of the following commands to read the contents of the zip file into your R code.

```

source("src/yourfile.R") # Reads a zipped R script
load("src/yourData.rdata") # Reads a zipped R data file

```

NOTE

Azure Machine Learning treats files in the zip as if they are in the `src/` directory, so you need to prefix your file names with this directory name. For example, if the zip contains the files `yourfile.R` and `yourData.rdata` in the root of the zip, you would address these as `src/yourfile.R` and `src/yourData.rdata` when using `source` and `load`.

We already discussed loading datasets in [Loading the dataset](#). Once you have created and tested the R script shown in the previous section, do the following:

1. Save the R script into a .R file. I call my script file "simpleplot.R". Here's the contents.

```

## Only one of the following two lines should be used
## If running in Machine Learning Studio, use the first line with maml.mapInputPort()
## If in RStudio, use the second line with read.csv()
cadairydata <- maml.mapInputPort(1)
# cadairydata <- read.csv("cadairydata.csv", header = TRUE, stringsAsFactors = FALSE)
str(cadairydata)
pairs(~ Cottagecheese.Prod + Icecream.Prod + Milk.Prod + N.CA.Fat.Price, data = cadairydata)
## The following line should be executed only when running in
## Azure Machine Learning Studio
maml.mapOutputPort('cadairydata')

```

2. Create a zip file and copy your script into this zip file. On Windows, you can right-click on the file and select **Send to**, and then **Compressed folder**. This will create a new zip file containing the "simpleplot.R" file.
3. Add your file to the **datasets** in Machine Learning Studio, specifying the type as **zip**. You should now see the zip

file in your datasets.

4. Drag and drop the zip file from **datasets** onto the **ML Studio canvas**.
5. Connect the output of the **zip data** icon to the **Script Bundle** input of the **Execute R Script** module.
6. Type the `source()` function with your zip file name into the code window for the **Execute R Script** module. In my case I typed `source("src/simpleplot.R")`.
7. Make sure you click **Save**.

Once these steps are complete, the **Execute R Script** module will execute the R script in the zip file when the experiment is run. At this point your experiment should look something like Figure 5.

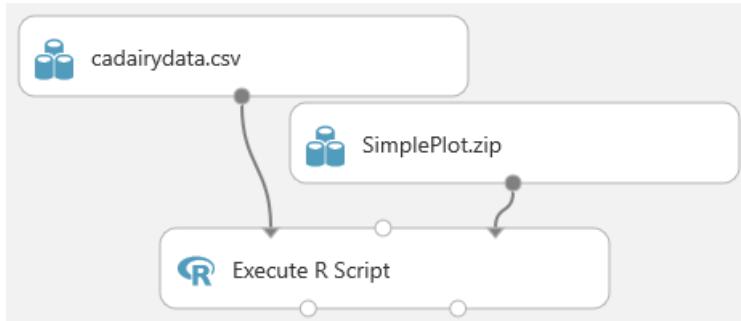


Figure 5. Experiment using zipped R script.

Dataset1

You can pass a rectangular table of data to your R code by using the **Dataset1** input. In our simple script the `maml.mapInputPort(1)` function reads the data from port 1. This data is then assigned to a dataframe variable name in your code. In our simple script the first line of code performs the assignment.

```
cadairydata <- maml.mapInputPort(1)
```

Execute your experiment by clicking on the **Run** button. When the execution finishes, click on the **Execute R Script** module and then click **View output log** on the properties pane. A new page should appear in your browser showing the contents of the `output.log` file. When you scroll down you should see something like the following.

```
[ModuleOutput] InputDataStructure
[ModuleOutput]
[ModuleOutput] {
[ModuleOutput] "InputName":Dataset1
[ModuleOutput] "Rows":228
[ModuleOutput] "Cols":9
[ModuleOutput] "ColumnTypes":System.Int32,3,System.Double,5,System.String,1
[ModuleOutput] }
```

Farther down the page is more detailed information on the columns, which will look something like the following.

```
[ModuleOutput] [1] "Loading variable port1..."  
[ModuleOutput]  
[ModuleOutput] 'data.frame': 228 obs. of 9 variables:  
[ModuleOutput]  
[ModuleOutput] $ Column 0      : int 1 2 3 4 5 6 7 8 9 10 ...  
[ModuleOutput]  
[ModuleOutput] $ Year.Month   : num 1995 1995 1995 1995 1995 ...  
[ModuleOutput]  
[ModuleOutput] $ Month.Number : int 1 2 3 4 5 6 7 8 9 10 ...  
[ModuleOutput]  
[ModuleOutput] $ Year       : int 1995 1995 1995 1995 1995 1995 1995 1995 1995 ...  
[ModuleOutput]  
[ModuleOutput] $ Month      : chr "Jan" "Feb" "Mar" "Apr" ...  
[ModuleOutput]  
[ModuleOutput] $ Cottagecheese.Prod: num 4.37 3.69 4.54 4.28 4.47 ...  
[ModuleOutput]  
[ModuleOutput] $ Icecream.Prod  : num 51.6 56.1 68.5 65.7 73.7 ...  
[ModuleOutput]  
[ModuleOutput] $ Milk.Prod    : num 2.11 1.93 2.16 2.13 2.23 ...  
[ModuleOutput]  
[ModuleOutput] $ N.CA.Fat.Price: num 0.98 0.892 0.892 0.897 0.897 ...
```

These results are mostly as expected, with 228 observations and 9 columns in the dataframe. We can see the column names, the R data type and a sample of each column.

NOTE

This same printed output is conveniently available from the R Device output of the [Execute R Script](#) module. We will discuss the outputs of the [Execute R Script](#) module in the next section.

Dataset2

The behavior of the Dataset2 input is identical to that of Dataset1. Using this input you can pass a second rectangular table of data into your R code. The function `maml.mapInputPort(2)`, with the argument 2, is used to pass this data.

Execute R Script outputs

Output a dataframe

You can output the contents of an R dataframe as a rectangular table through the Result Dataset1 port by using the `maml.mapOutputPort()` function. In our simple R script this is performed by the following line.

```
maml.mapOutputPort('cadairydata')
```

After running the experiment, click on the Result Dataset1 output port and then click on **Visualize**. You should see something like Figure 6.

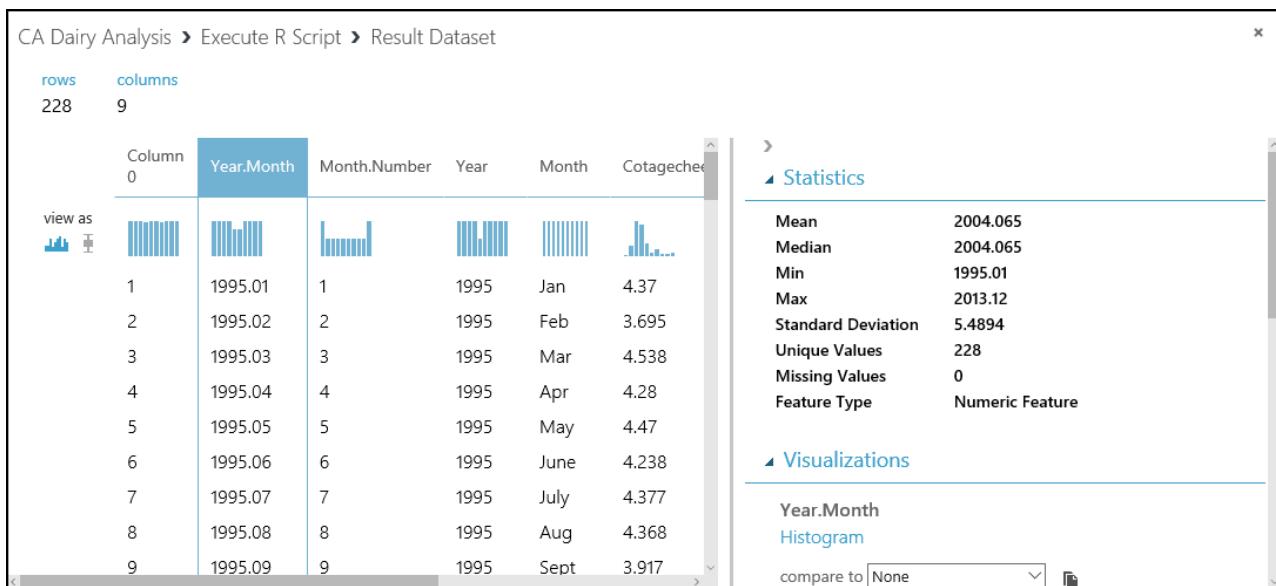


Figure 6. The visualization of the output of the California dairy data.

This output looks identical to the input, exactly as we expected.

R Device output

The Device output of the **Execute R Script** module contains messages and graphics output. Both standard output and standard error messages from R are sent to the R Device output port.

To view the R Device output, click on the port and then on **Visualize**. We see the standard output and standard error from the R script in Figure 7.

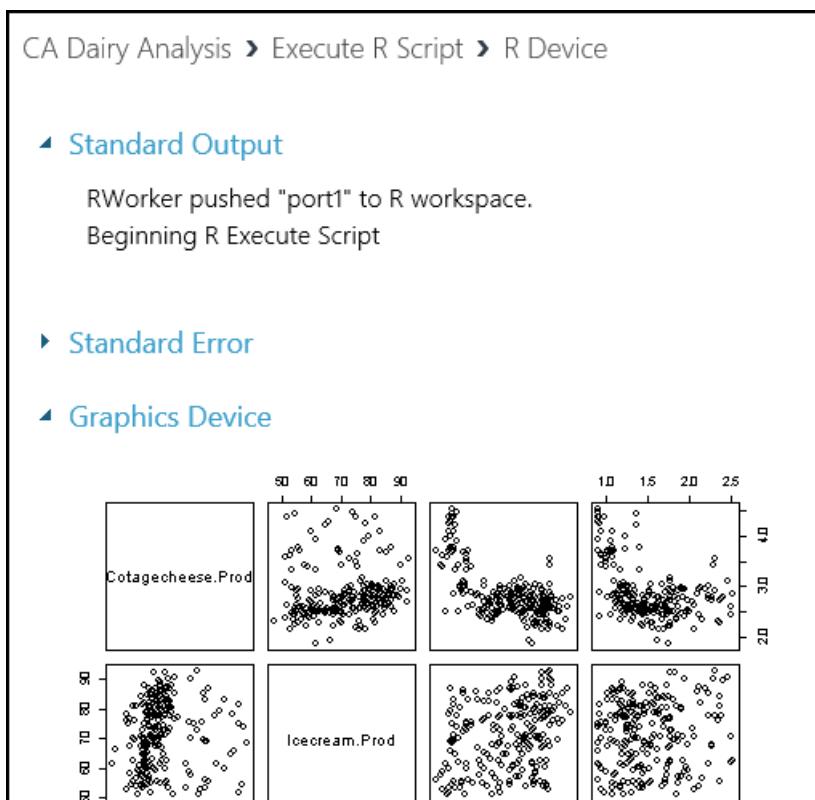


Figure 7. Standard output and standard error from the R Device port.

Scrolling down we see the graphics output from our R script in Figure 8.

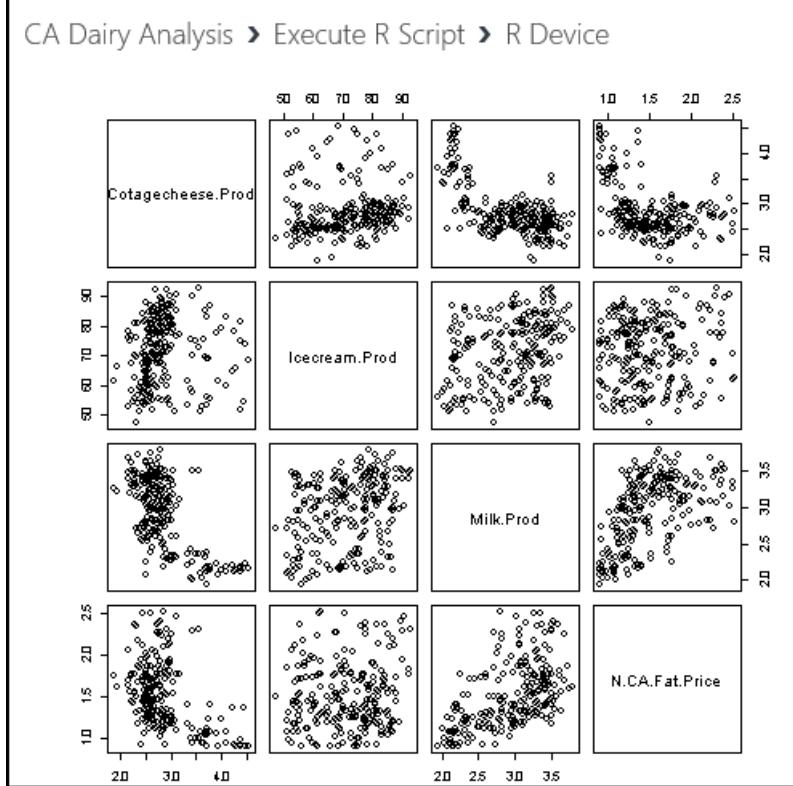


Figure 8. Graphics output from the R Device port.

Data filtering and transformation

In this section we will perform some basic data filtering and transformation operations on the California dairy data. By the end of this section we will have data in a format suitable for building an analytic model.

More specifically, in this section we will perform several common data cleaning and transformation tasks: type transformation, filtering on dataframes, adding new computed columns, and value transformations. This background should help you deal with the many variations encountered in real-world problems.

The complete R code for this section is available in the zip file you downloaded earlier.

Type transformations

Now that we can read the California dairy data into the R code in the [Execute R Script](#) module, we need to ensure that the data in the columns has the intended type and format.

R is a dynamically typed language, which means that data types are coerced from one to another as required. The atomic data types in R include numeric, logical and character. The factor type is used to compactly store categorical data. You can find much more information on data types in the references in [Appendix B - Further reading](#).

When tabular data is read into R from an external source, it is always a good idea to check the resulting types in the columns. You may want a column of type character, but in many cases this will show up as factor or vice versa. In other cases a column you think should be numeric is represented by character data, e.g. '1.23' rather than 1.23 as a floating point number.

Fortunately, it is easy to convert one type to another, as long as mapping is possible. For example, you cannot convert 'Nevada' into a numeric value, but you can convert it to a factor (categorical variable). As another example, you can convert a numeric 1 into a character '1' or a factor.

The syntax for any of these conversions is simple: `as.datatype()`. These type conversion functions include the following.

- `as.numeric()`
- `as.character()`

- `as.logical()`
- `as.factor()`

Looking at the data types of the columns we input in the previous section: all columns are of type numeric, except for the column labeled 'Month', which is of type character. Let's convert this to a factor and test the results.

I have deleted the line that created the scatterplot matrix and added a line converting the 'Month' column to a factor. In my experiment I will just cut and paste the R code into the code window of the [Execute R Script Module](#). You could also update the zip file and upload it to Azure Machine Learning Studio, but this takes several steps.

```
## Only one of the following two lines should be used
## If running in Machine Learning Studio, use the first line with maml.mapInputPort()
## If in RStudio, use the second line with read.csv()
cadairydata <- maml.mapInputPort(1)
# cadairydata <- read.csv("cadairydata.csv", header = TRUE, stringsAsFactors = FALSE)
## Ensure the coding is consistent and convert column to a factor
cadairydata$Month <- as.factor(cadairydata$Month)
str(cadairydata) # Check the result
## The following line should be executed only when running in
## Azure Machine Learning Studio
maml.mapOutputPort('cadairydata')
```

Let's execute this code and look at the output log for the R script. The relevant data from the log is shown in Figure 9.

```
[ModuleOutput] [1] "Loading variable port1..."
[ModuleOutput]
[ModuleOutput] 'data.frame': 228 obs. of 9 variables:
[ModuleOutput]
[ModuleOutput] $ Column 0      : int 1 2 3 4 5 6 7 8 9 10 ...
[ModuleOutput]
[ModuleOutput] $ Year.Month   : num 1995 1995 1995 1995 1995 ...
[ModuleOutput]
[ModuleOutput] $ Month.Number : int 1 2 3 4 5 6 7 8 9 10 ...
[ModuleOutput]
[ModuleOutput] $ Year        : int 1995 1995 1995 1995 1995 1995 1995 1995 ...
[ModuleOutput]
[ModuleOutput] $ Month       : Factor w/ 14 levels "Apr","April",..: 6 5 9 1 11 8 7 3 14 13 ...
[ModuleOutput]
[ModuleOutput] $ Cottagecheese.Prod: num 4.37 3.69 4.54 4.28 4.47 ...
[ModuleOutput]
[ModuleOutput] $ Icecream.Prod  : num 51.6 56.1 68.5 65.7 73.7 ...
[ModuleOutput]
[ModuleOutput] $ Milk.Prod     : num 2.11 1.93 2.16 2.13 2.23 ...
[ModuleOutput]
[ModuleOutput] $ N.CA.Fat.Price : num 0.98 0.892 0.892 0.897 0.897 ...
[ModuleOutput]
[ModuleOutput] [1] "Saving variable cadairydata ..."
[ModuleOutput]
[ModuleOutput] [1] "Saving the following item(s): .maml.oport1"
```

Figure 9. Summary of the dataframe with a factor variable.

The type for Month should now say '**Factor w/ 14 levels**'. This is a problem since there are only 12 months in the year. You can also check to see that the type in **Visualize** of the Result Dataset port is '**Categorical**'.

The problem is that the 'Month' column has not been coded systematically. In some cases a month is called April and in others it is abbreviated as Apr. We can solve this problem by trimming the string to 3 characters. The line of code now looks like the following:

```
## Ensure the coding is consistent and convert column to a factor  
cadairydata$Month <- as.factor(substr(cadairydata$Month, 1, 3))
```

Rerun the experiment and view the output log. The expected results are shown in Figure 10.

```
[ModuleOutput] [1] "Loading variable port1..."  
[ModuleOutput]  
[ModuleOutput] 'data.frame': 228 obs. of 9 variables:  
[ModuleOutput]  
[ModuleOutput] $ Column 0 : int 1 2 3 4 5 6 7 8 9 10 ...  
[ModuleOutput]  
[ModuleOutput] $ Year.Month : num 1995 1995 1995 1995 1995 ...  
[ModuleOutput]  
[ModuleOutput] $ Month.Number : int 1 2 3 4 5 6 7 8 9 10 ...  
[ModuleOutput]  
[ModuleOutput] $ Year : int 1995 1995 1995 1995 1995 1995 1995 1995 1995 ...  
[ModuleOutput]  
[ModuleOutput] $ Month : Factor w/ 12 levels "Apr","Aug","Dec",...: 5 4 8 1 9 7 6 2 12 11 ...  
[ModuleOutput]  
[ModuleOutput] $ Cottagecheese.Prod: num 4.37 3.69 4.54 4.28 4.47 ...  
[ModuleOutput]  
[ModuleOutput] $ Icecream.Prod : num 51.6 56.1 68.5 65.7 73.7 ...  
[ModuleOutput]  
[ModuleOutput] $ Milk.Prod : num 2.11 1.93 2.16 2.13 2.23 ...  
[ModuleOutput]  
[ModuleOutput] $ N.CA.Fat.Price : num 0.98 0.892 0.892 0.897 0.897 ...  
[ModuleOutput]  
[ModuleOutput] [1] "Saving variable cadairydata ..."  
[ModuleOutput]  
[ModuleOutput] [1] "Saving the following item(s): .maml.lopert1"
```

Figure 10. Summary of the dataframe with correct number of factor levels.

Our factor variable now has the desired 12 levels.

Basic data frame filtering

R dataframes support powerful filtering capabilities. Datasets can be subsetted by using logical filters on either rows or columns. In many cases, complex filter criteria will be required. The references in [Appendix B - Further reading](#) contain extensive examples of filtering dataframes.

There is one bit of filtering we should do on our dataset. If you look at the columns in the cadairydata dataframe, you will see two unnecessary columns. The first column just holds a row number, which is not very useful. The second column, Year.Month, contains redundant information. We can easily exclude these columns by using the following R code.

NOTE

From now on in this section, I will just show you the additional code I am adding in the [Execute R Script](#) module. I will add each new line **before** the `str()` function. I use this function to verify my results in Azure Machine Learning Studio.

I add the following line to my R code in the [Execute R Script](#) module.

```
# Remove two columns we do not need  
cadairydata <- cadairydata[, c(-1, -2)]
```

Run this code in your experiment and check the result from the output log. These results are shown in Figure 11.

```
[ModuleOutput] [1] "Loading variable port1..."  

[ModuleOutput]  

[ModuleOutput] 'data.frame': 228 obs. of 7 variables:  

[ModuleOutput]  

[ModuleOutput] $ Month.Number : int 1 2 3 4 5 6 7 8 9 10 ...  

[ModuleOutput]  

[ModuleOutput] $ Year : int 1995 1995 1995 1995 1995 1995 1995 1995 1995 ...  

[ModuleOutput]  

[ModuleOutput] $ Month : Factor w/ 12 levels "Apr","Aug","Dec",...: 5 4 8 1 9 7 6 2 12 11 ...  

[ModuleOutput]  

[ModuleOutput] $ Cotagecheese.Prod: num 4.37 3.69 4.54 4.28 4.47 ...  

[ModuleOutput]  

[ModuleOutput] $ Icecream.Prod : num 51.6 56.1 68.5 65.7 73.7 ...  

[ModuleOutput]  

[ModuleOutput] $ Milk.Prod : num 2.11 1.93 2.16 2.13 2.23 ...  

[ModuleOutput]  

[ModuleOutput] $ N.CA.Fat.Price : num 0.98 0.892 0.892 0.897 0.897 ...  

[ModuleOutput]  

[ModuleOutput] [1] "Saving variable cadairydata ..."  

[ModuleOutput]  

[ModuleOutput] [1] "Saving the following item(s): .maml.lopert1"
```

Figure 11. Summary of the dataframe with two columns removed.

Good news! We get the expected results.

Add a new column

To create time series models it will be convenient to have a column containing the months since the start of the time series. We will create a new column 'Month.Count'.

To help organize the code we will create our first simple function, `num.month()`. We will then apply this function to create a new column in the dataframe. The new code is as follows.

```
## Create a new column with the month count  

## Function to find the number of months from the first  

## month of the time series  

num.month <- function(Year, Month) {  

  ## Find the starting year  

  min.year <- min(Year)  

  ## Compute the number of months from the start of the time series  

  12 * (Year - min.year) + Month - 1  

}  

## Compute the new column for the dataframe  

cadairydata$Month.Count <- num.month(cadairydata$Year, cadairydata$Month.Number)
```

Now run the updated experiment and use the output log to view the results. These results are shown in Figure 12.

```
[ModuleOutput] [1] "Loading variable port1..."  

[ModuleOutput]  

[ModuleOutput] 'data.frame': 228 obs. of 8 variables:  

[ModuleOutput]  

[ModuleOutput] $ Month.Number : int 1 2 3 4 5 6 7 8 9 10 ...  

[ModuleOutput]  

[ModuleOutput] $ Year : int 1995 1995 1995 1995 1995 1995 1995 1995 1995 ...  

[ModuleOutput]  

[ModuleOutput] $ Month : Factor w/ 12 levels "Apr","Aug","Dec",...: 5 4 8 1 9 7 6 2 12 11 ...  

[ModuleOutput]  

[ModuleOutput] $ Cottagecheese.Prod: num 4.37 3.69 4.54 4.28 4.47 ...  

[ModuleOutput]  

[ModuleOutput] $ Icecream.Prod : num 51.6 56.1 68.5 65.7 73.7 ...  

[ModuleOutput]  

[ModuleOutput] $ Milk.Prod : num 2.11 1.93 2.16 2.13 2.23 ...  

[ModuleOutput]  

[ModuleOutput] $ N.CA.Fat.Price : num 0.98 0.892 0.892 0.897 0.897 ...  

[ModuleOutput]  

[ModuleOutput] $ Month.Count : num 0 1 2 3 4 5 6 7 8 9 ...  

[ModuleOutput]  

[ModuleOutput] [1] "Saving variable cadata ..."  

[ModuleOutput]  

[ModuleOutput] [1] "Saving the following item(s): .maml.oport!"
```

Figure 12. Summary of the dataframe with the additional column.

It looks like everything is working. We have the new column with the expected values in our dataframe.

Value transformations

In this section we will perform some simple transformations on the values in some of the columns of our dataframe. The R language supports nearly arbitrary value transformations. The references in [Appendix B - Further Reading](#) contain extensive examples.

If you look at the values in the summaries of our dataframe you should see something odd here. Is more ice cream than milk produced in California? No, of course not, as this makes no sense, sad as this fact may be to some of us ice cream lovers. The units are different. The price is in units of US pounds, milk is in units of 1 M US pounds, ice cream is in units of 1,000 US gallons, and cottage cheese is in units of 1,000 US pounds. Assuming ice cream weighs about 6.5 pounds per gallon, we can easily do the multiplication to convert these values so they are all in equal units of 1,000 pounds.

For our forecasting model we use a multiplicative model for trend and seasonal adjustment of this data. A log transformation allows us to use a linear model, simplifying this process. We can apply the log transformation in the same function where the multiplier is applied.

In the following code, I define a new function, `log.transform()`, and apply it to the rows containing the numerical values. The R `Map()` function is used to apply the `log.transform()` function to the selected columns of the dataframe. `Map()` is similar to `apply()` but allows for more than one list of arguments to the function. Note that a list of multipliers supplies the second argument to the `log.transform()` function. The `na.omit()` function is used as a bit of cleanup to ensure we do not have missing or undefined values in the dataframe.

```

log.transform<- function(invec, multiplier = 1) {
  ## Function for the transformation, which is the log
  ## of the input value times a multiplier

  warningmessages <- c("ERROR: Non-numeric argument encountered in function log.transform",
    "ERROR: Arguments to function log.transform must be greater than zero",
    "ERROR: Argument multiplier to function log.transform must be a scalar",
    "ERROR: Invalid time series value encountered in function log.transform"
  )

  ## Check the input arguments
  if(!is.numeric(invec) | !is.numeric(multiplier)) {warning(warningmessages[1]); return(NA)}
  if(any(invec < 0.0) | any(multiplier < 0.0)) {warning(warningmessages[2]); return(NA)}
  if(length(multiplier) != 1) {warning(warningmessages[3]); return(NA)}

  ## Wrap the multiplication in tryCatch
  ## If there is an exception, print the warning message to
  ## standard error and return NA
  tryCatch(log(multiplier * invec),
    error = function(e){warning(warningmessages[4]); NA})
}

## Apply the transformation function to the 4 columns
## of the dataframe with production data
multipliers <- list(1.0, 6.5, 1000.0, 1000.0)
cadairydata[, 4:7] <- Map(log.transform, cadairydata[, 4:7], multipliers)

## Get rid of any rows with NA values
cadairydata <- na.omit(cadairydata)

```

There is quite a bit happening in the `log.transform()` function. Most of this code is checking for potential problems with the arguments or dealing with exceptions, which can still arise during the computations. Only a few lines of this code actually do the computations.

The goal of the defensive programming is to prevent the failure of a single function that prevents processing from continuing. An abrupt failure of a long-running analysis can be quite frustrating for users. To avoid this situation, default return values must be chosen that will limit damage to downstream processing. A message is also produced to alert users that something has gone wrong.

If you are not used to defensive programming in R, all this code may seem a bit overwhelming. I will walk you through the major steps:

1. A vector of four messages is defined. These messages are used to communicate information about some of the possible errors and exceptions that can occur with this code.
2. I return a value of NA for each case. There are many other possibilities that might have fewer side effects. I could return a vector of zeroes, or the original input vector, for example.
3. Checks are run on the arguments to the function. In each case, if an error is detected, a default value is returned and a message is produced by the `warning()` function. I am using `warning()` rather than `stop()` as the latter will terminate execution, exactly what I am trying to avoid. Note that I have written this code in a procedural style, as in this case a functional approach seemed complex and obscure.
4. The log computations are wrapped in `tryCatch()` so that exceptions will not cause an abrupt halt to processing. Without `tryCatch()` most errors raised by R functions result in a stop signal, which does just that.

Execute this R code in your experiment and have a look at the printed output in the `output.log` file. You will now see the transformed values of the four columns in the log, as shown in Figure 13.

```
[ModuleOutput] [1] "Loading variable port1..."  

[ModuleOutput]  

[ModuleOutput] 'data.frame': 228 obs. of 8 variables:  

[ModuleOutput]  

[ModuleOutput] $ Month.Number : int 1 2 3 4 5 6 7 8 9 10 ...  

[ModuleOutput]  

[ModuleOutput] $ Year : int 1995 1995 1995 1995 1995 1995 1995 1995 1995 ...  

[ModuleOutput]  

[ModuleOutput] $ Month : Factor w/ 12 levels "Apr","Aug","Dec",...: 5 4 8 1 9 7 6 2 12 11 ...  

[ModuleOutput]  

[ModuleOutput] $ Cottagecheese.Prod: num 1.47 1.31 1.51 1.45 1.5 ...  

[ModuleOutput]  

[ModuleOutput] $ Icecream.Prod : num 5.82 5.9 6.1 6.06 6.17 ...  

[ModuleOutput]  

[ModuleOutput] $ Milk.Prod : num 7.66 7.57 7.68 7.66 7.71 ...  

[ModuleOutput]  

[ModuleOutput] $ N.CA.Fat.Price : num 6.89 6.79 6.79 6.8 6.8 ...  

[ModuleOutput]  

[ModuleOutput] $ Month.Count : num 0 1 2 3 4 5 6 7 8 9 ...  

[ModuleOutput]  

[ModuleOutput] [1] "Saving variable cadata ..."  

[ModuleOutput]  

[ModuleOutput] [1] "Saving the following item(s): .maml.oport!"
```

Figure 13. Summary of the transformed values in the dataframe.

We see the values have been transformed. Milk production now greatly exceeds all other dairy product production, recalling that we are now looking at a log scale.

At this point our data is cleaned up and we are ready for some modeling. Looking at the visualization summary for the Result Dataset output of our [Execute R Script](#) module, you will see the 'Month' column is 'Categorical' with 12 unique values, again, just as we wanted.

Time series objects and correlation analysis

In this section we will explore a few basic R time series objects and analyze the correlations between some of the variables. Our goal is to output a dataframe containing the pairwise correlation information at several lags.

The complete R code for this section is in the zip file you downloaded earlier.

Time series objects in R

As already mentioned, time series are a series of data values indexed by time. R time series objects are used to create and manage the time index. There are several advantages to using time series objects. Time series objects free you from the many details of managing the time series index values that are encapsulated in the object. In addition, time series objects allow you to use the many time series methods for plotting, printing, modeling, etc.

The POSIXct time series class is commonly used and is relatively simple. This time series class measures time from the start of the epoch, January 1, 1970. We will use POSIXct time series objects in this example. Other widely used R time series object classes include zoo and xts, extensible time series.

Time series object example

Let's get started with our example. Drag and drop a [new Execute R Script](#) module into your experiment. Connect the Result Dataset1 output port of the existing [Execute R Script](#) module to the Dataset1 input port of the new [Execute R Script](#) module.

As I did for the first examples, as we progress through the example, at some points I will show only the incremental additional lines of R code at each step.

Reading the dataframe

As a first step, let's read in a dataframe and make sure we get the expected results. The following code should do

the job.

```
# Comment the following if using RStudio
cadairydata <- maml.mapInputPort(1)
str(cadairydata) # Check the results
```

Now, run the experiment. The log of the new Execute R Script shape should look like Figure 14.

```
[ModuleOutput] [1] "Loading variable port1..."
[ModuleOutput]
[ModuleOutput] 'data.frame': 228 obs. of 8 variables:
[ModuleOutput]
[ModuleOutput] $ Month.Number   : int 1 2 3 4 5 6 7 8 9 10 ...
[ModuleOutput]
[ModuleOutput] $ Year        : int 1995 1995 1995 1995 1995 1995 1995 1995 1995 ...
[ModuleOutput]
[ModuleOutput] $ Month       : Factor w/ 12 levels "Apr","Aug","Dec",...: 5 4 8 1 9 7 6 2 12 11 ...
[ModuleOutput]
[ModuleOutput] $ Cottagecheese.Prod: num 1.47 1.31 1.51 1.45 1.5 ...
[ModuleOutput]
[ModuleOutput] $ Icecream.Prod  : num 5.82 5.9 6.1 6.06 6.17 ...
[ModuleOutput]
[ModuleOutput] $ Milk.Prod    : num 7.66 7.57 7.68 7.66 7.71 ...
[ModuleOutput]
[ModuleOutput] $ N.CA.Fat.Price: num 6.89 6.79 6.79 6.8 6.8 ...
[ModuleOutput]
[ModuleOutput] $ Month.Count  : num 0 1 2 3 4 5 6 7 8 9 ...
```

Figure 14. Summary of the dataframe in the Execute R Script module.

This data is of the expected types and format. Note that the 'Month' column is of type factor and has the expected number of levels.

Creating a time series object

We need to add a time series object to our dataframe. Replace the current code with the following, which adds a new column of class POSIXct.

```
# Comment the following if using RStudio
cadairydata <- maml.mapInputPort(1)

## Create a new column as a POSIXct object
Sys.setenv(TZ = "PST8PDT")
cadairydata$Time <- as.POSIXct(strptime(paste(as.character(cadairydata$Year), "-", as.character(cadairydata$Month.Number), "-01 00:00:00", sep = ""), "%Y-%m-%d %H:%M:%S"))

str(cadairydata) # Check the results
```

Now, check the log. It should look like Figure 15.

```
[ModuleOutput] [1] "Loading variable port1..."  

[ModuleOutput]  

[ModuleOutput] 'data.frame': 228 obs. of 9 variables:  

[ModuleOutput]  

[ModuleOutput] $ Month.Number : int 1 2 3 4 5 6 7 8 9 10 ...  

[ModuleOutput]  

[ModuleOutput] $ Year : int 1995 1995 1995 1995 1995 1995 1995 1995 1995 1995 ...  

[ModuleOutput]  

[ModuleOutput] $ Month : Factor w/ 12 levels "Apr","Aug","Dec",...: 5 4 8 1 9 7 6 2 12 11 ...  

[ModuleOutput]  

[ModuleOutput] $ Cottagecheese.Prod: num 1.47 1.31 1.51 1.45 1.5 ...  

[ModuleOutput]  

[ModuleOutput] $ Icecream.Prod : num 5.82 5.9 6.1 6.06 6.17 ...  

[ModuleOutput]  

[ModuleOutput] $ Milk.Prod : num 7.66 7.57 7.68 7.66 7.71 ...  

[ModuleOutput]  

[ModuleOutput] $ N.CA.Fat.Price : num 6.89 6.79 6.79 6.8 6.8 ...  

[ModuleOutput]  

[ModuleOutput] $ Month.Count : num 0 1 2 3 4 5 6 7 8 9 ...  

[ModuleOutput]  

[ModuleOutput] $ Time : POSIXct, format: "1995-01-01" "1995-02-01" ...
```

Figure 15. Summary of the dataframe with a time series object.

We can see from the summary that the new column is in fact of class `POSIXct`.

Exploring and transforming the data

Let's explore some of the variables in this dataset. A scatterplot matrix is a good way to produce a quick look. I am replacing the `str()` function in the previous R code with the following line.

```
pairs(~ Cottagecheese.Prod + Icecream.Prod + Milk.Prod + N.CA.Fat.Price, data = cadairydata, main = "Pairwise Scatterplots of dairy time series")
```

Run this code and see what happens. The plot produced at the R Device port should look like Figure 16.

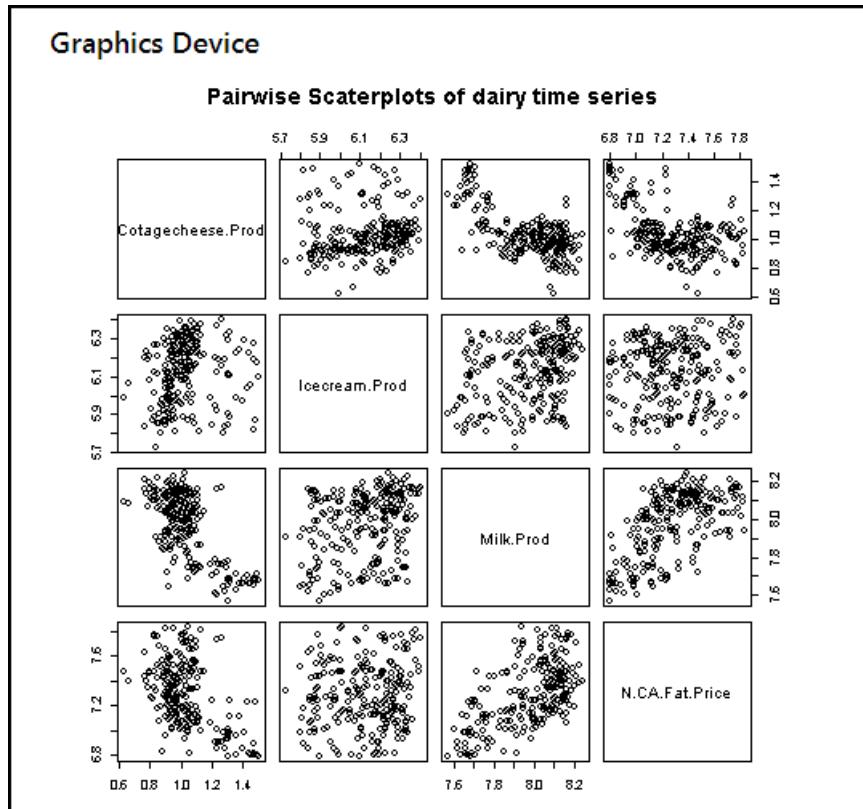


Figure 16. Scatterplot matrix of selected variables.

There is some odd-looking structure in the relationships between these variables. Perhaps this arises from trends in the data and from the fact that we have not standardized the variables.

Correlation analysis

To perform correlation analysis we need to both de-trend and standardize the variables. We could simply use the R `scale()` function, which both centers and scales variables. This function might well run faster. However, I want to show you an example of defensive programming in R.

The `ts.detrend()` function shown below performs both of these operations. The following two lines of code de-trend the data and then standardize the values.

```
ts.detrend <- function(ts, Time, min.length = 3){  
  ## Function to de-trend and standardize a time series  
  
  ## Define some messages if they are NULL  
  messages <- c('ERROR: ts.detrend requires arguments ts and Time to have the same length',  
            'ERROR: ts.detrend requires argument ts to be of type numeric',  
            paste('WARNING: ts.detrend has encountered a time series with length less than', as.character(min.length)),  
            'ERROR: ts.detrend has encountered a Time argument not of class POSIXct',  
            'ERROR: Detrend regression has failed in ts.detrend',  
            'ERROR: Exception occurred in ts.detrend while standardizing time series in function ts.detrend'  
  )  
  # Create a vector of zeros to return as a default in some cases  
  zerovec <- rep(length(ts), 0.0)  
  
  # The input arguments are not of the same length, return ts and quit  
  if(length(Time) != length(ts)) {warning(messages[1]); return(ts)}  
  
  # If the ts is not numeric, just return a zero vector and quit  
  if(!is.numeric(ts)) {warning(messages[2]); return(zerovec)}  
  
  # If the ts is too short, just return it and quit  
  if((ts.length <- length(ts)) < min.length) {warning(messages[3]); return(ts)}  
  
  ## Check that the Time variable is of class POSIXct  
  if(class(cadairydata$Time)[[1]] != "POSIXct") {warning(messages[4]); return(ts)}  
  
  ## De-trend the time series by using a linear model  
  ts.frame <- data.frame(ts = ts, Time = Time)  
  tryCatch({ts <- ts - fitted(lm(ts ~ Time, data = ts.frame))},  
          error = function(e){warning(messages[5]); zerovec})  
  
  tryCatch( {stdev <- sqrt(sum((ts - mean(ts))^2)/(ts.length - 1))  
            ts <- ts/stdev},  
          error = function(e){warning(messages[6]); zerovec})  
  
  ts  
}  
## Apply the detrend.ts function to the variables of interest  
df.detrend <- data.frame(lapply(cadairydata[, 4:7], ts.detrend, cadairydata$Time))  
  
## Plot the results to look at the relationships  
pairs(~ Cottagecheese.Prod + Icecream.Prod + Milk.Prod + N.CA.Fat.Price, data = df.detrend, main = "Pairwise Scatterplots of detrended  
standardized time series")
```

There is quite a bit happening in the `ts.detrend()` function. Most of this code is checking for potential problems with the arguments or dealing with exceptions, which can still arise during the computations. Only a few lines of this code actually do the computations.

We have already discussed an example of defensive programming in [Value transformations](#). Both computation blocks are wrapped in `tryCatch()`. For some errors it makes sense to return the original input vector, and in other cases, I return a vector of zeros.

Note that the linear regression used for de-trending is a time series regression. The predictor variable is a time series object.

Once `ts.detrend()` is defined we apply it to the variables of interest in our dataframe. We must coerce the resulting list created by `lapply()` to data datafarme by using `as.data.frame()`. Because of defensive aspects of `ts.detrend()`, failure to process one of the variables will not prevent correct processing of the others.

The final line of code creates a pairwise scatterplot. After running the R code, the results of the scatterplot are shown in Figure 17.

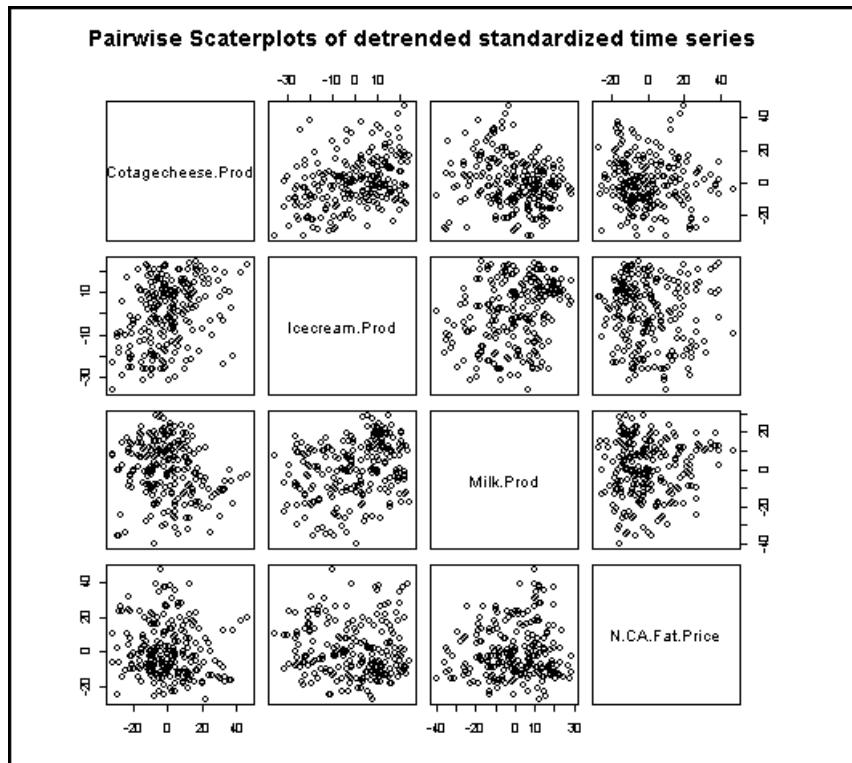


Figure 17. Pairwise scatterplot of de-trended and standardized time series.

You can compare these results to those shown in Figure 16. With the trend removed and the variables standardized, we see a lot less structure in the relationships between these variables.

The code to compute the correlations as R ccf objects is as follows.

```
## A function to compute pairwise correlations from a
## list of time series value vectors
pair.cor <- function(pair.ind, ts.list, lag.max = 1, plot = FALSE){
  ccf(ts.list[[pair.ind[1]]], ts.list[[pair.ind[2]]], lag.max = lag.max, plot = plot)
}

## A list of the pairwise indices
corpairs <- list(c(1,2), c(1,3), c(1,4), c(2,3), c(2,4), c(3,4))

## Compute the list of ccf objects
cadairycorrelations <- lapply(corpairs, pair.cor, df.detrend)

cadairycorrelations
```

Running this code produces the log shown in Figure 18.

```

[ModuleOutput] Loading objects:
[ModuleOutput] port1
[ModuleOutput] [1] "Loading variable port1..."
[ModuleOutput] [[1]]
[ModuleOutput]
[ModuleOutput]
[ModuleOutput] Autocorrelations of series 'X', by lag
[ModuleOutput]
[ModuleOutput]
[ModuleOutput] -1 0 1
[ModuleOutput] 0.148 0.358 0.317
[ModuleOutput]
[ModuleOutput]
[ModuleOutput] [[2]]
[ModuleOutput]
[ModuleOutput]
[ModuleOutput] Autocorrelations of series 'X', by lag
[ModuleOutput]
[ModuleOutput]
[ModuleOutput]
[ModuleOutput] -1 0 1
[ModuleOutput] -0.395 -0.186 -0.238
[ModuleOutput]
[ModuleOutput]
[ModuleOutput] [[3]]
[ModuleOutput]
[ModuleOutput]
[ModuleOutput]
[ModuleOutput] Autocorrelations of series 'X', by lag
[ModuleOutput]
[ModuleOutput]
[ModuleOutput]
[ModuleOutput] -1 0 1
[ModuleOutput] -0.059 -0.089 -0.127
[ModuleOutput]
[ModuleOutput]
[ModuleOutput]
[ModuleOutput] [[4]]
[ModuleOutput]
[ModuleOutput]
[ModuleOutput]
[ModuleOutput] Autocorrelations of series 'X', by lag
[ModuleOutput]
[ModuleOutput]
[ModuleOutput]
[ModuleOutput] -1 0 1
[ModuleOutput] 0.140 0.294 0.293
[ModuleOutput]
[ModuleOutput]
[ModuleOutput] [[5]]
[ModuleOutput]
[ModuleOutput]
[ModuleOutput]
[ModuleOutput] Autocorrelations of series 'X', by lag
[ModuleOutput]
[ModuleOutput]
[ModuleOutput]
[ModuleOutput] -1 0 1
[ModuleOutput] -0.002 -0.074 -0.124

```

Figure 18. List of ccf objects from the pairwise correlation analysis.

There is a correlation value for each lag. None of these correlation values is large enough to be significant. We can therefore conclude that we can model each variable independently.

Output a dataframe

We have computed the pairwise correlations as a list of R ccf objects. This presents a bit of a problem as the Result Dataset output port really requires a dataframe. Further, the ccf object is itself a list and we want only the values in the first element of this list, the correlations at the various lags.

The following code extracts the lag values from the list of ccf objects, which are themselves lists.

```

df.correlations <- data.frame(do.call(rbind, lapply(cadairycorrelations, "[[", 1)))

c.names <- c("correlation pair", "-1 lag", "0 lag", "+1 lag")
r.names <- c("Corr Cot Cheese - Ice Cream",
            "Corr Cot Cheese - Milk Prod",
            "Corr Cot Cheese - Fat Price",
            "Corr Ice Cream - Mik Prod",
            "Corr Ice Cream - Fat Price",
            "Corr Milk Prod - Fat Price")

## Build a dataframe with the row names column and the
## correlation data frame and assign the column names
outframe <- cbind(r.names, df.correlations)
colnames(outframe) <- c.names
outframe

## WARNING!
## The following line works only in Azure Machine Learning
## When running in RStudio, this code will result in an error
#maml.mapOutputPort('outframe')

```

The first line of code is a bit tricky, and some explanation may help you understand it. Working from the inside out we have the following:

1. The '`[[1]`' operator with the argument '`1`' selects the vector of correlations at the lags from the first element of the `ccf` object list.
2. The `do.call()` function applies the `rbind()` function over the elements of the list returns by `lapply()`.
3. The `data.frame()` function coerces the result produced by `do.call()` to a dataframe.

Note that the row names are in a column of the dataframe. Doing so preserves the row names when they are output from the [Execute R Script](#).

Running the code produces the output shown in Figure 19 when I **Visualize** the output at the Result Dataset port. The row names are in the first column, as intended.

CA Dairy Analysis > Execute R Script > Result Dataset				
rows	columns			
6	4			
view as	correlation pair	-1 lag	0 lag	+1 lag
	Corr Cot Cheese - Ice Cream	0.147775	0.358106	0.317
	Corr Cot Cheese - Milk Prod	-0.395444	-0.185777	-0.238486
	Corr Cot Cheese - Fat Price	-0.059431	-0.08864	-0.127372
	Corr Ice Cream - Mik Prod	0.139825	0.294231	0.293429
	Corr Ice Cream - Fat Price	-0.001878	-0.07351	-0.124443
	Corr Milk Prod - Fat Price	0.032936	0.075419	0.052024

Figure 19. Results output from the correlation analysis.

Time series example: seasonal forecasting

Our data is now in a form suitable for analysis, and we have determined there are no significant correlations between the variables. Let's move on and create a time series forecasting model. Using this model we will forecast California milk production for the 12 months of 2013.

Our forecasting model will have two components, a trend component and a seasonal component. The complete forecast is the product of these two components. This type of model is known as a multiplicative model. The alternative is an additive model. We have already applied a log transformation to the variables of interest, which makes this analysis tractable.

The complete R code for this section is in the zip file you downloaded earlier.

Creating the dataframe for analysis

Start by adding a new **Execute R Script** module to your experiment. Connect the **Result Dataset** output of the existing **Execute R Script** module to the **Dataset1** input of the new module. The result should look something like Figure 20.

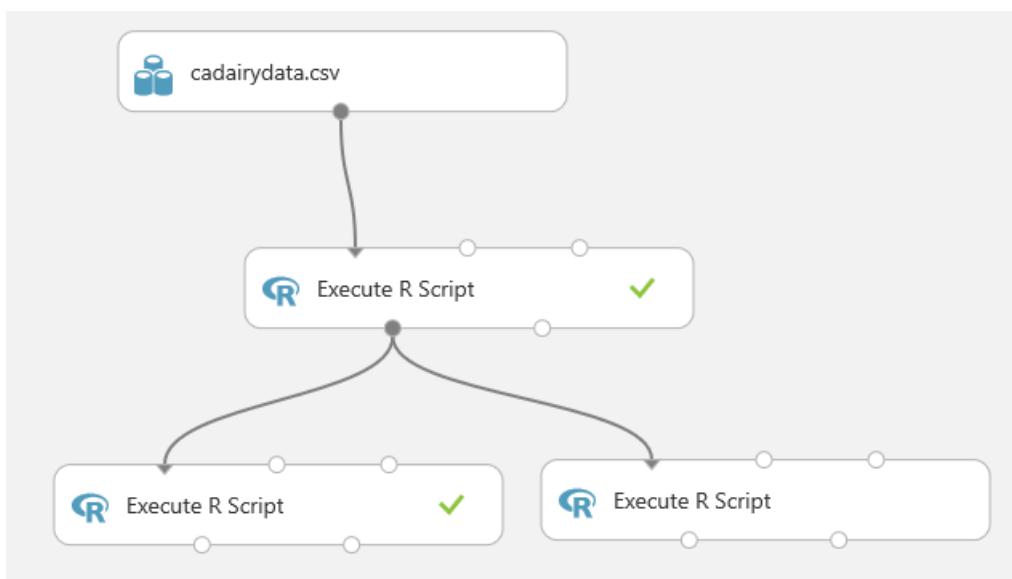


Figure 20. The experiment with the new Execute R Script module added.

As with the correlation analysis we just completed, we need to add a column with a POSIXct time series object. The following code will do just this.

```
# If running in Machine Learning Studio, uncomment the first line with maml.mapInputPort()
cadairydata <- maml.mapInputPort(1)

## Create a new column as a POSIXct object
Sys.setenv(TZ = "PST8PDT")
cadairydata$Time <- as.POSIXct(strptime(paste(as.character(cadairydata$Year), "-", as.character(cadairydata$Month.Number), "-01 00:00:00", sep = ""), "%Y-%m-%d %H:%M:%S"))

str(cadairydata)
```

Run this code and look at the log. The result should look like Figure 21.

```
[ModuleOutput] [1] "Loading variable port1..."  

[ModuleOutput]  

[ModuleOutput] 'data.frame': 228 obs. of 9 variables:  

[ModuleOutput]  

[ModuleOutput] $ Month.Number : int 1 2 3 4 5 6 7 8 9 10 ...  

[ModuleOutput]  

[ModuleOutput] $ Year : int 1995 1995 1995 1995 1995 1995 1995 1995 1995 ...  

[ModuleOutput]  

[ModuleOutput] $ Month : Factor w/ 12 levels "Apr","Aug","Dec",...: 5 4 8 1 9 7 6 2 12 11 ...  

[ModuleOutput]  

[ModuleOutput] $ Cotagecheese.Prod: num 1.47 1.31 1.51 1.45 1.5 ...  

[ModuleOutput]  

[ModuleOutput] $ Icecream.Prod : num 5.82 5.9 6.1 6.06 6.17 ...  

[ModuleOutput]  

[ModuleOutput] $ Milk.Prod : num 7.66 7.57 7.68 7.66 7.71 ...  

[ModuleOutput]  

[ModuleOutput] $ N.CA.Fat.Price : num 6.89 6.79 6.79 6.8 6.8 ...  

[ModuleOutput]  

[ModuleOutput] $ Month.Count : num 0 1 2 3 4 5 6 7 8 9 ...  

[ModuleOutput]  

[ModuleOutput] $ Time : POSIXct, format: "1995-01-01" "1995-02-01" ...
```

Figure 21. A summary of the dataframe.

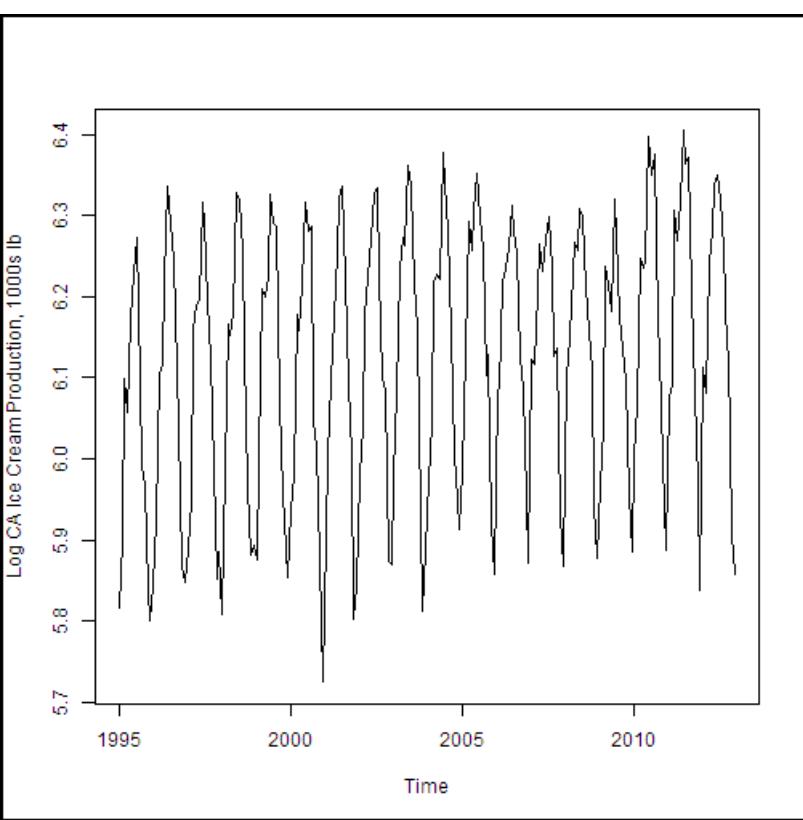
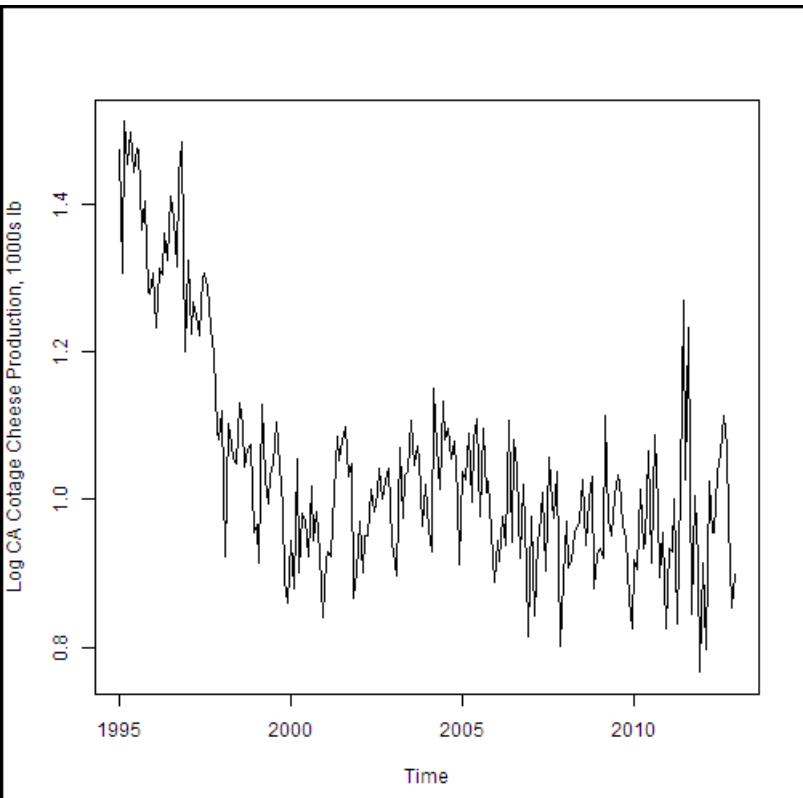
With this result, we are ready to start our analysis.

Create a training dataset

With the dataframe constructed we need to create a training dataset. This data will include all of the observations except the last 12, of the year 2013, which is our test dataset. The following code subsets the dataframe and creates plots of the dairy production and price variables. I then create plots of the four production and price variables. An anonymous function is used to define some augment for plot, and then iterate over the list of the other two arguments with `Map()`. If you are thinking that a for loop would have worked fine here, you are correct. But, since R is a functional language I am showing you a functional approach.

```
cadairytrain <- cadairydata[1:216, ]  
  
Ylabs <- list("Log CA Cotage Cheese Production, 1000s lb",  
             "Log CA Ice Cream Production, 1000s lb",  
             "Log CA Milk Production 1000s lb",  
             "Log North CA Milk Milk Fat Price per 1000 lb")  
  
Map(function(y, Ylabs){plot(cadairytrain$Time, y, xlab = "Time", ylab = Ylabs, type = "l")}, cadairytrain[, 4:7], Ylabs)
```

Running the code produces the series of time series plots from the R Device output shown in Figure 22. Note that the time axis is in units of dates, a nice benefit of the time series plot method.



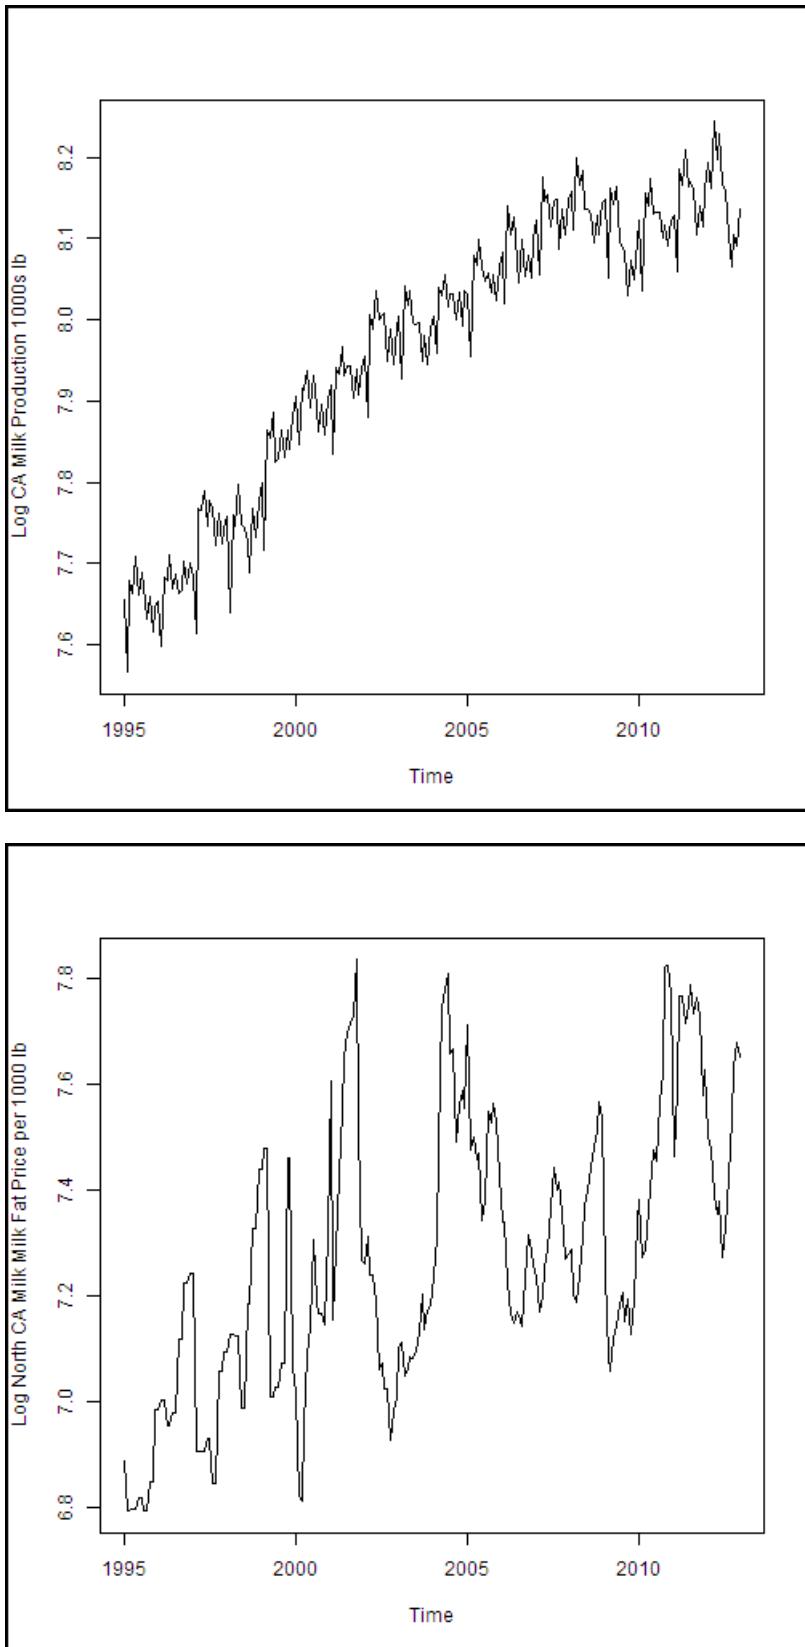


Figure 22. Time series plots of California dairy production and price data.

A trend model

Having created a time series object and having had a look at the data, let's start to construct a trend model for the California milk production data. We can do this with a time series regression. However, it is clear from the plot that we will need more than a slope and intercept to accurately model the observed trend in the training data.

Given the small scale of the data, I will build the model for trend in RStudio and then cut and paste the resulting model into Azure Machine Learning. RStudio provides an interactive environment for this type of interactive analysis.

As a first attempt, I will try a polynomial regression with powers up to 3. There is a real danger of over-fitting these kinds of models. Therefore, it is best to avoid high order terms. The `I()` function inhibits interpretation of the contents (interprets the contents 'as is') and allows you to write a literally interpreted function in a regression equation.

```
milk.lm<- lm(Milk.Prod ~ Time + I(Month.Count^2) + I(Month.Count^3), data = cadairytrain)
summary(milk.lm)
```

This generates the following.

```
##
## Call:
## lm(formula = Milk.Prod ~ Time + I(Month.Count^2) + I(Month.Count^3),
##   data = cadairytrain)
##
## Residuals:
##   Min     1Q Median     3Q    Max 
## -0.12667 -0.02730  0.00236  0.02943  0.10586
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 6.33e+00  1.45e-01  43.60 <2e-16 ***
## Time        1.63e-09 1.72e-10  9.47 <2e-16 ***
## I(Month.Count^2) -1.71e-06 4.89e-06 -0.35  0.726  
## I(Month.Count^3) -3.24e-08 1.49e-08 -2.17  0.031 *  
## ---            
## Signif. codes:  0 '****' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 '' 1
##
## Residual standard error: 0.0418 on 212 degrees of freedom
## Multiple R-squared:  0.941, Adjusted R-squared:  0.94 
## F-statistic: 1.12e+03 on 3 and 212 DF, p-value: <2e-16
```

From P values ($\text{Pr}(>|t|)$) in this output, we can see that the squared term may not be significant. I will use the `update()` function to modify this model by dropping the squared term.

```
milk.lm<- update(milk.lm, . ~ . - I(Month.Count^2))
summary(milk.lm)
```

This generates the following.

```
##
## Call:
## lm(formula = Milk.Prod ~ Time + I(Month.Count^3), data = cadairytrain)
##
## Residuals:
##   Min     1Q Median     3Q    Max 
## -0.12597 -0.02659  0.00185  0.02963  0.10696
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 6.38e+00  4.07e-02 156.6 <2e-16 ***
## Time        1.57e-09 4.32e-11  36.3 <2e-16 ***
## I(Month.Count^3) -3.76e-08 2.50e-09 -15.1 <2e-16 ***
## ---            
## Signif. codes:  0 '****' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 '' 1
##
## Residual standard error: 0.0417 on 213 degrees of freedom
## Multiple R-squared:  0.941, Adjusted R-squared:  0.94 
## F-statistic: 1.69e+03 on 2 and 213 DF, p-value: <2e-16
```

This looks better. All of the terms are significant. However, the 2e-16 value is a default value, and should not be taken too seriously.

As a sanity test, let's make a time series plot of the California dairy production data with the trend curve shown. I have added the following code in the Azure Machine Learning [Execute R Script](#) model (not RStudio) to create the model and make a plot. The result is shown in Figure 23.

```
milk.lm<- lm(Milk.Prod ~ Time + I(Month.Count^3), data = cadairytrain)  
  
plot(cadairytrain$Time, cadairytrain$Milk.Prod, xlab = "Time", ylab = "Log CA Milk Production 1000s lb", type = "l")  
lines(cadairytrain$Time, predict(milk.lm, cadairytrain), lty = 2, col = 2)
```

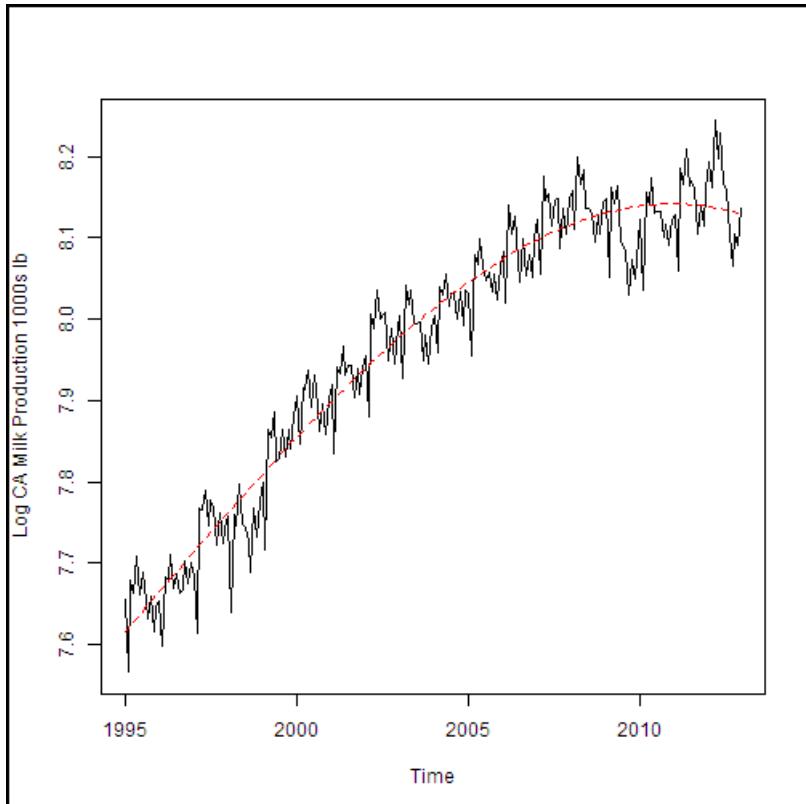


Figure 23. California milk production data with trend model shown.

It looks like the trend model fits the data fairly well. Further, there does not seem to be evidence of over-fitting, such as odd wiggles in the model curve.

Seasonal model

With a trend model in hand, we need to push on and include the seasonal effects. We will use the month of the year as a dummy variable in the linear model to capture the month-by-month effect. Note that when you introduce factor variables into a model, the intercept must not be computed. If you do not do this, the formula is over-specified and R will drop one of the desired factors but keep the intercept term.

Since we have a satisfactory trend model we can use the `update()` function to add the new terms to the existing model. The `-1` in the update formula drops the intercept term. Continuing in RStudio for the moment:

```
milk.lm2<- update(milk.lm, . ~ . + Month - 1)  
summary(milk.lm2)
```

This generates the following.

```

## 
## Call:
## lm(formula = Milk.Prod ~ Time + I(Month.Count^3) + Month - 1,
##   data = cadatairtrain)
##
## Residuals:
##    Min      1Q  Median      3Q     Max 
## -0.06879 -0.01693  0.00346  0.01543  0.08726 
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## Time        1.57e-09  2.72e-11  57.7 <2e-16 ***
## I(Month.Count^3) -3.74e-08  1.57e-09 -23.8 <2e-16 ***
## MonthApr    6.40e+00  2.63e-02  243.3 <2e-16 ***
## MonthAug    6.38e+00  2.63e-02  242.2 <2e-16 ***
## MonthDec    6.38e+00  2.64e-02  241.9 <2e-16 ***
## MonthFeb    6.31e+00  2.63e-02  240.1 <2e-16 *** 
## MonthJan    6.39e+00  2.63e-02  243.1 <2e-16 *** 
## MonthJul    6.39e+00  2.63e-02  242.6 <2e-16 *** 
## MonthJun    6.38e+00  2.63e-02  242.4 <2e-16 *** 
## MonthMar    6.42e+00  2.63e-02  244.2 <2e-16 *** 
## MonthMay    6.43e+00  2.63e-02  244.3 <2e-16 *** 
## MonthNov    6.34e+00  2.63e-02  240.6 <2e-16 *** 
## MonthOct    6.37e+00  2.63e-02  241.8 <2e-16 *** 
## MonthSep    6.34e+00  2.63e-02  240.6 <2e-16 *** 
## --- 
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1 
## 
## Residual standard error: 0.0263 on 202 degrees of freedom
## Multiple R-squared:  1, Adjusted R-squared:  1 
## F-statistic: 1.42e+06 on 14 and 202 DF, p-value: <2e-16

```

We see that the model no longer has an intercept term and has 12 significant month factors. This is exactly what we wanted to see.

Let's make another time series plot of the California dairy production data to see how well the seasonal model is working. I have added the following code in the Azure Machine Learning [Execute R Script](#) to create the model and make a plot.

```

milk.lm2<- lm(Milk.Prod ~ Time + I(Month.Count^3) + Month - 1, data = cadatairtrain)

plot(cadatairtrain$Time, cadatairtrain$Milk.Prod, xlab = "Time", ylab = "Log CA Milk Production 1000s lb", type = "l")
lines(cadatairtrain$Time, predict(milk.lm2, cadatairtrain), lty = 2, col = 2)

```

Running this code in Azure Machine Learning produces the plot shown in Figure 24.

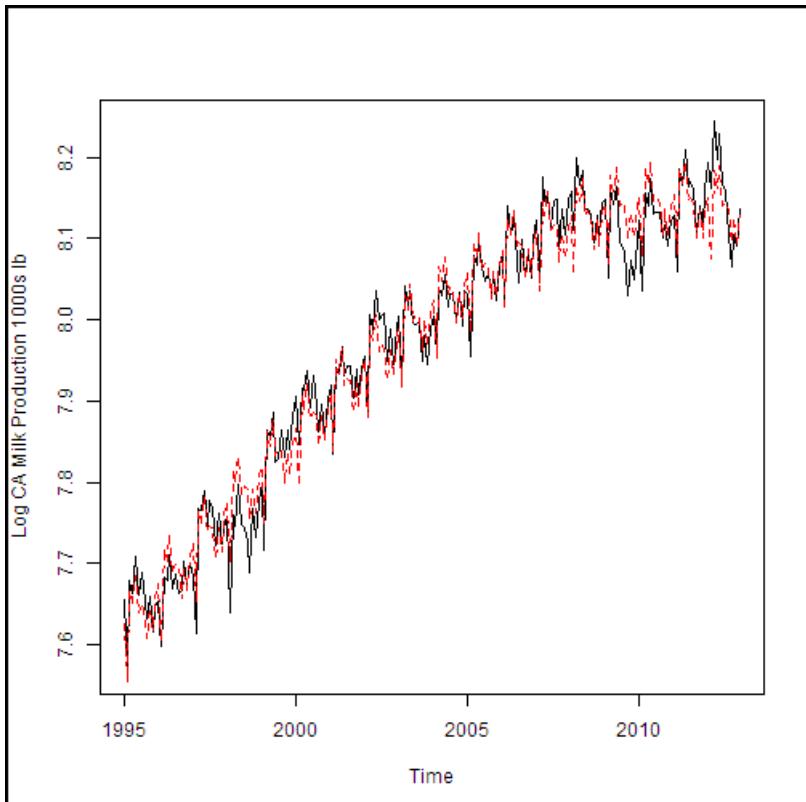


Figure 24. California milk production with model including seasonal effects.

The fit to the data shown in Figure 24 is rather encouraging. Both the trend and the seasonal effect (monthly variation) look reasonable.

As another check on our model, let's have a look at the residuals. The following code computes the predicted values from our two models, computes the residuals for the seasonal model, and then plots these residuals for the training data.

```

## Compute predictions from our models
predict1 <- predict(milk.lm, cadairydata)
predict2 <- predict(milk.lm2, cadairydata)

## Compute and plot the residuals
residuals <- cadairydata$Milk.Prod - predict2
plot(cadairytrain$Time, residuals[1:216], xlab = "Time", ylab ="Residuals of Seasonal Model")

```

The residual plot is shown in Figure 25.

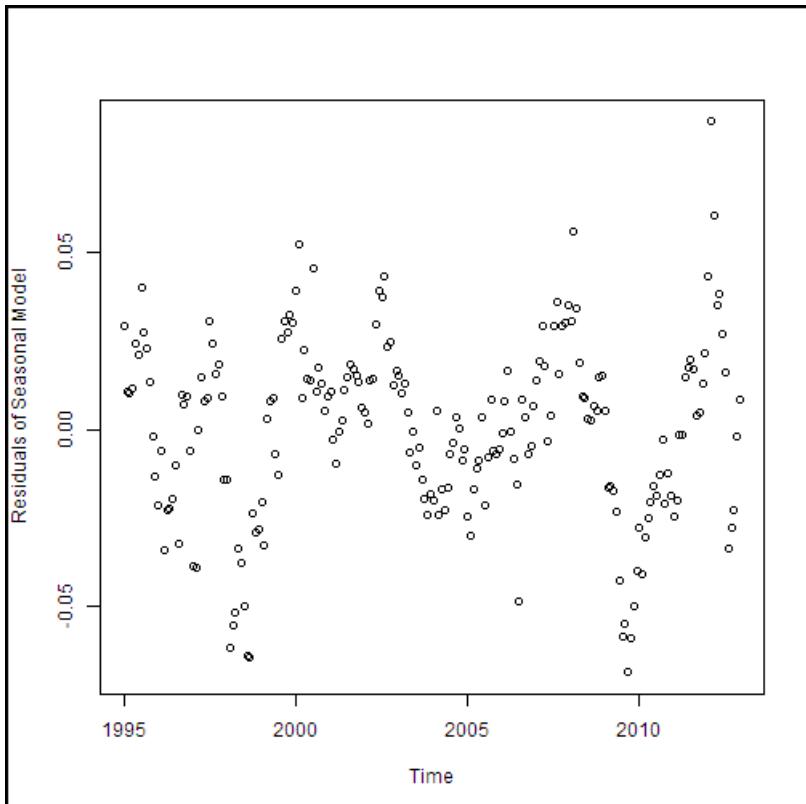


Figure 25. Residuals of the seasonal model for the training data.

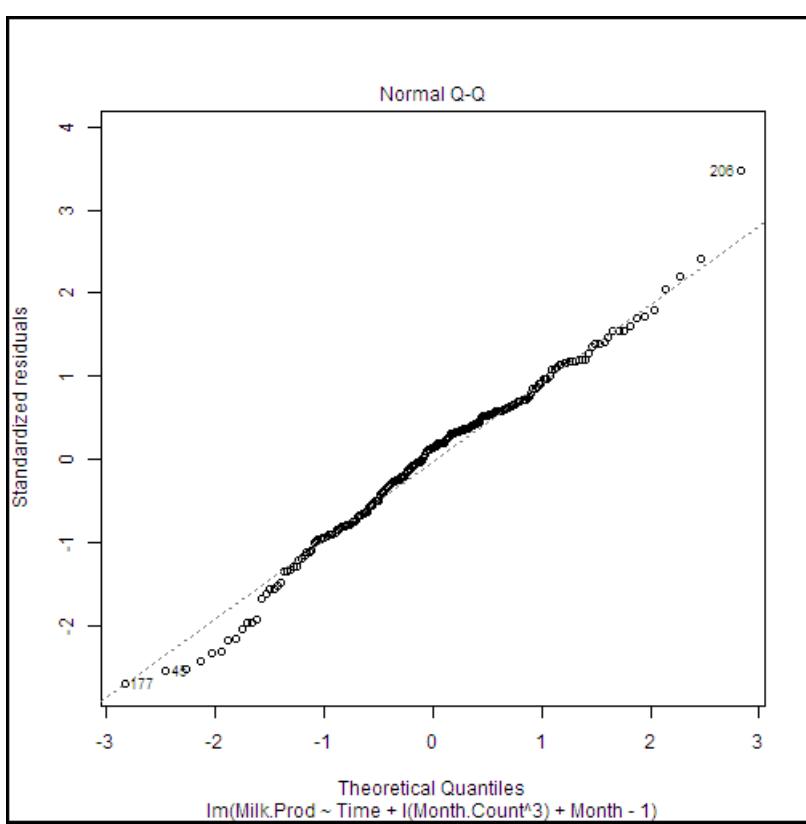
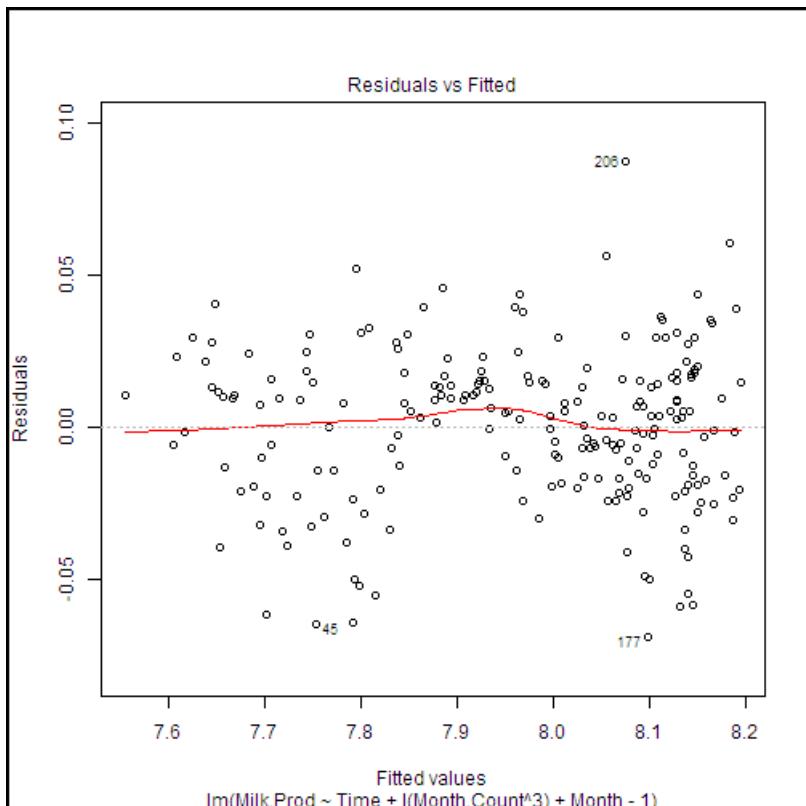
These residuals look reasonable. There is no particular structure, except the effect of the 2008-2009 recession, which our model does not account for particularly well.

The plot shown in Figure 25 is useful for detecting any time-dependent patterns in the residuals. The explicit approach of computing and plotting the residuals I used places the residuals in time order on the plot. If, on the other hand, I had plotted `milk.lm$residuals`, the plot would not have been in time order.

You can also use `plot.lm()` to produce a series of diagnostic plots.

```
## Show the diagnostic plots for the model
plot(milk.lm2, ask = FALSE)
```

This code produces a series of diagnostic plots shown in Figure 26.



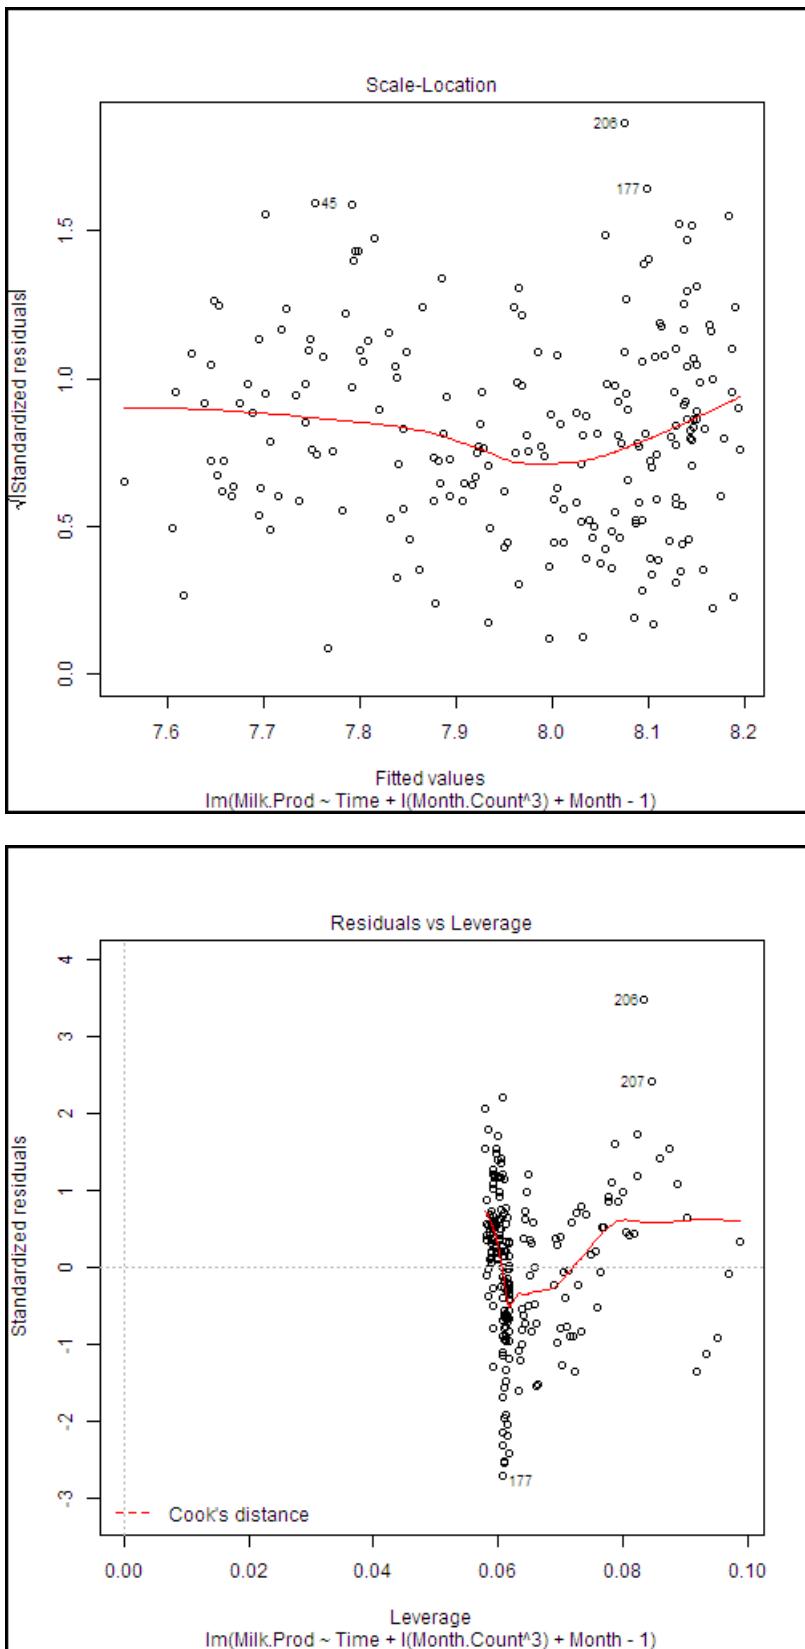


Figure 26. Diagnostic plots for the seasonal model.

There are a few highly influential points identified in these plots, but nothing to cause great concern. Further, we can see from the Normal Q-Q plot that the residuals are close to normally distributed, an important assumption for linear models.

Forecasting and model evaluation

There is just one more thing to do to complete our example. We need to compute forecasts and measure the error against the actual data. Our forecast will be for the 12 months of 2013. We can compute an error measure for this forecast to the actual data that is not part of our training dataset. Additionally, we can compare performance on the 18 years of training data to the 12 months of test data.

A number of metrics are used to measure the performance of time series models. In our case we will use the root mean square (RMS) error. The following function computes the RMS error between two series.

```
RMS.error <- function(series1, series2, is.log = TRUE, min.length = 2){  
  ## Function to compute the RMS error or difference between two  
  ## series or vectors  
  
  messages <- c("ERROR: Input arguments to function RMS.error of wrong type encountered",  
             "ERROR: Input vector to function RMS.error is too short",  
             "ERROR: Input vectors to function RMS.error must be of same length",  
             "WARNING: Function rms.error has received invalid input time series.")  
  
  ## Check the arguments  
  if(!is.numeric(series1) | !is.numeric(series2) | !is.logical(is.log) | !is.numeric(min.length)) {  
    warning(messages[1])  
    return(NA)}  
  
  if(length(series1) < min.length) {  
    warning(messages[2])  
    return(NA)}  
  
  if((length(series1) != length(series2))) {  
    warning(messages[3])  
    return(NA)}  
  
  ## If is.log is TRUE exponentiate the values, else just copy  
  if(is.log) {  
    tryCatch( {  
      temp1 <- exp(series1)  
      temp2 <- exp(series2) },  
      error = function(e){warning(messages[4]); NA}  
    )  
  } else {  
    temp1 <- series1  
    temp2 <- series2  
  }  
  
  ## Compute predictions from our models  
  predict1 <- predict(milk.lm, cadairydata)  
  predict2 <- predict(milk.lm2, cadairydata)  
  
  ## Compute the RMS error in a dataframe  
  tryCatch( {  
    sqrt(sum((temp1 - temp2)^2) / length(temp1))),  
    error = function(e){warning(messages[4]); NA}  
  })
```

As with the `log.transform()` function we discussed in the "Value transformations" section, there is quite a lot of error checking and exception recovery code in this function. The principles employed are the same. The work is done in two places wrapped in `tryCatch()`. First, the time series are exponentiated, since we have been working with the logs of the values. Second, the actual RMS error is computed.

Equipped with a function to measure the RMS error, let's build and output a dataframe containing the RMS errors. We will include terms for the trend model alone and the complete model with seasonal factors. The following code does the job by using the two linear models we have constructed.

```

## Compute the RMS error in a dataframe
## Include the row names in the first column so they will
## appear in the output of the Execute R Script
RMS.df <- data.frame(
  rowNames = c("Trend Model", "Seasonal Model"),
  Traing = c(
    RMS.error(predict1[1:216], cadairydata$Milk.Prod[1:216]),
    RMS.error(predict2[1:216], cadairydata$Milk.Prod[1:216])),
  Forecast = c(
    RMS.error(predict1[217:228], cadairydata$Milk.Prod[217:228]),
    RMS.error(predict2[217:228], cadairydata$Milk.Prod[217:228]))
)
RMS.df

## The following line should be executed only when running in
## Azure Machine Learning Studio
maml.mapOutputPort("RMS.df")

```

Running this code produces the output shown in Figure 27 at the Result Dataset output port.

CA Dairy Analysis > Execute R Script > Result Dataset			
rows	columns		
2	3		
	rowNames	Traing	Forecast
view as			
	Trend Model	122.238008	174.088492
	Seasonal Model	75.115958	94.725325

Figure 27. Comparison of RMS errors for the models.

From these results, we see that adding the seasonal factors to the model reduces the RMS error significantly. Not too surprisingly, the RMS error for the training data is a bit less than for the forecast.

APPENDIX A: Guide to RStudio

RStudio is quite well documented, so in this appendix I will provide some links to the key sections of the RStudio documentation to get you started.

1. Creating projects

You can organize and manage your R code into projects by using RStudio. The documentation that uses projects can be found at <https://support.rstudio.com/hc/articles/200526207-Using-Projects>.

I recommend that you follow these directions and create a project for the R code examples in this document.

2. Editing and executing R code

RStudio provides an integrated environment for editing and executing R code. Documentation can be found at <https://support.rstudio.com/hc/articles/200484448-Editing-and-Executing-Code>.

3. Debugging

RStudio includes powerful debugging capabilities. Documentation for these features is at <https://support.rstudio.com/hc/articles/200713843-Debugging-with-RStudio>.

The breakpoint troubleshooting features are documented at
<https://support.rstudio.com/hc/articles/200534337-Breakpoint-Troubleshooting>.

APPENDIX B: Further reading

This R programming tutorial covers the basics of what you need to use the R language with Azure Machine Learning Studio. If you are not familiar with R, two introductions are available on CRAN:

- R for Beginners by Emmanuel Paradis is a good place to start at http://cran.r-project.org/doc/contrib/Paradis-rdebuts_en.pdf.
- An Introduction to R by W. N. Venables et. al. goes into a bit more depth, at <http://cran.r-project.org/doc/manuals/R-intro.html>.

There are many books on R that can help you get started. Here are a few I find useful:

- The Art of R Programming: A Tour of Statistical Software Design by Norman Matloff is an excellent introduction to programming in R.
- R Cookbook by Paul Teator provides a problem and solution approach to using R.
- R in Action by Robert Kabacoff is another useful introductory book. The companion Quick R website is a useful resource at <http://www.statmethods.net/>.
- R Inferno by Patrick Burns is a surprisingly humorous book that deals with a number of tricky and difficult topics that can be encountered when programming in R. The book is available for free at <http://www.burns-stat.com/documents/books/the-r-inferno/>.
- If you want a deep dive into advanced topics in R, have a look at the book Advanced R by Hadley Wickham. The online version of this book is available for free at <http://adv-r.had.co.nz/>.

A catalogue of R time series packages can be found in the CRAN Task View for time series analysis: <http://cran.r-project.org/web/views/TimeSeries.html>. For information on specific time series object packages, you should refer to the documentation for that package.

The book Introductory Time Series with R by Paul Cowpertwait and Andrew Metcalfe provides an introduction to using R for time series analysis. Many more theoretical texts provide R examples.

Some great internet resources:

- DataCamp: DataCamp teaches R in the comfort of your browser with video lessons and coding exercises. There are interactive tutorials on the latest R techniques and packages. Take the free interactive R tutorial at <https://www.datacamp.com/courses/introduction-to-r>
- A guide on Getting started with R from Programiz <https://www.programiz.com/r-programming>
- A quick R tutorial by Kelly Black from Clarkson University <http://www.cyclismo.org/tutorial/R/>
- 60+ R resources listed at <http://www.computerworld.com/article/2497464/business-intelligence-60-r-resources-to-improve-your-data-skills.html>

Create and share an Azure Machine Learning workspace

12/7/2017 • 2 min to read • [Edit Online](#)

This menu links to topics that describe how to set up the various data science environments used by the Cortana Analytics Process (CAPS).

To use Azure Machine Learning Studio, you need to have a Machine Learning workspace. This workspace contains the tools you need to create, manage, and publish experiments.

NOTE

Try Azure Machine Learning for free

No credit card or Azure subscription required. [Get started now.](#)

To create a workspace

1. Sign in to the [Azure portal](#)

NOTE

To sign in and create a workspace, you need to be an Azure subscription administrator.

2. Click **+New**
3. In the search box, type **Machine Learning Studio Workspace** and select the matching item. Then, select click **Create** at the bottom of the page.
4. Enter your workspace information:
 - The *workspace name* may be up to 260 characters, not ending in a space. The name can't include these characters: <>* % & :\ ? +/
 - The *web service plan* you choose (or create), along with the associated *pricing tier* you select, is used if you deploy web services from this workspace.

Machine Learning Workspace □ X

Machine Learning Workspace

* Workspace name
My-workspace ✓

* Subscription
My-subscription

* Resource group ⓘ
 Create new Use existing
My-resource-group ✓

* Location
South Central US

* Storage account ⓘ
 Create new Use existing
storageformyworkspace ✓

Workspace pricing tier ⓘ
Standard

* Web service plan ⓘ
 Create new Use existing
My-web-service-plan ✓

* Web service plan pricing tier ⓘ >
S1 Standard

5. Click **Create**.

Once the workspace is deployed, you can open it in Machine Learning Studio.

1. Browse to Machine Learning Studio at <https://studio.azureml.net/>.
2. Select your workspace in the upper-right-hand corner.



3. Click **my experiments**.

Welcome back luisa!

MY RECENT WORKSPACES:

My-workspace

MY RECENT EXPERIMENTS:

my experiments →

For information about managing your workspace, see [Manage an Azure Machine Learning workspace](#). If you encounter a problem creating your workspace, see [Troubleshooting guide: Create and connect to a Machine Learning workspace](#).

Sharing an Azure Machine Learning workspace

Once a Machine Learning workspace is created, you can invite users to your workspace to share access to your workspace and all its experiments, datasets, notebooks, etc. You can add users in one of two roles:

- **User** - A workspace user can create, open, modify, and delete experiments, datasets, etc. in the workspace.
- **Owner** - An owner can invite and remove users in the workspace, in addition to what a user can do.

NOTE

The administrator account that creates the workspace is automatically added to the workspace as workspace Owner. However, other administrators or users in that subscription are not automatically granted access to the workspace - you need to invite them explicitly.

To share a workspace

1. Sign in to Machine Learning Studio at <https://studio.azureml.net/Home>
2. In the left panel, click **SETTINGS**
3. Click the **USERS** tab
4. Click **INVITE MORE USERS** at the bottom of the page

NAME	AUTHORIZATION TOKENS	USERS	DATA GATEWAYS
NAME	EMAIL	ROLE	STATUS
luisa	luisa@contoso.com	Owner	Active

5. Enter one or more email addresses. The users need a valid Microsoft account or an organizational account (from Azure Active Directory).
6. Select whether you want to add the users as Owner or User.
7. Click the **OK** checkmark button.

Each user you add will receive an email with instructions on how to sign in to the shared workspace.

NOTE

For users to be able to deploy or manage web services in this workspace, they must be a contributor or administrator in the Azure subscription.

Manage an Azure Machine Learning workspace

12/14/2017 • 1 min to read • [Edit Online](#)

NOTE

For information on managing Web services in the Machine Learning Web Services portal, see [Manage a Web service using the Azure Machine Learning Web Services portal](#).

You can manage Machine Learning workspaces in the Azure portal.

NOTE

Try Azure Machine Learning for free

No credit card or Azure subscription required. [Get started now](#).

Use the Azure portal

To manage a workspace in the Azure portal:

1. Sign in to the [Azure portal](#) using an Azure subscription administrator account.
2. In the search box at the top of the page, enter "machine learning workspaces" and then select **Machine Learning Workspaces**.
3. Click the workspace you want to manage.

In addition to the standard resource management information and options available, you can:

- View **Properties** - This page displays the workspace and resource information, and you can change the subscription and resource group that this workspace is connected with.
- **Resync Storage Keys** - The workspace maintains keys to the storage account. If the storage account changes keys, then you can click **Resync keys** to synchronize the keys with the workspace.

To manage the web services associated with this workspace, use the Machine Learning Web Services portal. See [Manage a Web service using the Azure Machine Learning Web Services portal](#) for complete information.

NOTE

To deploy or manage New web services you must be assigned a contributor or administrator role on the subscription to which the web service is deployed. If you invite another user to a machine learning workspace, you must assign them to a contributor or administrator role on the subscription before they can deploy or manage web services.

For more information on setting access permissions, see [View access assignments for users and groups in the Azure portal](#).

Next steps

- Learn more about [deploy Machine Learning with Azure Resource Manager Templates](#).

Troubleshooting guide: Create and connect to an Machine Learning workspace

9/25/2017 • 1 min to read • [Edit Online](#)

This guide provides solutions for some frequently encountered challenges when you are setting up Azure Machine Learning workspaces.

NOTE

Try Azure Machine Learning for free

No credit card or Azure subscription required. [Get started now.](#)

Workspace owner

To open a workspace in Machine Learning Studio, you must be signed in to the Microsoft Account you used to create the workspace, or you need to receive an invitation from the owner to join the workspace. From the Azure portal you can manage the workspace, which includes the ability to configure access.

For more information on managing a workspace, see [Manage an Azure Machine Learning workspace](#).

Allowed regions

Machine Learning is currently available in a limited number of regions. If your subscription does not include one of these regions, you may see the error message, "You have no subscriptions in the allowed regions."

To request that a region be added to your subscription, create a new Microsoft support request from the Azure portal, choose **Billing** as the problem type, and follow the prompts to submit your request.

Storage account

The Machine Learning service needs a storage account to store data. You can use an existing storage account, or you can create a new storage account when you create the new Machine Learning workspace (if you have quota to create a new storage account).

After the new Machine Learning workspace is created, you can sign in to Machine Learning Studio by using the Microsoft account you used to create the workspace. If you encounter the error message, "Workspace Not Found" (similar to the following screenshot), please use the following steps to delete your browser cookies.

Workspace Not Found

[SIGN OUT](#)

[ABOUT MICROSOFT AZURE ML](#)

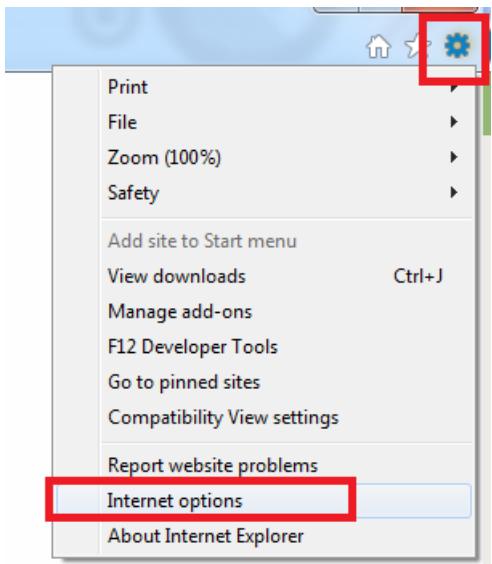
[MACHINE LEARNING CENTER](#)

[CONTACT MICROSOFT SUPPORT](#)

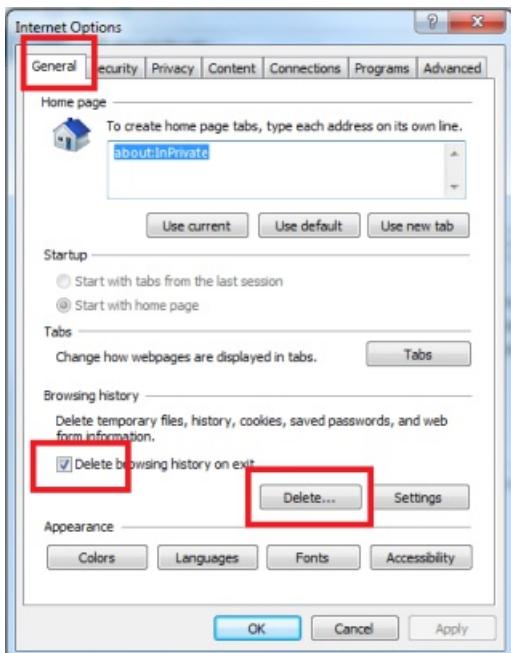
The user is not authorized for the workspace

To delete browser cookies

1. If you use Internet Explorer, click the **Tools** button in the upper-right corner and select **Internet options**.



1. Under the **General** tab, click **Delete...**



1. In the **Delete Browsing History** dialog box, make sure **Cookies and website data** is selected, and click **Delete**.



After the cookies are deleted, restart the browser and then go to the [Microsoft Azure Machine Learning](#) page. When you are prompted for a user name and password, enter the same Microsoft account you used to create the workspace.

Comments

Our goal is to make the Machine Learning experience as seamless as possible. Please post any comments and issues at the [Azure Machine Learning forum](#) to help us serve you better.

Deploy Machine Learning Workspace Using Azure Resource Manager

9/25/2017 • 3 min to read • [Edit Online](#)

Introduction

Using an Azure Resource Manager deployment template saves you time by giving you a scalable way to deploy interconnected components with a validation and retry mechanism. To set up Azure Machine Learning Workspaces, for example, you need to first configure an Azure storage account and then deploy your workspace. Imagine doing this manually for hundreds of workspaces. An easier alternative is to use an Azure Resource Manager template to deploy an Azure Machine Learning Workspace and all its dependencies. This article takes you through this process step-by-step. For a great overview of Azure Resource Manager, see [Azure Resource Manager overview](#).

Step-by-step: create a Machine Learning Workspace

We will create an Azure resource group, then deploy a new Azure storage account and a new Azure Machine Learning Workspace using a Resource Manager template. Once the deployment is complete, we will print out important information about the workspaces that were created (the primary key, the workspaceID, and the URL to the workspace).

Create an Azure Resource Manager template

A Machine Learning Workspace requires an Azure storage account to store the dataset linked to it. The following template uses the name of the resource group to generate the storage account name and the workspace name. It also uses the storage account name as a property when creating the workspace.

```
{
  "contentVersion": "1.0.0.0",
  "$schema": "http://schema.management.azure.com/schemas/2015-01-01/deploymentTemplate.json#",
  "variables": {
    "namePrefix": "[resourceGroup().name]",
    "location": "[resourceGroup().location]",
    "mlVersion": "2016-04-01",
    "stgVersion": "2015-06-15",
    "storageAccountName": "[concat(variables('namePrefix'),'stg')]",
    "mlWorkspaceName": "[concat(variables('namePrefix'),'mlwk')]",
    "mlResourceId": "[resourceId('Microsoft.MachineLearning/workspaces', variables('mlWorkspaceName'))]",
    "stgResourceId": "[resourceId('Microsoft.Storage/storageAccounts', variables('storageAccountName'))]",
    "storageAccountType": "Standard_LRS"
  },
  "resources": [
    {
      "apiVersion": "[variables('stgVersion')]",
      "name": "[variables('storageAccountName')]",
      "type": "Microsoft.Storage/storageAccounts",
      "location": "[variables('location')]",
      "properties": {
        "accountType": "[variables('storageAccountType')]"
      }
    },
    {
      "apiVersion": "[variables('mlVersion')]",
      "type": "Microsoft.MachineLearning/workspaces",
      "name": "[variables('mlWorkspaceName')]",
      "location": "[variables('location')]",
      "dependsOn": "[variables('stgResourceId')]",
      "properties": {
        "UserStorageAccountId": "[variables('stgResourceId')]"
      }
    }
  ],
  "outputs": {
    "mlWorkspaceObject": {"type": "object", "value": "[reference(variables('mlResourceId'), variables('mlVersion'))]"},
    "mlWorkspaceToken": {"type": "string", "value": "[listWorkspaceKeys(variables('mlResourceId'), variables('mlVersion')).primaryToken]"},
    "mlWorkspaceWorkspaceID": {"type": "string", "value": "[reference(variables('mlResourceId'), variables('mlVersion')).WorkspaceId]"},
    "mlWorkspaceWorkspaceLink": {"type": "string", "value": "[concat('https://studio.azureml.net/Home/ViewWorkspace/', reference(variables('mlResourceId'), variables('mlVersion')).WorkspaceId)]"}
  }
}
```

Save this template as mlworkspace.json file under c:\temp.

Deploy the resource group, based on the template

- Open PowerShell
- Install modules for Azure Resource Manager and Azure Service Management

```
# Install the Azure Resource Manager modules from the PowerShell Gallery (press "A")
Install-Module AzureRM -Scope CurrentUser

# Install the Azure Service Management modules from the PowerShell Gallery (press "A")
Install-Module Azure -Scope CurrentUser
```

These steps download and install the modules necessary to complete the remaining steps. This only needs to be done once in the environment where you are executing the PowerShell commands.

- Authenticate to Azure

```
# Authenticate (enter your credentials in the pop-up window)
Add-AzureRmAccount
```

This step needs to be repeated for each session. Once authenticated, your subscription information should be displayed.

```
Environment      : AzureCloud
Account         :
TenantId        :
SubscriptionId  :
SubscriptionName: Windows Azure MSDN - Visual Studio Ultimate
CurrentStorageAccount :
```

Now that we have access to Azure, we can create the resource group.

- Create a resource group

```
$rg = New-AzureRmResourceGroup -Name "uniquenamerequired523" -Location "South Central US"
$rg
```

Verify that the resource group is correctly provisioned. **ProvisioningState** should be "Succeeded." The resource group name is used by the template to generate the storage account name. The storage account name must be between 3 and 24 characters in length and use numbers and lower-case letters only.

```
ResourceGroupName : uniquenamerequired543
Location         : southcentralus
ProvisioningState : Succeeded
Tags             :
ResourceId       : /subscriptions/xxxxxxxx-xxxx-xxxx-xxxx-xxxxxxxxxxxx/resourceGroups/uniquenamerequired543
```

- Using the resource group deployment, deploy a new Machine Learning Workspace.

```
# Create a Resource Group, TemplateFile is the location of the JSON template.
$rgd = New-AzureRmResourceGroupDeployment -Name "demo" -TemplateFile "C:\temp\mlworkspace.json" -ResourceGroupName
$rg.ResourceGroupName
```

Once the deployment is completed, it is straightforward to access properties of the workspace you deployed. For example, you can access the Primary Key Token.

```
# Access Azure ML Workspace Token after its deployment.
$rgd.Outputs.mlWorkspaceToken.Value
```

Another way to retrieve tokens of existing workspace is to use the `Invoke-AzureRmResourceAction` command. For example, you can list the primary and secondary tokens of all workspaces.

```
# List the primary and secondary tokens of all workspaces
Get-AzureRmResource |? { $_.ResourceType -Like "*MachineLearning/workspaces*" } |% { Invoke-AzureRmResourceAction -ResourceId
$_._ResourceId -Action listworkspacekeys -Force }
```

After the workspace is provisioned, you can also automate many Azure Machine Learning Studio tasks using the [PowerShell Module for Azure Machine Learning](#).

Next Steps

- Learn more about [authoring Azure Resource Manager Templates](#).
- Have a look at the [Azure Quickstart Templates Repository](#).
- Watch this video about [Azure Resource Manager](#).

4 min to read •

Import your training data into Azure Machine Learning Studio from various data sources

11/30/2017 • 3 min to read • [Edit Online](#)

To use your own data in Machine Learning Studio to develop and train a predictive analytics solution, you can:

- upload data from a **local file** ahead of time from your hard drive to create a dataset module in your workspace
- access data from one of several **online data sources** while your experiment is running using the [Import Data](#) module
- use data from another Azure Machine learning **experiment** saved as a dataset
- use data from an on-premises **SQL Server database**

Each of these options is described in one of the topics on the menu below. These topics show you how to import data from these various data sources to use in Machine Learning Studio.

NOTE

There are a number of sample datasets available in Machine Learning Studio that you can use for training data. For information on these, see [Use the sample datasets in Azure Machine Learning Studio](#).

This introductory topic also discusses how to get data ready for use in Machine Learning Studio and describes which data formats and data types are supported.

NOTE

Try Azure Machine Learning for free

No credit card or Azure subscription required. [Get started now](#).

Get data ready for use in Azure Machine Learning Studio

Machine Learning Studio is designed to work with rectangular or tabular data, such as text data that's delimited or structured data from a database, though in some circumstances non-rectangular data may be used.

It's best if your data is relatively clean. That is, you'll want to take care of issues such as unquoted strings before you upload the data into your experiment.

However, there are modules available in Machine Learning Studio that enable some manipulation of data within your experiment. Depending on the machine learning algorithms you'll be using, you may need to decide how you'll handle data structural issues such as missing values and sparse data, and there are modules that can help with that. Look in the **Data Transformation** section of the module palette for modules that perform these functions.

At any point in your experiment you can view or download the data that's produced by a module by clicking the output port. Depending on the module, there may be different download options available, or you may be able to visualize the data within your web browser in Machine Learning Studio.

Data formats and data types supported

You can import a number of data types into your experiment, depending on what mechanism you use to import data and where it's coming from:

- Plain text (.txt)
- Comma-separated values (CSV) with a header (.csv) or without (.nh.csv)
- Tab-separated values (TSV) with a header (.tsv) or without (.nh.tsv)
- Excel file
- Azure table
- Hive table
- SQL database table
- OData values
- SVMLight data (.svmlight) (see the [SVMLight definition](#) for format information)
- Attribute Relation File Format (ARFF) data (.arff) (see the [ARFF definition](#) for format information)
- Zip file (.zip)
- R object or workspace file (.RData)

If you import data in a format such as ARFF that includes metadata, Machine Learning Studio uses this metadata to define the heading and data type of each column.

If you import data such as TSV or CSV format that doesn't include this metadata, Machine Learning Studio infers the data type for each column by sampling the data. If the data also doesn't have column headings, Machine Learning Studio provides default names.

You can explicitly specify or change the headings and data types for columns using the [Edit Metadata](#).

The following **data types** are recognized by Machine Learning Studio:

- String
- Integer
- Double
- Boolean
- DateTime
- TimeSpan

Machine Learning Studio uses an internal data type called **Data Table** to pass data between modules. You can explicitly convert your data into Data Table format using the [Convert to Dataset](#) module.

Any module that accepts formats other than Data Table will convert the data to Data Table silently before passing it to the next module.

If necessary, you can convert Data Table format back into CSV, TSV, ARFF, or SVMLight format using other conversion modules. Look in the **Data Format Conversions** section of the module palette for modules that perform these functions.

Import training data from a file on your hard drive into Machine Learning Studio

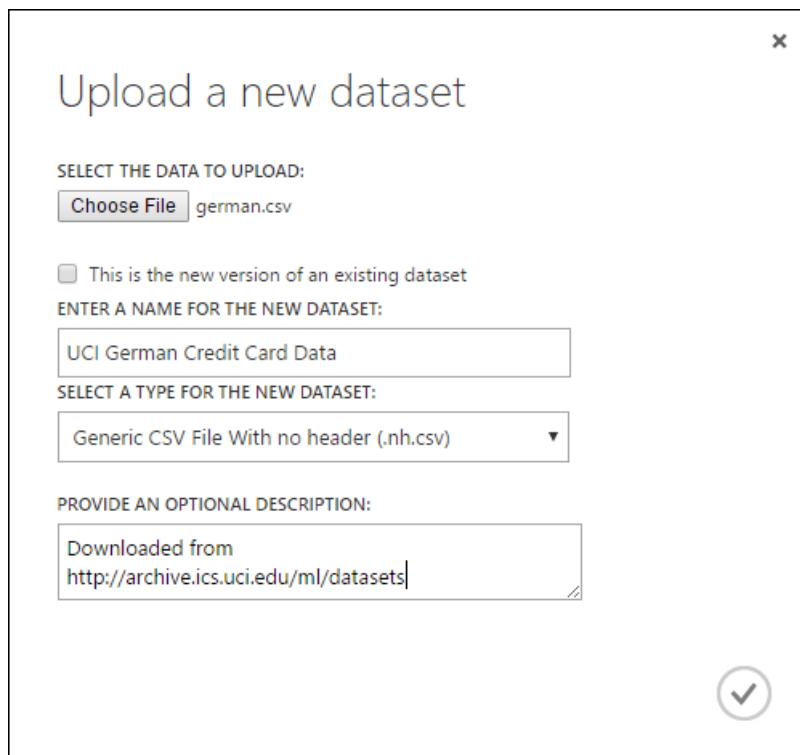
9/25/2017 • 1 min to read • [Edit Online](#)

Learn how to upload a data file from your hard drive to use as training data in Azure Machine Learning Studio. By importing the data file, you have a dataset module ready for use in your workspace.

Steps to import data from a local file

To import data from a local hard drive, do the following:

1. Click **+NEW** at the bottom of the Machine Learning Studio window.
2. Select **DATASET** and **FROM LOCAL FILE**.
3. In the **Upload a new dataset** dialog, browse to the file you want to upload
4. Enter a name, identify the data type, and optionally enter a description. A description is recommended - it allows you to record any characteristics about the data that you want to remember when using the data in the future.
5. The checkbox **This is the new version of an existing dataset** allows you to update an existing dataset with new data. Click this checkbox and then enter the name of an existing dataset.



During upload, you'll see a message that your file is being uploaded. Upload time depends on the size of your data and the speed of your connection to the service. If you know the file will take a long time, you can do other things inside Machine Learning Studio while you wait. However, closing the browser causes the data upload to fail.

Dataset module is ready for use

Once your data is uploaded, it's stored in a dataset module and is available to any experiment in your workspace.

When you're editing an experiment, you can find the datasets you've created in the **My Datasets** list under the

Saved Datasets list in the module palette. You can drag and drop the dataset onto the experiment canvas when you want to use the dataset for further analytics and machine learning.

Import data into Azure Machine Learning Studio from various online data sources with the Import Data module

12/11/2017 • 7 min to read • [Edit Online](#)

This article describes the support for importing online data from various sources and the information needed to move data from these sources into an Azure Machine Learning experiment.

NOTE

This article provides general information about the [Import Data](#) module. For more detailed information about the types of data you can access, formats, parameters, and answers to common questions, see the module reference topic for the [Import Data](#) module.

Introduction

By using the [Import Data](#) module, you can access data from one of several online data sources while your experiment is running in [Azure Machine Learning Studio](#):

- A Web URL using HTTP
- Hadoop using HiveQL
- Azure blob storage
- Azure table
- Azure SQL database or SQL Server on Azure VM
- On-premises SQL Server database
- A data feed provider, OData currently
- Azure Cosmos DB

To access online data sources in your Studio experiment, add the [Import Data](#) module to your, select the **Data source**, and then provide the parameters needed to access the data. The online data sources that are supported are itemized in the table below. This table also summarizes the file formats that are supported and parameters that are used to access the data.

Note that because this training data is accessed while your experiment is running, it's only available in that experiment. By comparison, data that has been stored in a dataset module are available to any experiment in your workspace.

IMPORTANT

Currently, the [Import Data](#) and [Export Data](#) modules can read and write data only from Azure storage created using the Classic deployment model. In other words, the new Azure Blob Storage account type that offers a hot storage access tier or cool storage access tier is not yet supported.

Generally, any Azure storage accounts that you might have created before this service option became available should not be affected. If you need to create a new account, select **Classic** for the Deployment model, or use Resource manager and select **General purpose** rather than **Blob storage** for **Account kind**.

For more information, see [Azure Blob Storage: Hot and Cool Storage Tiers](#).

Supported online data sources

Azure Machine Learning **Import Data** module supports the following data sources:

DATA SOURCE	DESCRIPTION	PARAMETERS
Web URL via HTTP	Reads data in comma-separated values (CSV), tab-separated values (TSV), attribute-relation file format (ARFF), and Support Vector Machines (SVM-light) formats, from any web URL that uses HTTP	<p>URL: Specifies the full name of the file, including the site URL and the file name, with any extension.</p> <p>Data format: Specifies one of the supported data formats: CSV, TSV, ARFF, or SVM-light. If the data has a header row, it is used to assign column names.</p>
Hadoop/HDFS	Reads data from distributed storage in Hadoop. You specify the data you want by using HiveQL, a SQL-like query language. HiveQL can also be used to aggregate data and perform data filtering before you add the data to Machine Learning Studio.	<p>Hive database query: Specifies the Hive query used to generate the data.</p> <p>HCatalog server URI : Specified the name of your cluster using the format <<i>your cluster name</i>>.azurehdinsight.net.</p> <p>Hadoop user account name: Specifies the Hadoop user account name used to provision the cluster.</p> <p>Hadoop user account password : Specifies the credentials used when provisioning the cluster. For more information, see Create Hadoop clusters in HDInsight.</p> <p>Location of output data: Specifies whether the data is stored in a Hadoop distributed file system (HDFS) or in Azure. If you store output data in HDFS, specify the HDFS server URI. (Be sure to use the HDInsight cluster name without the HTTPS:// prefix).</p> <p>If you store your output data in Azure, you must specify the Azure storage account name, Storage access key and Storage container name.</p>

DATA SOURCE	DESCRIPTION	PARAMETERS
SQL database	<p>Reads data that is stored in an Azure SQL database or in a SQL Server database running on an Azure virtual machine.</p>	<p>Database server name: Specifies the name of the server on which the database is running. In case of Azure SQL Database enter the server name that is generated. Typically it has the form <i><generated_identifier>.database.windows.net</i>.</p> <p>In case of a SQL server hosted on a Azure Virtual machine enter <i>tcp:<Virtual Machine DNS Name>,1433</i></p> <p>Database name : Specifies the name of the database on the server.</p> <p>Server user account name: Specifies a user name for an account that has access permissions for the database.</p> <p>Server user account password: Specifies the password for the user account.</p> <p>Database query: Enter a SQL statement that describes the data you want to read.</p>

DATA SOURCE	DESCRIPTION	PARAMETERS
On-premises SQL database	<p>Reads data that is stored in an on-premises SQL database.</p>	<p>Data gateway: Specifies the name of the Data Management Gateway installed on a computer where it can access your SQL Server database. For information about setting up the gateway, see Perform advanced analytics with Azure Machine Learning using data from an on-premises SQL server.</p> <p>Database server name: Specifies the name of the server on which the database is running.</p> <p>Database name : Specifies the name of the database on the server.</p> <p>Server user account name: Specifies a user name for an account that has access permissions for the database.</p> <p>User name and password: Click Enter values to enter your database credentials. You can use Windows Integrated Authentication or SQL Server Authentication depending upon how your on-premises SQL Server is configured.</p> <p>Database query: Enter a SQL statement that describes the data you want to read.</p>
Azure Table	<p>Reads data from the Table service in Azure Storage.</p> <p>If you read large amounts of data infrequently, use the Azure Table Service. It provides a flexible, non-relational (NoSQL), massively scalable, inexpensive, and highly available storage solution.</p>	<p>The options in the Import Data change depending on whether you are accessing public information or a private storage account that requires login credentials. This is determined by the Authentication Type which can have value of "PublicOrSAS" or "Account", each of which has its own set of parameters.</p> <p>Public or Shared Access Signature (SAS) URI: The parameters are:</p> <p>Table URI: Specifies the Public or SAS URL for the table.</p> <p>Specifies the rows to scan for property names: The values are <i>TopN</i> to scan the specified number of rows, or <i>ScanAll</i> to get all rows in the table.</p> <p>If the data is homogeneous and predictable, it is recommended that you select <i>TopN</i> and enter a number for N. For large tables, this can result in quicker reading times.</p> <p>If the data is structured with sets of</p>

DATA SOURCE	DESCRIPTION	
	<p>properties that vary based on the PARAMETERS depth and position of the table, choose the <i>ScanAll</i> option to scan all rows. This ensures the integrity of your resulting property and metadata conversion.</p> <p>Private Storage Account: The parameters are:</p> <p>Account name: Specifies the name of the account that contains the table to read.</p> <p>Account key: Specifies the storage key associated with the account.</p> <p>Table name : Specifies the name of the table that contains the data to read.</p> <p>Rows to scan for property names: The values are <i>TopN</i> to scan the specified number of rows, or <i>ScanAll</i> to get all rows in the table.</p> <p>If the data is homogeneous and predictable, we recommend that you select <i>TopN</i> and enter a number for N. For large tables, this can result in quicker reading times.</p> <p>If the data is structured with sets of properties that vary based on the depth and position of the table, choose the <i>ScanAll</i> option to scan all rows. This ensures the integrity of your resulting property and metadata conversion.</p>	

DATA SOURCE	DESCRIPTION	PARAMETERS
Azure Blob Storage	<p>Reads data stored in the Blob service in Azure Storage, including images, unstructured text, or binary data.</p> <p>You can use the Blob service to publicly expose data, or to privately store application data. You can access your data from anywhere by using HTTP or HTTPS connections.</p>	<p>The options in the Import Data module change depending on whether you are accessing public information or a private storage account that requires login credentials. This is determined by the Authentication Type which can have a value either of "PublicOrSAS" or of "Account".</p> <p>Public or Shared Access Signature (SAS) URI: The parameters are:</p> <ul style="list-style-type: none"> URI: Specifies the Public or SAS URL for the storage blob. File Format: Specifies the format of the data in the Blob service. The supported formats are CSV, TSV, and ARFF. <p>Private Storage Account: The parameters are:</p> <ul style="list-style-type: none"> Account name: Specifies the name of the account that contains the blob you want to read. Account key: Specifies the storage key associated with the account. Path to container, directory, or blob : Specifies the name of the blob that contains the data to read. Blob file format: Specifies the format of the data in the blob service. The supported data formats are CSV, TSV, ARFF, CSV with a specified encoding, and Excel. If the format is CSV or TSV, be sure to indicate whether the file contains a header row. You can use the Excel option to read data from Excel workbooks. In the <i>Excel data format</i> option, indicate whether the data is in an Excel worksheet range, or in an Excel table. In the <i>Excel sheet or embedded table</i> option, specify the name of the sheet or table that you want to read from.

DATA SOURCE	DESCRIPTION	PARAMETERS
Data Feed Provider	Reads data from a supported feed provider. Currently only the Open Data Protocol (OData) format is supported.	<p>Data content type: Specifies the OData format.</p> <p>Source URL: Specifies the full URL for the data feed. For example, the following URL reads from the Northwind sample database: http://services.odata.org/northwind/northwind.svc/</p>

Next steps

[Deploying Azure ML web services that use Data Import and Data Export modules](#)

Import your data into Azure Machine Learning Studio from another experiment

9/25/2017 • 1 min to read • [Edit Online](#)

There will be times when you'll want to take an intermediate result from one experiment and use it as part of another experiment. To do this, you save the module as a dataset:

1. Click the output of the module that you want to save as a dataset.
2. Click **Save as Dataset**.
3. When prompted, enter a name and a description that would allow you to identify the dataset easily.
4. Click the **OK** checkmark.

When the save finishes, the dataset will be available for use within any experiment in your workspace. You can find it in the **Saved Datasets** list in the module palette.

Perform advanced analytics with Azure Machine Learning using data from an on-premises SQL Server database

9/25/2017 • 9 min to read • [Edit Online](#)

Often enterprises that work with on-premises data would like to take advantage of the scale and agility of the cloud for their machine learning workloads. But they don't want to disrupt their current business processes and workflows by moving their on-premises data to the cloud. Azure Machine Learning now supports reading your data from an on-premises SQL Server database and then training and scoring a model with this data. You no longer have to manually copy and sync the data between the cloud and your on-premises server. Instead, the **Import Data** module in Azure Machine Learning Studio can now read directly from your on-premises SQL Server database for your training and scoring jobs.

This article provides an overview of how to ingress on-premises SQL server data into Azure Machine Learning. It assumes that you're familiar with Azure Machine Learning concepts like workspaces, modules, datasets, experiments, etc..

NOTE

This feature is not available for free workspaces. For more information about Machine Learning pricing and tiers, see [Azure Machine Learning Pricing](#).

NOTE

[Try Azure Machine Learning for free](#)

No credit card or Azure subscription required. [Get started now](#).

Install the Microsoft Data Management Gateway

To access an on-premises SQL Server database in Azure Machine Learning, you need to download and install the Microsoft Data Management Gateway. When you configure the gateway connection in Machine Learning Studio, you have the opportunity to download and install the gateway using the **Download and register data gateway** dialog described below.

You also can install the Data Management Gateway ahead of time by downloading and running the MSI setup package from the [Microsoft Download Center](#). Choose the latest version, selecting either 32-bit or 64-bit as appropriate for your computer. The MSI can also be used to upgrade an existing Data Management Gateway to the latest version, with all settings preserved.

The gateway has the following prerequisites:

- The supported Windows operating system versions are Windows 7, Windows 8/8.1, Windows 10, Windows Server 2008 R2, Windows Server 2012, and Windows Server 2012 R2.
- The recommended configuration for the gateway machine is at least 2 GHz, 4 cores, 8GB RAM, and 80GB disk.
- If the host machine hibernates, the gateway won't respond to data requests. Therefore, configure an appropriate power plan on the computer before installing the gateway. If the machine is configured to hibernate, the gateway installation displays a message.

- Because copy activity occurs at a specific frequency, the resource usage (CPU, memory) on the machine also follows the same pattern with peak and idle times. Resource utilization also depends heavily on the amount of data being moved. When multiple copy jobs are in progress, you'll observe resource usage go up during peak times. While the minimum configuration listed above is technically sufficient, you may want to have a configuration with more resources than the minimum configuration depending on your specific load for data movement.

Consider the following when setting up and using a Data Management Gateway:

- You can install only one instance of Data Management Gateway on a single computer.
- You can use a single gateway for multiple on-premises data sources.
- You can connect multiple gateways on different computers to the same on-premises data source.
- You configure a gateway for only one workspace at a time. Currently, gateways can't be shared across workspaces.
- You can configure multiple gateways for a single workspace. For example, you may want to use a gateway that's connected to your test data sources during development and a production gateway when you're ready to operationalize.
- The gateway does not need to be on the same machine as the data source. But staying closer to the data source reduces the time for the gateway to connect to the data source. We recommend that you install the gateway on a machine that's different from the one that hosts the on-premises data source so that the gateway and data source don't compete for resources.
- If you already have a gateway installed on your computer serving Power BI or Azure Data Factory scenarios, install a separate gateway for Azure Machine Learning on another computer.

NOTE

You can't run Data Management Gateway and Power BI Gateway on the same computer.

- You need to use the Data Management Gateway for Azure Machine Learning even if you are using Azure ExpressRoute for other data. You should treat your data source as an on-premises data source (that's behind a firewall) even when you use ExpressRoute. Use the Data Management Gateway to establish connectivity between Machine Learning and the data source.

You can find detailed information on installation prerequisites, installation steps, and troubleshooting tips in the article [Data Management Gateway](#).

Ingress data from your on-premises SQL Server database into Azure Machine Learning

In this walkthrough, you will set up a Data Management Gateway in an Azure Machine Learning workspace, configure it, and then read data from an on-premises SQL Server database.

TIP

Before you start, disable your browser's pop-up blocker for studio.azureml.net. If you're using the Google Chrome browser, download and install one of the several plug-ins available at Google Chrome WebStore [Click Once App Extension](#).

Step 1: Create a gateway

The first step is to create and set up the gateway to access your on-premises SQL database.

- Log in to [Azure Machine Learning Studio](#) and select the workspace that you want to work in.
- Click the **SETTINGS** blade on the left, and then click the **DATA GATEWAYS** tab at the top.

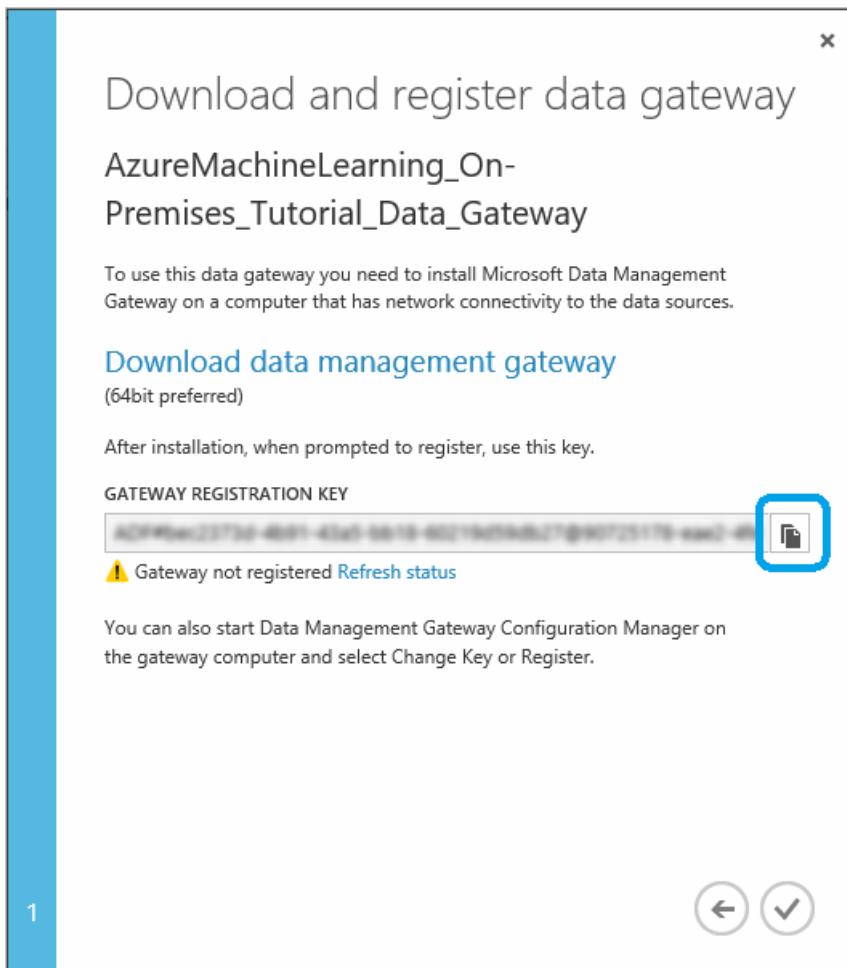
3. Click **NEW DATA GATEWAY** at the bottom of the screen.

The screenshot shows the Azure Machine Learning studio interface. On the left, there is a vertical sidebar with icons and labels: PROJECTS, EXPERIMENTS, WEB SERVICES, NOTEBOOKS, DATASETS, TRAINED MODELS, and SETTINGS. Below this is a dark bar with a '+ NEW' button and a 'NEW DATA GATEWAY' button, which is highlighted with a blue square. The main area is titled 'settings' and has tabs for NAME, AUTHORIZATION TOKENS, USERS, and DATA GATEWAYS. The 'DATA GATEWAYS' tab is selected and highlighted with a red rectangle. Below the tabs is a table header with columns for NAME, HOST, and STATUS, along with a search icon. At the bottom of the interface are buttons for DELETE, EDIT DESCRIPTION, DOWNLOAD AND REGISTER, and REFRESH.

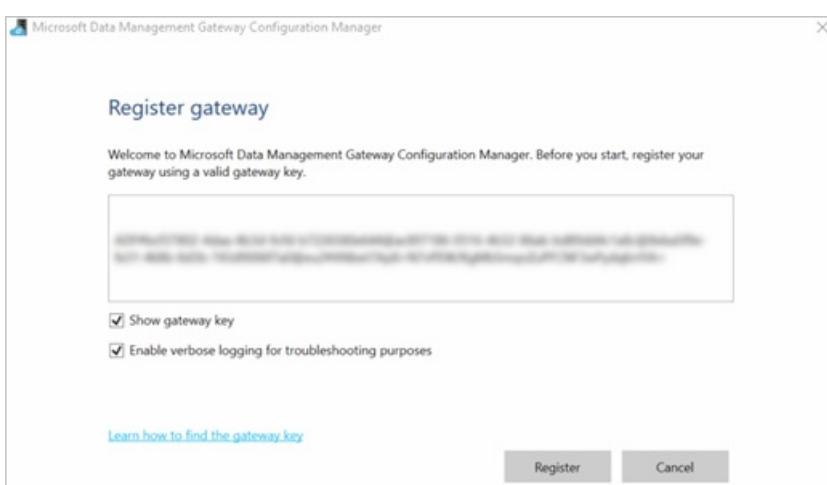
4. In the **New data gateway** dialog, enter the **Gateway Name** and optionally add a **Description**. Click the arrow on the bottom right-hand corner to go to the next step of the configuration.

The screenshot shows the 'New data gateway' configuration dialog. The title is 'New data gateway'. A descriptive text states: 'A data gateway provides access to on-premise data sources in your corporate environment.' There are two input fields: 'GATEWAY NAME' containing 'AzureMachineLearning_On-Premises_Tutorial_Data_Gateway' and 'DESCRIPTION' containing 'This gateway is used to access data from my on-premises SQL server.' At the bottom right is a large blue button with a white arrow pointing right, labeled '2' in the bottom right corner.

5. In the Download and register data gateway dialog, copy the GATEWAY REGISTRATION KEY to the clipboard.

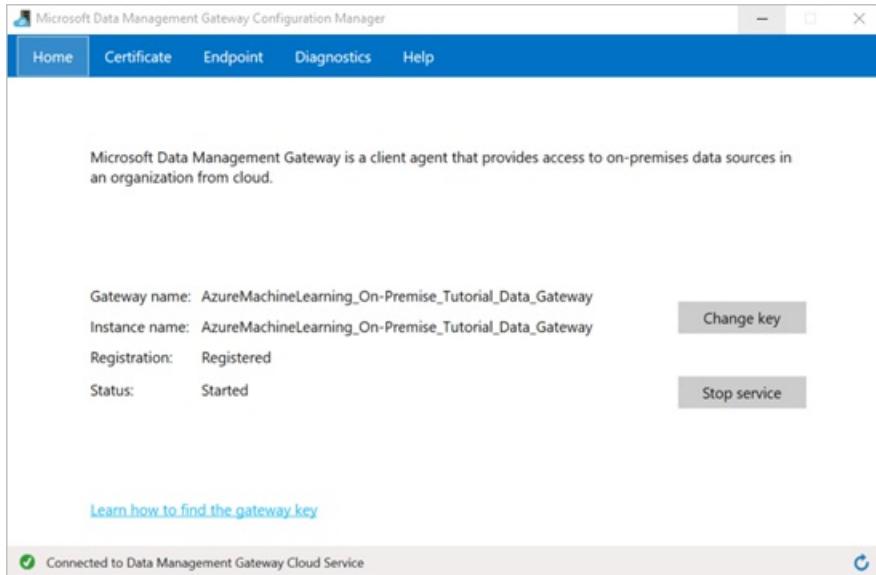


6. If you have not yet downloaded and installed the Microsoft Data Management Gateway, then click **Download data management gateway**. This takes you to the Microsoft Download Center where you can select the gateway version you need, download it, and install it. You can find detailed information on installation prerequisites, installation steps, and troubleshooting tips in the beginning sections of the article [Move data between on-premises sources and cloud with Data Management Gateway](#).
7. After the gateway is installed, the Data Management Gateway Configuration Manager will open and the **Register gateway** dialog is displayed. Paste the **Gateway Registration Key** that you copied to the clipboard and click **Register**.
8. If you already have a gateway installed, run the Data Management Gateway Configuration Manager. Click **Change key**, paste the **Gateway Registration Key** that you copied to the clipboard in the previous step, and click **OK**.
9. When the installation is complete, the **Register gateway** dialog for Microsoft Data Management Gateway Configuration Manager is displayed. Paste the GATEWAY REGISTRATION KEY that you copied to the clipboard in a previous step and click **Register**.

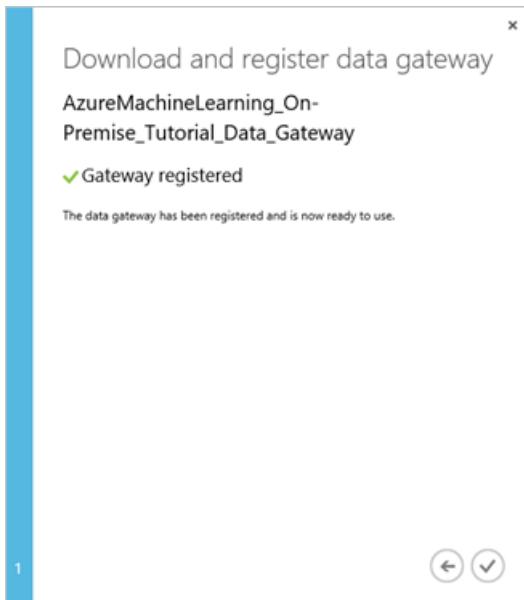


10. The gateway configuration is complete when the following values are set on the **Home** tab in Microsoft Data Management Gateway Configuration Manager:

- **Gateway name** and **Instance name** are set to the name of the gateway.
- **Registration** is set to **Registered**.
- **Status** is set to **Started**.
- The status bar at the bottom displays **Connected to Data Management Gateway Cloud Service** along with a green check mark.



Azure Machine Learning Studio also gets updated when the registration is successful.



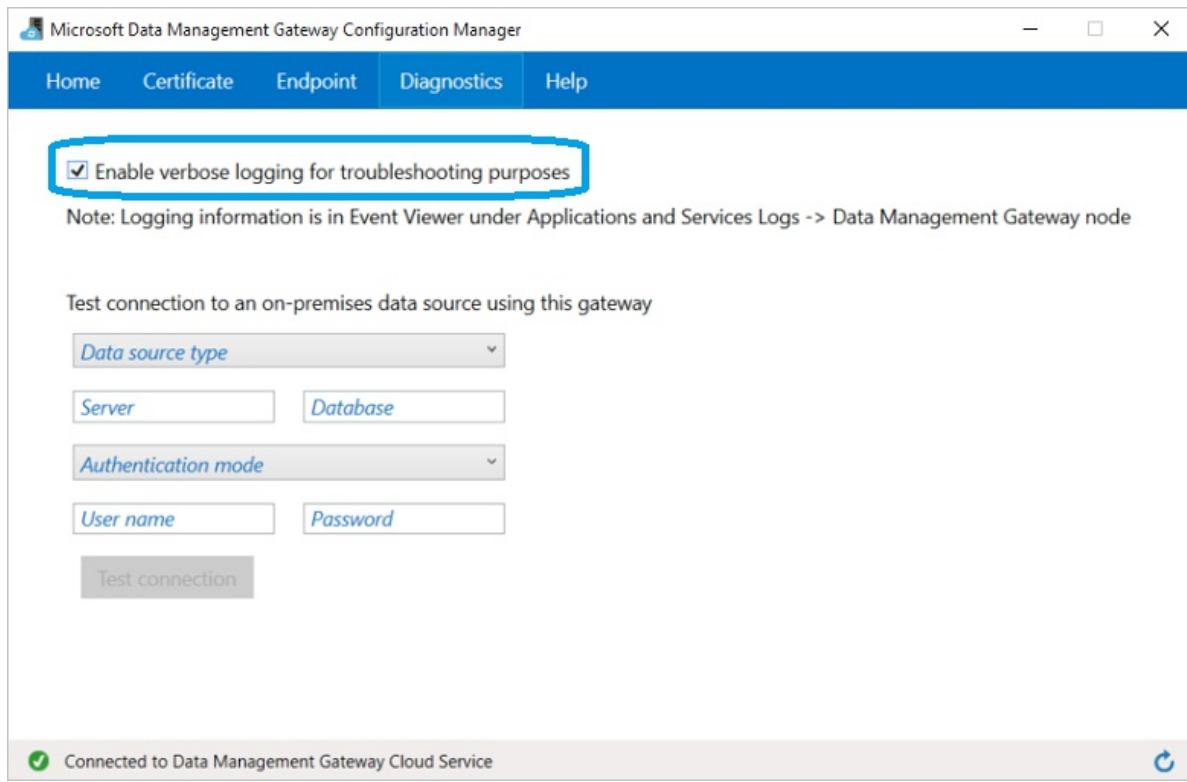
11. In the **Download and register data gateway** dialog, click the check mark to complete the setup. The **Settings** page displays the gateway status as "Online". In the right-hand pane, you'll find status and other useful information.

The screenshot shows the 'settings' page of the Microsoft Data Management Gateway Configuration Manager. On the left, there's a sidebar with icons for Experiments, Web Services, Notebooks, Datasets, Trained Models, and Settings. The main area has a title 'settings' and a table header with columns: NAME, AUTHORIZATION TOKENS, USERS, and DATA GATEWAYS. A single row is selected in the table, showing 'AzureMachineLearning_On-Premise_Tuto...' under NAME, 'krishnacarbon3r.redmond.corp.microsoft.com' under HOST, and 'Online - Up to date' under STATUS. To the right of the table, detailed information about the gateway is displayed, including its description ('This Gateway is used to access data from my on-premise sql server'), host ('krishnacarbon3r.redmond.corp.microsoft.com'), status ('Online'), last connected time ('11/17/2015 12:58:52 PM'), version ('1.7.5764.1'), and creation and registration times ('Created: 11/17/2015 11:28:10 AM; Registered: 11/17/2015 12:45:06 PM'). At the bottom of the table area are buttons for NEW, NEW DATA GATEWAY, DELETE, EDIT DESCRIPTION, DOWNLOAD AND REGISTER, and REFRESH.

12. In the Microsoft Data Management Gateway Configuration Manager switch to the **Certificate** tab. The certificate specified on this tab is used to encrypt/decrypt credentials for the on-premises data store that you specify in the portal. This certificate is the default certificate. Microsoft recommends changing this to your own certificate that you back up in your certificate management system. Click **Change** to use your own certificate instead.

The screenshot shows the Microsoft Data Management Gateway Configuration Manager with the 'Certificate' tab selected. A yellow warning bar at the top says 'Export and backup the certificate that you must use to restore the gateway in case of a failure.' Below this, a message states 'The certificate will be used to encrypt data source credentials before they are saved into the cloud credential store.' A section labeled 'Current certificate:' contains a text input field with 'Data Management Gateway Credential Certificate' and three buttons: 'View...', 'Change...', and 'Export...'. At the bottom of the window, a status bar says 'Connected to Data Management Gateway Cloud Service'.

13. (optional) If you want to enable verbose logging in order to troubleshoot issues with the gateway, in the Microsoft Data Management Gateway Configuration Manager switch to the **Diagnostics** tab and check the **Enable verbose logging for troubleshooting purposes** option. The logging information can be found in the Windows Event Viewer under the **Applications and Services Logs -> Data Management Gateway** node. You can also use the **Diagnostics** tab to test the connection to an on-premises data source using the gateway.



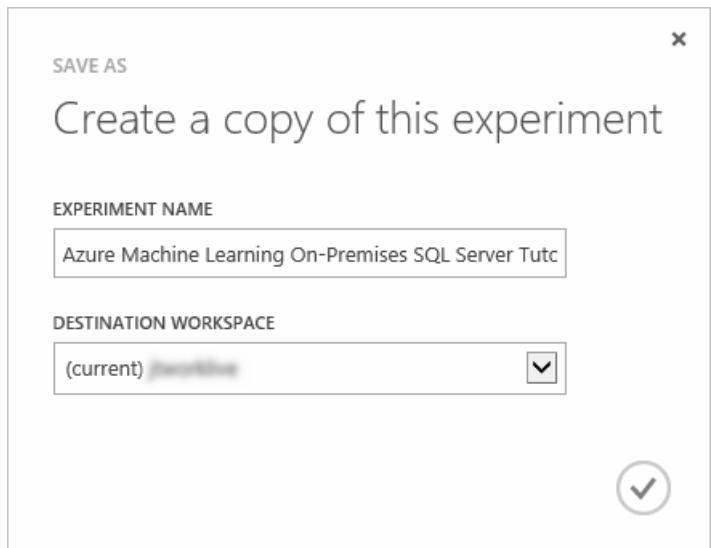
This completes the gateway setup process in Azure Machine Learning. You're now ready to use your on-premises data.

You can create and set up multiple gateways in Studio for each workspace. For example, you may have a gateway that you want to connect to your test data sources during development, and a different gateway for your production data sources. Azure Machine Learning gives you the flexibility to set up multiple gateways depending upon your corporate environment. Currently you can't share a gateway between workspaces and only one gateway can be installed on a single computer. For more information, see [Move data between on-premises sources and cloud with Data Management Gateway](#).

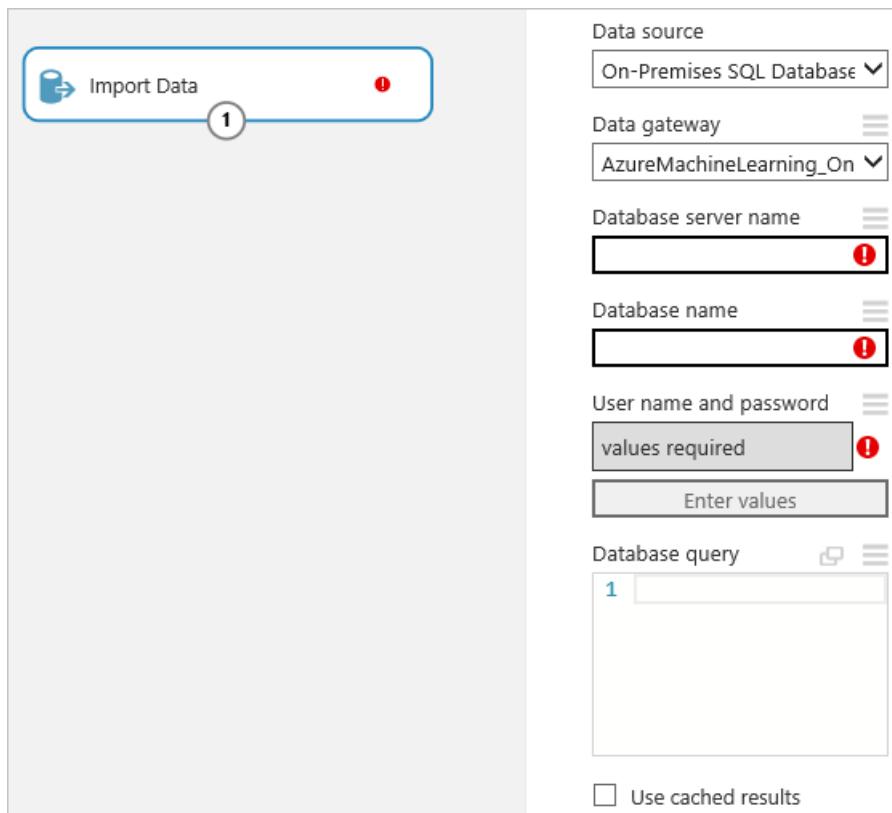
Step 2: Use the gateway to read data from an on-premises data source

After you set up the gateway, you can add an **Import Data** module to an experiment that inputs the data from the on-premises SQL Server database.

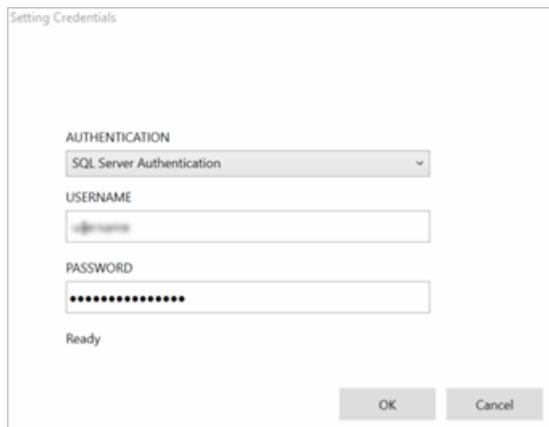
1. In Machine Learning Studio, select the **EXPERIMENTS** tab, click **+NEW** in the lower-left corner, and select **Blank Experiment** (or select one of several sample experiments available).
2. Find and drag the **Import Data** module to the experiment canvas.
3. Click **Save as** below the canvas. Enter "Azure Machine Learning On-Premises SQL Server Tutorial" for the experiment name, select the workspace, and click the **OK** check mark.



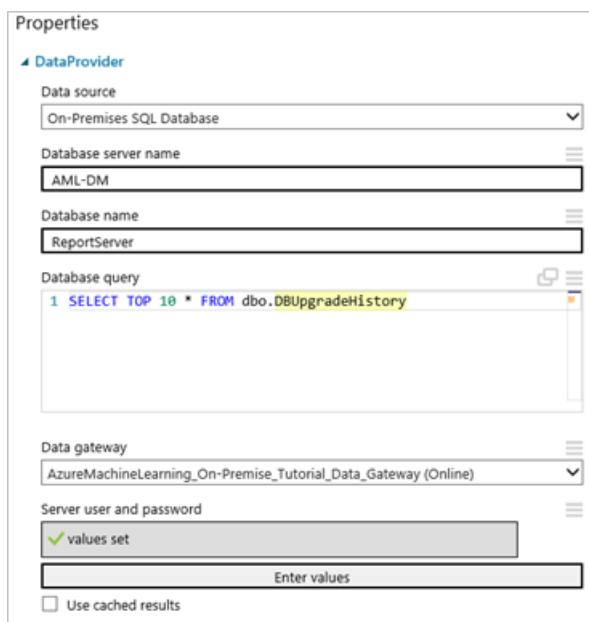
4. Click the **Import Data** module to select it, then in the **Properties** pane to the right of the canvas, select "On-Premises SQL Database" in the **Data source** dropdown list.
5. Select the **Data gateway** you installed and registered. You can set up another gateway by selecting "(add new Data Gateway...)".



6. Enter the SQL **Database server name** and **Database name**, along with the SQL **Database query** you want to execute.
7. Click **Enter values** under **User name and password** and enter your database credentials. You can use Windows Integrated Authentication or SQL Server Authentication depending upon how your on-premises SQL Server is configured.



The message "values required" changes to "values set" with a green check mark. You only need to enter the credentials once unless the database information or password changes. Azure Machine Learning uses the certificate you provided when you installed the gateway to encrypt the credentials in the cloud. Azure never stores on-premises credentials without encryption.



8. Click **RUN** to run the experiment.

Once the experiment finishes running, you can visualize the data you imported from the database by clicking the output port of the **Import Data** module and selecting **Visualize**.

Once you finish developing your experiment, you can deploy and operationalize your model. Using the Batch Execution Service, data from the on-premises SQL Server database configured in the **Import Data** module will be read and used for scoring. While you can use the Request Response Service for scoring on-premises data, Microsoft recommends using the [Excel Add-in](#) instead. Currently, writing to an on-premises SQL Server database through **Export Data** is not supported either in your experiments or published web services.

Application Lifecycle Management in Azure Machine Learning Studio

9/25/2017 • 6 min to read • [Edit Online](#)

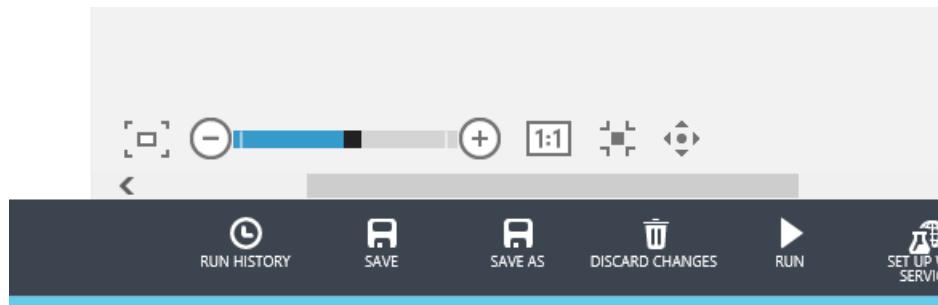
Azure Machine Learning Studio is a tool for developing machine learning experiments that are operationalized in the Azure cloud platform. It is like the Visual Studio IDE and scalable cloud service merged into a single platform. You can incorporate standard Application Lifecycle Management (ALM) practices, from versioning various assets to automated execution and deployment, into Azure Machine Learning Studio. This article discusses some of the options and approaches.

Versioning experiment

There are two recommended ways to version your experiments. You can either rely on built-in run history, or export the experiment in JavaScript Object Notation (JSON) format and manage it externally. Each approach comes with its pros and cons.

Experiment snapshots using Run History

In the execution model of the Azure Machine Learning Studio learning experiment, every time you click the **Run** button in the experiment editor, an immutable snapshot of the experiment is submitted to the job scheduler. You can view this list of snapshots by clicking the **Run History** button on the command bar in the experiment editor view.



You can then open the snapshot in Locked mode by clicking the name of the experiment at the time the experiment was submitted to run and the snapshot was taken. Notice that only the first item in the list, which represents the current experiment, is in an Editable state. Also notice that each snapshot can be in various Status states as well, including Finished (Partial run), Failed, Failed (Partial run), or Draft.

add many rows

NAME	STATE	STATUS	START TIME	END TIME
Add Many Rows	Editable	Draft	-	-
Add Many Rows	Locked	Finished	10/7/2016 12:14:34 AM	10/7/2016 12:15:06 AM
Add Many Rows	Locked	Finished (Partial run)	10/7/2016 12:13:45 AM	10/7/2016 12:13:48 AM
Add Many Rows	Locked	Finished	10/7/2016 12:13:04 AM	10/7/2016 12:13:15 AM
Add Many Rows	Locked	Finished	10/7/2016 12:11:59 AM	10/7/2016 12:12:23 AM
Add Many Rows	Locked	Finished	10/7/2016 12:10:35 AM	10/7/2016 12:10:52 AM

After it's opened, you can save the snapshot experiment as a new experiment and then modify it. If your experiment snapshot contains assets such as a trained model, transform, or dataset that have updated versions, the snapshot retains the references to the original version when the snapshot was taken. If you save the locked snapshot as a

new experiment, Azure Machine Learning Studio detects the existence of a newer version of these assets, and automatically updates them in the new experiment.

If you delete the experiment, all snapshots of that experiment are deleted.

Export/import experiment in JSON format

The run history snapshots keep an immutable version of the experiment in Azure Machine Learning Studio every time it is submitted to run. You can also save a local copy of the experiment and check it in to your favorite source control system, such as Team Foundation Server, and later on re-create an experiment from that local file. You can use the [Azure Machine Learning PowerShell](#) commandlets *Export-AmlExperimentGraph* and *Import-AmlExperimentGraph* to accomplish that.

The JSON file is a textual representation of the experiment graph, which might include a reference to assets in the workspace such as a dataset or a trained model. It doesn't contain a serialized version of the asset. If you attempt to import the JSON document back into the workspace, the referenced assets must already exist with the same asset IDs that are referenced in the experiment. Otherwise you will not be able to access the imported experiment.

Versioning trained model

A trained model in Azure Machine Learning is serialized into a format known as an .iLearner file, and is stored in the Azure Blob storage account associated with the workspace. One way to get a copy of the .iLearner file is through the retraining API. [This article](#) explains how the retraining API works. The high-level steps:

1. Set up your training experiment.
2. Add a web service output port to the Train Model module, or the module that produces the trained model, such as Tune Model Hyperparameter or Create R Model.
3. Run your training experiment and then deploy it as a model training web service.
4. Call the BES endpoint of the training web service, and specify the desired .iLearner file name and Blob storage account location where it will be stored.
5. Harvest the produced .iLearner file after the BES call finishes.

Another way to retrieve the .iLearner file is through the PowerShell commandlet *Download-AmlExperimentNodeOutput*. This might be easier if you just want to get a copy of the .iLearner file without the need to retrain the model programmatically.

After you have the .iLearner file containing the trained model, you can then employ your own versioning strategy. The strategy can be as simple as applying a pre/postfix as a naming convention and just leaving the .iLearner file in Blob storage, or copying/importing it into your version control system.

The saved .iLearner file can then be used for scoring through deployed web services.

Versioning web service

You can deploy two types of web services from an Azure Machine Learning experiment. The classic web service is tightly coupled with the experiment as well as the workspace. The new web service uses the Azure Resource Manager framework, and it is no longer coupled with the original experiment or the workspace.

Classic web service

To version a classic web service, you can take advantage of the web service endpoint construct. Here is a typical flow:

1. From your predictive experiment, you deploy a new classic web service, which contains a default endpoint.
2. You create a new endpoint named ep2, which exposes the current version of the experiment/trained model.
3. You go back and update your predictive experiment and trained model.
4. You redeploy the predictive experiment, which will then update the default endpoint. But this will not alter ep2.

5. You create an additional endpoint named ep3, which exposes the new version of the experiment and trained model.
6. Go back to step 3 if needed.

Over time, you might have many endpoints created in the same web service. Each endpoint represents a point-in-time copy of the experiment containing the point-in-time version of the trained model. You can then use external logic to determine which endpoint to call, which effectively means selecting a version of the trained model for the scoring run.

You can also create many identical web service endpoints, and then patch different versions of the .iLearner file to the endpoint to achieve similar effect. [This article](#) explains in more detail how to accomplish that.

New web service

If you create a new Azure Resource Manager-based web service, the endpoint construct is no longer available. Instead, you can generate web service definition (WSD) files, in JSON format, from your predictive experiment by using the [Export-AmlWebServiceDefinitionFromExperiment](#) PowerShell commandlet, or by using the [Export-AzureRmMlWebservice](#) PowerShell commandlet from a deployed Resource Manager-based web service.

After you have the exported WSD file and version control it, you can also deploy the WSD as a new web service in a different web service plan in a different Azure region. Just make sure you supply the proper storage account configuration as well as the new web service plan ID. To patch in different .iLearner files, you can modify the WSD file and update the location reference of the trained model, and deploy it as a new web service.

Automate experiment execution and deployment

An important aspect of ALM is to be able to automate the execution and deployment process of the application. In Azure Machine Learning, you can accomplish this by using the [PowerShell module](#). Here is an example of end-to-end steps that are relevant to a standard ALM automated execution/deployment process by using the [Azure Machine Learning Studio PowerShell module](#). Each step is linked to one or more PowerShell commandlets that you can use to accomplish that step.

1. [Upload a dataset](#).
2. Copy a training experiment into the workspace from a [workspace](#) or from [Gallery](#), or [import](#) an [exported](#) experiment from local disk.
3. [Update the dataset](#) in the training experiment.
4. [Run the training experiment](#).
5. [Promote the trained model](#).
6. [Copy a predictive experiment](#) into the workspace.
7. [Update the trained model](#) in the predictive experiment.
8. [Run the predictive experiment](#).
9. [Deploy a web service](#) from the predictive experiment.
10. Test the web service [RRS](#) or [BES](#) endpoint.

Next steps

- Download the [Azure Machine Learning Studio PowerShell](#) module and start to automate your ALM tasks.
- Learn how to [create and manage large number of ML models by using just a single experiment](#) through PowerShell and retraining API.
- Learn more about [deploying Azure Machine Learning web services](#).

Manage experiment iterations in Azure Machine Learning Studio

9/25/2017 • 4 min to read • [Edit Online](#)

Developing a predictive analysis model is an iterative process - as you modify the various functions and parameters of your experiment, your results converge until you are satisfied that you have a trained, effective model. Key to this process is tracking the various iterations of your experiment parameters and configurations.

NOTE

Try Azure Machine Learning for free

No credit card or Azure subscription required. [Get started now.](#)

You can review previous runs of your experiments at any time in order to challenge, revisit, and ultimately either confirm or refine previous assumptions. When you run an experiment, Machine Learning Studio keeps a history of the run, including dataset, module, and port connections and parameters. This history also captures results, runtime information such as start and stop times, log messages, and execution status. You can look back at any of these runs at any time to review the chronology of your experiment and intermediate results. You can even use a previous run of your experiment to launch into a new phase of inquiry and discovery on your path to creating simple, complex, or even ensemble modeling solutions.

NOTE

When you view a previous run of an experiment, that version of the experiment is locked and can't be edited. You can, however, save a copy of it by clicking **SAVE AS** and providing a new name for the copy. Machine Learning Studio opens the new copy, which you can then edit and run. This copy of your experiment is available in the **EXPERIMENTS** list along with all your other experiments.

Viewing the Prior Run

When you have an experiment open that you have run at least once, you can view the preceding run of the experiment by clicking **Prior Run** in the properties pane.

For example, suppose you create an experiment and run versions of it at 11:23, 11:42, and 11:55. If you open the last run of the experiment (11:55) and click **Prior Run**, the version ran at 11:42 is opened.

Viewing the Run History

You can view all the previous runs of an experiment by clicking **View Run History** in an open experiment.

For example, suppose you create an experiment with the [Linear Regression](#) module and you want to observe the effect of changing the value of **Learning rate** on your experiment results. You run the experiment multiple times with different values for this parameter, as follows:

LEARNING RATE VALUE	RUN START TIME
0.1	9/11/2014 4:18:58 pm

LEARNING RATE VALUE	RUN START TIME
0.2	9/11/2014 4:24:33 pm
0.4	9/11/2014 4:28:36 pm
0.5	9/11/2014 4:33:31 pm

If you click **VIEW RUN HISTORY**, you see a list of all these runs:

example experiment					
NAME	STATE	STATUS	START TIME	END TIME	
Example Experiment	Editable	Finished	9/11/2014 4:32:58 PM	9/11/2014 4:33:31 PM	
Example Experiment	Locked	Finished	9/11/2014 4:32:58 PM	9/11/2014 4:33:31 PM	
Example Experiment	Locked	Finished	9/11/2014 4:28:04 PM	9/11/2014 4:28:36 PM	
Example Experiment	Locked	Finished	9/11/2014 4:24:05 PM	9/11/2014 4:24:33 PM	
Example Experiment	Locked	Finished	9/11/2014 4:18:10 PM	9/11/2014 4:18:58 PM	

Click any of these runs to view a snapshot of the experiment at the time you ran it. The configuration, parameter values, comments, and results are all preserved to give you a complete record of that run of your experiment.

TIP

To document your iterations of the experiment, you can modify the title each time you run it, you can update the **Summary** of the experiment in the properties pane, and you can add or update comments on individual modules to record your changes. The title, summary, and module comments are saved with each run of the experiment.

The list of experiments in the **EXPERIMENTS** tab in Machine Learning Studio always displays the latest version of an experiment. If you open a previous run of the experiment (using **Prior Run** or **VIEW RUN HISTORY**), you can return to the draft version by clicking **VIEW RUN HISTORY** and selecting the iteration that has a **STATE** of **Editable**.

Iterating on a Previous Run

When you click **Prior Run** or **VIEW RUN HISTORY** and open a previous run, you can view a finished experiment in read-only mode.

If you want to begin an iteration of your experiment starting with the way you configured it for a previous run, you can do this by opening the run and clicking **SAVE AS**. This creates a new experiment, with a new title, an empty run history, and all the components and parameter values of the previous run. This new experiment is listed in the **EXPERIMENTS** tab in the Machine Learning Studio home page, and you can modify and run it, initiating a new run history for this iteration of your experiment.

For example, suppose you have the experiment run history shown in the previous section. You want to observe what happens when you set the **Learning rate** parameter to 0.4, and try different values for the **Number of training epochs** parameter.

1. Click **VIEW RUN HISTORY** and open the iteration of the experiment that you ran at 4:28:36 pm (in which you set the parameter value to 0.4).
2. Click **SAVE AS**.

3. Enter a new title and click the **OK** checkmark. A new copy of the experiment is created.
4. Modify the **Number of training epochs** parameter.
5. Click **RUN**.

You can now continue to modify and run this version of your experiment, building a new run history to record your work.

Create many Machine Learning models and web service endpoints from one experiment using PowerShell

9/25/2017 • 9 min to read • [Edit Online](#)

Here's a common machine learning problem: You want to create many models that have the same training workflow and use the same algorithm, but have different training datasets as input. This article shows you how to do this at scale in Azure Machine Learning Studio using just a single experiment.

For example, let's say you own a global bike rental franchise business. You want to build a regression model to predict the rental demand based on historic data. You have 1,000 rental locations across the world and you've collected a dataset for each location that includes important features such as date, time, weather, and traffic that are specific to each location.

You could train your model once using a merged version of all the datasets across all locations. But because each of your locations has a unique environment, a better approach would be to train your regression model separately using the dataset for each location. That way, each trained model could take into account the different store sizes, volume, geography, population, bike-friendly traffic environment, etc..

That may be the best approach, but you don't want to create 1,000 training experiments in Azure Machine Learning with each one representing a unique location. Besides being an overwhelming task, it's also seems pretty inefficient since each experiment would have all the same components except for the training dataset.

Fortunately, we can accomplish this by using the [Azure Machine Learning retraining API](#) and automating the task with [Azure Machine Learning PowerShell](#).

NOTE

To make our sample run faster, we'll reduce the number of locations from 1,000 to 10. But the same principles and procedures apply to 1,000 locations. The only difference is that if you want to train from 1,000 datasets you probably want to think of running the following PowerShell scripts in parallel. How to do that is beyond the scope of this article, but you can find examples of PowerShell multi-threading on the Internet.

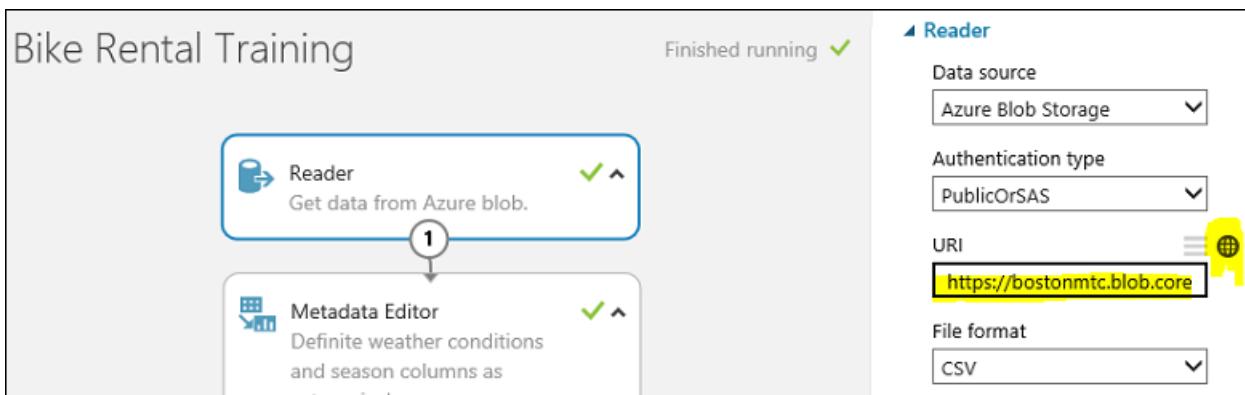
Set up the training experiment

We're going to use an example [training experiment](#) that we've already created in the [Cortana Intelligence Gallery](#). Open this experiment in your [Azure Machine Learning Studio](#) workspace.

NOTE

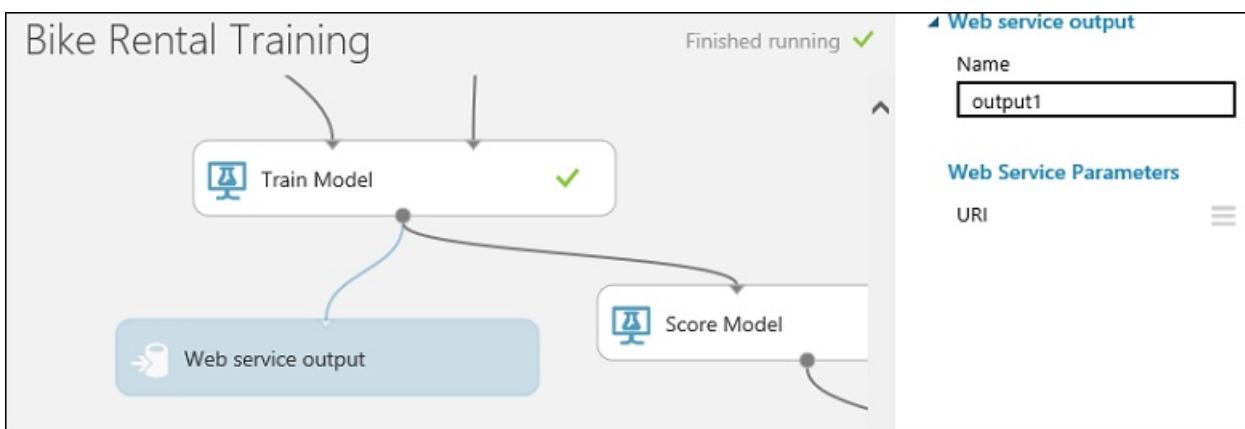
In order to follow along with this example, you may want to use a standard workspace rather than a free workspace. We'll be creating one endpoint for each customer - for a total of 10 endpoints - and that will require a standard workspace since a free workspace is limited to 3 endpoints. If you only have a free workspace, just modify the scripts below to allow for only 3 locations.

The experiment uses an **Import Data** module to import the training dataset *customer001.csv* from an Azure storage account. Let's assume we have collected training datasets from all bike rental locations and stored them in the same blob storage location with file names ranging from *rentalloc001.csv* to *rentalloc10.csv*.



Note that a **Web Service Output** module has been added to the **Train Model** module. When this experiment is deployed as a web service, the endpoint associated with that output will return the trained model in the format of a .ilearner file.

Also note that we set up a web service parameter for the URL that the **Import Data** module uses. This allows us to use the parameter to specify individual training datasets to train the model for each location. There are other ways we could have done this, such as using a SQL query with a web service parameter to get data from a SQL Azure database, or simply using a **Web Service Input** module to pass in a dataset to the web service.



Now, let's run this training experiment using the default value *rental001.csv* as the training dataset. If you view the output of the **Evaluate** module (click the output and select **Visualize**), you can see we get a decent performance of $AUC = 0.91$. At this point, we're ready to deploy a web service out of this training experiment.

Deploy the training and scoring web services

To deploy the training web service, click the **Set Up Web Service** button below the experiment canvas and select **Deploy Web Service**. Call this web service ""Bike Rental Training"".

Now we need to deploy the scoring web service. To do this, we can click **Set Up Web Service** below the canvas and select **Predictive Web Service**. This creates a scoring experiment. We'll need to make a few minor adjustments to make it work as a web service, such as removing the label column "cnt" from the input data and limiting the output to only the instance id and the corresponding predicted value.

To save yourself that work, you can simply open the [predictive experiment](#) in the Gallery that's already been prepared.

To deploy the web service, run the predictive experiment, then click the **Deploy Web Service** button below the canvas. Name the scoring web service "Bike Rental Scoring"".

Create 10 identical web service endpoints with PowerShell

This web service comes with a default endpoint. But we're not as interested in the default endpoint since it can't be updated. What we need to do is to create 10 additional endpoints, one for each location. We'll do this with

PowerShell.

First, we set up our PowerShell environment:

```
Import-Module .\AzureMLPS.dll
# Assume the default configuration file exists and is properly set to point to the valid Workspace.
$scoringSvc = Get-AmlWebService | where Name -eq 'Bike Rental Scoring'
$trainingSvc = Get-AmlWebService | where Name -eq 'Bike Rental Training'
```

Then, run the following PowerShell command:

```
# Create 10 endpoints on the scoring web service.
For ($i = 1; $i -le 10; $i++){
    $seq = $i.ToString().PadLeft(3, '0');
    $endpointName = 'rentalloc' + $seq;
    Write-Host ('adding endpoint ' + $endpointName + '...')
    Add-AmlWebServiceEndpoint -WebServiceId $scoringSvc.Id -EndpointName $endpointName -Description $endpointName
}
```

Now we've created 10 endpoints and they all contain the same trained model trained on *customer001.csv*. You can view them in the Azure Management Portal.

endpoints				
Name	Version	Update Available	Max Con...	🔍
default ➔	2/11/2016 12:34:...	false	20	
rentalloc002	2/11/2016 12:34:...	false	4	
rentalloc001	2/11/2016 12:34:...	false	4	
rentalloc003	2/11/2016 12:34:...	false	4	
rentalloc004	2/11/2016 12:34:...	false	4	
rentalloc005	2/11/2016 12:34:...	false	4	
rentalloc006	2/11/2016 12:34:...	false	4	
rentalloc007	2/11/2016 12:34:...	false	4	
rentalloc008	2/11/2016 12:34:...	false	4	
rentalloc009	2/11/2016 12:34:...	false	4	
rentalloc010	2/11/2016 12:34:...	false	4	

Update the endpoints to use separate training datasets using PowerShell

The next step is to update the endpoints with models uniquely trained on each customer's individual data. But first we need to produce these models from the **Bike Rental Training** web service. Let's go back to the **Bike Rental Training** web service. We need to call its BES endpoint 10 times with 10 different training datasets in order to produce 10 different models. We'll use the **InvokeAmlWebServiceBESEndpoint** PowerShell cmdlet to do this.

You will also need to provide credentials for your blob storage account into `$configContent`, namely, at the fields `AccountName`, `AccountKey` and `RelativeLocation`. The `AccountName` can be one of your account names, as seen in the **Classic Azure Management Portal** (*Storage* tab). Once you click on a storage account, its `AccountKey` can be found by pressing the **Manage Access Keys** button at the bottom and copying the *Primary Access Key*. The `RelativeLocation` is the path relative to your storage where a new model will be stored. For instance, the path

`hai/retrain/bike_rental/` in the script below points to a container named `hai`, and `/retrain/bike_rental/` are subfolders. Currently, you cannot create subfolders through the portal UI, but there are [several Azure Storage Explorers](#) that allow you to do so. It is recommended that you create a new container in your storage to store the new trained models (.learner files) as follows: from your storage page, click on the **Add** button at the bottom and name it `retrain`. In summary, the necessary changes to the script below pertain to `AccountName`, `AccountKey` and `RelativeLocation` (`: "retrain/model" + $seq + '.learner'`).

```
# Invoke the retraining API 10 times
# This is the default (and the only) endpoint on the training web service
$trainingSvcEp = (Get-AmlWebServiceEndpoint -WebServiceProviderId $trainingSvc.Id)[0];
$submitJobRequestUrl = $trainingSvcEp.ApiLocation + '/jobs?api-version=2.0';
$apiKey = $trainingSvcEp.PrimaryKey;
For ($i = 1; $i -le 10; $i++){
    $seq = $i.ToString().PadLeft(3, '0');
    $inputFileName = 'https://bostonmtc.blob.core.windows.net/ai/retrain/bike_rental/BikeRental' + $seq + '.csv';
    $configContent = '{ "GlobalParameters": { "URI": "' + $inputFileName + '" }, "Outputs": { "output1": { "ConnectionString": "DefaultEndpointsProtocol=https;AccountName=<myaccount>;AccountKey=<mykey>", "RelativeLocation": "ai/retrain/bike_rental/model" + $seq + '.learner" } } }';
    Write-Host ('Training regression model on ' + $inputFileName + ' for rental location ' + $seq + '...');
    Invoke-AmlWebServiceBESEndpoint -JobConfigString $configContent -SubmitJobRequestUrl $submitJobRequestUrl -ApiKey $apiKey
}
```

NOTE

The BES endpoint is the only supported mode for this operation. RRS cannot be used for producing trained models.

As you can see above, instead of constructing 10 different BES job configuration json files, we dynamically create the config string instead and feed it to the `JobConfigString` parameter of the **InvokeAmlWebServiceBESEndpoint** cmdlet, since there is really no need to keep a copy on disk.

If everything goes well, after a while you should see 10 .learner files, from `model001.learner` to `model010.learner`, in your Azure storage account. Now we're ready to update our 10 scoring web service endpoints with these models using the **Patch-AmlWebServiceEndpoint** PowerShell cmdlet. Remember again that we can only patch the non-default endpoints we programmatically created earlier.

```
# Patch the 10 endpoints with respective .learner models
$baseLoc = 'http://bostonmtc.blob.core.windows.net/'
$sasToken = '<my_blob_sas_token>'
For ($i = 1; $i -le 10; $i++){
    $seq = $i.ToString().PadLeft(3, '0');
    $endpointName = 'rentalloc' + $seq;
    $relativeLoc = 'ai/retrain/bike_rental/model' + $seq + '.learner';
    Write-Host ('Patching endpoint ' + $endpointName + '...');
    Patch-AmlWebServiceEndpoint -WebServiceProviderId $scoringSvc.Id -EndpointName $endpointName -ResourceName 'Bike Rental [trained model]' -BaseLocation $baseLoc -RelativeLocation $relativeLoc -SasBlobToken $sasToken
}
```

This should run fairly quickly. When the execution finishes, we'll have successfully created 10 predictive web service endpoints, each containing a trained model uniquely trained on the dataset specific to a rental location, all from a single training experiment. To verify this, you can try calling these endpoints using the **InvokeAmlWebServiceRRSEndpoint** cmdlet, providing them with the same input data, and you should expect to see different prediction results since the models are trained with different training sets.

Full PowerShell script

Here's the listing of the full source code:

```

Import-Module .\AzureMLPS.dll
# Assume the default configuration file exists and properly set to point to the valid workspace.
$scoringSvc = Get-AmlWebService | where Name -eq 'Bike Rental Scoring'
$trainingSvc = Get-AmlWebService | where Name -eq 'Bike Rental Training'

# Create 10 endpoints on the scoring web service
For ($i = 1; $i -le 10; $i++){
    $seq = $i.ToString().PadLeft(3, '0');
    $endpointName = 'rentalloc' + $seq;
    Write-Host ('adding endpoint ' + $endpointName + '...')
    Add-AmlWebServiceEndpoint -WebServiceProviderId $scoringSvc.Id -EndpointName $endpointName -Description $endpointName
}

# Invoke the retraining API 10 times to produce 10 regression models in .ilearner format
$trainingSvcEp = (Get-AmlWebServiceEndpoint -WebServiceProviderId $trainingSvc.Id)[0];
$submitJobRequestUrl = $trainingSvcEp.ApiLocation + '/jobs?api-version=2.0';
$apiKey = $trainingSvcEp.PrimaryKey;
For ($i = 1; $i -le 10; $i++){
    $seq = $i.ToString().PadLeft(3, '0');
    $inputFileName = 'https://bostonmtc.blob.core.windows.net/hai/retrain/bike_rental/BikeRental' + $seq + '.csv';
    $configContent = '{"GlobalParameters": {"URI": "' + $inputFileName + '"}, "Outputs": {"output1": {"ConnectionString": "DefaultEndpointsProtocol=https;AccountName=<myaccount>;AccountKey=<mykey>", "RelativeLocation": "hai/retrain/bike_rental/model" + $seq + ".learner" }}}';
    Write-Host ('training regression model on ' + $inputFileName + ' for rental location ' + $seq + '...');
    Invoke-AmlWebServiceBESEndpoint -JobConfigString $configContent -SubmitJobRequestUrl $submitJobRequestUrl -ApiKey $apiKey
}

# Patch the 10 endpoints with respective .ilearner models
$baseLoc = 'http://bostonmtc.blob.core.windows.net/'
$sasToken = "?test"
For ($i = 1; $i -le 10; $i++){
    $seq = $i.ToString().PadLeft(3, '0');
    $endpointName = 'rentalloc' + $seq;
    $relativeLoc = 'hai/retrain/bike_rental/model' + $seq + '.learner';
    Write-Host ('Patching endpoint ' + $endpointName + '...');
    Patch-AmlWebServiceEndpoint -WebServiceProviderId $scoringSvc.Id -EndpointName $endpointName -ResourceName 'Bike Rental [trained model]' -BaseLocation $baseLoc -RelativeLocation $relativeLoc -SasBlobToken $sasToken
}

```

How to choose algorithms for Microsoft Azure Machine Learning

9/25/2017 • 15 min to read • [Edit Online](#)

The answer to the question "What machine learning algorithm should I use?" is always "It depends." It depends on the size, quality, and nature of the data. It depends on what you want to do with the answer. It depends on how the math of the algorithm was translated into instructions for the computer you are using. And it depends on how much time you have. Even the most experienced data scientists can't tell which algorithm will perform best before trying them.

The Machine Learning Algorithm Cheat Sheet

The **Microsoft Azure Machine Learning Algorithm Cheat Sheet** helps you choose the right machine learning algorithm for your predictive analytics solutions from the Microsoft Azure Machine Learning library of algorithms. This article walks you through how to use it.

NOTE

To download the cheat sheet and follow along with this article, go to [Machine learning algorithm cheat sheet for Microsoft Azure Machine Learning Studio](#).

This cheat sheet has a very specific audience in mind: a beginning data scientist with undergraduate-level machine learning, trying to choose an algorithm to start with in Azure Machine Learning Studio. That means that it makes some generalizations and oversimplifications, but it points you in a safe direction. It also means that there are lots of algorithms not listed here. As Azure Machine Learning grows to encompass a more complete set of available methods, we'll add them.

These recommendations are compiled feedback and tips from many data scientists and machine learning experts. We didn't agree on everything, but I've tried to harmonize our opinions into a rough consensus. Most of the statements of disagreement begin with "It depends..."

How to use the cheat sheet

Read the path and algorithm labels on the chart as "For *<path label>*, use *<algorithm>*." For example, "For *speed*, use *two class logistic regression*." Sometimes more than one branch applies. Sometimes none of them are a perfect fit. They're intended to be rule-of-thumb recommendations, so don't worry about it being exact. Several data scientists I talked with said that the only sure way to find the very best algorithm is to try all of them.

Here's an example from the [Cortana Intelligence Gallery](#) of an experiment that tries several algorithms against the same data and compares the results: [Compare Multi-class Classifiers: Letter recognition](#).

TIP

To download and print a diagram that gives an overview of the capabilities of Machine Learning Studio, see [Overview diagram of Azure Machine Learning Studio capabilities](#).

Flavors of machine learning

Supervised

Supervised learning algorithms make predictions based on a set of examples. For instance, historical stock prices can be used to hazard guesses at future prices. Each example used for training is labeled with the value of interest—in this case the stock price. A supervised learning algorithm looks for patterns in those value labels. It can use any information that might be relevant—the day of the week, the season, the company's financial data, the type of industry, the presence of disruptive geopolitical events—and each algorithm looks for different types of patterns. After the algorithm has found the best pattern it can, it uses that pattern to make predictions for unlabeled testing data—tomorrow's prices.

Supervised learning is a popular and useful type of machine learning. With one exception, all the modules in Azure Machine Learning are supervised learning algorithms. There are several specific types of supervised learning that are represented within Azure Machine Learning: classification, regression, and anomaly detection.

- **Classification.** When the data are being used to predict a category, supervised learning is also called classification. This is the case when assigning an image as a picture of either a 'cat' or a 'dog'. When there are only two choices, it's called **two-class** or **binomial classification**. When there are more categories, as when predicting the winner of the NCAA March Madness tournament, this problem is known as **multi-class classification**.
- **Regression.** When a value is being predicted, as with stock prices, supervised learning is called regression.
- **Anomaly detection.** Sometimes the goal is to identify data points that are simply unusual. In fraud detection, for example, any highly unusual credit card spending patterns are suspect. The possible variations are so numerous and the training examples so few, that it's not feasible to learn what fraudulent activity looks like. The approach that anomaly detection takes is to simply learn what normal activity looks like (using a history of non-fraudulent transactions) and identify anything that is significantly different.

Unsupervised

In unsupervised learning, data points have no labels associated with them. Instead, the goal of an unsupervised learning algorithm is to organize the data in some way or to describe its structure. This can mean grouping it into clusters or finding different ways of looking at complex data so that it appears simpler or more organized.

Reinforcement learning

In reinforcement learning, the algorithm gets to choose an action in response to each data point. The learning algorithm also receives a reward signal a short time later, indicating how good the decision was. Based on this, the algorithm modifies its strategy in order to achieve the highest reward. Currently there are no reinforcement learning algorithm modules in Azure Machine Learning. Reinforcement learning is common in robotics, where the set of sensor readings at one point in time is a data point, and the algorithm must choose the robot's next action. It is also a natural fit for Internet of Things applications.

Considerations when choosing an algorithm

Accuracy

Getting the most accurate answer possible isn't always necessary. Sometimes an approximation is adequate, depending on what you want to use it for. If that's the case, you may be able to cut your processing time dramatically by sticking with more approximate methods. Another advantage of more approximate methods is that they naturally tend to avoid [overfitting](#).

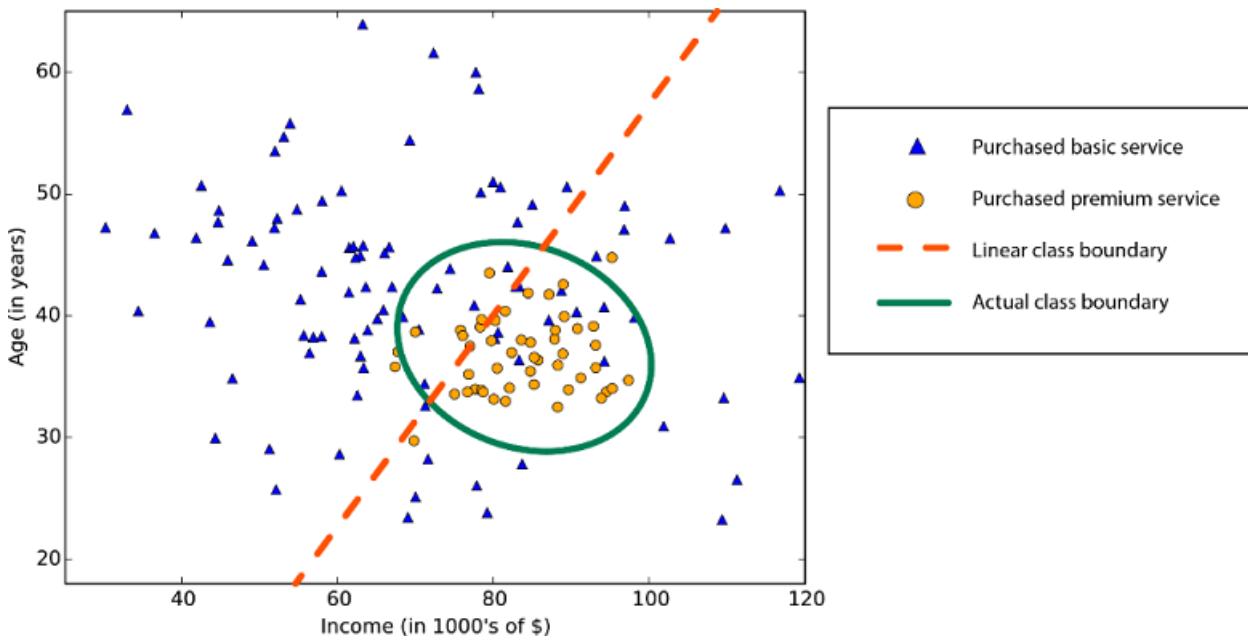
Training time

The number of minutes or hours necessary to train a model varies a great deal between algorithms. Training time is often closely tied to accuracy—one typically accompanies the other. In addition, some algorithms are more sensitive to the number of data points than others. When time is limited it can drive the choice of algorithm, especially when the data set is large.

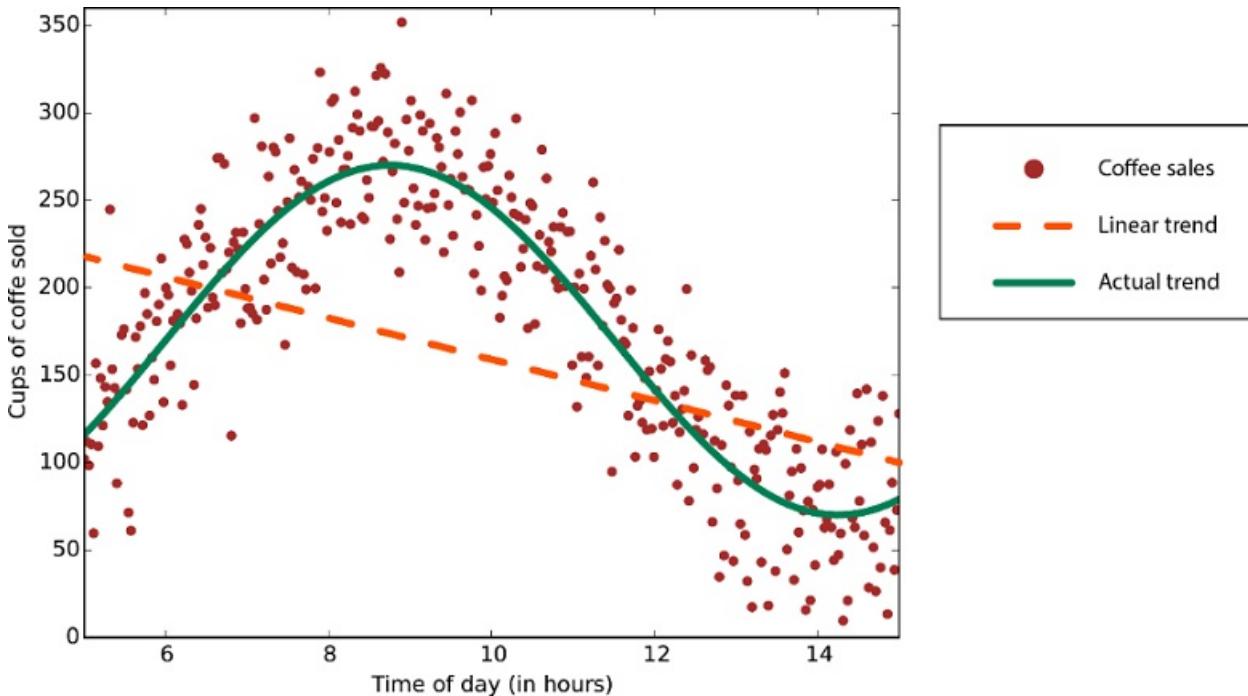
Linearity

Lots of machine learning algorithms make use of linearity. Linear classification algorithms assume that classes

can be separated by a straight line (or its higher-dimensional analog). These include logistic regression and support vector machines (as implemented in Azure Machine Learning). Linear regression algorithms assume that data trends follow a straight line. These assumptions aren't bad for some problems, but on others they bring accuracy down.



Non-linear class boundary - relying on a linear classification algorithm would result in low accuracy



Data with a nonlinear trend - using a linear regression method would generate much larger errors than necessary

Despite their dangers, linear algorithms are very popular as a first line of attack. They tend to be algorithmically simple and fast to train.

Number of parameters

Parameters are the knobs a data scientist gets to turn when setting up an algorithm. They are numbers that affect the algorithm's behavior, such as error tolerance or number of iterations, or options between variants of how the algorithm behaves. The training time and accuracy of the algorithm can sometimes be quite sensitive to getting just the right settings. Typically, algorithms with large numbers of parameters require the most trial and error to find a good combination.

Alternatively, there is a [parameter sweeping](#) module block in Azure Machine Learning that automatically tries all parameter combinations at whatever granularity you choose. While this is a great way to make sure you've spanned the parameter space, the time required to train a model increases exponentially with the number of parameters.

The upside is that having many parameters typically indicates that an algorithm has greater flexibility. It can often achieve very good accuracy. Provided you can find the right combination of parameter settings.

Number of features

For certain types of data, the number of features can be very large compared to the number of data points. This is often the case with genetics or textual data. The large number of features can bog down some learning algorithms, making training time unfeasibly long. Support Vector Machines are particularly well suited to this case (see below).

Special cases

Some learning algorithms make particular assumptions about the structure of the data or the desired results. If you can find one that fits your needs, it can give you more useful results, more accurate predictions, or faster training times.

ALGORITHM	ACCURACY	TRAINING TIME	LINEARITY	PARAMETERS	NOTES
Two-class classification					
logistic regression		●	●	5	
decision forest	●	○		6	
decision jungle	●	○		6	Low memory footprint
boosted decision tree	●	○		6	Large memory footprint
neural network	●			9	Additional customization is possible
averaged perceptron	○	○	●	4	
support vector machine		○	●	5	Good for large feature sets
locally deep support vector machine	○			8	Good for large feature sets
Bayes' point machine		○	●	3	
Multi-class classification					

ALGORITHM	ACCURACY	TRAINING TIME	LINEARITY	PARAMETERS	NOTES
logistic regression		●	●	5	
decision forest	●	○		6	
decision jungle	●	○		6	Low memory footprint
neural network	●			9	Additional customization is possible
one-v-all	-	-	-	-	See properties of the two-class method selected
Regression					
linear		●	●	4	
Bayesian linear		○	●	2	
decision forest	●	○		6	
boosted decision tree	●	○		5	Large memory footprint
fast forest quantile	●	○		9	Distributions rather than point predictions
neural network	●			9	Additional customization is possible
Poisson			●	5	Technically log-linear. For predicting counts
ordinal				0	For predicting rank-ordering
Anomaly detection					
support vector machine	○	○		2	Especially good for large feature sets
PCA-based anomaly detection		○	●	3	

ALGORITHM	ACCURACY	TRAINING TIME	LINEARITY	PARAMETERS	NOTES
K-means		○	●	4	A clustering algorithm

Algorithm properties:

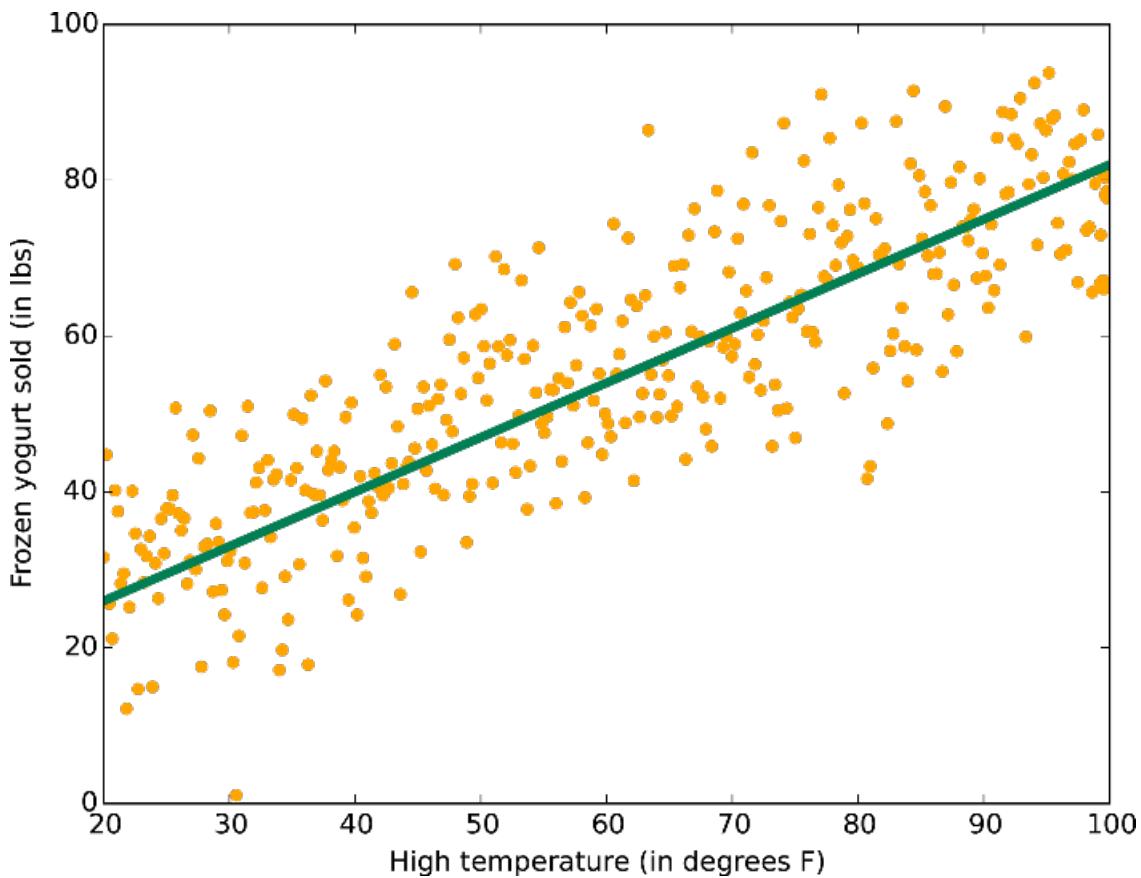
● - shows excellent accuracy, fast training times, and the use of linearity

○ - shows good accuracy and moderate training times

Algorithm notes

Linear regression

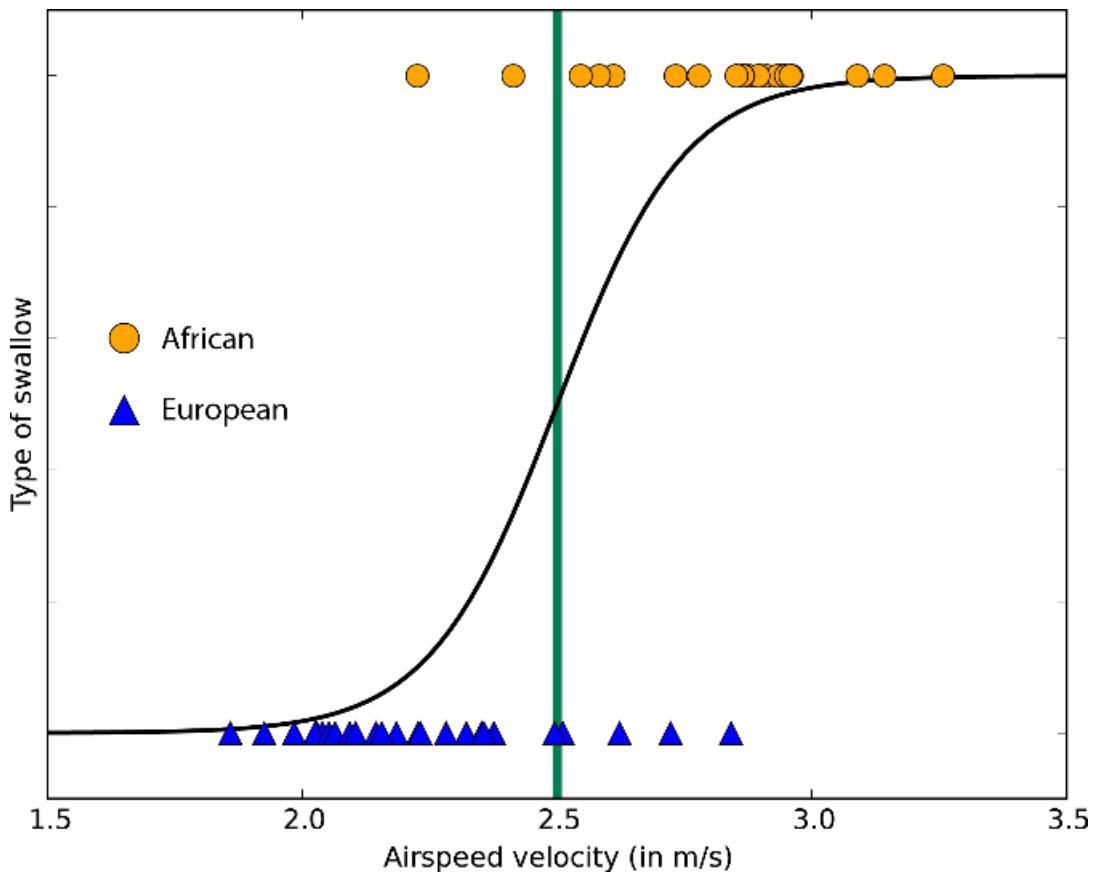
As mentioned previously, [linear regression](#) fits a line (or plane, or hyperplane) to the data set. It's a workhorse, simple and fast, but it may be overly simplistic for some problems. Check here for a [linear regression tutorial](#).



Data with a linear trend

Logistic regression

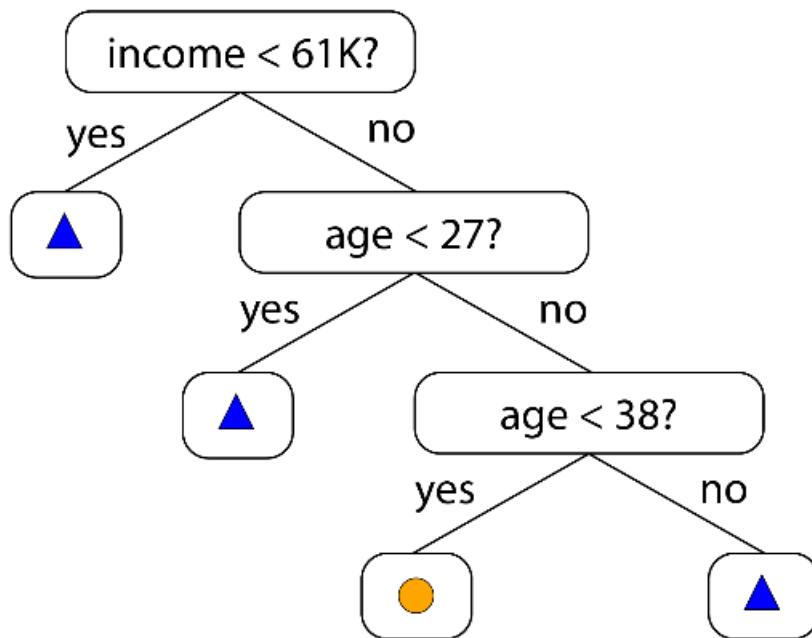
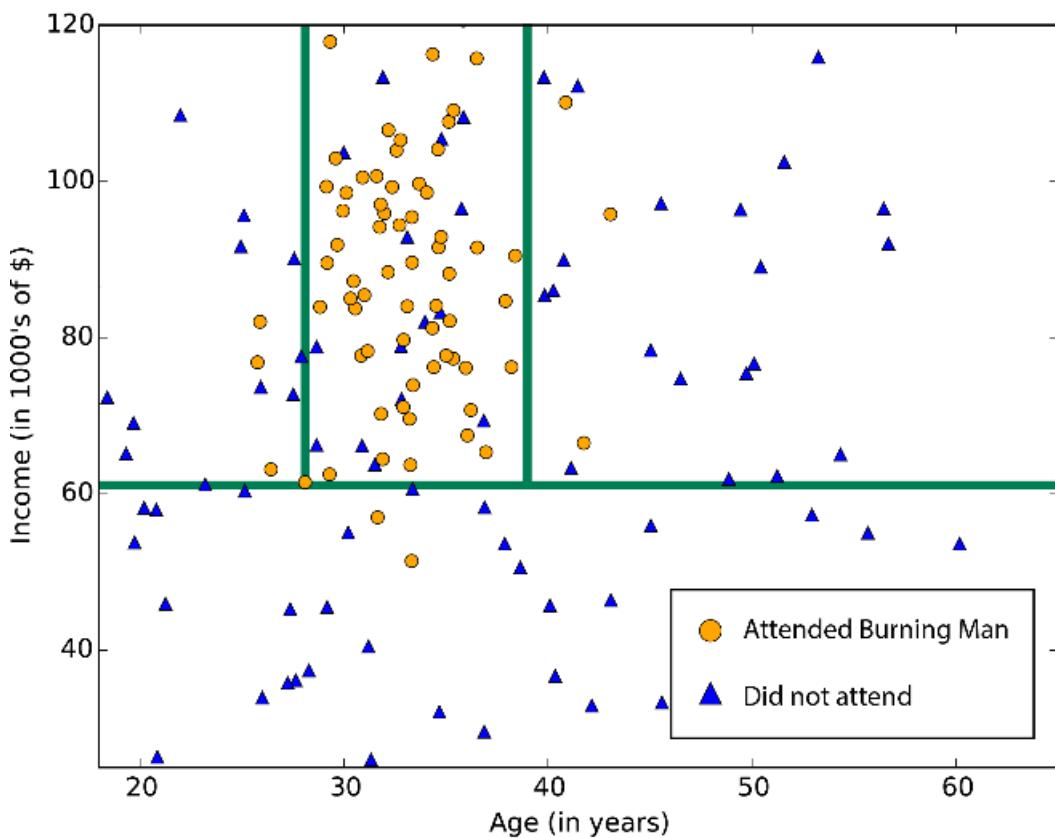
Although it confusingly includes 'regression' in the name, logistic regression is actually a powerful tool for [two-class](#) and [multiclass](#) classification. It's fast and simple. The fact that it uses an 'S'-shaped curve instead of a straight line makes it a natural fit for dividing data into groups. Logistic regression gives linear class boundaries, so when you use it, make sure a linear approximation is something you can live with.



A logistic regression to two-class data with just one feature - the class boundary is the point at which the logistic curve is just as close to both classes

Trees, forests, and jungles

Decision forests ([regression](#), [two-class](#), and [multiclass](#)), decision jungles ([two-class](#) and [multiclass](#)), and boosted decision trees ([regression](#) and [two-class](#)) are all based on decision trees, a foundational machine learning concept. There are many variants of decision trees, but they all do the same thing—subdivide the feature space into regions with mostly the same label. These can be regions of consistent category or of constant value, depending on whether you are doing classification or regression.



A decision tree subdivides a feature space into regions of roughly uniform values

Because a feature space can be subdivided into arbitrarily small regions, it's easy to imagine dividing it finely enough to have one data point per region. This is an extreme example of overfitting. In order to avoid this, a large set of trees are constructed with special mathematical care taken that the trees are not correlated. The average of this "decision forest" is a tree that avoids overfitting. Decision forests can use a lot of memory. Decision jungles are a variant that consumes less memory at the expense of a slightly longer training time.

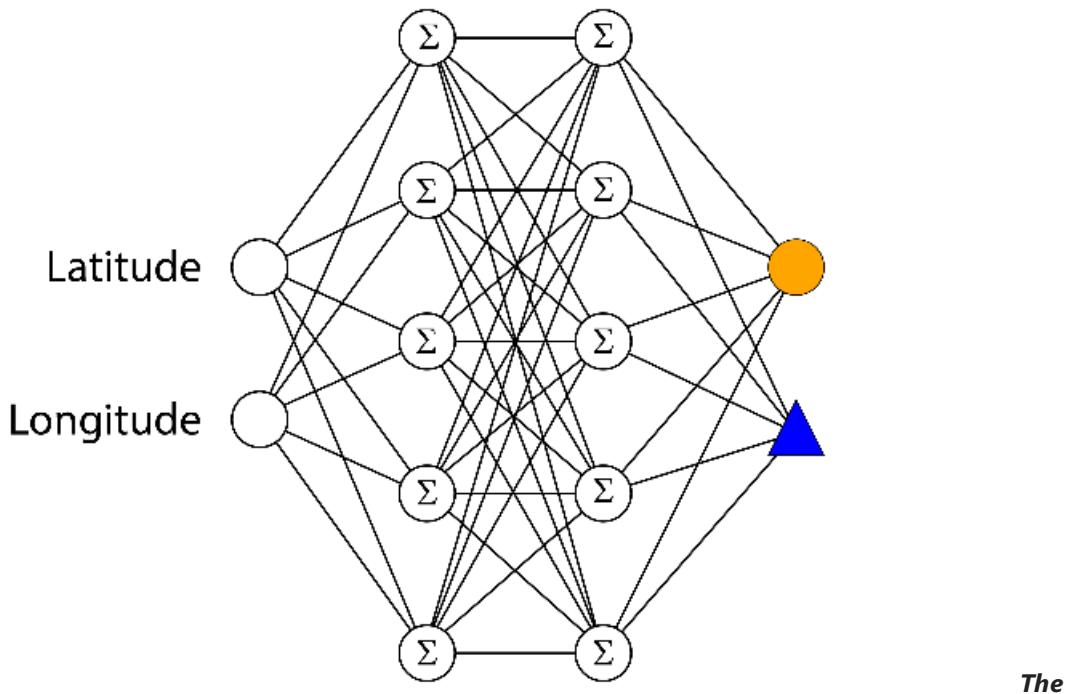
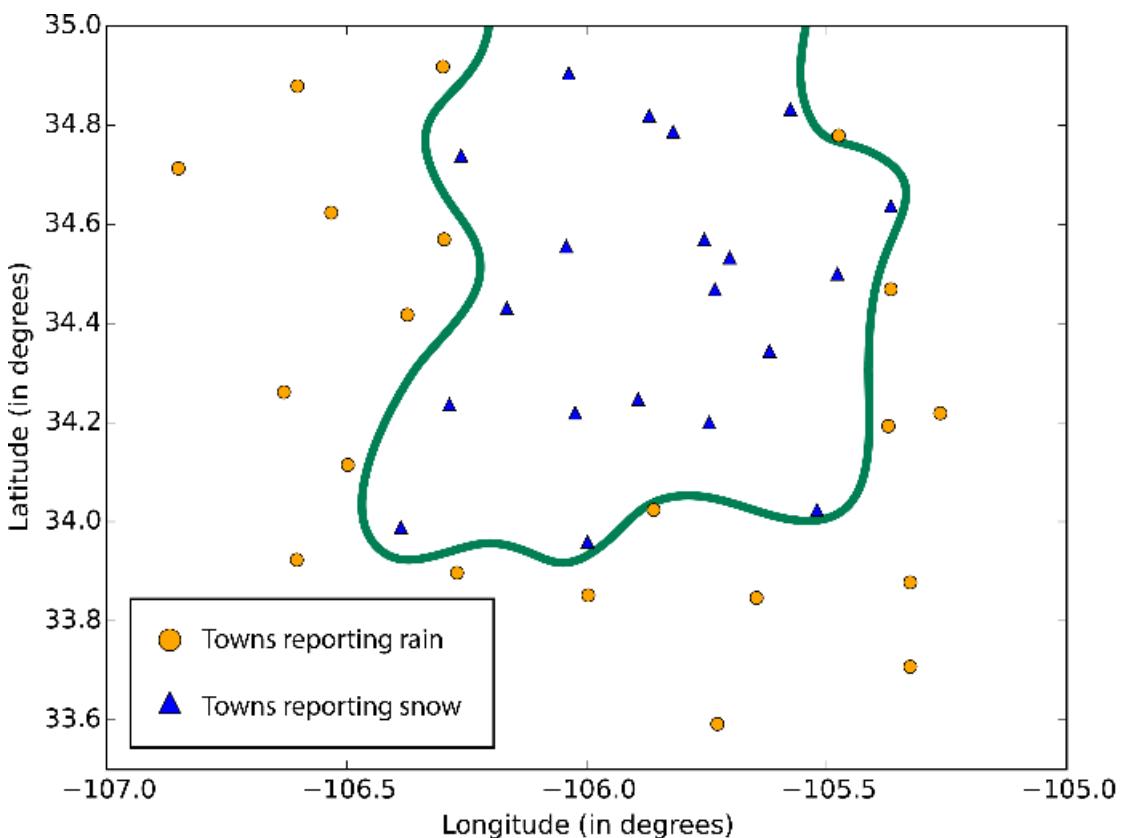
Boosted decision trees avoid overfitting by limiting how many times they can subdivide and how few data points are allowed in each region. The algorithm constructs a sequence of trees, each of which learns to compensate for the error left by the tree before. The result is a very accurate learner that tends to use a lot of memory. For the full technical description, check out [Friedman's original paper](#).

[Fast forest quantile regression](#) is a variation of decision trees for the special case where you want to know not only the typical (median) value of the data within a region, but also its distribution in the form of quantiles.

Neural networks and perceptrons

Neural networks are brain-inspired learning algorithms covering [multiclass](#), [two-class](#), and [regression](#) problems. They come in an infinite variety, but the neural networks within Azure Machine Learning are all of the form of directed acyclic graphs. That means that input features are passed forward (never backward) through a sequence of layers before being turned into outputs. In each layer, inputs are weighted in various combinations, summed, and passed on to the next layer. This combination of simple calculations results in the ability to learn sophisticated class boundaries and data trends, seemingly by magic. Many-layered networks of this sort perform the "deep learning" that fuels so much tech reporting and science fiction.

This high performance doesn't come for free, though. Neural networks can take a long time to train, particularly for large data sets with lots of features. They also have more parameters than most algorithms, which means that parameter sweeping expands the training time a great deal. And for those overachievers who wish to [specify their own network structure](#), the possibilities are inexhaustible.

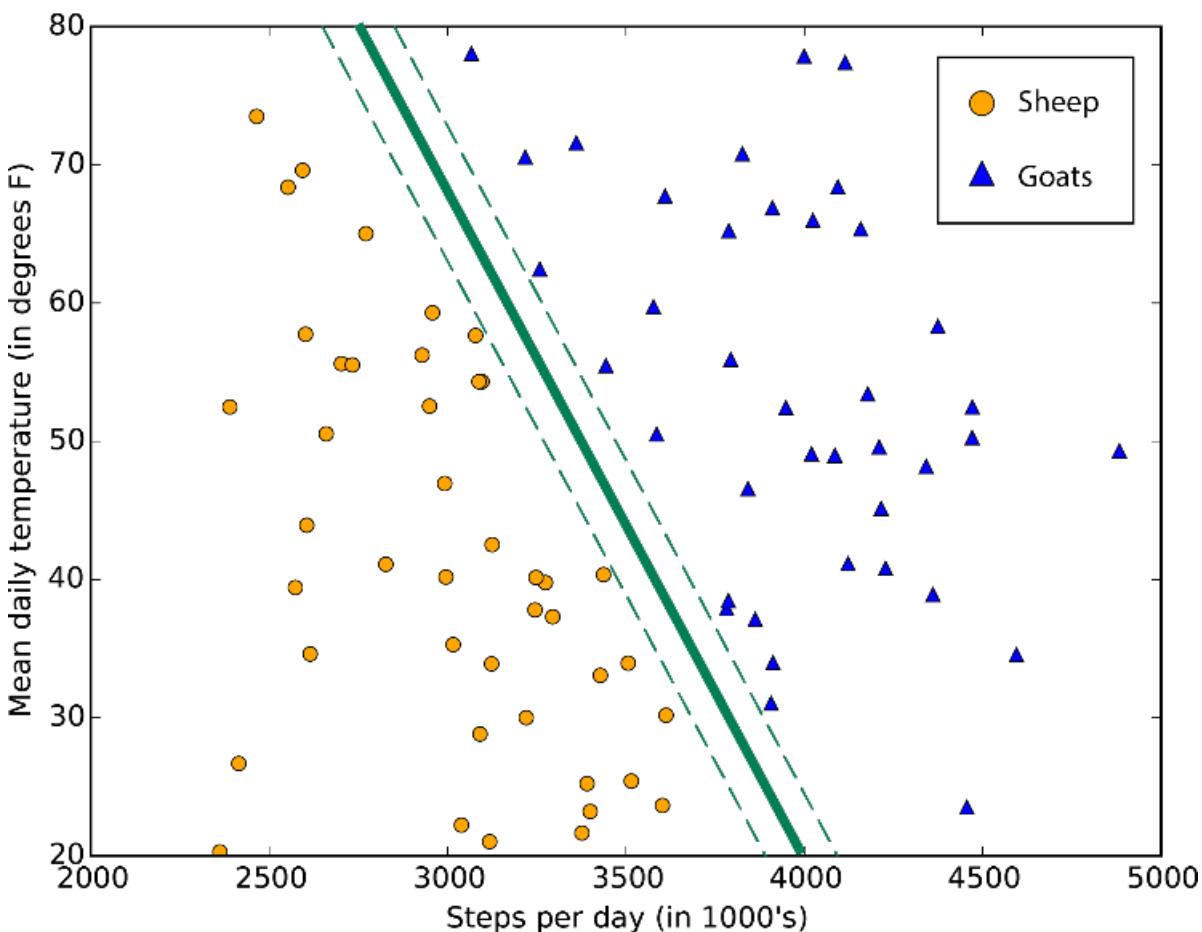


boundaries learned by neural networks can be complex and irregular

The [two-class averaged perceptron](#) is neural networks' answer to skyrocketing training times. It uses a network structure that gives linear class boundaries. It is almost primitive by today's standards, but it has a long history of working robustly and is small enough to learn quickly.

SVMs

Support vector machines (SVMs) find the boundary that separates classes by as wide a margin as possible. When the two classes can't be clearly separated, the algorithms find the best boundary they can. As written in Azure Machine Learning, the [two-class SVM](#) does this with a straight line only. (In SVM-speak, it uses a linear kernel.) Because it makes this linear approximation, it is able to run fairly quickly. Where it really shines is with feature-intense data, like text or genomic. In these cases SVMs are able to separate classes more quickly and with less overfitting than most other algorithms, in addition to requiring only a modest amount of memory.



A typical support vector machine class boundary maximizes the margin separating two classes

Another product of Microsoft Research, the [two-class locally deep SVM](#) is a non-linear variant of SVM that retains most of the speed and memory efficiency of the linear version. It is ideal for cases where the linear approach doesn't give accurate enough answers. The developers kept it fast by breaking down the problem into a bunch of small linear SVM problems. Read the [full description](#) for the details on how they pulled off this trick.

Using a clever extension of nonlinear SVMs, the [one-class SVM](#) draws a boundary that tightly outlines the entire data set. It is useful for anomaly detection. Any new data points that fall far outside that boundary are unusual enough to be noteworthy.

Bayesian methods

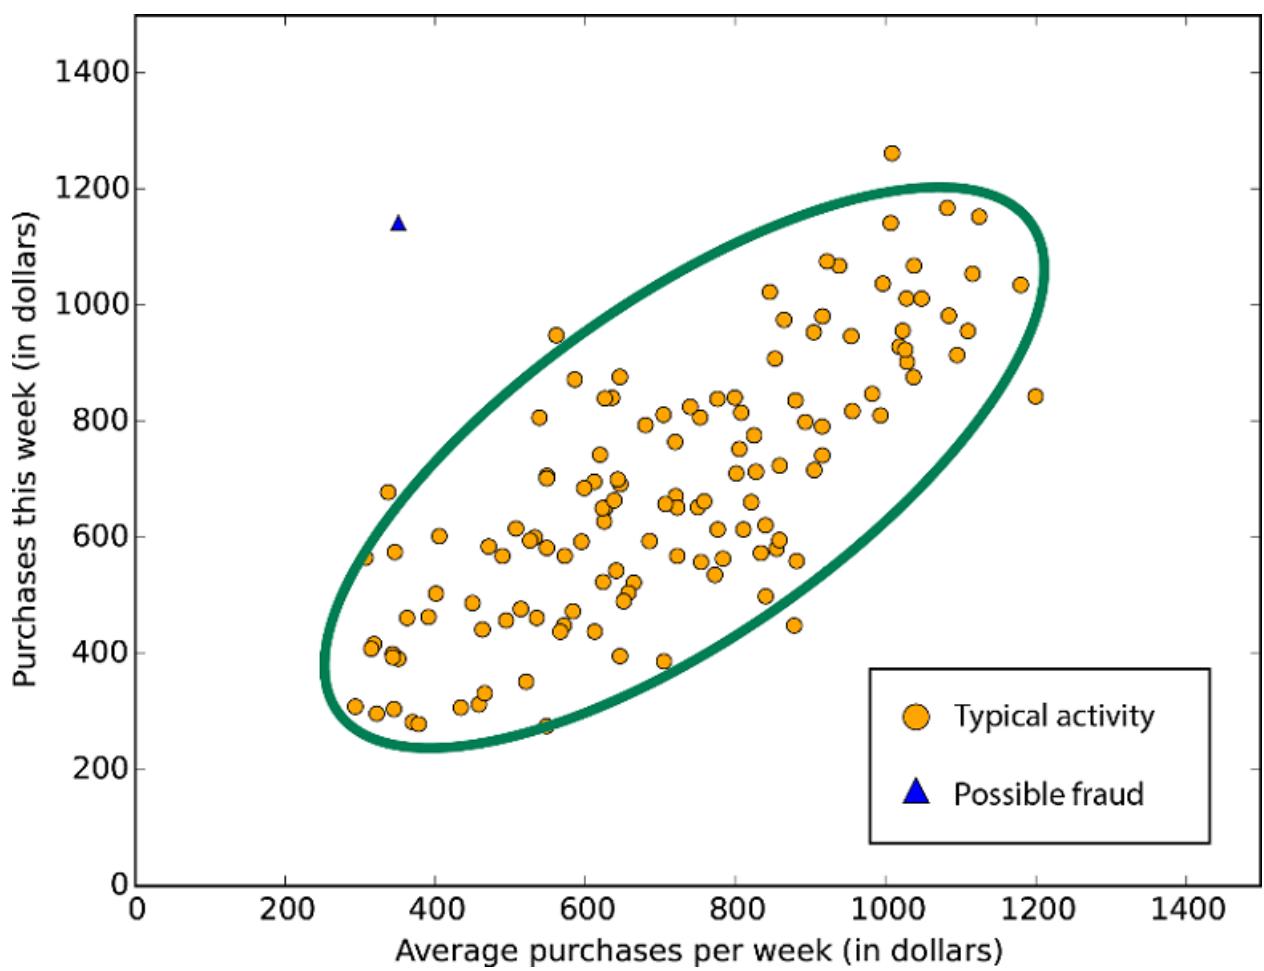
Bayesian methods have a highly desirable quality: they avoid overfitting. They do this by making some assumptions beforehand about the likely distribution of the answer. Another byproduct of this approach is that they have very few parameters. Azure Machine Learning has both Bayesian algorithms for both classification ([Two-class Bayes' point machine](#)) and regression ([Bayesian linear regression](#)). Note that these assume that the data can be split or fit with a straight line.

On a historical note, Bayes' point machines were developed at Microsoft Research. They have some exceptionally beautiful theoretical work behind them. The interested student is directed to the [original article in JMLR](#) and an [insightful blog by Chris Bishop](#).

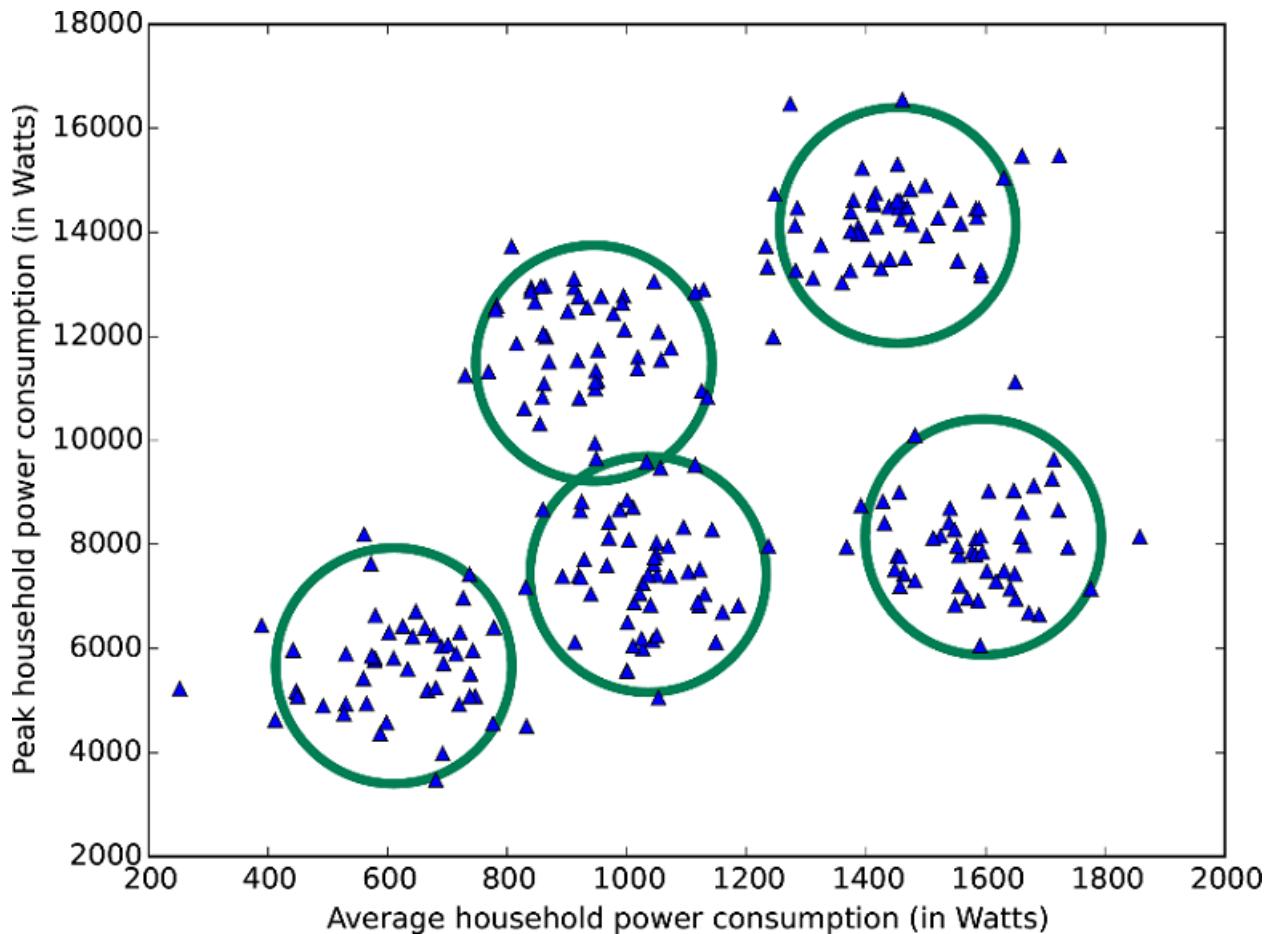
Specialized algorithms

If you have a very specific goal you may be in luck. Within the Azure Machine Learning collection, there are algorithms that specialize in:

- rank prediction ([ordinal regression](#)),
- count prediction ([Poisson regression](#)),
- anomaly detection (one based on [principal components analysis](#) and one based on [support vector machines](#))
- clustering ([K-means](#))

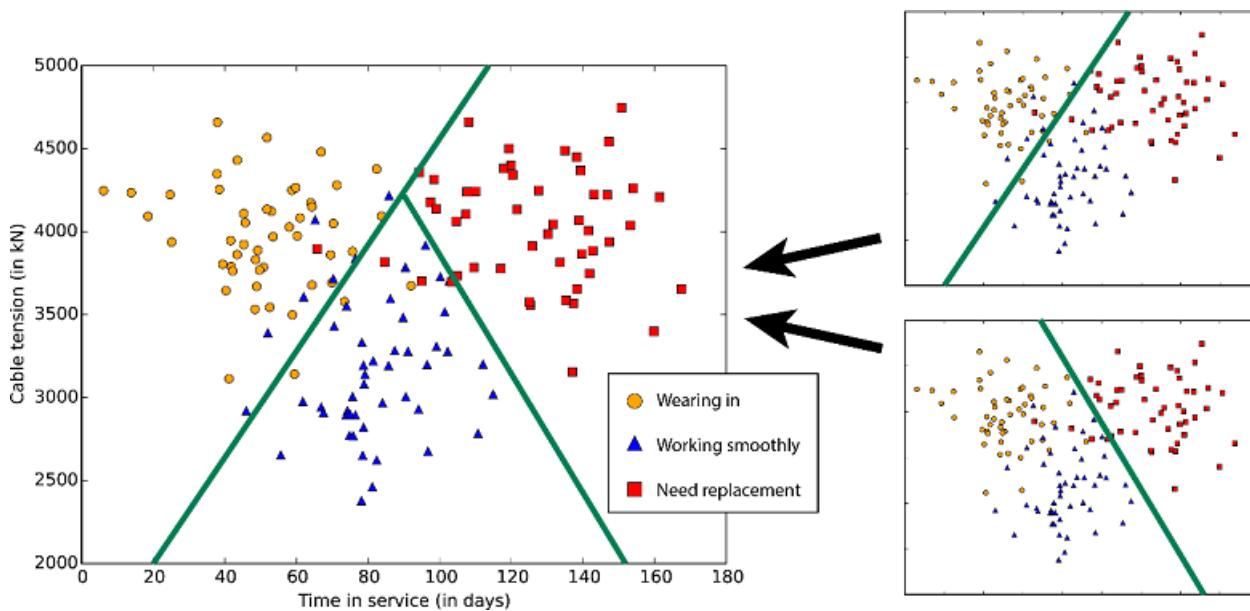


PCA-based anomaly detection - the vast majority of the data falls into a stereotypical distribution; points deviating dramatically from that distribution are suspect



A data set is grouped into five clusters using K-means

There is also an ensemble [one-v-all multiclass classifier](#), which breaks the N-class classification problem into N-1 two-class classification problems. The accuracy, training time, and linearity properties are determined by the two-class classifiers used.



A pair of two-class classifiers combine to form a three-class classifier

Azure Machine Learning also includes access to a powerful machine learning framework under the title of [Vowpal Wabbit](#). VW defies categorization here, since it can learn both classification and regression problems and can even learn from partially unlabeled data. You can configure it to use any one of a number of learning algorithms, loss functions, and optimization algorithms. It was designed from the ground up to be efficient, parallel, and extremely fast. It handles ridiculously large feature sets with little apparent effort. Started and led by Microsoft Research's own John Langford, VW is a Formula One entry in a field of stock car algorithms. Not every problem fits VW, but if yours does, it may be worth your while to climb the learning curve on its interface. It's also available as [stand-alone open source code](#) in several languages.

More help with algorithms

- For a downloadable infographic that describes algorithms and provides examples, see [Downloadable Infographic: Machine learning basics with algorithm examples](#).
- For a list by category of all the machine learning algorithms available in Azure Machine Learning Studio, see [Initialize Model](#) in the Machine Learning Studio Algorithm and Module Help.
- For a complete alphabetical list of algorithms and modules in Azure Machine Learning Studio, see [A-Z list of Machine Learning Studio modules](#) in Machine Learning Studio Algorithm and Module Help.
- To download and print a diagram that gives an overview of the capabilities of Azure Machine Learning Studio, see [Overview diagram of Azure Machine Learning Studio capabilities](#).

Machine learning algorithm cheat sheet for Microsoft Azure Machine Learning Studio

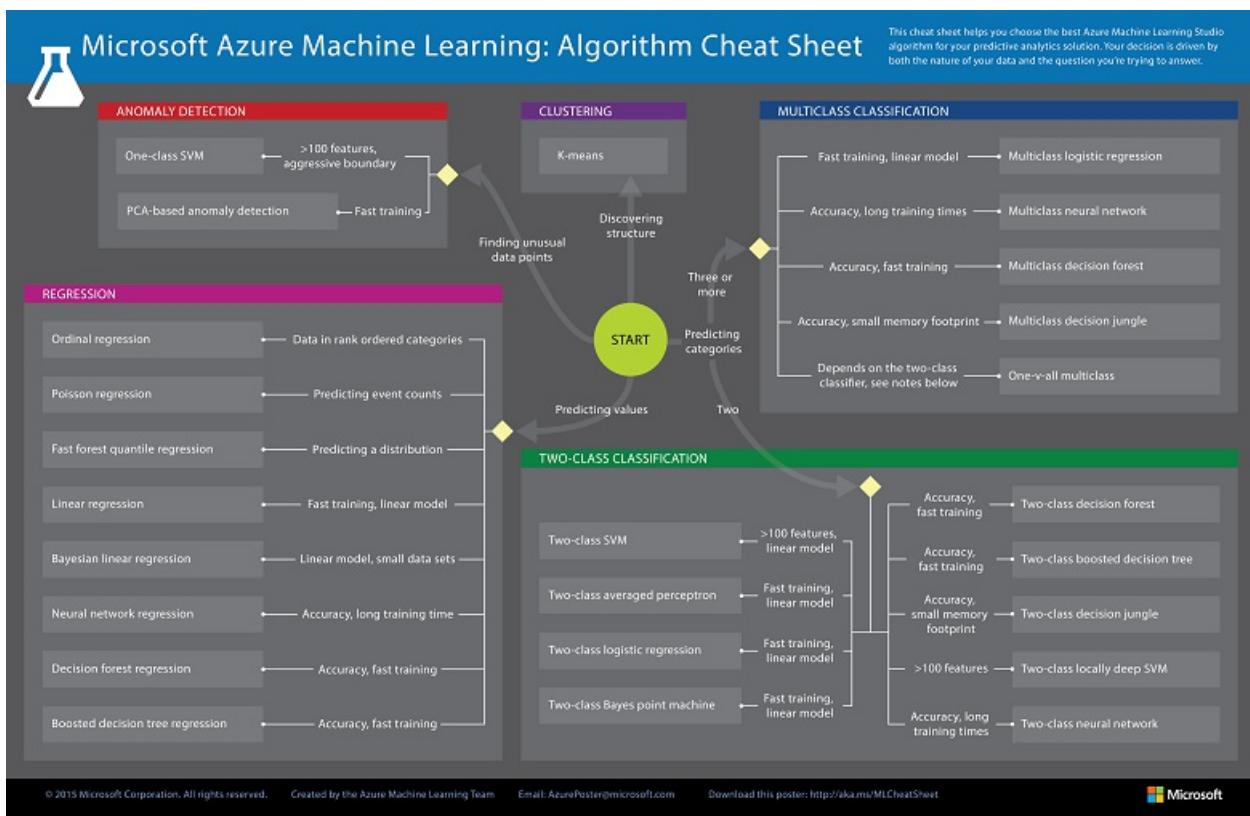
9/25/2017 • 5 min to read • [Edit Online](#)

The **Microsoft Azure Machine Learning Algorithm Cheat Sheet** helps you choose the right algorithm for a predictive analytics model.

Azure Machine Learning Studio has a large library of algorithms from the **regression**, **classification**, **clustering**, and **anomaly detection** families. Each is designed to address a different type of machine learning problem.

Download: Machine learning algorithm cheat sheet

Download the cheat sheet here: [Machine Learning Algorithm Cheat Sheet \(11x17 in.\)](#)



Download and print the Machine Learning Algorithm Cheat Sheet in tabloid size to keep it handy and get help choosing an algorithm.

NOTE

See the article [How to choose algorithms for Microsoft Azure Machine Learning](#) for a detailed guide to using this cheat sheet.

More help with algorithms

- For help in using this cheat sheet for choosing the right algorithm, plus a deeper discussion of the different types of machine learning algorithms and how they're used, see [How to choose algorithms for Microsoft Azure Machine Learning](#).

- For a downloadable infographic that describes algorithms and provides examples, see [Downloadable Infographic: Machine learning basics with algorithm examples](#).
- For a list by category of all the machine learning algorithms available in Machine Learning Studio, see [Initialize Model](#) in the Machine Learning Studio Algorithm and Module Help.
- For a complete alphabetical list of algorithms and modules in Machine Learning Studio, see [A-Z list of Machine Learning Studio modules](#) in Machine Learning Studio Algorithm and Module Help.
- To download and print a diagram that gives an overview of the capabilities of Machine Learning Studio, see [Overview diagram of Azure Machine Learning Studio capabilities](#).

NOTE

[Try Azure Machine Learning for free](#)

No credit card or Azure subscription required. [Get started now](#).

Notes and terminology definitions for the machine learning algorithm cheat sheet

- The suggestions offered in this algorithm cheat sheet are approximate rules-of-thumb. Some can be bent, and some can be flagrantly violated. This is intended to suggest a starting point. Don't be afraid to run a head-to-head competition between several algorithms on your data. There is simply no substitute for understanding the principles of each algorithm and understanding the system that generated your data.
- Every machine learning algorithm has its own style or *inductive bias*. For a specific problem, several algorithms may be appropriate and one algorithm may be a better fit than others. But it's not always possible to know beforehand which is the best fit. In cases like these, several algorithms are listed together in the cheat sheet. An appropriate strategy would be to try one algorithm, and if the results are not yet satisfactory, try the others. Here's an example from the [Cortana Intelligence Gallery](#) of an experiment that tries several algorithms against the same data and compares the results: [Compare Multi-class Classifiers: Letter recognition](#).
- There are three main categories of machine learning: **supervised learning**, **unsupervised learning**, and **reinforcement learning**.
 - In **supervised learning**, each data point is labeled or associated with a category or value of interest. An example of a categorical label is assigning an image as either a 'cat' or a 'dog'. An example of a value label is the sale price associated with a used car. The goal of supervised learning is to study many labeled examples like these, and then to be able to make predictions about future data points. For example, identifying new photos with the correct animal or assigning accurate sale prices to other used cars. This is a popular and useful type of machine learning. All of the modules in Azure Machine Learning are supervised learning algorithms except for [K-Means Clustering](#).
 - In **unsupervised learning**, data points have no labels associated with them. Instead, the goal of an unsupervised learning algorithm is to organize the data in some way or to describe its structure. This can mean grouping it into clusters, as K-means does, or finding different ways of looking at complex data so that it appears simpler.
 - In **reinforcement learning**, the algorithm gets to choose an action in response to each data point. It is a common approach in robotics, where the set of sensor readings at one point in time is a data point, and the algorithm must choose the robot's next action. It's also a natural fit for Internet of Things applications. The learning algorithm also receives a reward signal a short time later, indicating how good the decision was. Based on this, the algorithm modifies its strategy in order to achieve the highest reward. Currently there are no reinforcement learning algorithm modules in Azure ML.

- **Bayesian methods** make the assumption of statistically independent data points. This means that the unmodeled variability in one data point is uncorrelated with others, that is, it can't be predicted. For example, if the data being recorded is the number of minutes until the next subway train arrives, two measurements taken a day apart are statistically independent. However, two measurements taken a minute apart are not statistically independent - the value of one is highly predictive of the value of the other.
- **Boosted decision tree regression** takes advantage of feature overlap or interaction among features. That means that, in any given data point, the value of one feature is somewhat predictive of the value of another. For example, in daily high/low temperature data, knowing the low temperature for the day allows you to make a reasonable guess at the high. The information contained in the two features is somewhat redundant.
- Classifying data into more than two categories can be done by either using an inherently multi-class classifier, or by combining a set of two-class classifiers into an **ensemble**. In the ensemble approach, there is a separate two-class classifier for each class - each one separates the data into two categories: "this class" and "not this class." Then these classifiers vote on the correct assignment of the data point. This is the operational principle behind [One-vs-All Multiclass](#).
- Several methods, including logistic regression and the Bayes point machine, assume **linear class boundaries**. That is, they assume that the boundaries between classes are approximately straight lines (or hyperplanes in the more general case). Often this is a characteristic of the data that you don't know until after you've tried to separate it, but it's something that typically can be learned by visualizing beforehand. If the class boundaries look very irregular, stick with decision trees, decision jungles, support vector machines, or neural networks.
- Neural networks can be used with categorical variables by creating a **dummy variable** for each category, setting it to 1 in cases where the category applies, 0 where it doesn't.

Using linear regression in Azure Machine Learning

9/25/2017 • 7 min to read • [Edit Online](#)

Kate Baroni and Ben Boatman are enterprise solution architects in Microsoft's Data Insights Center of Excellence. In this article, they describe their experience migrating an existing regression analysis suite to a cloud-based solution using Azure Machine Learning.

NOTE

Try Azure Machine Learning for free

No credit card or Azure subscription required. [Get started now.](#)

Goal

Our project started with two goals in mind:

1. Use predictive analytics to improve the accuracy of our organization's monthly revenue projections
2. Use Azure Machine Learning to confirm, optimize, increase velocity, and scale of our results.

Like many businesses, our organization goes through a monthly revenue forecasting process. Our small team of business analysts was tasked with using Azure Machine Learning to support the process and improve forecast accuracy. The team spent several months collecting data from multiple sources and running the data attributes through statistical analysis identifying key attributes relevant to services sales forecasting. The next step was to begin prototyping statistical regression models on the data in Excel. Within a few weeks, we had an Excel regression model that was outperforming the current field and finance forecasting processes. This became the baseline prediction result.

We then took the next step to moving our predictive analytics over to Azure Machine Learning to find out how Machine Learning could improve on predictive performance.

Achieving predictive performance parity

Our first priority was to achieve parity between Machine Learning and Excel regression models. Given the same data, and the same split for training and testing data, we wanted to achieve predictive performance parity between Excel and Machine Learning. Initially we failed. The Excel model outperformed the Machine Learning model. The failure was due to a lack of understanding of the base tool setting in Machine Learning. After a sync with the Machine Learning product team, we gained a better understanding of the base setting required for our data sets, and achieved parity between the two models.

Create regression model in Excel

Our Excel Regression used the standard linear regression model found in the Excel Analysis ToolPak.

We calculated *Mean Absolute % Error* and used it as the performance measure for the model. It took 3 months to arrive at a working model using Excel. We brought much of the learning into the Machine Learning Studio experiment which ultimately was beneficial in understanding requirements.

Create comparable experiment in Azure Machine Learning

We followed these steps to create our experiment in Machine Learning Studio:

1. Uploaded the dataset as a csv file to Machine Learning Studio (very small file)
2. Created a new experiment and used the [Select Columns in Dataset](#) module to select the same data features used in Excel
3. Used the [Split Data](#) module (with *Relative Expression* mode) to divide the data into the same training datasets as had been done in Excel
4. Experimented with the [Linear Regression](#) module (default options only), documented, and compared the results to our Excel regression model

Review initial results

At first, the Excel model clearly outperformed the Machine Learning Studio model:

	EXCEL	STUDIO
Performance		
Adjusted R Square	0.96	N/A
Coefficient of Determination	N/A	0.78 (low accuracy)
Mean Absolute Error	\$9.5M	\$ 19.4M
Mean Absolute Error (%)	6.03%	12.2%

When we ran our process and results by the developers and data scientists on the Machine Learning team, they quickly provided some useful tips.

- When you use the [Linear Regression](#) module in Machine Learning Studio, two methods are provided:
 - Online Gradient Descent: May be more suitable for larger-scale problems
 - Ordinary Least Squares: This is the method most people think of when they hear linear regression. For small datasets, Ordinary Least Squares can be a more optimal choice.
- Consider tweaking the L2 Regularization Weight parameter to improve performance. It is set to 0.001 by default, but for our small data set we set it to 0.005 to improve performance.

Mystery solved!

When we applied the recommendations, we achieved the same baseline performance in Machine Learning Studio as with Excel:

	EXCEL	STUDIO (INITIAL)	STUDIO W/ LEAST SQUARES
Labeled value	Actuals (numeric)	same	same
Learner	Excel -> Data Analysis -> Regression	Linear Regression.	Linear Regression
Learner options	N/A	Defaults	ordinary least squares L2 = 0.005
Data Set	26 rows, 3 features, 1 label. All numeric.	same	same

	EXCEL	STUDIO (INITIAL)	STUDIO W/ LEAST SQUARES
Split: Train	Excel trained on the first 18 rows, tested on the last 8 rows.	same	same
Split: Test	Excel regression formula applied to the last 8 rows	same	same
Performance			
Adjusted R Square	0.96	N/A	
Coefficient of Determination	N/A	0.78	0.952049
Mean Absolute Error	\$9.5M	\$ 19.4M	\$9.5M
Mean Absolute Error (%)	6.03%	12.2%	6.03%

In addition, the Excel coefficients compared well to the feature weights in the Azure trained model:

	EXCEL COEFFICIENTS	AZURE FEATURE WEIGHTS
Intercept/Bias	19470209.88	19328500
Feature A	0.832653063	0.834156
Feature B	11071967.08	11007300
Feature C	25383318.09	25140800

Next Steps

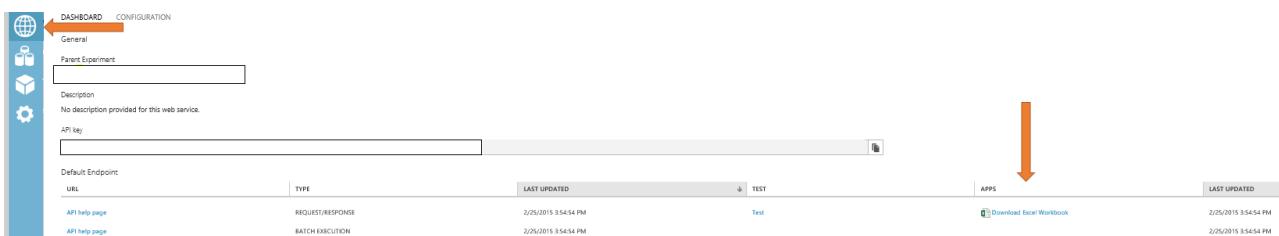
We wanted to consume the Machine Learning web service within Excel. Our business analysts rely on Excel and we needed a way to call the Machine Learning web service with a row of Excel data and have it return the predicted value to Excel.

We also wanted to optimize our model, using the options and algorithms available in Machine Learning Studio.

Integration with Excel

Our solution was to operationalize our Machine Learning regression model by creating a web service from the trained model. Within a few minutes, the web service was created and we could call it directly from Excel to return a predicted revenue value.

The *Web Services Dashboard* section includes a downloadable Excel workbook. The workbook comes pre-formatted with the web service API and schema information embedded. When you click *Download Excel Workbook*, the workbook opens and you can save it to your local computer.



With the workbook open, copy your predefined parameters into the blue Parameter section as shown below. Once the parameters are entered, Excel calls out to the Machine Learning web service and the predicted scored labels will display in the green Predicted Values section. The workbook will continue to create predictions for parameters based on your trained model for all row items entered under Parameters. For more information on how to use this feature, see [Consuming an Azure Machine Learning Web Service from Excel](#).

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S	T	U	V	W	X	Y	Z	AA	AB
1																												
2	This spreadsheet contains your Windows Azure Access Key. Please be mindful when sharing the spreadsheet or you may incur unexpected costs.																											
3	WEB SERVICE URL																											
4	ACCESS KEY																											
5	SCHEMA																											
6	If the web service you consume is hosted in a Free Workspace you may experience delay due to throttling. Upgrade to a Standard Workspace to have higher performance.																											
7	PARAMETERS																											
8	PREDICTED VALUES																											

Optimization and further experiments

Now that we had a baseline with our Excel model, we moved ahead to optimize our Machine Learning Linear Regression Model. We used the module [Filter-Based Feature Selection](#) to improve on our selection of initial data elements and it helped us achieve a performance improvement of 4.6% Mean Absolute Error. For future projects we will use this feature which could save us weeks in iterating through data attributes to find the right set of features to use for modelling.

Next we plan to include additional algorithms like [Bayesian](#) or [Boosted Decision Trees](#) in our experiment to compare performance.

If you want to experiment with regression, a good dataset to try is the Energy Efficiency Regression sample dataset, which has lots of numerical attributes. The dataset is provided as part of the sample datasets in Machine Learning Studio. You can use a variety of learning modules to predict either Heating Load or Cooling Load. The chart below is a performance comparison of different regression learns against the Energy Efficiency dataset predicting for the target variable Cooling Load:

MODEL	MEAN ABSOLUTE ERROR	ROOT MEAN SQUARED ERROR	RELATIVE ABSOLUTE ERROR	RELATIVE SQUARED ERROR	COEFFICIENT OF DETERMINATION
Boosted Decision Tree	0.930113	1.4239	0.106647	0.021662	0.978338
Linear Regression (Gradient Descent)	2.035693	2.98006	0.233414	0.094881	0.905119
Neural Network Regression	1.548195	2.114617	0.177517	0.047774	0.952226
Linear Regression (Ordinary Least Squares)	1.428273	1.984461	0.163767	0.042074	0.957926

Key Takeaways

We learned a lot by from running Excel regression and Azure Machine Learning experiments in parallel. Creating the baseline model in Excel and comparing it to models using Machine Learning [Linear Regression](#) helped us learn Azure Machine Learning, and we discovered opportunities to improve data selection and model performance.

We also found that it is advisable to use [Filter-Based Feature Selection](#) to accelerate future prediction projects. By applying feature selection to your data, you can create an improved model in Machine Learning with better overall performance.

The ability to transfer the predictive analytic forecasting from Machine Learning to Excel systemically allows a significant increase in the ability to successfully provide results to a broad business user audience.

Resources

Here are some resources for helping you work with regression:

- Regression in Excel. If you've never tried regression in Excel, this tutorial makes it easy: <http://www.excel-easy.com/examples/regression.html>
- Regression vs forecasting. Tyler Chessman wrote a blog article explaining how to do time series forecasting in Excel, which contains a good beginner's description of linear regression. <http://sqlmag.com/sql-server-analysis-services/understanding-time-series-forecasting-concepts>
- Ordinary Least Squares Linear Regression: Flaws, Problems and Pitfalls. For an introduction and discussion of Regression: <http://www.clockbackward.com/2009/06/18/ordinary-least-squares-linear-regression-flaws-problems-and-pitfalls/>

Create text analytics models in Azure Machine Learning Studio

9/25/2017 • 5 min to read • [Edit Online](#)

You can use Azure Machine Learning to build and operationalize text analytics models. These models can help you solve, for example, document classification or sentiment analysis problems.

In a text analytics experiment, you would typically:

1. Clean and preprocess text dataset
2. Extract numeric feature vectors from pre-processed text
3. Train classification or regression model
4. Score and validate the model
5. Deploy the model to production

In this tutorial, you learn these steps as we walk through a sentiment analysis model using Amazon Book Reviews dataset (see this research paper “Biographies, Bollywood, Boom-boxes and Blenders: Domain Adaptation for Sentiment Classification” by John Blitzer, Mark Dredze, and Fernando Pereira; Association of Computational Linguistics (ACL), 2007.) This dataset consists of review scores (1-2 or 4-5) and a free-form text. The goal is to predict the review score: low (1-2) or high (4-5).

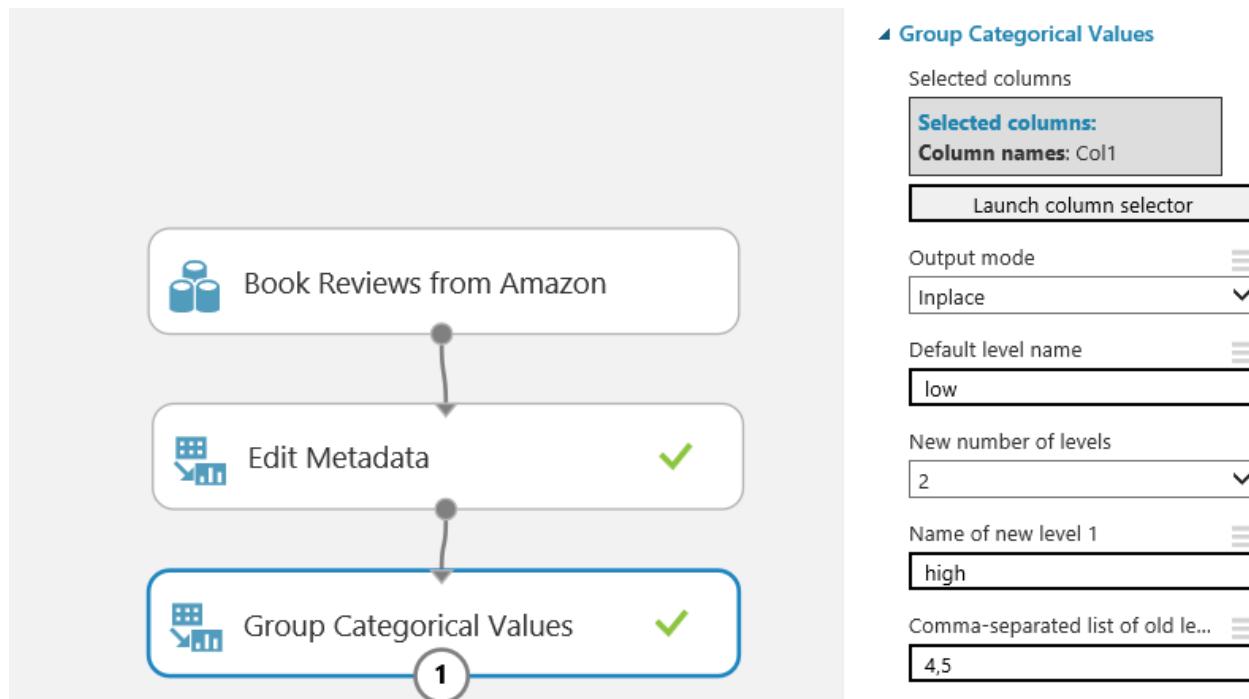
You can find experiments covered in this tutorial at Cortana Intelligence Gallery:

[Predict Book Reviews](#)

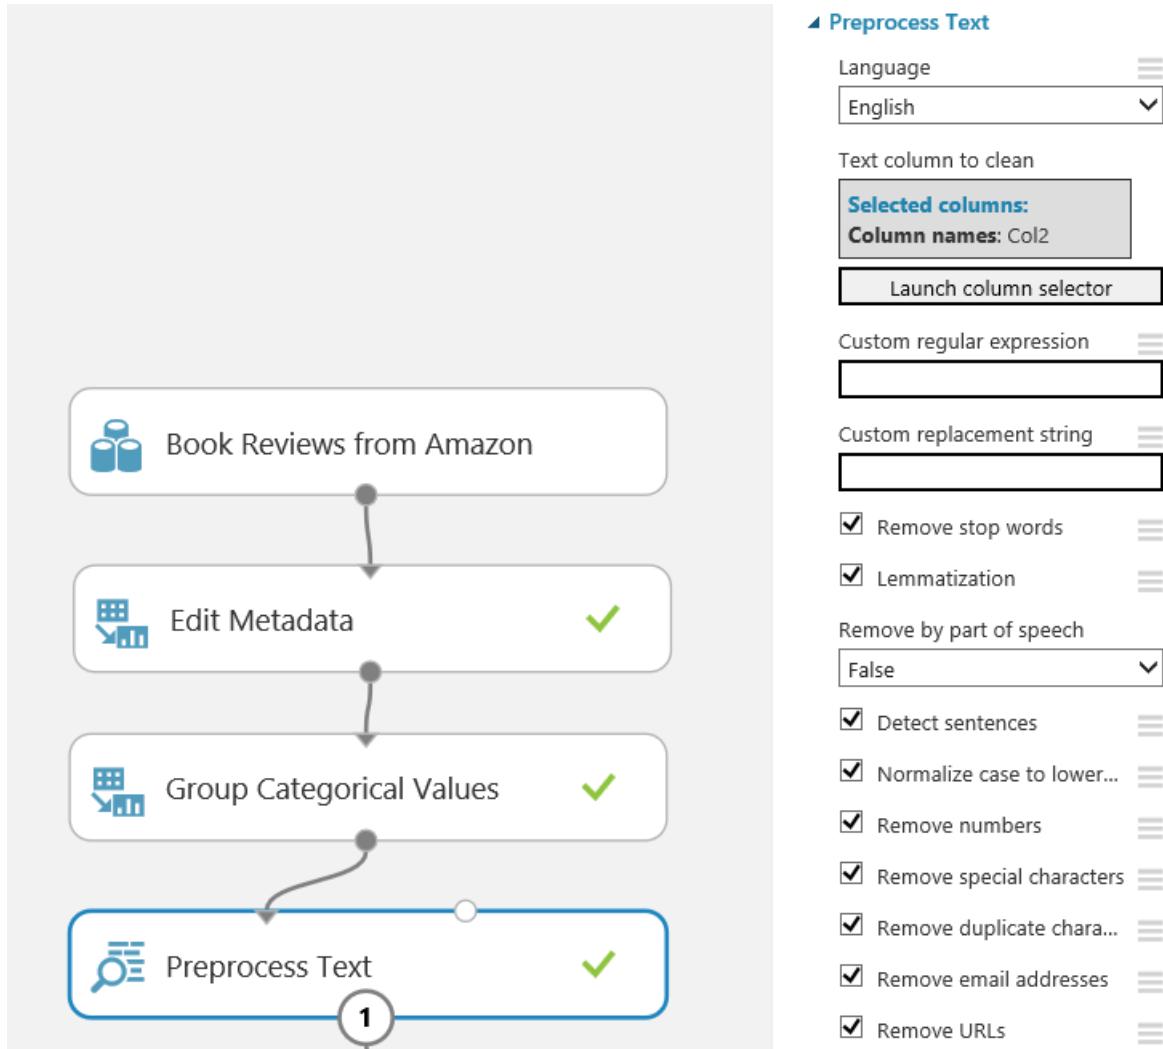
[Predict Book Reviews - Predictive Experiment](#)

Step 1: Clean and preprocess text dataset

We begin the experiment by dividing the review scores into categorical low and high buckets to formulate the problem as two-class classification. We use [Edit Metadata](#) and [Group Categorical Values](#) modules.



Then, we clean the text using [Preprocess Text](#) module. The cleaning reduces the noise in the dataset, help you find the most important features, and improve the accuracy of the final model. We remove stopwords - common words such as "the" or "a" - and numbers, special characters, duplicated characters, email addresses, and URLs. We also convert the text to lowercase, lemmatize the words, and detect sentence boundaries that are then indicated by "|||" symbol in pre-processed text.

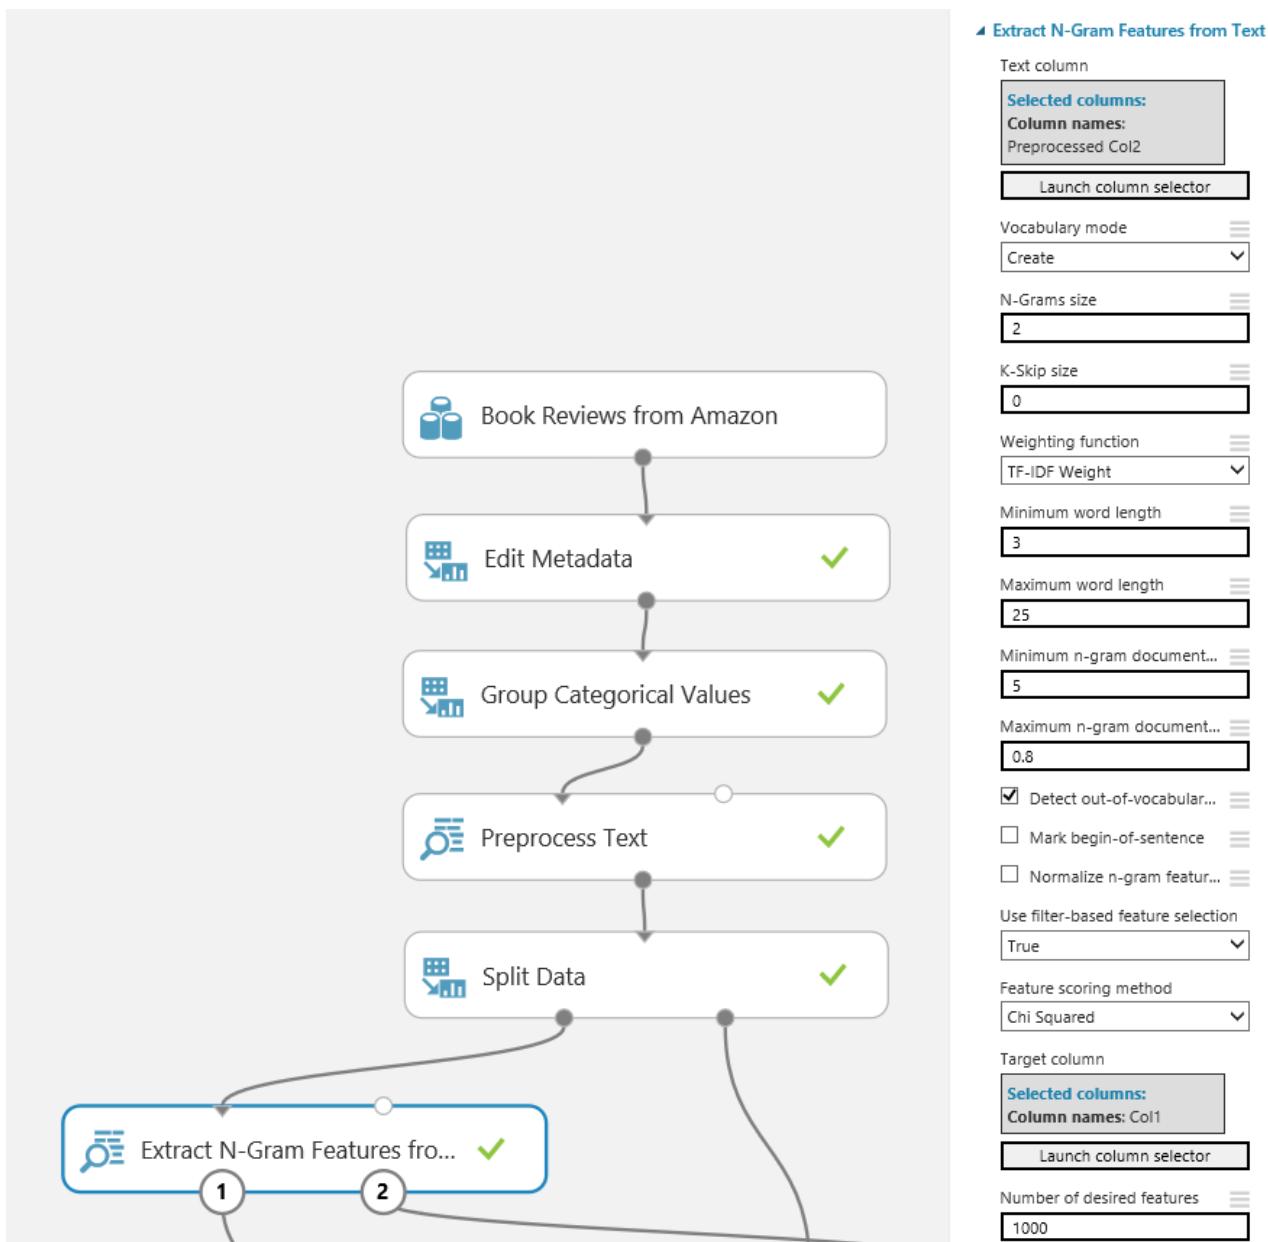


What if you want to use a custom list of stopwords? You can pass it in as optional input. You can also use custom C# syntax regular expression to replace substrings, and remove words by part of speech: nouns, verbs, or adjectives.

After the preprocessing is complete, we split the data into train and test sets.

Step 2: Extract numeric feature vectors from pre-processed text

To build a model for text data, you typically have to convert free-form text into numeric feature vectors. In this example, we use [Extract N-Gram Features from Text](#) module to transform the text data to such format. This module takes a column of whitespace-separated words and computes a dictionary of words, or N-grams of words, that appear in your dataset. Then, it counts how many times each word, or N-gram, appears in each record, and creates feature vectors from those counts. In this tutorial, we set N-gram size to 2, so our feature vectors include single words and combinations of two subsequent words.



We apply TF*IDF (Term Frequency Inverse Document Frequency) weighting to N-gram counts. This approach adds weight of words that appear frequently in a single record but are rare across the entire dataset. Other options include binary, TF, and graph weighing.

Such text features often have high dimensionality. For example, if your corpus has 100,000 unique words, your feature space would have 100,000 dimensions, or more if N-grams are used. The Extract N-Gram Features module gives you a set of options to reduce the dimensionality. You can choose to exclude words that are short or long, or too uncommon or too frequent to have significant predictive value. In this tutorial, we exclude N-grams that appear in fewer than 5 records or in more than 80% of records.

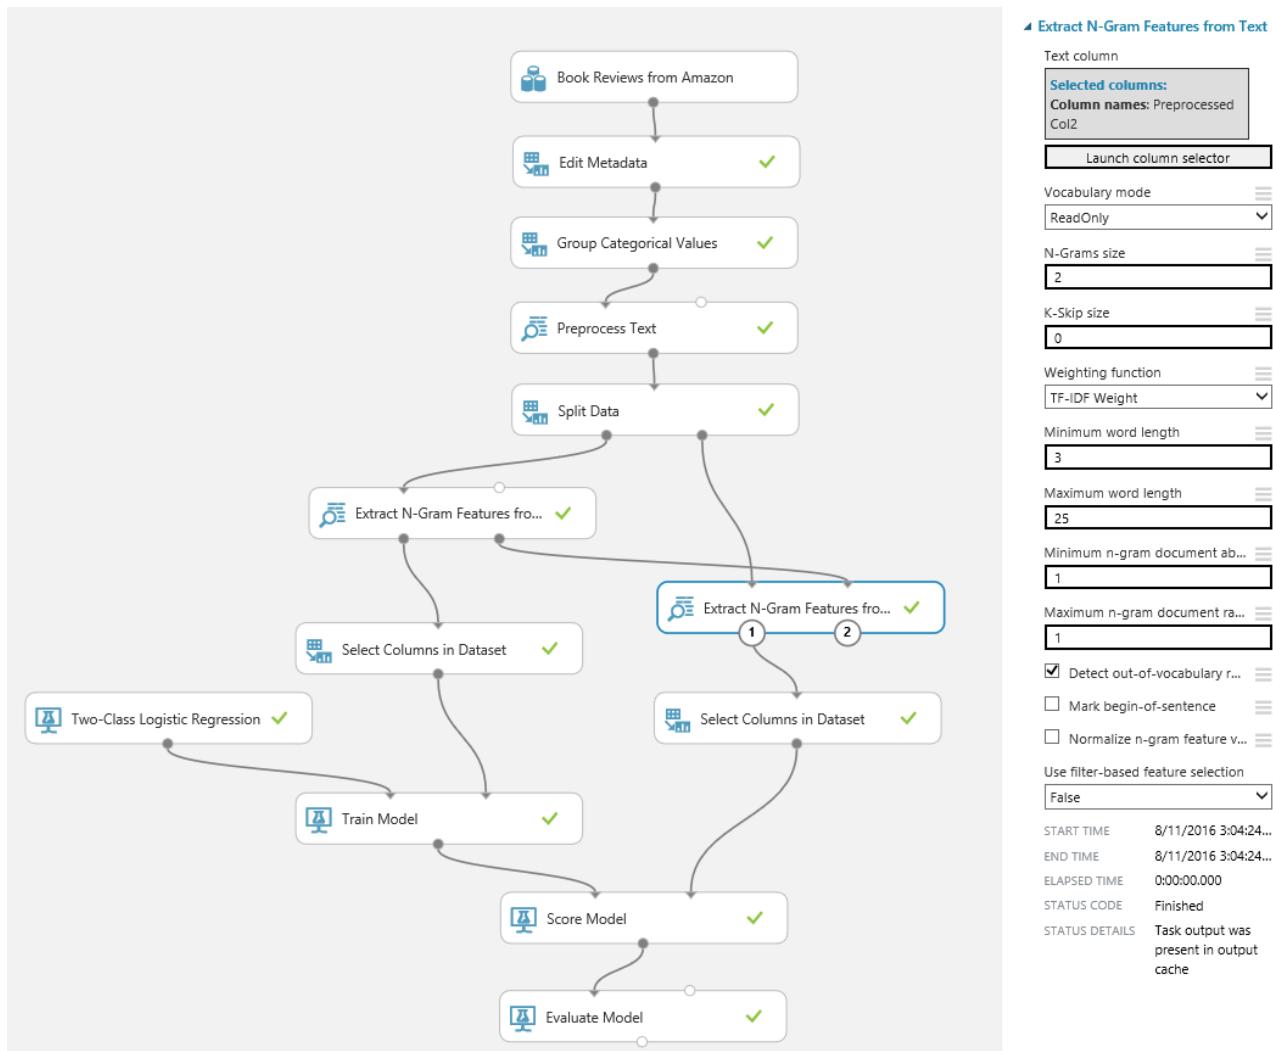
Also, you can use feature selection to select only those features that are the most correlated with your prediction target. We use Chi-Squared feature selection to select 1000 features. You can view the vocabulary of selected words or N-grams by clicking the right output of Extract N-grams module.

As an alternative approach to using Extract N-Gram Features, you can use Feature Hashing module. Note though that [Feature Hashing](#) does not have build-in feature selection capabilities, or TF*IDF weighing.

Step 3: Train classification or regression model

Now the text has been transformed to numeric feature columns. The dataset still contains string columns from previous stages, so we use Select Columns in Dataset to exclude them.

We then use [Two-Class Logistic Regression](#) to predict our target: high or low review score. At this point, the text analytics problem has been transformed into a regular classification problem. You can use the tools available in Azure Machine Learning to improve the model. For example, you can experiment with different classifiers to find out how accurate results they give, or use hyperparameter tuning to improve the accuracy.



Step 4: Score and validate the model

How would you validate the trained model? We score it against the test dataset and evaluate the accuracy. However, the model learned the vocabulary of N-grams and their weights from the training dataset. Therefore, we should use that vocabulary and those weights when extracting features from test data, as opposed to creating the vocabulary anew. Therefore, we add Extract N-Gram Features module to the scoring branch of the experiment, connect the output vocabulary from training branch, and set the vocabulary mode to read-only. We also disable the filtering of N-grams by frequency by setting the minimum to 1 instance and maximum to 100%, and turn off the feature selection.

After the text column in test data has been transformed to numeric feature columns, we exclude the string columns from previous stages like in training branch. We then use Score Model module to make predictions and Evaluate Model module to evaluate the accuracy.

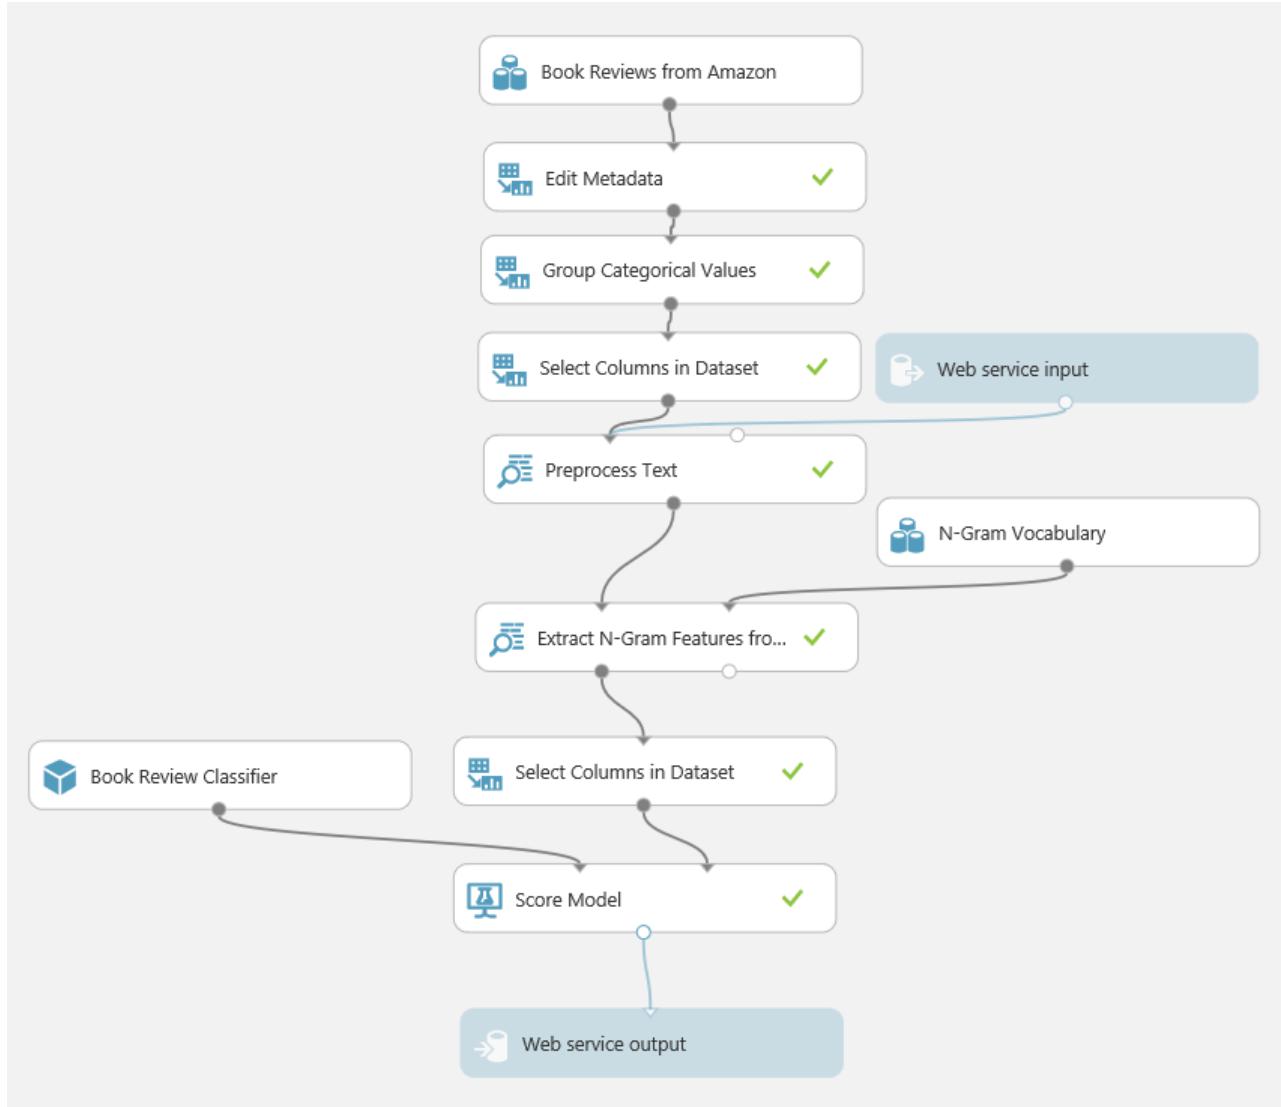
Step 5: Deploy the model to production

The model is almost ready to be deployed to production. When deployed as web service, it takes free-form text string as input, and return a prediction "high" or "low." It uses the learned N-gram vocabulary to transform the text to features, and trained logistic regression model to make a prediction from those features.

To set up the predictive experiment, we first save the N-gram vocabulary as dataset, and the trained logistic

regression model from the training branch of the experiment. Then, we save the experiment using "Save As" to create an experiment graph for predictive experiment. We remove the Split Data module and the training branch from the experiment. We then connect the previously saved N-gram vocabulary and model to Extract N-Gram Features and Score Model modules, respectively. We also remove the Evaluate Model module.

We insert Select Columns in Dataset module before Preprocess Text module to remove the label column, and unselect "Append score column to dataset" option in Score Module. That way, the web service does not request the label it is trying to predict, and does not echo the input features in response.



Now we have an experiment that can be published as a web service and called using request-response or batch execution APIs.

Next Steps

Learn about text analytics modules from [MSDN documentation](#).

How to evaluate model performance in Azure Machine Learning

9/25/2017 • 12 min to read • [Edit Online](#)

This article demonstrates how to evaluate the performance of a model in Azure Machine Learning Studio and provides a brief explanation of the metrics available for this task. Three common supervised learning scenarios are presented:

- regression
- binary classification
- multiclass classification

NOTE

Try Azure Machine Learning for free

No credit card or Azure subscription required. [Get started now.](#)

Evaluating the performance of a model is one of the core stages in the data science process. It indicates how successful the scoring (predictions) of a dataset has been by a trained model.

Azure Machine Learning supports model evaluation through two of its main machine learning modules: [Evaluate Model](#) and [Cross-Validate Model](#). These modules allow you to see how your model performs in terms of a number of metrics that are commonly used in machine learning and statistics.

Evaluation vs. Cross Validation

Evaluation and cross validation are standard ways to measure the performance of your model. They both generate evaluation metrics that you can inspect or compare against those of other models.

[Evaluate Model](#) expects a scored dataset as input (or 2 in case you would like to compare the performance of 2 different models). This means that you need to train your model using the [Train Model](#) module and make predictions on some dataset using the [Score Model](#) module, before you can evaluate the results. The evaluation is based on the scored labels/probabilities along with the true labels, all of which are output by the [Score Model](#) module.

Alternatively, you can use cross validation to perform a number of train-score-evaluate operations (10 folds) automatically on different subsets of the input data. The input data is split into 10 parts, where one is reserved for testing, and the other 9 for training. This process is repeated 10 times and the evaluation metrics are averaged. This helps in determining how well a model would generalize to new datasets. The [Cross-Validate Model](#) module takes in an untrained model and some labeled dataset and outputs the evaluation results of each of the 10 folds, in addition to the averaged results.

In the following sections, we will build simple regression and classification models and evaluate their performance, using both the [Evaluate Model](#) and the [Cross-Validate Model](#) modules.

Evaluating a Regression Model

Assume we want to predict a car's price using some features such as dimensions, horsepower, engine specs, and so on. This is a typical regression problem, where the target variable (*price*) is a continuous numeric value. We can fit a

simple linear regression model that, given the feature values of a certain car, can predict the price of that car. This regression model can be used to score the same dataset we trained on. Once we have the predicted prices for all of the cars, we can evaluate the performance of the model by looking at how much the predictions deviate from the actual prices on average. To illustrate this, we use the *Automobile price data (Raw)* dataset available in the **Saved Datasets** section in Azure Machine Learning Studio.

Creating the Experiment

Add the following modules to your workspace in Azure Machine Learning Studio:

- Automobile price data (Raw)
- Linear Regression
- Train Model
- Score Model
- Evaluate Model

Connect the ports as shown below in Figure 1 and set the Label column of the **Train Model** module to *price*.

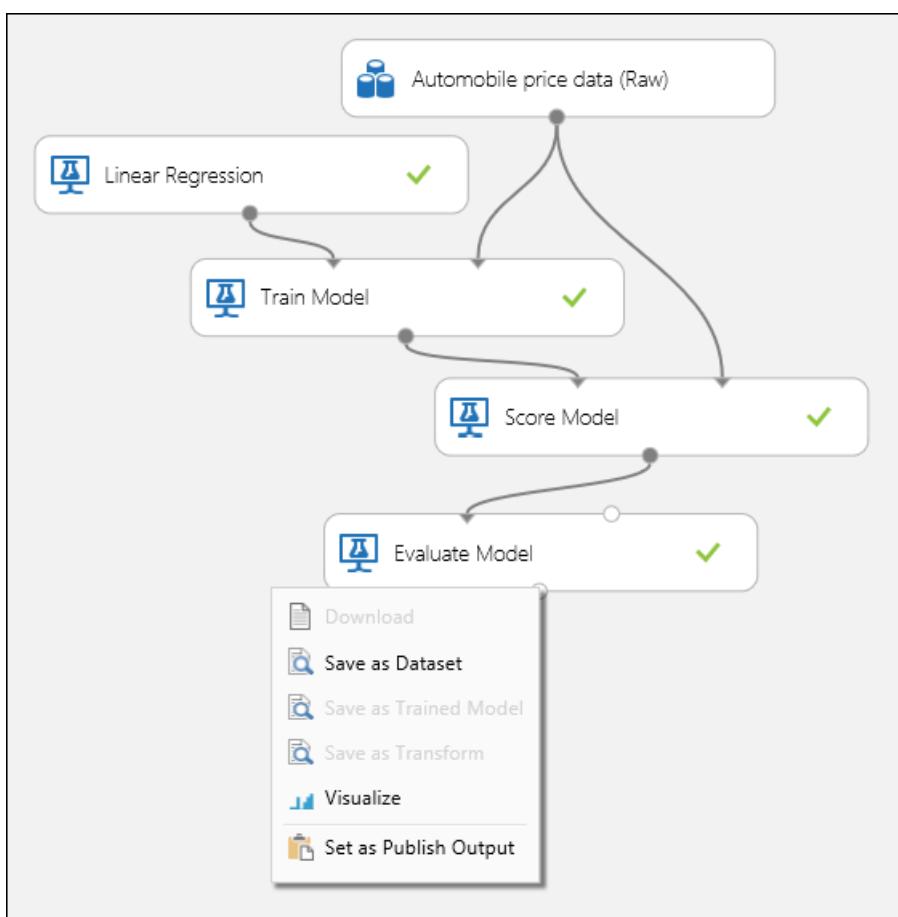


Figure 1. Evaluating a Regression Model.

Inspecting the Evaluation Results

After running the experiment, you can click on the output port of the **Evaluate Model** module and select *Visualize* to see the evaluation results. The evaluation metrics available for regression models are: *Mean Absolute Error*, *Root Mean Absolute Error*, *Relative Absolute Error*, *Relative Squared Error*, and the *Coefficient of Determination*.

The term "error" here represents the difference between the predicted value and the true value. The absolute value or the square of this difference are usually computed to capture the total magnitude of error across all instances, as the difference between the predicted and true value could be negative in some cases. The error metrics measure the predictive performance of a regression model in terms of the mean deviation of its predictions from the true values. Lower error values mean the model is more accurate in making predictions. An overall error metric of 0 means that the model fits the data perfectly.

The coefficient of determination, which is also known as R squared, is also a standard way of measuring how well the model fits the data. It can be interpreted as the proportion of variation explained by the model. A higher proportion is better in this case, where 1 indicates a perfect fit.

Metrics	
Mean Absolute Error	747.975254
Root Mean Squared Error	955.587783
Relative Absolute Error	0.163528
Relative Squared Error	0.026598
Coefficient of Determination	0.973402

Figure 2. Linear Regression Evaluation Metrics.

Using Cross Validation

As mentioned earlier, you can perform repeated training, scoring and evaluations automatically using the [Cross-Validate Model](#) module. All you need in this case is a dataset, an untrained model, and a [Cross-Validate Model](#) module (see figure below). Note that you need to set the label column to *price* in the [Cross-Validate Model](#) module's properties.

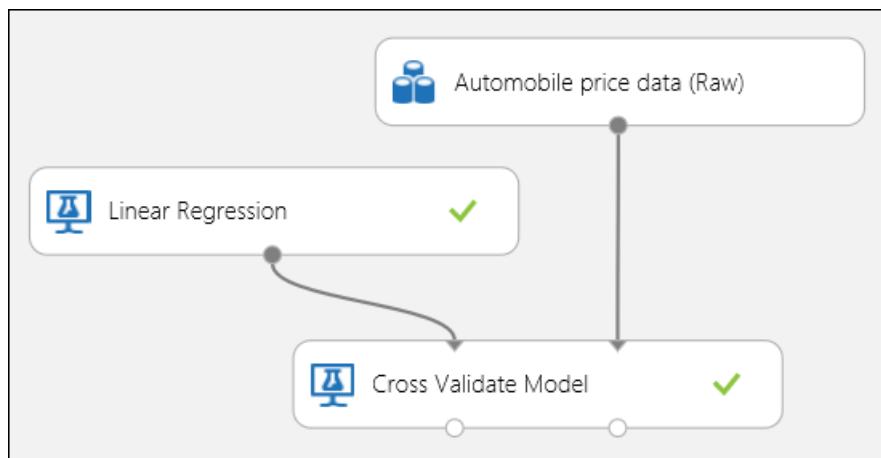


Figure 3. Cross-Validating a Regression Model.

After running the experiment, you can inspect the evaluation results by clicking on the right output port of the [Cross-Validate Model](#) module. This will provide a detailed view of the metrics for each iteration (fold), and the averaged results of each of the metrics (Figure 4).

rows	columns		Fold Number	Number of examples in fold	Model	Mean Absolute Error	Root Mean Squared Error	Relative Absolute Error	Relative Squared Error	Coefficient of Determination
12	8		0	20	Microsoft.Analytics.Machi nLearning.LocalBatchLin earRegressor	1975.690002	2868.853475	0.318462	0.261852	0.738148
			1	21	Microsoft.Analytics.Machi nLearning.LocalBatchLin earRegressor	1204.367667	1736.369322	0.259077	0.101772	0.898228
			2	20	Microsoft.Analytics.Machi nLearning.LocalBatchLin earRegressor	1275.525323	1565.686411	0.146945	0.021472	0.978528
			3	21	Microsoft.Analytics.Machi nLearning.LocalBatchLin earRegressor	1100.785629	1479.819120	0.150116	0.024304	0.975696
			4	20	Microsoft.Analytics.Machi nLearning.LocalBatchLin earRegressor	672.068144	902.187494	0.091132	0.010176	0.989824
			5	20	Microsoft.Analytics.Machi nLearning.LocalBatchLin earRegressor	1166.036215	1492.147079	0.227148	0.055168	0.944832
			6	21	Microsoft.Analytics.Machi nLearning.LocalBatchLin earRegressor	1544.169782	1956.495628	0.272313	0.07162	0.92838
			7	21	Microsoft.Analytics.Machi nLearning.LocalBatchLin earRegressor	1435.049593	2055.843829	0.17445	0.042546	0.957454
			8	20	Microsoft.Analytics.Machi nLearning.LocalBatchLin earRegressor	1583.894079	2192.621829	0.186784	0.042628	0.957372
			9	21	Microsoft.Analytics.Machi nLearning.LocalBatchLin earRegressor	1069.820402	1464.412889	0.113144	0.020426	0.979574
			Mean	205	Microsoft.Analytics.Machi nLearning.LocalBatchLin earRegressor	1310.740683	1771.443708	0.199957	0.065197	0.934803
			Standard Deviation	205	Microsoft.Analytics.Machi nLearning.LocalBatchLin earRegressor	350.085716	532.636933	0.073523	0.07438	0.07436

Figure 4. Cross-Validation Results of a Regression Model.

Evaluating a Binary Classification Model

In a binary classification scenario, the target variable has only two possible outcomes, for example: {0, 1} or {false, true}, {negative, positive}. Assume you are given a dataset of adult employees with some demographic and employment variables, and that you are asked to predict the income level, a binary variable with the values {"<=50K", ">50K"}. In other words, the negative class represents the employees who make less than or equal to 50K per year, and the positive class represents all other employees. As in the regression scenario, we would train a model, score some data, and evaluate the results. The main difference here is the choice of metrics Azure Machine Learning computes and outputs. To illustrate the income level prediction scenario, we will use the [Adult](#) dataset to create an Azure Machine Learning experiment and evaluate the performance of a two-class logistic regression model, a commonly used binary classifier.

Creating the Experiment

Add the following modules to your workspace in Azure Machine Learning Studio:

- Adult Census Income Binary Classification dataset
- [Two-Class Logistic Regression](#)
- Train Model
- Score Model
- Evaluate Model

Connect the ports as shown below in Figure 5 and set the Label column of the [Train Model](#) module to *income*.

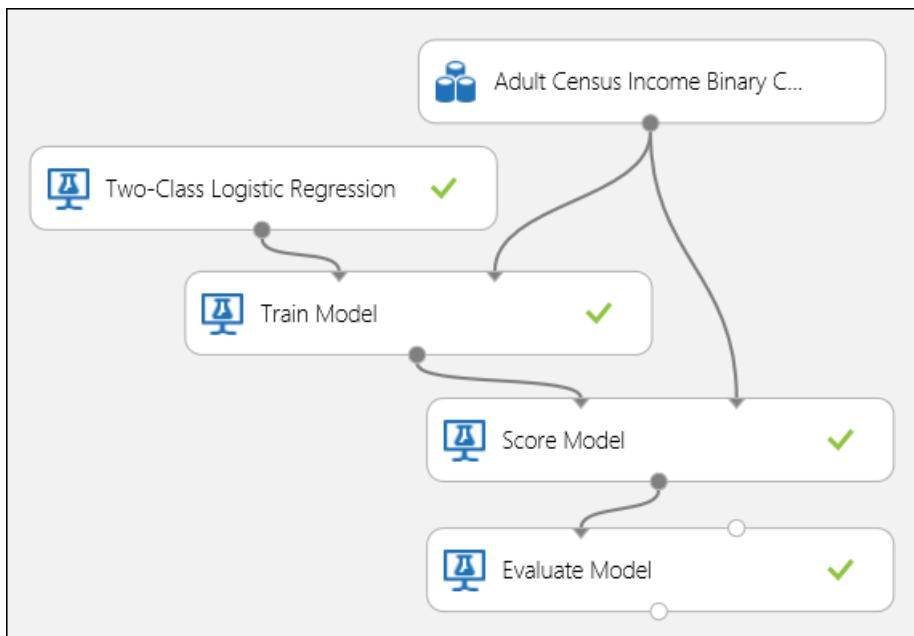


Figure 5. Evaluating a Binary Classification Model.

Inspecting the Evaluation Results

After running the experiment, you can click on the output port of the [Evaluate Model](#) module and select *Visualize* to see the evaluation results (Figure 7). The evaluation metrics available for binary classification models are: *Accuracy*, *Precision*, *Recall*, *F1 Score*, and *AUC*. In addition, the module outputs a confusion matrix showing the number of true positives, false negatives, false positives, and true negatives, as well as *ROC*, *Precision/Recall*, and *Lift* curves.

Accuracy is simply the proportion of correctly classified instances. It is usually the first metric you look at when evaluating a classifier. However, when the test data is unbalanced (where most of the instances belong to one of the classes), or you are more interested in the performance on either one of the classes, accuracy doesn't really capture the effectiveness of a classifier. In the income level classification scenario, assume you are testing on some

data where 99% of the instances represent people who earn less than or equal to 50K per year. It is possible to achieve a 0.99 accuracy by predicting the class " $<=50K$ " for all instances. The classifier in this case appears to be doing a good job overall, but in reality, it fails to classify any of the high-income individuals (the 1%) correctly.

For that reason, it is helpful to compute additional metrics that capture more specific aspects of the evaluation. Before going into the details of such metrics, it is important to understand the confusion matrix of a binary classification evaluation. The class labels in the training set can take on only 2 possible values, which we usually refer to as positive or negative. The positive and negative instances that a classifier predicts correctly are called true positives (TP) and true negatives (TN), respectively. Similarly, the incorrectly classified instances are called false positives (FP) and false negatives (FN). The confusion matrix is simply a table showing the number of instances that fall under each of these 4 categories. Azure Machine Learning automatically decides which of the two classes in the dataset is the positive class. If the class labels are Boolean or integers, then the 'true' or '1' labeled instances are assigned the positive class. If the labels are strings, as in the case of the income dataset, the labels are sorted alphabetically and the first level is chosen to be the negative class while the second level is the positive class.

		Predicted	
		Positive	Negative
Actual True	TP	FN	
	FP	TN	

Figure 6. Binary Classification Confusion Matrix.

Going back to the income classification problem, we would want to ask several evaluation questions that help us understand the performance of the classifier used. A very natural question is: 'Out of the individuals whom the model predicted to be earning $>50K$ (TP+FP), how many were classified correctly (TP)?' This question can be answered by looking at the **Precision** of the model, which is the proportion of positives that are classified correctly: $TP/(TP+FP)$. Another common question is "Out of all the high earning employees with income $>50k$ (TP+FN), how many did the classifier classify correctly (TP)". This is actually the **Recall**, or the true positive rate: $TP/(TP+FN)$ of the classifier. You might notice that there is an obvious trade-off between precision and recall. For example, given a relatively balanced dataset, a classifier that predicts mostly positive instances, would have a high recall, but a rather low precision as many of the negative instances would be misclassified resulting in a large number of false positives. To see a plot of how these two metrics vary, you can click on the 'PRECISION/RECALL' curve in the evaluation result output page (top left part of Figure 7).

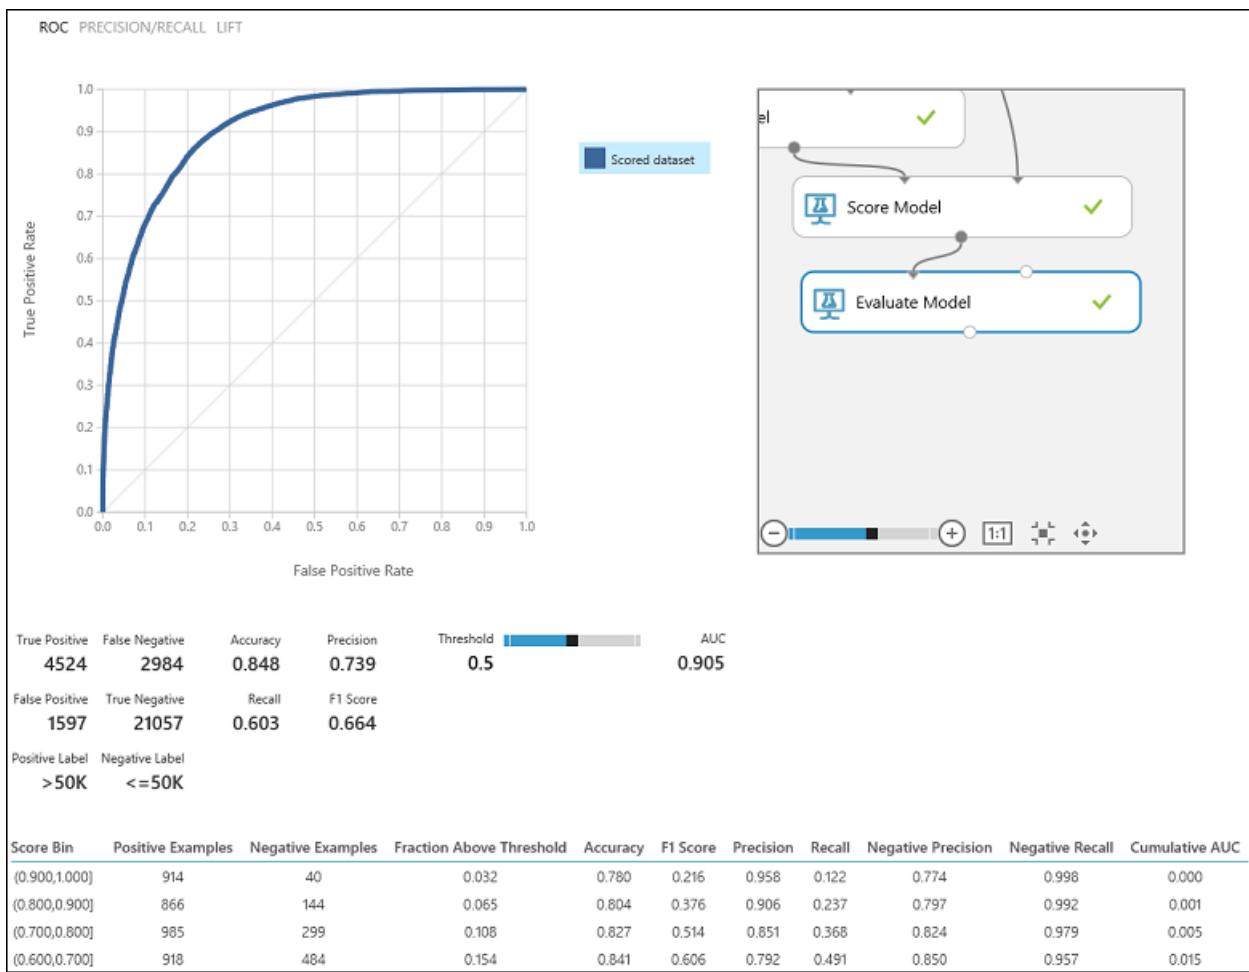


Figure 7. Binary Classification Evaluation Results.

Another related metric that is often used is the **F1 Score**, which takes both precision and recall into consideration. It is the harmonic mean of these 2 metrics and is computed as such: $F1 = 2 \times (\text{precision} \times \text{recall}) / (\text{precision} + \text{recall})$. The F1 score is a good way to summarize the evaluation in a single number, but it's always a good practice to look at both precision and recall together to better understand how a classifier behaves.

In addition, one can inspect the true positive rate vs. the false positive rate in the **Receiver Operating Characteristic (ROC)** curve and the corresponding **Area Under the Curve (AUC)** value. The closer this curve is to the upper left corner, the better the classifier's performance is (that is maximizing the true positive rate while minimizing the false positive rate). Curves that are close to the diagonal of the plot, result from classifiers that tend to make predictions that are close to random guessing.

Using Cross Validation

As in the regression example, we can perform cross validation to repeatedly train, score and evaluate different subsets of the data automatically. Similarly, we can use the **Cross-Validate Model** module, an untrained logistic regression model, and a dataset. The label column must be set to *income* in the **Cross-Validate Model** module's properties. After running the experiment and clicking on the right output port of the **Cross-Validate Model** module, we can see the binary classification metric values for each fold, in addition to the mean and standard deviation of each.

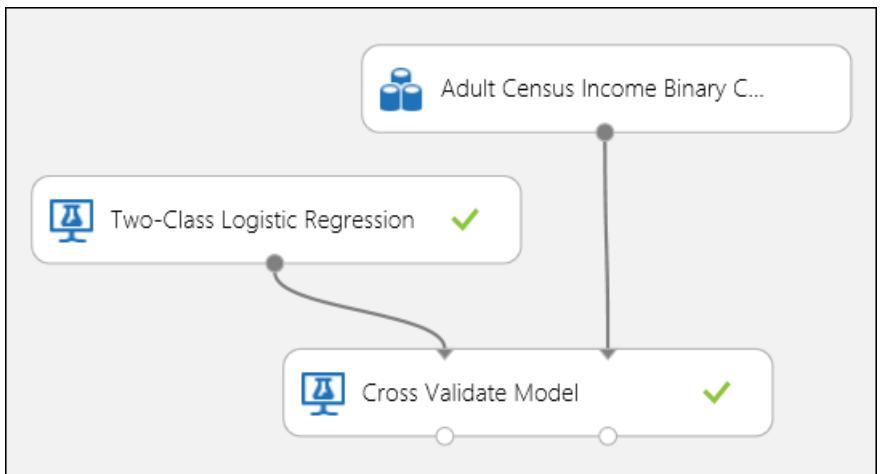


Figure 8. Cross-Validating a Binary Classification Model.

rows	columns								
12	10								
Fold Number	Number of examples in fold	Model	Accuracy	Precision	Recall	F-Score	AUC	Average Log Loss	Training Log Loss
view as									
0	3256	Logistic Regression	0.850399	0.737342	0.621333	0.674385	0.905809	0.325788	41.989626
1	3256	Logistic Regression	0.842245	0.707358	0.585062	0.640424	0.901871	0.32527	40.990733
2	3256	Logistic Regression	0.852583	0.767007	0.591864	0.668148	0.905203	0.32936	41.51439
3	3256	Logistic Regression	0.847333	0.729685	0.598639	0.657698	0.898848	0.333642	40.075913
4	3256	Logistic Regression	0.844305	0.741194	0.614213	0.671756	0.90568	0.332448	41.921163
5	3256	Logistic Regression	0.834601	0.716393	0.57199	0.636099	0.897829	0.339869	39.879498
6	3256	Logistic Regression	0.85	0.728188	0.598621	0.657078	0.902695	0.327404	40.596835
7	3257	Logistic Regression	0.846128	0.743464	0.599473	0.663749	0.90208	0.333052	41.153979
8	3256	Logistic Regression	0.846973	0.734861	0.601071	0.661267	0.902736	0.329349	41.259562
9	3256	Logistic Regression	0.849388	0.742718	0.607947	0.668609	0.903525	0.328219	41.623392
Mean	32561	Logistic Regression	0.846396	0.734821	0.599021	0.659921	0.902629	0.33044	41.100509
Standard Deviation	32561	Logistic Regression	0.005145	0.016276	0.014102	0.012778	0.00269	0.004408	0.725968

Figure 9. Cross-Validation Results of a Binary Classifier.

Evaluating a Multiclass Classification Model

In this experiment we will use the popular [Iris](#) dataset which contains instances of 3 different types (classes) of the iris plant. There are 4 feature values (sepal length/width and petal length/width) for each instance. In the previous experiments we trained and tested the models using the same datasets. Here, we will use the [Split Data](#) module to create 2 subsets of the data, train on the first, and score and evaluate on the second. The Iris dataset is publicly available on the [UCI Machine Learning Repository](#), and can be downloaded using an [Import Data](#) module.

Creating the Experiment

Add the following modules to your workspace in Azure Machine Learning Studio:

- [Import Data](#)
- [Multiclass Decision Forest](#)
- [Split Data](#)
- [Train Model](#)
- [Score Model](#)
- [Evaluate Model](#)

Connect the ports as shown below in Figure 10.

Set the Label column index of the [Train Model](#) module to 5. The dataset has no header row but we know that the class labels are in the fifth column.

Click on the [Import Data](#) module and set the *Data source* property to *Web URL via HTTP*, and the *URL* to <http://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data>.

Set the fraction of instances to be used for training in the [Split Data](#) module (0.7 for example).

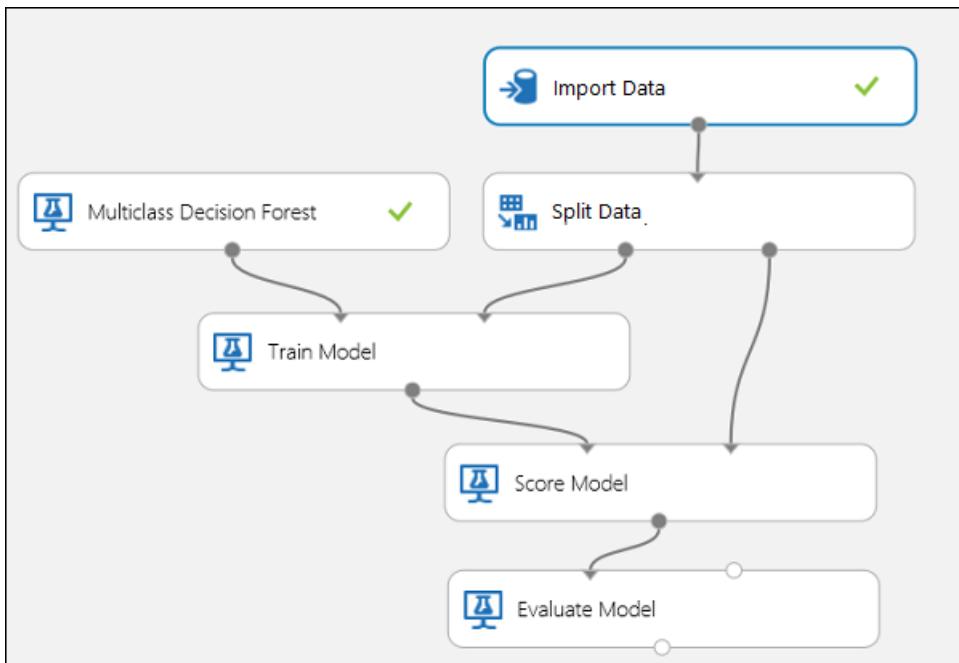


Figure 10. Evaluating a Multiclass Classifier

Inspecting the Evaluation Results

Run the experiment and click on the output port of [Evaluate Model](#). The evaluation results are presented in the form of a confusion matrix, in this case. The matrix shows the actual vs. predicted instances for all 3 classes.

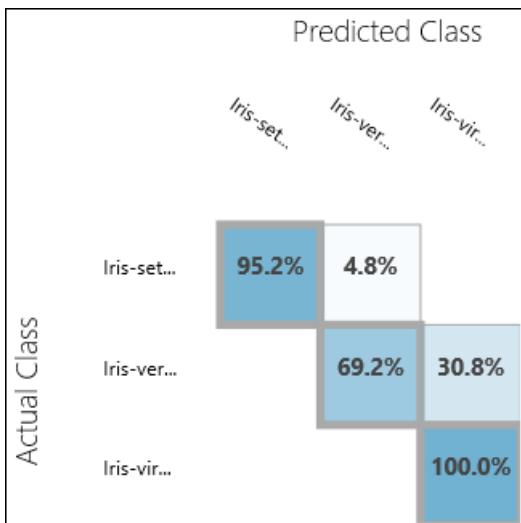


Figure 11. Multiclass Classification Evaluation Results.

Using Cross Validation

As mentioned earlier, you can perform repeated training, scoring and evaluations automatically using the [Cross-Validate Model](#) module. You would need a dataset, an untrained model, and a [Cross-Validate Model](#) module (see figure below). Again you need to set the label column of the [Cross-Validate Model](#) module (column index 5 in this case). After running the experiment and clicking the right output port of the [Cross-Validate Model](#), you can inspect the metric values for each fold as well as the mean and standard deviation. The metrics displayed here are the similar to the ones discussed in the binary classification case. However, note that in multiclass classification, computing the true positives/negatives and false positives/negatives is done by counting on a per-class basis, as there is no overall positive or negative class. For example, when computing the precision or recall of the 'Iris-setosa' class, it is assumed that this is the positive class and all others as negative.

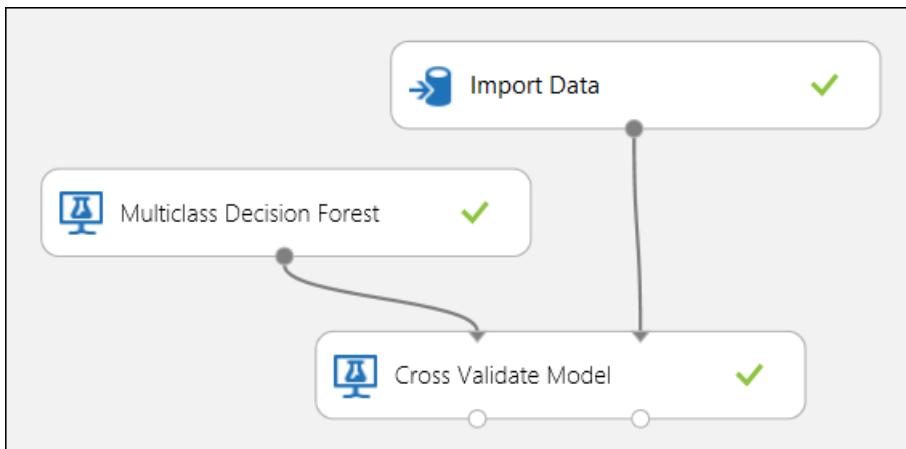


Figure 12. Cross-Validating a Multiclass Classification Model.

rows	columns										
Fold Number	Number of examples in fold	Model	Average Log Loss for Class "Iris-setosa"	Precision for Class "Iris-setosa"	Recall for Class "Iris-setosa"	Average Log Loss for Class "Iris-versicolor"	Precision for Class "Iris-versicolor"	Recall for Class "Iris-versicolor"	Average Log Loss for Class "Iris-virginica"	Precision for Class "Iris-virginica"	Recall for Class "Iris-virginica"
view as											
0	15	Microsoft.Analytics.Modules.Gemin.III.MulticlassGemin.DecisionForestClassifier	0	1	1	0.415888	1	0.8	0	0.857143	1
1	15	Microsoft.Analytics.Modules.Gemin.III.MulticlassGemin.DecisionForestClassifier	0	1	1	0.026706	0.833333	1	0.122604	1	0.875
2	15	Microsoft.Analytics.Modules.Gemin.III.MulticlassGemin.DecisionForestClassifier	0.057536	1	1	0.026706	1	1	0.057536	1	1

Figure 13. Cross-Validation Results of a Multiclass Classification Model.

Choose parameters to optimize your algorithms in Azure Machine Learning

11/30/2017 • 3 min to read • [Edit Online](#)

This topic describes how to choose the right hyperparameter set for an algorithm in Azure Machine Learning. Most machine learning algorithms have parameters to set. When you train a model, you need to provide values for those parameters. The efficacy of the trained model depends on the model parameters that you choose. The process of finding the optimal set of parameters is known as *model selection*.

NOTE

Try Azure Machine Learning for free

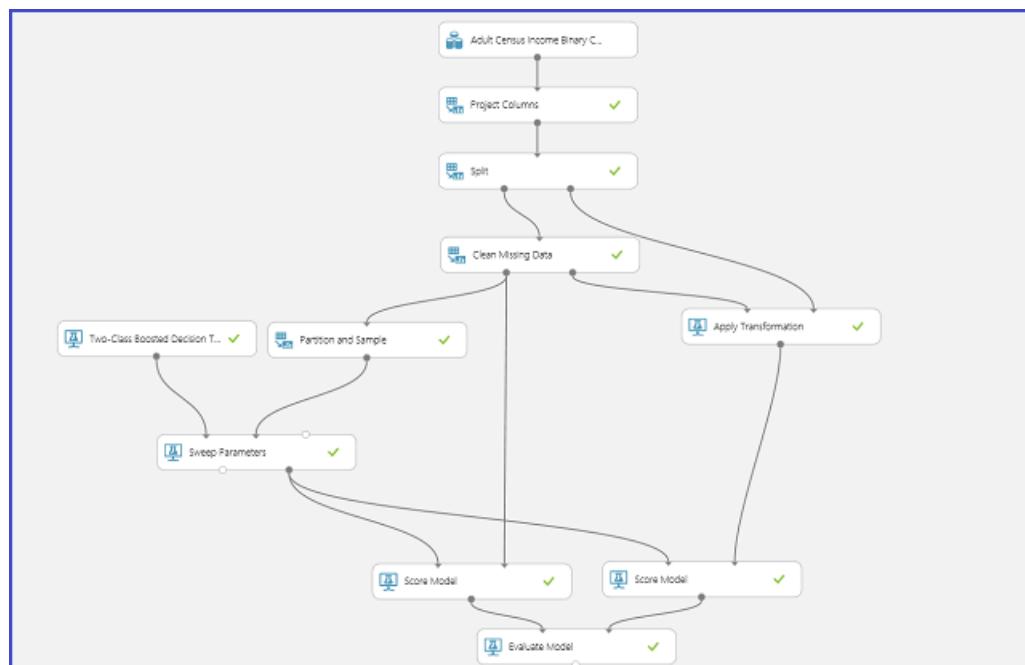
No credit card or Azure subscription required. [Get started now.](#)

There are various ways to do model selection. In machine learning, cross-validation is one of the most widely used methods for model selection, and it is the default model selection mechanism in Azure Machine Learning. Because Azure Machine Learning supports both R and Python, you can always implement their own model selection mechanisms by using either R or Python.

There are four steps in the process of finding the best parameter set:

1. **Define the parameter space:** For the algorithm, first decide the exact parameter values you want to consider.
2. **Define the cross-validation settings:** Decide how to choose cross-validation folds for the dataset.
3. **Define the metric:** Decide what metric to use for determining the best set of parameters, such as accuracy, root mean squared error, precision, recall, or f-score.
4. **Train, evaluate, and compare:** For each unique combination of the parameter values, cross-validation is carried out by and based on the error metric you define. After evaluation and comparison, you can choose the best-performing model.

The following image illustrates shows how this can be achieved in Azure Machine Learning.



Define the parameter space

You can define the parameter set at the model initialization step. The parameter pane of all machine learning algorithms has two trainer modes: *Single Parameter* and *Parameter Range*. Choose Parameter Range mode. In Parameter Range mode, you can enter multiple values for each parameter. You can enter comma-separated values in the text box.

▲ Two-Class Boosted Decision Tree

Create trainer mode

Single Parameter

Maximum number of leaves per tree
20

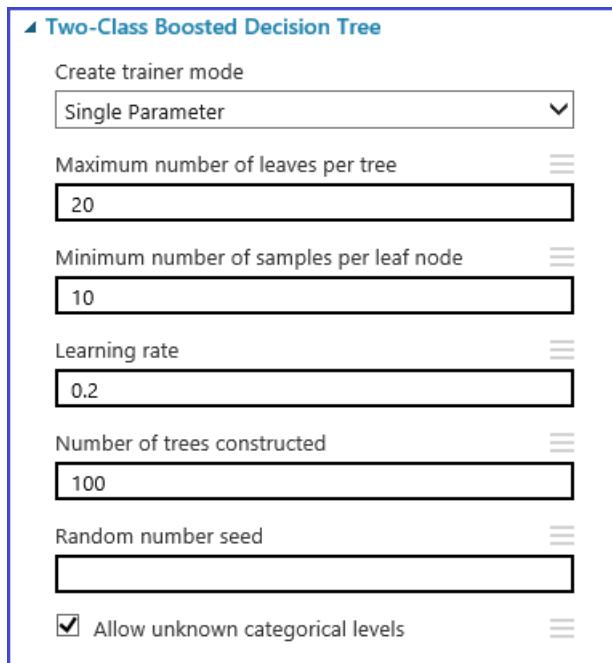
Minimum number of samples per leaf node
10

Learning rate
0.2

Number of trees constructed
100

Random number seed

Allow unknown categorical levels



Alternately, you can define the maximum and minimum points of the grid and the total number of points to be generated with **Use Range Builder**. By default, the parameter values are generated on a linear scale. But if **Log Scale** is checked, the values are generated in the log scale (that is, the ratio of the adjacent points is constant instead of their difference). For integer parameters, you can define a range by using a hyphen. For example, "1-10" means that all integers between 1 and 10 (both inclusive) form the parameter set. A mixed mode is also supported. For example, the parameter set "1-10, 20, 50" would include integers 1-10, 20, and 50.

▲ Two-Class Boosted Decision Tree

Create trainer mode

Parameter Range

Maximum number of leaves per tree
Use Range Builder
Parameter Range : 101 - 900


Number of points : 3

Log Scale

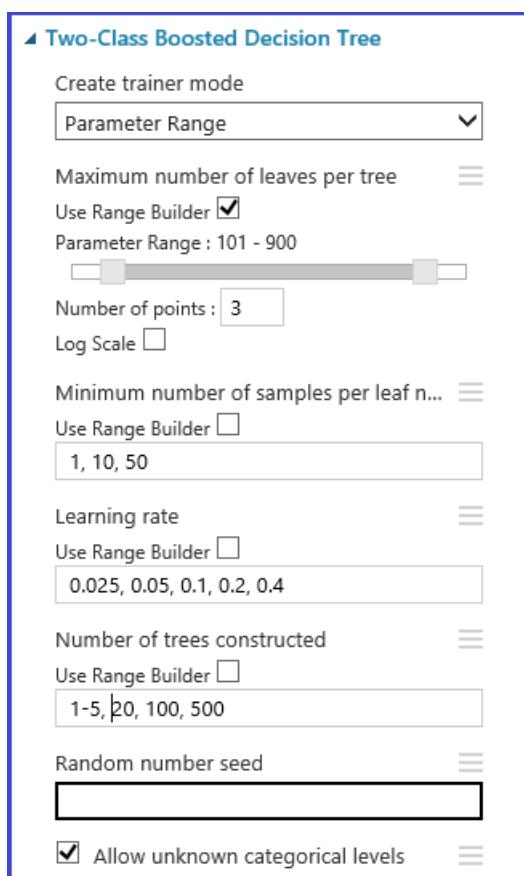
Minimum number of samples per leaf n...
Use Range Builder
1, 10, 50

Learning rate
Use Range Builder
0.025, 0.05, 0.1, 0.2, 0.4

Number of trees constructed
Use Range Builder
1-5, 20, 100, 500

Random number seed

Allow unknown categorical levels



Define cross-validation folds

The [Partition and Sample](#) module can be used to randomly assign folds to the data. In the following sample configuration for the module, we define five folds and randomly assign a fold number to the sample instances.

The screenshot shows the configuration interface for the Partition and Sample module. It includes fields for partition mode (set to 'Assign to Folds'), random seed (set to 0), partitioner method (set to 'Partition evenly'), and the number of folds (set to 5). The 'Randomized split' checkbox is checked. A 'Stratified split' dropdown is set to 'False'.

Define the metric

The [Tune Model Hyperparameters](#) module provides support for empirically choosing the best set of parameters for a given algorithm and dataset. In addition to other information regarding training the model, the **Properties** pane of this module includes the metric for determining the best parameter set. It has two different drop-down list boxes for classification and regression algorithms, respectively. If the algorithm under consideration is a classification algorithm, the regression metric is ignored and vice versa. In this specific example, the metric is **Accuracy**.

The screenshot shows the 'Sweep Parameters' section of the Tune Model Hyperparameters module. It includes a dropdown for 'Specify parameter sweeping mode' (set to 'Entire grid'), a 'Label column' field containing 'Selected columns: Column names: income' (with a 'Launch column selector' button), and dropdowns for 'Metric for measuring performance for classification' (set to 'Accuracy') and 'Metric for measuring performance for regression' (set to 'Mean absolute error').

Train, evaluate, and compare

The same [Tune Model Hyperparameters](#) module trains all the models that correspond to the parameter set, evaluates various metrics, and then creates the best-trained model based on the metric you choose. This module has two mandatory inputs:

- The untrained learner
- The dataset

The module also has an optional dataset input. Connect the dataset with fold information to the mandatory dataset input. If the dataset is not assigned any fold information, then a 10-fold cross-validation is automatically executed by default. If the fold assignment is not done and a validation dataset is provided at the optional dataset port, then a train-test mode is chosen and the first dataset is used to train the model for each parameter combination.

Boosted Decision Tree Classifier

Settings

Setting	Value
Number Of Leaves	8
Minimum Leaf Instances	10
Learning Rate	0.1
Number Of Trees	500
Allow Unknown Levels	true
Random Number Seed	

The model is then evaluated on the validation dataset. The left output port of the module shows different metrics as functions of parameter values. The right output port gives the trained model that corresponds to the best-performing model according to the chosen metric (**Accuracy** in this case).

rows	columns										
180	11										
view as	Number of leaves	Minimum leaf instances	Learning rate	Number of trees	Accuracy	Precision	Recall	F-Score	AUC	Average Log Loss	Training Log Loss
grid	8	10	0.1	500	0.866128	0.75785	0.652595	0.701295	0.922837	0.290561	47.363286
grid	8	1	0.05	500	0.866097	0.762796	0.644306	0.698562	0.922672	0.289389	47.575474
grid	8	1	0.2	100	0.865944	0.763813	0.641755	0.697484	0.921469	0.291591	47.176649
grid	32	1	0.05	100	0.865913	0.759367	0.648769	0.699725	0.921503	0.29273	46.970281
grid	8	10	0.2	100	0.865882	0.763821	0.641372	0.697262	0.922052	0.290024	47.46048
grid	32	1	0.025	500	0.865668	0.749604	0.663946	0.70418	0.923187	0.292484	47.014813
grid	8	10	0.05	500	0.865514	0.761322	0.643158	0.697269	0.922803	0.288929	47.658898
grid	32	1	0.1	100	0.865453	0.751234	0.659737	0.702519	0.922703	0.292779	46.961499
grid	8	1	0.025	500	0.865391	0.764778	0.636909	0.695011	0.921386	0.29142	47.207519
grid	8	50	0.05	500	0.865115	0.760932	0.641372	0.696055	0.92261	0.288991	47.647701
grid	32	10	0.025	500	0.865115	0.748308	0.662798	0.702962	0.923435	0.29145	47.202173
grid	8	50	0.025	500	0.865084	0.764336	0.635761	0.694145	0.921226	0.291554	47.183315
grid	32	10	0.05	100	0.865084	0.756395	0.648642	0.698387	0.92155	0.29239	47.031913
grid	8	10	0.025	500	0.864992	0.764066	0.635633	0.693957	0.921674	0.290844	47.312033
grid	32	10	0.2	20	0.864961	0.757939	0.645326	0.697114	0.920488	0.294732	46.607674
grid	8	1	0.1	500	0.864808	0.752756	0.653105	0.699399	0.922545	0.291764	47.145278
grid	32	10	0.1	100	0.864777	0.748626	0.66012	0.701593	0.92295	0.291274	47.234005
grid	8	1	0.1	100	0.864715	0.763255	0.635251	0.693395	0.920514	0.293092	46.904731
grid	8	10	0.1	100	0.864715	0.763416	0.634996	0.693309	0.920639	0.292815	46.954929
grid	8	10	0.4	100	0.864623	0.754863	0.648387	0.697585	0.921302	0.292915	46.936722
grid	32	1	0.2	20	0.864439	0.757089	0.64354	0.695712	0.920124	0.295428	46.481506

You can see the exact parameters chosen by visualizing the right output port. This model can be used in scoring a test set or in an operationalized web service after saving as a trained model.

Interpret model results in Azure Machine Learning

11/30/2017 • 13 min to read • [Edit Online](#)

This topic explains how to visualize and interpret prediction results in Azure Machine Learning Studio. After you have trained a model and done predictions on top of it ("scored the model"), you need to understand and interpret the prediction result.

NOTE

Try Azure Machine Learning for free

No credit card or Azure subscription required. [Get started now.](#)

There are four major kinds of machine learning models in Azure Machine Learning:

- Classification
- Clustering
- Regression
- Recommender systems

The modules used for prediction on top of these models are:

- [Score Model](#) module for classification and regression
- [Assign to Clusters](#) module for clustering
- [Score Matchbox Recommender](#) for recommendation systems

This document explains how to interpret prediction results for each of these modules. For an overview of these modules, see [How to choose parameters to optimize your algorithms in Azure Machine Learning](#).

This topic addresses prediction interpretation but not model evaluation. For more information about how to evaluate your model, see [How to evaluate model performance in Azure Machine Learning](#).

If you are new to Azure Machine Learning and need help creating a simple experiment to get started, see [Create a simple experiment in Azure Machine Learning Studio](#) in Azure Machine Learning Studio.

Classification

There are two subcategories of classification problems:

- Problems with only two classes (two-class or binary classification)
- Problems with more than two classes (multi-class classification)

Azure Machine Learning has different modules to deal with each of these types of classification, but the methods for interpreting their prediction results are similar.

Two-class classification

Example experiment

An example of a two-class classification problem is the classification of iris flowers. The task is to classify iris flowers based on their features. The Iris data set provided in Azure Machine Learning is a subset of the popular [Iris data set](#) containing instances of only two flower species (classes 0 and 1). There are four features for each flower (sepal length, sepal width, petal length, and petal width).

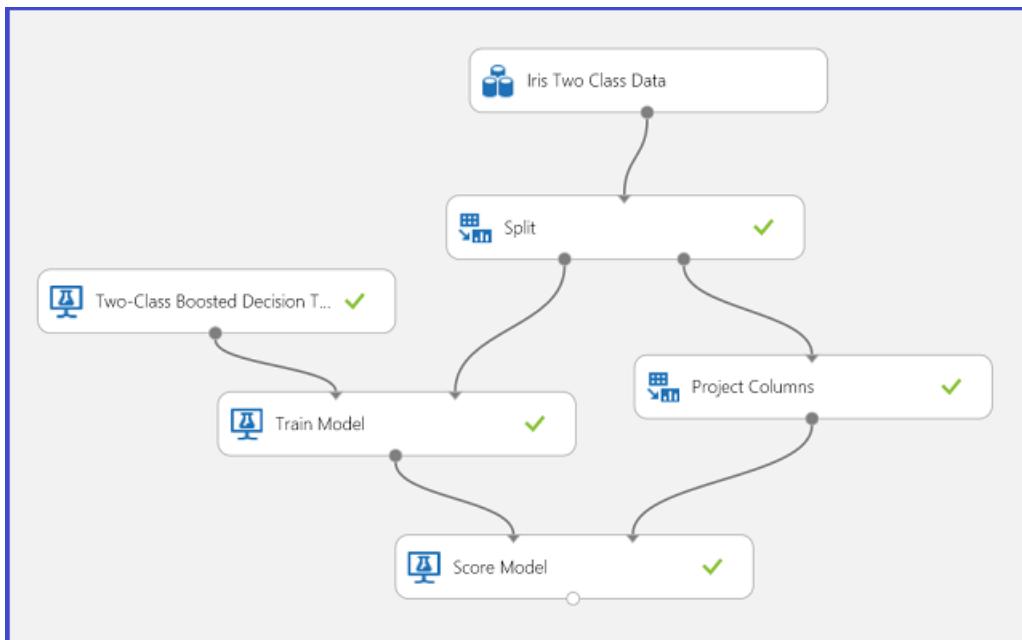
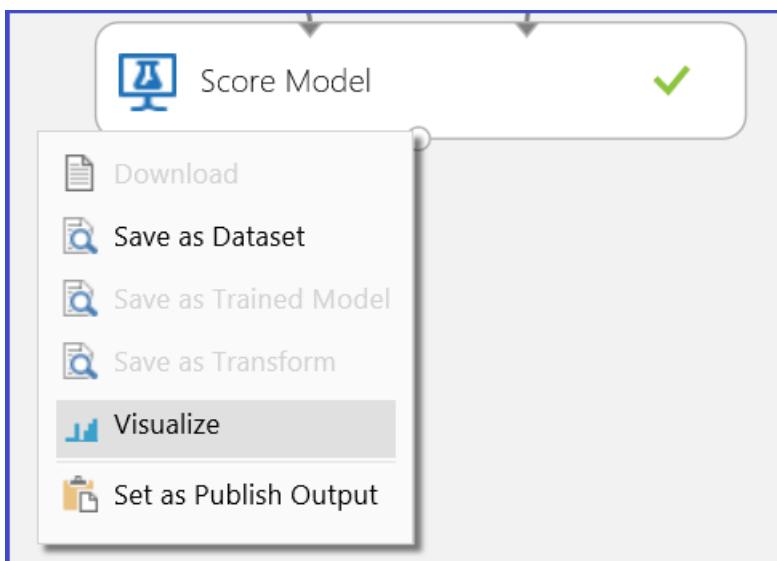


Figure 1. Iris two-class classification problem experiment

An experiment has been performed to solve this problem, as shown in Figure 1. A two-class boosted decision tree model has been trained and scored. Now you can visualize the prediction results from the **Score Model** module by clicking the output port of the **Score Model** module and then clicking **Visualize**.



This brings up the scoring results as shown in Figure 2.

Interpreting Score Module - Two-class > Score Model > Scored dataset							
row	column	sepal-length	sepal-width	petal-length	petal-width	Scored Labels	Scored Probabilities
40	sepal-length	4.4	3	1.3	0.2	0	0.028571
5.8	sepal-width	2.7	5.1	1.9	1	1	0.965517
6.4	petal-length	3.2	5.3	2.3	1	1	0.965517
5	petal-width	3	1.6	0.2	0	0	0.028571
4.7	Scored Labels	3.2	1.6	0.2	0	0	0.028571
5	Scored Probabilities	3.2	1.2	0.2	0	0	0.028571
6.3	sepal-length	2.9	5.6	1.8	1	1	0.965517
5.2	sepal-width	3.5	1.5	0.2	0	0	0.028571
6.1	petal-length	2.6	5.6	1.4	1	1	0.965517
6.3	petal-width	2.5	5	1.9	1	1	0.965517
5.8	Scored Labels	2.8	5.1	2.4	1	1	0.965517
6.4	Scored Probabilities	2.7	5.3	1.9	1	1	0.965517
6.5	sepal-length	3	5.2	2	1	1	0.965517
7.2	sepal-width	3.6	6.1	2.5	1	1	0.965517
6.3	petal-length	2.8	5.1	1.5	1	1	0.965517
6.4	petal-width	2.8	5.6	2.2	1	1	0.965517
5.1	Scored Labels	3.8	5.1	0.3	0	0	0.028571
6.9	Scored Probabilities	3.1	5.1	2.3	1	1	0.965517
4.4	sepal-length	3.2	1.3	0.2	0	0	0.028571
5.5	sepal-width	4.2	1.4	0.2	0	0	0.028571

Figure 2. Visualize a score model result in two-class classification

Result interpretation

There are six columns in the results table. The left four columns are the four features. The right two columns, Scored Labels and Scored Probabilities, are the prediction results. The Scored Probabilities column shows the probability that a flower belongs to the positive class (Class 1). For example, the first number in the column (0.028571) means there is 0.028571 probability that the first flower belongs to Class 1. The Scored Labels column shows the predicted class for each flower. This is based on the Scored Probabilities column. If the scored probability of a flower is larger than 0.5, it is predicted as Class 1. Otherwise, it is predicted as Class 0.

Web service publication

After the prediction results have been understood and judged sound, the experiment can be published as a web service so that you can deploy it in various applications and call it to obtain class predictions on any new iris flower. To learn how to change a training experiment into a scoring experiment and publish it as a web service, see [Publish the Azure Machine Learning web service](#). This procedure provides you with a scoring experiment as shown in Figure 3.

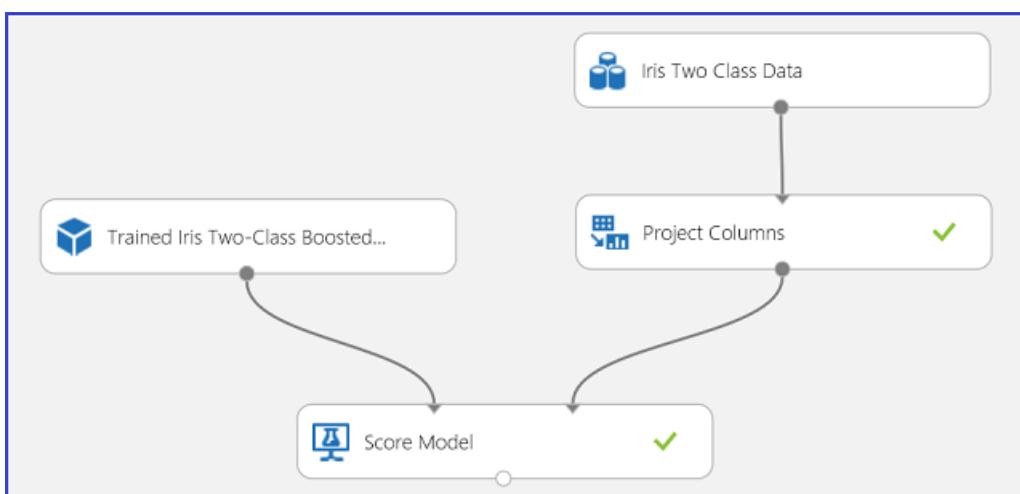


Figure 3. Scoring the iris two-class classification problem experiment

Now you need to set the input and output for the web service. The input is the right input port of [Score Model](#), which is the Iris flower features input. The choice of the output depends on whether you are interested in the predicted class (scored label), the scored probability, or both. In this example, it is assumed that you are interested in both. To select the desired output columns, use a [Select Columns in Data set](#) module. Click [Select Columns in Data set](#), click **Launch column selector**, and select **Scored Labels** and **Scored Probabilities**. After setting the output port of [Select Columns in Data set](#) and running it again, you should be ready to publish the scoring experiment as a web service by clicking **PUBLISH WEB SERVICE**. The final experiment looks like Figure 4.

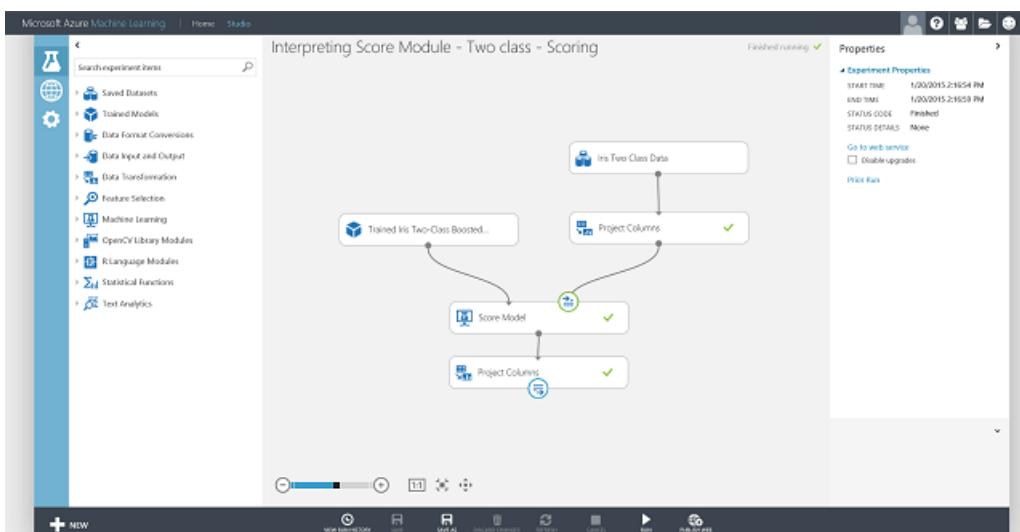


Figure 4. Final scoring experiment of an iris two-class classification problem

After you run the web service and enter some feature values of a test instance, the result returns two numbers. The first number is the scored label, and the second is the scored probability. This flower is predicted as Class 1 with 0.9655 probability.

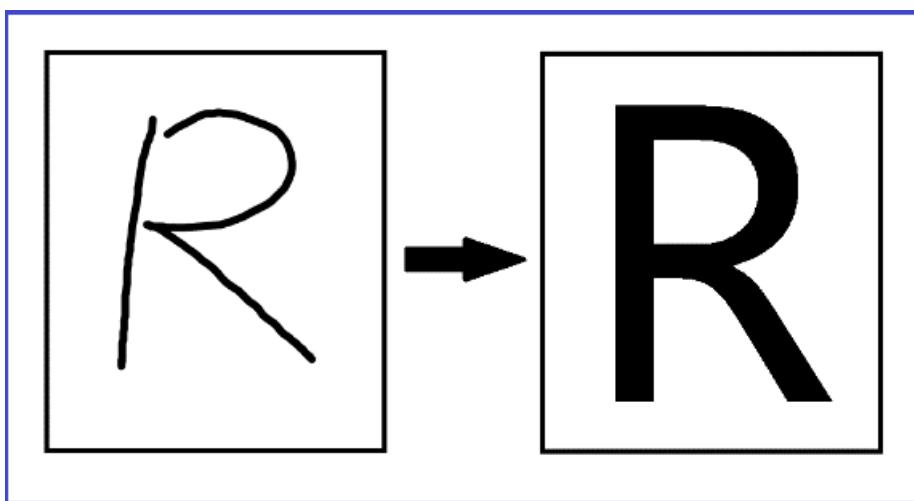
The screenshot shows a user interface for a 'Test Interpreting Score Module - Two class - Scoring Service'. At the top, it says 'Enter data to predict'. Below are four input fields for 'SEPAL-LENGTH' (value 6), 'SEPAL-WIDTH' (value 3), 'PETAL-LENGTH' (value 5), and 'PETAL-WIDTH' (value 2). A large green checkmark icon is located at the bottom right of the input area. Below the input form, a dark grey box displays the results: a left arrow followed by the text "'Interpreting Score Module - Two class - Scoring' test finished successfully", and a green checkmark followed by the text 'Result: 1,0.965517282485962'.

Figure 5. Web service result of iris two-class classification

Multi-class classification

Example experiment

In this experiment, you perform a letter-recognition task as an example of multiclass classification. The classifier attempts to predict a certain letter (class) based on some hand-written attribute values extracted from the hand-written images.



In the training data, there are 16 features extracted from hand-written letter images. The 26 letters form our 26 classes. Figure 6 shows an experiment that will train a multiclass classification model for letter recognition and predict on the same feature set on a test data set.

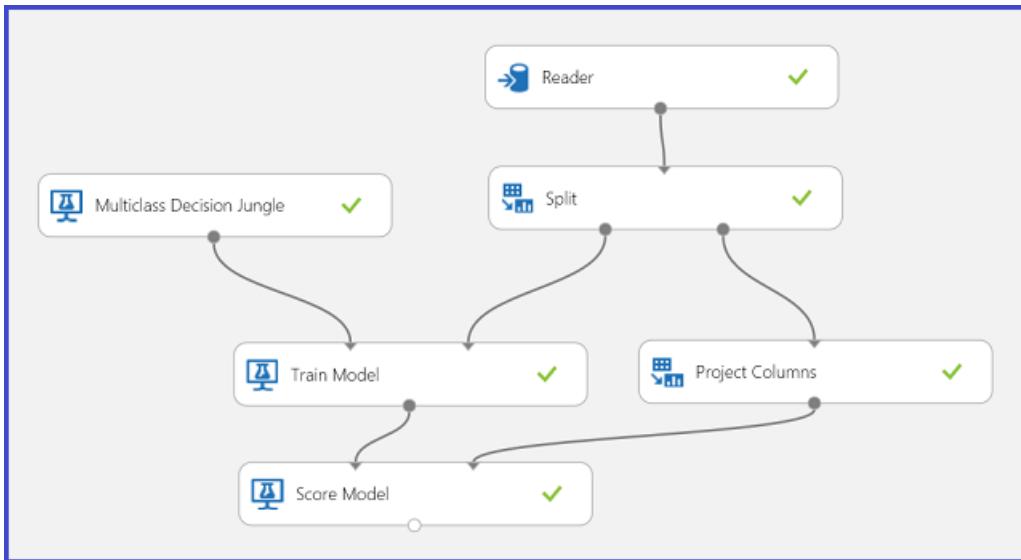


Figure 6. Letter recognition multiclass classification problem experiment

Visualizing the results from the [Score Model](#) module by clicking the output port of [Score Model](#) module and then clicking **Visualize**, you should see content as shown in Figure 7.

Col16	Col17	Scored Probabilities for Class "A"	Scored Probabilities for Class "B"	Scored Probabilities for Class "C"	Scored Probabilities for Class "D"	Scored Probabilities for Class "E"	Scored Probabilities for Class "F"	Scored Probabilities for Class "G"	Scored Probabilities for Class "H"	Scored Probabilities for Class "I"	Scored Probabilities for Class "J"	Scored Labels
2	6	0.003571	0.000451	0	0.00119	0	0.916995	0	0	0.01228	0.00169	F
8	5	0.077476	0	0	0	0.04205	0.035784	0	0.027099	0.029174	0.01563	
6	12	0	0.236149	0.006445	0.0731	0.005016	0.01138	0.038975	0.06729	0.008711	0.01713	
3	7	0.000984	0	0	0.000434	0	0.001416	0	0.02798	0.00623	0.0025	
0	8	0	0	0	0	0	0	0	0.000718	0.997776	0.00113	
8	6	0.031685	0.035017	0.002778	0.021153	0.008333	0.093765	0.002778	0.001265	0.005515	0.014991	
1	5	0.001012	0	0	0.000817	0	0.00458	0.001634	0	0	0	
4	8	0.000984	0	0	0.000434	0	0.000868	0	0.08663	0.00125	0	
8	9	0.004996	0.000273	0	0.003841	0	0.063475	0.009544	0.012195	0.090871	0.566711	
4	7	0	0	0	0.002725	0	0.005265	0	0.000883	0.00433	0	
2	5	0	0	0	0	0	0	0	0	0	0	
11	6	0.012967	0.072156	0.096037	0.240164	0	0.134389	0.010402	0.020935	0.061644	0.007225	
4	8	0	0	0	0.013801	0	0.009321	0.001582	0.002466	0.00433	0	
9	9	0.013273	0.084488	0	0.018118	0.006148	0.023945	0.026386	0.034539	0.06094	0.069495	
3	11	0.005247	0	0.124714	0.000012	0.012192	0	0.002435	0.053353	0	0.001046	
6	9	0.009623	0.447126	0.000534	0.123377	0.003	0.007448	0.005077	0.115194	0.004	0.003	
7	7	0.001147	0.024564	0.022777	0.004237	0.049034	0	0.546592	0.045465	0	0	
n	n	n	n	n	n	n	n	n	n	n	n	

Figure 7. Visualize score model results in a multi-class classification

Result interpretation

The left 16 columns represent the feature values of the test set. The columns with names like Scored Probabilities for Class "XX" are just like the Scored Probabilities column in the two-class case. They show the probability that the corresponding entry falls into a certain class. For example, for the first entry, there is 0.003571 probability that it is an "A," 0.000451 probability that it is a "B," and so forth. The last column (Scored Labels) is the same as Scored Labels in the two-class case. It selects the class with the largest scored probability as the predicted class of the corresponding entry. For example, for the first entry, the scored label is "F" since it has the largest probability to be an "F" (0.916995).

Web service publication

You can also get the scored label for each entry and the probability of the scored label. The basic logic is to find the largest probability among all the scored probabilities. To do this, you need to use the [Execute R Script](#) module. The R code is shown in Figure 8, and the result of the experiment is shown in Figure 9.

```

1 # Map 1-based optional input ports to variables
2 dataset <- maml.mapInputPort(1) # class: data.frame
3
4 # Get the Scored Labels and the corresponding probabilities
5 data.set = data.frame('Scored Labels'=dataset['Scored Labels'], 'Label Probability'=apply(dataset[,17:42], 1, max))
6
7 # Select data.frame to be sent to the output Dataset port
8 maml.mapOutputPort("data.set");

```

Figure 8. R code for extracting Scored Labels and the associated probabilities of the labels

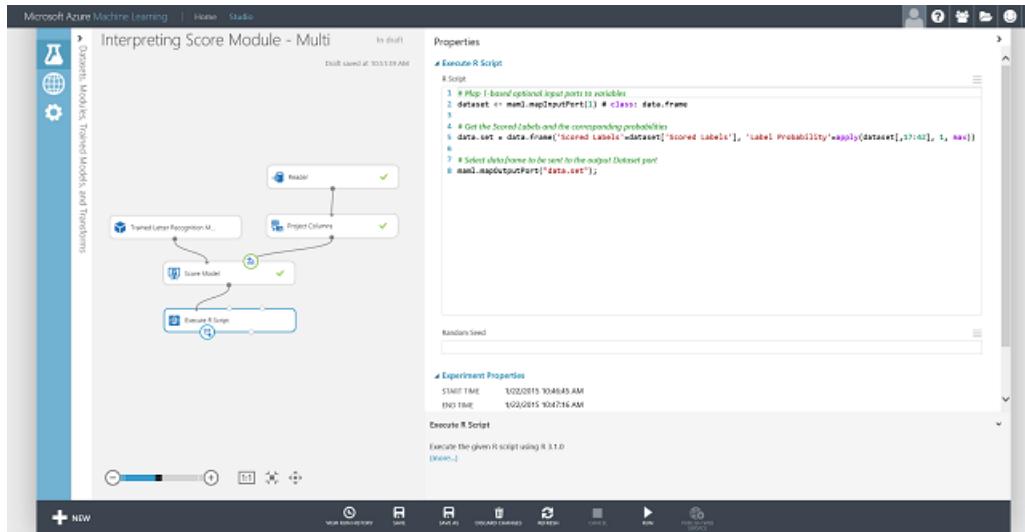


Figure 9. Final scoring experiment of the letter-recognition multiclass classification problem

After you publish and run the web service and enter some input feature values, the returned result looks like Figure 10. This hand-written letter, with its extracted 16 features, is predicted to be a "T" with 0.9715 probability.

Figure 10. Web service result of multiclass classification

Regression

Regression problems are different from classification problems. In a classification problem, you're trying to predict discrete classes, such as which class an iris flower belongs to. But as you can see in the following example of a

regression problem, you're trying to predict a continuous variable, such as the price of a car.

Example experiment

Use automobile price prediction as your example for regression. You are trying to predict the price of a car based on its features, including make, fuel type, body type, and drive wheel. The experiment is shown in Figure 11.

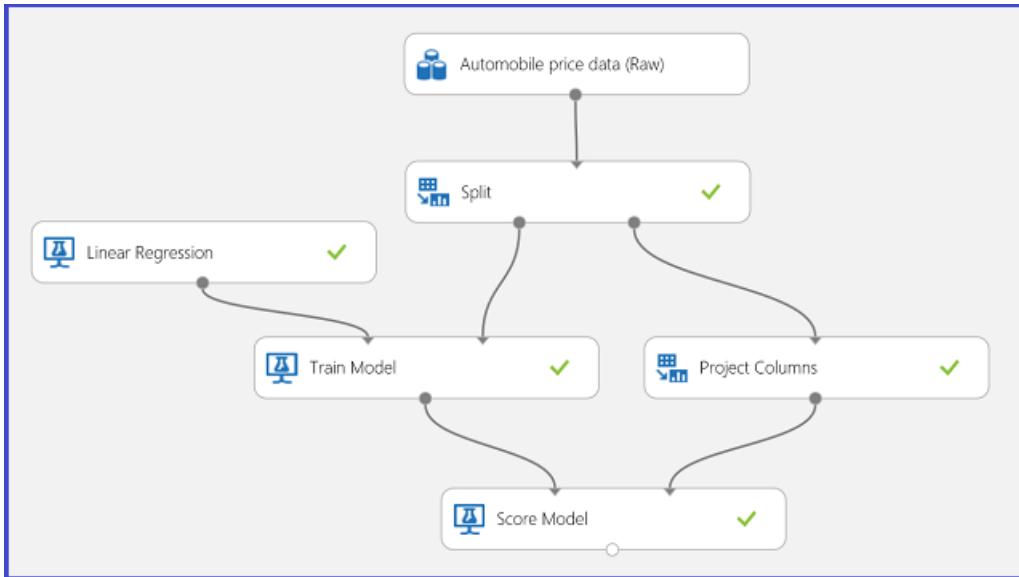


Figure 11. Automobile price regression problem experiment

Visualizing the [Score Model](#) module, the result looks like Figure 12.

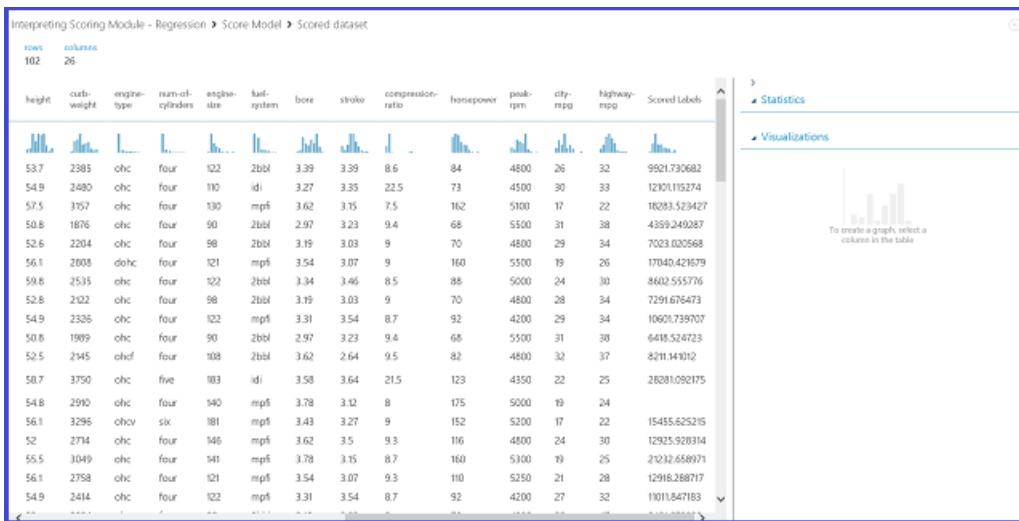


Figure 12. Scoring result for the automobile price prediction problem

Result interpretation

Scored Labels is the result column in this scoring result. The numbers are the predicted price for each car.

Web service publication

You can publish the regression experiment into a web service and call it for automobile price prediction in the same way as in the two-class classification use case.

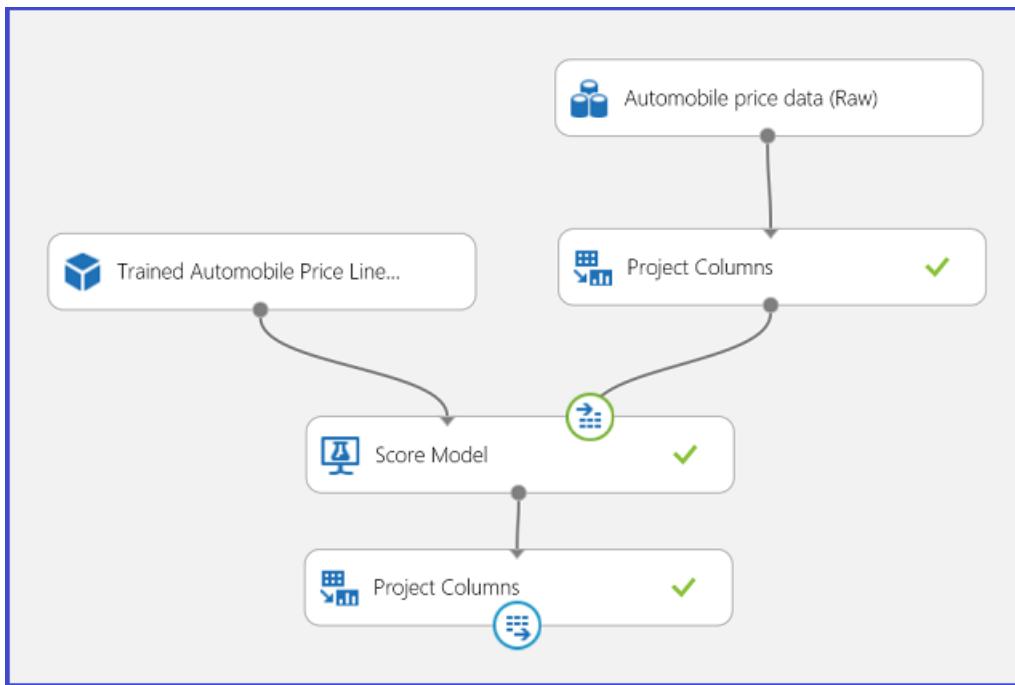


Figure 13. Scoring experiment of an automobile price regression problem

Running the web service, the returned result looks like Figure 14. The predicted price for this car is \$15,085.52.

The screenshot shows a user interface for testing a scoring module. The top part is a form titled "Test Interpreting Scoring Module - Regression - Scoring Service" with the sub-instruction "Enter data to predict". It contains several input fields for different car features:

- SYMBOLING**: Value: 2
- NORMALIZED-LOSSES**: Value: 164
- MAKE**: Value: audi
- FUEL-TYPE**: Value: gas
- ASPIRATION**: Value: std

Below the form is a message indicating the test was successful:

← 'Interpreting Scoring Module - Regression - Scoring' test finished successfully

Result: 15085.5173663723

Figure 14. Web service result of an automobile price regression problem

Clustering

Example experiment

Let's use the Iris data set again to build a clustering experiment. Here you can filter out the class labels in the data set so that it only has features and can be used for clustering. In this iris use case, specify the number of clusters to be two during the training process, which means you would cluster the flowers into two classes. The experiment is shown in Figure 15.

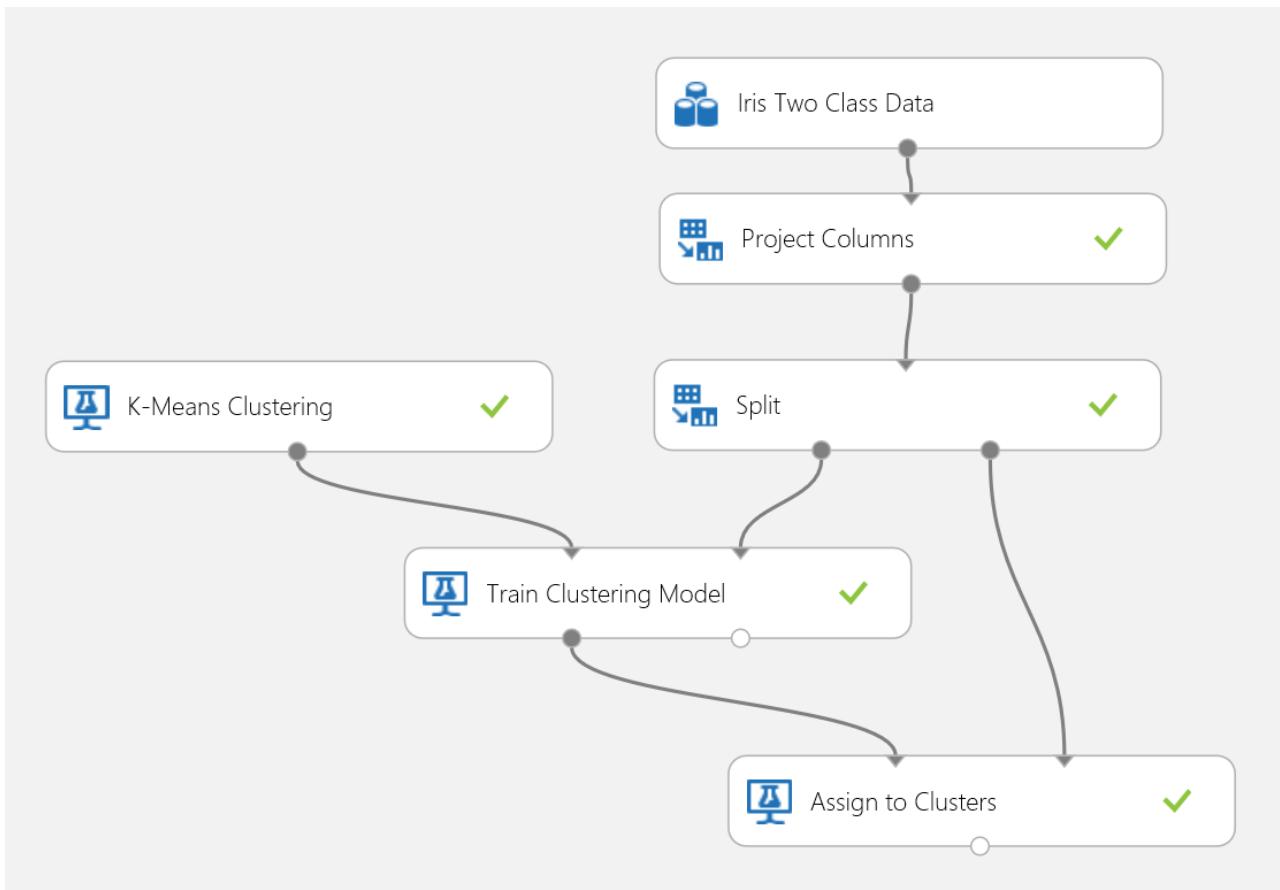


Figure 15. Iris clustering problem experiment

Clustering differs from classification in that the training data set doesn't have ground-truth labels by itself. Clustering groups the training data set instances into distinct clusters. During the training process, the model labels the entries by learning the differences between their features. After that, the trained model can be used to further classify future entries. There are two parts of the result we are interested in within a clustering problem. The first part is labeling the training data set, and the second is classifying a new data set with the trained model.

The first part of the result can be visualized by clicking the left output port of [Train Clustering Model](#) and then clicking **Visualize**. The visualization is shown in Figure 16.

Interpret Scoring Module - Clustering > Train Clustering Model > Results dataset					
rows	columns				
60	5				
view as	sepal-length	sepal-width	petal-length	petal-width	Assignments
histogram	4.9	3.1	1.5	0.1	1
histogram	5.7	4.4	1.5	0.4	1
histogram	4.8	3.4	1.6	0.2	1
histogram	6.1	3	4.9	1.8	0
histogram	6.4	2.8	5.6	2.1	0
histogram	6.2	3.4	5.4	2.3	0
histogram	5	3.5	1.6	0.6	1
histogram	5.4	3.7	1.5	0.2	1
histogram	6.3	2.7	4.9	1.8	0
histogram	7.6	3	6.6	2.1	0
histogram	7.9	3.8	6.4	2	0
histogram	6.7	3.3	5.7	2.5	0
histogram	4.9	3.6	1.4	0.1	1
histogram	4.9	3	1.4	0.2	1
histogram	7.7	3.8	6.7	2.2	0
histogram	4.5	2.3	1.3	0.3	1
histogram	5.1	3.5	1.4	0.2	1
histogram	5	3.4	1.6	0.4	1
histogram	4.3	3	1.1	0.1	1
histogram	4.9	3.1	1.5	0.2	1

Figure 16. Visualize clustering result for the training data set

The result of the second part, clustering new entries with the trained clustering model, is shown in Figure 17.

Interpret Scoring Module - Clustering > Assign to Clusters > Results dataset

rows	columns				
40	5				
view as	sepal-length	sepal-width	petal-length	petal-width	Assignments
4.4	3	1.3	0.2	1	
5.8	2.7	5.1	1.9	0	
6.4	3.2	5.3	2.3	0	
5	3	1.6	0.2	1	
4.7	3.2	1.6	0.2	1	
5	3.2	1.2	0.2	1	
6.3	2.9	5.6	1.8	0	
5.2	3.5	1.5	0.2	1	
6.1	2.6	5.6	1.4	0	
6.3	2.5	5	1.9	0	
5.8	2.8	5.1	2.4	0	
6.4	2.7	5.3	1.9	0	
6.5	3	5.2	2	0	
7.2	3.6	6.1	2.5	0	
6.3	2.8	5.1	1.5	0	
6.4	2.8	5.6	2.2	0	
5.1	3.8	1.5	0.3	1	
6.9	3.1	5.1	2.3	0	
4.4	3.2	1.3	0.2	1	
5.5	4.2	1.4	0.2	1	

To create a graph, select a column in the table

Figure 17. Visualize clustering result on a new data set

Result interpretation

Although the results of the two parts stem from different experiment stages, they look the same and are interpreted in the same way. The first four columns are features. The last column, Assignments, is the prediction result. The entries assigned the same number are predicted to be in the same cluster, that is, they share similarities in some way (this experiment uses the default Euclidean distance metric). Because you specified the number of clusters to be 2, the entries in Assignments are labeled either 0 or 1.

Web service publication

You can publish the clustering experiment into a web service and call it for clustering predictions the same way as in the two-class classification use case.

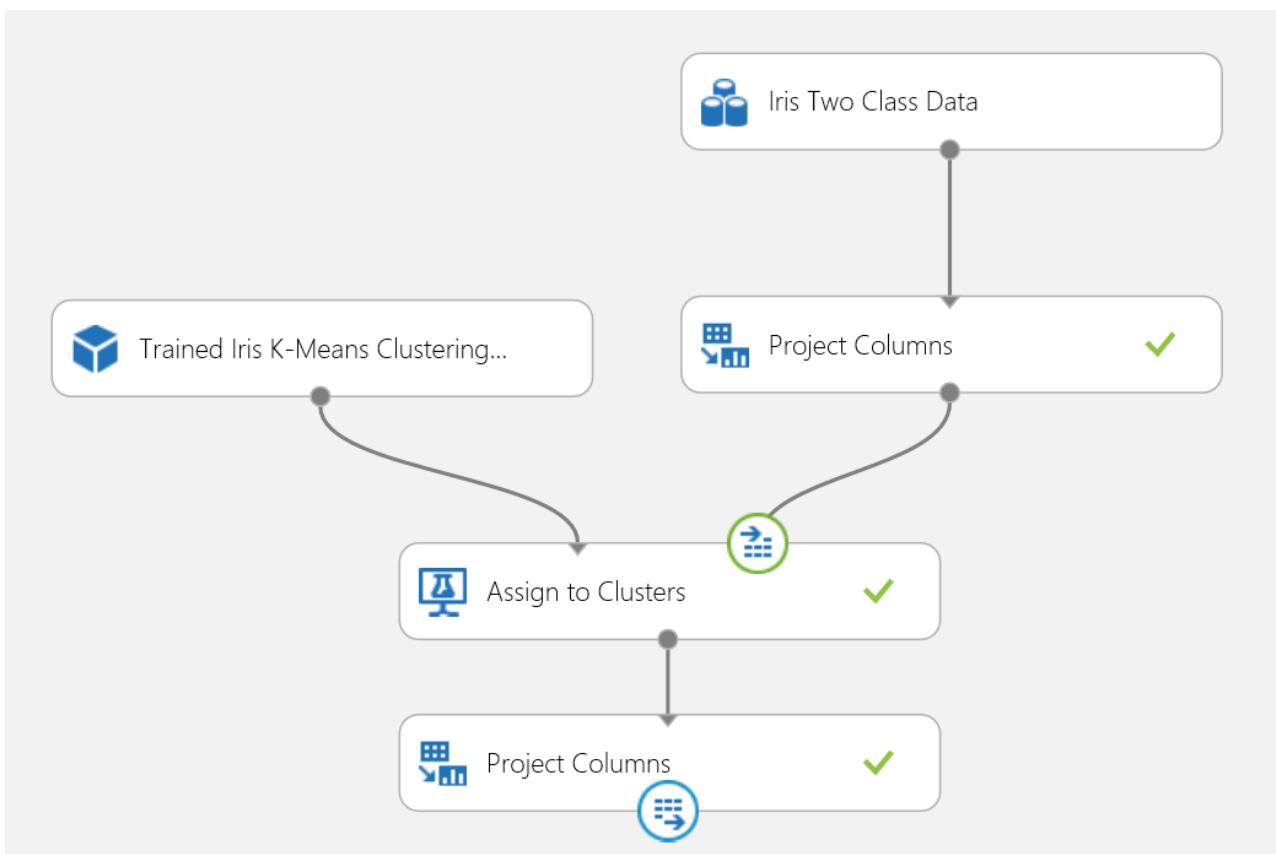


Figure 18. Scoring experiment of an iris clustering problem

After you run the web service, the returned result looks like Figure 19. This flower is predicted to be in cluster 0.

The screenshot shows a web application window titled "Test Interpret Scoring Module - Clustering - Scoring Service". At the top, there is a header "Enter data to predict". Below it, there are four input fields for the Iris flower features:

- SEPAL-LENGTH: Input value is 6.
- SEPAL-WIDTH: Input value is 3.
- PETAL-LENGTH: Input value is 5.
- PETAL-WIDTH: Input value is 2.

Below the input fields is a large green button with a checkmark icon. The main content area has a dark grey background and displays two messages:

- A left arrow icon followed by the text "'Interpret Scoring Module - Clustering - Scoring' test finished successfully".
- A green checkmark icon followed by the text "Result: 0".

Figure 19. Web service result of iris two-class classification

Recommender system

Example experiment

For recommender systems, you can use the restaurant recommendation problem as an example: you can recommend restaurants for customers based on their rating history. The input data consists of three parts:

- Restaurant ratings from customers
- Customer feature data
- Restaurant feature data

There are several things we can do with the [Train Matchbox Recommender](#) module in Azure Machine Learning:

- Predict ratings for a given user and item
- Recommend items to a given user
- Find users related to a given user
- Find items related to a given item

You can choose what you want to do by selecting from the four options in the **Recommender prediction kind** menu. Here you can walk through all four scenarios.

Properties

Score Matchbox Recommender

Recommender prediction kind

Rating Prediction

Item Recommendation

Related Users

Related Items

ELAPSED TIME 0:00:00.000

STATUS CODE Finished

STATUS DETAILS Task output was present in output cache

A typical Azure Machine Learning experiment for a recommender system looks like Figure 20. For information about how to use those recommender system modules, see [Train matchbox recommender](#) and [Score matchbox recommender](#).

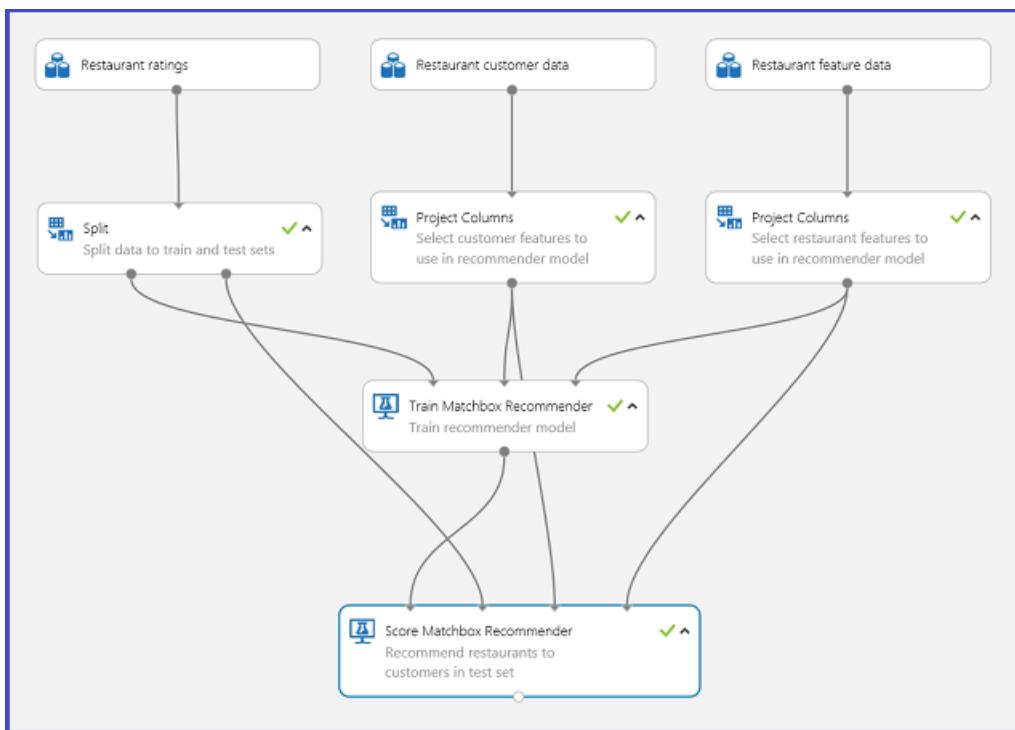


Figure 20. Recommender system experiment

Result interpretation

Predict ratings for a given user and item

By selecting **Rating Prediction** under **Recommender prediction kind**, you are asking the recommender system to predict the rating for a given user and item. The visualization of the **Score Matchbox Recommender** output looks like Figure 21.

Copy of - Sample Experiment: Recommender System > Score Matchbox Recommender > Scored dataset

rows	columns	
418	3	
User	Item	Rating
U1048	135026	2
U1048	132723	2
U1048	135065	2
U1048	135049	0
U1048	135034	2
U1117	135086	2
U1117	135018	2
U1049	132862	0
U1049	135042	0
U1049	135052	0
U1049	135032	0
U1049	135085	0
U1049	132821	0
U1049	135051	0
U1088	132830	1
U1088	135070	2
U1088	135069	2

Figure 21. Visualize the score result of the recommender system--rating prediction

The first two columns are the user-item pairs provided by the input data. The third column is the predicted rating of a user for a certain item. For example, in the first row, customer U1048 is predicted to rate restaurant 135026 as 2.

Recommend items to a given user

By selecting **Item Recommendation** under **Recommender prediction kind**, you're asking the recommender system to recommend items to a given user. The last parameter to choose in this scenario is *Recommended item selection*. The option **From Rated Items (for model evaluation)** is primarily for model evaluation during the training process. For this prediction stage, we choose **From All Items**. The visualization of the [Score Matchbox Recommender](#) output looks like Figure 22.

Copy of - Sample Experiment: Recommender System > Score Matchbox Recommender > Scored dataset

rows	columns				
User	Item 1	Item 2	Item 3	Item 4	Item 5
U1048	134986	135018	134975	135021	132862
U1117	134986	135018	132768	135075	132955
U1049	134986	135018	134976	135021	134975
U1088	135075	135057	132768	134996	132754
U1062	135018	134986	135057	134976	135075
U1035	132768	134986	135075	135018	134996
U1125	134986	135018	132768	135075	132955
U1013	134986	135018	134975	132955	132755
U1042	134986	135018	132768	135075	132922
U1123	134986	135018	132768	135075	132955
U1086	134986	135018	134975	132955	132768
U1006	134986	135018	132768	135075	135057
U1038	134986	135018	135025	132862	134975
U1053	134986	132768	135075	135018	135025
U1031	134986	135018	132768	135075	135057
U1005	134986	135018	132922	132755	132768
U1060	134986	135018	135052	134975	135025

Figure 22. Visualize score result of the recommender system--item recommendation

The first of the six columns represents the given user IDs to recommend items for, as provided by the input data. The other five columns represent the items recommended to the user in descending order of relevance. For example, in the first row, the most recommended restaurant for customer U1048 is 134986, followed by 135018, 134975, 135021, and 132862.

Find users related to a given user

By selecting **Related Users** under **Recommender prediction kind**, you're asking the recommender system to find related users to a given user. Related users are the users who have similar preferences. The last parameter to choose in this scenario is *Related user selection*. The option **From Users That Rated Items (for model evaluation)** is primarily for model evaluation during the training process. Choose **From All Users** for this prediction stage. The visualization of the [Score Matchbox Recommender](#) output looks like Figure 23.

Copy of - Sample Experiment: Recommender System > Score Matchbox Recommender > Scored dataset

rows	columns				
69	6				
User	Related User 1	Related User 2	Related User 3	Related User 4	Related User 5
U1048	U1051	U1066	U1044	U1017	U1072
U1117	U1079	U1136	U1102	U1103	U1001
U1049	U1069	U1070	U1025	U1089	U1013
U1088	U1110	U1058	U1033	U1036	U1093
U1062	U1114	U1015	U1082	U1027	U1129
U1035	U1121	U1020	U1042	U1034	U1103
U1125	U1123	U1043	U1057	U1067	U1060
U1013	U1064	U1086	U1089	U1057	U1043
U1042	U1079	U1083	U1117	U1035	U1102
U1123	U1125	U1043	U1067	U1057	U1060
U1086	U1064	U1089	U1013	U1070	U1057
U1006	U1080	U1107	U1028	U1024	U1031
U1038	U1066	U1087	U1044	U1051	U1092
U1053	U1011	U1099	U1107	U1090	U1092
U1031	U1023	U1029	U1028	U1006	U1089
U1005	U1040	U1087	U1066	U1044	U1052
U1060	U1057	U1057	U1043	U1123	U1125

view as:

Statistics

Visualizations

To create a graph, select a column in the table.

Figure 23. Visualize score results of the recommender system--related users

The first of the six columns shows the given user IDs needed to find related users, as provided by input data. The other five columns store the predicted related users of the user in descending order of relevance. For example, in the first row, the most relevant customer for customer U1048 is U1051, followed by U1066, U1044, U1017, and U1072.

Find items related to a given item

By selecting **Related Items** under **Recommender prediction kind**, you are asking the recommender system to find related items to a given item. Related items are the items most likely to be liked by the same user. The last parameter to choose in this scenario is *Related item selection*. The option **From Rated Items (for model evaluation)** is primarily for model evaluation during the training process. We choose **From All Items** for this prediction stage. The visualization of the **Score Matchbox Recommender** output looks like Figure 24.

Copy of - Sample Experiment: Recommender System > Score Matchbox Recommender > Scored dataset

rows	columns				
109	6				
User	Related Item 1	Related Item 2	Related Item 3	Related Item 4	Related Item 5
135026	135074	135035	132875	135055	134992
132723	135072	132861	135073	135000	132958
135065	132834	135054	135066	135051	132572
135049	135060	132717	135033	132951	135044
135034	134987	134999	132768	134996	135075
135088	132937	135011	135062	132958	135028
135010	132755	134986	134975	132955	134976
132862	134975	135088	135011	132955	135025
135042	132869	135039	135063	132717	132925
135052	135035	132854	135074	135032	135045
135032	135047	135026	135080	135052	135074
135085	135011	132937	132825	135088	132613
132921	135046	135058	132846	132847	135028
135051	135038	135066	135065	135085	132825
132830	135027	135108	132851	132937	135066
135070	135043	132584	132706	132613	135044
135069	132668	132715	132858	132732	132877

view as:

Statistics

Visualizations

To create a graph, select a column in the table.

Figure 24. Visualize score results of the recommender system--related items

The first of the six columns represents the given item IDs needed to find related items, as provided by the input data. The other five columns store the predicted related items of the item in descending order in terms of relevance. For example, in the first row, the most relevant item for item 135026 is 135074, followed by 135035, 132875, 135055, and 134992.

Web service publication

The process of publishing these experiments as web services to get predictions is similar for each of the four scenarios. Here we take the second scenario (recommend items to a given user) as an example. You can follow the same procedure with the other three.

Saving the trained recommender system as a trained model and filtering the input data to a single user ID column as requested, you can hook up the experiment as in Figure 25 and publish it as a web service.

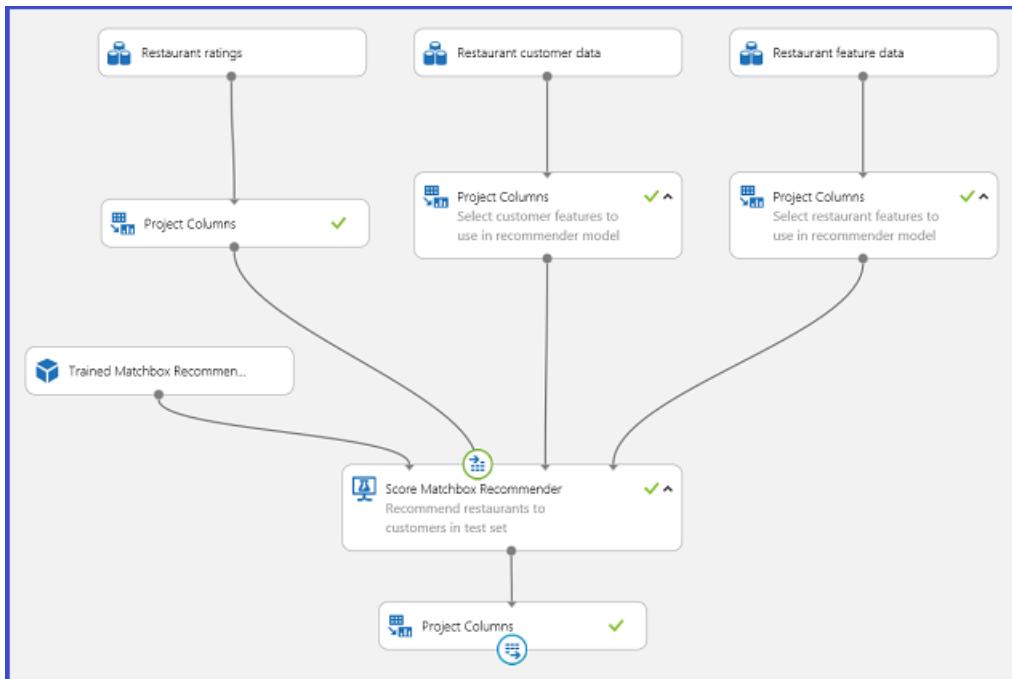


Figure 25. Scoring experiment of the restaurant recommendation problem

Running the web service, the returned result looks like Figure 26. The five recommended restaurants for user U1048 are 134986, 135018, 134975, 135021, and 132862.

The screenshot shows a web service interface. At the top, it says 'Test Copy of - Sample Experiment: Recommender System Service'. Below that, there's a form field labeled 'Enter data to predict' with a sub-label 'USERID' and a text input box containing 'U1048'. To the right of the input box is a checkmark icon. At the bottom of the page, a message box displays the results: '← 'Copy of - Sample Experiment: Recommender System' test finished successfully' and 'Result: 134986,135018,134975,135021,132862'.

Figure 26. Web service result of restaurant recommendation problem

Debug your Model in Azure Machine Learning

9/25/2017 • 2 min to read • [Edit Online](#)

This article explains the potential reasons why either of the following two failures might be encountered when running a model:

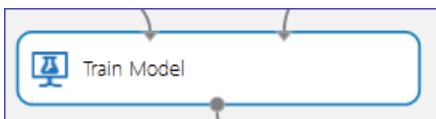
- the [Train Model](#) module produces an error
- the [Score Model](#) module produces incorrect results

NOTE

[Try Azure Machine Learning for free](#)

No credit card or Azure subscription required. [Get started now.](#)

Train Model Module produces an error



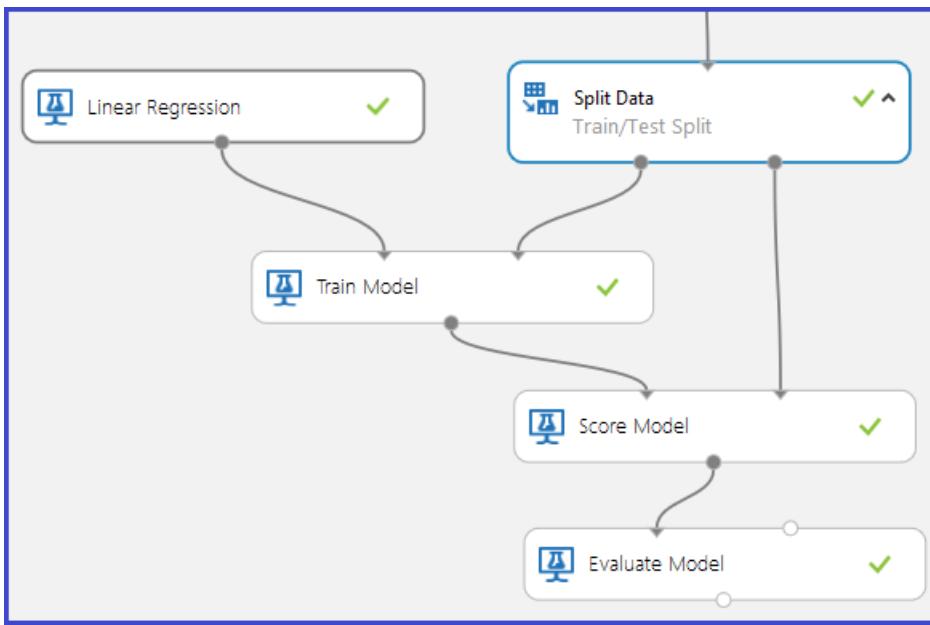
The [Train Model](#) Module expects two inputs:

1. The type of machine learning model from the collection of models provided by Azure Machine Learning.
2. The training data with a specified Label column which specifies the variable to predict (the other columns are assumed to be Features).

This module can produce an error in the following cases:

1. The Label column is specified incorrectly. This can happen if either more than one column is selected as the Label or an incorrect column index is selected. For example, the second case would apply if a column index of 30 is used with an input dataset which has only 25 columns.
2. The dataset does not contain any Feature columns. For example, if the input dataset has only one column, which is marked as the Label column, there would be no features with which to build the model. In this case, the [Train Model](#) module produces an error.
3. The input dataset (Features or Label) contains Infinity as a value.

Score Model Module produces incorrect results



In a typical training/testing experiment for supervised learning, the **Split Data** module divides the original dataset into two parts: one part is used to train the model and one part is used to score how well the trained model performs. The trained model is then used to score the test data, after which the results are evaluated to determine the accuracy of the model.

The **Score Model** module requires two inputs:

1. A trained model output from the **Train Model** module.
2. A scoring dataset that is different from the dataset used to train the model.

It's possible that even though the experiment succeeds, the **Score Model** module produces incorrect results. Several scenarios may cause this to happen:

1. If the specified Label is categorical and a regression model is trained on the data, an incorrect output would be produced by the **Score Model** module. This is because regression requires a continuous response variable. In this case, it would be more suitable to use a classification model.
2. Similarly, if a classification model is trained on a dataset having floating point numbers in the Label column, it may produce undesirable results. This is because classification requires a discrete response variable that only allows values that range over a finite, and usually somewhat small, set of classes.
3. If the scoring dataset does not contain all the features used to train the model, the **Score Model** produces an error.
4. If a row in the scoring dataset contains a missing value or an infinite value for any of its features, the **Score Model** will not produce any output corresponding to that row.
5. The **Score Model** may produce identical outputs for all rows in the scoring dataset. This could occur, for example, when attempting classification using Decision Forests if the minimum number of samples per leaf node is chosen to be more than the number of training examples available.

Extend your experiment with R

9/25/2017 • 1 min to read • [Edit Online](#)

You can extend the functionality of Azure Machine Learning Studio through the R language by using the [Execute R Script](#) module.

This module accepts multiple input datasets and yields a single dataset as output. You can type an R script into the **R Script** parameter of the [Execute R Script](#) module.

You access each input port of the module by using code similar to the following:

```
dataset1 <- maml.mapInputPort(1)
```

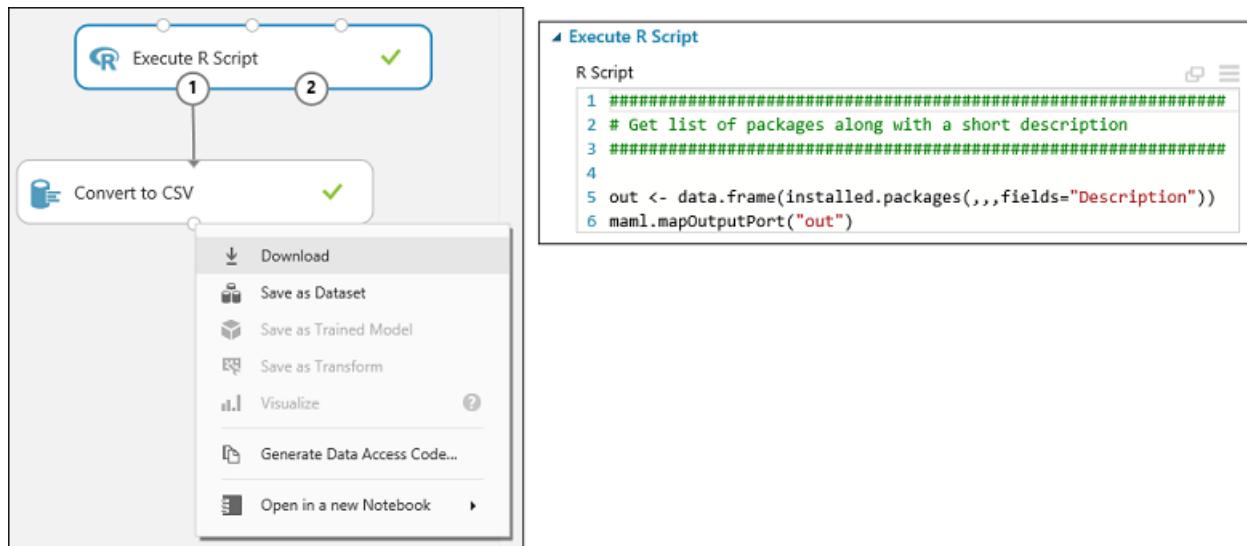
Listing all currently-installed packages

The list of installed packages can change. A list of currently installed packages can be found in [R Packages Supported by Azure Machine Learning](#).

You also can get the complete, current list of installed packages by entering the following code into the [Execute R Script](#) module:

```
out <- data.frame(installed.packages(,,fields="Description"))
maml.mapOutputPort("out")
```

This sends the list of packages to the output port of the [Execute R Script](#) module. To view the package list, connect a conversion module such as [Convert to CSV](#) to the left output of the [Execute R Script](#) module, run the experiment, then click the output of the conversion module and select **Download**.



Importing packages

You can import packages that are not already installed by using the following commands in the [Execute R Script](#) module:

```
install.packages("src/my_favorite_package.zip", lib = ".", repos = NULL, verbose = TRUE)
success <- library("my_favorite_package", lib.loc = ".", logical.return = TRUE, verbose = TRUE)
```

where the `my_favorite_package.zip` file contains your package.

NOTE

Try Azure Machine Learning for free

No credit card or Azure subscription required. [Get started now.](#)

Author custom R modules in Azure Machine Learning

11/30/2017 • 16 min to read • [Edit Online](#)

This topic describes how to author and deploy a custom R module in Azure Machine Learning. It explains what custom R modules are and what files are used to define them. It illustrates how to construct the files that define a module and how to register the module for deployment in a Machine Learning workspace. The elements and attributes used in the definition of the custom module are then described in more detail. How to use auxiliary functionality and files and multiple outputs is also discussed.

NOTE

[Try Azure Machine Learning for free](#)

No credit card or Azure subscription required. [Get started now.](#)

What is a custom R module?

A **custom module** is a user-defined module that can be uploaded to your workspace and executed as part of an Azure Machine Learning experiment. A **custom R module** is a custom module that executes a user-defined R function. R is a programming language for statistical computing and graphics that is widely used by statisticians and data scientists for implementing algorithms. Currently, R is the only language supported in custom modules, but support for additional languages is scheduled for future releases.

Custom modules have **first-class status** in Azure Machine Learning in the sense that they can be used just like any other module. They can be executed with other modules, included in published experiments or in visualizations. You have control over the algorithm implemented by the module, the input and output ports to be used, the modeling parameters, and other various runtime behaviors. An experiment that contains custom modules can also be published into the Cortana Intelligence Gallery for easy sharing.

Files in a custom R module

A custom R module is defined by a .zip file that contains, at a minimum, two files:

- A **source file** that implements the R function exposed by the module
- An **XML definition file** that describes the custom module interface

Additional auxiliary files can also be included in the .zip file that provides functionality that can be accessed from the custom module. This option is discussed in the **Arguments** part of the reference section **Elements in the XML definition file** following the quickstart example.

Quickstart example: define, package, and register a custom R module

This example illustrates how to construct the files required by a custom R module, package them into a zip file, and then register the module in your Machine Learning workspace. The example zip package and sample files can be downloaded from [Download CustomAddRows.zip file](#).

The source file

Consider the example of a **Custom Add Rows** module that modifies the standard implementation of the **Add Rows** module used to concatenate rows (observations) from two datasets (data frames). The standard **Add Rows**

module appends the rows of the second input dataset to the end of the first input dataset using the `rbind` algorithm. The customized `CustomAddRows` function similarly accepts two datasets, but also accepts a Boolean swap parameter as an additional input. If the swap parameter is set to **FALSE**, it returns the same data set as the standard implementation. But if the swap parameter is **TRUE**, the function appends rows of first input dataset to the end of the second dataset instead. The `CustomAddRows.R` file that contains the implementation of the R `CustomAddRows` function exposed by the **Custom Add Rows** module has the following R code.

```
CustomAddRows <- function(dataset1, dataset2, swap=FALSE)
{
  if(swap)
  {
    return (rbind(dataset2, dataset1));
  }
  else
  {
    return (rbind(dataset1, dataset2));
  }
}
```

The XML definition file

To expose this `CustomAddRows` function as an Azure Machine Learning module, an XML definition file must be created to specify how the **Custom Add Rows** module should look and behave.

```
<!-- Defined a module using an R Script -->
<Module name="Custom Add Rows">
  <Owner>Microsoft Corporation</Owner>
  <Description>Appends one dataset to another. Dataset 2 is concatenated to Dataset 1 when Swap is FALSE, and vice versa when Swap is TRUE.</Description>

  <!-- Specify the base language, script file and R function to use for this module. -->
  <Language name="R"
    sourceFile="CustomAddRows.R"
    entryPoint="CustomAddRows" />

  <!-- Define module input and output ports -->
  <!-- Note: The values of the id attributes in the Input and Arg elements must match the parameter names in the R Function CustomAddRows defined in CustomAddRows.R. -->
  <Ports>
    <Input id="dataset1" name="Dataset 1" type="DataTable">
      <Description>First input dataset</Description>
    </Input>
    <Input id="dataset2" name="Dataset 2" type="DataTable">
      <Description>Second input dataset</Description>
    </Input>
    <Output id="dataset" name="Dataset" type="DataTable">
      <Description>The combined dataset</Description>
    </Output>
  </Ports>

  <!-- Define module parameters -->
  <Arguments>
    <Arg id="swap" name="Swap" type="bool" >
      <Description>Swap input datasets.</Description>
    </Arg>
  </Arguments>
</Module>
```

It is critical to note that the value of the **id** attributes of the **Input** and **Arg** elements in the XML file must match the function parameter names of the R code in the `CustomAddRows.R` file EXACTLY: (`dataset1`, `dataset2`, and `swap` in the example). Similarly, the value of the **entryPoint** attribute of the **Language** element must match the name of the function in the R script EXACTLY: (`CustomAddRows` in the example).

In contrast, the **id** attribute for the **Output** element does not correspond to any variables in the R script. When more than one output is required, simply return a list from the R function with results placed *in the same order as Outputs* elements are declared in the XML file.

Package and register the module

Save these two files as *CustomAddRows.R* and *CustomAddRows.xml* and then zip the two files together into a *CustomAddRows.zip* file.

To register them in your Machine Learning workspace, go to your workspace in the Machine Learning Studio, click the **+NEW** button on the bottom and choose **MODULE -> FROM ZIP PACKAGE** to upload the new **Custom Add Rows** module.



The **Custom Add Rows** module is now ready to be accessed by your Machine Learning experiments.

Elements in the XML definition file

Module elements

The **Module** element is used to define a custom module in the XML file. Multiple modules can be defined in one XML file using multiple **module** elements. Each module in your workspace must have a unique name. Register a custom module with the same name as an existing custom module and it replaces the existing module with the new one. Custom modules can, however, be registered with the same name as an existing Azure Machine Learning module. If so, they appear in the **Custom** category of the module palette.

```
<Module name="Custom Add Rows" isDeterministic="false">
  <Owner>Microsoft Corporation</Owner>
  <Description>Appends one dataset to another...</Description>/>
```

Within the **Module** element, you can specify two additional optional elements:

- an **Owner** element that is embedded into the module
- a **Description** element that contains text that is displayed in quick help for the module and when you hover over the module in the Machine Learning UI.

Rules for characters limits in the Module elements:

- The value of the **name** attribute in the **Module** element must not exceed 64 characters in length.
- The content of the **Description** element must not exceed 128 characters in length.
- The content of the **Owner** element must not exceed 32 characters in length.

A module's results can be deterministic or nondeterministic.** By default, all modules are considered to be deterministic. That is, given an unchanging set of input parameters and data, the module should return the same results each time it is run. Given this behavior, Azure Machine Learning Studio only reruns modules marked as deterministic if a parameter or the input data has changed. Returning the cached results also provides much faster execution of experiments.

There are functions that are nondeterministic, such as RAND or a function that returns the current date or time. If your module uses a nondeterministic function, you can specify that the module is non-deterministic by setting the optional **isDeterministic** attribute to **FALSE**. This insures that the module is rerun whenever the experiment is run,

even if the module input and parameters have not changed.

Language Definition

The **Language** element in your XML definition file is used to specify the custom module language. Currently, R is the only supported language. The value of the **sourceFile** attribute must be the name of the R file that contains the function to call when the module is run. This file must be part of the zip package. The value of the **entryPoint** attribute is the name of the function being called and must match a valid function defined with in the source file.

```
<Language name="R" sourceFile="CustomAddRows.R" entryPoint="CustomAddRows" />
```

Ports

The input and output ports for a custom module are specified in child elements of the **Ports** section of the XML definition file. The order of these elements determines the layout experienced (UX) by users. The first child **input** or **output** listed in the **Ports** element of the XML file becomes the left-most input port in the Machine Learning UX. Each input and output port may have an optional **Description** child element that specifies the text shown when you hover the mouse cursor over the port in the Machine Learning UI.

Ports Rules:

- Maximum number of **input and output ports** is 8 for each.

Input elements

Input ports allow you to pass data to your R function and workspace. The **data types** that are supported for input ports are as follows:

DataTable: This type is passed to your R function as a data.frame. In fact, any types (for example, CSV files or ARFF files) that are supported by Machine Learning and that are compatible with **DataTable** are converted to a data.frame automatically.

```
<Input id="dataset1" name="Input 1" type="DataTable" isOptional="false">
  <Description>Input Dataset 1</Description>
</Input>
```

The **id** attribute associated with each **DataTable** input port must have a unique value and this value must match its corresponding named parameter in your R function. Optional **DataTable** ports that are not passed as input in an experiment have the value **NULL** passed to the R function and optional zip ports are ignored if the input is not connected. The **isOptional** attribute is optional for both the **DataTable** and **Zip** types and is *false* by default.

Zip: Custom modules can accept a zip file as input. This input is unpacked into the R working directory of your function

```
<Input id="zippedData" name="Zip Input" type="Zip" IsOptional="false">
  <Description>Zip files to be extracted to the R working directory.</Description>
</Input>
```

For custom R modules, the id for a Zip port does not have to match any parameters of the R function. This is because the zip file is automatically extracted to the R working directory.

Input Rules:

- The value of the **id** attribute of the **Input** element must be a valid R variable name.
- The value of the **id** attribute of the **Input** element must not be longer than 64 characters.
- The value of the **name** attribute of the **Input** element must not be longer than 64 characters.
- The content of the **Description** element must not be longer than 128 characters

- The value of the **type** attribute of the **Input** element must be *Zip* or *DataTable*.
- The value of the **isOptional** attribute of the **Input** element is not required (and is *false* by default when not specified); but if it is specified, it must be *true* or *false*.

Output elements

Standard output ports: Output ports are mapped to the return values from your R function, which can then be used by subsequent modules. *DataTable* is the only standard output port type supported currently. (Support for *Learners* and *Transforms* is forthcoming.) A *DataTable* output is defined as:

```
<Output id="dataset" name="Dataset" type="DataTable">
  <Description>Combined dataset</Description>
</Output>
```

For outputs in custom R modules, the value of the **id** attribute does not have to correspond with anything in the R script, but it must be unique. For a single module output, the return value from the R function must be a *data.frame*. In order to output more than one object of a supported data type, the appropriate output ports need to be specified in the XML definition file and the objects need to be returned as a list. The output objects are assigned to output ports from left to right, reflecting the order in which the objects are placed in the returned list.

For example, if you want to modify the **Custom Add Rows** module to output the original two datasets, *dataset1* and *dataset2*, in addition to the new joined dataset, *dataset*, (in an order, from left to right, as: *dataset*, *dataset1*, *dataset2*), then define the output ports in the CustomAddRows.xml file as follows:

```
<Ports>
  <Output id="dataset" name="Dataset Out" type="DataTable">
    <Description>New Dataset</Description>
  </Output>
  <Output id="dataset1_out" name="Dataset 1 Out" type="DataTable">
    <Description>First Dataset</Description>
  </Output>
  <Output id="dataset2_out" name="Dataset 2 Out" type="DataTable">
    <Description>Second Dataset</Description>
  </Output>
  <Input id="dataset1" name="Dataset 1" type="DataTable">
    <Description>First Input Table</Description>
  </Input>
  <Input id="dataset2" name="Dataset 2" type="DataTable">
    <Description>Second Input Table</Description>
  </Input>
</Ports>
```

And return the list of objects in a list in the correct order in 'CustomAddRows.R':

```
CustomAddRows <- function(dataset1, dataset2, swap=FALSE) {
  if(swap) { dataset <- rbind(dataset2, dataset1) }
  else { dataset <- rbind(dataset1, dataset2) }
  return (list(dataset, dataset1, dataset2))
}
```

Visualization output: You can also specify an output port of type *Visualization*, which displays the output from the R graphics device and console output. This port is not part of the R function output and does not interfere with the order of the other output port types. To add a visualization port to the custom modules, add an **Output** element with a value of *Visualization* for its **type** attribute:

```
<Output id="deviceOutput" name="View Port" type="Visualization">
<Description>View the R console graphics device output.</Description>
</Output>
```

Output Rules:

- The value of the **id** attribute of the **Output** element must be a valid R variable name.
- The value of the **id** attribute of the **Output** element must not be longer than 32 characters.
- The value of the **name** attribute of the **Output** element must not be longer than 64 characters.
- The value of the **type** attribute of the **Output** element must be *Visualization*.

Arguments

Additional data can be passed to the R function via module parameters which are defined in the **Arguments** element. These parameters appear in the rightmost properties pane of the Machine Learning UI when the module is selected. Arguments can be any of the supported types or you can create a custom enumeration when needed.

Similar to the **Ports** elements, **Arguments** elements can have an optional **Description** element that specifies the text that appears when you hover the mouse over the parameter name. Optional properties for a module, such as **defaultValue**, **minValue**, and **maxValue** can be added to any argument as attributes to a **Properties** element. Valid properties for the **Properties** element depend on the argument type and are described with the supported argument types in the next section. Arguments with the **isOptional** property set to "true" do not require the user to enter a value. If a value is not provided to the argument, then the argument is not passed to the entry point function. Arguments of the entry point function that are optional need to be explicitly handled by the function, e.g. assigned a default value of NULL in the entry point function definition. An optional argument will only enforce the other argument constraints, i.e. min or max, if a value is provided by the user. As with inputs and outputs, it is critical that each of the parameters have unique id values associated with them. In our quick start example the associated id/parameter was *swap*.

Arg element

A module parameter is defined using the **Arg** child element of the **Arguments** section of the XML definition file. As with the child elements in the **Ports** section, the ordering of parameters in the **Arguments** section defines the layout encountered in the UX. The parameters appear from top down in the UI in the same order in which they are defined in the XML file. The types supported by Machine Learning for parameters are listed here.

int – an Integer (32-bit) type parameter.

```
<Arg id="intValue1" name="Int Param" type="int">
<Properties min="0" max="100" default="0" />
<Description>Integer Parameter</Description>
</Arg>
```

- *Optional Properties:* **min**, **max**, **default** and **isOptional**

double – a double type parameter.

```
<Arg id="doubleValue1" name="Double Param" type="double">
<Properties min="0.000" max="0.999" default="0.3" />
<Description>Double Parameter</Description>
</Arg>
```

- *Optional Properties:* **min**, **max**, **default** and **isOptional**

bool – a Boolean parameter that is represented by a check-box in UX.

```
<Arg id="boolValue1" name="Boolean Param" type="bool">
  <Properties default="true" />
  <Description>Boolean Parameter</Description>
</Arg>
```

- *Optional Properties:* **default** - false if not set

string: a standard string

```
<Arg id="stringValue1" name="My string Param" type="string">
  <Properties isOptional="true" />
  <Description>String Parameter 1</Description>
</Arg>
```

- *Optional Properties:* **default** and **isOptional**

ColumnPicker: a column selection parameter. This type renders in the UX as a column chooser. The **Property** element is used here to specify the id of the port from which columns are selected, where the target port type must be *DataTable*. The result of the column selection is passed to the R function as a list of strings containing the selected column names.

```
<Arg id="colset" name="Column set" type="ColumnPicker">
  <Properties portId="datasetIn1" allowedTypes="Numeric" default="NumericAll"/>
  <Description>Column set</Description>
</Arg>
```

- *Required Properties:* **portId** - matches the id of an Input element with type *DataTable*.
- *Optional Properties:*
 - **allowedTypes** - Filters the column types from which you can pick. Valid values include:
 - Numeric
 - Boolean
 - Categorical
 - String
 - Label
 - Feature
 - Score
 - All
 - **default** - Valid default selections for the column picker include:
 - None
 - NumericFeature
 - NumericLabel
 - NumericScore
 - NumericAll
 - BooleanFeature
 - BooleanLabel
 - BooleanScore
 - BooleanAll
 - CategoricalFeature
 - CategoricalLabel

```
  <Properties portId="datasetIn1" allowedTypes="Numeric" default="NumericAll"/>
  <Description>Column set</Description>
</Arg>
```

- CategoricalScore
- CategoricalAll
- StringFeature
- StringLabel
- StringScore
- StringAll
- AllLabel
- AllFeature
- AllScore
- All

DropDown: a user-specified enumerated (dropdown) list. The dropdown items are specified within the **Properties** element using an **Item** element. The **id** for each **Item** must be unique and a valid R variable. The value of the **name** of an **Item** serves as both the text that you see and the value that is passed to the R function.

```
<Arg id="color" name="Color" type="DropDown">
<Properties default="red">
<Item id="red" name="Red Value"/>
<Item id="green" name="Green Value"/>
<Item id="blue" name="Blue Value"/>
</Properties>
<Description>Select a color.</Description>
</Arg>
```

- *Optional Properties:*
 - **default** - The value for the default property must correspond with an id value from one of the **Item** elements.

Auxiliary Files

Any file that is placed in your custom module ZIP file is going to be available for use during execution time. Any directory structures present are preserved. This means that file sourcing works the same locally and in Azure Machine Learning execution.

NOTE

Notice that all files are extracted to 'src' directory so all paths should have 'src/' prefix.

For example, say you want to remove any rows with NAs from the dataset, and also remove any duplicate rows, before outputting it into CustomAddRows, and you've already written an R function that does that in a file RemoveDupNARows.R:

```
RemoveDupNARows <- function(dataFrame) {
  #Remove Duplicate Rows:
  dataFrame <- unique(dataFrame)
  #Remove Rows with NAs:
  finalDataFrame <- dataFrame[complete.cases(dataFrame),]
  return(finalDataFrame)
}
```

You can source the auxiliary file RemoveDupNARows.R in the CustomAddRows function:

```
CustomAddRows <- function(dataset1, dataset2, swap=FALSE) {  
  source("src/RemoveDupNARows.R")  
  if(swap) {  
    dataset <- rbind(dataset2, dataset1)  
  } else {  
    dataset <- rbind(dataset1, dataset2)  
  }  
  dataset <- removeDupNARows(dataset)  
  return (dataset)  
}
```

Next, upload a zip file containing 'CustomAddRows.R', 'CustomAddRows.xml', and 'RemoveDupNARows.R' as a custom R module.

Execution Environment

The execution environment for the R script uses the same version of R as the **Execute R Script** module and can use the same default packages. You can also add additional R packages to your custom module by including them in the custom module zip package. Just load them in your R script as you would in your own R environment.

Limitations of the execution environment include:

- Non-persistent file system: Files written when the custom module is run are not persisted across multiple runs of the same module.
- No network access

Execute Python machine learning scripts in Azure Machine Learning Studio

11/30/2017 • 9 min to read • [Edit Online](#)

This topic describes the design principles underlying the current support for Python scripts in Azure Machine Learning. The main capabilities provided are also outlined, including:

- execute basic usage scenarios
- score an experiment in a web service
- support for importing existing code
- export visualizations
- perform supervised feature selection
- understand some limitations

Python is an indispensable tool in the tool chest of many data scientists. It has:

- an elegant and concise syntax,
- cross-platform support,
- a vast collection of powerful libraries, and
- mature development tools.

Python is being used in all phases of a workflow typically used in machine learning modeling:

- data ingest and processing
- feature construction
- model training
- model validation
- deployment of the models

Azure Machine Learning Studio supports embedding Python scripts into various parts of a machine learning experiment and also seamlessly publishing them as web services on Microsoft Azure.

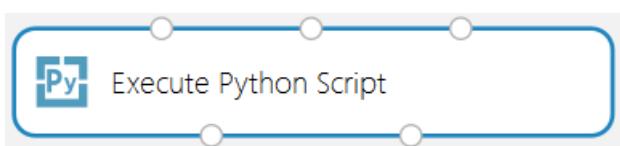
NOTE

[Try Azure Machine Learning for free](#)

No credit card or Azure subscription required. [Get started now.](#)

Design principles of Python scripts in Machine Learning

The primary interface to Python in Azure Machine Learning Studio is via the [Execute Python Script](#) module shown in Figure 1.



```

1 def azureml_main(dataframe1 = None, data
2 # Code to populate the result
3     result = pandas.DataFrame(...)
4     return result,
5

```

Figure 1. The **Execute Python Script** module.

The [Execute Python Script](#) module in Azure ML Studio accepts up to three inputs and produces up to two outputs (discussed in the following section), like its R analogue, the [Execute R Script](#) module. The Python code to be executed is entered into the parameter box as a specially named entry-point function called `azureml_main`. Here are the key design principles used to implement this module:

1. *Must be idiomatic for Python users.* Most Python users factor their code as functions inside modules. So putting a lot of executable statements in a top-level module is relatively rare. As a result, the script box also takes a specially named Python function as opposed to just a sequence of statements. The objects exposed in the function are standard Python library types such as [Pandas](#) data frames and [NumPy](#) arrays.
2. *Must have high-fidelity between local and cloud executions.* The backend used to execute the Python code is based on [Anaconda](#), a widely used cross-platform scientific Python distribution. It comes with close to 200 of the most common Python packages. Therefore, data scientists can debug and qualify their code on their local Azure Machine Learning-compatible Anaconda environment. Then use an existing development environment, such as [IPython](#) notebook or [Python Tools for Visual Studio](#), to run it as part of an Azure ML experiment. The `azureml_main` entry point is a vanilla Python function and so ****can be authored without Azure ML-specific code or the SDK installed.
3. *Must be seamlessly composable with other Azure Machine Learning modules.* The [Execute Python Script](#) module accepts, as inputs and outputs, standard Azure Machine Learning datasets. The underlying framework transparently and efficiently bridges the Azure ML and Python runtimes. So Python can be used in conjunction with existing Azure ML workflows, including those that call into R and SQLite. A result, data scientist could compose workflows that:
 - use Python and Pandas for data pre-processing and cleaning
 - feed the data to a SQL transformation, joining multiple datasets to form features
 - train models using the algorithms in Azure Machine Learning
 - evaluate and post-process the results using R.

Basic usage scenarios in ML for Python scripts

In this section, we survey some of the basic uses of the [Execute Python Script](#) module. Inputs to the Python module are exposed as Pandas data frames. The function must return a single Pandas data frame packaged inside of a Python [sequence](#) such as a tuple, list, or NumPy array. The first element of this sequence is then returned in the first output port of the module. This scheme is shown in Figure 2.

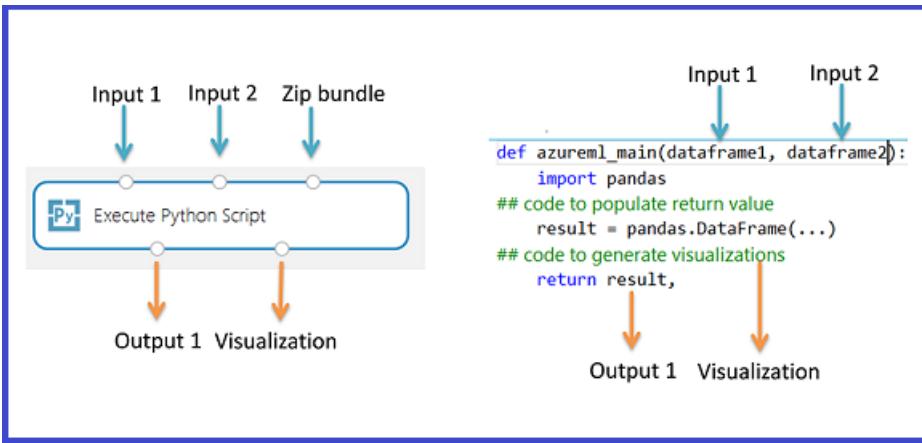


Figure 2. Mapping of input ports to parameters and return value to output port.

More detailed semantics of how the input ports get mapped to parameters of the `azureml_main` function are shown in Table 1:

Input port configuration	Python signature	Remarks
	<code>def azureml_main(): pass</code>	Can also be any function with <i>all optional</i> parameters
	<code>def azureml_main(dataframe1): pass</code>	Can also be any function with the <i>first parameter</i> being a data frame and all other parameters being optional.
	<code>def azureml_main(dataframe1): pass</code>	Can also be any function with the <i>first parameter</i> being a data frame and all other parameters being optional.
	<code>def azureml_main(dataframe1, dataframe2): pass</code>	Can also be any function with the <i>first two parameters</i> being data frames and all other parameters being optional.

Table 1. Mapping of input ports to function parameters.

The mapping between input ports and function parameters is positional:

- The first connected input port is mapped to the first parameter of the function.
- The second input (if connected) is mapped to the second parameter of the function.

See *Python for Data Analysis* (O'Reilly, 2012) by W. McKinney for more information on Python Pandas and on how it can be used to manipulate data effectively and efficiently.

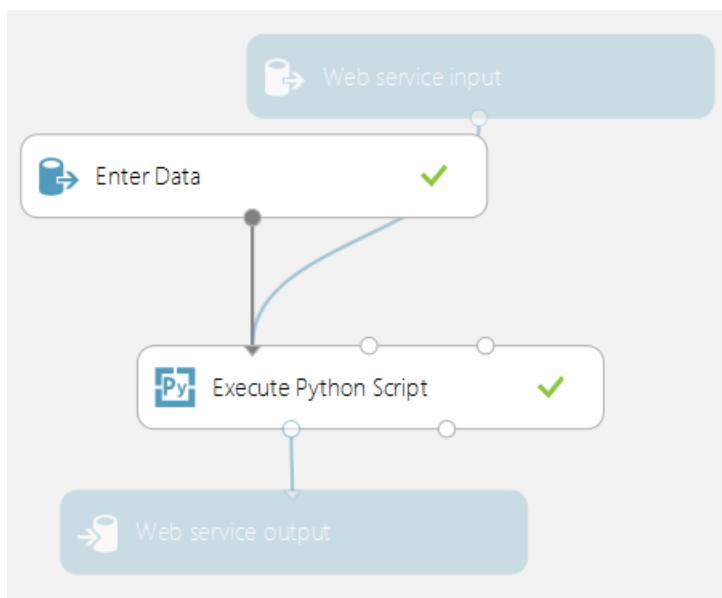
Translation of input and output types

Input datasets in Azure ML are converted to data frames in Pandas. Output data frames are converted back to Azure ML datasets. The following conversions are performed:

1. String and numeric columns are converted as-is and missing values in a dataset are converted to 'NA' values in Pandas. The same conversion happens on the way back (NA values in Pandas are converted to missing values in Azure ML).
2. Index vectors in Pandas are not supported in Azure ML. All input data frames in the Python function always have a 64-bit numerical index from 0 to the number of rows minus 1.
3. Azure ML datasets cannot have duplicate column names and column names that are not strings. If an output data frame contains non-numeric columns, the framework calls `str` on the column names. Likewise, any duplicate column names are automatically mangled to insure the names are unique. The suffix (2) is added to the first duplicate, (3) to the second duplicate, and so on.

Operationalizing Python scripts

Any [Execute Python Script](#) modules used in a scoring experiment are called when published as a web service. For example, Figure 3 shows a scoring experiment that contains the code to evaluate a single Python expression.



Python script

```
1 def azureml_main(expr_as_frame):  
2     import pandas as pd  
3     expr = expr_as_frame.iat[0,0]  
4     result = pd.DataFrame({'Expr': [expr], \  
5                           'Result': [eval(expr)]})  
6     return result,
```

Figure 3. Web service for evaluating a Python expression.

A web service created from this experiment:

- takes as input a Python expression (as a string)
- sends it to the Python interpreter
- returns a table containing both the expression and the evaluated result.

Importing existing Python script modules

A common use-case for many data scientists is to incorporate existing Python scripts into Azure ML experiments. Instead of requiring that all code be concatenated and pasted into a single script box, the [Execute Python Script](#) module accepts a zip file that contains Python modules at the third input port. The file is unzipped by the execution framework at runtime and the contents are added to the library path of the Python interpreter. The `azureml_main` entry point function can then import these modules directly.

As an example, consider the file Hello.py containing a simple "Hello, World" function.

```
def print_hello(name):  
    print "Hello ", name
```

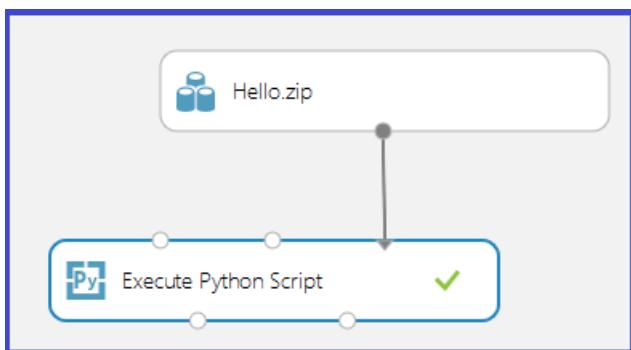
Figure 4. User-defined function in Hello.py file.

Next, we create a file Hello.zip that contains Hello.py:

Hello.zip					
Name	Type	Compressed size	Password ...	Size	Ratio
Hello.py	PY File	1 KB	No	1 KB	15%

Figure 5. Zip file containing user-defined Python code.

Upload the zip file as a dataset into Azure Machine Learning Studio. Then create and run an experiment that uses the Python code in the Hello.zip file by attaching it to the third input port of the **Execute Python Script** module, as shown in this figure.



Python script

```

1 # The script MUST include the following function,
2 # which is the entry point for this module:
3 # Param<dataframe1>: a pandas.DataFrame
4 # Param<dataframe2>: a pandas.DataFrame
5 def azureml_main():
6     import pandas as pd
7     import Hello
8     Hello.print_hello("World")
9     return pd.DataFrame([{"Output":}),
10
11
  
```

Figure 6. Sample experiment with user-defined Python code uploaded as a zip file.

The module output shows that the zip file has been unpackaged and that the function `print_hello` has been run.

[ModuleOutput]	Extracting Script Bundle.zip to .\Script Bundle	Modified	Size
[ModuleOutput]	File Name		
[ModuleOutput]	Hello.py	2015-02-09 14:26:26	48
[ModuleOutput]	Hello_World		

Figure 7. User-defined function in use inside the **Execute Python Script** module.

Working with visualizations

Plots created using Matplotlib that can be visualized on the browser can be returned by the **Execute Python Script**. But the plots are not automatically redirected to images as they are when using R. So the user must explicitly save any plots to PNG files if they are to be returned back to Azure Machine Learning.

To generate images from Matplotlib, you must complete the following procedure:

- switch the backend to "AGG" from the default Qt-based renderer
- create a new figure object
- get the axis and generate all plots into it
- save the figure to a PNG file

This process is illustrated in the following Figure 8 that creates a scatter plot matrix using the scatter_matrix function in Pandas.

```
def azureml_main(dataframe1):
    import matplotlib
    matplotlib.use("agg") ← Change backend

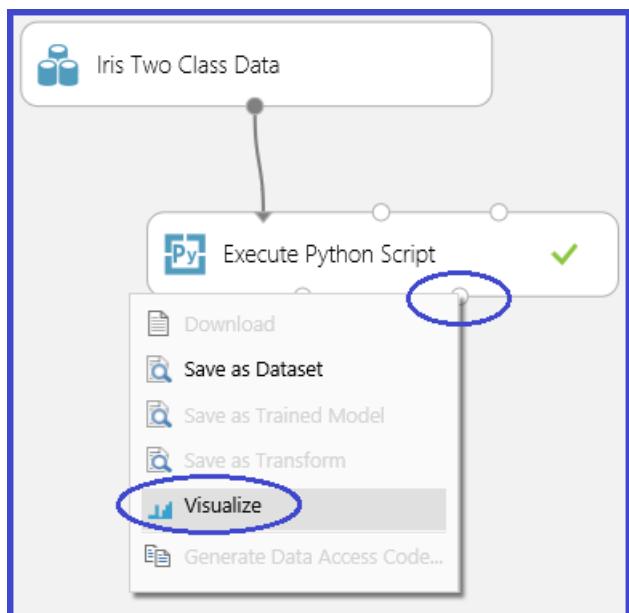
    from pandas.tools.plotting import scatter_matrix
    import matplotlib.pyplot as plt

    fig = plt.figure(); ← Create new figure
    ax = fig.gca()
    scatter_matrix(dataframe1, ax=ax) ← Plot into specified axis

    fig.savefig("scatter.png") ← Save figure to image
    return dataframe1,
```

Figure 8. Code to save Matplotlib figures to images.

Figure 9 shows an experiment that uses the script shown previously to return plots via the second output port.



Graphics

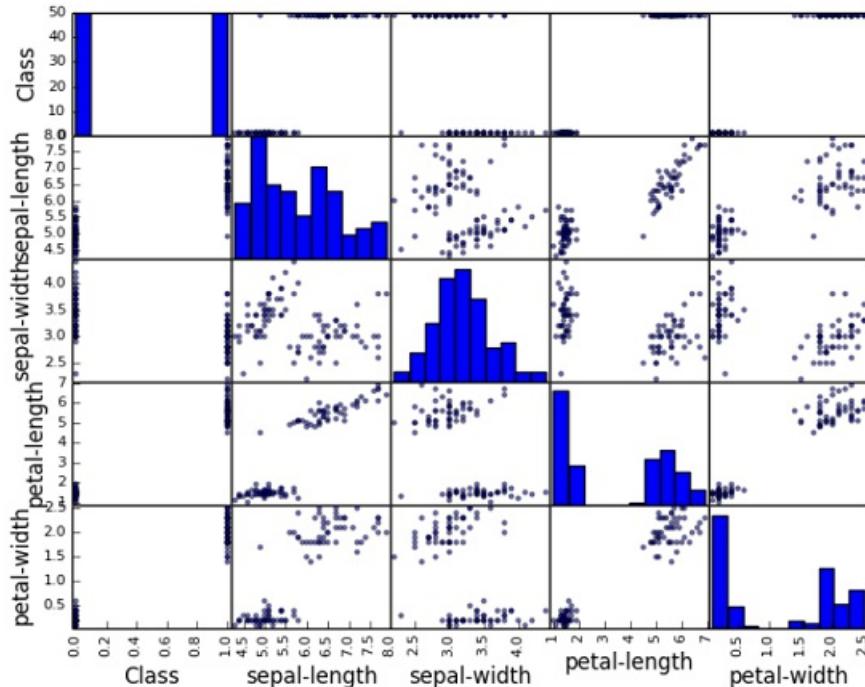


Figure 9. Visualizing plots generated from Python code.

It is possible to return multiple figures by saving them into different images, the Azure Machine Learning runtime picks up all images and concatenates them for visualization.

Advanced examples

The Anaconda environment installed in Azure Machine Learning contains common packages such as NumPy, SciPy, and Scikits-Learn. These packages can be effectively used for various data processing tasks in a machine learning pipeline. As an example, the following experiment and script illustrate the use of ensemble learners in Scikits-Learn to compute feature importance scores for a dataset. The scores can be used to perform supervised feature selection before being fed into another ML model.

Here is the Python function used to compute the importance scores and order the features based on the scores:

```
from sklearn.ensemble import GradientBoostingClassifier
import numpy as np, pandas as pd

def azureml_main(dataframe1):
    colnames = dataframe1.columns
    y = np.array(dataframe1[colnames[-1]])
    X = np.array(dataframe1.ix[:, :len(colnames)-1])
    clf = GradientBoostingClassifier(n_estimators=100, \
        learning_rate=1.0, max_depth=1, random_state=0).\
        fit(X, y)
    fint = clf.feature_importances_
    fnames = np.array(colnames[:-1])

    perm = fint.argsort()

    ret = pd.DataFrame()
    ret["Feature"] = fnames[perm[::-1]]
    ret["Score"] = fint[perm[::-1]]
    return ret,
```

Figure 10. Function to rank features by scores. The following experiment then computes and returns the importance scores of features in the "Pima Indian Diabetes" dataset in Azure Machine Learning:

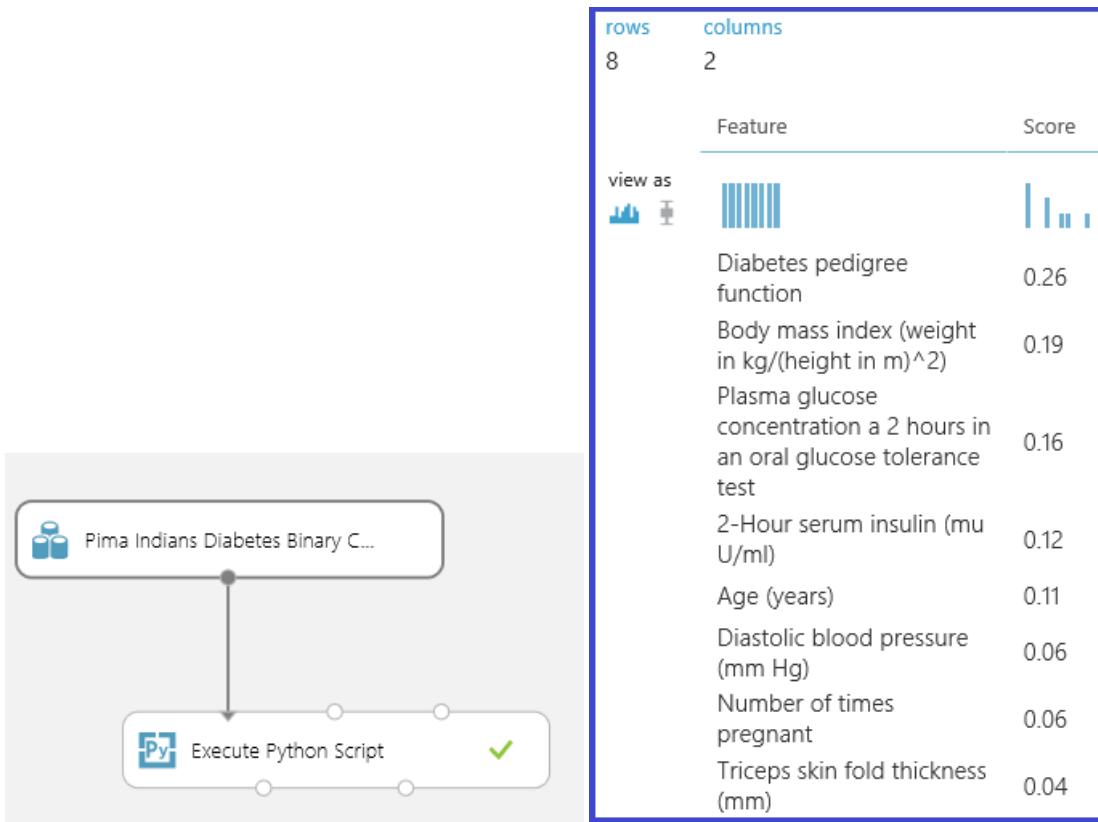


Figure 11. Experiment to rank features in the Pima Indian Diabetes dataset.

Limitations

The [Execute Python Script](#) currently has the following limitations:

- 1. Sandboxed execution.* The Python runtime is currently sandboxed and, as a result, does not allow access to the network or to the local file system in a persistent manner. All files saved locally are isolated and deleted once the module finishes. The Python code cannot access most directories on the machine it runs on, the exception being the current directory and its subdirectories.
- 2. Lack of sophisticated development and debugging support.* The Python module currently does not support IDE features such as intellisense and debugging. Also, if the module fails at runtime, the full Python stack trace is available. But it must be viewed in the output log for the module. We currently recommend that you develop and debug Python scripts in an environment such as IPython and then import the code into the module.
- 3. Single data frame output.* The Python entry point is only permitted to return a single data frame as output. It is not currently possible to return arbitrary Python objects such as trained models directly back to the Azure Machine Learning runtime. Like [Execute R Script](#), which has the same limitation, it is possible in many cases to pickle objects into a byte array and then return that inside of a data frame.
- 4. Inability to customize Python installation.* Currently, the only way to add custom Python modules is via the zip file mechanism described earlier. While this is feasible for small modules, it is cumbersome for large modules (especially those with native DLLs) or a large number of modules.

Conclusions

The [Execute Python Script](#) module allows a data scientist to incorporate existing Python code into cloud-hosted machine learning workflows in Azure Machine Learning and to seamlessly operationalize them as part of a web service. The Python script module interoperates naturally with other modules in Azure Machine Learning. The module can be used for a range of tasks from data exploration to pre-processing and feature extraction, and then to evaluation and post-processing of the results. The backend runtime used for execution is based on Anaconda, a well-tested and widely used Python distribution. This backend makes it simple for you to on-board existing code assets into the cloud.

We expect to provide additional functionality to the [Execute Python Script](#) module such as the ability to train and operationalize models in Python and to add better support for the development and debugging code in Azure Machine Learning Studio.

Next steps

For more information, see the [Python Developer Center](#).

Azure Machine Learning Web Services: Deployment and consumption

12/8/2017 • 3 min to read • [Edit Online](#)

You can use Azure Machine Learning to deploy machine-learning workflows and models as web services. These web services can then be used to call the machine-learning models from applications over the Internet to do predictions in real time or in batch mode. Because the web services are RESTful, you can call them from various programming languages and platforms, such as .NET and Java, and from applications, such as Excel.

The next sections provide links to walkthroughs, code, and documentation to help get you started.

Deploy a web service

With Azure Machine Learning Studio

Machine Learning Studio and the Microsoft Azure Machine Learning Web Services portal help you deploy and manage a web service without writing code.

The following links provide general information about how to deploy a new web service:

- For an overview about how to deploy a new web service that's based on Azure Resource Manager, see [Deploy a new web service](#).
- For a walkthrough about how to deploy a web service, see [Deploy an Azure Machine Learning web service](#).
- For a full walkthrough about how to create and deploy a web service, see [Walkthrough Step 1: Create a Machine Learning workspace](#).
- For specific examples that deploy a web service, see:
 - [Walkthrough Step 5: Deploy the Azure Machine Learning web service](#)
 - [How to deploy a web service to multiple regions](#)

With web services resource provider APIs (Azure Resource Manager APIs)

The Azure Machine Learning resource provider for web services enables deployment and management of web services by using REST API calls. For additional details, see the [Machine Learning Web Service \(REST\)](#) reference.

With PowerShell cmdlets

Azure Machine Learning resource provider for web services enables deployment and management of web services by using PowerShell cmdlets.

To use the cmdlets, you must first sign in to your Azure account from within the PowerShell environment by using the [Add-AzureRmAccount](#) cmdlet. If you are unfamiliar with how to call PowerShell commands that are based on Resource Manager, see [Using Azure PowerShell with Azure Resource Manager](#).

To export your predictive experiment, use [this sample code](#). After you create the .exe file from the code, you can type:

```
C:\<folder>\GetWSD <experiment-url> <workspace-auth-token>
```

Running the application creates a web service JSON template. To use the template to deploy a web service, you must add the following information:

- Storage account name and key

You can get the storage account name and key from the [Azure portal](#).

- Commitment plan ID

You can get the plan ID from the [Azure Machine Learning Web Services](#) portal by signing in and clicking a plan name.

Add them to the JSON template as children of the *Properties* node at the same level as the *MachineLearningWorkspace* node.

Here's an example:

```
"StorageAccount": {  
    "name": "YourStorageAccountName",  
    "key": "YourStorageAccountKey"  
},  
"CommitmentPlan": {  
    "id":  
    "subscriptions/YouSubscriptionID/resourceGroups/YourResourceGroupID/providers/Microsoft.MachineLearning/commitmentPlans/YourPlanNa  
me"  
}
```

See the following articles and sample code for additional details:

- [Azure Machine Learning Cmdlets](#) reference on MSDN
- Sample [walkthrough](#) on GitHub

Consume the web services

From the Azure Machine Learning Web Services UI (Testing)

You can test your web service from the Azure Machine Learning Web Services portal. This includes testing the Request-Response service (RRS) and Batch Execution service (BES) interfaces.

- [Deploy a new web service](#)
- [Deploy an Azure Machine Learning web service](#)
- [Walkthrough Step 5: Deploy the Azure Machine Learning web service](#)

From Excel

You can download an Excel template that consumes the web service:

- [Consuming an Azure Machine Learning web service from Excel](#)
- [Excel add-in for Azure Machine Learning Web Services](#)

From a REST-based client

Azure Machine Learning Web Services are RESTful APIs. You can consume these APIs from various platforms, such as .NET, Python, R, Java, etc. The **Consume** page for your web service on the [Microsoft Azure Machine Learning Web Services portal](#) has sample code that can help you get started. For more information, see [How to consume an Azure Machine Learning Web service](#).

Deploy an Azure Machine Learning web service

11/30/2017 • 9 min to read • [Edit Online](#)

Azure Machine Learning enables you to build, test, and deploy predictive analytic solutions.

From a high-level point-of-view, this is done in three steps:

- **Create a training experiment** - Azure Machine Learning Studio is a collaborative visual development environment that you use to train and test a predictive analytics model using training data that you supply.
- **Convert it to a predictive experiment** - Once your model has been trained with existing data and you're ready to use it to score new data, you prepare and streamline your experiment for predictions.
- **Deploy it as a web service** - You can deploy your predictive experiment as a [new](#) or [classic](#) Azure web service. Users can send data to your model and receive your model's predictions.

NOTE

Try Azure Machine Learning for free

No credit card or Azure subscription required. [Get started now.](#)

Create a training experiment

To train a predictive analytics model, you use Azure Machine Learning Studio to create a training experiment where you include various modules to load training data, prepare the data as necessary, apply machine learning algorithms, and evaluate the results. You can iterate on an experiment and try different machine learning algorithms to compare and evaluate the results.

The process of creating and managing training experiments is covered more thoroughly elsewhere. For more information, see these articles:

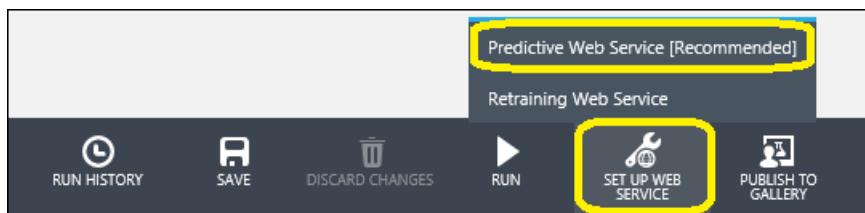
- [Create a simple experiment in Azure Machine Learning Studio](#)
- [Develop a predictive solution with Azure Machine Learning](#)
- [Import your training data into Azure Machine Learning Studio](#)
- [Manage experiment iterations in Azure Machine Learning Studio](#)

Convert the training experiment to a predictive experiment

Once you've trained your model, you're ready to convert your training experiment into a predictive experiment to score new data.

By converting to a predictive experiment, you're getting your trained model ready to be deployed as a scoring web service. Users of the web service can send input data to your model and your model will send back the prediction results. As you convert to a predictive experiment, keep in mind how you expect your model to be used by others.

To convert your training experiment to a predictive experiment, click **Run** at the bottom of the experiment canvas, click **Set Up Web Service**, then select **Predictive Web Service**.



For more information on how to perform this conversion, see [How to prepare your model for deployment in Azure Machine Learning Studio](#).

The following steps describe deploying a predictive experiment as a New web service. You can also deploy the experiment as Classic web service.

Deploy it as a web service

You can deploy the predictive experiment as a New web service or as a Classic web service.

Deploy the predictive experiment as a New web service

Now that the predictive experiment has been prepared, you can deploy it as a new Azure web service. Using the web service, users can send data to your model and the model will return its predictions.

To deploy your predictive experiment, click **Run** at the bottom of the experiment canvas. Once the experiment has finished running, click **Deploy Web Service** and select **Deploy Web Service New**. The deployment page of the Machine Learning Web Service portal opens.

NOTE

To deploy a New web service you must have sufficient permissions in the subscription to which you are deploying the web service. For more information see, [Manage a Web service using the Azure Machine Learning Web Services portal](#).

Machine Learning Web Service portal Deploy Experiment Page

On the Deploy Experiment page, enter a name for the web service. Select a pricing plan. If you have an existing pricing plan you can select it, otherwise you must create a new price plan for the service.

1. In the **Price Plan** drop down, select an existing plan or select the **Select new plan** option.
2. In **Plan Name**, type a name that will identify the plan on your bill.
3. Select one of the **Monthly Plan Tiers**. The plan tiers default to the plans for your default region and your web service is deployed to that region.

Click **Deploy** and the **Quickstart** page for your web service opens.

The web service Quickstart page gives you access and guidance on the most common tasks you will perform after creating a web service. From here, you can easily access both the Test page and Consume page.

Test your New web service

To test your new web service, click **Test web service** under common tasks. On the Test page, you can test your web service as a Request-Response Service (RRS) or a Batch Execution service (BES).

The RRS test page displays the inputs, outputs, and any global parameters that you have defined for the experiment. To test the web service, you can manually enter appropriate values for the inputs or supply a comma separated value (CSV) formatted file containing the test values.

To test using RRS, from the list view mode, enter appropriate values for the inputs and click **Test Request-Response**. Your prediction results display in the output column to the left.

Request-Response Batch

List View **CSV File Input**

input1

output1

Your prediction results will display here.

Class	1
sepal-length	1
sepal-width	1
petal-length	1
petal-width	1

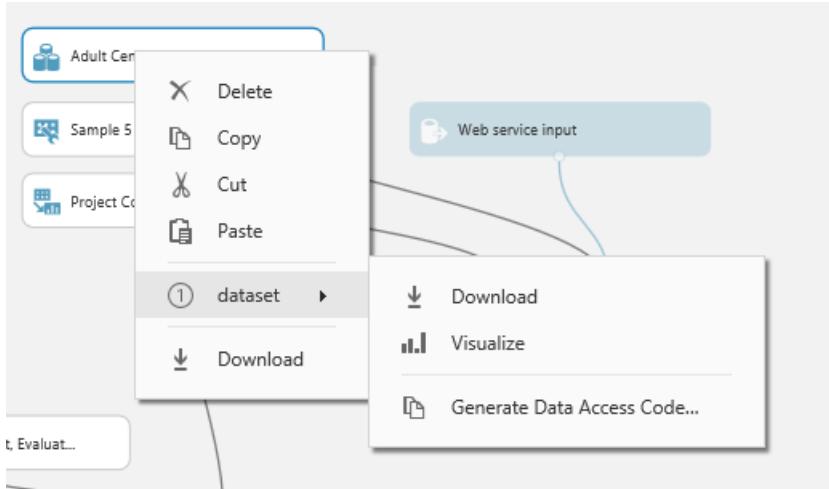
Global Parameters

Category	Trigonometric
Trigonometric Function	Sin
Column Set	{"isFilter":true,"rules":[{"ruleType":"ColumnTy
Output mode	ResultOnly

Test Request-Response

To test your BES, click **Batch**. On the Batch test page, click Browse under your input and select a CSV file containing appropriate sample values. If you don't have a CSV file, and you created your predictive experiment using Machine Learning Studio, you can download the data set for your predictive experiment and use it.

To download the data set, open Machine Learning Studio. Open your predictive experiment and right click the input for your experiment. From the context menu, select **dataset** and then select **Download**.



Click **Test**. The status of your Batch Execution job displays to the right under **Test Batch Jobs**.

Request-Response Batch

✓ input1

Browse your machine for local files... Browse...

✓ Global Parameters

Category	Trigonometric
Trigonometric Function	Sin
Column Set	{"isFilter":true,"rules":[{"ruleType":"ColumnTy}]}
Output mode	ResultOnly

Note: We will enable CORS on your storage account to upload this file

Test

Your batch job status will display here.

On the **CONFIGURATION** page, you can change the description, title, update the storage account key, and enable sample data for your web service.

☰ Microsoft Azure Machine Learning Web Services 👤

Quickstart Dashboard Batch Request Log Configure Consume Test Swagger API

← Web Services **ratings**

Description	Recommender: Restaurant ratings [Predictive Exp.]
Title	Restaurant Ratings Sample
Primary Key	
Secondary Key	
Storage Account Name	<input type="checkbox"/> Update Key for mitestaccount
Sample Data Enabled?	<input type="radio"/> Yes <input type="radio"/> No
Cancel Save	

Microsoft
[FAQ](#) [Privacy and Cookies](#) [Terms of Use](#) © Microsoft

Once you've deployed the web service, you can:

- **Access** it through the web service API.
- **Manage** it through Azure Machine Learning web services portal.
- **Update** it if your model changes.

Access your New web service

Once you deploy your web service from Machine Learning Studio, you can send data to the service and receive responses programmatically.

The **Consume** page provides all the information you need to access your web service. For example, the API key is provided to allow authorized access to the service.

For more information about accessing a Machine Learning web service, see [How to consume an Azure Machine](#)

Learning Web service.

Manage your New web service

You can manage your New web services Machine Learning Web Services portal. From the [main portal page](#), click **Web Services**. From the web services page, you can delete or copy a service. To monitor a specific service, click the service and then click **Dashboard**. To monitor batch jobs associated with the web service, click **Batch Request Log**.

Deploy the predictive experiment as a Classic web service

Now that the predictive experiment has been sufficiently prepared, you can deploy it as a Classic Azure web service. Using the web service, users can send data to your model and the model will return its predictions.

To deploy your predictive experiment, click **Run** at the bottom of the experiment canvas and then click **Deploy Web Service**. The web service is set up and you are placed in the web service dashboard.



Test your Classic web service

You can test the web service in either the Machine Learning Web Services portal or Machine Learning Studio.

To test the Request Response web service, click the **Test** button in the web service dashboard. A dialog pops up to ask you for the input data for the service. These are the columns expected by the scoring experiment. Enter a set of data and then click **OK**. The results generated by the web service are displayed at the bottom of the dashboard.

You can click the **Test** preview link to test your service in the Azure Machine Learning Web Services portal as shown previously in the New web service section.

To test the Batch Execution Service, click **Test** preview link . On the Batch test page, click Browse under your input and select a CSV file containing appropriate sample values. If you don't have a CSV file, and you created your predictive experiment using Machine Learning Studio, you can download the data set for your predictive experiment and use it.

Default Endpoint		API HELP PAGE	TEST	APPS	LAST UPDATED	SEARCH
REQUEST/RESPONSE	Test	Test preview		Excel 2013 or later Excel 2010 or earlier	10/12/2016 9:26:08 AM	
BATCH EXECUTION		Test preview		Excel 2013 or later workbook	10/12/2016 9:26:08 AM	

On the **CONFIGURATION** page, you can change the display name of the service and give it a description. The name and description is displayed in the [Azure portal](#) where you manage your web services.

You can provide a description for your input data, output data, and web service parameters by entering a string for each column under **INPUT SCHEMA**, **OUTPUT SCHEMA**, and **Web SERVICE PARAMETER**. These descriptions are used in the sample code documentation provided for the web service.

You can enable logging to diagnose any failures that you're seeing when your web service is accessed. For more information, see [Enable logging for Machine Learning web services](#).

The screenshot shows the Azure Machine Learning Web Services portal interface. On the left is a vertical sidebar with icons for Dashboard, Configuration, Settings, General, and a gear (Configuration). The main area has a title 'income level prediction'. Below it are tabs for 'DASHBOARD' and 'CONFIGURATION'. A 'settings' section is selected, indicated by a blue underline. Under 'GENERAL', there is a 'Display Name' field containing 'Income level prediction' and a description 'Predicts income level based on user data'. An 'ENABLE LOGGING' section has a 'NO' button highlighted in blue. Under 'INPUT SCHEMA', fields for 'age (Numeric)', 'fnlwgt (Numeric)', and 'education (String)' are listed. At the bottom is a dark bar with a 'NEW' button, a 'SAVE' icon, and a 'DISCARD' icon.

You can also configure the endpoints for the web service in the Azure Machine Learning Web Services portal similar to the procedure shown previously in the New web service section. The options are different, you can add or change the service description, enable logging, and enable sample data for testing.

Access your Classic web service

Once you deploy your web service from Machine Learning Studio, you can send data to the service and receive responses programmatically.

The dashboard provides all the information you need to access your web service. For example, the API key is provided to allow authorized access to the service, and API help pages are provided to help you get started writing your code.

For more information about accessing a Machine Learning web service, see [How to consume an Azure Machine Learning Web service](#).

Manage your Classic web service

There are various of actions you can perform to monitor a web service. You can update it, and delete it. You can also add additional endpoints to a Classic web service in addition to the default endpoint that is created when you deploy it.

For more information, see [Manage an Azure Machine Learning workspace](#) and [Manage a web service using the Azure Machine Learning Web Services portal](#).

Update the web service

You can make changes to your web service, such as updating the model with additional training data, and deploy it again, overwriting the original web service.

To update the web service, open the original predictive experiment you used to deploy the web service and make an editable copy by clicking **SAVE AS**. Make your changes and then click **Deploy Web Service**.

Because you've deployed this experiment before, you are asked if you want to overwrite (Classic Web Service)

or update (New web service) the existing service. Clicking **YES** or **Update** stops the existing web service and deploys the new predictive experiment in its place.

NOTE

If you made configuration changes in the original web service, for example, entering a new display name or description, you will need to enter those values again.

One option for updating your web service is to retrain the model programmatically. For more information, see [Retrain Machine Learning models programmatically](#).

How a Machine Learning model progresses from an experiment to an operationalized Web service

9/25/2017 • 7 min to read • [Edit Online](#)

Azure Machine Learning Studio provides an interactive canvas that allows you to develop, run, test, and iterate an **experiment** representing a predictive analysis model. There are a wide variety of modules available that can:

- Input data into your experiment
- Manipulate the data
- Train a model using machine learning algorithms
- Score the model
- Evaluate the results
- Output final values

Once you're satisfied with your experiment, you can deploy it as a **Classic Azure Machine Learning Web service** or a **New Azure Machine Learning Web service** so that users can send it new data and receive back results.

In this article, we give an overview of the mechanics of how your Machine Learning model progresses from a development experiment to an operationalized Web service.

NOTE

There are other ways to develop and deploy machine learning models, but this article is focused on how you use Machine Learning Studio. For example, to read a description of how to create a classic predictive Web service with R, see the blog post [Build & Deploy Predictive Web Apps Using RStudio and Azure ML](#).

While Azure Machine Learning Studio is designed to help you develop and deploy a *predictive analysis model*, it's possible to use Studio to develop an experiment that doesn't include a predictive analysis model. For example, an experiment might just input data, manipulate it, and then output the results. Just like a predictive analysis experiment, you can deploy this non-predictive experiment as a Web service, but it's a simpler process because the experiment isn't training or scoring a machine learning model. While it's not the typical to use Studio in this way, we'll include it in the discussion so that we can give a complete explanation of how Studio works.

Developing and deploying a predictive Web service

Here are the stages that a typical solution follows as you develop and deploy it using Machine Learning Studio:

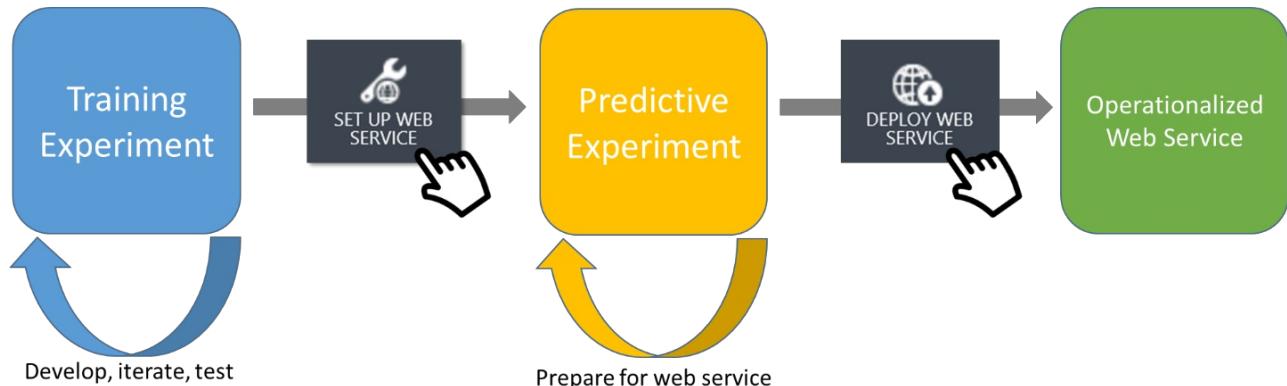


Figure 1 - Stages of a typical predictive analysis model

The training experiment

The **training experiment** is the initial phase of developing your Web service in Machine Learning Studio. The purpose of the training experiment is to give you a place to develop, test, iterate, and eventually train a machine learning model. You can even train multiple models simultaneously as you look for the best solution, but once you're done experimenting you'll select a single trained model and eliminate the rest from the experiment. For an example of developing a predictive analysis experiment, see [Develop a predictive analytics solution for credit risk assessment in Azure Machine Learning](#).

The predictive experiment

Once you have a trained model in your training experiment, click **Set Up Web Service** and select **Predictive Web Service** in Machine Learning Studio to initiate the process of converting your training experiment to a **predictive experiment**. The purpose of the predictive experiment is to use your trained model to score new data, with the goal of eventually becoming operationalized as an Azure Web service.

This conversion is done for you through the following steps:

- Convert the set of modules used for training into a single module and save it as a trained model
- Eliminate any extraneous modules not related to scoring
- Add input and output ports that the eventual Web service will use

There may be more changes you want to make to get your predictive experiment ready to deploy as a Web service. For example, if you want the Web service to output only a subset of results, you can add a filtering module before the output port.

In this conversion process, the training experiment is not discarded. When the process is complete, you have two tabs in Studio: one for the training experiment and one for the predictive experiment. This way you can make changes to the training experiment before you deploy your Web service and rebuild the predictive experiment. Or you can save a copy of the training experiment to start another line of experimentation.

NOTE

When you click **Predictive Web Service** you start an automatic process to convert your training experiment to a predictive experiment, and this works well in most cases. If your training experiment is complex (for example, you have multiple paths for training that you join together), you might prefer to do this conversion manually. For more information, see [How to prepare your model for deployment in Azure Machine Learning Studio](#).

The Web service

Once you're satisfied that your predictive experiment is ready, you can deploy your service as either a Classic Web service or a New Web service based on Azure Resource Manager. To operationalize your model by deploying it as a *Classic Machine Learning Web service*, click **Deploy Web Service** and select **Deploy Web Service [Classic]**. To deploy as *New Machine Learning Web service*, click **Deploy Web Service** and select **Deploy Web Service [New]**. Users can now send data to your model using the Web service REST API and receive back the results. For more information, see [How to consume an Azure Machine Learning Web service](#).

The non-typical case: creating a non-predictive Web service

If your experiment does not train a predictive analysis model, then you don't need to create both a training experiment and a scoring experiment - there's just one experiment, and you can deploy it as a Web service. Machine Learning Studio detects whether your experiment contains a predictive model by analyzing the modules you've used.

After you've iterated on your experiment and are satisfied with it:

1. Click **Set Up Web Service** and select **Retraining Web Service** - input and output nodes are added automatically

2. Click **Run**
3. Click **Deploy Web Service** and select **Deploy Web Service [Classic]** or **Deploy Web Service [New]** depending on the environment to which you want to deploy.

Your Web service is now deployed, and you can access and manage it just like a predictive Web service.

Updating your Web service

Now that you've deployed your experiment as a Web service, what if you need to update it?

That depends on what you need to update:

You want to change the input or output, or you want to modify how the Web service manipulates data

If you're not changing the model, but are just changing how the Web service handles data, you can edit the predictive experiment and then click **Deploy Web Service** and select **Deploy Web Service [Classic]** or **Deploy Web Service [New]** again. The Web service is stopped, the updated predictive experiment is deployed, and the Web service is restarted.

Here's an example: Suppose your predictive experiment returns the entire row of input data with the predicted result. You may decide that you want the Web service to just return the result. So you can add a **Project Columns** module in the predictive experiment, right before the output port, to exclude columns other than the result. When you click **Deploy Web Service** and select **Deploy Web Service [Classic]** or **Deploy Web Service [New]** again, the Web service is updated.

You want to retrain the model with new data

If you want to keep your machine learning model, but you would like to retrain it with new data, you have two choices:

1. **Retrain the model while the Web service is running** - If you want to retrain your model while the predictive Web service is running, you can do this by making a couple modifications to the training experiment to make it a **retraining experiment**, then you can deploy it as a **retraining web service**. For instructions on how to do this, see [Retrain Machine Learning models programmatically](#).
2. **Go back to the original training experiment and use different training data to develop your model**
 - Your predictive experiment is linked to the Web service, but the training experiment is not directly linked in this way. If you modify the original training experiment and click **Set Up Web Service**, it will create a *new* predictive experiment which, when deployed, will create a *new* Web service. It doesn't just update the original Web service.

If you need to modify the training experiment, open it and click **Save As** to make a copy. This will leave intact the original training experiment, predictive experiment, and Web service. You can now create a new Web service with your changes. Once you've deployed the new Web service you can then decide whether to stop the previous Web service or keep it running alongside the new one.

You want to train a different model

If you want to make changes to your original predictive experiment, such as selecting a different machine learning algorithm, trying a different training method, etc., then you need to follow the second procedure described above for retraining your model: open the training experiment, click **Save As** to make a copy, and then start down the new path of developing your model, creating the predictive experiment, and deploying the web service. This will create a new Web service unrelated to the original one - you can decide which one, or both, to keep running.

Next Steps

For more details on the process of developing and experiment, see the following articles:

- converting the experiment - [How to prepare your model for deployment in Azure Machine Learning Studio](#)
- deploying the Web service - [Deploy an Azure Machine Learning web service](#)
- retraining the model - [Retrain Machine Learning models programmatically](#)

For examples of the whole process, see:

- [Machine learning tutorial: Create your first experiment in Azure Machine Learning Studio](#)
- [Walkthrough: Develop a predictive analytics solution for credit risk assessment in Azure Machine Learning](#)

How to prepare your model for deployment in Azure Machine Learning Studio

9/25/2017 • 7 min to read • [Edit Online](#)

Azure Machine Learning Studio gives you the tools you need to develop a predictive analytics model and then operationalize it by deploying it as an Azure web service.

To do this, you use Studio to create an experiment - called a *training experiment* - where you train, score, and edit your model. Once you're satisfied, you get your model ready to deploy by converting your training experiment to a *predictive experiment* that's configured to score user data.

You can see an example of this process in [Walkthrough: Develop a predictive analytics solution for credit risk assessment in Azure Machine Learning](#).

This article takes a deep dive into the details of how a training experiment gets converted into a predictive experiment, and how that predictive experiment is deployed. By understanding these details, you can learn how to configure your deployed model to make it more effective.

NOTE

Try Azure Machine Learning for free

No credit card or Azure subscription required. [Get started now.](#)

Overview

The process of converting a training experiment to a predictive experiment involves three steps:

1. Replace the machine learning algorithm modules with your trained model.
2. Trim the experiment to only those modules that are needed for scoring. A training experiment includes a number of modules that are necessary for training but are not needed once the model is trained.
3. Define how your model will accept data from the web service user, and what data will be returned.

TIP

In your training experiment, you've been concerned with training and scoring your model using your own data. But once deployed, users will send new data to your model and it will return prediction results. So, as you convert your training experiment to a predictive experiment to get it ready for deployment, keep in mind how the model will be used by others.

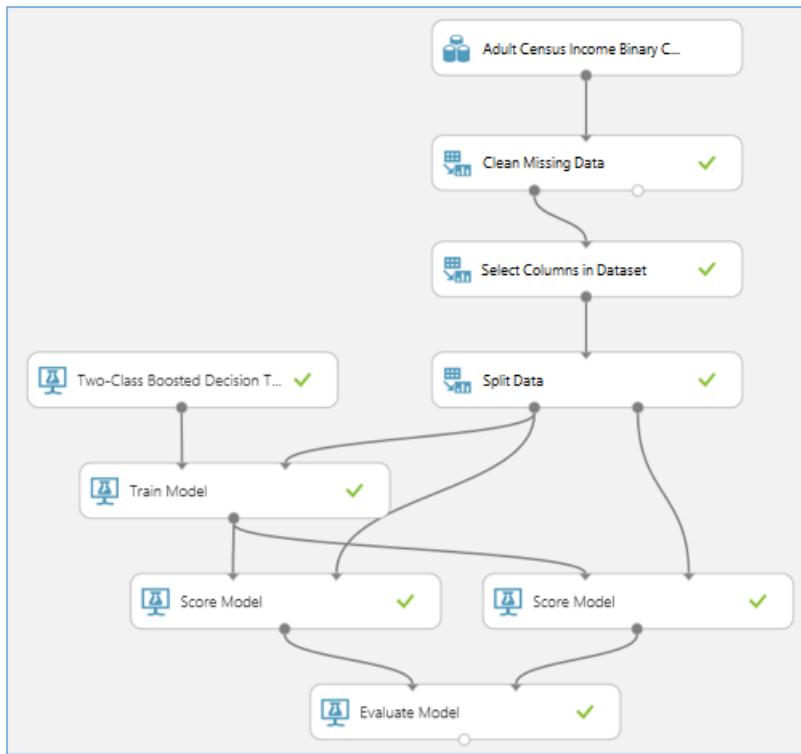
Set Up Web Service button

After you run your experiment (click **RUN** at the bottom of the experiment canvas), click the **Set Up Web Service** button (select the **Predictive Web Service** option). **Set Up Web Service** performs for you the three steps of converting your training experiment to a predictive experiment:

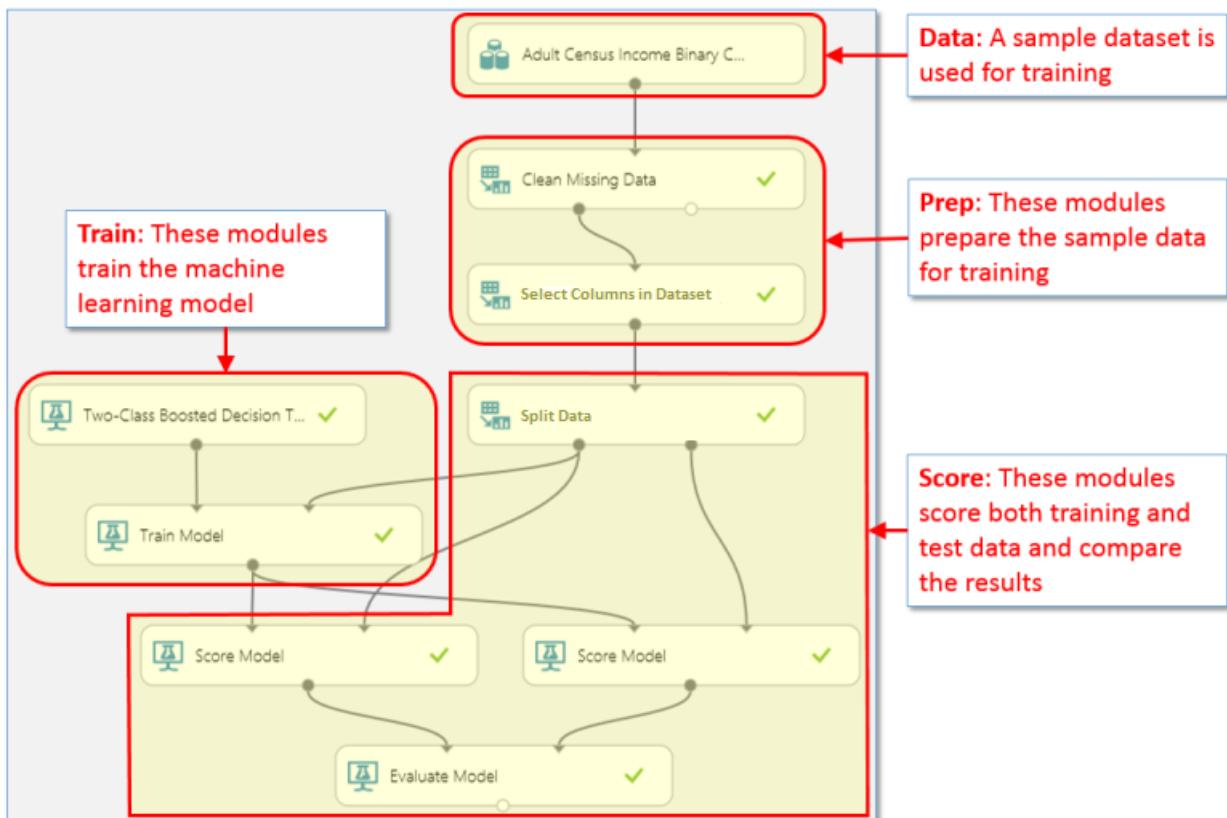
1. It saves your trained model in the **Trained Models** section of the module palette (to the left of the experiment canvas). It then replaces the machine learning algorithm and **Train Model** modules with the saved trained model.
2. It analyzes your experiment and removes modules that were clearly used only for training and are no longer needed.

3. It inserts *Web service input* and *output* modules into default locations in your experiment (these modules accept and return user data).

For example, the following experiment trains a two-class boosted decision tree model using sample census data:



The modules in this experiment perform basically four different functions:



When you convert this training experiment to a predictive experiment, some of these modules are no longer needed, or they now serve a different purpose:

- **Data** - The data in this sample dataset is not used during scoring - the user of the web service will supply the data to be scored. However, the metadata from this dataset, such as data types, is used by the trained

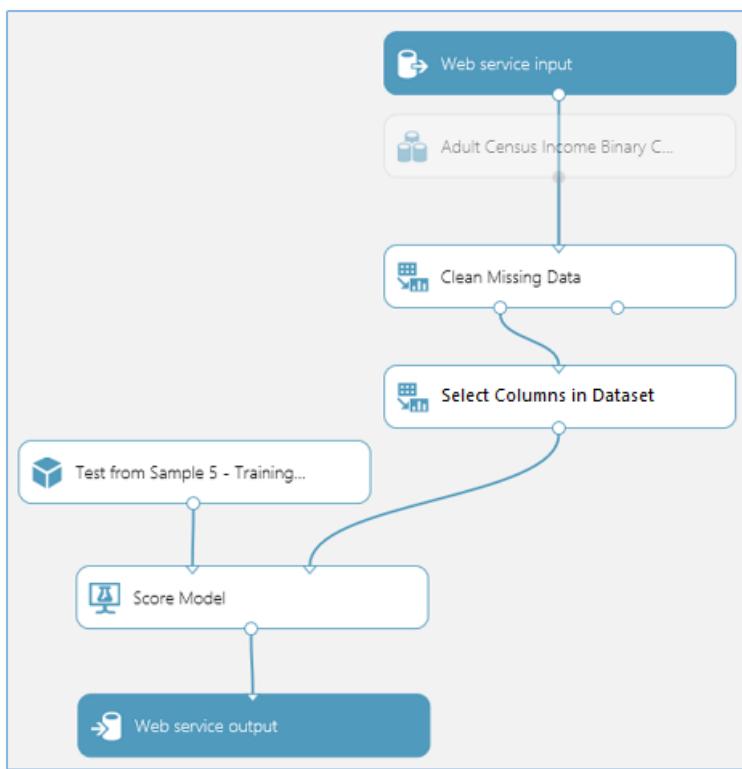
model. So you need to keep the dataset in the predictive experiment so that it can provide this metadata.

- **Prep** - Depending on the user data that will be submitted for scoring, these modules may or may not be necessary to process the incoming data. The **Set Up Web Service** button doesn't touch these - you need to decide how you want to handle them.

For instance, in this example the sample dataset may have missing values, so a [Clean Missing Data](#) module was included to deal with them. Also, the sample dataset includes columns that are not needed to train the model. So a [Select Columns in Dataset](#) module was included to exclude those extra columns from the data flow. If you know that the data that will be submitted for scoring through the web service will not have missing values, then you can remove the [Clean Missing Data](#) module. However, since the [Select Columns in Dataset](#) module helps define the columns of data that the trained model expects, that module needs to remain.

- **Train** - These modules are used to train the model. When you click **Set Up Web Service**, these modules are replaced with a single module that contains the model you trained. This new module is saved in the **Trained Models** section of the module palette.
- **Score** - In this example, the [Split Data](#) module is used to divide the data stream into test data and training data. In the predictive experiment, we're not training anymore, so [Split Data](#) can be removed. Similarly, the second [Score Model](#) module and the [Evaluate Model](#) module are used to compare results from the test data, so these modules are not needed in the predictive experiment. The remaining [Score Model](#) module, however, is needed to return a score result through the web service.

Here is how our example looks after clicking **Set Up Web Service**:

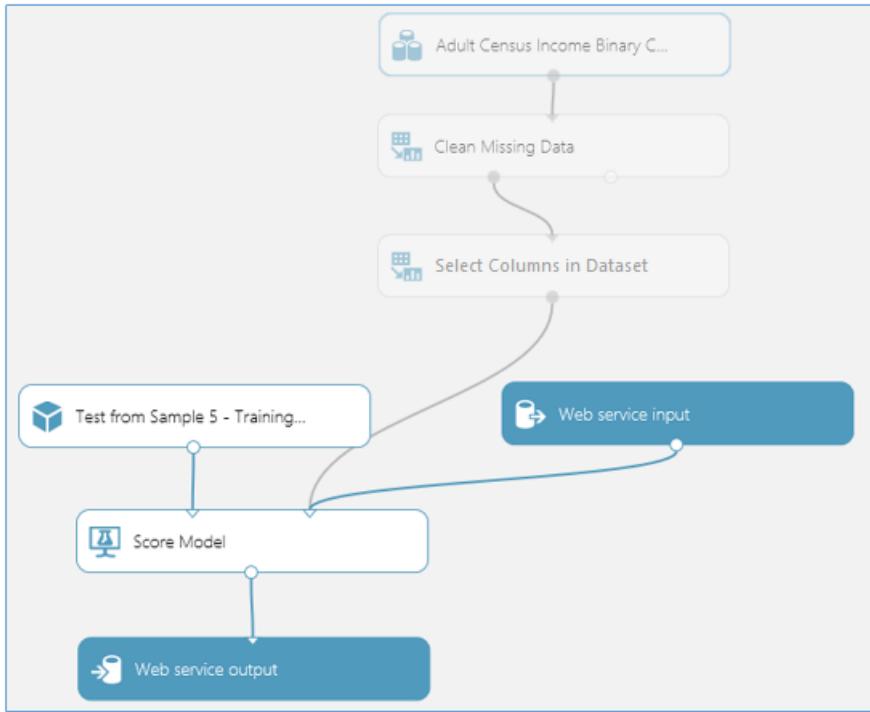


The work done by **Set Up Web Service** may be sufficient to prepare your experiment to be deployed as a web service. However, you may want to do some additional work specific to your experiment.

Adjust input and output modules

In your training experiment, you used a set of training data and then did some processing to get the data in a form that the machine learning algorithm needed. If the data you expect to receive through the web service will not need this processing, you can bypass it: connect the output of the **Web service input module** to a different module in your experiment. The user's data will now arrive in the model at this location.

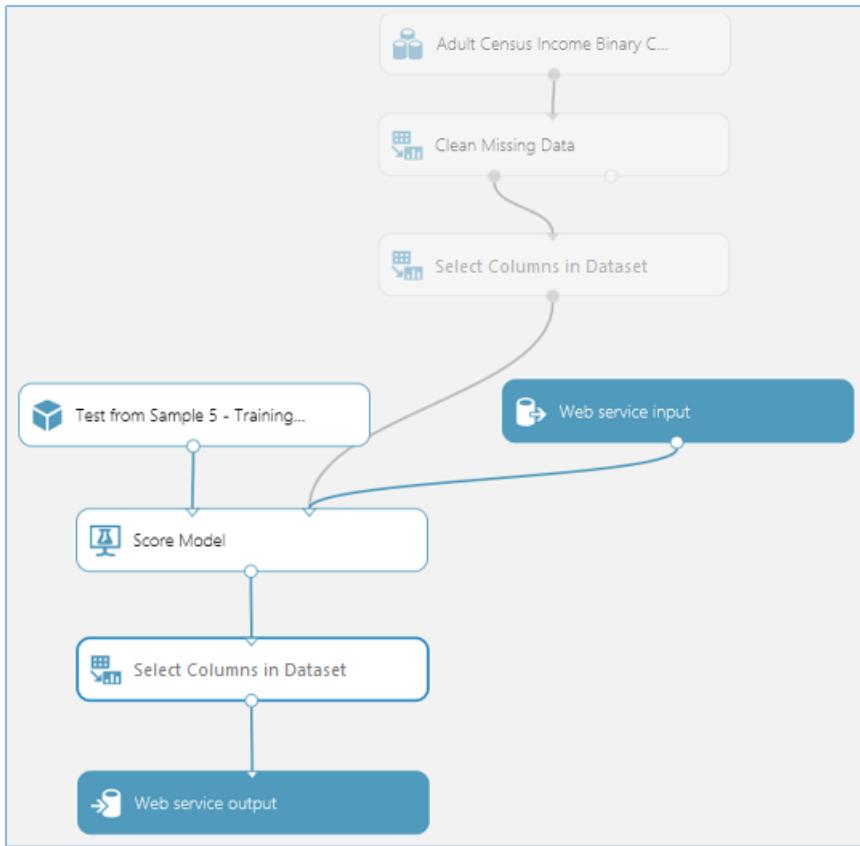
For example, by default **Set Up Web Service** puts the **Web service input** module at the top of your data flow, as shown in the figure above. But we can manually position the **Web service input** past the data processing modules:



The input data provided through the web service will now pass directly into the Score Model module without any preprocessing.

Similarly, by default **Set Up Web Service** puts the Web service output module at the bottom of your data flow. In this example, the web service will return to the user the output of the **Score Model** module, which includes the complete input data vector plus the scoring results. However, if you would prefer to return something different, then you can add additional modules before the **Web service output** module.

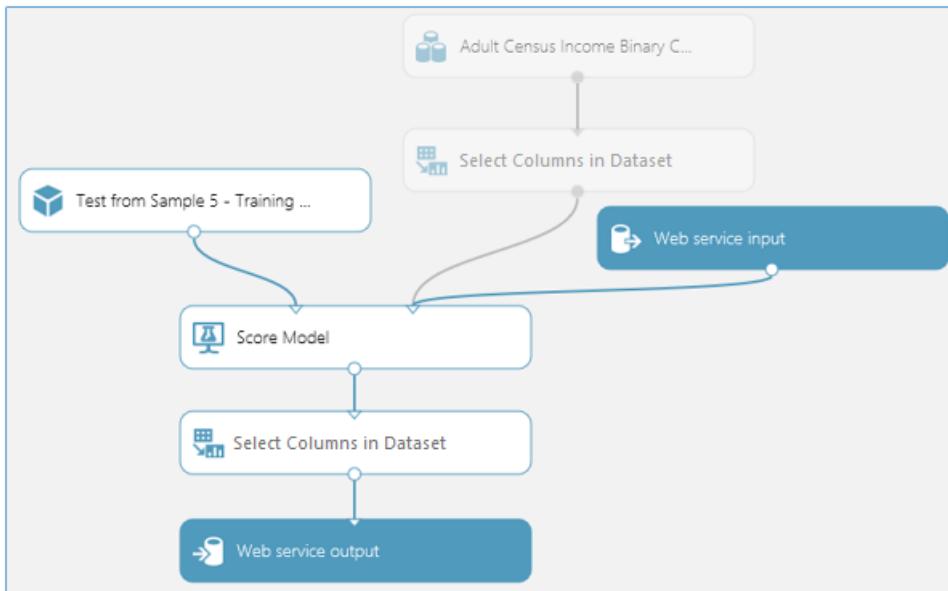
For example, to return only the scoring results and not the entire vector of input data, add a [Select Columns in Dataset](#) module to exclude all columns except the scoring results. Then move the **Web service output** module to the output of the [Select Columns in Dataset](#) module. The experiment looks like this:



Add or remove additional data processing modules

If there are more modules in your experiment that you know will not be needed during scoring, these can be removed. For example, because we moved the **Web service input** module to a point after the data processing modules, we can remove the **Clean Missing Data** module from the predictive experiment.

Our predictive experiment now looks like this:



Add optional Web Service Parameters

In some cases, you may want to allow the user of your web service to change the behavior of modules when the service is accessed. *Web Service Parameters* allow you to do this.

A common example is setting up an **Import Data** module so the user of the deployed web service can specify a different data source when the web service is accessed. Or configuring an **Export Data** module so that a different destination can be specified.

You can define Web Service Parameters and associate them with one or more module parameters, and you can specify whether they are required or optional. The user of the web service provides values for these parameters when the service is accessed, and the module actions are modified accordingly.

For more information about what Web Service Parameters are and how to use them, see [Using Azure Machine Learning Web Service Parameters](#).

Deploy the predictive experiment as a web service

Now that the predictive experiment has been sufficiently prepared, you can deploy it as an Azure web service. Using the web service, users can send data to your model and the model will return its predictions.

For more information on the complete deployment process, see [Deploy an Azure Machine Learning web service](#)

Deploying Azure ML web services that use Data Import and Data Export modules

12/8/2017 • 6 min to read • [Edit Online](#)

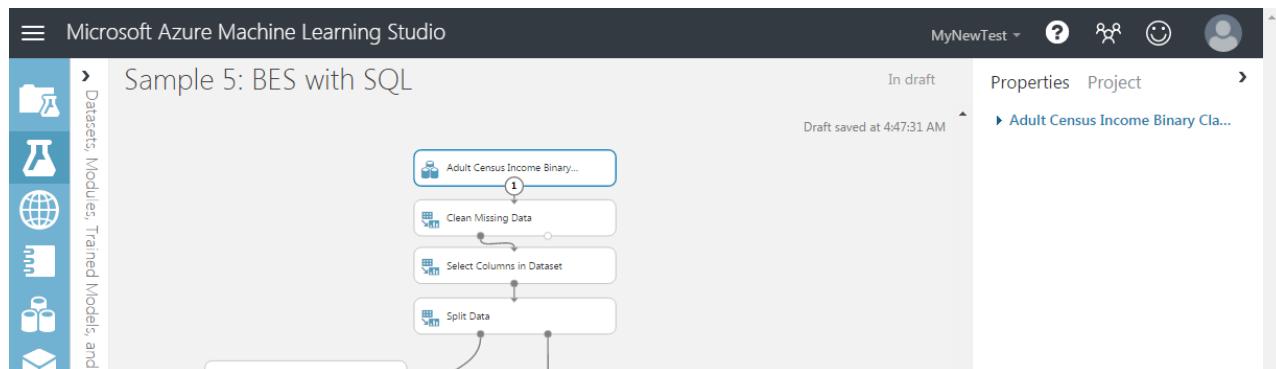
When you create a predictive experiment, you typically add a web service input and output. When you deploy the experiment, consumers can send and receive data from the web service through the inputs and outputs. For some applications, a consumer's data may be available from a data feed or already reside in an external data source such as Azure Blob storage. In these cases, they do not need read and write data using web service inputs and outputs. They can, instead, use the Batch Execution Service (BES) to read data from the data source using an Import Data module and write the scoring results to a different data location using an Export Data module.

The Import Data and Export data modules, can read from and write to various data locations such as a Web URL via HTTP, a Hive Query, an Azure SQL database, Azure Table storage, Azure Blob storage, a Data Feed provide, or an on-premises SQL database.

This topic uses the "Sample 5: Train, Test, Evaluate for Binary Classification: Adult Dataset" sample and assumes the dataset has already been loaded into an Azure SQL table named `censusdata`.

Create the training experiment

When you open the "Sample 5: Train, Test, Evaluate for Binary Classification: Adult Dataset" sample it uses the sample Adult Census Income Binary Classification dataset. And the experiment in the canvas will look similar to the following image:



To read the data from the Azure SQL table:

1. Delete the dataset module.
2. In the components search box, type import.
3. From the results list, add an *Import Data* module to the experiment canvas.
4. Connect output of the *Import Data* module the input of the *Clean Missing Data* module.
5. In properties pane, select **Azure SQL Database** in the **Data Source** dropdown.
6. In the **Database server name**, **Database name**, **User name**, and **Password** fields, enter the appropriate information for your database.
7. In the Database query field, enter the following query.

```
select [age],
```

```
[workclass],  
[fnlwgt],  
[education],  
[education-num],  
[marital-status],  
[occupation],  
[relationship],  
[race],  
[sex],  
[capital-gain],  
[capital-loss],  
[hours-per-week],  
[native-country],  
[income]
```

from dbo.censusdata;

8. At the bottom of the experiment canvas, click **Run**.

Create the predictive experiment

Next you set up the predictive experiment from which you deploy your web service.

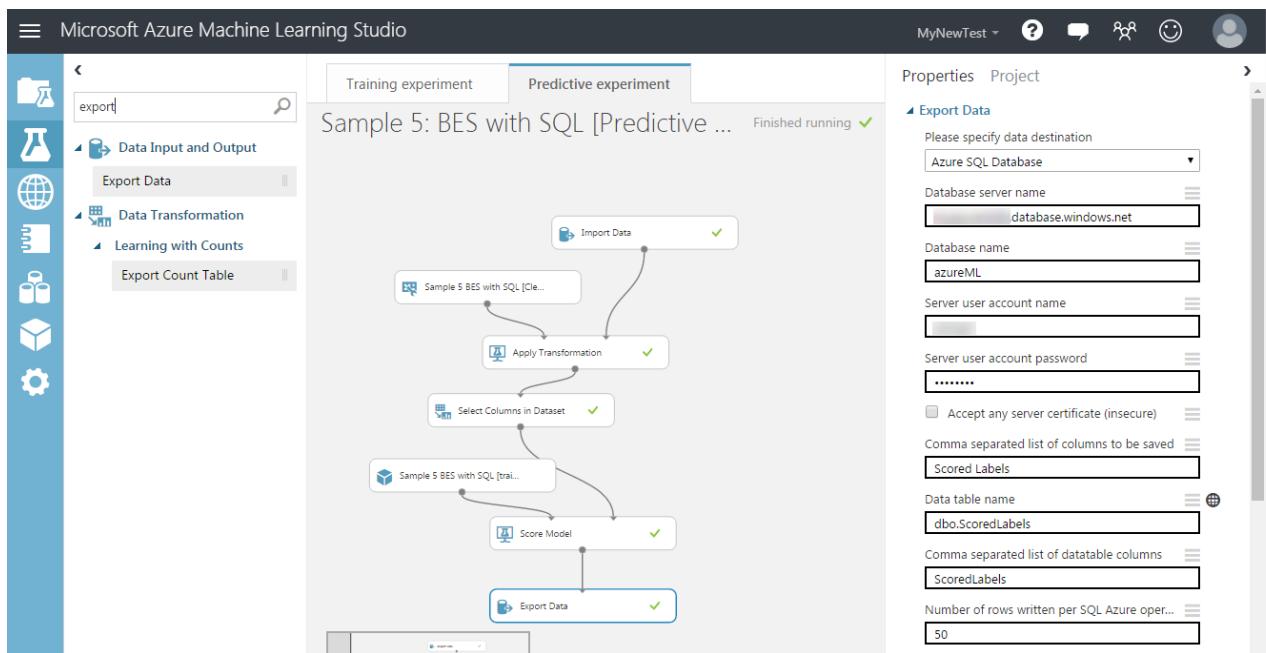
1. At the bottom of the experiment canvas, click **Set Up Web Service** and select **Predictive Web Service [Recommended]**.
2. Remove the *Web Service Input* and *Web Service Output* modules from the predictive experiment.
3. In the components search box, type export.
4. From the results list, add an *Export Data* module to the experiment canvas.
5. Connect output of the *Score Model* module the input of the *Export Data* module.
6. In properties pane, select **Azure SQL Database** in the data destination dropdown.
7. In the **Database server name**, **Database name**, **Server user account name**, and **Server user account password** fields, enter the appropriate information for your database.
8. In the **Comma separated list of columns to be saved** field, type Scored Labels.
9. In the **Data table name field**, type dbo.ScoredLabels. If the table does not exist, it is created when the experiment is run or the web service is called.
10. In the **Comma separated list of datatable columns** field, type ScoredLabels.

When you write an application that calls the final web service, you may want to specify a different input query or destination table at run time. To configure these inputs and outputs, use the Web Service Parameters feature to set the *Import Data* module *Data source* property and the *Export Data* mode data destination property. For more information on Web Service Parameters, see the [AzureML Web Service Parameters entry](#) on the Cortana Intelligence and Machine Learning Blog.

To configure the Web Service Parameters for the import query and the destination table:

1. In the properties pane for the *Import Data* module, click the icon at the top right of the **Database query** field and select **Set as web service parameter**.
2. In the properties pane for the *Export Data* module, click the icon at the top right of the **Data table name** field and select **Set as web service parameter**.
3. At the bottom of the *Export Data* module properties pane, in the **Web Service Parameters** section, click Database query and rename it Query.
4. Click **Data table name** and rename it Table.

When you are done, your experiment should look similar to the following image:



Now you can deploy the experiment as a web service.

Deploy the web service

You can deploy to either a Classic or New web service.

Deploy a Classic Web Service

To deploy as a Classic Web Service and create an application to consume it:

1. At the bottom of the experiment canvas, click Run.
2. When the run has completed, click **Deploy Web Service** and select **Deploy Web Service [Classic]**.
3. On the web service dashboard, locate your API key. Copy and save it to use later.
4. In the **Default Endpoint** table, click the **Batch Execution** link to open the API Help Page.
5. In Visual Studio, create a C# console application: **New > Project > Visual C# > Windows Classic Desktop > Console App (.NET Framework)**.
6. On the API Help Page, find the **Sample Code** section at the bottom of the page.
7. Copy and paste the C# sample code into your Program.cs file, and remove all references to the blob storage.
8. Update the value of the *apiKey* variable with the API key saved earlier.
9. Locate the request declaration and update the values of Web Service Parameters that are passed to the *Import Data* and *Export Data* modules. In this case, you use the original query, but define a new table name.

```
var request = new BatchExecutionRequest()
{
    GlobalParameters = new Dictionary<string, string>()
    {
        { "Query", @"select [age], [workclass], [fnlwgt], [education], [education-num], [marital-status], [occupation], [relationship], [race], [sex], [capital-gain], [capital-loss], [hours-per-week], [native-country], [income] from dbo.censusdata" },
        { "Table", "dbo.ScoredTable2" },
    }
};
```

10. Run the application.

On completion of the run, a new table is added to the database containing the scoring results.

Deploy a New Web Service

NOTE

To deploy a New web service you must have sufficient permissions in the subscription to which you deploying the web service. For more information, see [Manage a Web service using the Azure Machine Learning Web Services portal](#).

To deploy as a New Web Service and create an application to consume it:

1. At the bottom of the experiment canvas, click **Run**.
2. When the run has completed, click **Deploy Web Service** and select **Deploy Web Service [New]**.
3. On the Deploy Experiment page, enter a name for your web service, and select a pricing plan, then click **Deploy**.
4. On the **Quickstart** page, click **Consume**.
5. In the **Sample Code** section, click **Batch**.
6. In Visual Studio, create a C# console application: **New > Project > Visual C# > Windows Classic Desktop > Console App (.NET Framework)**.
7. Copy and paste the C# sample code into your Program.cs file.
8. Update the value of the *apiKey* variable with the **Primary Key** located in the **Basic consumption info** section.
9. Locate the *scoreRequest* declaration and update the values of Web Service Parameters that are passed to the *Import Data* and *Export Data* modules. In this case, you use the original query, but define a new table name.

```
var scoreRequest = new
{
    Inputs = new Dictionary<string, StringTable>()
    {
        },
    GlobalParameters = new Dictionary<string, string>()
    {
        {"Query", @"select [age], [workclass], [fnlwgt], [education], [education-num], [marital-status], [occupation], [relationship], [race],
        [sex], [capital-gain], [capital-loss], [hours-per-week], [native-country], [income] from dbo.censusdata" },
        {"Table", "dbo.ScoredTable3" },
    }
};
```

10. Run the application.

How to deploy a Web Service to multiple regions

12/8/2017 • 1 min to read • [Edit Online](#)

The New Azure Web Services allow you to easily deploy a web service to multiple regions without needing multiple subscriptions or workspaces.

Pricing is region specific, therefore you must define a billing plan for each region in which you will deploy the web service.

To create a plan in another region

1. Sign into [Microsoft Azure Machine Learning Web Services](#).
2. Click the **Plans** menu option.
3. On the Plans over view page, click **New**.
4. From the **Subscription** dropdown, select the subscription in which the new plan will reside.
5. From the **Region** dropdown, select a region for the new plan. The Plan Options for the selected region will display in the **Plan Options** section of the page.
6. From the **Resource Group** dropdown, select a resource group for the plan. From more information on resource groups, see [Azure Resource Manager overview](#).
7. In **Plan Name** type the name of the plan.
8. Under **Plan Options**, click the billing level for the new plan.
9. Click **Create**.

Deploying the web service to another region

1. Click the **Web Services** menu option.
2. Select the Web Service you are deploying to a new region.
3. Click **Copy**.
4. In **Web Service Name**, type a new name for the web service.
5. In **Web service description**, type a description for the web service.
6. From the **Subscription** dropdown, select the subscription in which the new web service will reside.
7. From the **Resource Group** dropdown, select a resource group for the web service. From more information on resource groups, see [Azure Resource Manager overview](#).
8. From the **Region** dropdown, select the region in which to deploy the web service.
9. From the **Storage account** dropdown, select a storage account in which to store the web service.
10. From the **Price Plan** dropdown, select a plan in the region you selected in step 8.
11. Click **Copy**.

Use Azure Machine Learning Web Service Parameters

9/25/2017 • 4 min to read • [Edit Online](#)

An Azure Machine Learning web service is created by publishing an experiment that contains modules with configurable parameters. In some cases, you may want to change the module behavior while the web service is running. *Web Service Parameters* allow you to do this task.

A common example is setting up the [Import Data](#) module so that the user of the published web service can specify a different data source when the web service is accessed. Or configuring the [Export Data](#) module so that a different destination can be specified. Some other examples include changing the number of bits for the [Feature Hashing](#) module or the number of desired features for the [Filter-Based Feature Selection](#) module.

You can set Web Service Parameters and associate them with one or more module parameters in your experiment, and you can specify whether they are required or optional. The user of the web service can then provide values for these parameters when they call the web service.

NOTE

[Try Azure Machine Learning for free](#)

No credit card or Azure subscription required. [Get started now.](#)

How to set and use Web Service Parameters

You define a Web Service Parameter by clicking the icon next to the parameter for a module and selecting "Set as web service parameter". This creates a new Web Service Parameter and connects it to that module parameter. Then, when the web service is accessed, the user can specify a value for the Web Service Parameter and it is applied to the module parameter.

Once you define a Web Service Parameter, it's available to any other module parameter in the experiment. If you define a Web Service Parameter associated with a parameter for one module, you can use that same Web Service Parameter for any other module, as long as the parameter expects the same type of value. For example, if the Web Service Parameter is a numeric value, then it can only be used for module parameters that expect a numeric value. When the user sets a value for the Web Service Parameter, it will be applied to all associated module parameters.

You can decide whether to provide a default value for the Web Service Parameter. If you do, then the parameter is optional for the user of the web service. If you don't provide a default value, then the user is required to enter a value when the web service is accessed.

The API documentation for the web service includes information for the web service user on how to specify the Web Service Parameter programmatically when accessing the web service.

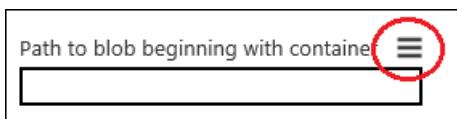
NOTE

The API documentation for a classic web service is provided through the [API help page](#) link in the web service **DASHBOARD** in Machine Learning Studio. The API documentation for a new web service is provided through the [Azure Machine Learning Web Services](#) portal on the [Consume](#) and [Swagger API](#) pages for your web service.

Example

As an example, let's assume we have an experiment with an [Export Data](#) module that sends information to Azure blob storage. We'll define a Web Service Parameter named "Blob path" that allows the web service user to change the path to the blob storage when the service is accessed.

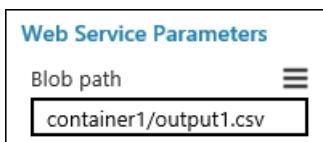
1. In Machine Learning Studio, click the [Export Data](#) module to select it. Its properties are shown in the Properties pane to the right of the experiment canvas.
2. Specify the storage type:
 - Under **Please specify data destination**, select "Azure Blob Storage".
 - Under **Please specify authentication type**, select "Account".
 - Enter the account information for the Azure blob storage.
3. Click the icon to the right of the **Path to blob beginning with container parameter**. It looks like this:



Select "Set as web service parameter".

An entry is added under **Web Service Parameters** at the bottom of the Properties pane with the name "Path to blob beginning with container". This is the Web Service Parameter that is now associated with this [Export Data](#) module parameter.

4. To rename the Web Service Parameter, click the name, enter "Blob path", and press the **Enter** key.
5. To provide a default value for the Web Service Parameter, click the icon to the right of the name, select "Provide default value", enter a value (for example, "container1/output1.csv"), and press the **Enter** key.



6. Click **Run**.
7. Click **Deploy Web Service** and select **Deploy Web Service [Classic]** or **Deploy Web Service [New]** to deploy the web service.

NOTE

To deploy a New web service you must have sufficient permissions in the subscription to which you deploying the web service. For more information see, [Manage a Web service using the Azure Machine Learning Web Services portal](#).

The user of the web service can now specify a new destination for the [Export Data](#) module when accessing the web service.

More information

For a more detailed example, see the [Web Service Parameters](#) entry in the [Machine Learning Blog](#).

For more information on accessing a Machine Learning web service, see [How to consume an Azure Machine Learning Web service](#).

Enable logging for Machine Learning web services

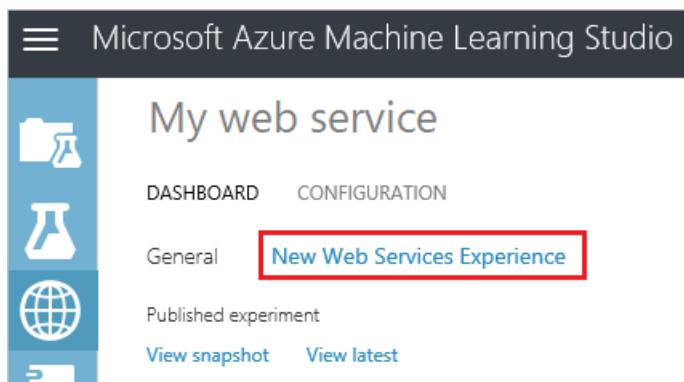
9/25/2017 • 2 min to read • [Edit Online](#)

This document provides information on the logging capability of Machine Learning web services. Logging provides additional information, beyond just an error number and a message, that can help you troubleshoot your calls to the Machine Learning APIs.

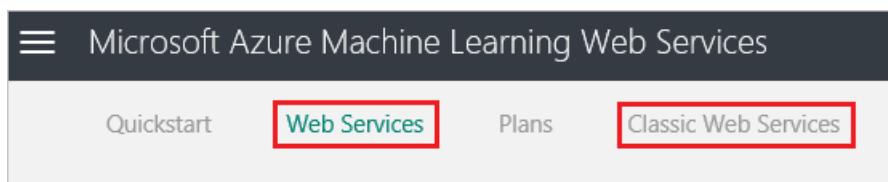
How to enable logging for a Web service

You enable logging from the [Azure Machine Learning Web Services](#) portal.

1. Sign in to the Azure Machine Learning Web Services portal at <https://services.azureml.net>. For a Classic web service, you can also get to the portal by clicking **New Web Services Experience** on the Machine Learning Web Services page in Machine Learning Studio.



2. On the top menu bar, click **Web Services** for a New web service, or click **Classic Web Services** for a Classic web service.



3. For a New web service, click the web service name. For a Classic web service, click the web service name and then on the next page click the appropriate endpoint.
4. On the top menu bar, click **Configure**.
5. Set the **Enable Logging** option to *Error* (to log only errors) or *All* (for full logging).

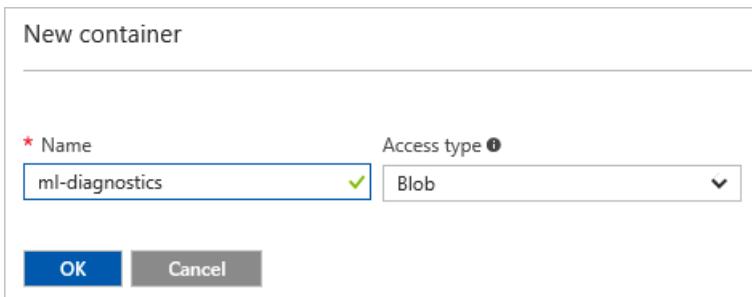


6. Click **Save**.
7. For Classic web services, create the **ml-diagnostics** container.

All web service logs are kept in a blob container named **ml-diagnostics** in the storage account associated with the web service. For New web services, this container is created the first time you access the web service. For Classic web services, you need to create the container if it doesn't already exist.

- a. In the [Azure portal](#), go to the storage account associated with the web service.

- b. Under **Blob Service**, click **Containers**.
- c. If the container **ml-diagnostics** doesn't exist, click **+Container**, give the container the name "ml-diagnostics", and select the **Access type** as "Blob". Click **OK**.



TIP

For a Classic web service, the Web Services Dashboard in Machine Learning Studio also has a switch to enable logging. However, because logging is now managed through the Web Services portal, you need to enable logging through the portal as described in this article. If you already enabled logging in Studio, then in the Web Services Portal, disable logging and enable it again.

The effects of enabling logging

When logging is enabled, the diagnostics and errors from the web service endpoint are logged in the **ml-diagnostics** blob container in the Azure Storage Account linked with the user's workspace. This container holds all the diagnostics information for all the web service endpoints for all the workspaces associated with this storage account.

The logs can be viewed using any of the several tools available to explore an Azure Storage Account. The easiest may be to navigate to the storage account in the Azure portal, click **Containers**, and then click the container **ml-diagnostics**.

Log blob detail information

Each blob in the container holds the diagnostics information for exactly one of the following actions:

- An execution of the Batch-Execution method
- An execution of the Request-Response method
- Initialization of a Request-Response container

The name of each blob has a prefix of the following form:

{Workspace Id}-{Web service Id}-{Endpoint Id}/{Log type}

Where *Log type* is one of the following values:

- batch
- score/requests
- score/init

Manage a Web service using the Azure Machine Learning Web Services portal

12/8/2017 • 7 min to read • [Edit Online](#)

You can manage your Machine Learning New and Classic Web services using the Microsoft Azure Machine Learning Web Services portal. Since Classic Web services and New Web services are based on different underlying technologies, you have slightly different management capabilities for each of them.

In the Machine Learning Web Services portal you can:

- Monitor how the web service is being used.
- Configure the description, update the keys for the web service (New only), update your storage account key (New only), enable logging, and enable or disable sample data.
- Delete the web service.
- Create, delete, or update billing plans (New only).
- Add and delete endpoints (Classic only)

NOTE

Try Azure Machine Learning for free

No credit card or Azure subscription required. [Get started now](#).

Permissions to manage New Resources Manager based web services

New web services are deployed as Azure resources. As such, you must have the correct permissions to deploy and manage New web services. To deploy or manage New web services you must be assigned a contributor or administrator role on the subscription to which the web service is deployed. If you invite another user to a machine learning workspace, you must assign them to a contributor or administrator role on the subscription before they can deploy or manage web services.

If the user does not have the correct permissions to access resources in the Azure Machine Learning Web Services portal, they will receive the following error when trying to deploy a web service:

Web Service deployment failed. This account does not have sufficient access to the Azure subscription that contains the Workspace. In order to deploy a Web Service to Azure, the same account must be invited to the Workspace and be given access to the Azure subscription that contains the Workspace.

For more information on creating a workspace, see [Create and share an Azure Machine Learning workspace](#).

For more information on setting access permissions, see [View access assignments for users and groups in the Azure portal - Public preview](#).

Manage New Web services

To manage your New Web services:

1. Sign in to the [Microsoft Azure Machine Learning Web Services](#) portal using your Microsoft Azure account - use the account that's associated with the Azure subscription.
2. On the menu, click **Web Services**.

This displays a list of deployed Web services for your subscription.

To manage a Web service, click Web Services. From the Web Services page you can:

- Click the web service to manage it.
- Click the Billing Plan for the web service to update it.
- Delete a web service.
- Copy a web service and deploy it to another region.

When you click a web service, the web service Quickstart page opens. The web service Quickstart page has two menu options that allow you to manage your web service:

- **DASHBOARD** - Allows you to view Web service usage.
- **CONFIGURE** - Allows you to add descriptive text, update the key for the storage account associated with the Web service, and enable or disable sample data.

Monitoring how the web service is being used

Click the **DASHBOARD** tab.

From the dashboard, you can view overall usage of your Web service over a period of time. You can select the period to view from the Period dropdown menu in the upper right of the usage charts. The dashboard shows the following information:

- **Requests Over Time** displays a step graph of the number of requests over the selected time period. It can help identify if you are experiencing spikes in usage.
- **Request-Response Requests** displays the total number of Request-Response calls the service has received over the selected time period and how many of them failed.
- **Average Request-Response Compute Time** displays an average of the time needed to execute the received requests.
- **Batch Requests** displays the total number of Batch Requests the service has received over the selected time period and how many of them failed.
- **Average Job Latency** displays an average of the time needed to execute the received requests.
- **Errors** displays the aggregate number of errors that have occurred on calls to the web service.
- **Services Costs** displays the charges for the billing plan associated with the service.

Configuring the web service

Click the **CONFIGURE** menu option.

You can update the following properties:

- **Description** allows you to enter a description for the Web service.
- **Title** allows you to enter a title for the Web service
- **Keys** allows you to rotate your primary and secondary API keys.
- **Storage account key** allows you to update the key for the storage account associated with the Web service changes.
- **Enable Sample data** allows you to provide sample data that you can use to test the Request-Response service. If you created the web service in Machine Learning Studio, the sample data is taken from the data you used to train your model. If you created the service programmatically, the data is taken from the example data you provided as part of the JSON package.

Managing billing plans

Click the **Plans** menu option from the web services Quickstart page. You can also click the plan associated with specific Web service to manage that plan.

- **New** allows you to create a new plan.

- **Add/Remove Plan instance** allows you to "scale out" an existing plan to add capacity.
- **Upgrade/DownGrade** allows you to "scale up" an existing plan to add capacity.
- **Delete** allows you to delete a plan.

Click a plan to view its dashboard. The dashboard gives you snapshot or plan usage over a selected period of time. To select the time period to view, click the **Period** dropdown at the upper right of dashboard.

The plan dashboard provides the following information:

- **Plan Description** displays information about the costs and capacity associated with the plan.
- **Plan Usage** displays the number of transactions and compute hours that have been charged against the plan.
- **Web Services** displays the number of Web services that are using this plan.
- **Top Web Service By Calls** displays the top four Web services that are making calls that are charged against the plan.
- **Top Web Services by Compute Hrs** displays the top four Web services that are using compute resources that are charged against the plan.

Manage Classic Web Services

NOTE

The procedures in this section are relevant to managing Classic web services through the Azure Machine Learning Web Services portal. For information on managing Classic Web services through the Machine Learning Studio and the Azure classic portal, see [Manage an Azure Machine Learning workspace](#).

To manage your Classic Web services:

1. Sign in to the [Microsoft Azure Machine Learning Web Services](#) portal using your Microsoft Azure account - use the account that's associated with the Azure subscription.
2. On the menu, click **Classic Web Services**.

To manage a Classic Web Service, click **Classic Web Services**. From the Classic Web Services page you can:

- Click the web service to view the associated endpoints.
- Delete a web service.

When you manage a Classic Web service, you manage each of the endpoints separately. When you click a web service in the Web Services page, the list of endpoints associated with the service opens.

On the Classic Web Service endpoint page, you can add and delete endpoints on the service. For more information on adding endpoints, see [Creating Endpoints](#).

Click one of the endpoints to open the web service Quickstart page. On the Quickstart page, there are two menu options that allow you to manage your web service:

- **DASHBOARD** - Allows you to view Web service usage.
- **CONFIGURE** - Allows you to add descriptive text, turn error logging on and off, update the key for the storage account associated with the Web service, and enable and disable sample data.

Monitoring how the web service is being used

Click the **DASHBOARD** tab.

From the dashboard, you can view overall usage of your Web service over a period of time. You can select the period to view from the Period dropdown menu in the upper right of the usage charts. The dashboard shows the

following information:

- **Requests Over Time** displays a step graph of the number of requests over the selected time period. It can help identify if you are experiencing spikes in usage.
- **Request-Response Requests** displays the total number of Request-Response calls the service has received over the selected time period and how many of them failed.
- **Average Request-Response Compute Time** displays an average of the time needed to execute the received requests.
- **Batch Requests** displays the total number of Batch Requests the service has received over the selected time period and how many of them failed.
- **Average Job Latency** displays an average of the time needed to execute the received requests.
- **Errors** displays the aggregate number of errors that have occurred on calls to the web service.
- **Services Costs** displays the charges for the billing plan associated with the service.

Configuring the web service

Click the **CONFIGURE** menu option.

You can update the following properties:

- **Description** allows you to enter a description for the Web service. Description is a required field.
- **Logging** allows you to enable or disable error logging on the endpoint. For more information on Logging, see [Enable logging for Machine Learning web services](#).
- **Enable Sample data** allows you to provide sample data that you can use to test the Request-Response service. If you created the web service in Machine Learning Studio, the sample data is taken from the data you used to train your model. If you created the service programmatically, the data is taken from the example data you provided as part of the JSON package.

Learn how to manage AzureML web services using API Management

11/30/2017 • 11 min to read • [Edit Online](#)

Overview

This guide shows you how to quickly get started using API Management to manage your AzureML web services.

What is Azure API Management?

Azure API Management is an Azure service that lets you manage your REST API endpoints by defining user access, usage throttling, and dashboard monitoring. Click [here](#) for details on Azure API Management. Click [here](#) for a guide on how to get started with Azure API Management. This other guide, which this guide is based on, covers more topics, including notification configurations, tier pricing, response handling, user authentication, creating products, developer subscriptions, and usage dashboarding.

What is AzureML?

AzureML is an Azure service for machine learning that enables you to easily build, deploy, and share advanced analytics solutions. Click [here](#) for details on AzureML.

Prerequisites

To complete this guide, you need:

- An Azure account. If you don't have an Azure account, click [here](#) for details on how to create a free trial account.
- An AzureML account. If you don't have an AzureML account, click [here](#) for details on how to create a free trial account.
- The workspace, service, and api_key for an AzureML experiment deployed as a web service. Click [here](#) for details on how to create an AzureML experiment. Click [here](#) for details on how to deploy an AzureML experiment as a web service. Alternately, Appendix A has instructions for how to create and test a simple AzureML experiment and deploy it as a web service.

Create an API Management instance

You can manage your Azure Machine Learning web service with an API Management instance.

1. Sign in to the [Azure portal](#).
2. Select **+ Create a resource**.
3. In the search box, type "API management", then select the "API management" resource.
4. Click **Create**.
5. The **Name** value will be used to create a unique URL (this example uses "demoazureml").
6. Select a **Subscription**, **Resource group**, and **Location** for your service instance.
7. Specify a value for **Organization name** (this example uses "demoazureml").
8. Enter your **Administrator email** - this email will be used for notifications from the API Management system.
9. Click **Create**.

It may take up to 30 minutes for a new service to be created.

API Management service

* Name
demoazureml .azure-api.net

* Subscription
my_subscription

* Resource group
 Create new Use existing
demoazureml

* Location
West Central US

* Organization name ⓘ
demoazureml

* Administrator email ⓘ
admin@contoso.com

Pricing tier (View full pricing details)
Developer (1M API calls, 5GB data transfer)

Pin to dashboard

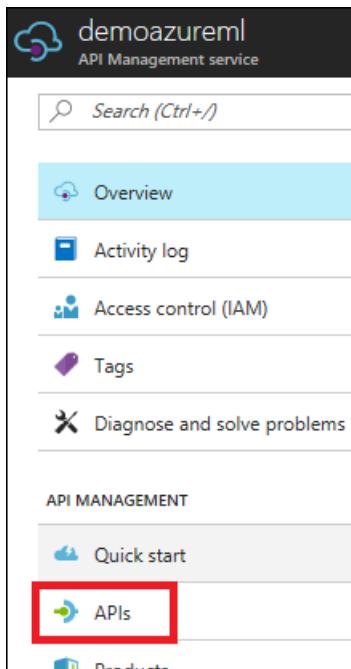
Create Automation options

Create the API

Once the service instance is created, the next step is to create the API. An API consists of a set of operations that can be invoked from a client application. API operations are proxied to existing web services. This guide creates APIs that proxy to the existing AzureML RRS and BES web services.

To create the API:

1. In the Azure portal, open the service instance you just created.
2. In the left navigation pane, select **APIs**.



3. Click **Add API**.
4. Enter a **Web API name** (this example uses "AzureML Demo API").
5. For **Web service URL**, enter "`https://ussouthcentral.services.azureml.net`".
6. Enter a **Web API URL suffix". This will become the last part of the URL that customers will use for sending requests to the service instance (this example uses "azureml-demo").
7. For **Web API URL scheme**, select **HTTPS**.
8. For **Products**, select **Starter**.
9. Click **Save**.

Add the operations

Operations are added and configured to an API in the publisher portal. To access the publisher portal, click **Publisher portal** in the Azure portal for your API Management service, select **APIs, Operations**, then click **Add operation**.

The screenshot shows the Azure API Management Publisher portal interface. On the left, there's a sidebar with 'API MANAGEMENT' and several menu items: Dashboard, APIs (which is highlighted with a red box), Products, Policies, Analytics, Users, Groups, and Notifications. The main content area is titled 'APIs - AzureML Demo API'. At the top, there are tabs: Summary, Settings, Operations (which is highlighted with a red box), Security, Issues, and Products. Below the tabs, the word 'Operations' is displayed. A descriptive text says 'Define service operations to enable service documentation, interactive API console and operation-level statistics.' At the bottom of this section, there's a button labeled '+ ADD OPERATION' with a red box around it. At the very bottom of the main content area, there are links for 'POST' and 'Create resource'.

The **New operation** window will be displayed and the **Signature** tab will be selected by default.

Add RRS Operation

First create an operation for the AzureML RRS service:

1. For the **HTTP verb**, select **POST**.
2. For the **URL template**, type "`/workspaces/{workspace}/services/{service}/execute?api-version={apiversion}&details={details}`".
3. Enter a **Display name** (this example uses "RRS Execute").

Signature	
Caching	HTTP verb* POST
REQUEST	URL template* /workspaces/{workspace}/services/{service}/execute?api-version={apiversion}&details Example: GET Example: customers/{cid}/orders/{oid}/?date={date}
Parameters	Rewrite URL template
Body	<i>When specified, rewrite template will be used to make requests to the web service. It can contain all of the template parameters specified in the original template.</i>
RESPONSES	Display name* RRS Execute
+ ADD	

4. Click **Responses > ADD** on the left and select **200 OK**.

5. Click **Save** to save this operation.

Operation - RRS Execute	
Signature	Code 200 OK
Caching	✖ DELETE THIS RESPONSE
REQUEST	Description
Parameters	Enter description
Body	
RESPONSES	+ ADD REPRESENTATION
Code 200	<i>If operation generates responses in one or more representation formats (e.g. JSON and/or XML), you may describe them in this section.</i>
+ ADD	
Save	

Add BES Operations

NOTE

Screenshots are not included here for the BES operations as they are very similar to those for adding the RRS operation.

Submit (but not start) a Batch Execution job

1. Click **add operation** to add a BES operation to the API.
2. For the **HTTP verb**, select **POST**.
3. For the **URL template**, type "`/workspaces/{workspace}/services/{service}/jobs?api-version={apiversion}`".
4. Enter a **Display name** (this example uses "BES Submit").
5. Click **Responses > ADD** on the left and select **200 OK**.
6. Click **Save**.

Start a Batch Execution job

1. Click **add operation** to add a BES operation to the API.
2. For the **HTTP verb**, select **POST**.

3. For the **HTTP verb**, type "`/workspaces/{workspace}/services/{service}/jobs/{jobid}/start?api-version={apiversion}`".
4. Enter a **Display name** (this example uses "BES Start").
5. Click **Responses > ADD** on the left and select **200 OK**.
6. Click **Save**.

Get the status or result of a Batch Execution job

1. Click **add operation** to add a BES operation to the API.
2. For the **HTTP verb**, select **GET**.
3. For the **URL template**, type "`/workspaces/{workspace}/services/{service}/jobs/{jobid}?api-version={apiversion}`".
4. Enter a **Display name** (this example uses "BES Status").
5. Click **Responses > ADD** on the left and select **200 OK**.
6. Click **Save**.

Delete a Batch Execution job

1. Click **add operation** to add a BES operation to the API.
2. For the **HTTP verb**, select **DELETE**.
3. For the **URL template**, type "`/workspaces/{workspace}/services/{service}/jobs/{jobid}?api-version={apiversion}`".
4. Enter a **Display name** (this example uses "BES Delete").
5. Click **Responses > ADD** on the left and select **200 OK**.
6. Click **Save**.

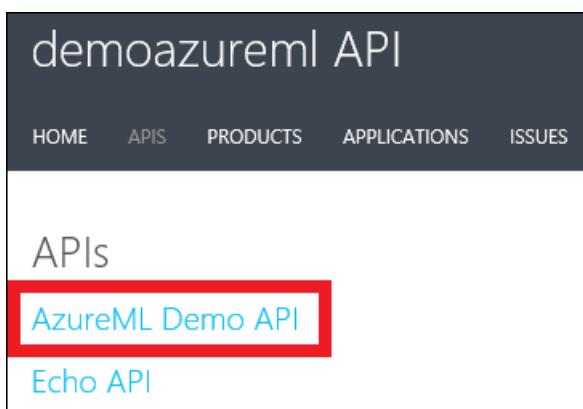
Call an operation from the Developer portal

Operations can be called directly from the Developer portal, which provides a convenient way to view and test the operations of an API. In this step you will call the **RRS Execute** method that was added to the **AzureML Demo API**.

1. Click **Developer portal**.



2. Click **APIs** from the top menu, and then click **AzureML Demo API** to see the operations available.



3. Select **RRS Execute** for the operation. Click **Try It**.



4. For **Request parameters**, type your **workspace** and **service**, type "2.0" for the **apiversion**, and "true" for the **details**. You can find your **workspace** and **service** in the AzureML web service dashboard (see **Test the web service** in Appendix A).

For **Request headers**, click **Add header** and type "Content-Type" and "application/json". Click **Add header** again and type "Authorization" and "Bearer <your service API-KEY>". You can find your API-KEY in the AzureML web service dashboard (see **Test the web service** in Appendix A).

For **Request body**, type

```
{"Inputs": {"input1": {"ColumnNames": ["Col2"], "Values": [["This is a good day"]]}, "GlobalParameters": {}}
```

AzureML Demo API

</workspaces/{workspace}/services/{service}/execute?api-version={apiversion}&details={details}>

Request parameters

workspace	<input type="text"/>
service	<input type="text"/>
apiversion	2.0
details	true

[+ Add parameter](#)

Request headers

Ocp-Apim-Trace	true	
Ocp-Apim-Subscription-Key	*****	
Content-Type	application/json	✖ Remove header
Authorization	Bearer /	✖ Remove header

[+ Add header](#)

Request body

```
{"Inputs": {"input1": {"ColumnNames": ["Col2"], "Values": [["This is a good day"]]}, "GlobalParameters": {}}
```

Authorization

Subscription key	<input type="text"/>	x v
------------------	----------------------	-----

5. Click **Send**.

[Send](#)

After an operation is invoked, the developer portal displays the **Requested URL** from the back-end service, the **Response status**, the **Response headers**, and any **Response content**.

Response status
200 OK

Response latency
438 ms

Response headers

```

1. x-ms-request-id: 3b648c32-9c35-40b6-a309-138396f59e12
2. Ocp-Apim-Trace-Location: https://apimgmtstsfa4wt270nq7f1.blob.core.windows.net/apiinspectorcontainer/k4m0GuwG51EdW
   1uyP4rCvA2-2?sv=2013-08-15&sr=b&sig=N%2B%2By6qu%2Fi1pDs0WI0X1zwqArAyUFS05aooy2gNPL0FA%3D&se=2015-06-03T19%3A16%3A27
   Z&sp=r&traceId=5be3eddac41342e789b5a6c656b94730
3. Date: Tue, 02 Jun 2015 19:16:27 GMT
4. Content-Length: 428
5. Content-Type: application/json; charset=utf-8

```

Response content

```
[
  "This is a good day",
  "1",
  "1",
  "2",
  "2",
  "2",
  "0",
  "2",
  "0",
  "1"
]
```

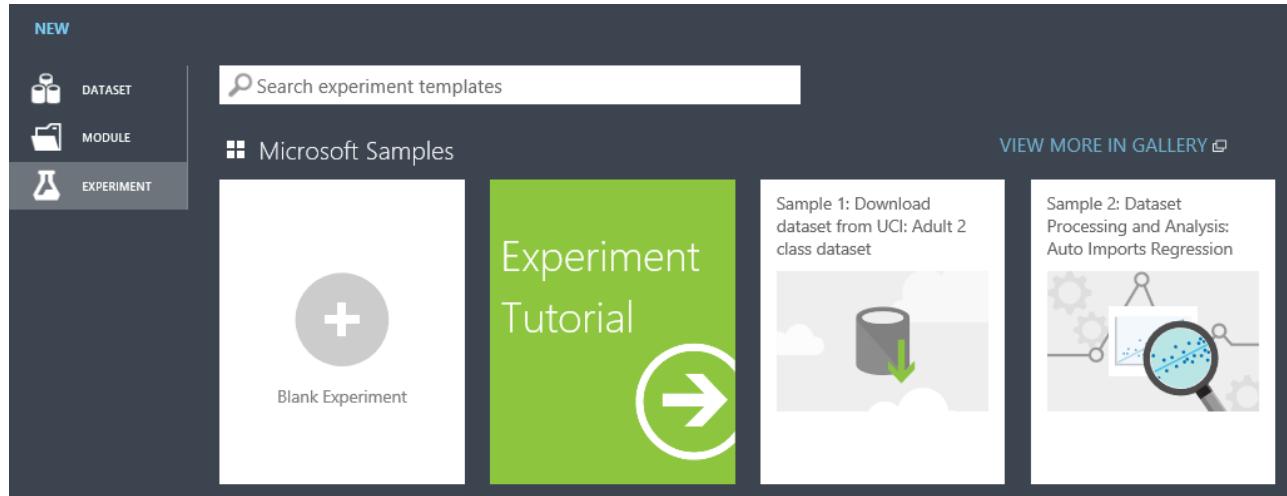
Appendix A - Creating and testing a simple AzureML web service

Creating the experiment

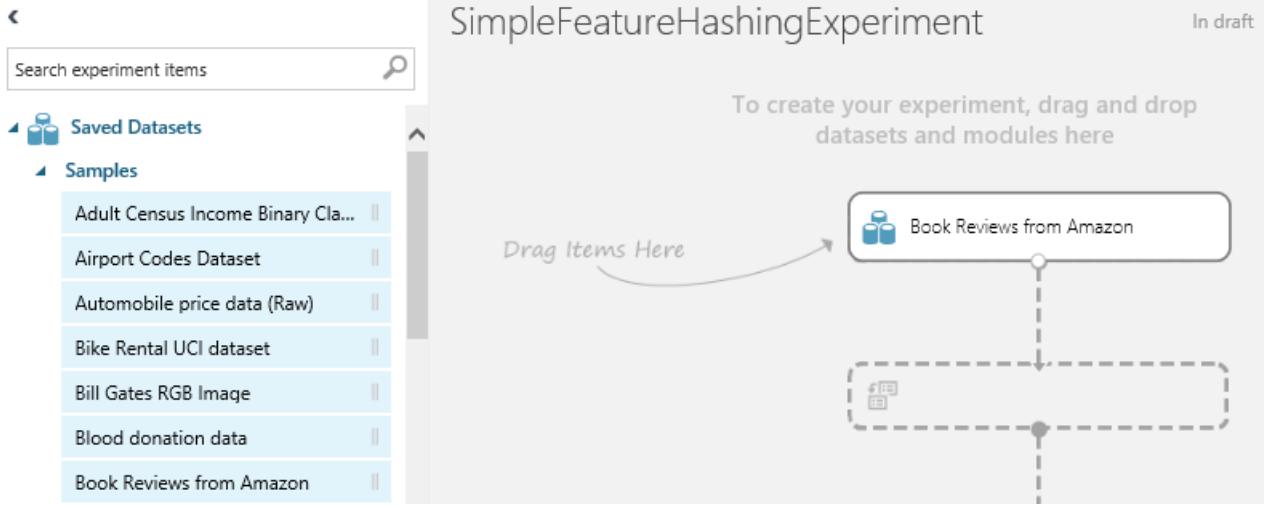
Below are the steps for creating a simple AzureML experiment and deploying it as a web service. The web service takes as input a column of arbitrary text and returns a set of features represented as integers. For example:

TEXT	HASHED TEXT
This is a good day	1 1 2 2 0 2 0 1

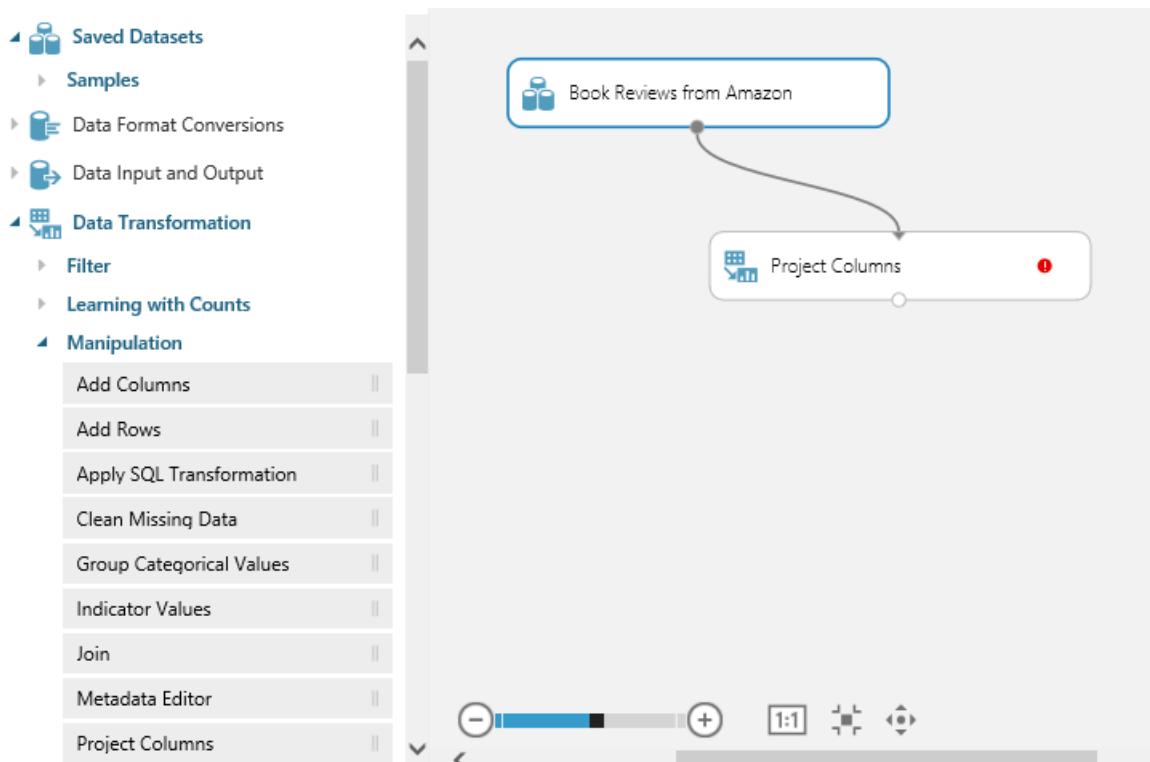
First, using a browser of your choice, navigate to: <https://studio.azureml.net/> and enter your credentials to log in. Next, create a new blank experiment.



Rename it to **SimpleFeatureHashingExperiment**. Expand **Saved Datasets** and drag **Book Reviews from Amazon** onto your experiment.



Expand **Data Transformation** and **Manipulation** and drag **Select Columns in Dataset** onto your experiment. Connect **Book Reviews from Amazon** to **Select Columns in Dataset**.



Click **Select Columns in Dataset** and then click **Launch column selector** and select **Col2**. Click the checkmark to apply these changes.

Select columns ×

Allow duplicates and preserve column order in selection

Begin With ▼

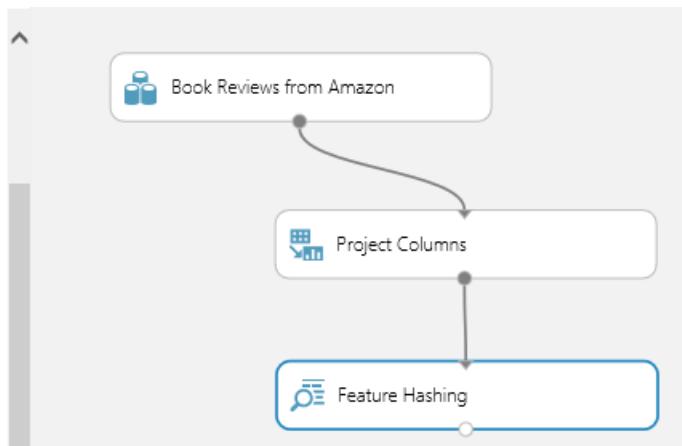
Include ▼ ▼ Col2 X + -

✓

Expand **Text Analytics** and drag **Feature Hashing** onto the experiment. Connect **Select Columns in Dataset** to

Feature Hashing.

- ▶ Data Transformation
- ▶ Feature Selection
- ▶ Machine Learning
- ▶ OpenCV Library Modules
- ▶ Python Language Modules
- ▶ R Language Modules
- ▶ Statistical Functions
- ◀ **Text Analytics**
- Feature Hashing



Type **3** for the **Hashing bitsize**. This will create 8 (23) columns.

Properties

◀ Feature Hashing

Target column(s)

Selected columns:

Column type: String,
Feature

Launch column selector

Hashing bitsize

3

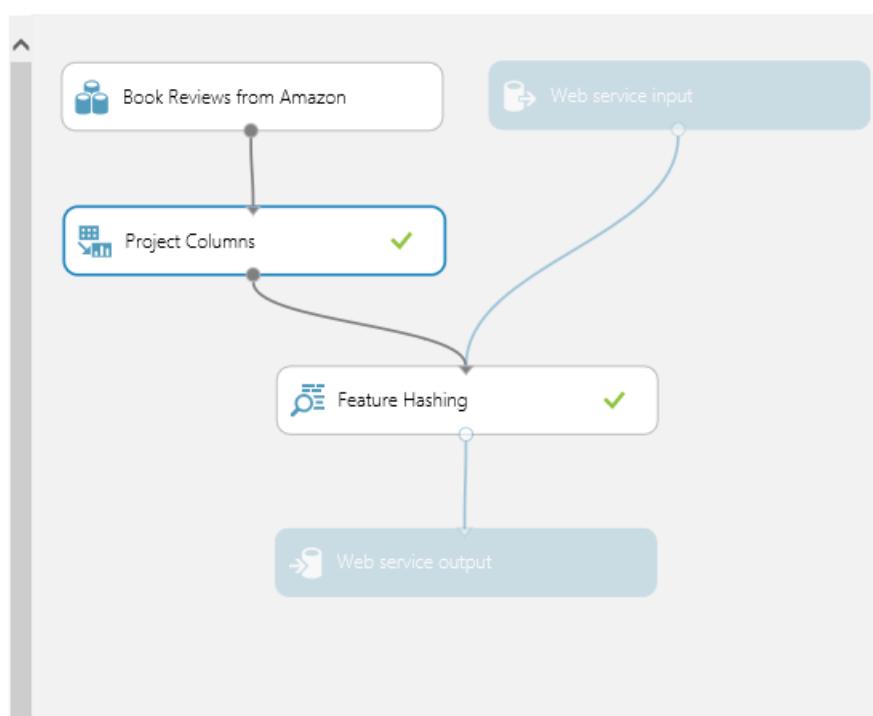
At this point, you may want to click **Run** to test the experiment.



Create a web service

Now create a web service. Expand **Web Service** and drag **Input** onto your experiment. Connect **Input** to **Feature Hashing**. Also drag **output** onto your experiment. Connect **Output** to **Feature Hashing**.

- ◀ Saved Datasets
 - ▶ Samples
- ▶ Data Format Conversions
- ▶ Data Input and Output
- ▶ Data Transformation
- ▶ Feature Selection
- ▶ Machine Learning
- ▶ OpenCV Library Modules
- ▶ Python Language Modules
- ▶ R Language Modules
- ▶ Statistical Functions
- ▶ Text Analytics
- ◀ **Web Service**
 - Input
 - Output



Click **Publish web service**.



Click **Yes** to publish the experiment.



Test the web service

An AzureML web service consists of RSS (request/response service) and BES (batch execution service) endpoints. RSS is for synchronous execution. BES is for asynchronous job execution. To test your web service with the sample Python source below, you may need to download and install the Azure SDK for Python (see: [How to install Python](#)).

You will also need the **workspace**, **service**, and **api_key** of your experiment for the sample source below. You can find the workspace and service by clicking either **Request/Response** or **Batch Execution** for your experiment in the web service dashboard.

Request

Method Request URI

POST https://ussouthcentral.services.azureml.net/workspaces/bf4161ba/services/c9129f01/ex
api-version=2.0&details=true

workspace service

You can find the **api_key** by clicking your experiment in the web service dashboard.

DASHBOARD CONFIGURATION

General

Published experiment

[View snapshot](#) [View latest](#)

Description

No description provided for this web service.

API key

/EHr+Icv vvdeJl

Default Endpoint

REQUEST/RESPONSE TEST APPS LAST UPDATED

BATCH EXECUTION

Download Excel Workbook 5/5/2015 10:13:51 AM

5/5/2015 10:13:51 AM

Test RRS endpoint

Test button

An easy way to test the RRS endpoint is to click **Test** on the web service dashboard.

DASHBOARD CONFIGURATION

General

Published experiment

[View snapshot](#) [View latest](#)

Description

No description provided for this web service.

API key

/EHr+Icv vvdeJl

Default Endpoint

REQUEST/RESPONSE TEST APPS LAST UPDATED

BATCH EXECUTION

Type **This is a good day** for **col2**. Click the checkmark.



Enter data to predict

COL2



You will see something like

'SimpleFeatureHashingExperiment' test returned ["This is a good day","1","1","2","2","0","2","0","1"]...

Sample Code

Another way to test your RRS is from your client code. If you click **Request/response** on the dashboard and scroll to the bottom, you will see sample code for C#, Python, and R. You will also see the syntax of the RRS request, including the request URI, headers, and body.

This guide shows a working Python example. You will need to modify it with the **workspace**, **service**, and **api_key** of your experiment.

```
import urllib2
import json
workspace = "<REPLACE WITH YOUR EXPERIMENT'S WEB SERVICE WORKSPACE ID>"
service = "<REPLACE WITH YOUR EXPERIMENT'S WEB SERVICE SERVICE ID>"
api_key = "<REPLACE WITH YOUR EXPERIMENT'S WEB SERVICE API KEY>"
data = {
    "Inputs": {
        "input1": {
            "ColumnNames": ["Col2"],
            "Values": [ ["This is a good day"] ]
        },
    },
    "GlobalParameters": {}
}
url = "https://ussouthcentral.services.azureml.net/workspaces/" + workspace + "/services/" + service + "/execute?api-version=2.0&details=true"
headers = {'Content-Type':'application/json', 'Authorization':('Bearer ' + api_key)}
body = str.encode(json.dumps(data))
try:
    req = urllib2.Request(url, body, headers)
    response = urllib2.urlopen(req)
    result = response.read()
    print "result:" + result
except urllib2.HTTPError, error:
    print("The request failed with status code: " + str(error.code))
    print(error.info())
    print(json.loads(error.read()))
```

Test BES endpoint

Click **Batch execution** on the dashboard and scroll to the bottom. You will see sample code for C#, Python, and R. You will also see the syntax of the BES requests to submit a job, start a job, get the status or results of a job, and delete a job.

This guide shows a working Python example. You need to modify it with the **workspace**, **service**, and **api_key** of your experiment. Additionally, you need to modify the **storage account name**, **storage account key**, and **storage container name**. Lastly, you will need to modify the location of the **input file** and the location of the **output file**.

```

import urllib2
import json
import time
from azure.storage import *
workspace = "<REPLACE WITH YOUR WORKSPACE ID>"
service = "<REPLACE WITH YOUR SERVICE ID>"
api_key = "<REPLACE WITH THE API KEY FOR YOUR WEB SERVICE>"
storage_account_name = "<REPLACE WITH YOUR AZURE STORAGE ACCOUNT NAME>"
storage_account_key = "<REPLACE WITH YOUR AZURE STORAGE KEY>"
storage_container_name = "<REPLACE WITH YOUR AZURE STORAGE CONTAINER NAME>"
input_file = "<REPLACE WITH THE LOCATION OF YOUR INPUT FILE>" # Example: C:\\mydata.csv
output_file = "<REPLACE WITH THE LOCATION OF YOUR OUTPUT FILE>" # Example: C:\\myresults.csv
input_blob_name = "mydatablob.csv"
output_blob_name = "myresultsblob.csv"
def printHttpError(httpError):
    print("The request failed with status code: " + str(httpError.code))
    print(httpError.info())
    print(json.loads(httpError.read()))
    return
def saveBlobToFile(blobUrl, resultsLabel):
    print("Reading the result from " + blobUrl)
    try:
        response = urllib2.urlopen(blobUrl)
    except urllib2.HTTPError, error:
        printHttpError(error)
        return
    with open(output_file, "w+") as f:
        f.write(response.read())
    print(resultsLabel + " have been written to the file " + output_file)
    return
def processResults(result):
    first = True
    results = result["Results"]
    for outputName in results:
        result_blob_location = results[outputName]
        sas_token = result_blob_location["SasBlobToken"]
        base_url = result_blob_location["BaseLocation"]
        relative_url = result_blob_location["RelativeLocation"]
        print("The results for " + outputName + " are available at the following Azure Storage location:")
        print("BaseLocation: " + base_url)
        print("RelativeLocation: " + relative_url)
        print("SasBlobToken: " + sas_token)
        if(first):
            first = False
            url3 = base_url + relative_url + sas_token
            saveBlobToFile(url3, "The results for " + outputName)
    return
def invokeBatchExecutionService():
    url = "https://ussouthcentral.services.azureml.net/workspaces/" + workspace + "/services/" + service + "/jobs"
    blob_service = BlobService(account_name=storage_account_name, account_key=storage_account_key)
    print("Uploading the input to blob storage...")
    data_to_upload = open(input_file, "r").read()
    blob_service.put_blob(storage_container_name, input_blob_name, data_to_upload, x_ms_blob_type="BlockBlob")
    print "Uploaded the input to blob storage"
    input_blob_path = "/" + storage_container_name + "/" + input_blob_name
    connection_string = "DefaultEndpointsProtocol=https;AccountName=" + storage_account_name + ";AccountKey=" + storage_account_key
    payload = {
        "Input": {
            "ConnectionString": connection_string,
            "RelativeLocation": input_blob_path
        },
        "Outputs": {
            "output1": { "ConnectionString": connection_string, "RelativeLocation": "/" + storage_container_name + "/" + output_blob_name },
        },
        "GlobalParameters": {
        }
    }

```

```

}

body = str.encode(json.dumps(payload))
headers = { "Content-Type":"application/json", "Authorization":("Bearer " + api_key)}
print("Submitting the job...")
# submit the job
req = urllib2.Request(url + "?api-version=2.0", body, headers)
try:
    response = urllib2.urlopen(req)
except urllib2.HTTPError, error:
    printHttpError(error)
    return
result = response.read()
job_id = result[1:-1] # remove the enclosing double-quotes
print("Job ID: " + job_id)
# start the job
print("Starting the job...")
req = urllib2.Request(url + "/" + job_id + "/start?api-version=2.0", "", headers)
try:
    response = urllib2.urlopen(req)
except urllib2.HTTPError, error:
    printHttpError(error)
    return
url2 = url + "/" + job_id + "?api-version=2.0"

while True:
    print("Checking the job status...")
    # If you are using Python 3+, replace urllib2 with urllib.request in the following code
    req = urllib2.Request(url2, headers = { "Authorization":("Bearer " + api_key) })
    try:
        response = urllib2.urlopen(req)
    except urllib2.HTTPError, error:
        printHttpError(error)
        return
    result = json.loads(response.read())
    status = result["StatusCode"]
    print "status:" + status
    if(status == 0 or status == "NotStarted"):
        print("Job " + job_id + " not yet started...")
    elif(status == 1 or status == "Running"):
        print("Job " + job_id + " running...")
    elif(status == 2 or status == "Failed"):
        print("Job " + job_id + " failed!")
        print("Error details: " + result["Details"])
        break
    elif(status == 3 or status == "Cancelled"):
        print("Job " + job_id + " cancelled!")
        break
    elif(status == 4 or status == "Finished"):
        print("Job " + job_id + " finished!")
        processResults(result)
        break
    time.sleep(1) # wait one second
return
invokeBatchExecutionService()

```

Creating endpoints

11/30/2017 • 1 min to read • [Edit Online](#)

NOTE

This topic describes techniques applicable to a **Classic** Machine Learning Web service.

When you create Web services that you sell forward to your customers, you need to provide trained models to each customer that are still linked to the experiment from which the Web service was created. In addition, any updates to the experiment should be applied selectively to an endpoint without overwriting the customizations.

To accomplish this, Azure Machine Learning allows you to create multiple endpoints for a deployed Web service. Each endpoint in the Web service is independently addressed, throttled, and managed. Each endpoint is a unique URL and authorization key that you can distribute to your customers.

NOTE

[Try Azure Machine Learning for free](#)

No credit card or Azure subscription required. [Get started now](#).

Adding endpoints to a Web service

There are two ways to add an endpoint to a Web service.

- Programmatically
- Through the Azure Machine Learning Web Services portal

Once the endpoint is created, you can consume it through synchronous APIs, batch APIs, and excel worksheets. In addition to adding endpoints through this UI, you can also use the Endpoint Management APIs to programmatically add endpoints.

NOTE

If you have added additional endpoints to the Web service, you cannot delete the default endpoint.

Adding an endpoint programmatically

You can add an endpoint to your Web service programmatically using the [AddEndpoint](#) sample code.

Adding an endpoint using the Azure Machine Learning Web Services portal

1. In Machine Learning Studio, on the left navigation column, click Web Services.
2. At the bottom of the Web service dashboard, click **Manage endpoints**. The Azure Machine Learning Web Services portal opens to the endpoints page for the Web service.
3. Click **New**.
4. Type a name and description for the new endpoint. Endpoint names must be 24 character or less in length, and

must be made up of lower-case alphabets or numbers. Select the logging level and whether sample data is enabled. For more information on logging, see [Enable logging for Machine Learning Web services](#).

Next steps

[How to consume an Azure Machine Learning Web service.](#)

Scaling an Azure Machine Learning web service by adding additional endpoints

11/30/2017 • 1 min to read • [Edit Online](#)

NOTE

This topic describes techniques applicable to a **Classic** Machine Learning Web service.

By default, each published Web service is configured to support 20 concurrent requests and can be as high as 200 concurrent requests. Azure Machine Learning automatically optimizes the setting to provide the best performance for your web service and the portal value is ignored.

If you plan to call the API with a higher load than a Max Concurrent Calls value of 200 will support, you should create multiple endpoints on the same Web service. You can then randomly distribute your load across all of them.

The scaling of a Web service is a common task. Some reasons to scale are to support more than 200 concurrent requests, increase availability through multiple endpoints, or provide separate endpoints for the web service. You can increase the scale by adding additional endpoints for the same Web service through the [Azure Machine Learning Web Service](#) portal.

For more information on adding new endpoints, see [Creating Endpoints](#).

Keep in mind that using a high concurrency count can be detrimental if you're not calling the API with a correspondingly high rate. You might see sporadic timeouts and/or spikes in the latency if you put a relatively low load on an API configured for high load.

The synchronous APIs are typically used in situations where a low latency is desired. Latency here implies the time it takes for the API to complete one request, and doesn't account for any network delays. Let's say you have an API with a 50-ms latency. To fully consume the available capacity with throttle level High and Max Concurrent Calls = 20, you need to call this API $20 * 1000 / 50 = 400$ times per second. Extending this further, a Max Concurrent Calls of 200 allows you to call the API 4000 times per second, assuming a 50-ms latency.

Retrain a Machine Learning Model

12/8/2017 • 3 min to read • [Edit Online](#)

As part of the process of operationalization of machine learning models in Azure Machine Learning, your model is trained and saved. You then use it to create a predictive Web service. The Web service can then be consumed in web sites, dashboards, and mobile apps.

Models you create using Machine Learning are typically not static. As new data becomes available or when the consumer of the API has their own data the model needs to be retrained.

Retraining may occur frequently. With the Programmatic Retraining API feature, you can programmatically retrain the model using the Retraining APIs and update the Web service with the newly trained model.

This document describes the retraining process, and shows you how to use the Retraining APIs.

Why retrain: defining the problem

As part of the machine learning training process, a model is trained using a set of data. Models you create using Machine Learning are typically not static. As new data becomes available or when the consumer of the API has their own data the model needs to be retrained.

In these scenarios, a programmatic API provides a convenient way to allow you or the consumer of your APIs to create a client that can, on a one-time or regular basis, retrain the model using their own data. They can then evaluate the results of retraining, and update the Web service API to use the newly trained model.

NOTE

If you have an existing Training Experiment and New Web service, you may want to check out Retrain an existing Predictive Web service instead of following the walkthrough mentioned in the following section.

End-to-end workflow

The process involves the following components: A Training Experiment and a Predictive Experiment published as a Web service. To enable retraining of a trained model, the Training Experiment must be published as a Web service with the output of a trained model. This enables API access to the model for retraining.

The following steps apply to both New and Classic Web services:

Create the initial Predictive Web service:

- Create a training experiment
- Create a predictive web experiment
- Deploy a predictive web service

Retrain the Web service:

- Update training experiment to allow for retraining
- Deploy the retraining web service
- Use the Batch Execution Service code to retrain the model

For a walkthrough of the preceding steps, see [Retrain Machine Learning models programmatically](#).

NOTE

To deploy a New web service you must have sufficient permissions in the subscription to which you deploying the web service. For more information see, [Manage a Web service using the Azure Machine Learning Web Services portal](#).

If you deployed a Classic Web Service:

- Create a new Endpoint on the Predictive Web service
- Get the PATCH URL and code
- Use the PATCH URL to point the new Endpoint at the retrained model

For a walkthrough of the preceding steps, see [Retrain a Classic Web service](#).

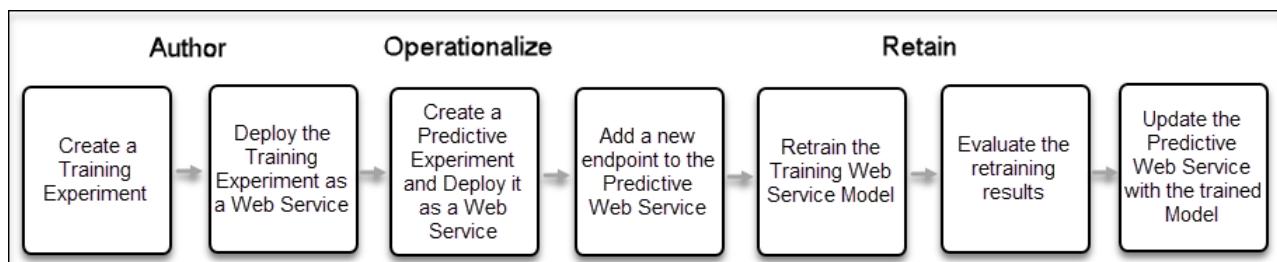
If you run into difficulties retraining a Classic Web service, see [Troubleshooting the retraining of an Azure Machine Learning Classic Web service](#).

If you deployed a New Web service:

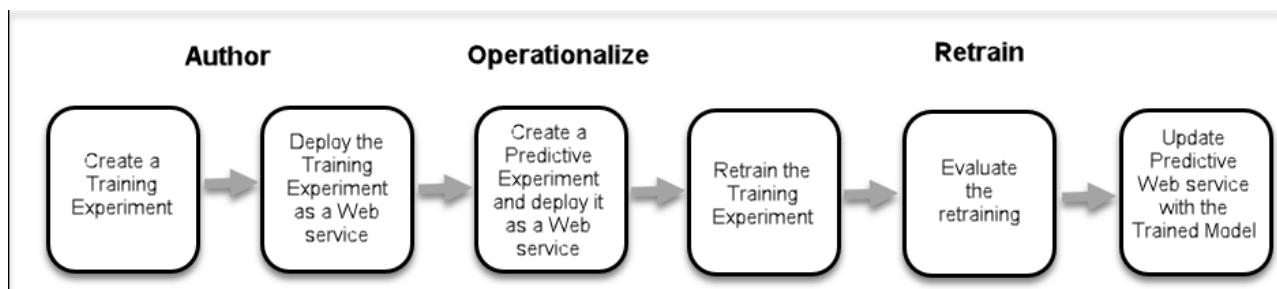
- Sign in to your Azure Resource Manager account
- Get the Web service definition
- Export the Web Service Definition as JSON
- Update the reference to the `ilearner` blob in the JSON
- Import the JSON into a Web Service Definition
- Update the Web service with new Web Service Definition

For a walkthrough of the preceding steps, see [Retrain a New Web service using the Machine Learning Management PowerShell cmdlets](#).

The process for setting up retraining for a Classic Web service involves the following steps:



The process for setting up retraining for a New Web service involves the following steps:



Other Resources

- [Retraining and Updating Azure Machine Learning models with Azure Data Factory](#)
- [Create many Machine Learning models and web service endpoints from one experiment using PowerShell](#)
- The [AML Retraining Models Using APIs](#) video shows you how to retrain Machine Learning models created in Azure Machine Learning using the Retraining APIs and PowerShell.

Retrain Machine Learning models programmatically

9/25/2017 • 7 min to read • [Edit Online](#)

In this walkthrough, you will learn how to programmatically retrain an Azure Machine Learning Web Service using C# and the Machine Learning Batch Execution service.

Once you have retrained the model, the following walkthroughs show how to update the model in your predictive web service:

- If you deployed a Classic web service in the Machine Learning Web Services portal, see [Retrain a Classic web service](#).
- If you deployed a New web service, see [Retrain a New web service using the Machine Learning Management cmdlets](#).

For an overview of the retraining process, see [Retrain a Machine Learning Model](#).

If you want to start with your existing New Azure Resource Manager based web service, see [Retrain an existing Predictive web service](#).

Create a training experiment

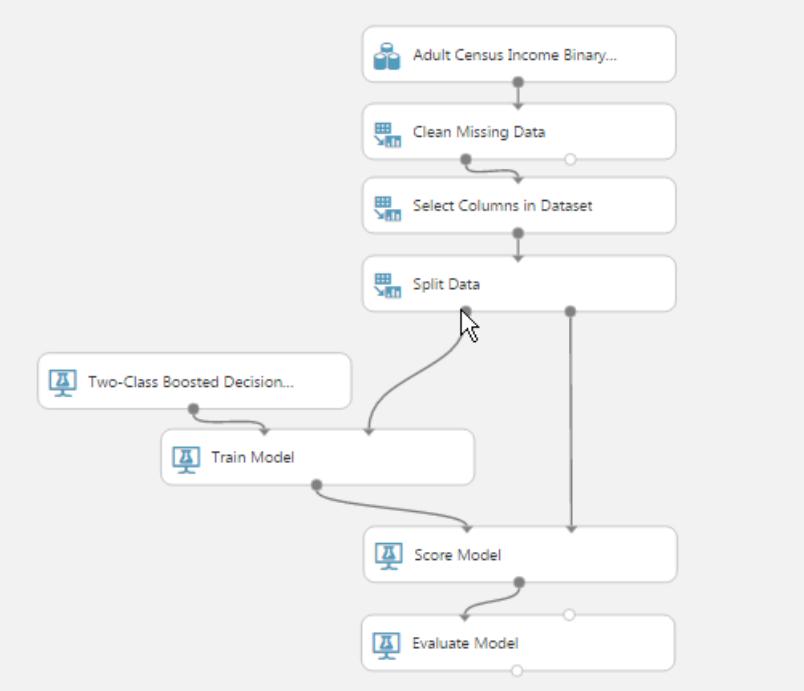
For this example, you will use "Sample 5: Train, Test, Evaluate for Binary Classification: Adult Dataset" from the Microsoft Azure Machine Learning samples.

To create the experiment:

1. Sign into to Microsoft Azure Machine Learning Studio.
2. On the bottom right corner of the dashboard, click **New**.
3. From the Microsoft Samples, select Sample 5.
4. To rename the experiment, at the top of the experiment canvas, select the experiment name "Sample 5: Train, Test, Evaluate for Binary Classification: Adult Dataset".
5. Type Census Model.
6. At the bottom of the experiment canvas, click **Run**.
7. Click **Set Up web service** and select **Retraining web service**.

The following shows the initial experiment.

Census Model



Create a predictive experiment and publish as a web service

Next you create a Predictive Experiment.

1. At the bottom of the experiment canvas, click **Set Up Web Service** and select **Predictive Web Service**. This saves the model as a Trained Model and adds web service Input and Output modules.
2. Click **Run**.
3. After the experiment has finished running, click **Deploy Web Service [Classic]** or **Deploy Web Service [New]**.

NOTE

To deploy a New web service you must have sufficient permissions in the subscription to which you deploying the web service. For more information see, [Manage a Web service using the Azure Machine Learning Web Services portal](#).

Deploy the training experiment as a Training web service

To retrain the trained model, you must deploy the training experiment that you created as a Retraining web service. This web service needs a *Web Service Output* module connected to the *Train Model* module, to be able to produce new trained models.

1. To return to the training experiment, click the Experiments icon in the left pane, then click the experiment named Census Model.
2. In the Search Experiment Items search box, type web service.
3. Drag a *Web Service Input* module onto the experiment canvas and connect its output to the *Clean Missing Data* module. This ensures that your retraining data is processed the same way as your original training data.
4. Drag two *web service Output* modules onto the experiment canvas. Connect the output of the *Train Model* module to one and the output of the *Evaluate Model* module to the other. The web service output for **Train Model** gives us the new trained model. The output attached to **Evaluate Model** returns that module's output, which is the performance results.
5. Click **Run**.

Next you must deploy the training experiment as a web service that produces a trained model and model evaluation results. To accomplish this, your next set of actions are dependent on whether you are working with a Classic web service or a New web service.

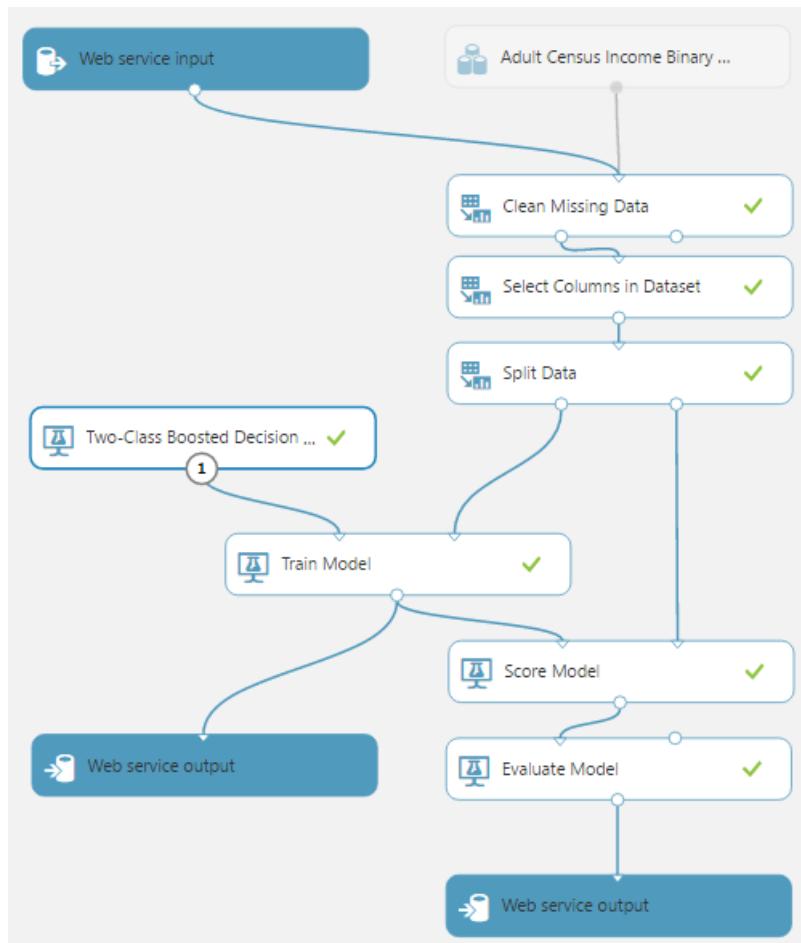
Classic web service

At the bottom of the experiment canvas, click **Set Up Web Service** and select **Deploy Web Service [Classic]**. The Web Service **Dashboard** is displayed with the API Key and the API help page for Batch Execution. Only the Batch Execution method can be used for creating Trained Models.

New web service

At the bottom of the experiment canvas, click **Set Up Web Service** and select **Deploy Web Service [New]**. The Web Service Azure Machine Learning Web Services portal opens to the Deploy web service page. Type a name for your web service and choose a payment plan, then click **Deploy**. Only the Batch Execution method can be used for creating Trained Models.

In either case, after experiment has completed running, the resulting workflow should look as follows:



Retrain the model with new data using BES

For this example, you are using C# to create the retraining application. You can also use the Python or R sample code to accomplish this task.

To call the Retraining APIs:

1. Create a C# console application in Visual Studio: **New > Project > Visual C# > Windows Classic Desktop > Console App (.NET Framework)**.
2. Sign in to the Machine Learning Web Service portal.
3. If you are working with a Classic web service, click **Classic Web Services**.
 - a. Click the web service you are working with.

- b. Click the default endpoint.
 - c. Click **Consume**.
 - d. At the bottom of the **Consume** page, in the **Sample Code** section, click **Batch**.
 - e. Continue to step 5 of this procedure.
4. If you are working with a New web service, click **Web Services**.
- a. Click the web service you are working with.
 - b. Click **Consume**.
 - c. At the bottom of the Consume page, in the **Sample Code** section, click **Batch**.
5. Copy the sample C# code for batch execution and paste it into the Program.cs file, making sure the namespace remains intact.

Add the Nuget package Microsoft.AspNet.WebApi.Client as specified in the comments. To add the reference to Microsoft.WindowsAzure.Storage.dll, you might first need to install the client library for Microsoft Azure storage services. For more information, see [Windows Storage Services](#).

Update the apikey declaration

Locate the **apikey** declaration.

```
const string apiKey = "abc123"; // Replace this with the API key for the web service
```

In the **Basic consumption info** section of the **Consume** page, locate the primary key and copy it to the **apikey** declaration.

Update the Azure Storage information

The BES sample code uploads a file from a local drive (For example "C:\temp\CensusIpnput.csv") to Azure Storage, processes it, and writes the results back to Azure Storage.

To accomplish this task, you must retrieve the Storage account name, key, and container information for your Storage account from the classic Azure portal and the update corresponding values in the code.

1. Sign in to the classic Azure portal.
2. In the left navigation column, click **Storage**.
3. From the list of storage accounts, select one to store the retrained model.
4. At the bottom of the page, click **Manage Access Keys**.
5. Copy and save the **Primary Access Key** and close the dialog.
6. At the top of the page, click **Containers**.
7. Select an existing container or create a new one and save the name.

Locate the *StorageAccountName*, *StorageAccountKey*, and *StorageContainerName* declarations and update the values you saved from the Azure portal.

```
const string StorageAccountName = "mystorageacct"; // Replace this with your Azure Storage Account name
const string StorageAccountKey = "a_storage_account_key"; // Replace this with your Azure Storage Key
const string StorageContainerName = "mycontainer"; // Replace this with your Azure Storage Container name
```

You also must ensure the input file is available at the location you specify in the code.

Specify the output location

When specifying the output location in the Request Payload, the extension of the file specified in *RelativeLocation* must be specified as ilearner.

See the following example:

```

Outputs = new Dictionary<string, AzureBlobDataReference>()
{
    "output1",
    new AzureBlobDataReference()
    {
        ConnectionString = storageConnectionString,
        RelativeLocation = string.Format("{0}/output1/results.learner", StorageContainerName) /*Replace this with the location you would like
        to use for your output file, and valid file extension (usually .csv for scoring results, or .learner for trained models)*/
    }
},

```

NOTE

The names of your output locations may be different from the ones in this walkthrough based on the order in which you added the web service output modules. Since you set up this training experiment with two outputs, the results include storage location information for both of them.

```

Running...
Finished!
The result 'output1' is available at the following Azure Storage location:
BaseLocation: https://esintussouthsus.blob.core.windows.net/
RelativeLocation: experimentoutput/9fcefacd-b375-45ab-bcc6-2bf69c491b0d/9fcefacd
-b375-45ab-bcc6-2bf69c491b0d.learner
SasBlobToken: ?sv=2013-08-15&sr=c&sig=%2BMREp%2Fo1%2Bbtv6MpQfm9HzsZ65HQTgd4DHpwh
WoJ1UdI%3D&st=2015-02-03T18%3A41%3A13Z&se=2015-02-04T18%3A46%3A13Z&sp=r1

The result 'output2' is available at the following Azure Storage location:
BaseLocation: https://esintussouthsus.blob.core.windows.net/
RelativeLocation: experimentoutput/1cea1d59-fa6a-45d5-83f5-b3ea60a8475b/1cea1d59
-fa6a-45d5-83f5-b3ea60a8475b.csv
SasBlobToken: ?sv=2013-08-15&sr=c&sig=%2BMREp%2Fo1%2Bbtv6MpQfm9HzsZ65HQTgd4DHpwh
WoJ1UdI%3D&st=2015-02-03T18%3A41%3A13Z&se=2015-02-04T18%3A46%3A13Z&sp=r1

```

Diagram 4: Retraining output.

Evaluate the Retraining Results

When you run the application, the output includes the URL and SAS token necessary to access the evaluation results.

You can see the performance results of the retrained model by combining the *BaseLocation*, *RelativeLocation*, and *SasBlobToken* from the output results for *output2* (as shown in the preceding retraining output image) and pasting the complete URL in the browser address bar.

Examine the results to determine whether the newly trained model performs well enough to replace the existing one.

Copy the *BaseLocation*, *RelativeLocation*, and *SasBlobToken* from the output results, you will use them during the retraining process.

Next steps

If you deployed the predictive web service by clicking **Deploy Web Service [Classic]**, see [Retrain a Classic web service](#).

If you deployed the predictive web service by clicking **Deploy Web Service [New]**, see [Retrain a New web service using the Machine Learning Management cmdlets](#).

Retrain a Classic web service

12/8/2017 • 3 min to read • [Edit Online](#)

The Predictive Web Service you deployed is the default scoring endpoint. Default endpoints are kept in sync with the original training and scoring experiments, and therefore the trained model for the default endpoint cannot be replaced. To retrain the web service, you must add a new endpoint to the web service.

Prerequisites

You must have set up a training experiment and a predictive experiment as shown in [Retrain Machine Learning models programmatically](#).

IMPORTANT

The predictive experiment must be deployed as a Classic machine learning web service.

For additional information on Deploying web services, see [Deploy an Azure Machine Learning web service](#).

Add a new endpoint

The Predictive Web Service that you deployed contains a default scoring endpoint that is kept in sync with the original training and scoring experiments trained model. To update your web service to with a new trained model, you must create a new scoring endpoint.

To create a new scoring endpoint, on the Predictive Web Service that can be updated with the trained model:

NOTE

Be sure you are adding the endpoint to the Predictive Web Service, not the Training Web Service. If you have correctly deployed both a Training and a Predictive Web Service, you should see two separate web services listed. The Predictive Web Service should end with "[predictive exp.]".

There are two ways in which you can add a new end point to a web service:

1. Programmatically
2. Use the Microsoft Azure Web Services portal

Programmatically add an endpoint

You can add scoring endpoints using the sample code provided in this [github repository](#).

Use the Microsoft Azure Web Services portal to add an endpoint

1. In Machine Learning Studio, on the left navigation column, click Web Services.
2. At the bottom of the web service dashboard, click **Manage endpoints preview**.
3. Click **Add**.
4. Type a name and description for the new endpoint. Select the logging level and whether sample data is enabled.
For more information on logging, see [Enable logging for Machine Learning web services](#).

Update the added endpoint's trained model

To complete the retraining process, you must update the trained model of the new endpoint that you added.

If you added the endpoint using the sample code, this includes location of the help URL identified by the *HelpLocationURL* value in the output.

To retrieve the path URL:

1. Copy and paste the URL into your browser.
2. Click the Update Resource link.
3. Copy the POST URL of the PATCH request. For example:

PATCH URL:

<https://management.azureml.net/workspaces/00bf70534500b34rebfa1843d6/webservices/af3er32ad393852f9b30ac9a35b/endpoints/newendpoint2>

You can now use the trained model to update the scoring endpoint that you created previously.

The following sample code shows you how to use the *BaseLocation*, *RelativeLocation*, *SasBlobToken*, and PATCH URL to update the endpoint.

```
private async Task OverwriteModel()
{
    var resourceLocations = new
    {
        Resources = new[]
        {
            new
            {
                Name = "Census Model [trained model]",
                Location = new AzureBlobDataReference()
                {
                    BaseLocation = "https://esintussouthsus.blob.core.windows.net/",
                    RelativeLocation = "your endpoint relative location", //from the output, for example: "experimentoutput/8946abfd-79d6-4438-89a9-3e5d109183/8946abfd-79d6-4438-89a9-3e5d109183.ileamer"
                    SasBlobToken = "your endpoint SAS blob token" //from the output, for example: "?sv=2013-08-15&sr=c&sig=37lTTfngRwxCcf94%3D&st=2015-01-30T22%3A53%3A06Z&se=2015-01-31T22%3A58%3A06Z&sp=r"
                }
            }
        };
    };

    using (var client = new HttpClient())
    {
        client.DefaultRequestHeaders.Authorization = new AuthenticationHeaderValue("Bearer", apiKey);

        using (var request = new HttpRequestMessage(new HttpMethod("PATCH"), endpointUrl))
        {
            request.Content = new StringContent(JsonConvert.SerializeObject(resourceLocations), System.Text.Encoding.UTF8, "application/json");
            HttpResponseMessage response = await client.SendAsync(request);

            if (!response.IsSuccessStatusCode)
            {
                await WriteFailedResponse(response);
            }

            // Do what you want with a successful response here.
        }
    }
}
```

The *apiKey* and the *endpointUrl* for the call can be obtained from endpoint dashboard.

The value of the *Name* parameter in *Resources* should match the Resource Name of the saved Trained Model in the predictive experiment. To get the Resource Name:

1. Sign in to the [classic Azure portal](#).
2. In the left menu, click **Machine Learning**.
3. Under Name, click your workspace and then click **Web Services**.
4. Under Name, click **Census Model [predictive exp.]**.
5. Click the new endpoint you added.
6. On the endpoint dashboard, click **Update Resource**.
7. On the Update Resource API Documentation page for the web service, you can find the **Resource Name** under **Updatable Resources**.

If your SAS token expires before you finish updating the endpoint, you must perform a GET with the Job Id to obtain a fresh token.

When the code has successfully run, the new endpoint should start using the retrained model in approximately 30 seconds.

Summary

Using the Retraining APIs, you can update the trained model of a predictive Web Service enabling scenarios such as:

- Periodic model retraining with new data.
- Distribution of a model to customers with the goal of letting them retrain the model using their own data.

Next steps

[Troubleshooting the retraining of an Azure Machine Learning classic web service](#)

Retrain a New Resource Manager based web service using the Machine Learning Management PowerShell cmdlets

12/8/2017 • 3 min to read • [Edit Online](#)

When you retrain a New web service, you update the predictive web service definition to reference the new trained model.

Prerequisites

You must set up a training experiment and a predictive experiment as shown in [Retrain Machine Learning models programmatically](#).

IMPORTANT

The predictive experiment must be deployed as an Azure Resource Manager (New) based machine learning web service. To deploy a New web service you must have sufficient permissions in the subscription to which you deploying the web service. For more information, see [Manage a Web service using the Azure Machine Learning Web Services portal](#).

For additional information on Deploying web services, see [Deploy an Azure Machine Learning web service](#).

This process requires that you have installed the Azure Machine Learning Cmdlets. For information installing the Machine Learning cmdlets, see the [Azure Machine Learning Cmdlets](#) reference on MSDN.

Copied the following information from the retraining output:

- BaseLocation
- RelativeLocation

The steps you take are:

1. Sign in to your Azure Resource Manager account.
2. Get the web service definition
3. Export the Web Service Definition as JSON
4. Update the reference to the ilearner blob in the JSON.
5. Import the JSON into a Web Service Definition
6. Update the web service with new Web Service Definition

Sign in to your Azure Resource Manager account

You must first sign in to your Azure account from within the PowerShell environment using the [Add-AzureRmAccount](#) cmdlet.

Get the Web Service Definition

Next, get the Web Service by calling the [Get-AzureRmMIWebService](#) cmdlet. The Web Service Definition is an internal representation of the trained model of the web service and is not directly modifiable. Make sure that you are retrieving the Web Service Definition for your Predictive experiment and not your training experiment.

```
$wsd = Get-AzureRmMIWebService -Name 'RetrainSamplePre.2016.8.17.0.3.51.237' -ResourceGroupName 'Default-MachineLearning-SouthCentralUS'
```

To determine the resource group name of an existing web service, run the `Get-AzureRmMIWebService` cmdlet without any parameters to display the web services in your subscription. Locate the web service, and then look at its web service ID. The name of the resource group is the fourth element in the ID, just after the *resourceGroups* element. In the following example, the resource group name is Default-MachineLearning-SouthCentralUS.

```
Properties : Microsoft.Azure.Management.MachineLearning.WebServices.Models.WebServicePropertiesForGraph
Id : /subscriptions/<subscription ID>/resourceGroups/Default-MachineLearning-SouthCentralUS/providers/Microsoft.MachineLearning/webServices/RetrainSamplePre.2016.8.17.0.3.51.237
Name : RetrainSamplePre.2016.8.17.0.3.51.237
Location : South Central US
Type : Microsoft.MachineLearning/webServices
Tags : {}
```

Alternatively, to determine the resource group name of an existing web service, log on to the Microsoft Azure Machine Learning Web Services portal. Select the web service. The resource group name is the fifth element of the URL of the web service, just after the *resourceGroups* element. In the following example, the resource group name is Default-MachineLearning-SouthCentralUS.

```
https://services.azureml.net/subscriptions/<subscription ID>/resourceGroups/Default-MachineLearning-SouthCentralUS/providers/Microsoft.MachineLearning/webServices/RetrainSamplePre.2016.8.17.0.3.51.237
```

Export the Web Service Definition as JSON

To modify the definition to the trained model to use the newly Trained Model, you must first use the [Export-AzureRmMIWebService](#) cmdlet to export it to a JSON format file.

```
Export-AzureRmMIWebService -WebService $wsd -OutputFile "C:\temp\mlservice_export.json"
```

Update the reference to the ilearner blob in the JSON.

In the assets, locate the [trained model], update the *uri* value in the *locationInfo* node with the URI of the ilearner blob. The URI is generated by combining the *BaseLocation* and the *RelativeLocation* from the output of the BES retraining call. This updates the path to reference the new trained model.

```
"asset3": {
  "name": "Retrain Sample [trained model]",
  "type": "Resource",
  "locationInfo": {
    "uri": "https://mltestaccount.blob.core.windows.net/azuremlassetscontainer/baca7bca650f46218633552c0bcbba0e.learner"
  },
  "outputPorts": {
    "Results dataset": {
      "type": "Dataset"
    }
  }
},
```

Import the JSON into a Web Service Definition

You must use the [Import-AzureRmMIWebService](#) cmdlet to convert the modified JSON file back into a Web

Service Definition that you can use to update the Web Service Definition.

```
$wsd = Import-AzureRmMIWebService -InputFile "C:\temp\mlservice_export.json"
```

Update the web service with new Web Service Definition

Finally, you use [Update-AzureRmMIWebService](#) cmdlet to update the Web Service Definition.

```
Update-AzureRmMIWebService -Name 'RetrainSamplePre.2016.8.17.0.3.51.237' -ResourceGroupName 'Default-MachineLearning-SouthCentralUS'  
-ServiceUpdates $wsd
```

Summary

Using the Machine Learning PowerShell management cmdlets, you can update the trained model of a predictive Web Service enabling scenarios such as:

- Periodic model retraining with new data.
- Distribution of a model to customers with the goal of letting them retrain the model using their own data.

Retrain an existing predictive web service

12/8/2017 • 7 min to read • [Edit Online](#)

This document describes the retraining process for the following scenario:

- You have a training experiment and a predictive experiment that you have deployed as an operationalized web service.
- You have new data that you want your predictive web service to use to perform its scoring.

NOTE

To deploy a New web service you must have sufficient permissions in the subscription to which you deploying the web service. For more information see, [Manage a Web service using the Azure Machine Learning Web Services portal](#).

Starting with your existing web service and experiments, you need to follow these steps:

1. Update the model.
 - a. Modify your training experiment to allow for web service inputs and outputs.
 - b. Deploy the training experiment as a retraining web service.
 - c. Use the training experiment's Batch Execution Service (BES) to retrain the model.
2. Use the Azure Machine Learning PowerShell cmdlets to update the predictive experiment.
 - a. Sign in to your Azure Resource Manager account.
 - b. Get the web service definition.
 - c. Export the web service definition as JSON.
 - d. Update the reference to the ilearner blob in the JSON.
 - e. Import the JSON into a web service definition.
 - f. Update the web service with a new web service definition.

Deploy the training experiment

To deploy the training experiment as a retraining web service, you must add web service inputs and outputs to the model. By connecting a *Web Service Output* module to the experiment's *Train Model* module, you enable the training experiment to produce a new trained model that you can use in your predictive experiment. If you have an *Evaluate Model* module, you can also attach web service output to get the evaluation results as output.

To update your training experiment:

1. Connect a *Web Service Input* module to your data input (for example, a *Clean Missing Data* module). Typically, you want to ensure that your input data is processed in the same way as your original training data.
2. Connect a *Web Service Output* module to the output of your *Train Model* module.
3. If you have an *Evaluate Model* module and you want to output the evaluation results, connect a *Web Service Output* module to the output of your *Evaluate Model* module.

Run your experiment.

Next, you must deploy the training experiment as a web service that produces a trained model and model evaluation results.

At the bottom of the experiment canvas, click **Set Up Web Service**, and then select **Deploy Web Service [New]**. The Azure Machine Learning Web Services portal opens to the **Deploy Web Service** page. Type a name for your

web service, choose a payment plan, and then click **Deploy**. You can only use the Batch Execution method for creating trained models.

Retrain the model with new data by using BES

For this example, we're using C# to create the retraining application. You can also use Python or R sample code to accomplish this task.

To call the retraining APIs:

1. Create a C# console application in Visual Studio: **New > Project > Visual C# > Windows Classic Desktop > Console App (.NET Framework)**.
2. Sign in to the Machine Learning Web Services portal.
3. Click the web service that you're working with.
4. Click **Consume**.
5. At the bottom of the **Consume** page, in the **Sample Code** section, click **Batch**.
6. Copy the sample C# code for batch execution and paste it into the Program.cs file. Make sure that the namespace remains intact.

Add the NuGet package Microsoft.AspNet.WebApi.Client, as specified in the comments. To add the reference to Microsoft.WindowsAzure.Storage.dll, you might first need to install the [client library for Azure Storage services](#).

The following screenshot shows the **Consume** page in the Azure Machine Learning Web Services portal.

The screenshot shows the 'Consume' tab selected in the Azure Machine Learning Web Services portal. The main content area displays the 'Retraining Example [Predictive Exp.]' experiment details. It includes sections for 'Web service consumption options' (Excel 2013 or later and Excel 2010 or earlier), 'Basic consumption info' (Primary Key and Secondary Key fields, both circled in red), 'Request-Response' (URL: https://ussouthcentral.services.azureml.net/subscriptions/789bd9c77b03463c8a5474a6ad134e35/services/94c201566e074c109bb554066b5c0f3a/execute?api-version=2.0&format=swagger, Documentation link), 'Batch Requests' (URL: https://ussouthcentral.services.azureml.net/subscriptions/789bd9c77b03463c8a5474a6ad134e35/services/94c201566e074c109bb554066b5c0f3a/jobs?api-version=2.0, Documentation link), and 'Sample Code' (Request-Response tab, Batch tab circled in red, C#, Python, Python 3+, R buttons). The Sample Code section contains C# code for batch execution, which requires the Microsoft.AspNet.WebApi.Client NuGet package.

```
// This code requires the Nuget package Microsoft.AspNet.WebApi.Client to be installed.  
// Instructions for doing this in Visual Studio:  
// Tools -> Nuget Package Manager -> Package Manager Console  
// Install-Package Microsoft.AspNet.WebApi.Client  
//  
// Also, add a reference to Microsoft.WindowsAzure.Storage.dll for reading from and writing to the Azure blob storage
```

Update the apikey declaration

Locate the **apikey** declaration:

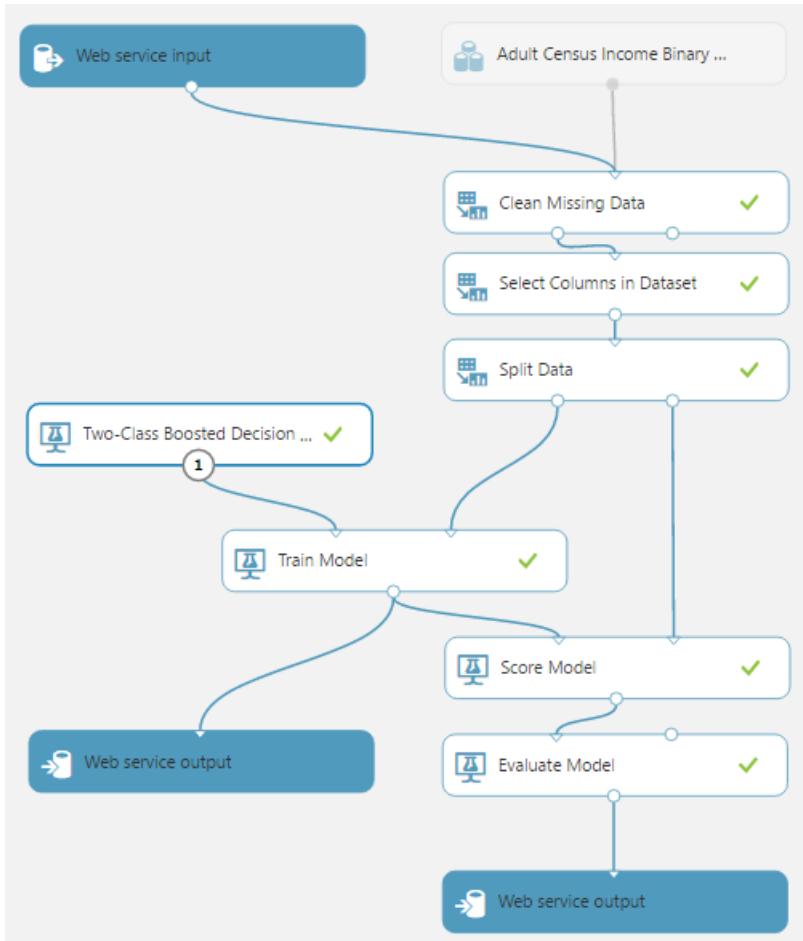
```
const string apiKey = "abc123"; // Replace this with the API key for the web service
```

In the **Basic consumption info** section of the **Consume** page, locate the primary key and copy it to the **apikey** declaration.

Update the Azure Storage information

The BES sample code uploads a file from a local drive (for example, "C:\temp\Census\pnput.csv") to Azure Storage, processes it, and writes the results back to Azure Storage.

After running your experiment, the resulting workflow should be similar to the following:



1. Sign in to the Azure portal.
2. In the left navigation column, click **More services**, search for **Storage accounts**, and select it.
3. From the list of storage accounts, select one to store the retrained model.
4. In the left navigation column, click **Access keys**.
5. Copy and save the **Primary Access Key**.
6. In the left navigation column, click **Containers**.
7. Select an existing container, or create a new one and save the name.

Locate the *StorageAccountName*, *StorageAccountKey*, and *StorageContainerName* declarations, and update the values that you saved from the portal.

```
const string StorageAccountName = "mystorageacct"; // Replace this with your Azure storage account name  
const string StorageAccountKey = "a_storage_account_key"; // Replace this with your Azure Storage key  
const string StorageContainerName = "mycontainer"; // Replace this with your Azure Storage container name
```

You also must ensure that the input file is available at the location that you specify in the code.

Specify the output location

When you specify the output location in the Request Payload, the extension of the file that is specified in *RelativeLocation* must be specified as `.ilearner`. See the following example:

```
Outputs = new Dictionary<string, AzureBlobDataReference>() {
{
    "output1",
    new AzureBlobDataReference()
    {
        ConnectionString = storageConnectionString,
        RelativeLocation = string.Format("{0}/output1/results.ilearner", StorageContainerName) /*Replace this with the location you want to use
for your output file and a valid file extension (usually .csv for scoring results or .ilearner for trained models)*/
    }
},
```

The following is an example of retraining output:

```
Running...
Finished!
The result 'output1' is available at the following Azure Storage location:
BaseLocation: https://esintussouthsus.blob.core.windows.net/
RelativeLocation: experimentoutput/9fcefacd-b375-45ab-bcc6-2bf69c491b0d/9fcefacd
-b375-45ab-bcc6-2bf69c491b0d.ilearner
SasBlobToken: ?sv=2013-08-15&r=c&sig=%2BMREp%2Fo1%2Bbty6MpQfm9HzsZ65HQTgd4DHpwH
WoJ1UdI%3D&st=2015-02-03T18%3A41%3A13Z&se=2015-02-04T18%3A46%3A13Z&sp=r1

The result 'output2' is available at the following Azure Storage location:
BaseLocation: https://esintussouthsus.blob.core.windows.net/
RelativeLocation: 1ceaid59-fa6a-45d5-83f5-b3ea60a8475b/1ceaid59
-fa6a-45d5-83f5-b3ea60a8475b.csv
SasBlobToken: ?sv=2013-08-15&r=c&sig=%2BMREp%2Fo1%2Bbty6MpQfm9HzsZ65HQTgd4DHpwH
WoJ1UdI%3D&st=2015-02-03T18%3A41%3A13Z&se=2015-02-04T18%3A46%3A13Z&sp=r1
```

Evaluate the retraining results

When you run the application, the output includes the URL and shared access signatures token that are necessary to access the evaluation results.

You can see the performance results of the retrained model by combining the *BaseLocation*, *RelativeLocation*, and *SasBlobToken* from the output results for *output2* (as shown in the preceding retraining output image) and pasting the complete URL into the browser address bar.

Examine the results to determine whether the newly trained model performs well enough to replace the existing one.

Copy the *BaseLocation*, *RelativeLocation*, and *SasBlobToken* from the output results.

Retrain the web service

When you retrain a new web service, you update the predictive web service definition to reference the new trained model. The web service definition is an internal representation of the trained model of the web service and is not directly modifiable. Make sure that you are retrieving the web service definition for your predictive experiment and not your training experiment.

Sign in to Azure Resource Manager

You must first sign in to your Azure account from within the PowerShell environment by using the [Add-AzureRmAccount](#) cmdlet.

Get the Web Service Definition object

Next, get the Web Service Definition object by calling the [Get-AzureRmMIWebService](#) cmdlet.

```
$wsd = Get-AzureRmMIWebService -Name 'RetrainSamplePre.2016.8.17.0.3.51.237' -ResourceGroupName 'Default-MachineLearning-SouthCentralUS'
```

To determine the resource group name of an existing web service, run the `Get-AzureRmMIWebService` cmdlet without any parameters to display the web services in your subscription. Locate the web service, and then look at its web service ID. The name of the resource group is the fourth element in the ID, just after the `resourceGroups` element. In the following example, the resource group name is Default-MachineLearning-SouthCentralUS.

```
Properties : Microsoft.Azure.Management.MachineLearning.WebServices.Models.WebServicePropertiesForGraph
Id : /subscriptions/<subscription ID>/resourceGroups/Default-MachineLearning-SouthCentralUS/providers/Microsoft.MachineLearning/webServices/RetrainSamplePre.2016.8.17.0.3.51.237
Name : RetrainSamplePre.2016.8.17.0.3.51.237
Location : South Central US
Type : Microsoft.MachineLearning/webServices
Tags : {}
```

Alternatively, to determine the resource group name of an existing web service, sign in to the Azure Machine Learning Web Services portal. Select the web service. The resource group name is the fifth element of the URL of the web service, just after the `resourceGroups` element. In the following example, the resource group name is Default-MachineLearning-SouthCentralUS.

```
https://services.azureml.net/subscriptions/<subscription ID>/resourceGroups/Default-MachineLearning-SouthCentralUS/providers/Microsoft.MachineLearning/webServices/RetrainSamplePre.2016.8.17.0.3.51.237
```

Export the Web Service Definition object as JSON

To modify the definition of the trained model to use the newly trained model, you must first use the [Export-AzureRmMIWebService](#) cmdlet to export it to a JSON-format file.

```
Export-AzureRmMIWebService -WebService $wsd -OutputFile "C:\temp\mlservice_export.json"
```

Update the reference to the iLearner blob

In the assets, locate the [trained model], update the `uri` value in the `locationInfo` node with the URI of the iLearner blob. The URI is generated by combining the `BaseLocation` and the `RelativeLocation` from the output of the BES retraining call.

```
"asset3": {
  "name": "Retrain Sample [trained model]",
  "type": "Resource",
  "locationInfo": {
    "uri": "https://mltestaccount.blob.core.windows.net/azuremlassetscontainer/baca7bca650f46218633552c0bcbba0e.ilearner"
  },
  "outputPorts": {
    "Results dataset": {
      "type": "Dataset"
    }
  }
},
```

Import the JSON into a Web Service Definition object

You must use the [Import-AzureRmMIWebService](#) cmdlet to convert the modified JSON file back into a Web Service

Definition object that you can use to update the predicative experiment.

```
$wsd = Import-AzureRmMIWebService -InputFile "C:\temp\mlservice_export.json"
```

Update the web service

Finally, use the [Update-AzureRmMIWebService](#) cmdlet to update the predictive experiment.

```
Update-AzureRmMIWebService -Name 'RetrainSamplePre.2016.8.17.0.3.51.237' -ResourceGroupName 'Default-MachineLearning-SouthCentralUS'
```

Troubleshooting the retraining of an Azure Machine Learning Classic web service

11/30/2017 • 3 min to read • [Edit Online](#)

Retraining overview

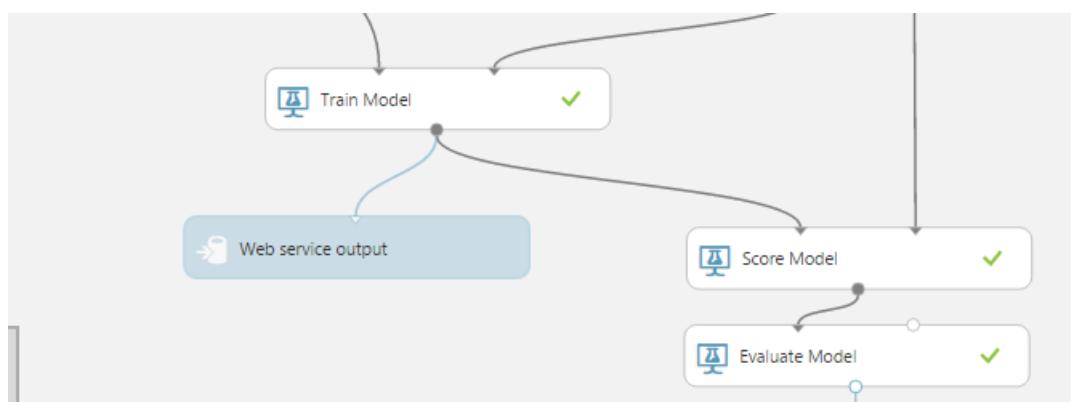
When you deploy a predictive experiment as a scoring web service it is a static model. As new data becomes available or when the consumer of the API has their own data, the model needs to be retrained.

For a complete walkthrough of the retraining process of a Classic web service, see [Retrain Machine Learning Models Programmatically](#).

Retraining process

When you need to retrain the web service, you must add some additional pieces:

- A web service deployed from the Training Experiment. The experiment must have a **Web Service Output** module attached to the output of the **Train Model** module.



- A new endpoint added to your scoring web service. You can add the endpoint programmatically using the sample code referenced in the Retrain Machine Learning models programmatically topic or through the Azure Machine Learning Web Services portal.

You can then use the sample C# code from the Training Web Service's API help page to retrain model. Once you have evaluated the results and are satisfied with them, you update the trained model scoring web service using the new endpoint that you added.

With all the pieces in place, the major steps you must take to retrain the model are as follows:

1. Call the Training Web Service: The call is to the Batch Execution Service (BES), not the Request Response Service (RRS). You can use the sample C# code on the API help page to make the call.
2. Find the values for the *BaseLocation*, *RelativeLocation*, and *SasBlobToken*: These values are returned in the output from your call to the Training Web Service.

```
Checking the job status...
Job af13598f64c34b1ea2520e0a0a28af26 running...
Checking the job status...
Job af13598f64c34b1ea2520e0a0a28af26 finished!
The result 'output2' is available at the following Azure Storage location:
BaseLocation: https://mltestaccount.blob.core.windows.net/
RelativeLocation: uploadedresources/output2results.ilearnr
SasBlobToken: ?sv=2015-02-21&sr=b&sig=JMbhu2g1ja947sJKUHxR%2FrcG86ABkManpn79418
%2Bcc%3D&st=2016-08-08T18%3A03%3A52Z&se=2016-08-09T18%3A08%3A52Z&sp=r
```

3. Update the added endpoint from the scoring web service with the new trained model: Using the sample code provided in the Retrain Machine Learning models programmatically, update the new endpoint you added to the

scoring model with the newly trained model from the Training Web Service.

Common obstacles

Check to see if you have the correct PATCH URL

The PATCH URL you are using must be the one associated with the new scoring endpoint you added to the scoring web service. There are a number of ways to obtain the PATCH URL:

Option 1: Programmatically

To get the correct PATCH URL:

1. Run the [AddEndpoint](#) sample code.
2. From the output of AddEndpoint, find the *HelpLocation* value and copy the URL.

```
Delete endpoint? y/n
y
Starting delete endpoint.

Endpoint deleted.
Create endpoint? y/n
y
Starting create endpoint
{
    "Name": "newendpoint2",
    "Description": "New end point",
    "CreationTime": "2016-08-05T21:51:28.363Z",
    "WorkspaceId": "8aa30f5a07e74e66bb767e25cc98876b",
    "WebServiceId": "710a704638de424191f9a0d0aa8ea48a",
    "HelpLocation": "https://studio.azureml.net/api/help/worksheets/8aa30f5a07e74e66bb767e25cc98876b/webservices/710a704638de424191f9a0d0aa8ea48a/endpoints/d8d040e7ae2044968178ea6ef149d923",
    "PrimaryKey": "25vzzUHphXVkmzPr2D77twUWEIKUEN8SIs9sAZCeq3P0Xrmc69+g6rX8KHzDgkpU0vI0qNPKYI9dSDLyIDXXYA==",
    "SecondaryKey": "F32EUhuvC7uwFZVDz9BXFGau8ZSb1mVxJ1V1PIoSuZi6uthhoDgI2AAAt0tCU9trLJtvZpSxZ39t51e38CmYT6g==",
    "ApiLocation": "https://ussouthcentral.services.azureml.net/worksheets/8aa30f5a07e74e66bb767e25cc98876b/services/d8d040e7ae2044968178ea6ef149d923",
    "ExperimentLocation": "https://studio.azureml.net/Home/ViewWorkspace/8aa30f5a07e74e66bb767e25cc98876b#Workspaces/Experiments/Experiment/8aa30f5a07e74e66bb767e25cc98876b.s-id.16104de92a164c70b35f67c610f9a32f.9953-06-04t06.17.00/ViewExperiment",
    "Resources": [
        {
            "Name": "Retrain Sample [trained model]",
            "Kind": "TrainedModel"
        }
    ],
    "Version": "2016-07-29T18:04:36.934Z",
    "PreventUpdate": false,
    "GlobalParameters": [],
    "MaxConcurrentCalls": 4,
    "DiagnosticsTraceLevel": "None",
    "ThrottleLevel": "Low"
}
Endpoint created.
```

3. Paste the URL into a browser to navigate to a page that provides help links for the web service.
4. Click the **Update Resource** link to open the patch help page.

Option 2: Use the Azure Machine Learning Web Services portal

1. Sign in to the [Azure Machine Learning Web Services](#) portal.
2. Click **Web Services** or **Classic Web Services** at the top.
3. Click the scoring web service you are working with (if you didn't modify the default name of the web service, it will end in "[Scoring Exp.]").
4. Click **+NEW**.
5. After the endpoint is added, click the endpoint name.
6. Under the **Patch** URL, click **API Help** to open the patching help page.

NOTE

If you have added the endpoint to the Training Web Service instead of the Predictive Web Service, you will receive the following error when you click the **Update Resource** link: "Sorry, but this feature is not supported or available in this context. This Web Service has no updatable resources. We apologize for the inconvenience and are working on improving this workflow."

The PATCH help page contains the PATCH URL you must use and provides sample code you can use to call it.

Request

Method	Request URI	HTTP Version
PATCH	<code>https://management.azureml.net/workspaces/c79283d685684d8488fc73fd3ec49a7f/webservices/0cf85bf373574e77a2918942de25beb9/endpoints/test</code>	HTTP/1.1

Check to see that you are updating the correct scoring endpoint

- Do not patch the training web service: The patch operation must be performed on the scoring web service.
- Do not patch the default endpoint on the web service: The patch operation must be performed on the new scoring web service endpoint that you added.

You can verify which web service the endpoint is on by visiting the Web Services portal.

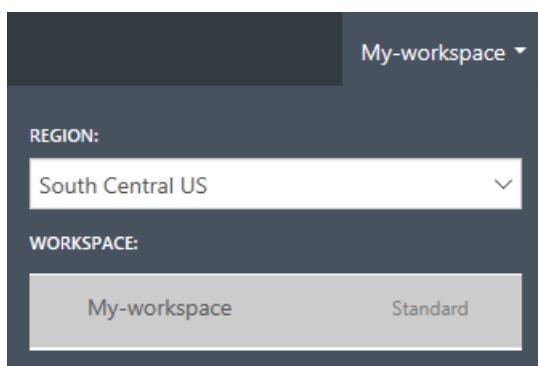
NOTE

Be sure you are adding the endpoint to the Predictive Web Service, not the Training Web Service. If you have correctly deployed both a Training and a Predictive Web Service, you should see two separate web services listed. The Predictive Web Service should end with "[predictive exp.]".

1. Sign in to the [Azure Machine Learning Web Services](#) portal.
2. Click **Web Services** or **Classic Web Services**.
3. Select your Predictive Web Service.
4. Verify that your new endpoint was added to the web service.

Check that your workspace is in the same region as the web service

1. Sign in to [Machine Learning Studio](#).
2. At the top, click the drop-down list of your workspaces.



3. Verify the region that your workspace is in.

How to consume an Azure Machine Learning Web service

11/30/2017 • 8 min to read • [Edit Online](#)

Once you deploy an Azure Machine Learning predictive model as a Web service, you can use a REST API to send it data and get predictions. You can send the data in real-time or in batch mode.

You can find more information about how to create and deploy a Machine Learning Web service using Machine Learning Studio here:

- For a tutorial on how to create an experiment in Machine Learning Studio, see [Create your first experiment](#).
- For details on how to deploy a Web service, see [Deploy a Machine Learning Web service](#).
- For more information about Machine Learning in general, visit the [Machine Learning Documentation Center](#).

NOTE

Try Azure Machine Learning for free

No credit card or Azure subscription required. [Get started now](#).

Overview

With the Azure Machine Learning Web service, an external application communicates with a Machine Learning workflow scoring model in real time. A Machine Learning Web service call returns prediction results to an external application. To make a Machine Learning Web service call, you pass an API key that is created when you deploy a prediction. The Machine Learning Web service is based on REST, a popular architecture choice for web programming projects.

Azure Machine Learning has two types of services:

- Request-Response Service (RRS) – A low latency, highly scalable service that provides an interface to the stateless models created and deployed from the Machine Learning Studio.
- Batch Execution Service (BES) – An asynchronous service that scores a batch for data records.

For more information about Machine Learning Web services, see [Deploy a Machine Learning Web service](#).

Get an Azure Machine Learning authorization key

When you deploy your experiment, API keys are generated for the Web service. You can retrieve the keys from several locations.

From the Microsoft Azure Machine Learning Web Services portal

Sign in to the [Microsoft Azure Machine Learning Web Services](#) portal.

To retrieve the API key for a New Machine Learning Web service:

1. In the Azure Machine Learning Web Services portal, click **Web Services** the top menu.
2. Click the Web service for which you want to retrieve the key.
3. On the top menu, click **Consume**.
4. Copy and save the **Primary Key**.

To retrieve the API key for a Classic Machine Learning Web service:

1. In the Azure Machine Learning Web Services portal, click **Classic Web Services** the top menu.
2. Click the Web service with which you are working.
3. Click the endpoint for which you want to retrieve the key.
4. On the top menu, click **Consume**.
5. Copy and save the **Primary Key**.

Classic Web service

You can also retrieve a key for a Classic Web service from Machine Learning Studio.

Machine Learning Studio

1. In Machine Learning Studio, click **WEB SERVICES** on the left.
2. Click a Web service. The **API key** is on the **DASHBOARD** tab.

Connect to a Machine Learning Web service

You can connect to a Machine Learning Web service using any programming language that supports HTTP request and response. You can view examples in C#, Python, and R from a Machine Learning Web service help page.

Machine Learning API help Machine Learning API help is created when you deploy a Web service. See [Azure Machine Learning Walkthrough- Deploy Web Service](#). The Machine Learning API help contains details about a prediction Web service.

1. Click the Web service with which you are working.
2. Click the endpoint for which you want to view the API Help Page.
3. On the top menu, click **Consume**.
4. Click **API help page** under either the Request-Response or Batch Execution endpoints.

To view Machine Learning API help for a New Web service

In the [Azure Machine Learning Web Services Portal](#):

1. Click **WEB SERVICES** on the top menu.
2. Click the Web service for which you want to retrieve the key.

Click **Use Web Service** to get the URIs for the Request-Reposonse and Batch Execution Services and Sample code in C#, R, and Python.

Click **Swagger API** to get Swagger based documentation for the APIs called from the supplied URIs.

C# Sample

To connect to a Machine Learning Web service, use an **HttpClient** passing ScoreData. ScoreData contains a FeatureVector, an n-dimensional vector of numerical features that represents the ScoreData. You authenticate to the Machine Learning service with an API key.

To connect to a Machine Learning Web service, the **Microsoft.AspNet.WebApi.Client** NuGet package must be installed.

Install Microsoft.AspNet.WebApi.Client NuGet in Visual Studio

1. Publish the Download dataset from UCI: Adult 2 class dataset Web Service.
2. Click **Tools > NuGet Package Manager > Package Manager Console**.
3. Choose **Install-Package Microsoft.AspNet.WebApi.Client**.

To run the code sample

1. Publish "Sample 1: Download dataset from UCI: Adult 2 class dataset" experiment, part of the Machine Learning sample collection.
2. Assign apiKey with the key from a Web service. See **Get an Azure Machine Learning authorization key** above.
3. Assign serviceUri with the Request URI.

Here is what a complete request will look like.

```

using System;
using System.Collections.Generic;
using System.IO;
using System.Net.Http;
using System.Net.Http.Formatting;
using System.Net.Http.Headers;
using System.Text;
using System.Threading.Tasks;

namespace CallRequestResponseService
{
    class Program
    {
        static void Main(string[] args)
        {
            InvokeRequestResponseService().Wait();
        }

        static async Task InvokeRequestResponseService()
        {
            using (var client = new HttpClient())
            {
                var scoreRequest = new
                {
                    Inputs = new Dictionary<string, List<Dictionary<string, string>>>() {
                        {
                            "input1",
                            // Replace columns labels with those used in your dataset
                            new List<Dictionary<string, string>>() {new Dictionary<string, string>() {
                                {
                                    "column1", "value1"
                                },
                                {
                                    "column2", "value2"
                                },
                                {
                                    "column3", "value3"
                                }
                            }
                        },
                        {
                            GlobalParameters = new Dictionary<string, string>()
                        }
                    };
                    // Replace these values with your API key and URI found on https://services.azureml.net/
                    const string apiKey = "<your-api-key>";
                    const string apiUri = "<your-api-uri>";

                    client.DefaultRequestHeaders.Authorization = new AuthenticationHeaderValue("Bearer", apiKey);
                    client.BaseAddress = new Uri(apiUri);

                    // WARNING: The 'await' statement below can result in a deadlock
                    // if you are calling this code from the UI thread of an ASP.NET application.
                    // One way to address this would be to call ConfigureAwait(false)
                    // so that the execution does not attempt to resume on the original context.
                    // For instance, replace code such as:
                    //     result = await DoSomeTask()
                }
            }
        }
    }
}

```

```
// with the following:  
//     result = await DoSomeTask().ConfigureAwait(false)  
  
HttpResponseMessage response = await client.PostAsJsonAsync("", scoreRequest);  
  
if(response.IsSuccessStatusCode)  
{  
    string result = await response.Content.ReadAsStringAsync();  
    Console.WriteLine("Result: {0}", result);  
}  
else  
{  
    Console.WriteLine(string.Format("The request failed with status code: {0}", response.StatusCode));  
  
    // Print the headers - they include the request ID and the timestamp,  
    // which are useful for debugging the failure  
    Console.WriteLine(response.Headers.ToString());  
  
    string responseContent = await response.Content.ReadAsStringAsync();  
    Console.WriteLine(responseContent);  
}  
}  
}  
}  
}
```

Python Sample

To connect to a Machine Learning Web service, use the **urllib2** library for Python 2.X and **urllib.request** library for Python 3.X. You will pass ScoreData, which contains a FeatureVector, an n-dimensional vector of numerical features that represents the ScoreData. You authenticate to the Machine Learning service with an API key.

To run the code sample

1. Deploy "Sample 1: Download dataset from UCI: Adult 2 class dataset" experiment, part of the Machine Learning sample collection.
2. Assign apiKey with the key from a Web service. See the [Get an Azure Machine Learning authorization key](#) section near the beginning of this article.
3. Assign serviceUri with the Request URI.

Here is what a complete request will look like.

```

import urllib2 # urllib.request for Python 3.X
import json

data = {
    "Inputs": {
        "input1": [
            {
                'column1': "value1",
                'column2': "value2",
                'column3': "value3"
            }
        ],
    },
    "GlobalParameters": {}
}

body = str.encode(json.dumps(data))

# Replace this with the URI and API Key for your web service
url='<your-api-uri>'
api_key='<your-api-key>'
headers = {'Content-Type':'application/json', 'Authorization':('Bearer '+ api_key)}

# "urllib.request.Request(uri, body, headers)" for Python 3.X
req = urllib2.Request(url, body, headers)

try:
    # "urllib.request.urlopen(req)" for Python 3.X
    response = urllib2.urlopen(req)

    result = response.read()
    print(result)
except urllib2.HTTPError as error:
    print("The request failed with status code: " + str(error.code))

    # Print the headers - they include the request ID and the timestamp, which are useful for debugging the failure
    print(error.info())
    print(json.loads(error.read()))

```

R Sample

To connect to a Machine Learning Web Service, use the **RCurl** and **rjson** libraries to make the request and process the returned JSON response. You will pass ScoreData, which contains a FeatureVector, an n-dimensional vector of numerical features that represents the ScoreData. You authenticate to the Machine Learning service with an API key.

Here is what a complete request will look like.

```

library("RCurl")
library("rjson")

# Accept SSL certificates issued by public Certificate Authorities
options(RCurlOptions = list(cainfo = system.file("CurlSSL", "cacert.pem", package = "RCurl")))

h = basicTextGatherer()
hdr = basicHeaderGatherer()

req = list(
  Inputs = list(
    "input1" = list(
      list(
        'column1' = "value1",
        'column2' = "value2",
        'column3' = "value3"
      )
    )
  ),
  GlobalParameters = setNames(fromJSON('{}'), character(0))
)

body = enc2utf8(toJSON(req))
api_key = "<your-api-key>" # Replace this with the API key for the web service
authz_hdr = paste('Bearer', api_key, sep='')

h$reset()
curlPerform(url = "<your-api-uri>",
httpheader=c('Content-Type' = "application/json", 'Authorization' = authz_hdr),
postfields=body,
writefunction = h$update,
headerfunction = hdr$update,
verbose = TRUE
)

headers = hdr$value()
HttpStatus = headers["status"]
if(HttpStatus >= 400)
{
  print(paste("The request failed with status code:", HttpStatus, sep=" "))
}

# Print the headers - they include the request ID and the timestamp, which are useful for debugging the failure
print(headers)
}

print("Result:")
result = h$value()
print(fromJSON(result))

```

JavaScript Sample

To connect to a Machine Learning Web Service, use the **request** npm package in your project. You will also use the `JSON` object to format your input and parse the result. Install using `npm install request --save`, or add `"request": "*"` to your `package.json` under `dependencies` and run `npm install`.

Here is what a complete request will look like.

```
let req = require("request");

const uri = "<your-api-uri>";
const apiKey = "<your-api-key>";

let data = {
  "Inputs": {
    "input1": [
      {
        'column1': "value1",
        'column2': "value2",
        'column3': "value3"
      }
    ],
    "GlobalParameters": {}
  }
}

const options = {
  uri,
  method: "POST",
  headers: {
    "Content-Type": "application/json",
    "Authorization": "Bearer " + apiKey,
  },
  body: JSON.stringify(data)
}

req(options, (err, res, body) => {
  if (!err && res.statusCode == 200) {
    console.log(body);
  } else {
    console.log("The request failed with status code: " + res.statusCode);
  }
});
```

Consuming an Azure Machine Learning Web Service from Excel

9/25/2017 • 1 min to read • [Edit Online](#)

Azure Machine Learning Studio makes it easy to call web services directly from Excel without the need to write any code.

If you are using Excel 2013 (or later) or Excel Online, then we recommend that you use the Excel [Excel add-in](#).

NOTE

Try Azure Machine Learning for free

No credit card or Azure subscription required. [Get started now](#).

Steps

Publish a web service. [This page](#) explains how to do it. Currently the Excel workbook feature is only supported for Request/Response services that have a single output (that is, a single scoring label).

Once you have a web service, click on the **WEB SERVICES** section on the left of the studio, and then select the web service to consume from Excel.

Classic Web Service

1. On the **DASHBOARD** tab for the web service is a row for the **REQUEST/RESPONSE** service. If this service had a single output, you should see the **Download Excel Workbook** link in that row.

Default Endpoint					
URL	TYPE	LAST UPDATED	TEST	APPS	LAST UPDATED
API help page	REQUEST/RESPONSE	2/6/2015 12:47:31 PM	Test	 Download Excel Workbook	2/6/2015 12:47:31 PM
API help page	BATCH EXECUTION	2/6/2015 12:47:31 PM			2/6/2015 12:47:31 PM

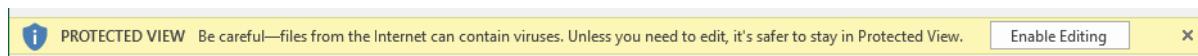
2. Click on **Download Excel Workbook**.

New Web Service

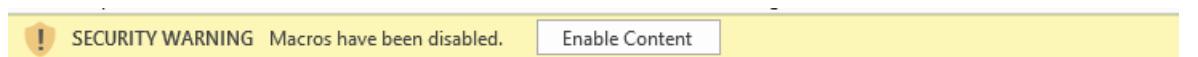
1. In the Azure Machine Learning Web Service portal, select **Consume**.
2. On the Consume page, in the **Web service consumption options** section, click the Excel icon.

Using the workbook

1. Open the workbook.
2. A Security Warning appears; click on the **Enable Editing** button.



3. A Security Warning appears. Click on the **Enable Content** button to run macros on your spreadsheet.



4. Once macros are enabled, a table is generated. Columns in blue are required as input into the RRS web

service, or **PARAMETERS**. Note the output of the RRS service, **PREDICTED VALUES** in green. When all columns for a given row are filled, the workbook automatically calls the scoring API, and displays the scored results.

PARAMETERS												PREDICTED VALUES	
age	workclass	education	education-num	marital-status	occupation	relationship	race	sex	hours-per-week	native-country	ScoredLabels	ScoredProbabilities	
0	0	0	0	0	0	0	0	0	0	0	0 <=50K	0.003267037	

5. To score more than one row, fill the second row with data and the predicted values are produced. You can even paste several rows at once.

You can use any of the Excel features (graphs, power map, conditional formatting, etc.) with the predicted values to help visualize the data.

Sharing your workbook

For the macros to work, your API Key must be part of the spreadsheet. That means that you should share the workbook only with entities/individuals you trust.

Automatic updates

An RRS call is made in these two situations:

1. The first time a row has content in all of its **PARAMETERS**
2. Any time any of the **PARAMETERS** changes in a row that had all of its **PARAMETERS** entered.

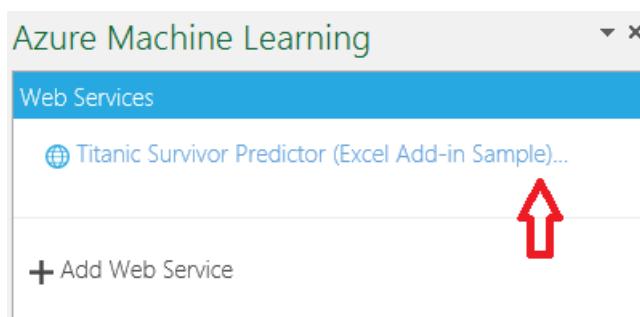
Excel Add-in for Azure Machine Learning web services

9/25/2017 • 2 min to read • [Edit Online](#)

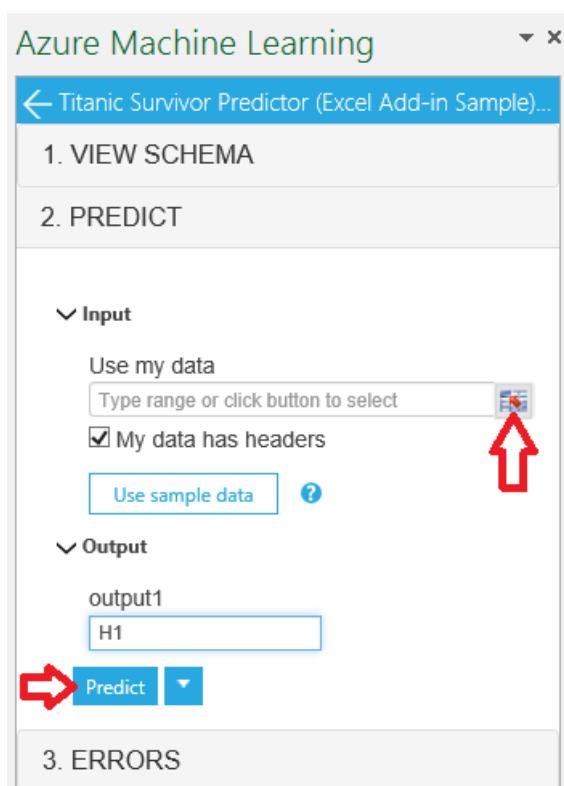
Excel makes it easy to call web services directly without the need to write any code.

Steps to Use an Existing web service in the Workbook

1. Open the [sample Excel file](#), which contains the Excel add-in and data about passengers on the Titanic.
2. Choose the web service by clicking it - "Titanic Survivor Predictor (Excel Add-in Sample) [Score]" in this example.



3. This takes you to the **Predict** section. This workbook already contains sample data, but for a blank workbook you can select a cell in Excel and click **Use sample data**.
4. Select the data with headers and click the input data range icon. Make sure the "My data has headers" box is checked.
5. Under **Output**, enter the cell number where you want the output to be, for example "H1" here.
6. Click **Predict**.

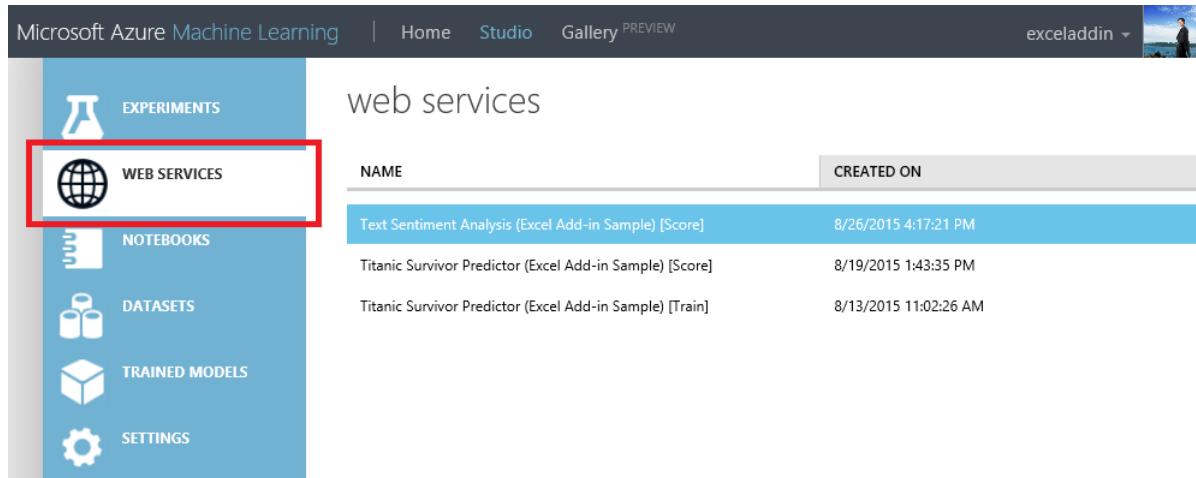


Deploy a web service or use an existing Web service. For more information on deploying a web service, see [Walkthrough Step 5: Deploy the Azure Machine Learning Web service](#).

Get the API key for your web service. Where you perform this action depends on whether you published a Classic Machine Learning web service or a New Machine Learning web service.

Use a Classic web service

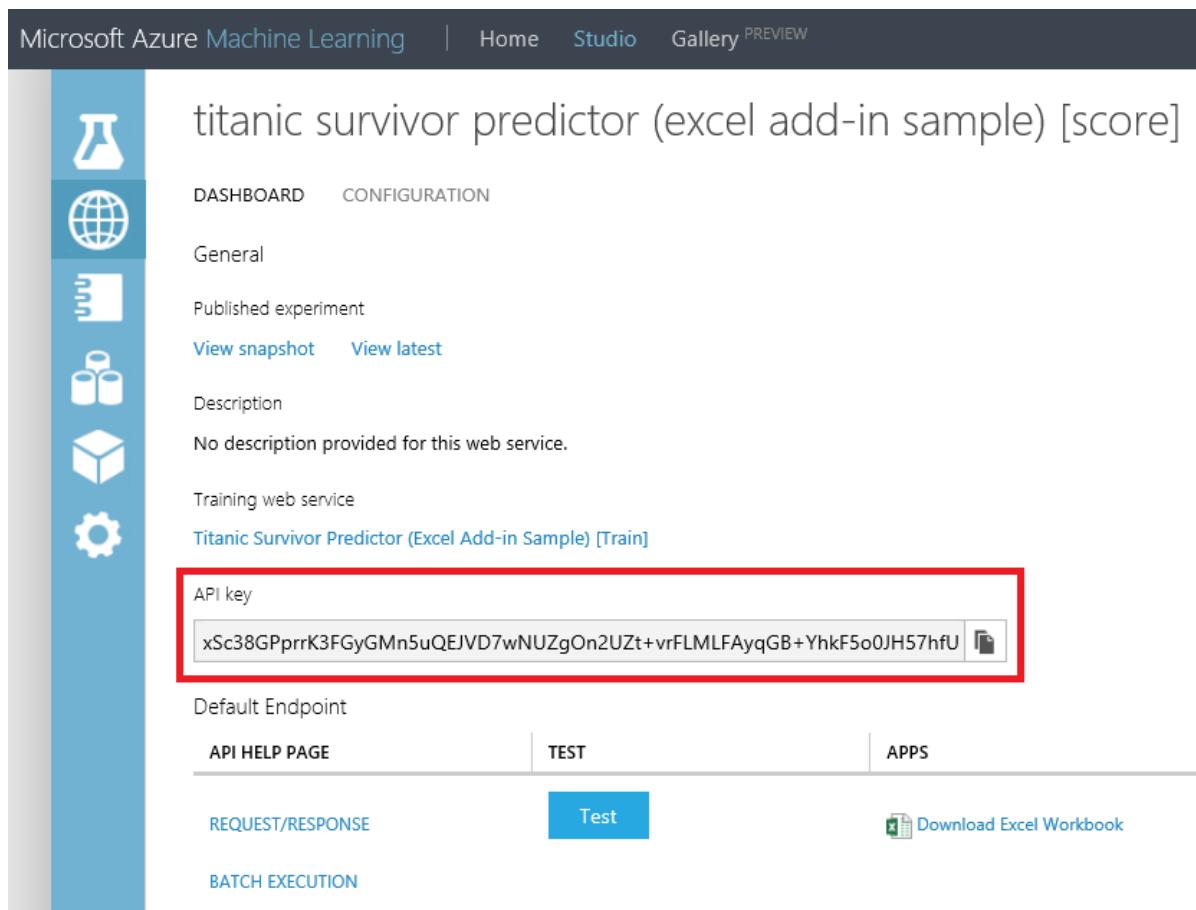
1. In Machine Learning Studio, click the **WEB SERVICES** section in the left pane, and then select the web service.



The screenshot shows the Microsoft Azure Machine Learning Studio interface. On the left, there is a vertical sidebar with icons for EXPERIMENTS, WEB SERVICES (which is highlighted with a red box), NOTEBOOKS, DATASETS, TRAINED MODELS, and SETTINGS. The main content area is titled "web services" and lists three entries:

NAME	CREATED ON
Text Sentiment Analysis (Excel Add-in Sample) [Score]	8/26/2015 4:17:21 PM
Titanic Survivor Predictor (Excel Add-in Sample) [Score]	8/19/2015 1:43:35 PM
Titanic Survivor Predictor (Excel Add-in Sample) [Train]	8/13/2015 11:02:26 AM

2. Copy the API key for the web service.



The screenshot shows the "titanic survivor predictor (excel add-in sample) [score]" web service dashboard. On the left, there is a vertical sidebar with icons for DASHBOARD (which is selected and highlighted with a red box), CONFIGURATION, General, Published experiment, View snapshot, View latest, Description, Training web service, and Titanic Survivor Predictor (Excel Add-in Sample) [Train]. The main content area has a section titled "API key" which contains the value "xSc38GPrrK3FGyGMn5uQEJVD7wNUZgOn2UZt+vrFLMLFAyqGB+YhkF5o0JH57hfU" and a copy icon. Below this, there is a "Default Endpoint" section and tabs for API HELP PAGE, TEST, and APPS. Under the TEST tab, there are REQUEST/RESPONSE and BATCH EXECUTION sections, along with a "Test" button and a "Download Excel Workbook" link.

3. On the **DASHBOARD** tab for the web service, click the **REQUEST/RESPONSE** link.
4. Look for the **Request URI** section. Copy and save the URL.

NOTE

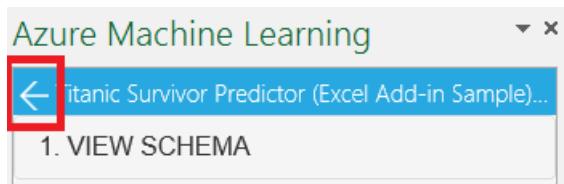
It is now possible to sign into the [Azure Machine Learning Web Services](#) portal to obtain the API key for a Classic Machine Learning web service.

Use a New web service

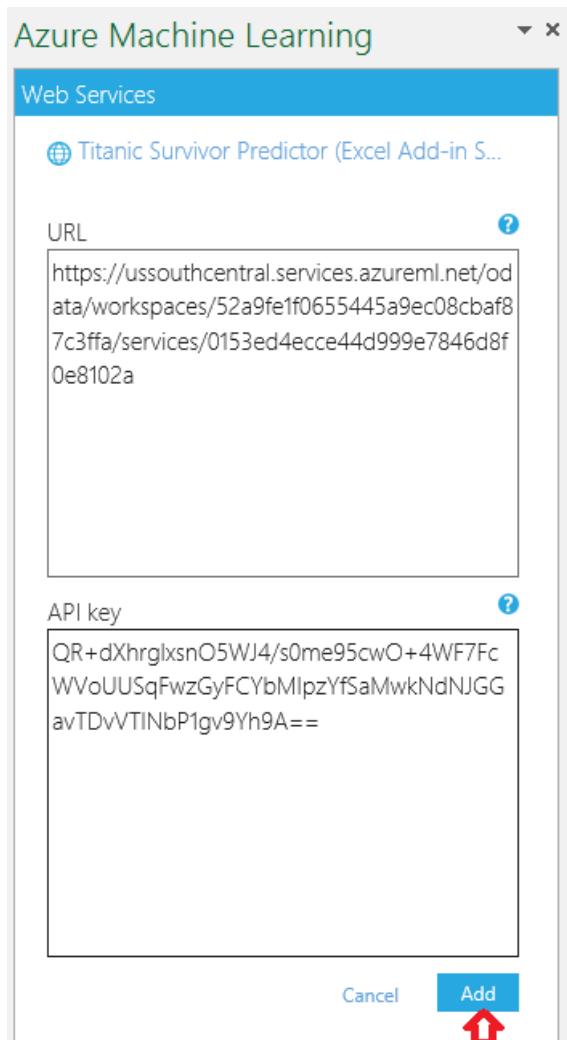
1. In the [Azure Machine Learning Web Services](#) portal, click **Web Services**, then select your web service.
2. Click **Consume**.
3. Look for the **Basic consumption info** section. Copy and save the **Primary Key** and the **Request-Response URL**.

Steps to Add a New web service

1. Deploy a web service or use an existing Web service. For more information on deploying a web service, see [Walkthrough Step 5: Deploy the Azure Machine Learning Web service](#).
2. Click **Consume**.
3. Look for the **Basic consumption info** section. Copy and save the **Primary Key** and the **Request-Response URL**.
4. In Excel, go to the **Web Services** section (if you are in the **Predict** section, click the back arrow to go to the list of web services).



5. Click **Add Web Service**.
6. Paste the URL into the Excel add-in text box labeled **URL**.
7. Paste the API/Primary key into the text box labeled **API key**.
8. Click **Add**.



9. To use the web service, follow the preceding directions, "Steps to Use an Existing web Service."

Sharing Your Workbook

If you save your workbook, then the API/Primary key for the web services you have added is also saved. That means you should only share the workbook with individuals you trust.

Ask any questions in the following comment section or on our [forum](#).

Consume an Azure Machine Learning web service with a web app template

9/25/2017 • 5 min to read • [Edit Online](#)

Once you've developed your predictive model and deployed it as an Azure web service using Machine Learning Studio, or using tools such as R or Python, you can access the operationalized model using a REST API.

There are a number of ways to consume the REST API and access the web service. For example, you can write an application in C#, R, or Python using the sample code generated for you when you deployed the web service (available in the [Machine Learning Web Services Portal](#) or in the web service dashboard in Machine Learning Studio). Or you can use the sample Microsoft Excel workbook created for you at the same time.

But the quickest and easiest way to access your web service is through the Web App Templates available in the [Azure Web App Marketplace](#).

NOTE

[Try Azure Machine Learning for free](#)

No credit card or Azure subscription required. [Get started now.](#)

The Azure Machine Learning Web App Templates

The web app templates available in the Azure Marketplace can build a custom web app that knows your web service's input data and expected results. All you need to do is give the web app access to your web service and data, and the template does the rest.

Two templates are available:

- [Azure ML Request-Response Service Web App Template](#)
- [Azure ML Batch Execution Service Web App Template](#)

Each template creates a sample ASP.NET application, using the API URI and Key for your web service, and deploys it as a web site to Azure. The Request-Response Service (RRS) template creates a web app that allows you to send a single row of data to the web service to get a single result. The Batch Execution Service (BES) template creates a web app that allows you to send many rows of data to get multiple results.

No coding is required to use these templates. You just supply the API Key and URI, and the template builds the application for you.

To get the API key and Request URI for a web service:

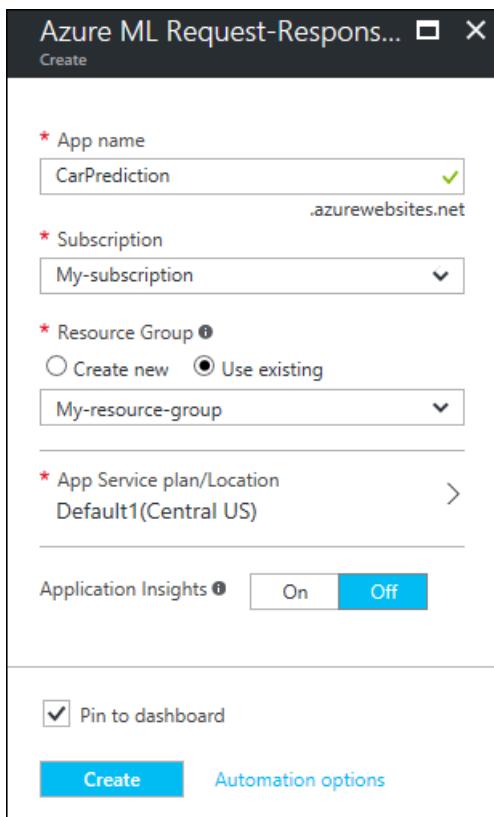
1. In the [Web Services Portal](#), for a New web service, click **Web Services** at the top. Or for a Classic web service click **Classic Web Services**.
2. Click the web service you want to access.
3. For a Classic web service, click the endpoint you want to access.
4. Click **Consume** at the top.
5. Copy the **Primary or Secondary Key** and save it.
6. If you're creating a Request-Response Service (RRS) template, copy the **Request-Response** URI and save it. If you're creating a Batch Execution Service (BES) template, copy the **Batch Requests** URI and save it.

How to use the Request-Response Service (RRS) template

Follow these steps to use the RRS web app template, as shown in the following diagram.



1. Go to the [Azure portal](#), **Login**, click **New**, search for and select **Azure ML Request-Response Service Web App**, then click **Create**.
 - Give your web app a unique name. The URL of the web app will be this name followed by `.azurewebsites.net`. For example, `http://carprediction.azurewebsites.net`.
 - Select the Azure subscription and services under which your web service is running.
 - Click **Create**.



- When Azure has finished deploying the web app, click the **URL** on the web app settings page in Azure, or enter the URL in a web browser. For example, <http://carprediction.azurewebsites.net>.
- When the web app first runs it will ask you for the **API Post URL** and **API Key**. Enter the values you saved earlier (**Request URI** and **API Key**, respectively).

Click **Submit**.

The dialog shows the 'Web Service Info' section with two fields: 'API Post URL:' containing the URL `https://ussouthcentral.services.azureml.net/workspaces/.../services/.../execute?api-version=2.0&details=true`, and 'API Key:' containing a redacted API key. A note at the bottom says 'Press Submit to load input parameters' next to a 'Submit' button.

- The web app displays its **Web App Configuration** page with the current web service settings. Here you can make changes to the settings used by the web app.

NOTE

Changing the settings here only changes them for this web app. It doesn't change the default settings of your web service. For example, if you change the **Description** here it doesn't change the description shown on the web service dashboard in Machine Learning Studio.

When you're done, click **Save changes**, and then click **Go to Home Page**.

5. From the home page you can enter values to send to your web service. Click **Submit** when you're done, and the result will be returned.

If you want to return to the **Configuration** page, go to the `setting.aspx` page of the web app. For example:

`http://carprediction.azurewebsites.net/setting.aspx` You will be prompted to enter the API key again - you need that to access the page and update the settings.

You can stop, restart, or delete the web app in the Azure portal like any other web app. As long as it is running you can browse to the home web address and enter new values.

How to use the Batch Execution Service (BES) template

You can use the BES web app template in the same way as the RRS template, except that the web app that's created will allow you to submit multiple rows of data and receive multiple results.

The input values for a batch execution web service can come from Azure storage or a local file; the results are stored in an Azure storage container. So, you'll need an Azure storage container to hold the results returned by the web app, and you'll need to get your input data ready.



1. Follow the same procedure to create the BES web app as for the RRS template, except go to [Azure ML Batch Execution Service Web App Template](#) to open the BES template on Azure Marketplace and click **Create Web App**.

2. To specify where you want the results stored, enter the destination container information on the web app home page. Also specify where the web app can get the input values, either in a local file or an Azure storage container. Click **Submit**.

The screenshot shows a web application titled "Car price prediction [Scoring Exp.]". At the top left is a logo of a head with gears. The main form is divided into sections:

- Azure Storage Info**: Fields for "Account Name" (with placeholder "account-name"), "Account Key" (with placeholder "*****"), and "Container Name" (with placeholder "container-name").
- Input File Info**: A radio button group where the first option, "Load from local machine", is selected. It includes a file input field showing "car-models.dataset" and a "Choose file" button. The second option, "Load from Azure Storage Blob", is shown with a "Same as above" link.
- Buttons at the bottom right: "Job Status" and "Submit".

At the bottom left of the form area, it says "Powered by Azure Machine Learning".

The web app will display a page with job status. When the job has completed you'll be given the location of the results in Azure blob storage. You also have the option of downloading the results to a local file.

For more information

To learn more about...

- creating a machine learning experiment with Machine Learning Studio, see [Create your first experiment in Azure Machine Learning Studio](#)
- how to deploy your machine learning experiment as a web service, see [Deploy an Azure Machine Learning web service](#)
- other ways to access your web service, see [How to consume an Azure Machine Learning Web service](#)

Azure Batch service for Machine Learning jobs

12/8/2017 • 4 min to read • [Edit Online](#)

Machine Learning Batch Pool processing provides customer-managed scale for the Azure Machine Learning Batch Execution Service. Classic batch processing for machine learning takes place in a multi-tenant environment, which limits the number of concurrent jobs you can submit, and jobs are queued on a first-in-first-out basis. This uncertainty means that you can't accurately predict when your job will run.

Batch Pool processing allows you to create pools on which you can submit batch jobs. You control the size of the pool, and to which pool the job is submitted. Your BES job runs in its own processing space providing predictable processing performance and the ability to create resource pools that correspond to the processing load that you submit.

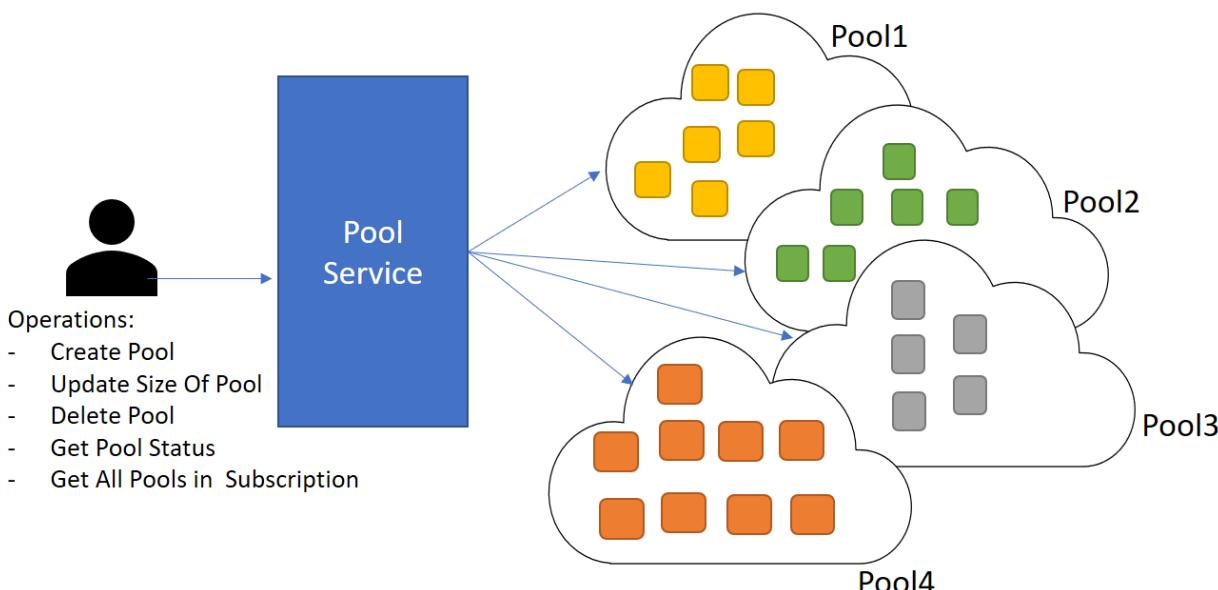
How to use Batch Pool processing

Configuration of Batch Pool Processing is not currently available through the Azure portal. To use Batch Pool processing, you must:

- Call CSS to create a Batch Pool Account and obtain a Pool Service URL and an authorization key
- Create a New Resource Manager based web service and billing plan

To create your account, call Microsoft Customer Service and Support (CSS) and provide your Subscription ID. CSS will work with you to determine the appropriate capacity for your scenario. CSS then configures your account with the maximum number of pools you can create and the maximum number of virtual machines (VMs) that you can place in each pool. Once your account is configured, you are provided your Pool Service URL and an authorization key.

After your account is created, you use the Pool Service URL and authorization key to perform pool management operations on your Batch Pool.



You create pools by calling the Create Pool operation on the pool service URL that CSS provided to you. When you create a pool, specify the number of VMs and the URL of the swagger.json of a New Resource Manager based Machine Learning web service. This web service is provided to establish the billing association. The Batch Pool service uses the swagger.json to associate the pool with a billing plan. You can run any BES web service, both New

Resource Manager based and classic, you choose on the pool.

You can use any New Resource Manager based web service, but be aware that the billing for the jobs are charged against the billing plan associated with that service. You may want to create a web service and new billing plan specifically for running Batch Pool jobs.

For more information on creating web services, see [Deploy an Azure Machine Learning web service](#).

Once you have created a pool, you submit the BES job using the Batch Requests URL for the web service. You can choose to submit it to a pool or to classic batch processing. To submit a job to Batch Pool processing, you add the following parameter to the job submission request body:

```
"AzureBatchPoolId": "<pool ID>"
```

If you do not add the parameter, the job is run in the classic batch process environment. If the pool has available resources, the job starts running immediately. If the pool does not have free resources, your job is queued until a resource is available.

If you find that you regularly reach the capacity of your pools, and you need increased capacity, you can call CSS and work with a representative to increase your quotas.

Example Request:

```
https://ussouthcentral.services.azureml.net/subscriptions/80c77c7674ba4c8c82294c3b2957990c/services/9fe659022c9747e3b9b7b923c3830623/jobs?api-version=2.0
```

```
{
  "Input":{

    "ConnectionString":"DefaultEndpointsProtocol=https;BlobEndpoint=https://sampleaccount.blob.core.windows.net;/TableEndpoint
    =https://sampleaccount.table.core.windows.net;/QueueEndpoint=https://sampleaccount.queue.core.windows.net;/FileEndpoint=https://zhuim
    lfile.core.windows.net;/AccountName=sampleaccount;AccountKey=&lt;Key&gt;;",

    "BaseLocation":null,
    "RelativeLocation":"testint/AdultCensusIncomeBinaryClassificationDataset.csv",
    "SasBlobToken":null
  },
  "GlobalParameters":{ },
  "Outputs":{

    "output1":{

      "ConnectionString":"DefaultEndpointsProtocol=https;BlobEndpoint=https://sampleaccount.blob.core.windows.net;/TableEndpo
      int=https://sampleaccount.table.core.windows.net;/QueueEndpoint=https://sampleaccount.queue.core.windows.net;/FileEndpoint=https://sample
      account.file.core.windows.net;/AccountName=sampleaccount;AccountKey=&lt;Key&gt;;",
      "BaseLocation":null,
      "RelativeLocation":"testintoutput/testint\_results.csv",
      "SasBlobToken":null
    }
  },
  "AzureBatchPoolId":"8dfc151b0d3e446497b845f3b29ef53b"
}
}
```

Considerations when using Batch Pool processing

Batch Pool Processing is an always-on billable service and that requires you to associate it with a Resource Manager based billing plan. You are only billed for the number of compute hours the pool is running; regardless of the number of jobs run during that time pool. If you create a pool, you are billed for the compute hours of each virtual machine in the pool until the pool is deleted, even if no batch jobs are running in the pool. Billing for the virtual machines starts when they have finished provisioning and stops when they have been deleted. You can use any of the plans found on the [Machine Learning Pricing page](#).

Billing example:

If you create a Batch Pool with 2 virtual machines and delete it after 24 hours your billing plan is debited 48 compute hours; regardless of how many jobs were run during that period.

If you create a Batch Pool with 4 virtual machines and delete it after 12 hours, your billing plan is also debited 48 compute hours.

We recommend that you poll the job status to determine when jobs complete. When all your jobs have finished running, call the Resize Pool operation to set the number of virtual machines in the pool to zero. If you are short on pool resources and you need to create a new pool, for example to bill against a different billing plan, you can delete the pool instead when all your jobs have finished running.

USE BATCH POOL PROCESSING WHEN

You need to run a large number of jobs
Or
You need to know that your jobs will run immediately
Or
You need guaranteed throughput. For example, you need to run a number of jobs in a given time frame, and want to scale out your compute resources to meet your needs.

USE CLASSIC BATCH PROCESSING WHEN

You are running just a few jobs
And
You don't need the jobs to run immediately

Copy example experiments to create new machine learning experiments

9/25/2017 • 2 min to read • [Edit Online](#)

Learn how to start with example experiments from [Cortana Intelligence Gallery](#) instead of creating machine learning experiments from scratch. You can use the examples to build your own machine learning solution.

The gallery has example experiments by the Microsoft Azure Machine Learning team as well as examples shared by the Machine Learning community. You also can ask questions or post comments about experiments.

To see how to use the gallery, watch the 3-minute video [Copy other people's work to do data science](#) from the series [Data Science for Beginners](#).

NOTE

Try Azure Machine Learning for free

No credit card or Azure subscription required. [Get started now](#).

Find an experiment to copy in Cortana Intelligence Gallery

To see what experiments are available, go to the [Gallery](#) and click **Experiments** at the top of the page.

Find the newest or most popular experiments

On this page, you can see **Recently added** experiments, or scroll down to look at **What's popular** or the latest **Popular Microsoft experiments**.

Look for an experiment that meets specific requirements

To browse all experiments:

1. Click **Browse all** at the top of the page.
2. On the left-hand side, under **Refine by** in the **Categories** section, select **Experiment** to see all the experiments in the Gallery.
3. You can find experiments that meet your requirements a couple different ways:
 - **Select filters on the left.** For example, to browse experiments that use a PCA-based anomaly detection algorithm: With **Experiment** selected under **Categories**, click **Show all**. Then, under **Algorithms Used**, choose **PCA-Based Anomaly Detection**.

Refine by

CATEGORIES

- Solution Template
- Experiment
- Machine Learning API
- Competition
- Tutorial
- Collection
- Notebook

SHOW

- Microsoft content only

TAGS

ALGORITHMS USED

- One-Class Support Vector Machine
- PCA-Based Anomaly Detection
- K-Means Clustering
- Multiclass Neural Network
- Bayesian Linear Regression

[Show all](#)

LANGUAGE

- **Use the search box.** For example, to find experiments contributed by Microsoft related to digit recognition that use a two-class support vector machine algorithm, enter "digit recognition" in the search box. Then, select the filters **Experiment**, **Microsoft content only**, and **Two-Class Support Vector Machine**:

Machine:

The screenshot shows the Cortana Intelligence Gallery interface. At the top, there's a search bar with 'digit recognition' typed in, and a magnifying glass icon. To the right of the search bar are user profile icons and a 'Sign in' link. Below the search bar is a navigation menu with 'Browse all' (which is underlined in green), 'Solution Templates', 'Experiments', 'Machine Learning APIs', and 'More'. On the left side, there's a 'Refine by' sidebar with sections for 'CATEGORIES', 'SHOW', 'TAGS', and 'ALGORITHMS USED'. Under 'CATEGORIES', 'Experiment' is checked. Under 'SHOW', 'Microsoft content only' is checked. Under 'ALGORITHMS USED', 'Two-Class Support Vector Machine' is checked. The main content area displays search results for 'digit recognition'. It says 'We've found 2 results for digit recognition' and 'Sort by: Relevance'. There are two experiment cards shown:

- Compare Multi-class Classifiers: Letter recogn...**: This card has a small image showing a table with 'Letter' and 'Probability' columns, listing 'a' (0.7), 'q' (0.2), and 'o' (0.1). Below the image is a brief description: 'This sample demonstrates how to compare multiple multi-class classifiers using the letter recogn...'. It includes a 'One-vs-All Multiclass, Two-Class Support Vector Machine, Multicla...' tag, 3909 views, 1195 likes, and was posted 9 months ago.
- Sample 7: Train, Test, Evaluate for Multiclass ...**: This card has an image showing three boxes labeled 'A', 'B', and 'C' with a green gear icon between them. Below the image is a brief description: 'Sample experiment that uses multiclass classification to predict the letter category as one of the...'. It includes a 'Multiclass Decision Jungle, One-vs-All Multiclass, Two-Class Support ...' tag, 2108 views, 2685 likes, and was posted 9 months ago.

 Both cards have a 'Microsoft' logo at the bottom.

4. Click an experiment to learn more about it.
5. To run and/or modify the experiment, click **Open in Studio** on the experiment's page.

EXPERIMENT

Compare Multi-class Classifiers: Letter recognition

Microsoft • published on September 2, 2014

Summary

This sample demonstrates how to compare multiple multi-class classifiers using the letter recognition dataset.

Description

Compare Multi-class Classifiers: Letter Recognition

This sample demonstrates how to create multiclass classifiers and evaluate and compare the performance of multiple models.

Data

For this experiment, we use the letter image recognition data from the [UCI repository](#). The first column is the label, which identifies each row as one of 26 letters, A-Z. The remaining 16 columns are feature columns. The dataset contains 20000 instances.

Description and other details about the data can be found at <http://archive.ics.uci.edu/ml/machine-learning-databases/letter-recognition/letter-recognition.names>.

Letter	Probability
a	0.7
q	0.2
o	0.1

[Open in Studio](#)

+ Add to Collection

3909 views

1195 downloads

[Tweet](#) [Share](#)

ALGORITHMS

[One-vs-All Multiclass](#) , [Two-Class Support Vector Machine](#) , [Multiclass Decision Forest](#) , [Multiclass Decision Jungle](#) , [Multiclass Logistic Regression](#) , [Multiclass Neural Network](#)

NOTE

When you open an experiment in Machine Learning Studio for the first time, you can try it for free or buy an Azure subscription. [Learn about the Machine Learning Studio free trial vs. paid service](#)

Create a new experiment using an example as a template

You also can create a new experiment in Machine Learning Studio using a Gallery example as a template.

1. Sign in with your Microsoft account credentials to the [Studio](#), and then click **New** to create an experiment.
2. Browse through the example content and click one.

A new experiment is created in your Machine Learning Studio workspace using the example experiment as a template.

Next steps

- [Import data from various sources](#)
- [Quickstart tutorial for the R language in Machine Learning](#)
- [Deploy a Machine Learning web service](#)

Use the sample datasets in Azure Machine Learning Studio

9/25/2017 • 15 min to read • [Edit Online](#)

When you create a new workspace in Azure Machine Learning, a number of sample datasets and experiments are included by default. Many of these sample datasets are used by the sample models in the [Azure Cortana Intelligence Gallery](#). Others are included as examples of various types of data typically used in machine learning.

Some of these datasets are available in Azure Blob storage. For these datasets, the following table provides a direct link. You can use these datasets in your experiments by using the [Import Data](#) module.

The rest of these sample datasets are available in your workspace under **Saved Datasets** in the module palette to the left of the experiment canvas when you open or create a new experiment in Machine Learning Studio. You can use any of these datasets in your own experiment by dragging it to your experiment canvas.

NOTE

Try Azure Machine Learning for free

No credit card or Azure subscription required. [Get started now.](#)

<th align=left>Dataset name <th align=left>Dataset description <td valign=top>Adult Census Income Binary Classification dataset <td valign=top> A subset of the 1994 Census database, using working adults over the age of 16 with an adjusted income index of > 100.

Usage: Classify people using demographics to predict whether a person earns over 50K a year.

Related Research: Kohavi, R., Becker, B., (1996). UCI Machine Learning Repository <http://archive.ics.uci.edu/ml>. Irvine, CA: University of California, School of Information and Computer Science <tr ID=airport-codes-dataset> <td valign=top>Airport Codes Dataset <td valign=top> U.S. airport codes.

This dataset contains one row for each U.S. airport, providing the airport ID number and name along with the location city and state. <td valign=top>Automobile price data (Raw) <td valign=top> Information about automobiles by make and model, including the price, features such as the number of cylinders and MPG, as well as an insurance risk score.

The risk score is initially associated with auto price and then adjusted for actual risk in a process known to actuaries as symboling. A value of +3 indicates that the auto is risky, and a value of -3 that it is probably safe.

Usage: Predict the risk score by features, using regression or multivariate classification.

Related Research: Schlimmer, J.C. (1987). UCI Machine Learning Repository <http://archive.ics.uci.edu/ml>. Irvine, CA: University of California, School of Information and Computer Science <tr ID=bike-rental-uci-dataset> <td valign=top>Bike Rental UCI dataset <td valign=top> UCI Bike Rental dataset that is based on real data from Capital Bikeshare company that maintains a bike rental network in Washington DC.

The dataset has one row for each hour of each day in 2011 and 2012, for a total of 17,379 rows. The range of hourly bike rentals is from 1 to 977. <tr ID=bill-gates-rgb-image> <td valign=top>Bill Gates RGB Image <td valign=top> Publicly available image file converted to CSV data.

The code for converting the image is provided in the [Color quantization using K-Means clustering](#) model detail page. <td valign=top>Blood donation data <td valign=top> A subset of data from the blood donor database of the Blood Transfusion Service Center of Hsin-Chu City, Taiwan.

Donor data includes the months since last donation), and frequency, or the total number of donations, time since last donation, and amount of blood donated.

Usage: The goal is to predict via classification whether the donor donated blood in March 2007, where 1 indicates a donor during the target period, and 0 a non-donor.

Related Research: Yeh, I.C., (2008). UCI Machine Learning Repository <http://archive.ics.uci.edu/ml>. Irvine, CA: University of California, School of Information and Computer Science

Yeh, I-Cheng, Yang, King-Jang, and Ting, Tao-Ming, "Knowledge discovery on RFM model using Bernoulli sequence, "Expert Systems with Applications, 2008, <http://dx.doi.org/10.1016/j.eswa.2008.07.018> <tr ID=book-reviews-from-amazon> <td valign=top> Book Reviews from Amazon <td valign=top> Reviews of books in Amazon, taken from the amazon.com website by University of Pennsylvania researchers ([sentiment](#)). See the research paper, "Biographies, Bollywood, Boom-boxes and Blenders: Domain Adaptation for Sentiment Classification" by John Blitzer, Mark Dredze, and Fernando Pereira; Association of Computational Linguistics (ACL), 2007.

The original dataset has 975K reviews with rankings 1, 2, 3, 4, or 5. The reviews were written in English and are from the time period 1997-2007. This dataset has been down-sampled to 10K reviews. <td valign=top> Breast cancer data <td valign=top> One of three cancer-related datasets provided by the Oncology Institute that appears frequently in machine learning literature. Combines diagnostic information with features from laboratory analysis of about 300 tissue samples.

Usage: Classify the type of cancer, based on 9 attributes, some of which are linear and some are categorical.

Related Research: Wohlberg, W.H., Street, W.N., & Mangasarian, O.L. (1995). UCI Machine Learning Repository <http://archive.ics.uci.edu/ml>. Irvine, CA: University of California, School of Information and Computer Science <tr ID=breast-cancer-features> <td valign=top> Breast Cancer Features <td valign=top> The dataset contains information for 102K suspicious regions (candidates) of X-ray images, each described by 117 features. The features are proprietary and their meaning is not revealed by the dataset creators (Siemens Healthcare). <tr ID=breast-cancer-info> <td valign=top> Breast Cancer Info <td valign=top> The dataset contains additional information for each suspicious region of X-ray image. Each example provides information (e.g., label, patient ID, coordinates of patch relative to the whole image) about the corresponding row number in the Breast Cancer Features dataset. Each patient has a number of examples. For patients who have a cancer, some examples are positive and some are negative. For patients who don't have a cancer, all examples are negative. The dataset has 102K examples. The dataset is biased, 0.6% of the points are positive, the rest are negative. The dataset was made available by Siemens Healthcare. <tr ID=crm-appetency-labels-shared> <td valign=top> CRM Appetency Labels Shared <td valign=top> Labels from the KDD Cup 2009 customer relationship prediction challenge ([orange_small_train_appetency.labels](#)). <tr ID=crm-churn-labels-shared> <td valign=top> CRM Churn Labels Shared <td valign=top> Labels from the KDD Cup 2009 customer relationship prediction challenge ([orange_small_train_churn.labels](#)). <tr ID=crm-dataset-shared> <td valign=top> CRM Dataset Shared <td valign=top> This data comes from the KDD Cup 2009 customer relationship prediction challenge ([orange_small_train.data.zip](#)).

The dataset contains 50K customers from the French Telecom company Orange. Each customer has 230 anonymized features, 190 of which are numeric and 40 are categorical. The features are very sparse. <tr ID=crm-upselling-labels-shared> <td valign=top> CRM Upselling Labels Shared <td valign=top> Labels from the KDD Cup 2009 customer relationship prediction challenge ([orange_large_train_upselling.labels](#)). <td valign=top> Energy Efficiency Regression data <td valign=top> A collection of simulated energy profiles, based on 12 different building shapes. The buildings are differentiated by 8 features, such as glazing area, the glazing area distribution, and orientation.

Usage: Use either regression or classification to predict the energy-efficiency rating based as one of two real valued responses. For multi-class classification, is round the response variable to the nearest integer.

Related Research: Xifara, A. & Tsanas, A. (2012). UCI Machine Learning Repository <http://archive.ics.uci.edu/ml>. Irvine, CA: University of California, School of Information and Computer Science <tr ID=flight-delays-data> <td valign=top> Flight Delays Data <td valign=top> Passenger flight on-time performance data taken from the TranStats data collection of the U.S. Department of Transportation ([On-Time](#)).

The dataset covers the time period April–October 2013. Before uploading to Azure Machine Learning Studio, the dataset was processed as follows:

- The dataset was filtered to cover only the 70 busiest airports in the continental US
- Canceled flights were labeled as delayed by more than 15 minutes
- Diverted flights were filtered out
- The following columns were selected: Year, Month, DayofMonth, DayOfWeek, Carrier, OriginAirportID, DestAirportID, CRSDepTime, DepDelay, DepDel15, CRSArrTime, ArrDelay, ArrDel15, Canceled

<td valign=top>Flight on-time performance (Raw) <td valign=top> Records of airplane flight arrivals and departures within United States from October 2011.

Usage: Predict flight delays.

Related Research: From US Dept. of Transportation http://www.transtats.bts.gov/DL_SelectFields.asp

Table_ID=236&DB_Short_Name=On-Time. <td valign=top>Forest fires data <td valign=top> Contains weather data, such as temperature and humidity indices and wind speed, from an area of northeast Portugal, combined with records of forest fires.

Usage: This is a difficult regression task, where the aim is to predict the burned area of forest fires.

Related Research: Cortez, P., & Morais, A. (2008). UCI Machine Learning Repository <http://archive.ics.uci.edu/ml>. Irvine, CA: University of California, School of Information and Computer Science

[Cortez and Morais, 2007] P. Cortez and A. Morais. A Data Mining Approach to Predict Forest Fires using Meteorological Data. In J. Neves, M. F. Santos and J. Machado Eds., New Trends in Artificial Intelligence, Proceedings of the 13th EPIA 2007 - Portuguese Conference on Artificial Intelligence, December, Guimarães, Portugal, pp. 512-523, 2007. APPIA, ISBN-13 978-989-95618-0-9. Available at: <http://www.dsi.uminho.pt/~pcortez/fires.pdf>. <tr ID=german-credit-card-uci-dataset> <td valign=top>German Credit Card UCI dataset <td valign=top> The UCI Statlog (German Credit Card) dataset ([Statlog+German+Credit+Data](#)), using the german.data file.

The dataset classifies people, described by a set of attributes, as low or high credit risks. Each example represents a person. There are 20 features, both numerical and categorical, and a binary label (the credit risk value). High credit risk entries have label = 2, low credit risk entries have label = 1. The cost of misclassifying a low risk example as high is 1, whereas the cost of misclassifying a high risk example as low is 5. <tr ID=imdb-movie-titles> <td valign=top>IMDB Movie Titles <td valign=top> The dataset contains information about movies that were rated in Twitter tweets: IMDB movie ID, movie name, genre, and production year. There are 17K movies in the dataset. The dataset was introduced in the paper "S. Dooms, T. De Pessemier and L. Martens. MovieTweetings: a Movie Rating Dataset Collected From Twitter. Workshop on Crowdsourcing and Human Computation for Recommender Systems, CrowdRec at RecSys 2013." <td valign=top>Iris two class data <td valign=top> This is perhaps the best known database to be found in the pattern recognition literature. The dataset is relatively small, containing 50 examples each of petal measurements from three iris varieties.

Usage: Predict the iris type from the measurements.

Related Research: Fisher, R.A. (1988). UCI Machine Learning Repository <http://archive.ics.uci.edu/ml>. Irvine, CA: University of California, School of Information and Computer Science <tr ID=movie-tweets> <td valign=top>Movie Tweets <td valign=top> The dataset is an extended version of the Movie Tweetings dataset. The dataset has 170K ratings for movies, extracted from well-structured tweets on Twitter. Each instance represents a tweet and is a tuple: user ID, IMDB movie ID, rating, timestamp, number of favorites for this tweet, and number of retweets of this tweet. The dataset was made available by A. Said, S. Dooms, B. Loni and D. Tikk for Recommender Systems Challenge 2014. <td valign=top>MPG data for various automobiles <td valign=top> This dataset is a slightly modified version of the dataset provided by the StatLib library of Carnegie Mellon University. The dataset was used in the 1983 American Statistical Association Exposition.

The data lists fuel consumption for various automobiles in miles per gallon, along with information such as the number of cylinders, engine displacement, horsepower, total weight, and acceleration.

Usage: Predict fuel economy based on 3 multivalued discrete attributes and 5 continuous attributes.

Related Research: StatLib, Carnegie Mellon University, (1993). UCI Machine Learning Repository

<http://archive.ics.uci.edu/ml>. Irvine, CA: University of California, School of Information and Computer Science <td valign=top> Pima Indians Diabetes Binary Classification dataset <td valign=top> A subset of data from the National Institute of Diabetes and Digestive and Kidney Diseases database. The dataset was filtered to focus on female patients of Pima Indian heritage. The data includes medical data such as glucose and insulin levels, as well as lifestyle factors.

Usage: Predict whether the subject has diabetes (binary classification).

Related Research: Sigillito, V. (1990). UCI Machine Learning Repository <http://archive.ics.uci.edu/ml>". Irvine, CA: University of California, School of Information and Computer Science <td valign=top> Restaurant customer data <td valign=top> A set of metadata about customers, including demographics and preferences.

Usage: Use this dataset, in combination with the other two restaurant datasets, to train and test a recommender system.

Related Research: Bache, K. and Lichman, M. (2013). UCI Machine Learning Repository <http://archive.ics.uci.edu/ml>. Irvine, CA: University of California, School of Information and Computer Science. <td valign=top> Restaurant feature data <td valign=top> A set of metadata about restaurants and their features, such as food type, dining style, and location.

Usage: Use this dataset, in combination with the other two restaurant datasets, to train and test a recommender system.

Related Research: Bache, K. and Lichman, M. (2013). UCI Machine Learning Repository <http://archive.ics.uci.edu/ml>. Irvine, CA: University of California, School of Information and Computer Science. <td valign=top> Restaurant ratings <td valign=top> Contains ratings given by users to restaurants on a scale from 0 to 2.

Usage: Use this dataset, in combination with the other two restaurant datasets, to train and test a recommender system.

Related Research: Bache, K. and Lichman, M. (2013). UCI Machine Learning Repository <http://archive.ics.uci.edu/ml>. Irvine, CA: University of California, School of Information and Computer Science. <td valign=top> Steel Annealing multi-class dataset <td valign=top> This dataset contains a series of records from steel annealing trials with the physical attributes (width, thickness, type (coil, sheet, etc.) of the resulting steel types.

Usage: Predict any of two numeric class attributes; hardness or strength. You might also analyze correlations among attributes.

Steel grades follow a set standard, defined by SAE and other organizations. You are looking for a specific 'grade' (the class variable) and want to understand the values needed.

Related Research: Sterling, D. & Buntine, W. (NA). UCI Machine Learning Repository <http://archive.ics.uci.edu/ml>. Irvine, CA: University of California, School of Information and Computer Science

A useful guide to steel grades can be found here: <http://www.outokumpu.com/SiteCollectionDocuments/Outokumpu-steel-grades-properties-global-standards.pdf> <td valign=top> Telescope data <td valign=top> Records of high energy gamma particle bursts along with background noise, both simulated using a Monte Carlo process.

The intent of the simulation was to improve the accuracy of ground-based atmospheric Cherenkov gamma telescopes, using statistical methods to differentiate between the desired signal (Cherenkov radiation showers) and background noise (hadronic showers initiated by cosmic rays in the upper atmosphere).

The data has been pre-processed to create an elongated cluster with the long axis oriented towards the camera center. The characteristics of this ellipse (often called Hillas parameters) are among the image parameters that can be used for discrimination.

Usage: Predict whether image of a shower represents signal or background noise.

Notes: Simple classification accuracy is not meaningful for this data, since classifying a background event as signal is worse than classifying a signal event as background. For comparison of different classifiers the ROC graph should be used. The probability of accepting a background event as signal must be below one of the following thresholds: 0.01 , 0.02 , 0.05 , 0.1 , or 0.2.

Also, note that the number of background events (h, for hadronic showers) is underestimated, whereas in real measurements, the h or noise class represents the majority of events.

Related Research: Bock, R.K. (1995). UCI Machine Learning Repository <http://archive.ics.uci.edu/ml>. Irvine, CA: University of California, School of Information <tr ID=weather-dataset> <td valign=top> Weather Dataset <td valign=top> Hourly land-based weather observations from NOAA ([merged data from 201304 to 201310](#)).

The weather data covers observations made from airport weather stations, covering the time period April-October 2013. Before uploading to Azure Machine Learning Studio, the dataset was processed as follows:

- Weather station IDs were mapped to corresponding airport IDs
- Weather stations not associated with the 70 busiest airports were filtered out
- The Date column was split into separate Year, Month, and Day columns
- The following columns were selected: AirportID, Year, Month, Day, Time, TimeZone, SkyCondition, Visibility, WeatherType, DryBulbFarenheit, DryBulbCelsius, WetBulbFarenheit, WetBulbCelsius, DewPointFarenheit, DewPointCelsius, RelativeHumidity, WindSpeed, WindDirection, ValueForWindCharacter, StationPressure, PressureTendency, PressureChange, SeaLevelPressure, RecordType, HourlyPrecip, Altimeter

<tr ID=wikipedia-sp-500-dataset> <td valign=top> Wikipedia SP 500 Dataset <td valign=top> Data is derived from Wikipedia (<http://www.wikipedia.org/>) based on articles of each S&P 500 company, stored as XML data.

Before uploading to Azure Machine Learning Studio, the dataset was processed as follows:

- Extract text content for each specific company
- Remove wiki formatting
- Remove non-alphanumeric characters
- Convert all text to lowercase
- Known company categories were added

Note that for some companies an article could not be found, so the number of records is less than 500. <tr ID=direct-marketing> <td valign=top>[direct_marketing.csv](#) <td valign=top> The dataset contains customer data and indications about their response to a direct mailing campaign. Each row represents a customer. The dataset contains 9 features about user demographics and past behavior, and 3 label columns (visit, conversion, and spend). Visit is a binary column that indicates that a customer visited after the marketing campaign, conversion indicates a customer purchased something, and spend is the amount that was spent. The dataset was made available by Kevin Hillstrom for MineThatData E-Mail Analytics And Data Mining Challenge. <tr ID=lyrl2004-tokens-test> <td valign=top>[lyrl2004_tokens_test.csv](#) <td valign=top> Features of test examples in the RCV1-V2 Reuters news dataset. The dataset has 781K news articles along with their IDs (first column of the dataset). Each article is tokenized, stopworded, and stemmed. The dataset was made available by David. D. Lewis. <tr ID=lyrl2004-tokens-train> <td valign=top>[lyrl2004_tokens_train.csv](#) <td valign=top> Features of training examples in the RCV1-V2 Reuters news dataset. The dataset has 23K news articles along with their IDs (first column of the dataset). Each article is tokenized, stopworded, and stemmed. The dataset was made available by David. D. Lewis. <tr ID=intrusion-detection> <td valign=top>[network_intrusion_detection.csv](#) <td valign=top> Dataset from the KDD Cup 1999 Knowledge Discovery and Data Mining Tools Competition ([kddcup99.html](#)).

The dataset was downloaded and stored in Azure Blob storage ([network_intrusion_detection.csv](#)) and includes both training and testing datasets. The training dataset has approximately 126K rows and 43 columns, including the labels. Three columns are part of the label information, and 40 columns, consisting of numeric and string/categorical features, are available for training the model. The test data has approximately 22.5K test examples with the same 43 columns as in the training data. <tr ID=rcv1-v2-topics-qrels> <td valign=top>[rcv1-v2.topics.qrels.csv](#) <td

valign=top> Topic assignments for news articles in the RCV1-V2 Reuters news dataset. A news article can be assigned to several topics. The format of each row is "<topic name> <document id> 1". The dataset contains 2.6M topic assignments. The dataset was made available by David. D. Lewis. <tr ID=student-performance> <td valign=top>[student_performance.txt](#) <td valign=top> This data comes from the KDD Cup 2010 Student performance evaluation challenge ([student performance evaluation](#)). The data used is the Algebra_2008_2009 training set (Stamper, J., Niculescu-Mizil, A., Ritter, S., Gordon, G.J., & Koedinger, K.R. (2010). Algebra I 2008-2009. Challenge dataset from KDD Cup 2010 Educational Data Mining Challenge. Find it at [downloads.jsp](#) or [algebra_2008_2009.zip](#).

The dataset was downloaded and stored in Azure Blob storage ([student_performance.txt](#)) and contains log files from a student tutoring system. The supplied features include problem ID and its brief description, student ID, timestamp, and how many attempts the student made before solving the problem in the right way. The original dataset has 8.9M records; this dataset has been down-sampled to the first 100K rows. The dataset has 23 tab-separated columns of various types: numeric, categorical, and timestamp.

Analyzing Customer Churn by using Azure Machine Learning

9/25/2017 • 12 min to read • [Edit Online](#)

Overview

This article presents a reference implementation of a customer churn analysis project that is built by using Azure Machine Learning. In this article, we discuss associated generic models for holistically solving the problem of industrial customer churn. We also measure the accuracy of models that are built by using Machine Learning, and we assess directions for further development.

Acknowledgements

This experiment was developed and tested by Serge Berger, Principal Data Scientist at Microsoft, and Roger Barga, formerly Product Manager for Microsoft Azure Machine Learning. The Azure documentation team gratefully acknowledges their expertise and thanks them for sharing this white paper.

NOTE

The data used for this experiment is not publicly available. For an example of how to build a machine learning model for churn analysis, see: [Retail churn model template in Cortana Intelligence Gallery](#)

NOTE

Try Azure Machine Learning for free

No credit card or Azure subscription required. [Get started now.](#)

The problem of customer churn

Businesses in the consumer market and in all enterprise sectors have to deal with churn. Sometimes churn is excessive and influences policy decisions. The traditional solution is to predict high-propensity churners and address their needs via a concierge service, marketing campaigns, or by applying special dispensations. These approaches can vary from industry to industry and even from a particular consumer cluster to another within one industry (for example, telecommunications).

The common factor is that businesses need to minimize these special customer retention efforts. Thus, a natural methodology would be to score every customer with the probability of churn and address the top N ones. The top customers might be the most profitable ones; for example, in more sophisticated scenarios, a profit function is employed during the selection of candidates for special dispensation. However, these considerations are only a part of the holistic strategy for dealing with churn. Businesses also have to take into account risk (and associated risk tolerance), the level and cost of the intervention, and plausible customer segmentation.

Industry outlook and approaches

Sophisticated handling of churn is a sign of a mature industry. The classic example is the telecommunications industry where subscribers are known to frequently switch from one provider to another. This voluntary churn is a prime concern. Moreover, providers have accumulated significant knowledge about *churn drivers*, which are the factors that drive customers to switch.

For instance, handset or device choice is a well-known driver of churn in the mobile phone business. As a result, a popular policy is to subsidize the price of a handset for new subscribers and charging a full price to existing customers for an upgrade. Historically, this policy has led to customers hopping from one provider to another to get a new discount, which in turn, has prompted providers to refine their strategies.

High volatility in handset offerings is a factor that very quickly invalidates models of churn that are based on current handset models. Additionally, mobile phones are not only telecommunication devices; they are also fashion statements (consider the iPhone), and these social predictors are outside the scope of regular telecommunications data sets.

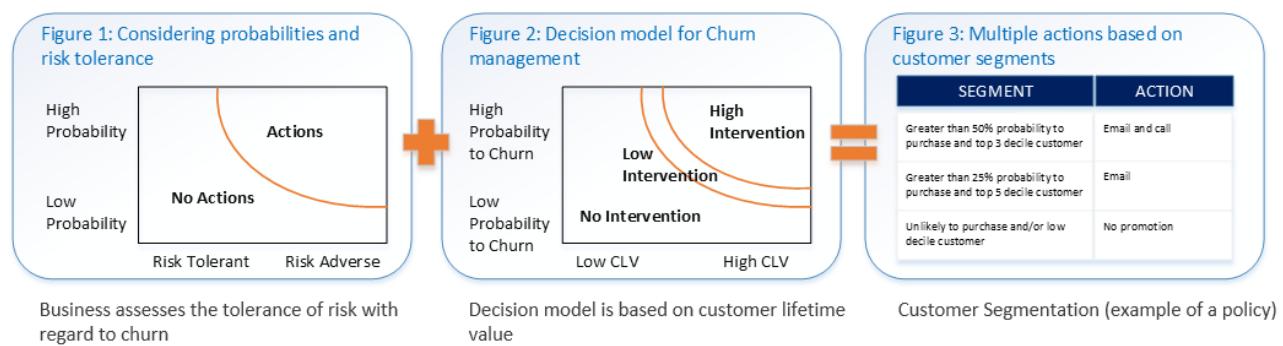
The net result for modeling is that you cannot devise a sound policy simply by eliminating known reasons for churn. In fact, a continuous modeling strategy, including classic models that quantify categorical variables (such as decision trees), is **mandatory**.

Using big data sets on their customers, organizations are performing big data analytics (in particular, churn detection based on big data) as an effective approach to the problem. You can find more about the big data approach to the problem of churn in the Recommendations on ETL section.

Methodology to model customer churn

A common problem-solving process to solve customer churn is depicted in Figures 1-3:

1. A risk model allows you to consider how actions affect probability and risk.
2. An intervention model allows you to consider how the level of intervention could affect the probability of churn and the amount of customer lifetime value (CLV).
3. This analysis lends itself to a qualitative analysis that is escalated to a proactive marketing campaign that targets customer segments to deliver the optimal offer.



This forward looking approach is the best way to treat churn, but it comes with complexity: we have to develop a multi-model archetype and trace dependencies between the models. The interaction among models can be encapsulated as shown in the following diagram:

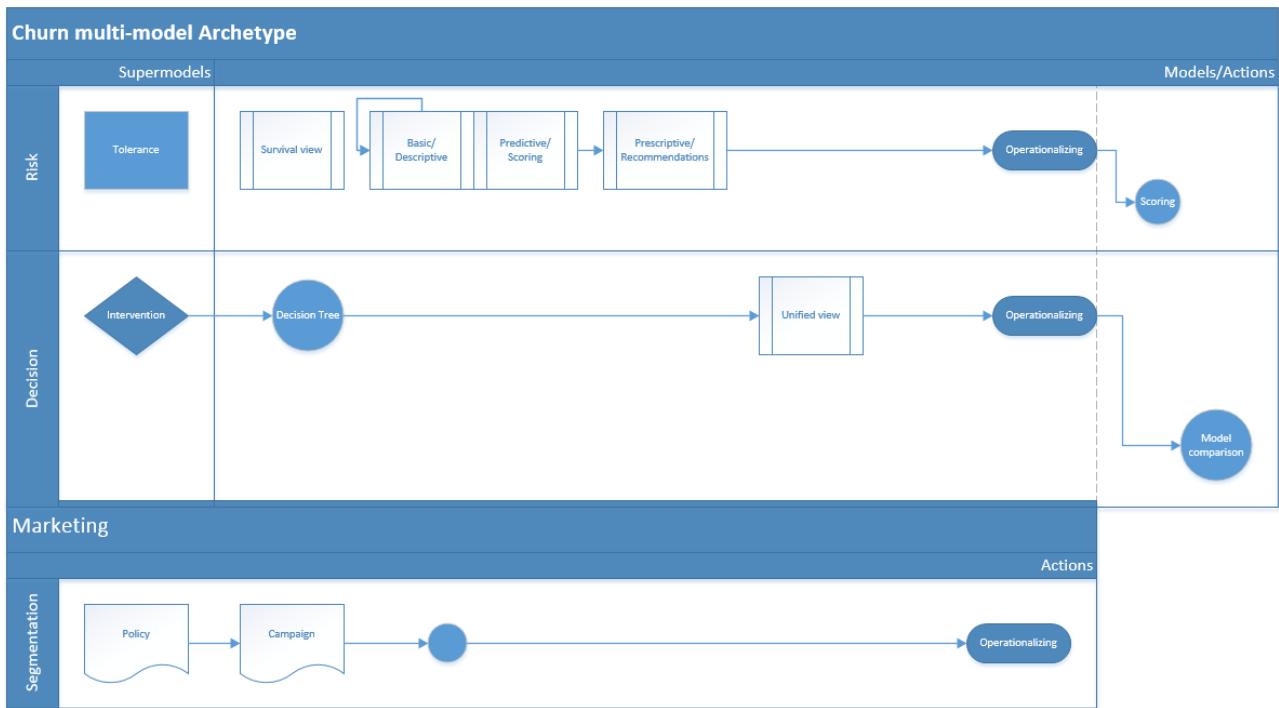


Figure 4: Unified multi-model archetype

Interaction between the models is key if we are to deliver a holistic approach to customer retention. Each model necessarily degrades over time; therefore, the architecture is an implicit loop (similar to the archetype set by the CRISP-DM data mining standard, [3]).

The overall cycle of risk-decision-marketing segmentation/decomposition is still a generalized structure, which is applicable to many business problems. Churn analysis is simply a strong representative of this group of problems because it exhibits all the traits of a complex business problem that does not allow a simplified predictive solution. The social aspects of the modern approach to churn are not particularly highlighted in the approach, but the social aspects are encapsulated in the modeling archetype, as they would be in any model.

An interesting addition here is big data analytics. Today's telecommunication and retail businesses collect exhaustive data about their customers, and we can easily foresee that the need for multi-model connectivity will become a common trend, given emerging trends such as the Internet of Things and ubiquitous devices, which allow business to employ smart solutions at multiple layers.

Implementing the modeling archetype in Machine Learning Studio

Given the problem just described, what is the best way to implement an integrated modeling and scoring approach? In this section, we will demonstrate how we accomplished this by using Azure Machine Learning Studio.

The multi-model approach is a must when designing a global archetype for churn. Even the scoring (predictive) part of the approach should be multi-model.

The following diagram shows the prototype we created, which employs four scoring algorithms in Machine Learning Studio to predict churn. The reason for using a multi-model approach is not only to create an ensemble classifier to increase accuracy, but also to protect against over-fitting and to improve prescriptive feature selection.

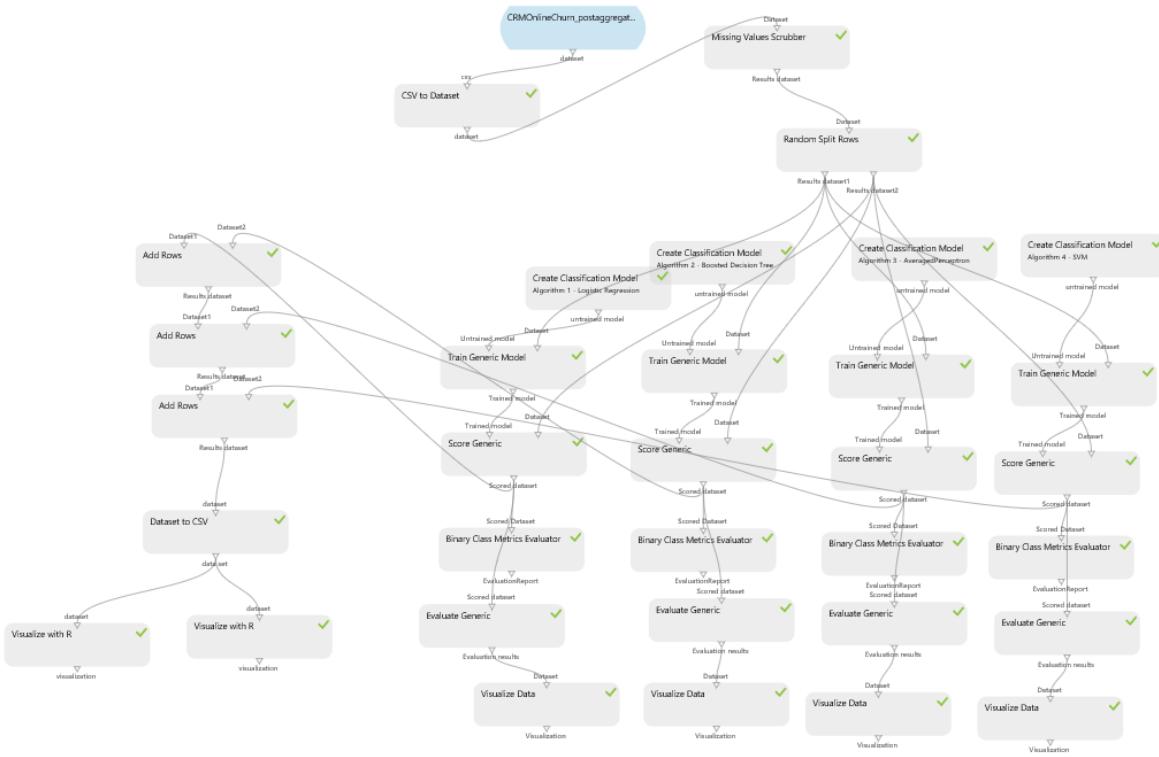


Figure 5: Prototype of a churn modeling approach

The following sections provide more details about the prototype scoring model that we implemented by using Machine Learning Studio.

Data selection and preparation

The data used to build the models and score customers was obtained from a CRM vertical solution, with the data obfuscated to protect customer privacy. The data contains information about 8,000 subscriptions in the U.S., and it combines three sources: provisioning data (subscription metadata), activity data (usage of the system), and customer support data. The data does not include any business related information about the customers; for example, it does not include loyalty metadata or credit scores.

For simplicity, ETL and data cleansing processes are out of scope because we assume that data preparation has already been done elsewhere.

Feature selection for modeling is based on preliminary significance scoring of the set of predictors, included in the process that uses the random forest module. For the implementation in Machine Learning Studio, we calculated the mean, median, and ranges for representative features. For example, we added aggregates for the qualitative data, such as minimum and maximum values for user activity.

We also captured temporal information for the most recent six months. We analyzed data for one year and we established that even if there were statistically significant trends, the effect on churn is greatly diminished after six months.

The most important point is that the entire process, including ETL, feature selection, and modeling was implemented in Machine Learning Studio, using data sources in Microsoft Azure.

The following diagrams illustrate the data that was used.

Churn Binary	GroupCustomerDBID	Mean(Act)	Impute_I	IMP_Mean	Mean(Activ)	Impute_	IMP	Mean(Cases)	Imp	IMP_Mean(I)	Mean(Co)	Impute_
1	Commerce.3766384	2.833213	0	2.833213	7.628518	0	7.6285	0	0	0	4.430817	0
1	Commerce.3648614	2.197225	0	2.197225	5.147494	0	5.1475	0	0	0	3.401197	0
1	Commerce.1701750	4.682131	0	4.682131	7.451242	0	7.4512	2.197225	0	2.197225	5.056246	0
1	Commerce.1451565	4.787492	0	4.787492	7.960803	0	7.9608	0	0	0	5.398163	0
1	Commerce.2115489	3.433987	0	3.433987	3.7612	0	3.7612	1.791759	0	1.791759	3.091042	0
1	Commerce.6205107	3.988984	0	3.988984	7.467873	0	7.4679	0.693147	0	0.693147	4.89784	0
1	Commerce.7125701	4.356709	0	4.356709	4.077537	0	4.0775	0	0	0	6.603944	0
1	Commerce.2808747	0.693147	0	0.693147	7.427045	0	7.427	0	0	0	7.029973	0
0	Commerce.3213575	7.17012	0	7.17012	5.590987	0	5.591	0	0	0	7.284135	0

Figure 6: Excerpt of data source (obfuscated)

Columns (292/0)
Churn Binary
GroupCustomerDBID
Mean(Accounts) etc. (3/0)
Mean(Activities) etc. (3/0)
Mean(Cases) etc. (3/0)
Mean(Contacts) etc. (3/0)
Mean>EmailAttachments) etc. (3/0)
Mean(KBArticles) etc. (3/0)
Mean(Leads) etc. (3/0)
Mean(MarketingCampaigns) etc. (3/0)
Mean(MarketingLists) etc. (3/0)
Mean(Notes) etc. (3/0)
Mean(Orders) etc. (3/0)
Mean(Quotes) etc. (3/0)
Std Dev(Accounts) etc. (3/0)
Std Dev(Activities) etc. (3/0)
Std Dev(Cases) etc. (3/0)
Std Dev(Contacts) etc. (3/0)
Std Dev>EmailAttachments) etc. (3/0)
Std Dev(KBArticles) etc. (3/0)
Std Dev(Leads) etc. (3/0)
Std Dev(MarketingCampaigns) etc. (3/0)
Std Dev(MarketingLists) etc. (3/0)
Std Dev(Notes) etc. (3/0)
Std Dev(Orders) etc. (3/0)
Std Dev(Quotes) etc. (3/0)
Min(Accounts) etc. (3/0)
Min(Activities) etc. (3/0)
Min(Cases) etc. (3/0)
Min(Contacts) etc. (3/0)
Min>EmailAttachments) etc. (3/0)

Figure 7: Features extracted from data source

Note that this data is private and therefore the model and data cannot be shared. However, for a similar model using publicly available data, see this sample experiment in the [Cortana Intelligence Gallery: Telco Customer Churn](#).

To learn more about how you can implement a churn analysis model using Cortana Intelligence Suite, we also recommend [this video](#) by Senior Program Manager Wee Hyong Tok.

Algorithms used in the prototype

We used the following four machine learning algorithms to build the prototype (no customization):

1. Logistic regression (LR)
2. Boosted decision tree (BT)
3. Averaged perceptron (AP)
4. Support vector machine (SVM)

The following diagram illustrates a portion of the experiment design surface, which indicates the sequence in which the models were created:

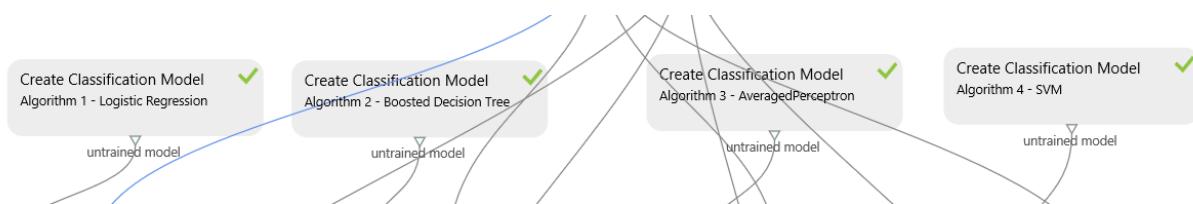


Figure 8: Creating models in Machine Learning Studio

Scoring methods

We scored the four models by using a labeled training dataset.

We also submitted the scoring dataset to a comparable model built by using the desktop edition of SAS Enterprise Miner 12. We measured the accuracy of the SAS model and all four Machine Learning Studio models.

Results

In this section, we present our findings about the accuracy of the models, based on the scoring dataset.

Accuracy and precision of scoring

Generally, the implementation in Azure Machine Learning is behind SAS in accuracy by about 10-15% (Area Under Curve or AUC).

However, the most important metric in churn is the misclassification rate: that is, of the top N churners as predicted by the classifier, which of them actually did **not** churn, and yet received special treatment? The following diagram compares this misclassification rate for all the models:

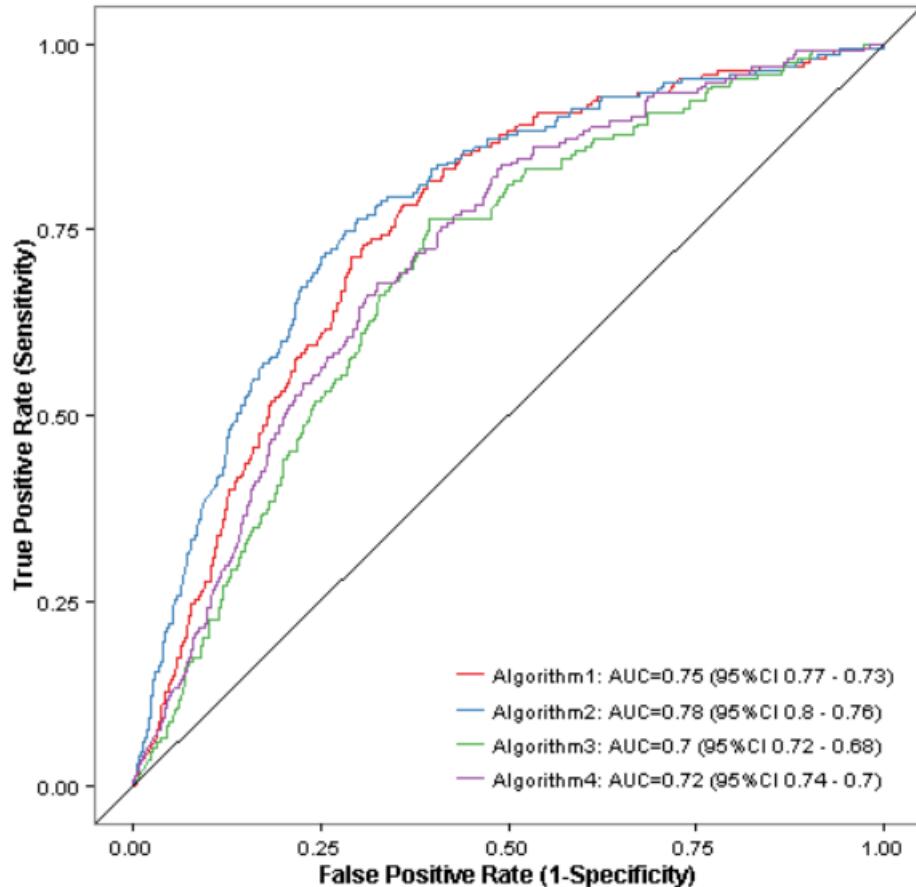


Figure 9: Passau prototype area under curve

Using AUC to compare results

Area Under Curve (AUC) is a metric that represents a global measure of *separability* between the distributions of scores for positive and negative populations. It is similar to the traditional Receiver Operator Characteristic (ROC) graph, but one important difference is that the AUC metric does not require you to choose a threshold value. Instead, it summarizes the results over **all** possible choices. In contrast, the traditional ROC graph shows the positive rate on the vertical axis and the false positive rate on the horizontal axis, and the classification threshold varies.

AUC is generally used as a measure of worth for different algorithms (or different systems) because it allows models to be compared by means of their AUC values. This is a popular approach in industries such as meteorology and biosciences. Thus, AUC represents a popular tool for assessing classifier performance.

Comparing misclassification rates

We compared the misclassification rates on the dataset in question by using the CRM data of approximately 8,000 subscriptions.

- The SAS misclassification rate was 10-15%.
- The Machine Learning Studio misclassification rate was 15-20% for the top 200-300 churners.

In the telecommunications industry, it is important to address only those customers who have the highest risk to churn by offering them a concierge service or other special treatment. In that respect, the Machine Learning Studio implementation achieves results on par with the SAS model.

By the same token, accuracy is more important than precision because we are mostly interested in correctly classifying potential churners.

The following diagram from Wikipedia depicts the relationship in a lively, easy-to-understand graphic:

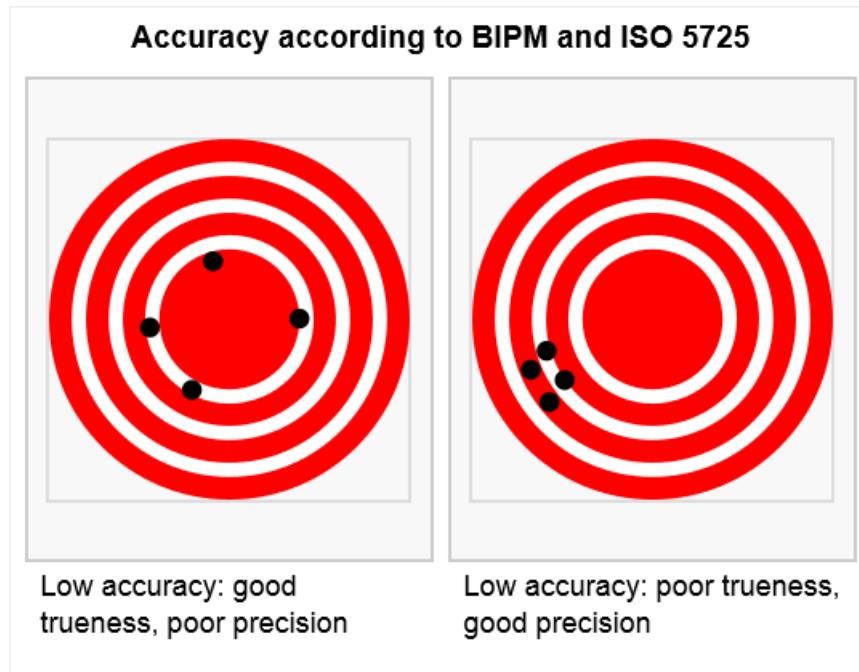


Figure 10: Tradeoff between accuracy and precision

Accuracy and precision results for boosted decision tree model

The following chart displays the raw results from scoring using the Machine Learning prototype for the boosted decision tree model, which happens to be the most accurate among the four models:

Accuracy	Precision	Recall	F-Score	AUC	Average Log Loss	Training Log Loss
0.895685	0.282258	0.179487	0.219436	0.778342	0.0306191	89.1761
0.895685	0.282258	0.179487	0.219436	0.778342	0.0306191	89.1761
0.895685	0.282258	0.179487	0.219436	0.778342	0.0306191	89.1761
0.895685	0.282258	0.179487	0.219436	0.778342	0.0306191	89.1761

Figure 11: Boosted decision tree model characteristics

Performance comparison

We compared the speed at which data was scored using the Machine Learning Studio models and a comparable model created by using the desktop edition of SAS Enterprise Miner 12.1.

The following table summarizes the performance of the algorithms:

Table 1. General performance (accuracy) of the algorithms

LR	BT	AP	SVM
Average Model	The Best Model	Underperforming	Average Model

The models hosted in Machine Learning Studio outperformed SAS by 15-25% for speed of execution, but accuracy was largely on par.

Discussion and recommendations

In the telecommunications industry, several practices have emerged to analyze churn, including:

- Derive metrics for four fundamental categories:
 - **Entity (for example, a subscription).** Provision basic information about the subscription and/or customer that is the subject of churn.
 - **Activity.** Obtain all possible usage information that is related to the entity, for example, the number of logins.
 - **Customer support.** Harvest information from customer support logs to indicate whether the subscription had issues or interactions with customer support.
 - **Competitive and business data.** Obtain any information possible about the customer (for example, can be unavailable or hard to track).
- Use importance to drive feature selection. This implies that the boosted decision tree model is always a promising approach.

The use of these four categories creates the illusion that a simple *deterministic* approach, based on indexes formed on reasonable factors per category, should suffice to identify customers at risk for churn. Unfortunately, although this notion seems plausible, it is a false understanding. The reason is that churn is a temporal effect and the factors contributing to churn are usually in transient states. What leads a customer to consider leaving today might be different tomorrow, and it certainly will be different six months from now. Therefore, a *probabilistic* model is a necessity.

This important observation is often overlooked in business, which generally prefers a business intelligence-oriented approach to analytics, mostly because it is an easier sell and admits straightforward automation.

However, the promise of self-service analytics by using Machine Learning Studio is that the four categories of information, graded by division or department, become a valuable source for machine learning about churn.

Another exciting capability coming in Azure Machine Learning is the ability to add a custom module to the repository of predefined modules that are already available. This capability, essentially, creates an opportunity to select libraries and create templates for vertical markets. It is an important differentiator of Azure Machine Learning in the market place.

We hope to continue this topic in the future, especially related to big data analytics.

Conclusion

This paper describes a sensible approach to tackling the common problem of customer churn by using a generic framework. We considered a prototype for scoring models and implemented it by using Azure Machine Learning. Finally, we assessed the accuracy and performance of the prototype solution with regard to comparable algorithms in SAS.

For more information:

Did this paper help you? Please give us your feedback. Tell us on a scale of 1 (poor) to 5 (excellent), how would you

rate this paper and why have you given it this rating? For example:

- Are you rating it high due to having good examples, excellent screen shots, clear writing, or another reason?
- Are you rating it low due to poor examples, fuzzy screen shots, or unclear writing?

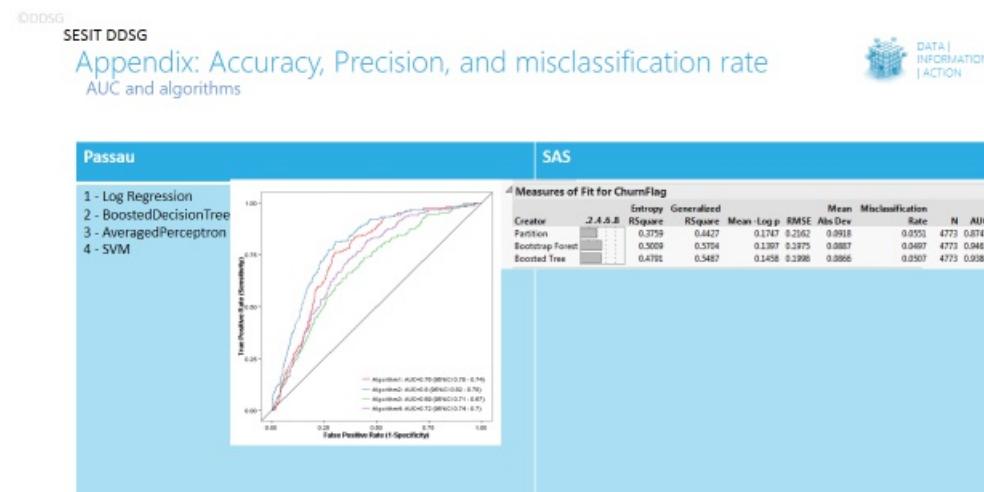
This feedback will help us improve the quality of white papers we release.

[Send feedback.](#)

References

- 1 Predictive Analytics: Beyond the Predictions, W. McKnight, Information Management, July/August 2011, p.18-20.
- 2 Wikipedia article: [Accuracy and precision](#)
- [3] [CRISP-DM 1.0: Step-by-Step Data Mining Guide](#)
- [4] [Big Data Marketing: Engage Your Customers More Effectively and Drive Value](#)
- [5] [Telco churn model template in Cortana Intelligence Gallery](#)

Appendix



Microsoft Confidential | PPT Internal Use Only MICROSOFT IT | DATA & DECISION SCIENCES GROUP

Figure 12: Snapshot of a presentation on churn prototype

PowerShell module for Microsoft Azure Machine Learning

9/25/2017 • 2 min to read • [Edit Online](#)

The PowerShell module for Azure Machine Learning is a powerful tool that allows you to use Windows PowerShell to manage workspaces, experiments, datasets, Classic web services, and more.

You can view the documentation and download the module, along with the full source code, at <https://aka.ms/amlps>.

NOTE

The Azure Machine Learning PowerShell module is currently in preview mode. The module will continue to be improved and expanded during this preview period. Keep an eye on the [Cortana Intelligence and Machine Learning Blog](#) for news and information.

What is the Machine Learning PowerShell module?

The Machine Learning PowerShell module is a .NET-based DLL module that allows you to fully manage Azure Machine Learning workspaces, experiments, datasets, Classic web services, and Classic web service endpoints from Windows PowerShell.

Along with the module, you can download the full source code which includes a cleanly separated [C# API layer](#). You can reference this DLL from your own .NET project and manage Azure Machine Learning through .NET code. In addition, the DLL depends on underlying REST APIs that you can use directly from your favorite client.

What can I do with the PowerShell module?

Here are some of the tasks you can perform with this PowerShell module. Check out the [full documentation](#) for these and many more functions.

- Provision a new workspace using a management certificate ([New-AmlWorkspace](#))
- Export and import a JSON file representing an experiment graph ([Export-AmlExperimentGraph](#) and [Import-AmlExperimentGraph](#))
- Run an experiment ([Start-AmlExperiment](#))
- Create a web service out of a predictive experiment ([New-AmlWebService](#))
- Create an endpoint on a published web service ([Add-AmlWebServiceEndpoint](#))
- Invoke an RRS and/or BES web service endpoint ([Invoke-AmlWebServiceRRSEndpoint](#) and [Invoke-AmlWebServiceBESEndpoint](#))

Here's a quick example of using PowerShell to run an existing experiment:

```
#Find the first Experiment named "xyz"
$exp = (Get-AmlExperiment | where Description -eq 'xyz')[0]
#Run the Experiment
Start-AmlExperiment -ExperimentId $exp.ExperimentId
```

For a more in-depth use case, see this article on using the PowerShell module to automate a commonly-requested task: [Create many Machine Learning models and web service endpoints from one experiment using PowerShell](#).

How do I get started?

To get started with Machine Learning PowerShell, download the [release package](#) from GitHub and follow the [instructions for installation](#). The instructions explain how to unblock the downloaded/unzipped DLL and then import it into your PowerShell environment. Most of the cmdlets require that you supply the workspace ID, the workspace authorization token, and the Azure region that the workspace is in. The simplest way to provide the values is through a default config.json file. The instructions also explain how to configure this file.

And if you want, you can clone the git tree, modify the code, and compile it locally using Visual Studio.

Next steps

You can find the full documentation for the PowerShell module at <https://aka.ms/amlps>.

For an extended example of how to use the module in a real-world scenario, check out the in-depth use case, [Create many Machine Learning models and web service endpoints from one experiment using PowerShell](#).

Machine Learning REST API Error Codes

9/25/2017 • 7 min to read • [Edit Online](#)

The following error codes could be returned by an operation on an Azure Machine Learning web service.

BadArgument (HTTP status code 400)

Invalid argument provided.

This class of errors means an argument provided somewhere was invalid. This could be a credential or location of Azure storage to something passed to the web service. Please look at the error "code" field in the "details" section to diagnose which specific argument was invalid.

ERROR CODE	USER MESSAGE
BadParameterValue	The parameter value supplied does not satisfy the parameter rule on the parameter
BadSubscriptionId	The subscription Id that is used to score is not the one present in the resource
BadVersionCall	Invalid version parameter was passed during the API call: {0}. Check the API help page for passing the correct version and try again.
BatchJobInputsNotSpecified	The following required input(s) were not specified with the request: {0}. Please ensure all input data is specified and try again.
BatchJobInputsTooManySpecified	The request specified more inputs than defined in the service. List of accepted input(s): {0}. Please ensure all input data is specified correctly and try again.
BlobNameTooLong	Azure blob storage path provided for diagnostic output is too long: {0}. Shorten the path and try again.
BlobNotFound	Unable to access the provided Azure blob - {0}. Azure error message: {1}.
ContainerIsEmpty	No Azure storage container name was provided. Provide a valid container name and try again.
ContainerSegmentInvalid	Invalid container name. Provide a valid container name and try again.
ContainerValidationFailed	Blob container validation failed with this error: {0}.
DataTypeNotSupported	Unsupported data type provided. Provide valid data type(s) and try again.

Error Code	User Message
DuplicateInputInBatchCall	The batch request is invalid. Cannot specify both single and multiple input at the same time. Remove one of these items from the request and try again.
ExpiryTimeInThePast	Expiry time provided is in the past: {0}. Provide a future expiry time in UTC and try again. To never expire, set expiry time to NULL.
IncompleteSettings	Diagnostics settings are incomplete.
InputBlobRelativeLocationInvalid	No Azure storage blob name provided. Provide a valid blob name and try again.
InvalidBlob	Invalid blob specification for blob: {0}. Verify that connection string / relative path or SAS token specification is correct and try again.
InvalidBlobConnectionString	The connection string specified for one of the input/output blobs is invalid: {0}. Please correct this and try again.
InvalidBlobExtension	The blob reference: {0} has an invalid or missing file extension. Supported file extensions for this output type are: "{1}".
InvalidInputNames	Invalid service input name(s) specified in the request: {0}. Please map the input data to the correct service inputs and try again.
InvalidOutputOverrideName	Invalid output override name: {0}. The service does not have an output node with this name. Please pass in a correct output node name to override (case sensitivity applies).
InvalidQueryParameter	Invalid query parameter '{0}'. {1}
MissingInputBlobInformation	Missing Azure storage blob information. Provide a valid connection string and relative path or URI and try again.
MissingJobId	No job Id provided. A job Id is returned when a job was submitted for the first time. Verify the job Id is correct and try again.
MissingKeys	No Keys provided or one of Primary or Secondary Key is not provided.
MissingModelPackage	No model package Id or model package provided. Provide a valid model package Id or model package and try again.
MissingOutputOverrideSpecification	The request is missing the blob specification for output override {0}. Please specify a valid blob location with the request, or remove the output specification if no location override is desired.
MissingRequestInput	The web service expects an input, but no input was provided. Ensure valid inputs are provided based on the published input ports in the model and try again.

ERROR CODE	USER MESSAGE
MissingRequiredGlobalParameters	Not all required web service parameter(s) provided. Verify the parameter(s) expected for the module(s) are correct and try again.
MissingRequiredOutputOverrides	When calling an encrypted service endpoint it is mandatory to pass in output overrides for all the service's outputs. Missing overrides at this time for these outputs: {0}
MissingWebServiceGroupId	No web service group Id provided. Provide a valid web service group Id and try again.
MissingWebServiceId	No web service Id provided. Provide a valid web service Id and try again.
MissingWebServicePackage	No web Service package provided. Provide a valid web service package and try again.
MissingWorkspaceId	No workspace Id provided. Provide a valid workspace Id and try again.
ModelConfigurationInvalid	Invalid model configuration in the model package. Ensure the model configuration contains output endpoint(s) definition, std error endpoint, and std out endpoint and try again.
ModelPackageIdInvalid	Invalid model package Id. Verify that the model package Id is correct and try again.
RequestBodyInvalid	No request body provided or error in deserializing the request body.
RequestIsEmpty	No request provided. Provide a valid request and try again.
UnexpectedParameter	Unexpected parameters provided. Verify all parameter names are spelled correctly, only expected parameters are passed, and try again.
UnknownError	Unknown error.
UserParameterInvalid	{0}
WebServiceConcurrentRequestRequirementInvalid	Cannot change concurrent requests requirements for {0} web service.
WebServiceIdInvalid	Invalid web service id provided. Web service id should be a valid guid.
WebServiceTooManyConcurrentRequestRequirement	Cannot set concurrent request requirement to more than {0}.
WebServiceTypeInvalid	Invalid web service type provided. Verify the valid web service type is correct and try again. Valid web service types: {0}.

BadUserArgument (HTTP status code 400)

Invalid user argument provided.

ERROR CODE	USER MESSAGE
InputMismatchError	Input data does not match input port schema.
InputParseError	Parsing of input vector failed. Verify the input vector has the correct number of columns and data types. Additional details: {0}.
MissingRequiredGlobalParameters	Parameter(s) expected by the web service are missing. Verify all the required parameters expected by the web service are correct and try again.
UnexpectedParameter	Verify only the required parameters expected by the web service are passed and try again.
UserParameterInvalid	{0}

InvalidOperationException (HTTP status code 400)

The request is invalid in the current context.

ERROR CODE	USER MESSAGE
CannotStartJob	The job cannot be started because it is in {0} state.
IncompatibleModel	The model is incompatible with the request version. The request version only supports single datatable output models.
MultipleInputsNotAllowed	The model does not allow multiple inputs.

LibraryExecutionError (HTTP status code 400)

Module execution encountered an internal library error.

ModuleExecutionError (HTTP status code 400)

Module execution encountered an error.

WebServicePackageError (HTTP status code 400)

Invalid web service package. Verify the web service package provided is correct and try again.

ERROR CODE	USER MESSAGE
FormatError	The web service package is malformed. Details: {0}
RuntimesError	The web service package graph is invalid. Details: {0}
ValidationErrors	The web service package graph is invalid. Details: {0}

Unauthorized (HTTP status code 401)

Request is unauthorized to access resource.

ERROR CODE	USER MESSAGE
AdminRequestUnauthorized	Unauthorized
ManagementRequestUnauthorized	Unauthorized
ScoreRequestUnauthorized	Invalid credentials provided.

NotFound (HTTP status code 404)

Resource not found.

ERROR CODE	USER MESSAGE
ModelPackageNotFound	Model package not found. Verify the model package Id is correct and try again.
WebServiceIdNotFoundInWorkspace	Web service under this workspace not found. There is a mismatch between the webServiceId and the workspaceId. Verify the web service provided is part of the workspace and try again.
WebServiceNotFound	Web service not found. Verify the web service Id is correct and try again.
WorkspaceNotFound	Workspace not found. Verify the workspace Id is correct and try again.

RequestTimeout (HTTP status code 408)

The operation could not be completed within the permitted time.

ERROR CODE	USER MESSAGE
RequestCanceled	Request was canceled by the client.
ScoreRequestTimeout	Execution request timed out.

Conflict (HTTP status code 409)

Resource already exists.

ERROR CODE	USER MESSAGE
ModelErrorMetadataMismatch	Invalid output parameter name. Try using the metadata editor module to rename columns and try again.

MemoryQuotaViolation (HTTP status code 413)

The model had exceeded the memory quota assigned to it.

ERROR CODE	USER MESSAGE
OutOfMemoryLimit	The model consumed more memory than was appropriated for it. Maximum allowed memory for the model is {0} MB. Please check your model for issues.

InternalError (HTTP status code 500)

Execution encountered an internal error.

ERROR CODE	USER MESSAGE
AdminAuthenticationFailed	
BackendArgumentError	
BackendBadRequest	
ClusterConfigBlobMisconfigured	
ContainerProcessTerminatedWithSystemError	The container process crashed with system error
ContainerProcessTerminatedWithUnknownError	The container process crashed with unknown error
ContainerValidationFailed	Blob container validation failed with this error: {0}.
DeleteWebServiceResourceFailed	
ExceptionDeserializationError	
FailedGettingApiDocument	
FailedStoringWebService	
InvalidMemoryConfiguration	InvalidMemoryConfiguration, ConfigValue: {0}
InvalidResourceCacheConfiguration	
InvalidResourceDownloadConfiguration	
InvalidWebServiceResources	
MissingTaskInstance	No arguments provided. Verify that valid arguments are passed and try again.
ModelPackageInvalid	
ModuleExecutionFailed	
ModuleLoadFailed	
ModuleObjectCloneFailed	

ERROR CODE	USER MESSAGE
OutputConversionFailed	
PortDataTypeNotSupported	Port id={0} has an unsupported data type: {1}.
ResourceDownload	
ResourceLoadFailed	
ServiceUrisNotFound	
SwaggerGeneration	Swagger generation failed, Details: {0}
UnexpectedScoreStatus	
UnknownBackendErrorResponse	
UnknownError	
UnknownJobStatusCode	Unknown job status code {0}.
UnknownModuleError	
UpdateWebServiceResourceFailed	
WebServiceGroupNotFound	
WebServicePackageInvalid	InvalidWebServicePackage, Details: {0}
WorkerAuthorizationFailed	
WorkerUnreachable	

InternalServerErrorSystemLowOnMemory (HTTP status code 500)

Execution encountered an internal error. System low on memory. Please try again.

ModelPackageFormatError (HTTP status code 500)

Invalid model package. Verify the model package provided is correct and try again.

WebServicePackageInternalError (HTTP status code 500)

Invalid web service package. Verify the web package provided is correct and try again.

ERROR CODE	USER MESSAGE
ModuleError	The web service package graph is invalid. Details: {0}

InitializingContainers (HTTP status code 503)

The request cannot execute as the containers are being initialized.

ServiceUnavailable (HTTP status code 503)

Service is temporarily unavailable.

ERROR CODE	USER MESSAGE
NoMoreResources	No resources available for request.
RequestThrottled	Request was throttled for {0} endpoint. The maximum concurrency for the endpoint is {1}.
TooManyConcurrentRequests	Too many concurrent requests sent.
TooManyHostsBeingInitialized	Too many hosts being initialized at the same time. Consider throttling / retrying.
TooManyHostsBeingInitializedPerModel	Too many hosts being initialized at the same time. Consider throttling / retrying.

GatewayTimeout (HTTP status code 504)

The operation could not be completed within the permitted time.

ERROR CODE	USER MESSAGE
BackendInitializationTimeout	The web service initialization could not be completed within the permitted time.
BackendScoreTimeout	The web service request execution could not be completed within the permitted time.

Share and discover resources in the Cortana Intelligence Gallery

9/25/2017 • 3 min to read • [Edit Online](#)

Cortana Intelligence Gallery is a community-driven site for discovering and sharing solutions built with Cortana Intelligence Suite. The Gallery has a variety of resources that you can use to develop your own analytics solutions.

NOTE

Try Azure Machine Learning for free

No credit card or Azure subscription required. [Get started now.](#)

What can I find in the Gallery?

The Cortana Intelligence Gallery contains a variety of resources that you can use to develop your own analytics solutions.

- **Experiments** - The Gallery contains a wide variety of experiments that have been developed in Azure Machine Learning Studio. These range from quick proof-of-concept experiments that demonstrate a specific machine learning technique, to fully-developed solutions for complex machine learning problems.
- **Jupyter Notebooks** - Jupyter Notebooks include code, data visualizations, and documentation in a single, interactive canvas. Notebooks in the Gallery provide tutorials and detailed explanations of advanced machine learning techniques and solutions.
- **Solutions** - Quickly build Cortana Intelligence Solutions from preconfigured solutions, reference architectures, and design patterns. Make them your own with the included instructions or with a featured partner.
- **Tutorials** - A number of tutorials are available to walk you through machine learning technologies and concepts, or to describe advanced methods for solving various machine learning problems.

These basic Gallery resources can be grouped together logically in a couple different ways:

- **Collections** - A collection allows you to group together experiments, APIs, and other Gallery items that address a specific solution or concept.
- **Industries** - The Industries section of the Gallery brings together various resources that are specific to such industries as retail, manufacturing, banking, and healthcare.

Finally, **Competitions** provide an exciting opportunity to compete with the community of data scientists to solve complex problems using Cortana Intelligence Suite.

Cortana Intelligence Gallery

Browse all Industries Solutions Experiments Machine Learning APIs Competitions Notebooks More

Cortana Intelligence Gallery enables our growing community of developers and data scientists to share their analytics solutions. [Learn how to contribute.](#)

EXPERIMENT
Winning Solution Decoding Brain Signals challenge
ALB

SOLUTION
IT Anomaly Insights
Microsoft

SOLUTION
Data Warehousing and Data Science with SQL Data Warehouse
Microsoft

TUTORIAL
Linux Data Science Virtual Machine
Microsoft

SOLUTION
Stream Analysis with Azure ML
Microsoft

Recently added

See all

EXPERIMENT

TUTORIAL

EXPERIMENT

EXPERIMENT

Discover and learn

Anyone can browse and search the different types of resources in the Gallery that have been contributed by Microsoft and the advanced analytics community. Use these to learn more and get a head start on solving your own data analysis problems. You can also download experiments and Jupyter notebooks to your own Machine Learning Studio workspace.

You can easily find recently published and popular resources in the Gallery, or you can search by name, tags, algorithms, and other attributes. Click **Browse all** in the Gallery header, and then select search refinements on the left of the page and enter search terms at the top.

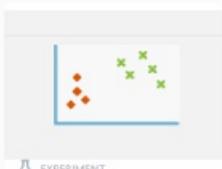
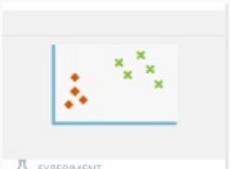
View contributions from a particular author by clicking the author name from within any of the tiles:

Cortana Intelligence Gallery

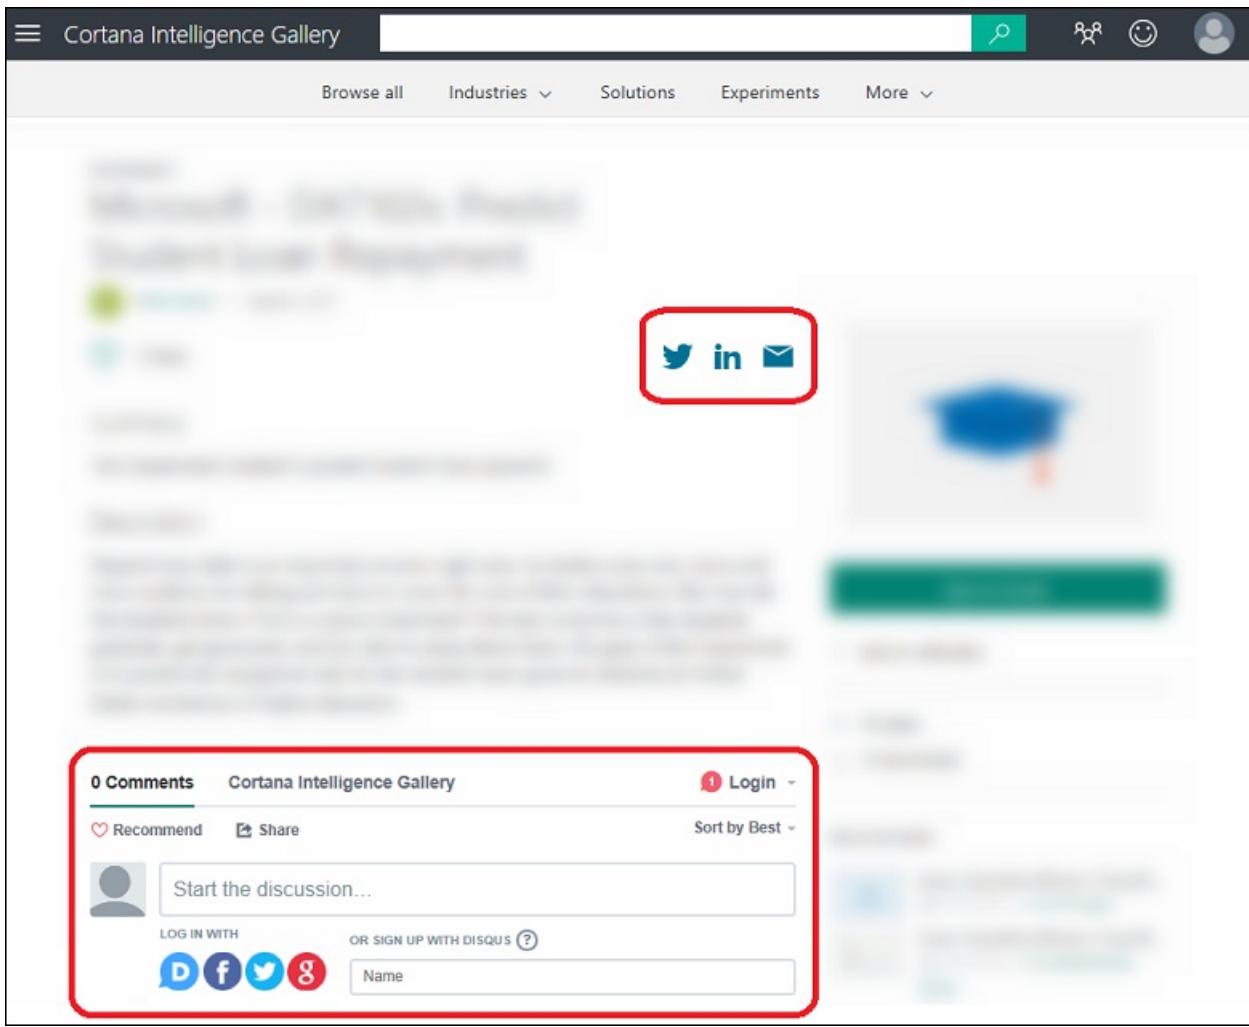
Browse all Industries Solutions Experiments Machine Learning APIs Competitions Notebooks More

12 ITEMS 4 COLLECTIONS

Sort by: Popular

 EXPERIMENT Predictive Experiment - Mini Twitter sentiment analysis This experiment demonstrates the use of the Execute R Script, Feature Selection, Feature Hashing modules to train a text sentiment classification engine. 🕒 2604 ⚡ 1701 11 months ago	 EXPERIMENT Time Series Forecasting using Custom Modules Time Series Forecasting with Azure ML R custom modules for Arima and ETS 🕒 280 ⚡ 116 13 days ago	 EXPERIMENT Training Experiment Mini Twitter sentiment analysis This experiment demonstrates the use of the Execute R Script, Feature Selection, Feature Hashing modules to train a text sentiment classification engine. 🕒 674 ⚡ 213 11 months ago	 EXPERIMENT Demand Estimation Using Lag Features This experiment demonstrates demand estimation using regression with UCI bike rental data. 🕒 190 ⚡ 91 13 days ago	 EXPERIMENT Training Experiment for Twitter sentiment analysis This experiment demonstrates the use of the Execute R Script, Feature Selection, Feature Hashing modules to train a text sentiment classification engine. 🕒 249 ⚡ 104 5 months ago
 EXPERIMENT Compare Datasets This module compares two datasets and summarizes the comparison results.	 EXPERIMENT Predictive Experiment for Twitter sentiment analysis This experiment demonstrates the use of the Execute R Script, Feature Selection, Feature Hashing modules to train a text sentiment classification engine.	 EXPERIMENT Using XGBoost to build Graduation Admit Model This experiment shows a sample of using XGBoost to find a probability of graduate admission.	 EXPERIMENT Retrieve CPI from data.gov This experiment dynamically retrieves XML data from data.gov for CPI and converts XML to JSON.	 EXPERIMENT Retrieve data.gov datasets This experiment dynamically retrieves XML data from data.gov for your data.gov data catalog with.

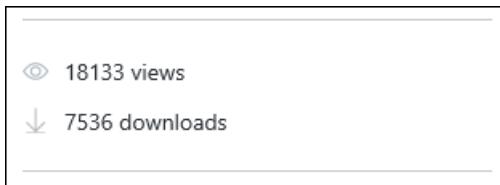
You can comment, provide feedback, or ask questions through the comments section on each resource page. You can even share a resource of interest with friends or colleagues using the share capabilities of LinkedIn or Twitter. You may also email links to these resources to invite other users to view the pages.



Contribute to the Gallery

When you sign in you become a member of the Gallery community. This allows you to contribute your own Gallery items so that others can benefit from the solutions you've discovered.

As others come across your contribution in the Gallery, you can follow the number of views and downloads of your contribution:



Users can also add comments and share your contribution with other members of the data science community. You can log in with a discussion tool such as Disqus and receive notifications for comments on your contributions.

8 Comments

Cortana Intelligence Gallery

♥ Recommend 8

Share



Join the discussion...



disqus_lolC7Ftsg · 2 months ago

After setting input output deploying web service gives error: ports selected for input and output are not valid.

^ | v · Reply · Share >



Brandon Rohrer → disqus_lolC7Ftsg · 2 months ago

Hi disqus_lolC7Ftsg,

Without seeing your workspace, it's tough to know exactly what correction to make. I recommend opening up this finalized version of the experiment (<https://gallery.cortanaintelli...> and comparing yours. This version runs and deploys for me, so I'm guessing there is some subtle difference hiding between yours and it.

Brandon

^ | v · Reply · Share >



Charlie Loveland tbloveland@MSETI

Brandon Rohrer · 2 months ago

You can contribute the following items to the Gallery - follow these links for more information:

- [Collections](#)
- [Experiments](#)
- [Tutorials](#)

We want to hear from you!

We want the Gallery to be driven by our users and for our users. Use the smiley on the right to tell us what you love or hate about the Gallery.

Sign in

Send a smile Send a frown

What did you like?

<Provide your feedback here>...

969 character(s) left

[TAKE ME TO THE GALLERY >>](#)

Discover industry-specific solutions in the Cortana Intelligence Gallery

9/25/2017 • 1 min to read • [Edit Online](#)

Cortana Intelligence Gallery is a community-driven site for discovering and sharing solutions built with Cortana Intelligence Suite. The Gallery has a variety of resources that you can use to develop your own analytics solutions.

Industry-specific Cortana Intelligence solutions

The **Industries** section of the Gallery brings together various resources that are specific to such industries as

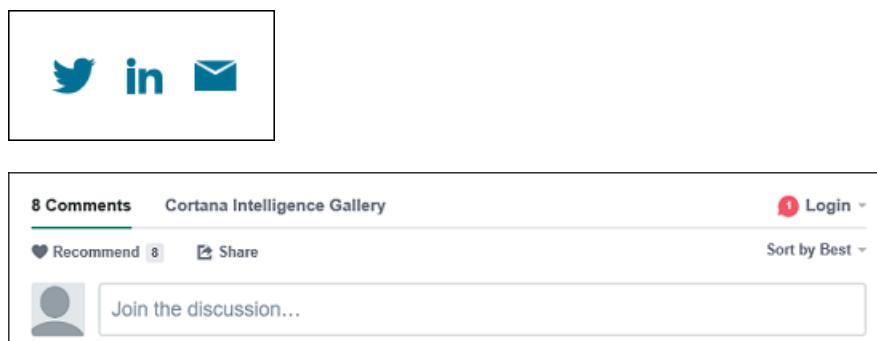
- **Retail** - Find retail solutions such as sales forecasting, predicting customer churn, and developing pricing models.
- **Manufacturing** - Find manufacturing solutions such as anticipating equipment maintenance and forecasting energy prices.
- **Banking** - Find banking solutions such as predicting credit risk and monitoring for online fraud.
- **Healthcare** - Find healthcare solutions such as detecting disease, and predicting hospital readmissions.

These resources include experiments, custom modules, APIs, collections, and any other Gallery items that can help you develop solutions specific to the industry you're working in.

Discover

To browse through the industry-specific solutions in the Gallery, open the [Gallery](#), point your mouse at **Industries** at the top of the Gallery home page, select a specific industry segment, or select **View All** to see an overview page for all industries.

Each industry page displays a list of the most popular Gallery items associated with that industry. Click **See all** to view all the industry-specific resources in the Gallery. From this page, you can browse all the resources in the Gallery. You also can search by selecting filter criteria on the left of the page and entering search terms at the top.



The screenshot shows the Cortana Intelligence Gallery interface. At the top, there are three social sharing icons: Twitter, LinkedIn, and Email. Below them is a comment section header with "8 Comments" and "Cortana Intelligence Gallery". To the right are "Login" and "Sort by Best" buttons. At the bottom is a text input field with a placeholder "Join the discussion..." and a user profile icon.

Click any Gallery item to open the item's details page for more information.

[TAKE ME TO THE GALLERY >>](#)

NOTE

Try Azure Machine Learning for free

No credit card or Azure subscription required. [Get started now.](#)

Discover Solutions in the Cortana Intelligence Gallery

9/25/2017 • 1 min to read • [Edit Online](#)

Cortana Intelligence Gallery is a community-driven site for discovering and sharing solutions built with Cortana Intelligence Suite. The Gallery has a variety of resources that you can use to develop your own analytics solutions.

Solutions

A **Solution** in the Gallery provides a jumpstart to quickly build Cortana Intelligence Solutions from preconfigured solutions, reference architectures, and design patterns. Make them your own with the included instructions or with a featured partner.

Discover

To browse for solutions in the Gallery, open the [Gallery](#) and click **Solutions** at the top of the Gallery home page.

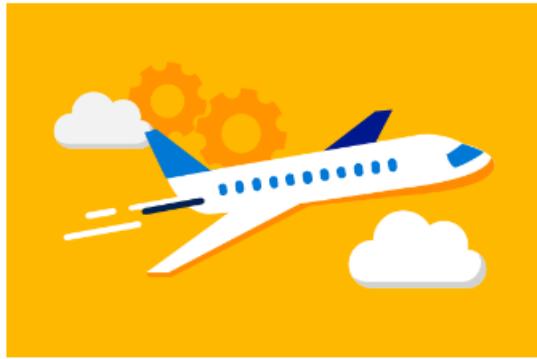
The [Solutions](#) page displays a list of the most recently added Solutions. Click **See all** to view all Solutions. From this page, you can browse all the Solutions in the Gallery. You also can search by selecting filter criteria on the left of the page and entering search terms at the top.

Click any Solution to open its details page and read more information. On this page you can comment, provide feedback, or ask questions through the comments section. You can even share it with friends or colleagues using the share capabilities of LinkedIn or Twitter. You can also email a link to the Solution to invite other users to view the page.

The screenshot shows a user interface for a solution page. At the top, there are three social sharing icons: Twitter, LinkedIn, and Email. Below these are sections for '8 Comments' and 'Cortana Intelligence Gallery'. There are buttons for 'Recommend' (with a count of 8) and 'Share'. On the right, there's a 'Login' button and a 'Sort by Best' dropdown. At the bottom, there's a placeholder for a profile picture and a text input field with the placeholder 'Join the discussion...'. The overall layout is clean and modern, typical of a web-based application.

Deploy

If you want to use a Solution, click **Deploy**. Follow the steps presented to configure and deploy the Solution in your workspace.



Deploy ↗

[TAKE ME TO THE GALLERY >>](#)

NOTE

Try Azure Machine Learning for free

No credit card or Azure subscription required. [Get started now.](#)

Discover experiments in Cortana Intelligence Gallery

9/25/2017 • 8 min to read • [Edit Online](#)

Cortana Intelligence Gallery is a community-driven site for discovering and sharing solutions built with Cortana Intelligence Suite. The Gallery has a variety of resources that you can use to develop your own analytics solutions.

Experiments for Machine Learning Studio

The Gallery has a wide variety of [experiments](#) that have been developed in [Azure Machine Learning Studio](#). Experiments range from quick proof-of-concept experiments that demonstrate a specific machine learning technique, to fully developed solutions for complex machine learning problems.

NOTE

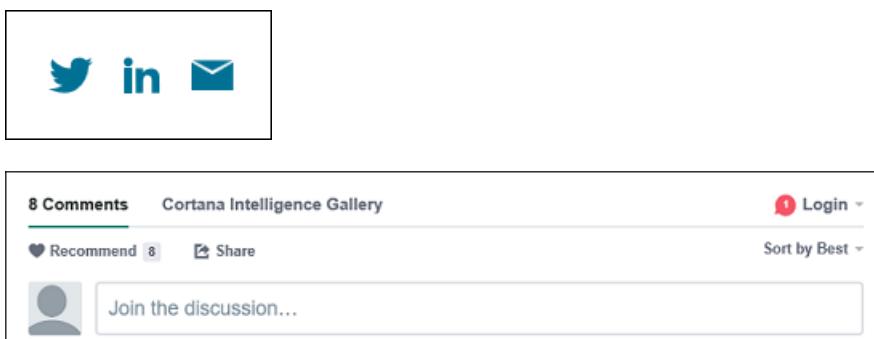
An **experiment** is a canvas in Machine Learning Studio that you can use to construct a predictive analysis model. You create the model by connecting data with various analytical modules. You can try different ideas, do trial runs, and eventually deploy your model as a web service in Azure. For an example of how to create a basic experiment, see [Machine learning tutorial: Create your first experiment in Azure Machine Learning Studio](#). For a more comprehensive walkthrough of how to create a predictive analytics solution, see [Walkthrough: Develop a predictive analytics solution for credit risk assessment in Azure Machine Learning](#).

Discover

To browse experiments [in the Gallery](#), at the top of the Gallery home page, select **Experiments**.

The **Experiments** page displays a list of recently added and popular experiments. To see all experiments, select the **See all** button. To search for a specific experiment, select **See all**, and then select filter criteria. You also can enter search terms in the **Search** box at the top of the Gallery page.

You can get more information about an experiment on the experiment details page. To open an experiment details page, select the experiment. On an experiment details page, in the **Comments** section, you can comment, provide feedback, or ask questions about the experiment. You can even share the experiment with friends or colleagues on Twitter or LinkedIn. You also can mail a link to the experiment details page, to invite other users to view the page.



The screenshot shows the Cortana Intelligence Gallery interface. At the top, there are three social sharing icons: Twitter, LinkedIn, and Email. Below them is a navigation bar with '8 Comments' on the left, the 'Cortana Intelligence Gallery' logo in the center, and a 'Login' button with a dropdown arrow on the right. Underneath the navigation bar, there are two buttons: 'Recommend 8' and 'Share'. To the right of these buttons is a 'Sort by Best' dropdown. At the bottom of the screenshot is a comment input field with a placeholder 'Join the discussion...' and a user profile icon to its left.

Download

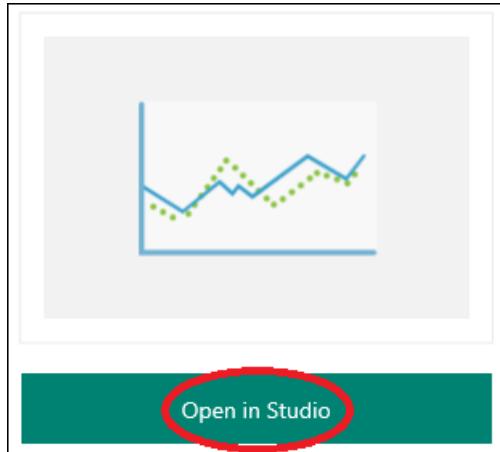
You can download a copy of any experiment in the Gallery to your Machine Learning Studio workspace. Then, you can modify your copy to create your own solutions.

Cortana Intelligence Gallery offers two ways to import a copy of an experiment:

- **From the Gallery.** If you find an experiment that you like in the Gallery, you can download a copy and then open it in your Machine Learning Studio workspace.
- **From within Machine Learning Studio.** In Machine Learning Studio, you can use any experiment in the Gallery as a template to create a new experiment.

From the Gallery

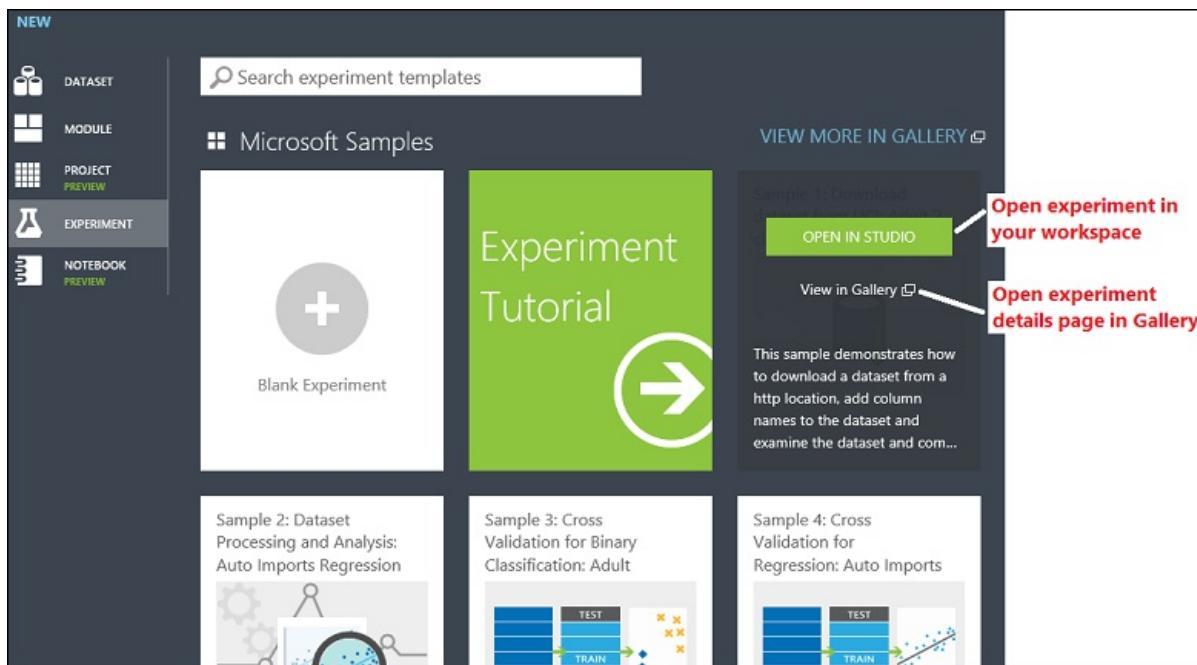
1. In the Gallery, open the experiment details page.
2. Select **Open in Studio**.



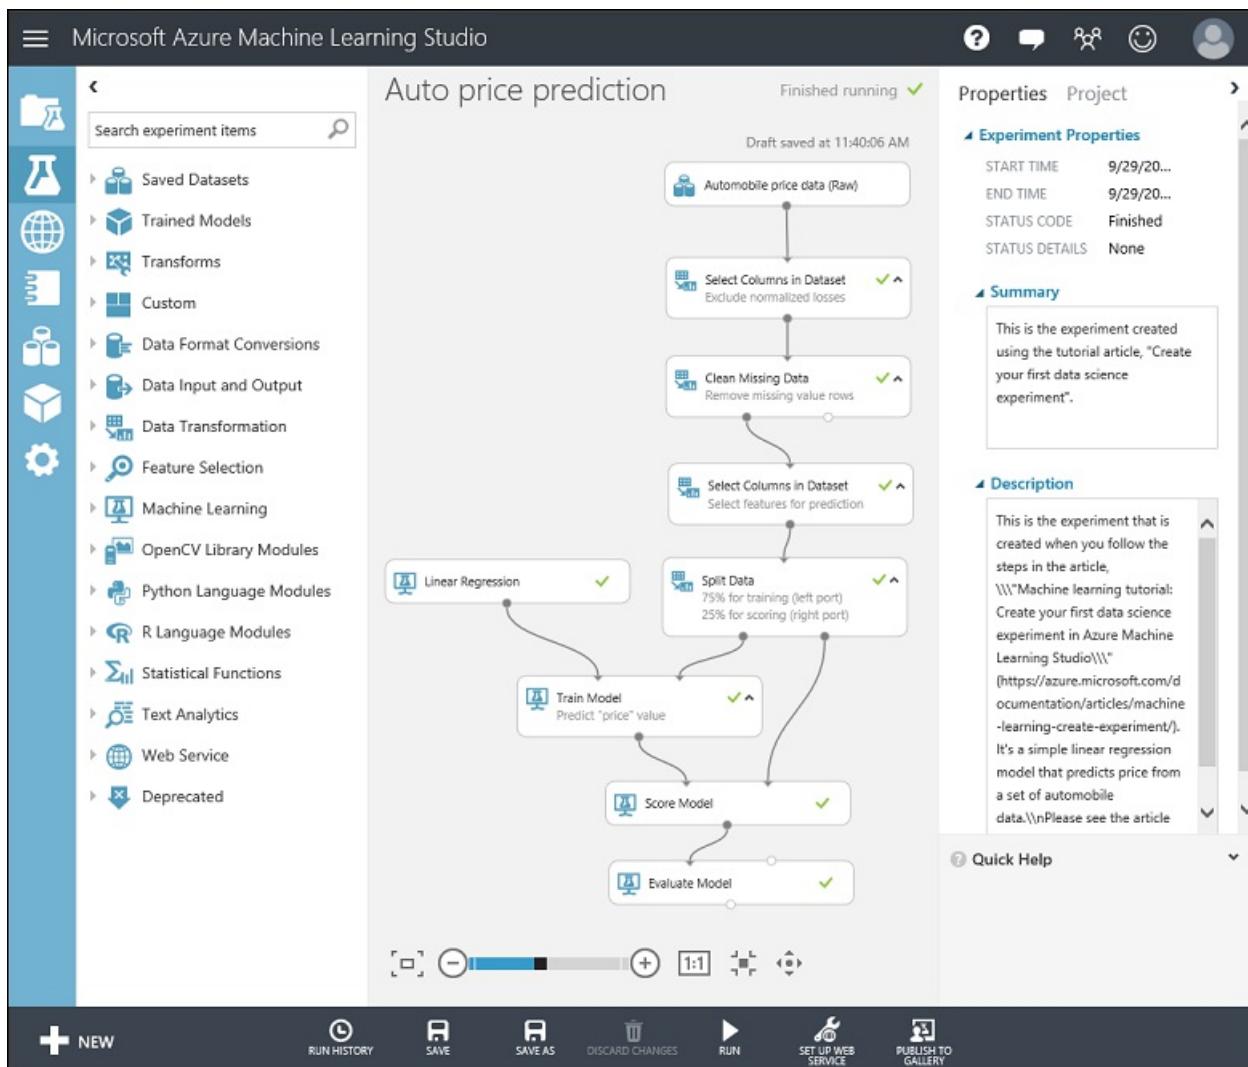
When you select **Open in Studio**, the experiment opens in your Machine Learning Studio workspace. (If you're not already signed in to Machine Learning Studio, you are prompted to first sign in by using your Microsoft account.)

From within Machine Learning Studio

1. In Machine Learning Studio, select **NEW**.
2. Select **Experiment**. You can choose from a list of Gallery experiments, or find a specific experiment by using the **Search** box.
3. Point your mouse at an experiment, and then select **Open in Studio**. (To see information about the experiment, select **View in Gallery**. This takes you to the experiment details page in the Gallery.)



You can customize, iterate, and deploy a downloaded experiment like any other experiment that you create in Machine Learning Studio.



Contribute

When you sign in to the Gallery, you become a member of the Gallery community. As a member of the community, you can contribute your own experiments, so other users can benefit from the solutions that you've discovered.

Publish your experiment to the Gallery

1. Sign in to Machine Learning Studio by using your Microsoft account.
2. Create your experiment, and then run it.
3. When you're ready to publish your experiment in the Gallery, in the list of actions below the experiment canvas, select **PUBLISH TO GALLERY**.



4. On the **Experiment Description** page, enter a title and tags. Make the title and tags descriptive. Highlight the techniques you used or the real-world problem you are solving. An example of a descriptive experiment title is "Binary Classification: Twitter Sentiment Analysis."

PUBLISH TO GALLERY X

Experiment Description

EXPERIMENT NAME

Binary Classification: Twitter Sentiment Analysis

TAG YOUR CONTENT

Press the enter key to create a tag

SUMMARY

This experiment demonstrates the use of the Execute R Script, Feature Selection and Feature Hashing modules to train a text sentiment classification engine.

DETAILED DESCRIPTION

M+ Markdown Preview

B **I** | | |

Enter a detailed description of your experiment

→ 2 3

5. In the **SUMMARY** box, enter a summary of your experiment. Briefly describe the problem the experiment solves, and how you approached it.
6. In the **DETAILED DESCRIPTION** box, describe the steps you took in each part of your experiment. Some useful topics to include are:
 - Experiment graph screenshot
 - Data sources and explanation
 - Data processing
 - Feature engineering
 - Model description
 - Results and evaluation of model performance

You can use markdown to format your description. To see how your entries on the experiment description page will look when the experiment is published, select **Preview**.

M+ Markdown Preview

B **I** | | |

Enter a detailed description of your experiment

TIP

The text boxes provided for markdown editing and preview are small. We recommend that you write your experiment documentation in a markdown editor, copy it, and then paste the completed documentation in the text box in the Gallery. After you publish your experiment, you can make any corrections by using standard web-based tools that use markdown for edit and preview.

7. On the **Image Selection** page, choose a thumbnail image for your experiment. The thumbnail image appears at the top of the experiment details page and in the experiment tile. Other users will see the thumbnail image when they browse the Gallery. You can upload an image from your computer, or select a stock image from the Gallery.

PUBLISH TO GALLERY X

Image Selection

Use a JPG, GIF, or PNG. A resolution of 960 by 640 pixels works best. Load image from Gallery



UPLOAD AN IMAGE **SELECT STOCK IMAGE**

 BROWSE YOUR COMPUTER...

8. On the **Settings** page, under **Visibility**, choose whether to publish your content publicly (**Public**) or to have it accessible only to people who have a link to the page (**Unlisted**).

PUBLISH TO GALLERY

Settings

Visibility

PUBLIC
Anyone can view and copy the experiment. Once public, it cannot be unlisted.

UNLISTED
Only people with the direct link can view and copy the experiment.

TIP

If you want to make sure your documentation looks correct before you release it publicly, you can first publish the experiment as **Unlisted**. Later, you can change the visibility setting to **Public** on the experiment details page.

9. To publish the experiment to the Gallery, select the **OK** check mark.



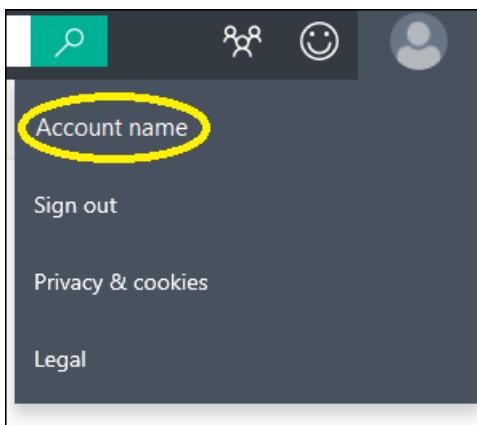
For tips on how to publish a high-quality Gallery experiment, see [Tips for documenting and publishing your](#)

experiment.

That's it--you're all done.

You can now view your experiment in the Gallery, and share the link with others. If you published your experiment by using the **Public** visibility setting, the experiment shows up in browse and search results in the Gallery. You can edit your experiment documentation on the experiment details page any time you are signed in to the Gallery.

To see the list of your contributions, select your image in the upper-right corner of any Gallery page. Then, select your name to open your account page.



Update your experiment

If you need to, you can make changes to the workflow (modules, parameters, and so on) in an experiment that you published to the Gallery. In Machine Learning Studio, make any changes you'd like to make to the experiment, and then publish again. Your published experiment will be updated with your changes.

You can change any of the following information for your experiment directly in the Gallery:

- Experiment name
- Summary or description
- Tags
- Image
- Visibility setting (**Public** or **Unlisted**)

You also can delete the experiment from the Gallery.

You can make these changes, or delete the experiment, from the experiment details page or from your profile page in the Gallery.

From the experiment details page

On the experiment details page, to change the details for your experiment, select **Edit**.



The details page enters edit mode. To make changes, select **Edit** next to the experiment name, summary, or tags. When you're finished making changes, select **Done**.

EXPERIMENT

Your first data science experiment - Automobile price prediction



Gary Ericson • September 29, 2016

Done



Summary



This is the experiment created using the tutorial article, "Create your first data science experiment".

Description



This is the experiment that's created when you follow the steps in the article, **Machine learning tutorial: Create your first data science experiment in Azure Machine Learning Studio**.

It's a simple linear regression model that predicts price from a set of automobile data.

Please see the article for more information on how this experiment was developed.

To change the visibility settings for the experiment (**Public** or **Unlisted**), or to delete the experiment from the Gallery, select the **Settings** icon.

EXPERIMENT

Your first data science experiment - Automobile price prediction



Gary Ericson • September 29, 2016



Summary

Change Visibility

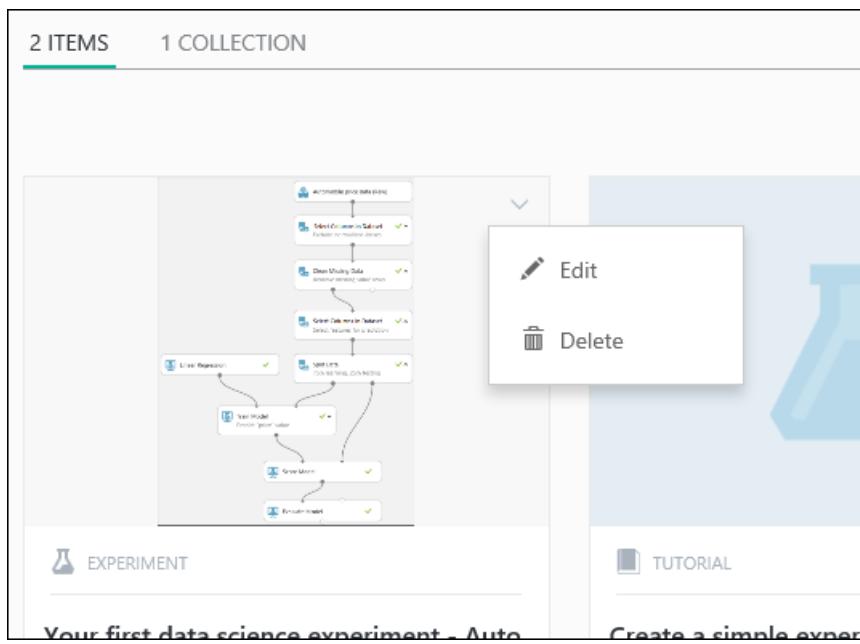
Delete

This is the experiment created using the tutorial article, "Create your first

From your profile page

On your profile page, select the down arrow for the experiment, and then select **Edit**. This takes you to the details page for your experiment, in edit mode. When you are finished making changes, select **Done**.

To delete the experiment from the Gallery, select **Delete**.



Tips for documenting and publishing your experiment

- You can assume that the reader has prior data science experience, but it can be helpful to use simple language. Explain things in detail whenever possible.
- Cortana Intelligence Suite is relatively new. Not all readers are familiar with how to use it. Provide enough information and step-by-step explanations to help readers navigate your experiment.
- Visuals can be helpful for readers to interpret and use your experiment documentation correctly. Visuals include experiment graphs and screenshots of data. For more information about how to include images in your experiment documentation, see the [Publishing Guidelines and Examples collection](#).
- If you include a data set in your experiment (that is, you're not importing the dataset through the Import Data module), the data set is part of your experiment and is published in the Gallery. Make sure that the data set you publish has licensing terms that allow sharing and downloading by anyone. Gallery contributions are covered under the Azure [Terms of Use](#).

Frequently asked questions

What are the requirements for submitting or editing an image for my experiment?

Images that you submit with your experiment are used to create an experiment tile for your contribution. We recommend that images be smaller than 500 KB, with an aspect ratio of 3:2, and a resolution of 960 × 640.

What happens to the data set I used in the experiment? Is the data set also published in the Gallery?

If your data set is part of your experiment and is not being imported through the Import Data module, the data set is published in the Gallery as part of your experiment. Make sure that the data set that you publish with your experiment has the appropriate licensing terms. The licensing terms should allow anyone to share and download the data.

I have an experiment that uses an Import Data module to pull data from Azure HDInsight or SQL Server. It uses my credentials to retrieve the data. Can I publish this kind of experiment? How can I be assured that my credentials won't be shared?

Currently, you cannot publish an experiment that uses credentials in the Gallery.

How do I enter multiple tags?

After you enter a tag, to enter another tag, press the Tab key.

[Go to the Gallery](#)

NOTE

Try Azure Machine Learning for free

No credit card or Azure subscription required. [Get started now.](#)

Discover Jupyter Notebooks in the Cortana Intelligence Gallery

9/25/2017 • 2 min to read • [Edit Online](#)

Cortana Intelligence Gallery is a community-driven site for discovering and sharing solutions built with Cortana Intelligence Suite. The Gallery has a variety of resources that you can use to develop your own analytics solutions.

Jupyter Notebooks

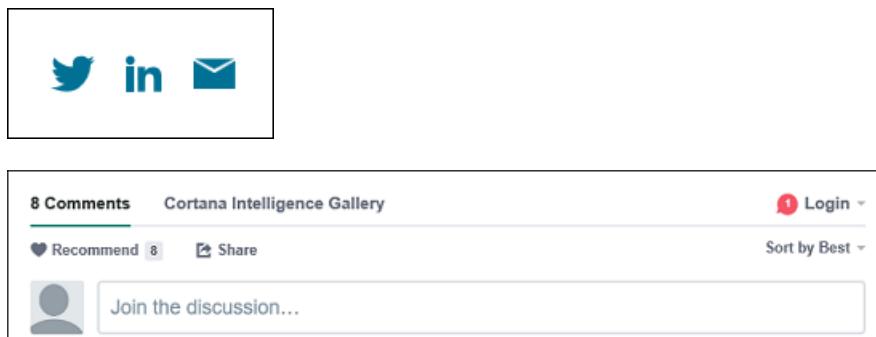
Jupyter Notebooks include code, data visualizations, and documentation in a single, interactive canvas. Notebooks in the Gallery provide tutorials and detailed explanations of advanced machine learning techniques and solutions.

Discover

To browse for notebooks in the Gallery, open the [Gallery](#) and click **Notebooks** at the top of the Gallery home page.

The [Jupyter Notebooks](#) page displays a list of the most popular notebooks. Click **See all** to view all notebooks. From this page, you can browse all the notebooks in the Gallery. You also can search by selecting filter criteria on the left of the page and entering search terms at the top.

Click any notebook to open the notebook's details page and read more information. On this page you can comment, provide feedback, or ask questions through the comments section. You can even share it with friends or colleagues using the share capabilities of LinkedIn or Twitter. You can also email a link to the notebook to invite other users to view the page.



A screenshot of a web browser showing a notebook details page. At the top, there are social sharing icons for Twitter, LinkedIn, and Email. Below that, the page header includes '8 Comments', 'Cortana Intelligence Gallery', and a 'Login' button. There are also 'Recommend' and 'Share' buttons. A large text area contains the notebook's content, with a placeholder 'Join the discussion...' at the bottom. The overall layout is clean and modern.

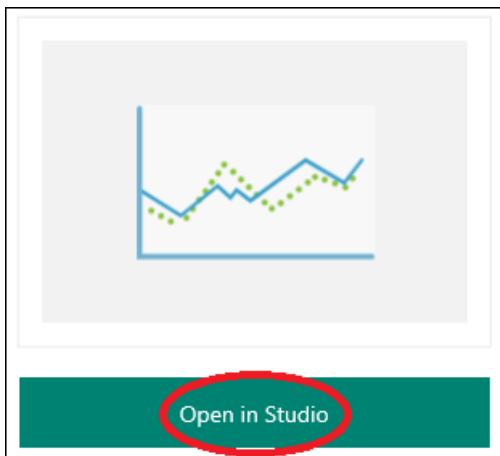
Download

You can download a copy of any notebook from the Gallery into your Machine Learning Studio workspace.

Download from the Gallery

To download a copy of a notebook from the Gallery:

1. Open the notebook's details page in the Gallery
2. Click **Open in Studio**
3. Select the Region and Workspace you want to use

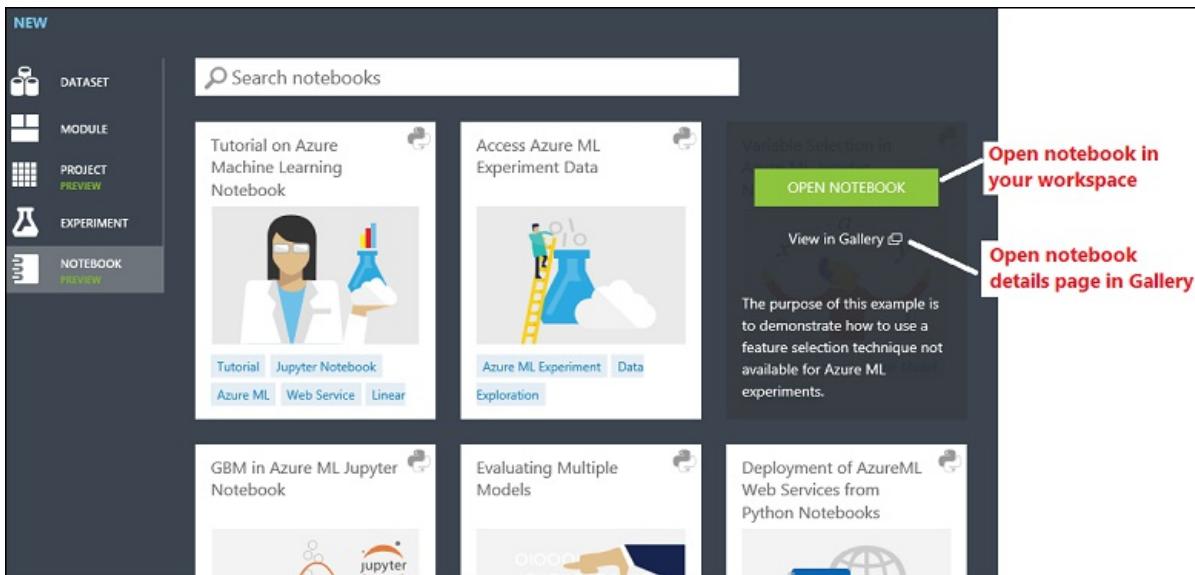


When you click **Open in Studio**, the notebook is loaded into your Machine Learning Studio workspace and opened (if you're not already signed in to Studio, you will be prompted to sign in using your Microsoft account before the notebook is copied to your workspace). You can find the notebook again later by clicking **Notebooks** on the left side of the Studio home page.

Download in Machine Learning Studio

You can also open any of the same Gallery notebooks while you're working in Machine Learning Studio:

1. In Machine Learning Studio, click **+NEW**
2. Select **Notebook** - a list of notebooks from the Gallery is displayed that you can choose from, or you can find a specific notebook using the search box
3. Point your mouse at a notebook and select **Open Notebook**



A copy of the notebook is downloaded and opened in your workspace in the Jupyter Notebooks section of Machine Learning Studio. The notebook will be listed, along with your other notebooks, on the **Notebooks** page of Studio (on the [Studio home page](#), click **Notebooks** on the left).

[TAKE ME TO THE GALLERY >>](#)

NOTE

[Try Azure Machine Learning for free](#)

No credit card or Azure subscription required. [Get started now.](#)

Discover competitions in the Cortana Intelligence Gallery

9/25/2017 • 1 min to read • [Edit Online](#)

Cortana Intelligence Gallery is a community-driven site for discovering and sharing solutions built with Cortana Intelligence Suite. The Gallery has a variety of resources that you can use to develop your own analytics solutions.

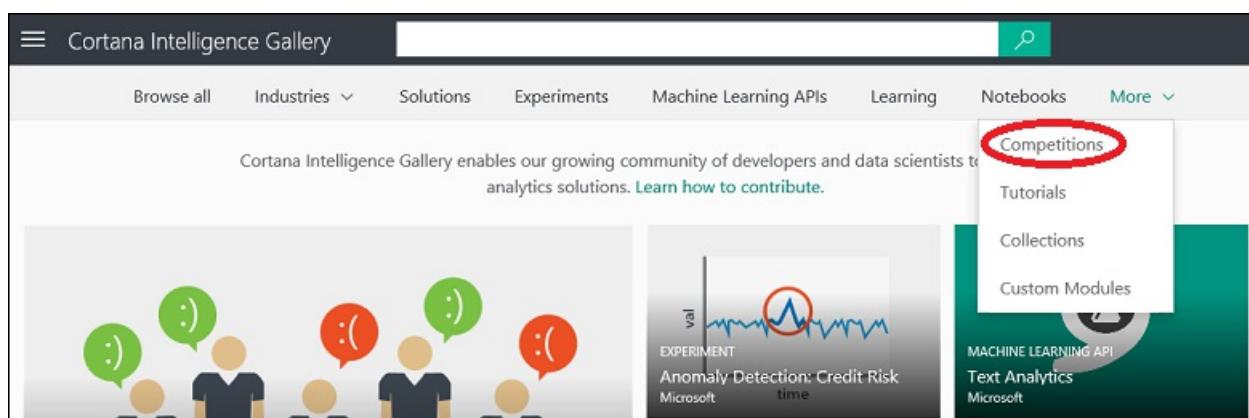
Cortana Intelligence Competitions

Competitions provide an exciting opportunity to compete with the community of data scientists to solve complex problems using Cortana Intelligence Suite.

For answers to frequently-asked questions about Competitions, see [Microsoft Cortana Intelligence Competitions FAQ](#).

Discover

To browse for competitions in the Gallery, open the [Gallery](#), point your mouse to **More** at the top of the Gallery home page, and select **Competitions**.



The **Competitions** page displays a list of the most popular competitions. Click **See all** to view all competitions. From this page, you can browse all the competitions in the Gallery. You also can search by selecting filter criteria on the left of the page and entering search terms at the top.

Click any competition to open the competition's details page and read more information. On this page you can comment, provide feedback, or ask questions through the comments section. You can even share it with friends or colleagues using the share capabilities of LinkedIn or Twitter. You can also email a link to the competition to invite other users to view the page.



Enter a competition

If the competition is open, the status on the details page is **Active**. To enter the competition, click **Enter Competition** and you will be given instructions on the resources you need and the steps you need to take.

COMPETITION

Women's Health Risk Assessment

Sponsored by Microsoft • ends 10/5/2016

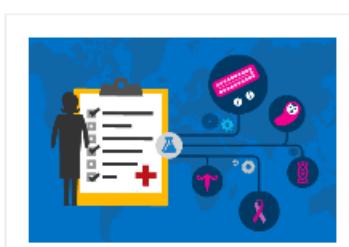
ACTIVE Microsoft Band 2 48 participants

WINNERS

1st	Farhad Ghassemi	2nd	Subhojit Som
	Xbox		worse xbox

Started 9/28/2016 12:32:50 PM (Pacific Daylight Time)
Ends 10/5/2016 12:32:50 PM (Pacific Daylight Time)

[Summary](#) [Frequently Asked Questions](#)



Enter Competition

2258 views
48 participants
322 submissions

If the competition has been completed, then a list of winners is shown.

[TAKE ME TO THE GALLERY >>](#)

NOTE

Try Azure Machine Learning for free

No credit card or Azure subscription required. [Get started now](#).

Microsoft Cortana Intelligence Competitions FAQ

9/25/2017 • 7 min to read • [Edit Online](#)

What is Cortana Intelligence Competitions?

The Microsoft Cortana Intelligence Competitions allow us to unite a global community of data enthusiasts by collectively solving some of the world's most complex data science problems. Cortana Intelligence Competitions allow data enthusiasts from across the world to compete and build highly accurate and intelligent data science models. Our hosted competitions are based on unique data sets that have been made available publically for the first time. Participants can win rewards or get recognition via our top 10 public leaderboard. Please go [here](#) to access the Competitions home page.

How often will Microsoft release new competitions?

We will be launching 1st-party, Microsoft-owned competitions on a regular basis, approximately every 8-12 weeks.

Where can I ask general questions about data science?

Please use our [Microsoft Azure Machine Learning forum](#).

How do I enter a competition?

Access the [Competitions](#) home page in the [Cortana Intelligence Gallery](#), or go to <http://aka.ms/CIComp>. The home page lists all competitions that are currently running. Each competition has detailed instructions and participation rules, prizes, and duration on its sign-up page.

1. Find the competition you'd like to participate in, read all the instructions and watch the tutorial video, then click the **Enter Competition** button to copy the Starter Experiment into your existing Azure Machine Learning workspace. If you don't already have access to a workspace, you must create one beforehand. Run the Starter Experiment, observe the performance metric, then use your creativity to improve the performance of the model. You'll likely spend the majority of your time in this step.
2. Create a Predictive Experiment with the trained model out of your Starter Experiment. Then carefully adjust the input and output schema of the web service to ensure they conform to the requirements specified in the competition documentation. The tutorial document generally will have detailed instructions for this. You can also watch the tutorial video, if available.
3. Deploy a web service out of your Predictive Experiment. Test your web service using the **Test** button or the Excel template automatically created for you to ensure it's working properly.
4. Submit your web service as the competition entry and see your public score in the Cortana Intelligence Gallery competition page. And celebrate if you make it onto the leaderboard!

After you successfully submit an entry, you can go back to your copied Starter Experiment, iterate, and update your Predictive Experiment, update the web service, and submit a new entry.

Can I use open source tools for participating in these competitions?

The competition participants leverage Azure Machine Learning Studio, a cloud-based service within Cortana Intelligence Suite for development of the data science models and to create competition entries for submission. Machine Learning Studio not only provides a GUI interface for constructing machine learning experiments, it also allows you to bring your own R and/or Python scripts for native execution. The R and Python runtimes in Studio come with a rich set of open source R/Python packages, and you can import your own packages as part of the experiment as well. Studio also has a built-in Jupyter Notebook service for you to do free style data exploration. Of

course, you can always download the datasets used in the competition and explore it in your favorite tool outside of Machine Learning Studio.

Do I need to be a data scientist to enter?

No. In fact, we encourage data enthusiasts, those curious about data science, and other aspiring data scientists to enter our contest. We have designed supporting documents to allow everyone to compete. Our target audience is:

- **Data Developers, Data Scientists, BI and Analytics Professionals:** those who are responsible for producing data and analytics content for others to consume
- **Data Stewards:** those who have the knowledge about the data, what it means, and how it's intended to be used and for which purpose
- **Students & Researchers:** those who are learning and gaining data related skills via academic programs in universities, or participants in Massive Open Online Courses (MOOCs)

Can I enter with my colleagues as a team?

The competition platform currently doesn't support team participation. Each competition entry is tied to a single user identity.

Do I need to pay to participate in a competition?

Competitions are free to participate in. You do, however, need access to an Azure Machine Learning workspace to participate. You can create a Free workspace without a credit card by simply logging in with a valid Microsoft account, or an Office 365 account. If you're already an Azure or Cortana Intelligence Suite customer, you can create and use a Standard workspace under the same Azure subscription. If you'd like to purchase an Azure subscription, go to the [Azure pricing](#) page. Note that the standard rates will apply when using a Standard workspace to construct experiments. See [Azure Machine Learning pricing information](#) for more information.

NOTE

[Try Azure Machine Learning for free](#)

No credit card or Azure subscription required. [Get started now.](#)

What are public and private scores?

In most competitions, you'll notice that you'll receive a public score for every submission you make, generally within 10-20 minutes. But after the competition ends, you'll receive a private score which is used for final ranking.

Here's what happens:

- The entire dataset used in the competition is randomly split with stratification into training and testing (the remaining) data. The random split is stratified to ensure that the distributions of labels in both training and testing data are consistent.
- The training data is uploaded and given to you as part of the Starter Experiment in the Import Data module configuration.
- The testing data is further split into public and private testing data, using the same stratification.
- The public testing data is used for the initial round of scoring. The result is referred to as the public score and it's what you see in your submission history when you submit your entry. This score is calculated for every entry you submit. This public score is used to rank you on the public leaderboard.
- The private testing data is used for the final round of scoring after the competition ends. This is referred to as private score.
- For each participant, a fixed number of entries with the highest public scores are automatically selected to enter the private scoring round (this number can vary depending on the competition). The entry with the highest private score is then selected to enter the final ranking, which ultimately determines the prize winners.

Can customers host a competition on our platform?

We welcome 3rd-party organizations to partner with us and host both public and private competitions on our platform. We have a competition onboarding team who will be happy to discuss the onboarding process for such competitions. Please get in touch with us at compsupport@microsoft.com for more details.

What are the limitations for submissions?

A typical competition may choose to limit the number of entries you can submit within a 24-hour span (UTC time 00:00:00 to 23:59:59), and the total number of entries you can submit over the duration of the competition. You'll receive appropriate error messages when a limitation is exceeded.

What happens if I win a competition?

Microsoft will verify the results of the private leaderboard, and then we'll contact you. Please make sure that your email address in your user profile is up to date.

How will I get the prize money if I win a competition?

If you are a competition winner, you will need to sign a declaration of eligibility, license, and release from. This form reiterates the Competition Rules. Winners need to fill out a US Tax form W-9, or a Form W-8BEN if they are not US taxpayers. We will contact all winners as part of the rewards disbursement process by using their registration email. Please refer to our [Terms and Conditions](#) for additional details.

What if more than one entry receives the same score?

The submission time is the tie-breaker. The entry submitted earlier outranks the entry submitted later.

Can I participate using guest workspace?

No. You must use a Free or a Standard workspace to participate. You can open the competition starter experiment in a Guest workspace, but you'll not be able to create a valid entry for submission from that workspace.

Can I participate using a workspace in any Azure region?

Currently, the competition platform only supports entries submitted from a workspace in the **South Central US** Azure region. All Free workspaces reside in South Central US, so you can submit an entry from any Free workspace. If you choose to use a Standard workspace, just ensure that it resides in the South Central US Azure region, otherwise your submission won't succeed.

Do we keep users' competitions solutions and entries?

User entries are only retained for evaluation purposes to identify the winning solutions. Please refer to our [Terms and Conditions](#) for specifics.

Discover and share tutorials in Cortana Intelligence Gallery

9/25/2017 • 2 min to read • [Edit Online](#)

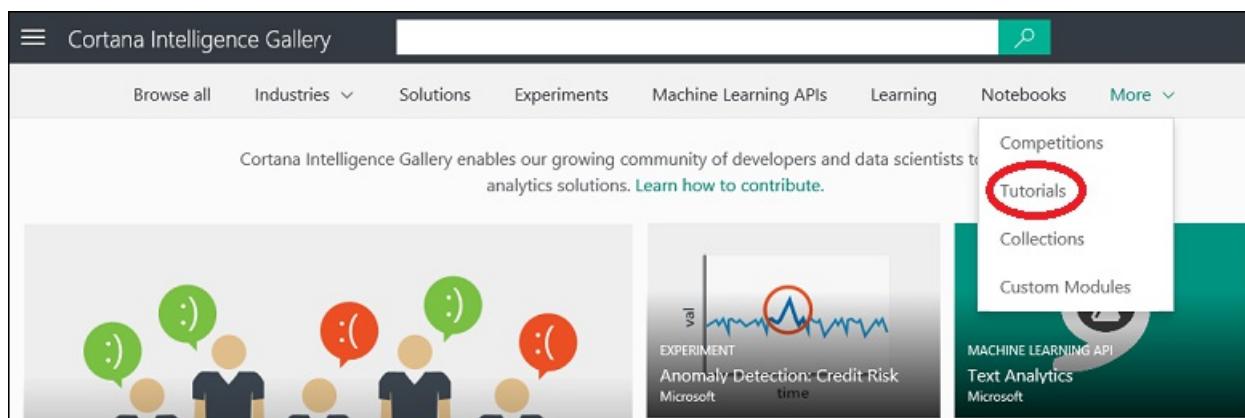
Cortana Intelligence Gallery is a community-driven site for discovering and sharing solutions built with Cortana Intelligence Suite. The Gallery has a variety of resources that you can use to develop your own analytics solutions.

Tutorials

Cortana Intelligence Gallery offers several [tutorials](#) that can walk you through machine learning technologies and concepts. Tutorials also might describe advanced methods for solving machine learning problems.

Discover

To browse tutorials [in the Gallery](#), under **More**, select **Tutorials**.



The [Tutorials](#) page displays a list of recently added and popular tutorials. To see all tutorials, select the **See all** button. To search for a specific tutorial, select **See all**, and then select filter criteria. You also can enter search terms in the **Search** box at the top of the Gallery page.

You can get more information about a tutorial on the tutorial details page. To open a tutorial details page, select the tutorial. In the **Comments** section on a tutorial details page, you can comment, provide feedback, or ask questions about the tutorial. You can even share the tutorial with friends or colleagues on Twitter or LinkedIn. You also can email a link to the tutorial details page, to invite other users to view the page.

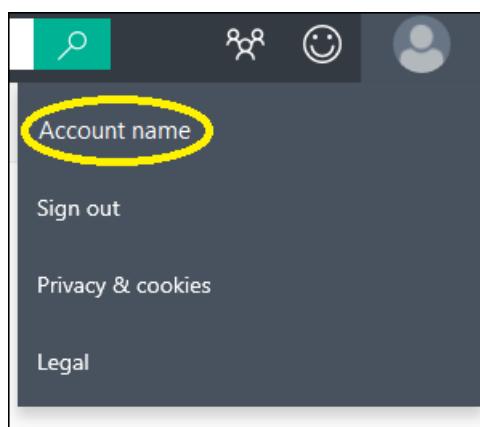
A screenshot of a tutorial details page. At the top, there are social sharing icons for Twitter, LinkedIn, and Email. Below that, it shows '8 Comments' and the 'Cortana Intelligence Gallery' logo. There are buttons for 'Login' and 'Sort by Best'. Underneath, there are sections for 'Recommend' (with 8 likes) and 'Share'. A large input field at the bottom is labeled 'Join the discussion...'. The overall layout is clean and modern.

Contribute

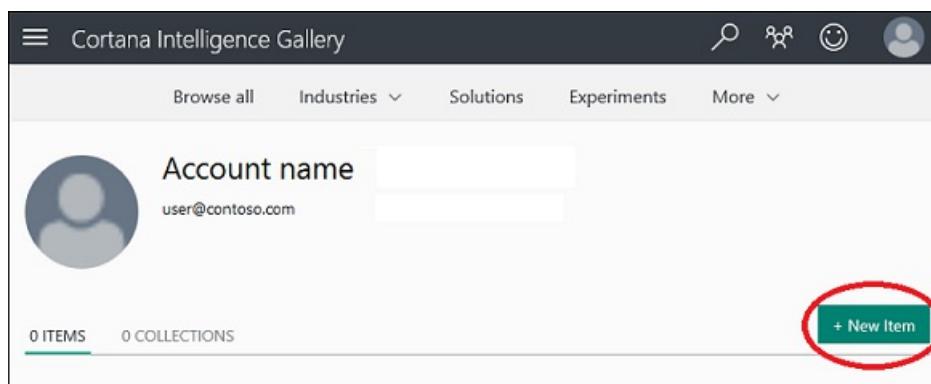
You can add a tutorial to the Gallery to help other users solve a problem or learn a concept.

Create a tutorial

1. Sign in to the Gallery by using your Microsoft account.
2. Select your image at the top of the page, and then select your name.



3. Select **New Item**.



4. On the **Description** page, for **ITEM TYPE**, select **Tutorial**. Enter a name for your tutorial, a brief summary, and any tags that might help other users find your tutorial.

CREATE NEW GALLERY ITEM

Description

ITEM TYPE

Tutorial

NAME

Create a simple experiment in Azure Machine Learning Studio

SUMMARY

Follow a simple step-by-step process to create your first experiment in Azure Machine Learning Studio.

DETAILED DESCRIPTION

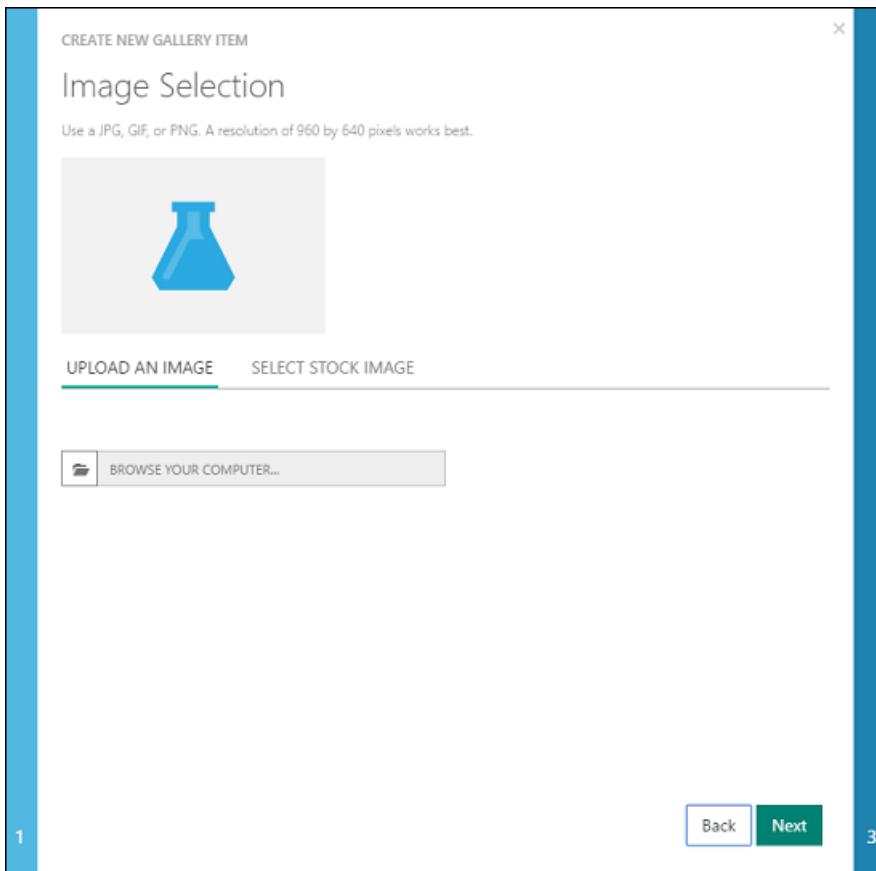
Markdown Preview

B I | | | | |

TAGS

Next 2 3

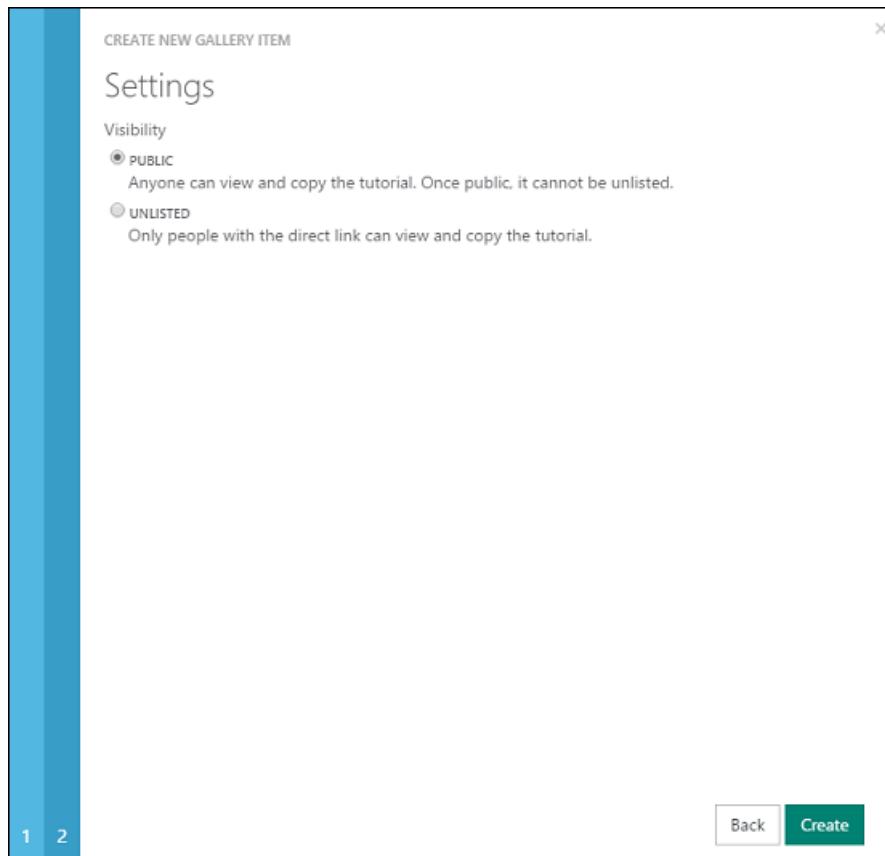
5. On the **Image Selection** page, select an image that is displayed with the collection. You can upload your own image file, or select a stock image. Choose an image that might help users identify the content and purpose of your tutorial.



6. On the **Settings** page, for **Visibility**, select whether the collection is **Public** (anyone can view the tutorial) or **Unlisted** (only people with a direct link to the tutorial can view it).

IMPORTANT

After you set a tutorial to **Public** and then select **Create**, you cannot later change the tutorial to **Unlisted**.



7. Select **Create**.

Your tutorial is now in Cortana Intelligence Gallery. Your tutorials are listed on your account page, on the **Items** tab.

[Go to the Gallery](#)

NOTE

[Try Azure Machine Learning for free](#)

No credit card or Azure subscription required. [Get started now.](#)

Discover and share collections in Cortana Intelligence Gallery

9/25/2017 • 3 min to read • [Edit Online](#)

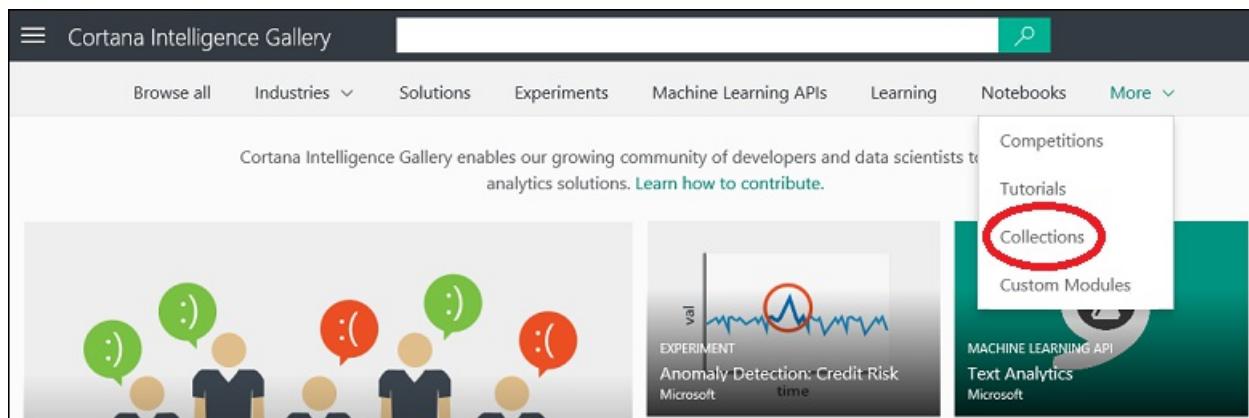
Cortana Intelligence Gallery is a community-driven site for discovering and sharing solutions built with Cortana Intelligence Suite. The Gallery has a variety of resources that you can use to develop your own analytics solutions.

Collections

You can use a **collection** to group experiments, APIs, and other Gallery items that you use for a specific solution or concept. You can group them to use, to save for later reference, or to share.

Discover

To browse collections [in the Gallery](#), under **More**, select **Collections**.



The [Collections](#) page displays a list of recently added and popular collections. To view all collections, select **See all**. To search for a specific collection, select **See all**, and then select filter criteria. You also can enter search terms in the **Search** box at the top of the Gallery page.

You can get more information about a collection on the collection details page. To open a collection details page, select the collection. On a collection details page, in the **Comments** section, you can comment, provide feedback, or ask questions about the collection. You can even share the collection with friends or colleagues on Twitter or LinkedIn. You also can email a link to the collection details page, to invite other users to view the page.

Two screenshots of the Cortana Intelligence Gallery interface. The top screenshot shows a light gray box containing social sharing icons for Twitter, LinkedIn, and Email. The bottom screenshot shows a collection details page with a header showing '8 Comments' and 'Cortana Intelligence Gallery', a login button, and sorting options. Below the header are 'Recommend' and 'Share' buttons, followed by a list of comments and a 'Join the discussion...' input field.

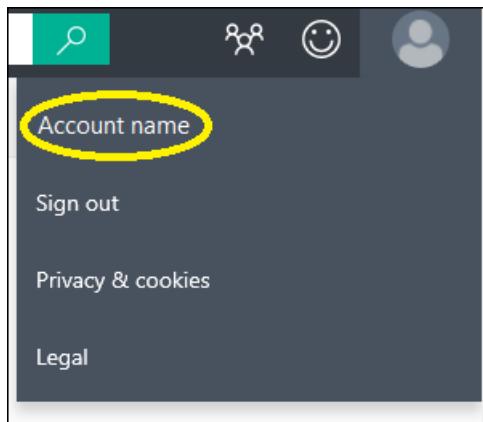
Contribute

You can create a collection that has items that you own or items that have been contributed by the community. An item in a collection can be any Gallery item that addresses a specific solution or concept.

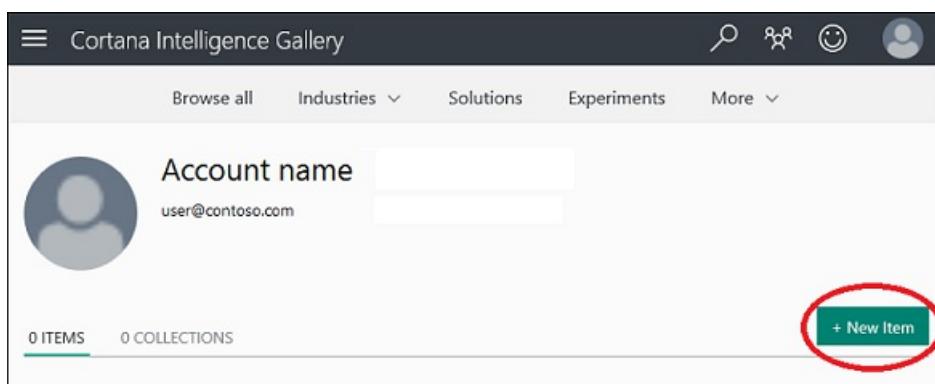
For example, you can use a collection to group items on a specific topic, or you can group a multi-step experiment that solves a complex problem. The initial collections contributed by Microsoft consist of multi-step machine learning experiment templates. You can use the templates to help solve real-world problems, such as online fraud detection, text classification, retail forecasting, or predictive maintenance.

Create a collection

1. Sign in to the Gallery by using your Microsoft account.
2. Select your image at the top of the page, and then select your name.



3. Select **New Item**.



4. On the **Description** page, for **ITEM TYPE**, select **Collection**. Enter a name for your collection, a brief summary, a detailed description, and any tags that might help other users find your collection.

CREATE NEW GALLERY ITEM

Description

ITEM TYPE

Collection

NAME

Linear regression examples

SUMMARY

Collection of various experiments, tutorials, etc. related to linear regression

DETAILED DESCRIPTION

Markdown Preview

B *I* | “ | | |

TAGS

Next **2**

5. On the **Image Selection** page, select an image that is displayed with the collection. You can upload your own image file, or select a stock image. Choose an image that might help users identify the content and purpose of the collection.

CREATE NEW GALLERY ITEM

Image Selection

Use a JPG, GIF, or PNG. A resolution of 960 by 640 pixels works best.

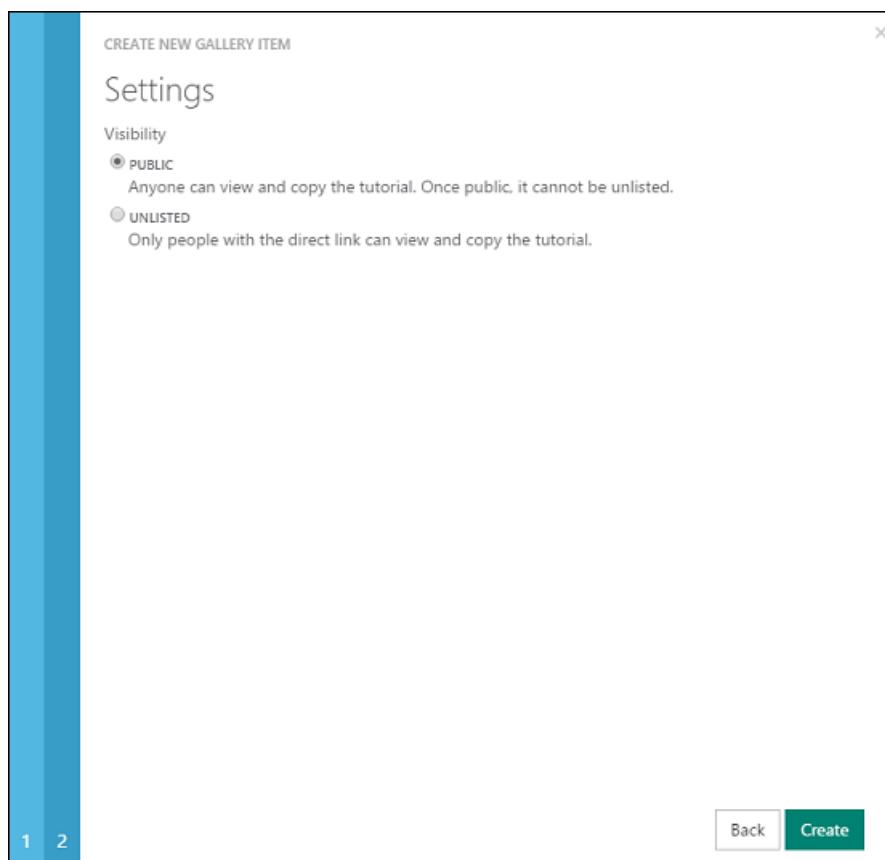
UPLOAD AN IMAGE [SELECT STOCK IMAGE](#)

[Back](#) **Next** **3**

6. On the **Settings** page, for **Visibility**, select whether the collection is **Public** (anyone can view the collection) or **Unlisted** (only people with a direct link to the collection can view it).

IMPORTANT

After you set a collection to **Public** and then select **Create**, you cannot later change the collection to **Unlisted**.



7. Select **Create**.

Your collection is now in Cortana Intelligence Gallery. Your collections are listed on your account page, on the **Collection** tab.

Add items to a collection

To add items to your collection, open the collection, select **Edit**, and then select **Add Item**.

A random collection of things [Edit](#)

Gary Ericson • September 29, 2016

[Done](#) [Settings](#)

[Add Item](#)

[Summary](#) [Edit](#)

A list of items you've contributed to the Gallery is shown on your account page. You also can search the Gallery for items to add. Select the items that you want to add. Every item that you select is included in the set of items to add. The **Add** button indicates how many items you have selected.

Add to collection

YOUR ITEMS SEARCH GALLERY

Your first data science experiment - Automobile price prediction
EXPERIMENT by Gary Ericson

Create a simple experiment in Azure Machine Learning Studio
TUTORIAL by Gary Ericson

A random collection of things
COLLECTION by Gary Ericson

CANCEL ADD

If you find an item that you want to add while you are browsing the Gallery, to open the item, select **Add to collection**. Select the collection that you want to add the item to.

Open in Studio

+ Add to Collection

12915 views

To change the summary, description, or tags for your collection, open the collection, and then select **Edit**.

While you are editing your collection, you can change the order of the items in the collection. Use the arrow buttons next to an item to move it up or down in the order list.

To add notes to the items in your collection, select the upper-right corner of an item, and then select **Add/Edit note**.

To remove an item from your collection, select **Remove**.

A random collection of things [Edit](#)

Gary Ericson • September 29, 2016

Done [Edit](#)

Add Item

Summary [Edit](#)

Description [Edit](#)

4 Items

Your first data science experiment - Automobile price prediction
EXPERIMENT by Gary Ericson

Womens Health Risk Assessment
COMPETITION by Microsoft

HEART DISEASE

Add/Edit note

Use as collection image

Remove

RELATED ITEMS

Go to the Gallery

NOTE

Try Azure Machine Learning for free

No credit card or Azure subscription required. [Get started now.](#)

Discover custom machine learning modules in Cortana Intelligence Gallery

9/25/2017 • 3 min to read • [Edit Online](#)

Cortana Intelligence Gallery is a community-driven site for discovering and sharing solutions built with Cortana Intelligence Suite. The Gallery has a variety of resources that you can use to develop your own analytics solutions.

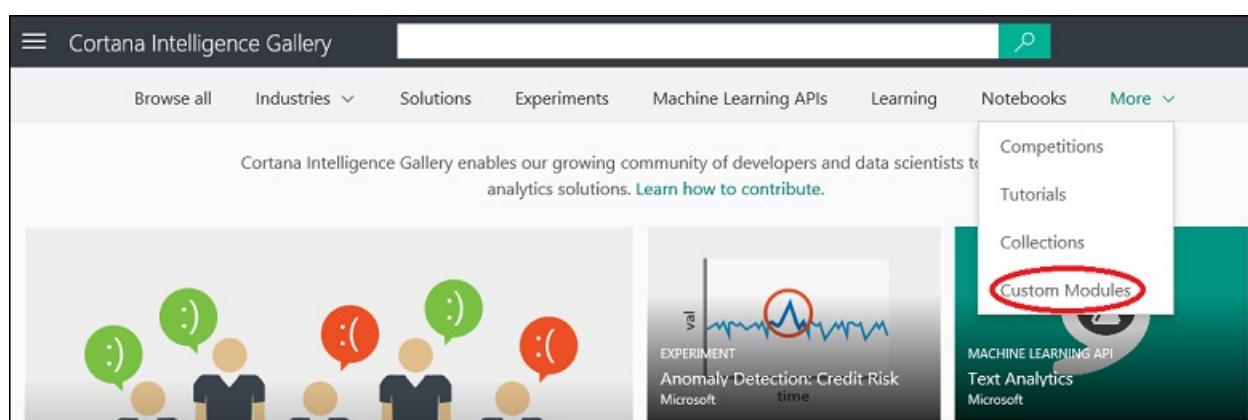
Custom modules for Machine Learning Studio

Cortana Intelligence Gallery offers several [custom modules](#) that expand the capabilities of Azure Machine Learning Studio. You can import the modules to use in your experiments, so you can develop even more advanced predictive analytics solutions.

Currently, the Gallery offers modules on *time series analytics*, *association rules*, *clustering algorithms* (beyond k-means), and *visualizations*, and other workhorse utility modules.

Discover

To browse custom modules [in the Gallery](#), under **More**, select **Custom Modules**.



The [Custom Modules](#) page displays a list of recently added and popular modules. To view all custom modules, select the **See all** button. To search for a specific custom module, select **See all**, and then select filter criteria. You also can enter search terms in the **Search** box at the top of the Gallery page.

Custom Modules

Download these custom Azure Machine Learning modules and deploy them into your experiments. Custom modules expand the capabilities of Azure Machine Learning Studio to allow you to develop even more advanced predictive analytics solutions.



Recently added

CUSTOM MODULE	CUSTOM MODULE	CUSTOM MODULE	CUSTOM MODULE
Train Score Timeseries	Discover Association Rules	Add mean and deviation	Visualize Grouped Data
This module trains a model for forecasting timeseries data using	Discover a set of association rules or frequent itemsets along with	This module adds new feature for moving average and standard	Visualize data grouped by key column

Understand

To understand how a published custom module works, select the custom module to open the module details page. The details page delivers a consistent and informative learning experience. For example, the details page highlights the purpose of the module, and it lists expected inputs, outputs, and parameters. The details page also has a link to the underlying source code, which you can examine and customize.

Comment and share

On a custom module details page, in the **Comments** section, you can comment, provide feedback, or ask questions about the module. You can even share the module with friends or colleagues on Twitter or LinkedIn. You also can email a link to the module details page, to invite other users to view the page.

8 Comments Cortana Intelligence Gallery [Login](#)

Heart Recommend 8 Share Sort by Best

Join the discussion...

Import

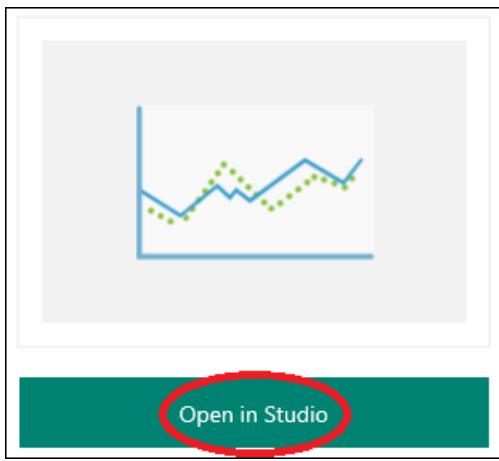
You can import any custom module from the Gallery to your own experiments.

Cortana Intelligence Gallery offers two ways to import a copy of the module:

- **From the Gallery.** When you import a custom module from the Gallery, you also get a sample experiment that gives you an example of how to use the module.
- **From within Machine Learning Studio.** You can import any custom module while you're working in Machine Learning Studio (in this case, you don't get the sample experiment).

From the Gallery

1. In the Gallery, open the module details page.
2. Select **Open in Studio**.

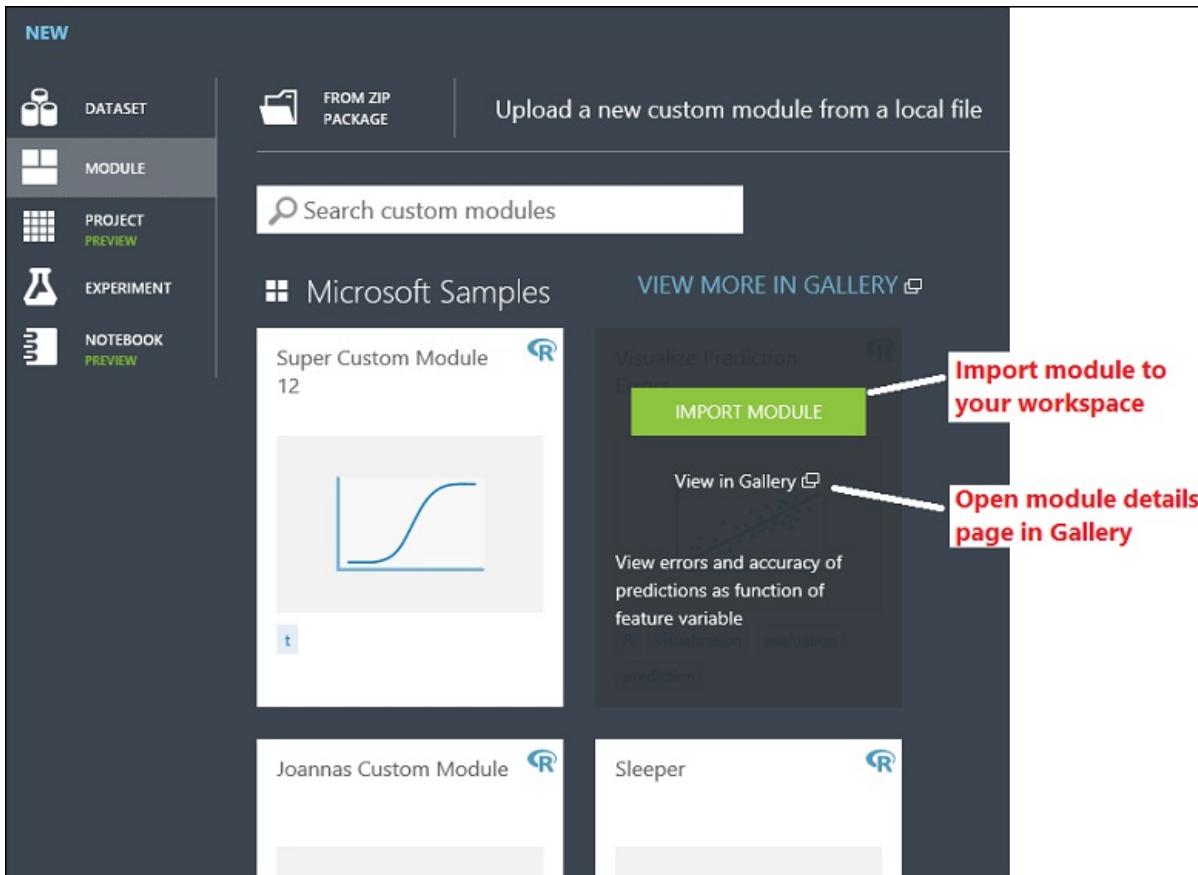


Each custom module includes a sample experiment that demonstrates how to use the module. When you select **Open in Studio**, the sample experiment opens in your Machine Learning Studio workspace. (If you're not already signed in to Studio, you are prompted to first sign in by using your Microsoft account.)

In addition to the sample experiment, the custom module is copied to your workspace. It's also placed in your module palette, with all your built-in or custom Machine Learning Studio modules. You can now use it in your own experiments, like any other module in your workspace.

From within Machine Learning Studio

1. In Machine Learning Studio, select **NEW**.
2. Select **Module**. You can choose from a list of Gallery modules, or find a specific module by using the **Search** box.
3. Point your mouse at a module, and then select **Import Module**. (To get information about the module, select **View in Gallery**. This takes you to the module details page in the Gallery.)



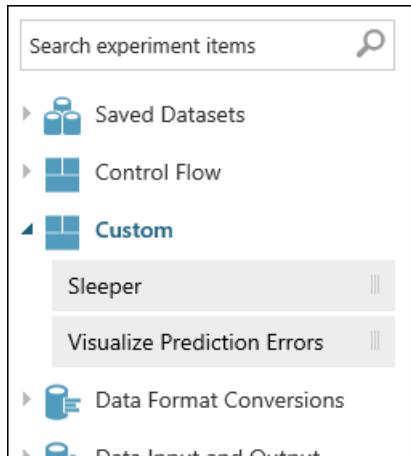
The custom module is copied to your workspace and placed in your module palette, with your built-in or custom Machine Learning Studio modules. You can now use it in your own experiments, like any other module in your workspace.

Use

Regardless of which method you choose to import a custom module, when you import the module, the module is placed in your module palette in Machine Learning Studio. From your module palette, you can use the custom module in any experiment in your workspace, just like any other module.

To use an imported module:

1. Create an experiment, or open an existing experiment.
2. To expand the list of custom modules in your workspace, in the module palette, select **Custom**. The module palette is to the left of the experiment canvas.



3. Select the module that you imported, and drag it to your experiment.

Go to the Gallery

NOTE

Try Azure Machine Learning for free

No credit card or Azure subscription required. [Get started now.](#)

Guide to Net# neural network specification language for Azure Machine Learning

9/25/2017 • 23 min to read • [Edit Online](#)

Overview

Net# is a language developed by Microsoft that is used to define neural network architectures. You can use Net# in neural network modules in Microsoft Azure Machine Learning.

In this article, you will learn basic concepts needed to develop a custom neural network:

- Neural network requirements and how to define the primary components
- The syntax and keywords of the Net# specification language
- Examples of custom neural networks created using Net#

NOTE

[Try Azure Machine Learning for free](#)

No credit card or Azure subscription required. [Get started now](#).

Neural network basics

A neural network structure consists of **nodes** that are organized in **layers**, and weighted **connections** (or **edges**) between the nodes. The connections are directional, and each connection has a **source** node and a **destination** node.

Each **trainable layer** (a hidden or an output layer) has one or more **connection bundles**. A connection bundle consists of a source layer and a specification of the connections from that source layer. All the connections in a given bundle share the same **source layer** and the same **destination layer**. In Net#, a connection bundle is considered as belonging to the bundle's destination layer.

Net# supports various kinds of connection bundles, which lets you customize the way inputs are mapped to hidden layers and mapped to the outputs.

The default or standard bundle is a **full bundle**, in which each node in the source layer is connected to every node in the destination layer.

Additionally, Net# supports the following four kinds of advanced connection bundles:

- **Filtered bundles.** The user can define a predicate by using the locations of the source layer node and the destination layer node. Nodes are connected whenever the predicate is True.
- **Convolutional bundles.** The user can define small neighborhoods of nodes in the source layer. Each node in the destination layer is connected to one neighborhood of nodes in the source layer.
- **Pooling bundles** and **Response normalization bundles**. These are similar to convolutional bundles in that the user defines small neighborhoods of nodes in the source layer. The difference is that the weights of the edges in these bundles are not trainable. Instead, a predefined function is applied to the source node values to determine the destination node value.

Using Net# to define the structure of a neural network makes it possible to define complex structures such as deep neural networks or convolutions of arbitrary dimensions, which are known to improve learning on data such as

image, audio, or video.

Supported customizations

The architecture of neural network models that you create in Azure Machine Learning can be extensively customized by using Net#. You can:

- Create hidden layers and control the number of nodes in each layer.
- Specify how layers are to be connected to each other.
- Define special connectivity structures, such as convolutions and weight sharing bundles.
- Specify different activation functions.

For details of the specification language syntax, see [Structure Specification](#).

For examples of defining neural networks for some common machine learning tasks, from simplex to complex, see [Examples](#).

General requirements

- There must be exactly one output layer, at least one input layer, and zero or more hidden layers.
- Each layer has a fixed number of nodes, conceptually arranged in a rectangular array of arbitrary dimensions.
- Input layers have no associated trained parameters and represent the point where instance data enters the network.
- Trainable layers (the hidden and output layers) have associated trained parameters, known as weights and biases.
- The source and destination nodes must be in separate layers.
- Connections must be acyclic; in other words, there cannot be a chain of connections leading back to the initial source node.
- The output layer cannot be a source layer of a connection bundle.

Structure specifications

A neural network structure specification is composed of three sections: the **constant declaration**, the **layer declaration**, the **connection declaration**. There is also an optional **share declaration** section. The sections can be specified in any order.

Constant declaration

A constant declaration is optional. It provides a means to define values used elsewhere in the neural network definition. The declaration statement consists of an identifier followed by an equal sign and a value expression.

For example, the following statement defines a constant **x**:

```
Const X = 28;
```

To define two or more constants simultaneously, enclose the identifier names and values in braces, and separate them by using semicolons. For example:

```
Const { X = 28; Y = 4; }
```

The right-hand side of each assignment expression can be an integer, a real number, a Boolean value (True or False), or a mathematical expression. For example:

```
Const { X = 17 * 2; Y=true; }
```

Layer declaration

The layer declaration is required. It defines the size and source of the layer, including its connection bundles and attributes. The declaration statement starts with the name of the layer (input, hidden, or output), followed by the dimensions of the layer (a tuple of positive integers). For example:

```
input Data auto;
hidden Hidden[5,20] from Data all;
output Result[2] from Hidden all;
```

- The product of the dimensions is the number of nodes in the layer. In this example, there are two dimensions [5,20], which means there are 100 nodes in the layer.
- The layers can be declared in any order, with one exception: If more than one input layer is defined, the order in which they are declared must match the order of features in the input data.

To specify that the number of nodes in a layer be determined automatically, use the **auto** keyword. The **auto** keyword has different effects, depending on the layer:

- In an input layer declaration, the number of nodes is the number of features in the input data.
- In a hidden layer declaration, the number of nodes is the number that is specified by the parameter value for **Number of hidden nodes**.
- In an output layer declaration, the number of nodes is 2 for two-class classification, 1 for regression, and equal to the number of output nodes for multiclass classification.

For example, the following network definition allows the size of all layers to be automatically determined:

```
input Data auto;
hidden Hidden auto from Data all;
output Result auto from Hidden all;
```

A layer declaration for a trainable layer (the hidden or output layers) can optionally include the output function (also called an activation function), which defaults to **sigmoid** for classification models, and **linear** for regression models. (Even if you use the default, you can explicitly state the activation function, if desired for clarity.)

The following output functions are supported:

- sigmoid
- linear
- softmax
- rlinear
- square
- sqrt
- srlinear
- abs
- tanh
- brlinear

For example, the following declaration uses the **softmax** function:

```
output Result [100] softmax from Hidden all;
```

Connection declaration

Immediately after defining the trainable layer, you must declare connections among the layers you have defined. The connection bundle declaration starts with the keyword **from**, followed by the name of the bundle's source layer and the kind of connection bundle to create.

Currently, five kinds of connection bundles are supported:

- **Full** bundles, indicated by the keyword **all**
- **Filtered** bundles, indicated by the keyword **where**, followed by a predicate expression
- **Convolutional** bundles, indicated by the keyword **convolve**, followed by the convolution attributes
- **Pooling** bundles, indicated by the keywords **max pool** or **mean pool**
- **Response normalization** bundles, indicated by the keyword **response norm**

Full bundles

A full connection bundle includes a connection from each node in the source layer to each node in the destination layer. This is the default network connection type.

Filtered bundles

A filtered connection bundle specification includes a predicate, expressed syntactically, much like a C# lambda expression. The following example defines two filtered bundles:

```
input Pixels [10, 20];
hidden ByRow[10, 12] from Pixels where (s,d) => s[0] == d[0];
hidden ByCol[5, 20] from Pixels where (s,d) => abs(s[1] - d[1]) <= 1;
```

- In the predicate for *ByRow*, **s** is a parameter representing an index into the rectangular array of nodes of the input layer, *Pixels*, and **d** is a parameter representing an index into the array of nodes of the hidden layer, *ByRow*. The type of both **s** and **d** is a tuple of integers of length two. Conceptually, **s** ranges over all pairs of integers with $0 \leq s[0] < 10$ and $0 \leq s[1] < 20$, and **d** ranges over all pairs of integers, with $0 \leq d[0] < 10$ and $0 \leq d[1] < 12$.
- On the right-hand side of the predicate expression, there is a condition. In this example, for every value of **s** and **d** such that the condition is True, there is an edge from the source layer node to the destination layer node. Thus, this filter expression indicates that the bundle includes a connection from the node defined by **s** to the node defined by **d** in all cases where **s[0]** is equal to **d[0]**.

Optionally, you can specify a set of weights for a filtered bundle. The value for the **Weights** attribute must be a tuple of floating point values with a length that matches the number of connections defined by the bundle. By default, weights are randomly generated.

Weight values are grouped by the destination node index. That is, if the first destination node is connected to K source nodes, the first K elements of the **Weights** tuple are the weights for the first destination node, in source index order. The same applies for the remaining destination nodes.

It's possible to specify weights directly as constant values. For example, if you learned the weights previously, you can specify them as constants using this syntax:

```
const Weights_1 = [0.0188045055, 0.130500451, ...]
```

Convolutional bundles

When the training data has a homogeneous structure, convolutional connections are commonly used to learn high-level features of the data. For example, in image, audio, or video data, spatial or temporal dimensionality can be fairly uniform.

Convolutional bundles employ rectangular **kernels** that are slid through the dimensions. Essentially, each kernel defines a set of weights applied in local neighborhoods, referred to as **kernel applications**. Each kernel application corresponds to a node in the source layer, which is referred to as the **central node**. The weights of a kernel are shared among many connections. In a convolutional bundle, each kernel is rectangular and all kernel applications are the same size.

Convolutional bundles support the following attributes:

InputShape defines the dimensionality of the source layer for the purposes of this convolutional bundle. The value must be a tuple of positive integers. The product of the integers must equal the number of nodes in the source layer, but otherwise, it does not need to match the dimensionality declared for the source layer. The length of this tuple becomes the **arity** value for the convolutional bundle. (Typically arity refers to the number of arguments or operands that a function can take.)

To define the shape and locations of the kernels, use the attributes **KernelShape**, **Stride**, **Padding**, **LowerPad**, and **UpperPad**:

- **KernelShape**: (required) Defines the dimensionality of each kernel for the convolutional bundle. The value must be a tuple of positive integers with a length that equals the arity of the bundle. Each component of this tuple must be no greater than the corresponding component of **InputShape**.
- **Stride**: (optional) Defines the sliding step sizes of the convolution (one step size for each dimension), that is the distance between the central nodes. The value must be a tuple of positive integers with a length that is the arity of the bundle. Each component of this tuple must be no greater than the corresponding component of **KernelShape**. The default value is a tuple with all components equal to one.
- **Sharing**: (optional) Defines the weight sharing for each dimension of the convolution. The value can be a single Boolean value or a tuple of Boolean values with a length that is the arity of the bundle. A single Boolean value is extended to be a tuple of the correct length with all components equal to the specified value. The default value is a tuple that consists of all True values.
- **MapCount**: (optional) Defines the number of feature maps for the convolutional bundle. The value can be a single positive integer or a tuple of positive integers with a length that is the arity of the bundle. A single integer value is extended to be a tuple of the correct length with the first components equal to the specified value and all the remaining components equal to one. The default value is one. The total number of feature maps is the product of the components of the tuple. The factoring of this total number across the components determines how the feature map values are grouped in the destination nodes.
- **Weights**: (optional) Defines the initial weights for the bundle. The value must be a tuple of floating point values with a length that is the number of kernels times the number of weights per kernel, as defined later in this article. The default weights are randomly generated.

There are two sets of properties that control padding, the properties being mutually exclusive:

- **Padding**: (optional) Determines whether the input should be padded by using a **default padding scheme**. The value can be a single Boolean value, or it can be a tuple of Boolean values with a length that is the arity of the bundle. A single Boolean value is extended to be a tuple of the correct length with all components equal to the specified value. If the value for a dimension is True, the source is logically padded in that dimension with zero-valued cells to support additional kernel applications, such that the central nodes of the first and last kernels in that dimension are the first and last nodes in that dimension in the source layer. Thus, the number of "dummy" nodes in each dimension is determined automatically, to fit exactly $(InputShape[d] - 1) / Stride[d] + 1$ kernels into the padded source layer. If the value for a dimension is False, the kernels are defined so that the number of nodes on each side that are left out is the same (up to a difference of 1). The default value of this attribute is a

tuple with all components equal to False.

- **UpperPad** and **LowerPad**: (optional) Provide greater control over the amount of padding to use. **Important:**

These attributes can be defined if and only if the **Padding** property above is **not** defined. The values should be integer-valued tuples with lengths that are the arity of the bundle. When these attributes are specified, "dummy" nodes are added to the lower and upper ends of each dimension of the input layer. The number of nodes added to the lower and upper ends in each dimension is determined by **LowerPad[i]** and **UpperPad[i]** respectively. To ensure that kernels correspond only to "real" nodes and not to "dummy" nodes, the following conditions must be met:

- Each component of **LowerPad** must be strictly less than $\text{KernelShape}[d]/2$.
- Each component of **UpperPad** must be no greater than $\text{KernelShape}[d]/2$.
- The default value of these attributes is a tuple with all components equal to 0.

The setting **Padding** = true allows as much padding as is needed to keep the "center" of the kernel inside the "real" input. This changes the math a bit for computing the output size. Generally, the output size D is computed as $D = (I - K) / S + 1$, where I is the input size, K is the kernel size, S is the stride, and $/$ is integer division (round toward zero). If you set **UpperPad** = [1, 1], the input size I is effectively 29, and thus $D = (29 - 5) / 2 + 1 = 13$. However, when **Padding** = true, essentially I gets bumped up by $K - 1$; hence $D = ((28 + 4) - 5) / 2 + 1 = 27 / 2 + 1 = 13 + 1 = 14$. By specifying values for **UpperPad** and **LowerPad** you get much more control over the padding than if you just set **Padding** = true.

For more information about convolutional networks and their applications, see these articles:

- <http://deeplearning.net/tutorial/lenet.html>
- <http://research.microsoft.com/pubs/68920/icdar03.pdf>
- http://people.csail.mit.edu/jvb/papers/cnn_tutorial.pdf

Pooling bundles

A **pooling bundle** applies geometry similar to convolutional connectivity, but it uses predefined functions to source node values to derive the destination node value. Hence, pooling bundles have no trainable state (weights or biases). Pooling bundles support all the convolutional attributes except **Sharing**, **MapCount**, and **Weights**.

Typically, the kernels summarized by adjacent pooling units do not overlap. If $\text{Stride}[d]$ is equal to $\text{KernelShape}[d]$ in each dimension, the layer obtained is the traditional local pooling layer, which is commonly employed in convolutional neural networks. Each destination node computes the maximum or the mean of the activities of its kernel in the source layer.

The following example illustrates a pooling bundle:

```
hidden P1 [5, 12, 12]
from C1 maxpool {
    InputShape = [ 5, 24, 24];
    KernelShape = [ 1, 2, 2];
    Stride = [ 1, 2, 2];
}
```

- The arity of the bundle is 3 (the length of the tuples **InputShape**, **KernelShape**, and **Stride**).
- The number of nodes in the source layer is $5 * 24 * 24 = 2880$.
- This is a traditional local pooling layer because **KernelShape** and **Stride** are equal.
- The number of nodes in the destination layer is $5 * 12 * 12 = 1440$.

For more information about pooling layers, see these articles:

- <http://www.cs.toronto.edu/~hinton/absps/imagenet.pdf> (Section 3.4)
- <http://cs.nyu.edu/~koray/publis/lecun-iscas-10.pdf>

- <http://cs.nyu.edu/~koray/publis/jarrett-iccv-09.pdf>

Response normalization bundles

Response normalization is a local normalization scheme that was first introduced by Geoffrey Hinton, et al, in the paper [ImageNet Classification with Deep Convolutional Neural Networks](#). Response normalization is used to aid generalization in neural nets. When one neuron is firing at a very high activation level, a local response normalization layer suppresses the activation level of the surrounding neurons. This is done by using three parameters (α , β , and k) and a convolutional structure (or neighborhood shape). Every neuron in the destination layer y corresponds to a neuron x in the source layer. The activation level of y is given by the following formula, where f is the activation level of a neuron, and Nx is the kernel (or the set that contains the neurons in the neighborhood of x), as defined by the following convolutional structure:

$$\frac{f(x)}{\left(k + \frac{\alpha}{|N_x|} \sum_{z \in N_x} (f(z))^2 \right)^\beta}$$

Response normalization bundles support all the convolutional attributes except **Sharing**, **MapCount**, and **Weights**.

- If the kernel contains neurons in the same map as x , the normalization scheme is referred to as **same map normalization**. To define same map normalization, the first coordinate in **InputShape** must have the value 1.
- If the kernel contains neurons in the same spatial position as x , but the neurons are in other maps, the normalization scheme is called **across maps normalization**. This type of response normalization implements a form of lateral inhibition inspired by the type found in real neurons, creating competition for big activation levels amongst neuron outputs computed on different maps. To define across maps normalization, the first coordinate must be an integer greater than one and no greater than the number of maps, and the rest of the coordinates must have the value 1.

Because response normalization bundles apply a predefined function to source node values to determine the destination node value, they have no trainable state (weights or biases).

Alert: The nodes in the destination layer correspond to neurons that are the central nodes of the kernels. For example, if *KernelShape[d]* is odd, then *KernelShape[d]/2* corresponds to the central kernel node. If *KernelShape[d]* is even, the central node is at *KernelShape[d]/2 - 1*. Therefore, if **Padding[d]** is False, the first and the last *KernelShape[d]/2* nodes do not have corresponding nodes in the destination layer. To avoid this situation, define **Padding** as [true, true, ..., true].

In addition to the four attributes described earlier, response normalization bundles also support the following attributes:

- **Alpha**: (required) Specifies a floating-point value that corresponds to α in the previous formula.
- **Beta**: (required) Specifies a floating-point value that corresponds to β in the previous formula.
- **Offset**: (optional) Specifies a floating-point value that corresponds to k in the previous formula. It defaults to 1.

The following example defines a response normalization bundle using these attributes:

```
hidden RN1 [5, 10, 10]
from P1 response norm {
  InputShape = [ 5, 12, 12];
  KernelShape = [ 1, 3, 3];
  Alpha = 0.001;
  Beta = 0.75;
}
```

- The source layer includes five maps, each with aof dimension of 12x12, totaling in 1440 nodes.
- The value of **KernelShape** indicates that this is a same map normalization layer, where the neighborhood is a 3x3 rectangle.
- The default value of **Padding** is False, thus the destination layer has only 10 nodes in each dimension. To include one node in the destination layer that corresponds to every node in the source layer, add **Padding = [true, true, true]**; and change the size of RN1 to [5, 12, 12].

Share declaration

Net# optionally supports defining multiple bundles with shared weights. The weights of any two bundles can be shared if their structures are the same. The following syntax defines bundles with shared weights:

```

share-declaration:
  share { layer-list }
  share { bundle-list }
  share { bias-list }

layer-list:
  layer-name , layer-name
  layer-list , layer-name

bundle-list:
  bundle-spec , bundle-spec
  bundle-list , bundle-spec

bundle-spec:
  layer-name => layer-name

bias-list:
  bias-spec , bias-spec
  bias-list , bias-spec

bias-spec:
  1 => layer-name

layer-name:
  identifier

```

For example, the following share-declaration specifies the layer names, indicating that both weights and biases should be shared:

```

Const {
  InputSize = 37;
  HiddenSize = 50;
}
input {
  Data1 [InputSize];
  Data2 [InputSize];
}
hidden {
  H1 [HiddenSize] from Data1 all;
  H2 [HiddenSize] from Data2 all;
}
output Result [2] {
  fromH1 all;
  fromH2 all;
}
share { H1, H2 } // share both weights and biases

```

- The input features are partitioned into two equal sized input layers.
- The hidden layers then compute higher level features on the two input layers.

- The share-declaration specifies that $H1$ and $H2$ must be computed in the same way from their respective inputs.

Alternatively, this could be specified with two separate share-declarations as follows:

```
share { Data1 => H1, Data2 => H2 } // share weights
```

```
share { 1 => H1, 1 => H2 } // share biases
```

You can use the short form only when the layers contain a single bundle. In general, sharing is possible only when the relevant structure is identical, meaning that they have the same size, same convolutional geometry, and so forth.

Examples of Net# usage

This section provides some examples of how you can use Net# to add hidden layers, define the way that hidden layers interact with other layers, and build convolutional networks.

Define a simple custom neural network: "Hello World" example

This simple example demonstrates how to create a neural network model that has a single hidden layer.

```
input Data auto;
hidden H [200] from Data all;
output Out [10] sigmoid from H all;
```

The example illustrates some basic commands as follows:

- The first line defines the input layer (named *Data*). When you use the **auto** keyword, the neural network automatically includes all feature columns in the input examples.
- The second line creates the hidden layer. The name *H* is assigned to the hidden layer, which has 200 nodes. This layer is fully connected to the input layer.
- The third line defines the output layer (named *O*), which contains 10 output nodes. If the neural network is used for classification, there is one output node per class. The keyword **sigmoid** indicates that the output function is applied to the output layer.

Define multiple hidden layers: computer vision example

The following example demonstrates how to define a slightly more complex neural network, with multiple custom hidden layers.

```

// Define the input layers
input Pixels [10, 20];
input MetaData [7];

// Define the first two hidden layers, using data only from the Pixels input
hidden ByRow [10, 12] from Pixels where (s,d) => s[0] == d[0];
hidden ByCol [5, 20] from Pixels where (s,d) => abs(s[1] - d[1]) <= 1;

// Define the third hidden layer, which uses as source the hidden layers ByRow and ByCol
hidden Gather [100]
{
    fromByRow all;
    fromByCol all;
}

// Define the output layer and its sources
output Result [10]
{
    from Gather all;
    from MetaData all;
}

```

This example illustrates several features of the neural networks specification language:

- The structure has two input layers, *Pixels* and *MetaData*.
- The *Pixels* layer is a source layer for two connection bundles, with destination layers, *ByRow* and *ByCol*.
- The layers *Gather* and *Result* are destination layers in multiple connection bundles.
- The output layer, *Result*, is a destination layer in two connection bundles; one with the second level hidden (*Gather*) as a destination layer, and the other with the input layer (*MetaData*) as a destination layer.
- The hidden layers, *ByRow* and *ByCol*, specify filtered connectivity by using predicate expressions. More precisely, the node in *ByRow* at [x, y] is connected to the nodes in *Pixels* that have the first index coordinate equal to the node's first coordinate, x. Similarly, the node in *ByCol* at [x, y] is connected to the nodes in *_Pixels* that have the second index coordinate within one of the node's second coordinate, y.

Define a convolutional network for multiclass classification: digit recognition example

The definition of the following network is designed to recognize numbers, and it illustrates some advanced techniques for customizing a neural network.

```

input Image [29, 29];
hidden Conv1 [5, 13, 13] from Image convolve
{
    InputShape = [29, 29];
    KernelShape = [5, 5];
    Stride = [2, 2];
    MapCount = 5;
}
hidden Conv2 [50, 5, 5]
from Conv1 convolve
{
    InputShape = [5, 13, 13];
    KernelShape = [1, 5, 5];
    Stride = [1, 2, 2];
    Sharing = [false, true, true];
    MapCount = 10;
}
hidden Hid3 [100] from Conv2 all;
output Digit [10] from Hid3 all;

```

- The structure has a single input layer, *Image*.
- The keyword **convolve** indicates that the layers named *Conv1* and *Conv2* are convolutional layers. Each of

these layer declarations is followed by a list of the convolution attributes.

- The net has a third hidden layer, *Hid3*, which is fully connected to the second hidden layer, *Conv2*.
- The output layer, *Digit*, is connected only to the third hidden layer, *Hid3*. The keyword **all** indicates that the output layer is fully connected to *Hid3*.
- The arity of the convolution is three (the length of the tuples **InputShape**, **KernelShape**, **Stride**, and **Sharing**).
- The number of weights per kernel is $1 + \text{KernelShape}[0] * \text{KernelShape}[1] * \text{KernelShape}[2] = 1 + 1 * 5 * 5 = 26$. Or $26 * 50 = 1300$.
- You can calculate the nodes in each hidden layer as follows:
 - **NodeCount**[0] = $(5 - 1) / 1 + 1 = 5$.
 - **NodeCount**[1] = $(13 - 5) / 2 + 1 = 5$.
 - **NodeCount**[2] = $(13 - 5) / 2 + 1 = 5$.
- The total number of nodes can be calculated by using the declared dimensionality of the layer, [50, 5, 5], as follows: **MapCount*** *****NodeCount**[0] * **NodeCount**[1] * **NodeCount**[2] = $10 * 5 * 5 * 5 * 5 = 50^*$.
- Because **Sharing**[d] is False only for $d == 0$, the number of kernels is **MapCount*** *****NodeCount**[0] = $10 * 5 = 50^*$.

Acknowledgements

The Net# language for customizing the architecture of neural networks was developed at Microsoft by Shon Katzenberger (Architect, Machine Learning) and Alexey Kamenev (Software Engineer, Microsoft Research). It is used internally for machine learning projects and applications ranging from image detection to text analytics. For more information, see [Neural Nets in Azure ML - Introduction to Net#](#)

Get help from Machine Learning Live Chat Support

9/25/2017 • 1 min to read • [Edit Online](#)

Azure Machine Learning Studio provides an intuitive interface for building machine learning models. There is a [Gallery](#) of examples contributed by the community to help you get started. And there is a [Forum](#) where you can ask questions to help keep you going.

But sometimes you just need to ask a quick question to unblock you. In [Machine Learning Studio](#) look for the chat icon in the top navigation. If you see this icon, that means a member of the product team is online to help you in real time.



Type in your question and get your answers!

A screenshot of a 'Chat with us' interface. At the top, it says 'Chat with us'. Below that, it states 'Zendesk privacy terms apply. See: http://www.zendesk.com/company/privacy'. A message from 'Olga' says 'Hello there!'. Below that, a message input field contains 'Hello! I have questions related to Multiclass Decision Forest algorithm....'. At the bottom, a blue button says 'Start Chatting' and the word 'zopim' is at the bottom right.

Don't see the Live Chat Icon?

The Live Chat is staffed by members of the Machine Learning team. If you don't see the live chat icon it's because the team is not currently available. For example, it may be outside of normal working hours.

You can always post your questions on the [Machine Learning Forum](#), where they will be seen by members of the Machine Learning team and other knowledgeable data scientists and developers.

NOTE

Try Azure Machine Learning for free

No credit card or Azure subscription required. [Get started now.](#)